Microsoft Fabric governance and administration documentation

Manage, protect, and unlock trusted reuse across your organization's Fabric data estate.

Fabric in your organization

OVERVIEW

Govern and manage with OneLake Catalog

What is administration in Fabric?

GET STARTED

Enable Fabric for your organization

Region availability

Find your Fabric home region

HOW-TO GUIDE

Understand Fabric admin roles

REFERENCE

Governance documentation

Security documentation

Tools and settings

OVERVIEW

About tenant settings

HOW-TO GUIDE

Set up git integration

Set up item certification

Configure notifications

Set up metadata scanning

Enable content certification

Enable service principal authentication

Configure Multi-Geo support

Domain administration

CONCEPT

Domains overview

HOW-TO GUIDE

Understand domain roles

Create and manage domains

Capacity administration

CONCEPT

Manage your Fabric capacity

HOW-TO GUIDE

Understand capacity administrator roles

Configure capacity contributor permissions

Monitor capacity with the Capacity Metrics app

Workspace administration

CONCEPT

Manage workspaces

HOW-TO GUIDE

Workspace roles

Delegate settings to workspaces

What is Microsoft Fabric administration?

Fabric administration is the set of responsibilities, roles, and tools used to govern, secure, operate, and optimize the Microsoft Fabric environment across an organization. Administration isn't a single settings screen—it's how you manage your organization's Fabric data estate so teams can build, share, and consume data with confidence.

Administrators focus on three connected responsibilities that keep the platform configured correctly and compliant with organizational policies:

- **Administration**—This *Fabric governance and administration documentation* covers how to manage the Fabric platform, configure tenant and workspace features, and monitor usage and activity. It also covers how to define and enforce policies for data access, sharing, classification, and auditing.
- **Security**—See the *Security documentation* to learn how to help safeguard data with identity, access, encryption, and network protection settings.

This article introduces Fabric administration as an operating model and shows how administration fits into managing the overall Fabric data estate.

Administer the Fabric data estate

Fabric administrators are responsible for a set of outcomes across the data estate. Use these outcomes—not individual settings screens—as the mental model for administration. Each outcome group links to the publicly documented capabilities you use to achieve it.

Expand table

Outcome	Purpose	Capabilities
Govern	Establish structure and policy so data is organized, discoverable, and compliant.	Tenant settings, Domains, Workspace controls
Secure	Protect data and control who can access it.	Security controls, Access controls, Compliance capabilities
Operate	Keep the platform running reliably at tenant scale.	Capacities, Tenant-wide configuration, Monitoring and health
Optimize	Get more value from your Fabric investment.	Govern report, Capacity planning

Administrative roles

Fabric provides multiple administrative roles so organizations can distribute responsibilities and avoid central bottlenecks. Administration is a delegated operating model: you assign scoped roles so the people closest to a capacity, workspace, or domain can manage it without needing full tenant access.

Depending on organizational structure, administrators may hold responsibilities across multiple scopes such as tenant administration, capacity administration, workspace administration, and governance ownership.

For details about which types of admins can perform specific tasks, see Understand Fabric admin roles, Microsoft Entra built-in roles, and Microsoft 365 admin roles.

Administrative tools: Governance and administration together

Governance defines the policies, structure, and insights that keep data trustworthy, while administration keeps the platform configured, secure, and operational. The **Govern** section in the OneLake catalog brings these capabilities together. Use it to review governance insights and access tenant settings, workspace and capacity management, policies, and other administrative controls.

The following table summarizes the tools used to manage the Fabric data estate. The entries you see in Govern depend on your permissions:

- Users without admin permissions see **Configurations**, with Help and support entries only, when this experience is enabled for their tenant.
- Domain admins see **Configurations**, with Help and support entries, and **Domains**.
- Capacity admins see **Configurations**, with Help and support entries, and **Capacities**.

Expand table

Task	Tool	Purpose
Manage Fabric	OneLake catalog Govern section	Govern and administer your organization's Fabric data estate from a single location.
Manage identity and access	Microsoft Entra ID	Configure conditional access to Fabric resources.

Task	Tool	Purpose
Manage identity and access	Microsoft 365 admin center	Manage users and groups, purchase and assign licenses, and control access to Fabric.
Manage security and compliance	Microsoft 365 Security & Microsoft Purview compliance portal	Review and manage auditing, data classification, data loss prevention policies, and data lifecycle management.
Automate administration	PowerShell cmdlets	Manage workspaces and other aspects of Fabric using scripts.
Automate administration	Administrative APIs and SDK	Build custom admin tools.

The OneLake catalog Govern section

Use the OneLake catalog **Govern** section as the primary destination for governing and administering your organization's Fabric data estate.

To access the OneLake catalog, sign in to Fabric. Select **OneLake catalog**, and then select the **Govern** section from the top menu.

You can also access Govern from the **Settings** (gear) menu by selecting **OneLake Catalog | Govern**.

> **Note**
>
> OneLake catalog and Govern are rolling out by region. If they're not available in your region, use the **Admin portal** during the regional rollout. In Fabric, select the **Settings** (gear) icon, and then select **Admin portal**.

Govern also provides **Quick access** cards for common administration tasks:

- Admin monitoring workspace
- Manage users
- Audit logs

Related content

- Understand Fabric admin roles
- Fabric tenant settings
- OneLake catalog overview

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Governance overview and guidance

Microsoft Fabric gives you the capabilities to govern, protect, discover, and operate your organization's data estate. Use these capabilities to organize ownership and administration, secure sensitive data and meet compliance requirements, help people find and trust the right data, understand the governance health of your estate, and monitor operations across Fabric. Many of these capabilities are built in and included with your Microsoft Fabric license, while some others require additional licensing from Microsoft Purview.

This article introduces the main outcomes you can achieve when you govern and manage your Fabric data estate, and it links to more detailed information for each one. Many of these outcomes come together in **the Govern section of the OneLake catalog**, alongside the Explore and Secure sections.

Manage and govern your Fabric estate in the OneLake catalog

The Govern section of the OneLake catalog is the centralized experience for managing and governing organizational content in Microsoft Fabric.

In addition to governance capabilities such as domains, endorsements, metadata scanning, policies, and governance insights, Govern also provides access to administrative experiences including workspace management, capacity management, tenant settings, and other organization-wide controls.

Manage the data estate

Organize ownership, domains, and administration so that teams can manage their data according to their specific needs. A single person often wears more than one hat. For example, a Capacity Admin can also be a Fabric Admin, Domain Admin, Workspace Admin, or Data Steward. The following capabilities are described by what you're trying to accomplish rather than by role.

Administrative experiences such as tenant configuration (OneLake catalog → Govern → Configurations), workspace management (OneLake catalog → Govern → Workspaces), and capacity management (OneLake catalog → Govern → Capacities) are available in the Govern section.

You can also manage policies for your Fabric estate from the Govern section of the OneLake catalog.

> **Note**
>
> Start in Govern for governance and administration. If the OneLake catalog and Govern aren't yet available in your region, use the **Admin portal** to access the existing administration capabilities during the regional rollout.

Tenant, domain, and workspace settings

Tenant, domain, and workspace admins each have settings within their scope that they can configure to control who has access to certain functionalities at different levels. You can delegate some tenant-level settings to domain and capacity admins.

For more information, see About tenant settings, Configure domain settings, and Workspace settings.

Guidance: Fabric admins should define tenant-wide settings, and domain admins should override delegated settings as needed. Individual teams (workspace owners) define their own more granular workspace-level controls and settings.

Domains

Use domains to logically group all the data in an organization that's relevant to particular areas or fields, such as by business unit. One of the most common uses for domains is to group data by business department, so each department can manage its data according to its specific regulations, restrictions, and needs.

Grouping data into domains and subdomains enables better discoverability and governance. For instance, in the OneLake catalog, users can filter content by domain to find content that is relevant to them. With respect to governance, you can delegate some tenant-level settings for managing and governing data to the domain level, which allows domain-specific configuration of those settings.

For more information, see Domains.

Guidance: Business and enterprise architects should design the organization's domain setup, while Fabric admins should implement this design by creating domains and subdomains and assigning domain owners. Preferably, center of excellence (COE) teams should be part of this discussion to align the domains with the overall strategy of the organization.

Workspaces

Teams in organizations use workspaces to create Fabric items and collaborate with each other. Assign these workspaces to teams or departments based on governance requirements and data boundaries. How you assign workspaces depends on internal team structure and how the teams want to handle their Fabric items (for example, do they need one or many workspaces).

Guidance: For development purposes, a best practice is to have isolated workspaces per developer, so each developer can work on their own without interfering with the shared workspace. Fabric admins define who has permission to create workspaces. Workspace admins define Spark environments that users can reuse. For more information about best practices, see Best practices for lifecycle management in Fabric.

Capacities

Capacities are the compute resources used by all Fabric workloads. Based on organizational requirements, use capacities as isolation boundaries for compute, chargebacks, and more.

Guidance: Split up capacities based on the requirements of the environment, such as development, test, acceptance, and production (DTAP). This approach provides better workload isolation and chargeback.

Delegate administration

You can delegate some aspects of administration and governance to capacities, domains, and workspaces so the respective admins can manage them in their scope.

For an overview of administration experiences and tools, see What is Microsoft Fabric administration?. For Govern tasks and navigation, see Govern and manage your Fabric data estate with the OneLake catalog.

Guidance: Platform/IT owners should have access to Fabric administration experiences. They can define domains, and delegate domain and capacity management to domain and capacity owners as best suits your organizational needs.

Metadata scanning

Metadata scanning helps your organization govern Microsoft Fabric data by enabling cataloging tools to catalog and report on the metadata of all your organization's Fabric items. It uses a set of

admin REST APIs, known as the *scanner APIs*. The scanner APIs extract metadata such as item name, ID, sensitivity, endorsement status, and more.

For more information, see Metadata scanning.

Encourage trusted discovery

Help people find, evaluate, and trust the right data across your estate.

OneLake catalog

The OneLake catalog makes it easy to find, explore, and use the Fabric data items in your organization that you have access to. It provides information about the items and entry points for working with them. Filtering and search options make it easier to get to relevant data. The catalog is also where data owners govern the data they own.

For more information, see the OneLake catalog overview.

Guidance: Carefully defining and setting up domains is essential for creating an efficient experience in the catalog. Carefully defined domains help set the context for teams and make for better definition of boundaries and ownership. Mapping workspaces to domains is key to helping implement this in Fabric.

Endorsement

Endorsement makes trustworthy, quality data easier to discover. Organizations often have large numbers of Microsoft Fabric items—data, processes, and content—available for sharing and reuse by their Fabric users. Endorsement helps users identify and find the trustworthy, high-quality items they need. With endorsement, item owners can promote their quality items, and organizations can certify items that meet their quality standards. Endorsed items are clearly labeled, both in Fabric and in other places where users look for Fabric items. Endorsed items are given priority in some searches, and you can sort for endorsed items in some lists.

For more information, see Endorsement.

Guidance: Delegate certification enablement to domain admins, and have the domain admins authorize data owners and producers to certify the items they create. Data owners and producers should then always certify their items that they test and are ready for use by other teams. This practice helps separate low-quality, untrusted items from trusted, ready-to-use items. It also makes these trusted items easier to find. In addition, educate data consumers about how to find

trusted items, and encourage them to use only certified items in their reports and other downstream processing.

Tags

Tags are configurable text labels that you can apply to Fabric items to enhance item discoverability and use. Fabric administrators can define a set of tags that data owners can use to categorize their items. Once you apply tags to items, data consumers can view, search, and filter by the applied tags across the various Fabric experiences.

For more information, see Tags in Microsoft Fabric.

Data lineage and impact analysis

In modern business intelligence projects, understanding the flow of data from a data source to its destination is a complex task. Questions like "What happens if I change this data?" or "Why isn't this report up to date?" can be hard to answer. They might require a team of experts or deep investigation to understand. Lineage helps users understand the flow of data by providing a visualization that shows the relations between all the items in a workspace. For each item in the lineage view, you can display an impact analysis that shows what downstream items are affected if you make changes to the item.

For more information, see Lineage and Impact analysis.

Guidance: Use proper and consistent naming conventions for items. This practice helps when looking at lineage information.

Protect and comply

Secure your data and meet compliance requirements. Some of these capabilities are included with Microsoft Fabric, while others require additional licensing from Microsoft Purview. Capabilities that require additional Purview licensing are noted in this section. For details about network security, access control, and encryption, see the Security overview.

Privacy

The first phase of any data protection strategy is to identify where your private data sits. This step is one of the most challenging but important steps toward making sure you can protect your data at the source. The following sections describe capabilities Fabric provides to help your organization meet this challenge.

Data security

To ensure data in Fabric is secure from unauthorized access and stays compliant with data privacy requirements, use sensitivity labels from Microsoft Purview Information Protection in combination with built-in Fabric capabilities to manually or automatically tag your organization's data. Purview Audit then captures audit trails on activities performed in Fabric. This process includes capturing user activities in the Fabric tenant, such as lakehouse access, Power BI access, Spark activities, data factory activities, sign-ins, and more.

Securing items in a workspace

Organizational teams can have individual workspaces where different personas collaborate and work on generating content. Workspace admins regulate access to the items in the workspace by assigning workspace roles to users.

Guidance: Data owners should recommend users who could be workspace administrators. These users could be team leads in your organization, for example. These workspace administrators should then govern access to the items in their workspace by assigning appropriate workspace roles to users and consumers of the items.

Securing data in Fabric items

Along with the broad security that gets applied at the tenant or workspace level, individual teams can deploy other data-level controls to manage access to individual tables, rows, and columns. Fabric currently provides such data-level control for SQL analytics endpoints, warehouses, Direct Lake, and KQL Database.

Guidance: Individual teams are expected to apply these additional controls at the item and data level.

Purview Information Protection

Information protection in Fabric requires additional licensing from Microsoft Purview. It enables you to discover, classify, and protect Fabric data by using sensitivity labels from Microsoft Purview Information Protection. Fabric provides multiple capabilities, such as default labeling, label inheritance, and programmatic labeling, to help achieve maximal sensitivity label coverage across your entire Fabric data estate. Once labeled, data remains protected even when it's exported out of Fabric via supported export paths. Compliance admins can monitor activities on sensitivity labels in Microsoft Purview Audit.

For more information, see Information Protection in Microsoft Fabric.

Guidance: Specify sensitivity labels from Microsoft Purview Information Protection and their associated label policies at an organizational level. They should be valid for the whole organization.

Purview Data Loss Prevention

Purview Data Loss Prevention (DLP) requires additional licensing from Microsoft Purview. DLP policies for Fabric and Power BI automatically detect sensitive information as you upload it into DLP-supported item types in your Fabric tenant. They help you take risk remediation actions so that your organization stays compliant with governmental and industry regulations.

Compliance and security administrators receive audit logs for every DLP detection. The audit logs give them further visibility into business-critical data and its location within the tenant. They can set up alerts that are automatically generated whenever sensitive information is detected in a DLP-supported item. They can also create customized messages to users to help guide them about how to deal with sensitive data. For example, admins could configure a message that is sent to the Fabric data owner whenever proprietary information is detected in their data, explaining that this information is internal and shouldn't be shared externally.

For more information, see Get started with Data loss prevention policies for Fabric and Power BI.

Auditing

To mitigate the risks of unauthorized access and use of your Fabric data, Fabric administrators and compliance teams in your organizations can track and investigate user activity on Fabric items by using Purview Audit, which is available in the Purview compliance portal. Many companies also need these audit logs for regulatory requirements, which often mandate storing audit logs for forensic investigation and potential data regulation violations.

Guidance: Fabric administrators and compliance teams should be aware that Fabric item-level audits are logged in Purview Audit and can be used for analysis.

Purview governance across your organization

Microsoft Purview offers solutions for protecting and governing data across an organization's entire data estate, and it requires Microsoft Purview licensing. The integration between Purview and Fabric makes it possible to use some of Purview's capabilities to govern your Fabric data in the context of your organization's entire data estate. The data governance capabilities offered on

Fabric via Purview's live view (preview) let data consumers view Fabric workspaces they have access to, and let you run manual scans that make item-level metadata available in Purview.

For more information, see Use Microsoft Purview to govern Microsoft Fabric.

Certifications

Microsoft Fabric has HIPAA BAA, ISO/IEC 27017, ISO/IEC 27018, ISO/IEC 27001, and ISO/IEC 27701 compliance certifications. To learn more, see Fabric compliance offerings.

Understand and improve governance health

Understand the governance state of your estate and take action to improve it. In the **Govern** section, the **All data estate** view gives governance administrators an organization-wide view, and the **My data** view gives data owners a view of the data they own.

- **Governance Health**: the overall governance state of your estate, so you can see where you stand at a glance.
- **Governance Insights**: explanations of your governance status that help you understand what's working and what needs attention.
- **Recommended Actions**: concrete next steps that guide you to improve governance, with links to supporting resources.
- **Governance Coverage**: how broadly governance capabilities are applied across your items and domains.
- **Governance Visibility**: how well you can see the governance state of the data across your estate.

For more information, see Govern with the OneLake catalog.

Monitor Fabric operations

Monitor capacity, activity, and system operations across Fabric. This section covers platform operations, which are distinct from the governance state and actions described in Understand and improve governance health.

Monitoring hub

The Microsoft Fabric monitoring hub enables users to monitor Fabric activities from a central location. Any Fabric user can use the monitoring hub; however, the monitoring hub displays activities only for Fabric items the user has permission to view.

For more information, see Use the Monitoring hub.

Guidance: Expose this capability to developers and team members for monitoring scheduled workloads, such as a dataflow or pipeline refresh, a Spark run, or a data warehouse query.

Capacity metrics

The Microsoft Fabric Capacity Metrics app helps you monitor capacity usage and consumption across your organization.

Guidance: Platform owners and users with platform administrator roles should use this feature to monitor usage and consumption. For more information, see What is the Microsoft Fabric Capacity Metrics app?.

Admin monitoring

The Govern report in the OneLake catalog gives Fabric administrators tenant-wide inventory, usage, sharing, protection, and curation insights. It consolidates and expands on information previously divided among separate administration and security reports. The administrator report and semantic model are stored in the Admin monitoring workspace. For more information, see Explore the Govern report and What is the admin monitoring workspace?

Guidance: Use this feature to gain an overall view of the Fabric platform.

Related content

- OneLake catalog overview
- Domains
- Fabric administration overview
- Fabric security overview
- Microsoft Purview permissions
- Fabric governance documentation

ⓘ**Note:** The author created this article with assistance from AI. Learn more

Last updated on 09/29/2026

Standards compliance in Microsoft Fabric

This article describes Microsoft Fabric's adherence to compliance standards and provides links to relevant resources. The target audience is Fabric administrators and others who need to know how and whether Fabric meets their organization's security, privacy, and reliability standards.

Introduction

Adherence to compliance standards transcends mere rule-following. It represents a commitment to building software with a foundation of security, privacy, and quality. These frameworks, encompassing legal requirements, industry standards, and internal best practices, serve as guiding principles throughout the development process.

Microsoft Fabric is governed by the Microsoft Online Services Terms and the Microsoft Enterprise Privacy Statement.

For the location of data processing, refer to the *Location of Data Processing* terms in the Microsoft Online Services Terms and to the Data Protection Addendum.

Business continuity and disaster recovery

When it comes to business continuity and disaster recovery (BCDR), Microsoft uses the shared responsibility model for disaster recovery (DR). In a shared responsibility model, Microsoft ensures that the baseline infrastructure and platform services are available. At the same time, many Azure services don't automatically replicate data or fall back from a failed region to cross-replicate to another enabled region. For those services, the customer is responsible for setting up a disaster recovery plan that works for their workload. Most services that run on Azure platform as a service (PaaS) offerings provide features and guidance to support DR, and you can use service-specific features to support fast recovery to help develop your DR plan.

DR is about recovering from high-impact events, such as natural disasters or failed deployments that result in downtime and data loss. Regardless of the cause, the best remedy for a disaster is a well-defined and tested DR plan and an application design that actively supports DR. Before you begin to think about creating your disaster recovery plan, see Recommendations for designing a disaster recovery strategy.

Security development lifecycle

Fabric follows the Security Development Lifecycle (SDL), which consists of a set of strict security practices that support security assurance and compliance requirements. The SDL helps developers build more secure software by reducing the number and severity of vulnerabilities in software, while reducing development cost. For more information, see Microsoft Security Development Lifecycle Practices.

Compliance offerings

Compliance offerings are grouped into four segments: **globally applicable**, **US government**, **industry specific**, and **region/country specific**. Compliance offerings are based on various types of assurances, including formal certifications, attestations, validations, authorizations, and assessments produced by independent third-party auditing firms, as well as contractual amendments, self-assessments, and customer guidance documents produced by Microsoft. Each offering description provides links to downloadable resources to assist you with your own compliance obligations. You can access Azure and other Microsoft cloud services audit documentation via the Service Trust Portal (STP).

While certifications typically occur after a product launch (Generally Available or GA), Microsoft integrates compliance best practices throughout the development lifecycle. This proactive approach ensures a strong foundation for future certifications, even though they follow established audit cycles. In simpler terms, Microsoft prioritizes building compliance in from the start, even if formal certification comes later.

Compliance is a shared responsibility. To comply with laws and regulations, cloud service providers and their customers enter a shared responsibility to ensure that each does their part.

- Compliance Offerings - Compliance offerings for Microsoft 365, Azure, and other Microsoft services. | Microsoft Learn
- Audit Reports - Service Trust Portal Home Page (microsoft.com)
- Compliance is a Shared Responsibility - Shared responsibility in the cloud - Microsoft Azure | Microsoft Learn
- Azure compliance - Azure - Compliance Offerings

For more compliance information, the Microsoft Trust Center is the primary resource for Fabric. For more information about compliance, see Microsoft compliance offerings.

Related content

- Fabric governance documentation
- Fabric security documentation
- Fabric admin documentation

Last updated on 05/23/2024

Fabric region availability

This article lists the region availability of the Fabric F SKUs, which are available in the Azure public cloud regions. Some of the Fabric workloads might not be immediately available in new regions, or regions where data centers become constrained.

For details about buying Fabric capacity, see Buy Fabric capacity in Azure.

Fabric workload and feature availability

Fabric workloads and features are available in most Azure public cloud regions, but some workloads and features might have limited availability in specific regions.

Home region and Fabric availability

Your home region is associated with your tenant. Fabric workload availability varies by tenant region, as shown in the next section. If you want to access all Fabric features in regions where they're not yet available, you can create a capacity in a region where Fabric is available. For more information, see Multi-Geo support for Fabric.

Regional Fabric and Power BI workload availability

The following table lists all Azure regions where Power BI or Fabric is available. Geography names follow the Azure regions list. Some regions support only Power BI, while others support all Fabric workloads. Unavailable features are noted where applicable.

Expand table

Geography	Region	Power BI	All Fabric workloads	Unavailable Fabric features
Americas	Brazil South	✅	✅	Not available: Fabric App (preview)
Americas	Canada Central	✅	✅	Not available: Fabric App (preview)
Americas	Canada East	✅	✅	Not available: Fabric App (preview)
Americas	Chile Central	✅	❌	Power BI only region

Geography	Region	Power BI	All Fabric workloads	Unavailable Fabric features
Americas	Mexico Central	✅	✅	Not available: Fabric App (preview)
Americas	US - Central US	✅	✅	
Americas	US - East US	✅	✅	Not available: Operations agent (preview)
Americas	US - East US 2	✅	✅	
Americas	US - North Central US	✅	✅	
Americas	US - South Central US	✅	✅	Not available: Digital twin builder (preview) Ontology (preview) Operations agent Schema Registry (preview)
Americas	US - West US	✅	✅	
Americas	US - West US 2	✅	✅	
Americas	US - West US 3	✅	✅	
Geography	**Region**	**Power BI**	**All Fabric workloads**	**Unavailable Fabric features**
Austria	Austria East	✅	❌	Power BI only region
Belgium	Belgium Central	✅	❌	Power BI only region
Denmark	Denmark East	✅	❌	Power BI only region
Europe	Europe - North Europe	✅	✅	Not available: Digital twin builder (preview) Fabric Database Hub
Europe	Europe - West Europe	✅	✅	Fabric Database Hub
Europe	France Central	✅	✅	
Europe	France South	✅	❌	Power BI only region
Europe	Germany North	✅	❌	Power BI only region
Europe	Germany West Central	✅	✅	

Geography	Region	Power BI	All Fabric workloads	Unavailable Fabric features
Europe	Italy North	✅	✅	
Europe	Norway East	✅	✅	
Europe	Norway West	✅	❌	Power BI only region
Europe	Poland Central	✅	✅	Not available: Fabric App (preview)
Europe	Spain Central	✅	✅	Not available: Fabric App (preview)
Europe	Sweden Central	✅	✅	Not available by default for Power BI: Business Continuity Disaster Recovery (BCDR). Power BI is supported in the region but not in its paired region
Europe	Switzerland North	✅	✅	
Europe	Switzerland West	✅	✅	Not available: Fabric App (preview)
Europe	UK South	✅	✅	Not available: Fabric App (preview)
Europe	UK West	✅	✅	Not available: Fabric App (preview)
Geography	**Region**	**Power BI**	**All Fabric workloads**	**Unavailable Fabric features**
Qatar	Qatar Central	✅	❌	Power BI only region
UAE	UAE Central	✅	❌	Power BI only region
UAE	UAE North	✅	✅	
South Africa	South Africa North	✅	✅	
South Africa	South Africa West	✅	❌	Power BI only region
Asia Pacific	Asia - East Asia	✅	✅	
Asia Pacific	Asia - Southeast Asia	✅	✅	

Geography	Region	Power BI	All Fabric workloads	Unavailable Fabric features
Asia Pacific	Australia East	✅	✅	
Asia Pacific	Australia Southeast	✅	✅	Not available: Fabric App (preview)
Asia Pacific	India - Central India	✅	✅	
Asia Pacific	India - India West	✅	❌	Power BI only region
Asia Pacific	India - South India	✅	✅	Not available: Fabric App (preview)
Asia Pacific	Indonesia Central	✅	✅	Not available: Fabric App (preview)
Asia Pacific	Israel Central	✅	✅	Not available: Digital twin builder (preview) Fabric App (preview)
Asia Pacific	Japan East	✅	✅	
Asia Pacific	Japan West	✅	✅	Not available: Digital twin builder (preview) Fabric App (preview)
Asia Pacific	Korea Central	✅	✅	
Asia Pacific	Korea South	✅	❌	Power BI only region
Asia Pacific	Malaysia West	✅	✅	Not available: Fabric App (preview)
Asia Pacific	New Zealand North	✅	✅	Not available: Fabric App (preview)
Asia Pacific	Taiwan North	✅	✅	Not available: Fabric App (preview)
Asia Pacific	Taiwan Northwest	✅	✅	Not available: Fabric App (preview)

Related content

- Buy Fabric capacity in Azure
- Find your Fabric home region

Last updated on 10/01/2026

Enable Microsoft Fabric for your organization

When you enable Microsoft Fabric for your organization, users can create Fabric items such as lakehouses, notebooks, and warehouses. Two things make this possible: a tenant configuration that allows users to create Fabric items, and capacity that supports Fabric workloads. A Fabric administrator controls the tenant configuration through the **Users can create Fabric items** setting.

You also need capacity that supports Fabric workloads. Fabric workloads can run on an Azure F SKU capacity, an existing Power BI Premium P SKU capacity, or a Fabric trial capacity. Per-user licenses such as Power BI Pro and PPU don't provide Fabric capacity. If your organization needs an Azure F SKU capacity, see buy Fabric capacity in Azure.

This article explains how to configure the tenant so users can create Fabric items.

> **Note**
>
> - Power BI is part of Fabric. The Power BI portal mentioned in this article is now the Fabric portal.
> - Fabric availability is restricted in some regions. For more information, see Fabric region availability.

Prerequisites

To enable Fabric, you need to have the *Fabric administrator* Microsoft Entra role.

Enable Fabric for your tenant

When you enable Fabric using the tenant setting, users can create Fabric items in that tenant, unless capacity admins turn it off for a specific capacity.

In your tenant, you can enable Fabric for:

- **The entire organization** - In most cases your organization has one tenant, so selecting this option enables it for the entire organization. In organizations that have several tenants, if you want to enable Fabric for the entire organization, you need to enable it in each tenant.

- **Specific security groups** - Use this option to enable Fabric for specific users. You can either specify the security groups that Fabric will be enabled for, or the security groups that Fabric won't be available for.

Follow these steps to enable Fabric for your tenant.

1. In Fabric, select **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. Under *Microsoft Fabric*, expand **Users can create Fabric items**.
2. Enable the **Users can create Fabric items** switch.
3. (Optional) Use the **Specific security groups** option to enable Fabric for specific users. You can also use the **Except specific security groups** option, to exclude specific users.
4. Select **Apply**.

> **Note**
>
> The *Delegate settings to other admins* option, isn't available because it's automatically delegated to capacity admins.

Enable Fabric for a capacity

Consider the Fabric setting at the tenant level a recommendation for the entire organization. Capacity admins can override this setting, depending on their needs. For example, Fabric can be enabled for all the users in your organization. However, for security reasons your organization decided to disable Fabric for a specific capacity. In such cases, Fabric can be disabled for that capacity.

Follow these steps to enable Fabric for a specific capacity.

1. In Fabric, select **OneLake catalog** > **Govern** > **Capacities**.
2. Next to the capacity name, select the **More options** menu > **Settings**.
3. Go to **Delegated tenant settings** > **Microsoft Fabric** > **Users can create Fabric items**.
4. Check the **Override tenant admin selection** checkbox and verify that the **Users can create Fabric items** setting is enabled.
5. (Optional) Use the **Specific security groups** option to enable Fabric for specific users. You can also use the **Except specific security groups** option, to enable Fabric for the capacity,

and exclude specific users.

6. Select **Apply**.

Can I disable Microsoft Fabric?

To disable Fabric, you can turn off the *Microsoft Fabric* admin switch. After disabling Fabric, users will have view permissions for Fabric items. If you disable Fabric for a specific capacity while Fabric is available in your organization, your selection will only affect that capacity.

Considerations for Fabric items

In some cases, users that don't have Fabric enabled will be able to view Fabric items and icons.

Users that don't have Fabric enabled, can:

- View Fabric items created by other users in the same workspace, as long as they have at least read-only access to that workspace.
- View Fabric icons in capacities where other users have Fabric enabled, as long as they have at least read-only access to that capacity.

Related content

- Admin overview

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Understand Microsoft Fabric admin roles

Several roles work together to administer Fabric for your organization. Assign most admin roles in the Microsoft 365 admin portal or by using PowerShell. Assign capacity and domain admin roles when you create the corresponding capacity or domain. To learn more about each admin role, see About admin roles. To learn how to assign admin roles, see Assign admin roles.

Microsoft 365 admin roles

This section lists the Microsoft 365 admin roles and the tasks they can perform.

- **Global administrator**
  - Unlimited access to all management features for the organization
  - Assign roles to other users
- **Billing administrator**
  - Manage subscriptions
  - Purchase licenses
- **License administrator**
  - Assign or remove licenses for users
- **User administrator**
  - Create and manage users and groups
  - Reset user passwords

Power Platform and Fabric admin roles

As a Power Platform or a Fabric admin, you have full access to all the Fabric management tasks.

- **Power Platform administrator** or **Fabric administrator**
  - Enable and disable Fabric features
  - Report on usage and performance
  - Review and manage auditing

These roles provide access to the admin experience in Govern and control over organization-wide Fabric settings, usage metrics, and admin features, except for licensing. By assigning users to dedicated Fabric admin roles, organizations can grant the permissions needed to administer Fabric without granting full Microsoft 365 admin rights.

Capacity admin roles

As a capacity admin, you can perform these tasks on the capacity you're an admin of.

- **Capacity administrator**
  - Assign workspaces to the capacity
  - Manage user permission to the capacity
  - Manage workloads to configure memory usage

Domain admin roles

Domain administration distributes governance responsibilities to business owners or designated experts for a specific domain. Fabric administrators assign domain admins when they create or edit a domain.

- **Domain admin**
  - View and edit the domains and subdomains they're assigned to
  - Update domain descriptions and images
  - Assign and update domain contributors
  - Associate workspaces with the domain
  - Override tenant settings that a Fabric administrator delegates to the domain

Domain admins can't create or delete top-level domains, rename domains, or add or remove other domain admins. For complete permissions and procedures, see Domain roles.

Assign users to Fabric admin roles

Microsoft 365 user admins assign users to the Fabric administrator or Power Platform administrator roles in the Microsoft 365 admin portal, or by using a PowerShell script.

Microsoft 365 admin portal

To assign users to an admin role in the Microsoft 365 admin portal, follow the instructions in Add an admin.

PowerShell

You can also assign users to roles by using PowerShell. To assign users to an admin role using PowerShell, follow the instructions in Assign admin roles to Microsoft 365 user accounts with PowerShell.

Related content

- Administration overview
- What is the admin monitoring workspace?

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Get started with Fabric governance

Microsoft Fabric governance helps you find, organize, protect, and track your organization's data so the right people can use the right data safely. Governance spans many capabilities, and when you're new to it, it's hard to know where to begin. Start in the OneLake catalog, where you can see everything your organization already has before you decide what to control.

When you finish, you'll have explored your organization's data in the OneLake catalog and know which Fabric governance capability to use next for each job.

Getting started with governance follows these high-level stages:

1. **Explore** your data estate in the OneLake catalog.
2. **Organize** your estate into domains and tags.
3. **Classify and protect** sensitive data.
4. **Trace** how data flows and what depends on it.
5. **Share** data safely outside your organization.
6. **Extend** governance to your broader security stack.

For background on what Fabric governance covers and why it matters, see Governance and compliance in Fabric.

Prerequisites

- A Microsoft Fabric tenant that already contains data you can browse.
- Access to at least one workspace or item, so items appear for you in the OneLake catalog. You see only the items you have permission to access.
- Purview-dependent capabilities—sensitivity labels, protection policies, and DLP—require extra Microsoft Purview licensing. Everything else works with your Fabric license.

Start in the OneLake catalog

Before you add any controls, browse what your organization already has in the OneLake catalog.

1. Open the **OneLake catalog** from the Fabric navigation pane.
2. Browse the items you can access, and filter by workspace or domain to focus on a business area.
3. Select an item for a comprehensive view of its metadata, lineage, permissions, and more.

4. Note the endorsement badges as you browse: subject matter experts review *promoted* items, and *certified* items meet a higher organizational bar.

For more information about these capabilities, see OneLake catalog and Endorsement.

After you know what's in your estate, pick the next job to tackle. Each capability in the following sections is where to go next for a specific goal.

Organize your data estate

Structure your estate so users can find the right data and you can delegate ownership.

Expand table

When you want to...	Use this capability
Group data by business area (finance, HR, sales) and delegate governance to domain owners	Domains
Apply a custom taxonomy your organization defines (project codes, data sensitivity tiers, cost centers)	Tags
Search across the whole tenant for a specific item, filtered by domain, workspace, or endorsement	OneLake catalog

Classify and protect sensitive data

Classify sensitive data so users and systems recognize it, then apply policies that act on the label. These capabilities require Microsoft Purview licensing.

Expand table

When you want to...	Use this capability
Label data as public, general, confidential, or restricted so users and automated systems can identify what needs extra care	Sensitivity labels
Restrict who can access an item based on the sensitivity label it carries	Protection policies
Block or alert on risky data movement, for example exporting an item that carries a specific label	Data loss prevention (DLP) policies

Trace how data flows and what depends on it

Find out where data comes from and what breaks when something changes.

Expand table

When you want to...	Use this capability
See every upstream source and every downstream item for a piece of data	Lineage
Preview what will break before you change or delete an item	Impact analysis

Share data safely outside your organization

Give partners, suppliers, or subsidiaries access to Fabric data without copying it or losing control of it.

Expand table

When you want to...	Use this capability
Share OneLake data with users in a different Microsoft Entra tenant without moving the data	External data sharing

Extend governance to your broader security stack

Connect Fabric to the security and auditing tools you already use.

Expand table

When you want to...	Use this capability
Feed a third-party catalog or governance tool with Microsoft Fabric metadata	Metadata scanning
Detect and respond to risky user activity in Power BI content with your existing cloud app security tool	Microsoft Defender for Cloud Apps controls
Audit who did what across your Microsoft Fabric tenant	Track user activity
Verify Microsoft Fabric meets a specific regulatory, industry, or internal standard	Standards compliance

Capabilities that require Microsoft Purview

Most Fabric governance capabilities work out of the box with your Fabric license. These capabilities require Microsoft Purview licensing:

- Sensitivity labels and information protection
- Protection policies
- Data loss prevention (DLP) policies

For licensing details and how these capabilities fit into an overall governance strategy, see Governance and compliance in Fabric.

Related content

- Governance and compliance in Fabric
- OneLake catalog
- Security in Microsoft Fabric
- Admin documentation

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

OneLake catalog overview

OneLake catalog is the central experience in Microsoft Fabric for discovering data, understanding governance status, and managing security across your Fabric environment. It's a centralized place that helps you find, explore, and use the Fabric items you need, and govern the data you own.

You access the catalog from the Fabric navigation pane. The catalog is also embedded in Microsoft Teams, Microsoft Excel, and Microsoft Copilot Studio so you can discover and act on items directly from these applications. You can also programmatically discover catalog metadata across workspaces by using the Microsoft Fabric Catalog Search REST API.

OneLake catalog is organized into three experiences—Explore, Govern, and Secure.

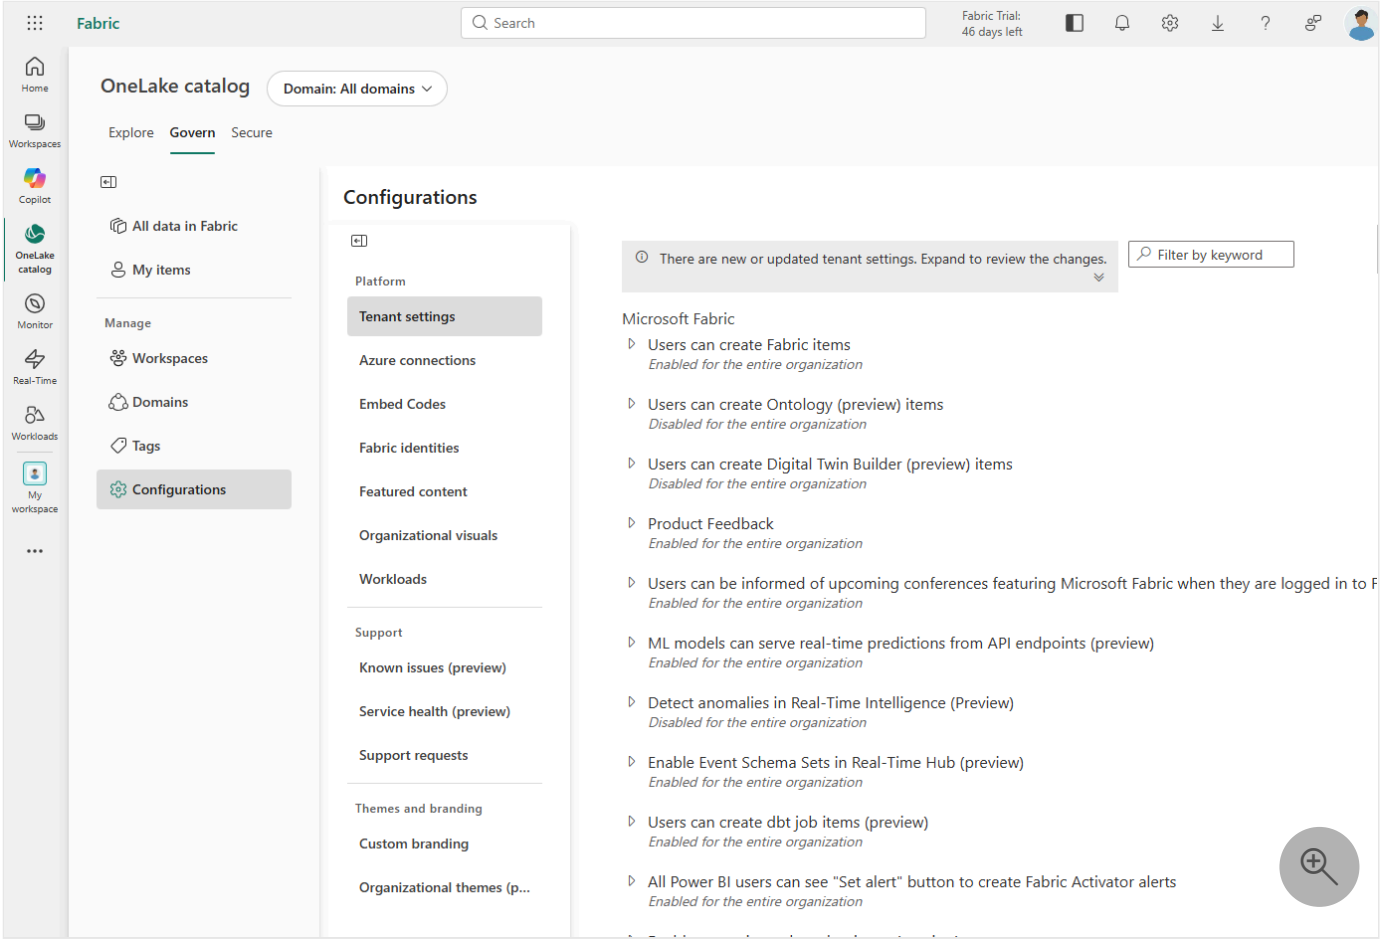

The Govern tab in the OneLake catalog is where you understand your governance status, act on recommendations, and manage administrative settings for the data you own. Explore and Secure are the two other tabs in the OneLake catalog.

Manage and govern your Fabric estate in Govern

Govern is the centralized experience for managing and governing organizational content in Microsoft Fabric.

In addition to governance capabilities such as domains, endorsements, metadata scanning, policies, and governance insights, Govern also provides access to administrative experiences including workspace management, capacity management, tenant settings, and other organization-wide controls.

As administrative capabilities continue to move from the Admin portal into Fabric, administrators can use Govern as their primary destination for managing and governing their organization's data estate.

Find and understand data

The **Explore tab** is where you find and understand Fabric items. It has an items list with an in-context item details view that makes it possible to browse through and explore items without losing your list context. It also provides selectors and filters to narrow down and focus the list, making it easier to find what you need. By default, OneLake catalog opens on the Explore tab. You can also open and work across multiple workspaces side by side using the object explorer.

For more information, see Discover and explore Fabric items in the OneLake catalog and the Microsoft Fabric Catalog Search REST API.

Understand Governance Health

The **Govern tab** helps you understand and improve the **Governance Health** of the data you own in Fabric.

Govern provides a unified experience for governing and managing your organization's Fabric estate.

In the Govern tab, **Governance Insights** explain the current governance status of your data and highlight the areas that need attention, so you always know where you stand.

Administrative experiences are now surfaced within Govern. Fabric administrators can access governance insights, tenant configuration, workspace and capacity management, policies, and other administrative controls from a single location.

Take action

After you understand your Governance Health, the Govern tab presents **Recommended Actions** you can take to improve the governance status of your data, along with guidance resources that help you act.

Govern serves both governance and administrative personas, helping organizations manage, secure, monitor, and govern their Fabric data estate from a centralized experience. A single person often holds more than one of these responsibilities—a Capacity Admin might also be a Fabric Admin, Domain Admin, Workspace Admin, or data owner—so Govern focuses on the outcome you want rather than on a single role.

For more information, see Govern your data in Fabric.

Secure data access

The **Secure tab** centralizes security management in Microsoft Fabric by providing a unified view of workspace roles and OneLake security roles across items. From a single location, you can review access, understand permissions, and manage security roles. Admins can audit permissions, view user access, and create, edit, or delete security roles to keep data access consistent.

For more information, see Secure your data.

Choose the right experience

Expand table

I want to...	Use	Where
Find, browse, and inspect Fabric items	Explore tab	OneLake catalog → Explore
Understand my Governance Health and act on Recommended Actions	Govern tab	OneLake catalog → Govern
Manage workspace and OneLake security roles, audit permissions, review user access	Secure tab	OneLake catalog → Secure
Search and explore catalog items and metadata programmatically	REST API	OneLake catalog REST API overview

Open OneLake catalog

To open OneLake catalog, select the OneLake icon in the Fabric navigation pane. Select the tab you're interested in if it isn't displayed by default.

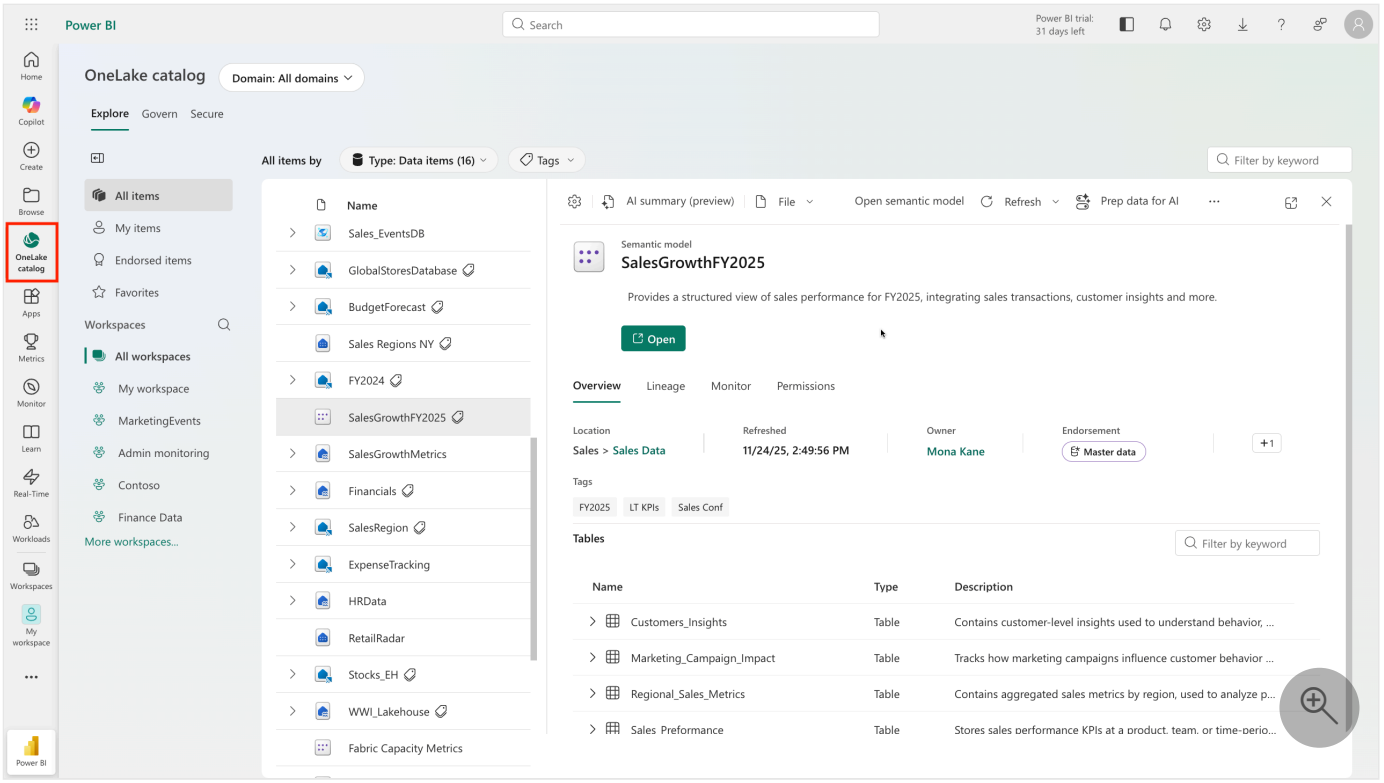

Related content

Discover

- Discover and explore Fabric items in the OneLake catalog
- Microsoft Fabric Catalog Search REST API

Govern

- Govern your data in Fabric
- Fabric domains

Secure

- Secure your data

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Find and explore data in the OneLake catalog

OneLake catalog's explore tab helps you find, explore, and use the Fabric items you need. It features an items list with an in-context item details view that makes it possible to browse through and explore items without losing your list context. The explore tab also features selectors and filters to narrow down and focus the list, making it easier to find what you need.

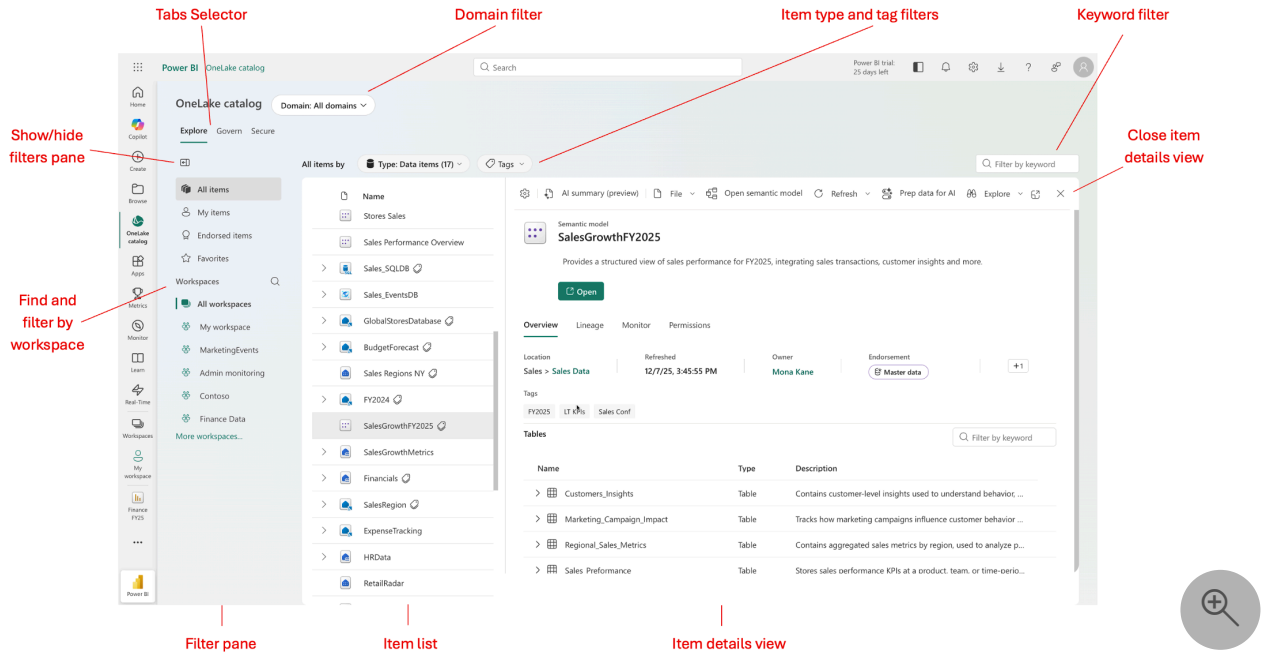

The explore tab provides

- A list of all Fabric items you have access to in your organization.
- An item details view that enables you to drill down on an item without leaving the context of the explore tab.
- Filters and selectors to help you find the content you're looking for.
- A way to scope the catalog to display only items of a particular domain.
- An options menu for item actions.

This article explains what you see in the explore tab and describes how to use it.

Open the explore tab

To open the explore tab, select the OneLake catalog icon in the Fabric navigation pane and then select the explore tab. The Explore tab surfaces lakehouses, warehouses, Fabric databases,

mirrored items, and other supported item types.

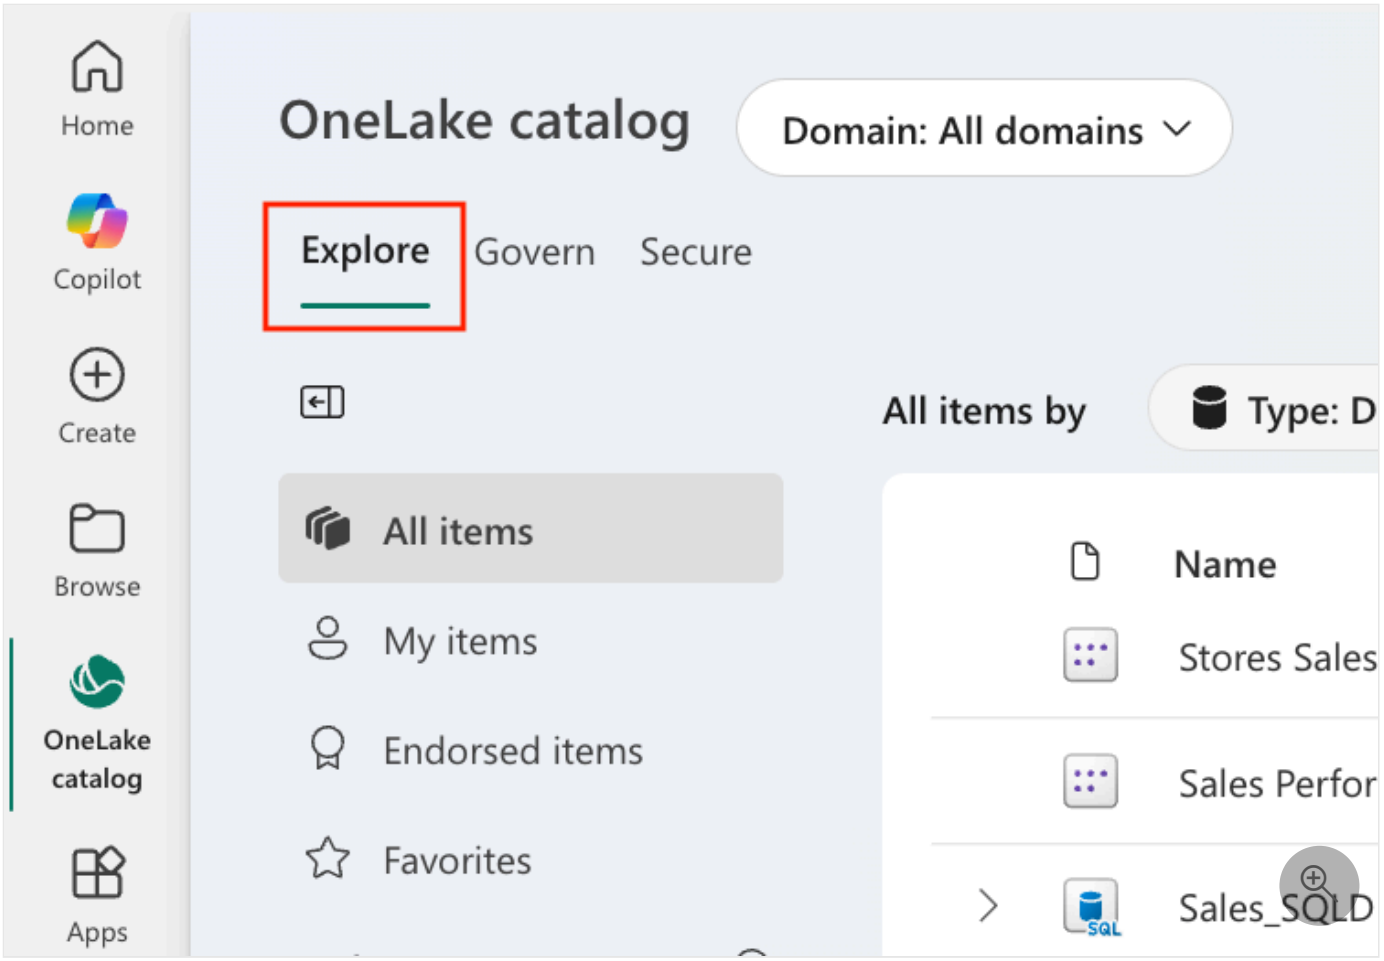

Find items in the items list

The items list displays all the Fabric items you have access to, or can directly request access to. See NOTE for details.

- To shorten the list, use the explore tab's filters and selectors.
- To view item details, select the item.
- To view an item's options menu, hover over the item and select the three dots that appear.

The following table describes the list columns.

Expand table

Column	Description
Name	The item name. To explore item details, select the name.
Type	The item type.

Column	Description
Owner	Item owner.
Refreshed	Last refresh time of data items (rounded to hour, day, month, and year. See the details section in item's details for the exact time of the last refresh).
Location	The workspace the item is located in.
Endorsement	Endorsement status.
Sensitivity	Sensitivity, if set. Select the info icon to view the sensitivity label description.

> **Note**
>
> The explore tab also lists the following items even if you don't have access to them. This listing enables you to request access.
>
> - Semantic models that are configured as discoverable.
> - Power BI reports whose associated semantic model violates a data loss prevention policy that restricts access.

Scope the catalog to a particular domain

If your organization defines domains, use the OneLake catalog's domain selector to select a domain or subdomain. The explore tab displays only workspaces and items belonging to the selected domain or subdomain. Your selection persists for subsequent sessions.

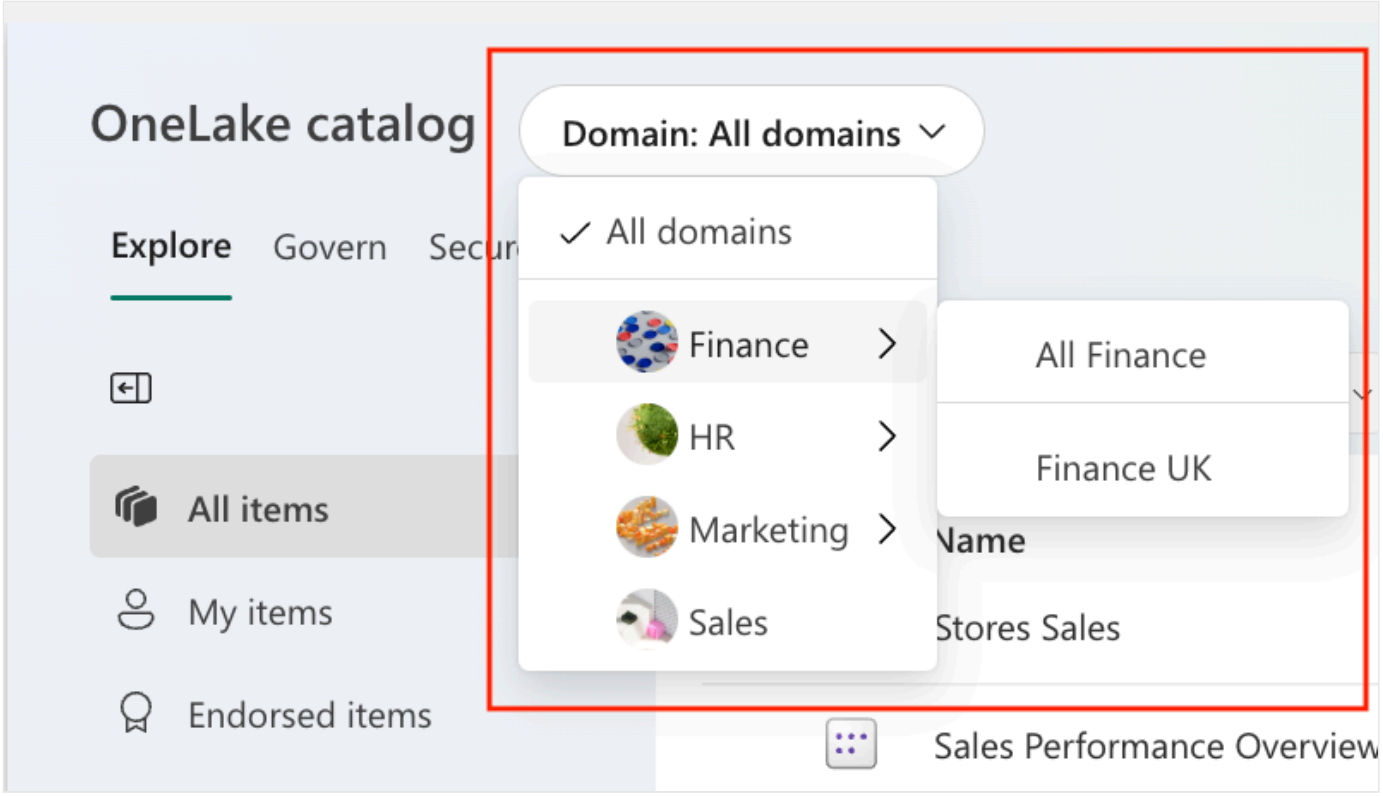

For more information about domains, see Fabric domains.

Filter the items list

Use the filters in the filter pane and the item type and tag selectors at the top of the items list to narrow down the items list. The list displays just the types of items you're interested in. Select items to explore them in more detail.

The filters pane contains several predefined filters. It also enables you to filter by workspace.

Filter items by predefined filters

The filters pane contains several predefined filters. These filters operate across workspaces. They show all the applicable items from across your tenant or domain, as long as you select **All workspaces** as the workspaces filter. They're described in the following table.

Expand table

Filter	Description
All Items	Items that you're allowed to find.
My items	Items that you own.
Endorsed items	Endorsed items in your organization that you're allowed to find. Certified data items are listed first, followed by promoted data items. For more information about endorsement, see the

Filter	Description
	Endorsement overview
Favorites	Items you marked as favorites.

Find items by item type category

Fabric items are categorized into buckets based on what they are for and where they fit into the task flow. For example, items that contain data, such as lakehouses and semantic models, are categorized as *Data types*, while reports and dashboards are categorized as *Insight types*. Use the item type category selector to select the item type category you're interested in.

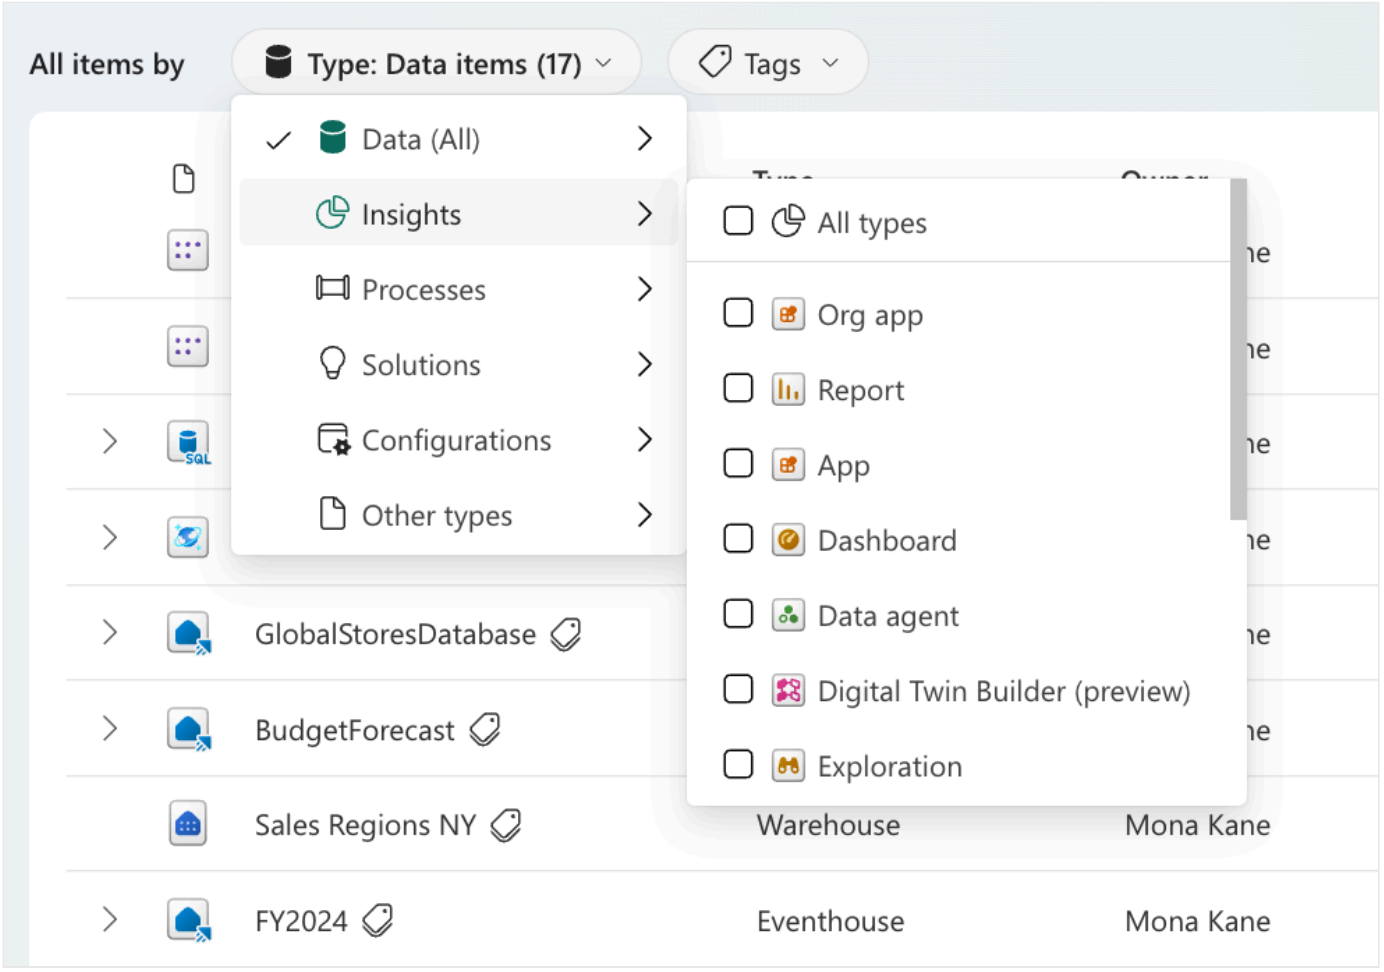

Find items by tag

Use the tags selector to display items tagged with the tags you select.

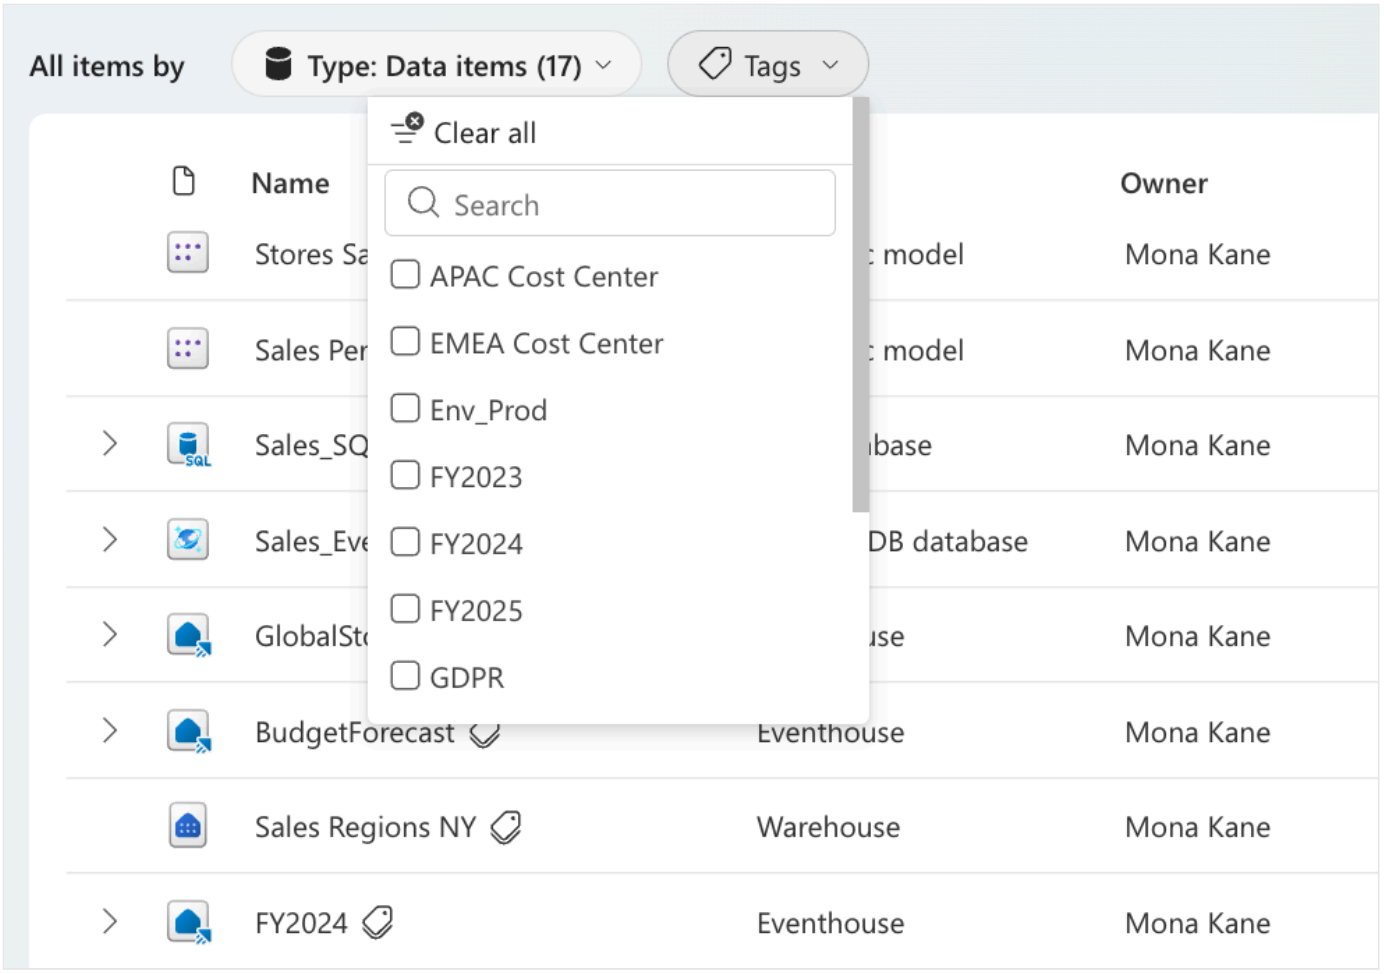

> **Note**
>
> The tag selector is only visible if tags are enabled in your organization.

Filter items by workspace

Related items are often grouped together in a workspace. To find items by workspace, find and select the workspace you're interested in under the **Workspaces** heading in the filters pane. The items you're allowed to see in that workspace appear in the items list.

> **Tip**
>
> The workspace filter can help you find the workspace you're looking for. Select the magnifying glass icon to open the filter.

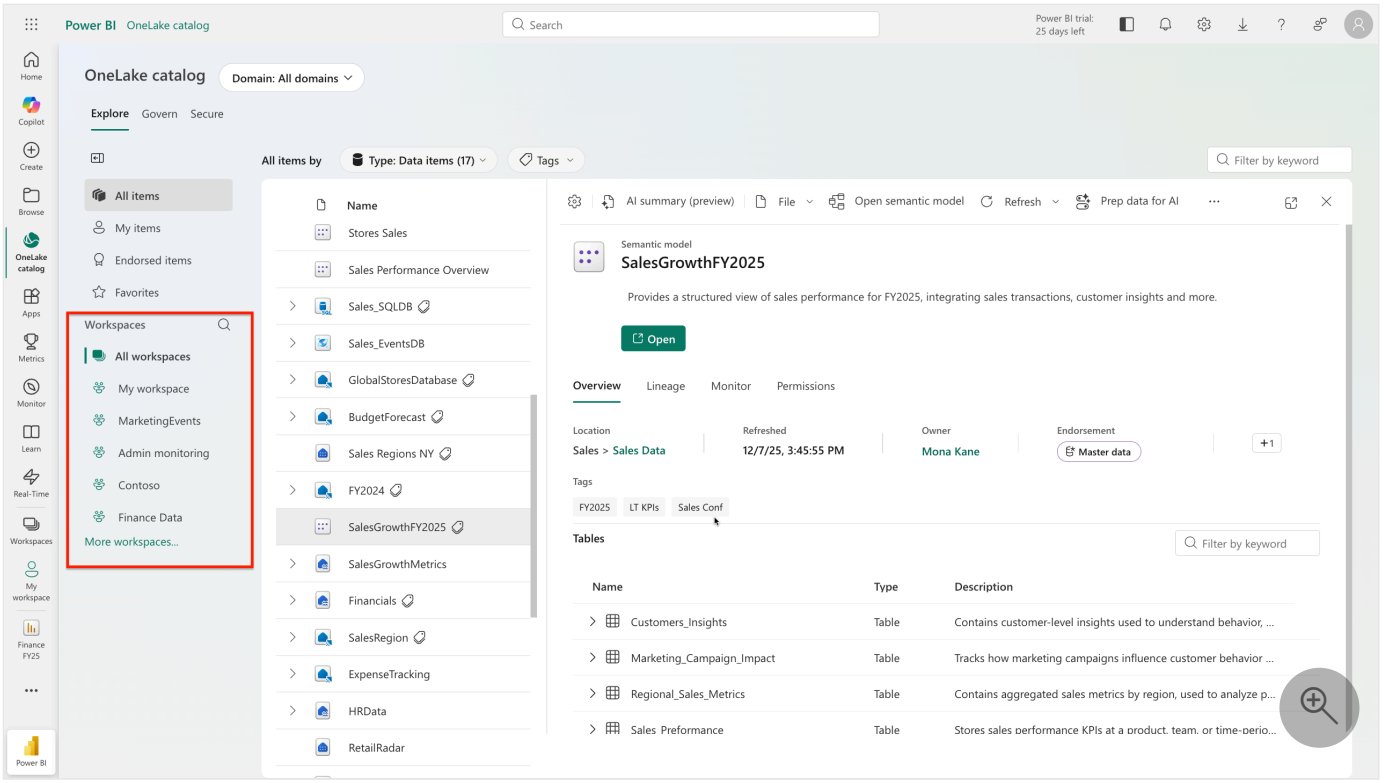

The preceding image shows the *Sales Org* workspace selected in the Workspaces section. Because the item classification selector is set to *Data types*, only semantic models are shown (as no other type of data items are present in this workspace). The semantic model *Contoso FY21 goals* is selected, and its details are shown in the item details view. To explore the details of other items in the list, just select them. The item details view stays open until you dismiss it. This makes it easy to browse the details of items one after another, without losing your list context.

> **Note**
>
> Generally, the **Workspaces** section only displays workspaces you have access to. However, workspaces you don't have access to might be listed if the workspace contains items that you do have access to (through explicitly granted permissions, for example) or that are configured as discoverable. If you select such a workspace, only the items you have access to and any discoverable items appear in the items list.

Open an item's options menu

Each item in the items list has an options menu that enables you to do things, such as open the item's settings, manage item permissions, and more. The options available depend on the item and your permissions on the item.

To display the options menu, hover over the item and select **More options (...)**.

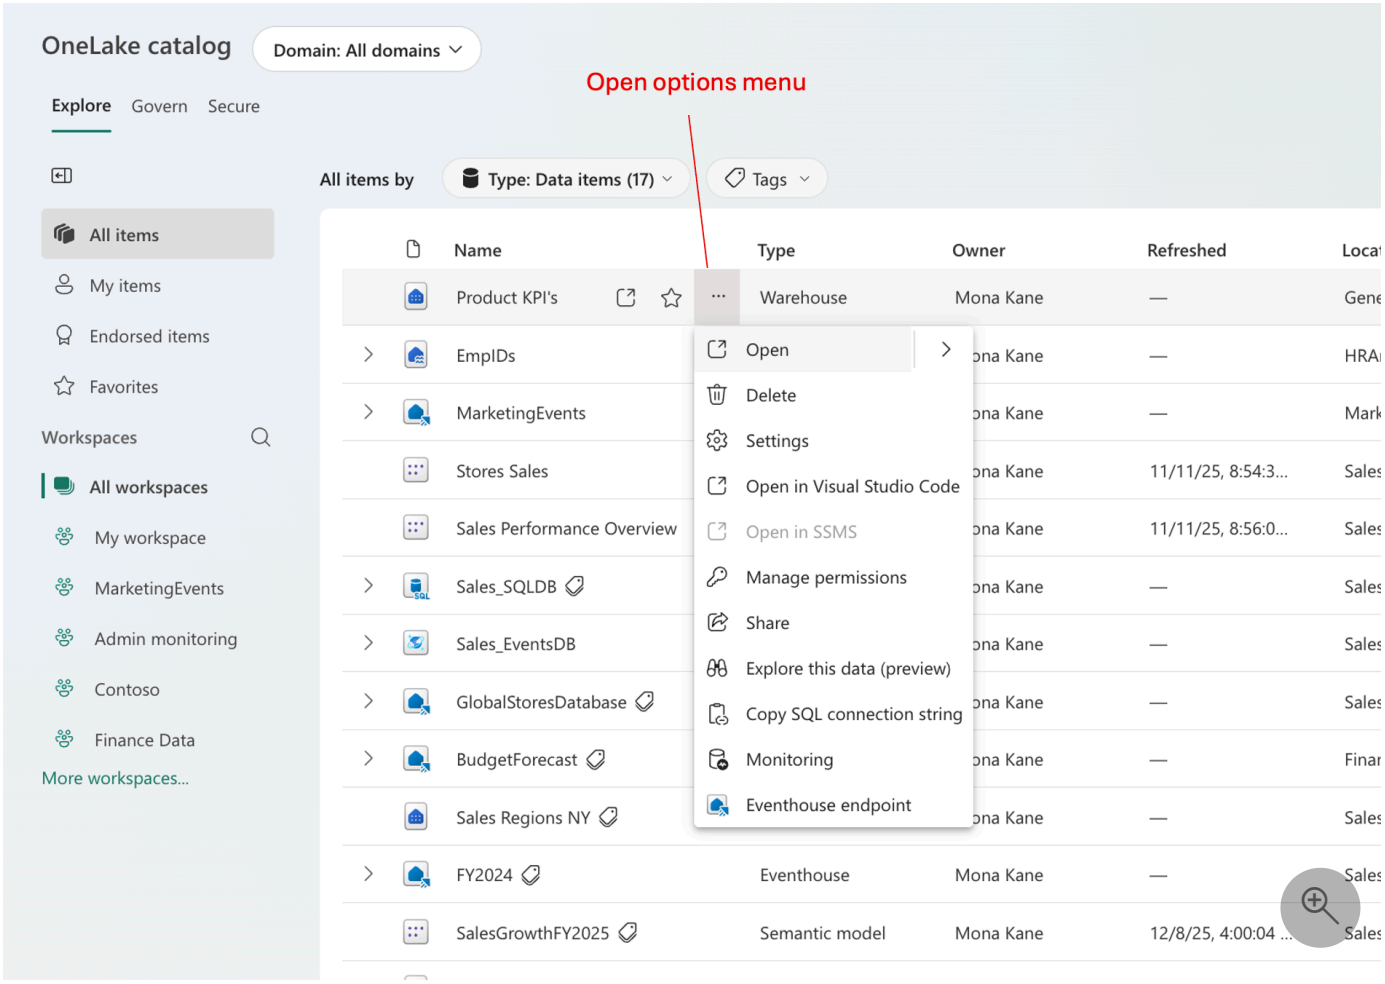

Open an item for editing or viewing

You can open the item for editing or viewing (depending on the item type) from both the items list and from the item details view.

- To open the item for editing or viewing from the items list, hover over the item and select the **Open** icon that appears.

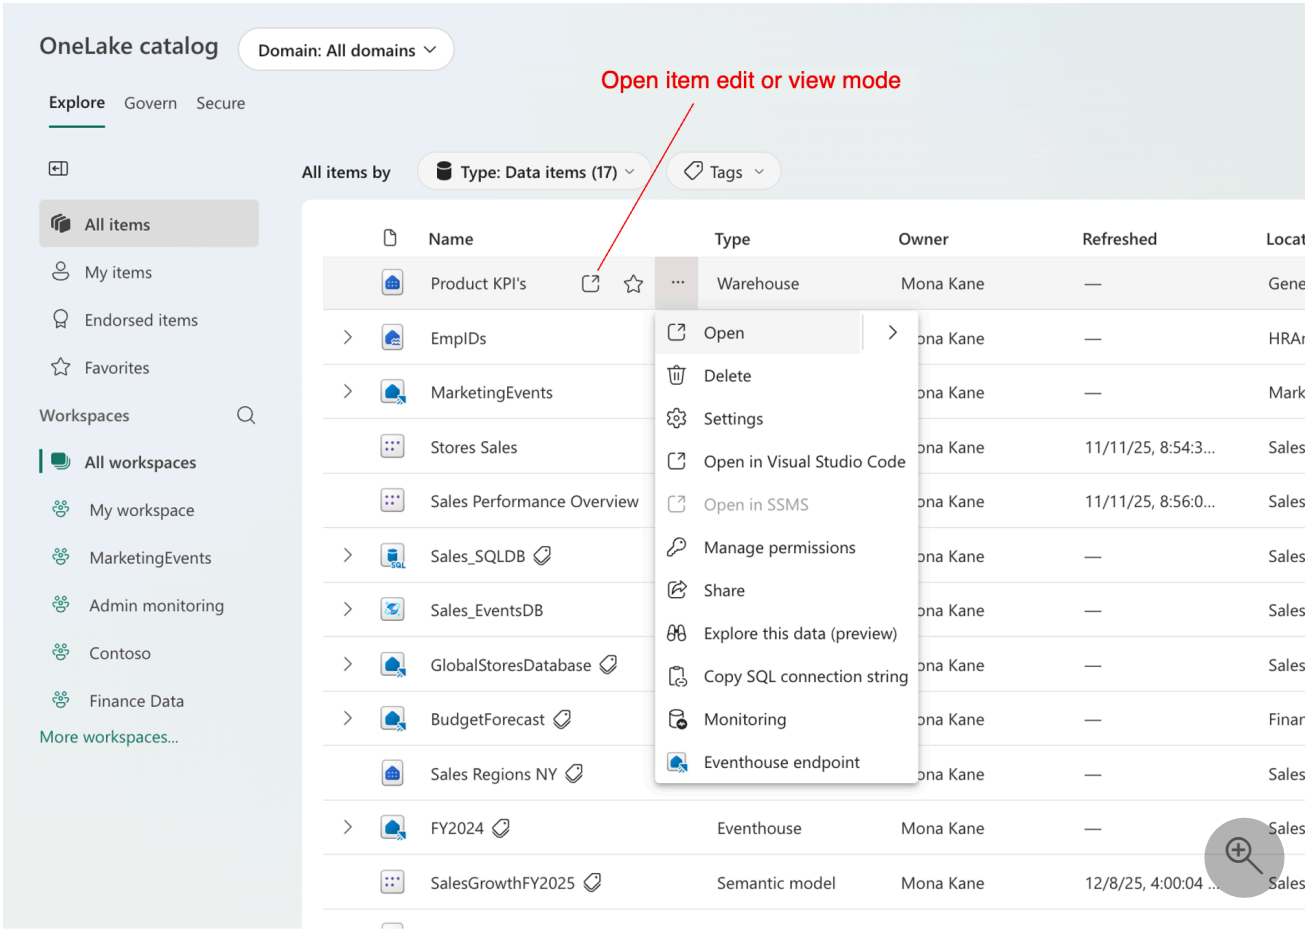

- To open the item for editing or viewing from the item details view, select the **Open** button in the view.

Create AI Auto-Summary for Semantic Models (Preview)

The Auto-Summary for semantic models uses AI to generate a high-level summary that helps you quickly understand an item's purpose and main characteristics without opening the item or reviewing its full metadata. It makes it easier to understand unfamiliar items and compare them directly in the OneLake catalog explorer.

The summary is created based on the item's metadata and structure. Users with the appropriate Copilot capacity and permissions can generate the summary from the quick actions in the main explore tab or directly from the semantic model's item details page. Each time you return to the catalog, a new summary can be generated so you always see the most up-to-date version.

After a summary is generated, you can generate another version, copy the text for use elsewhere, or provide feedback on the quality.

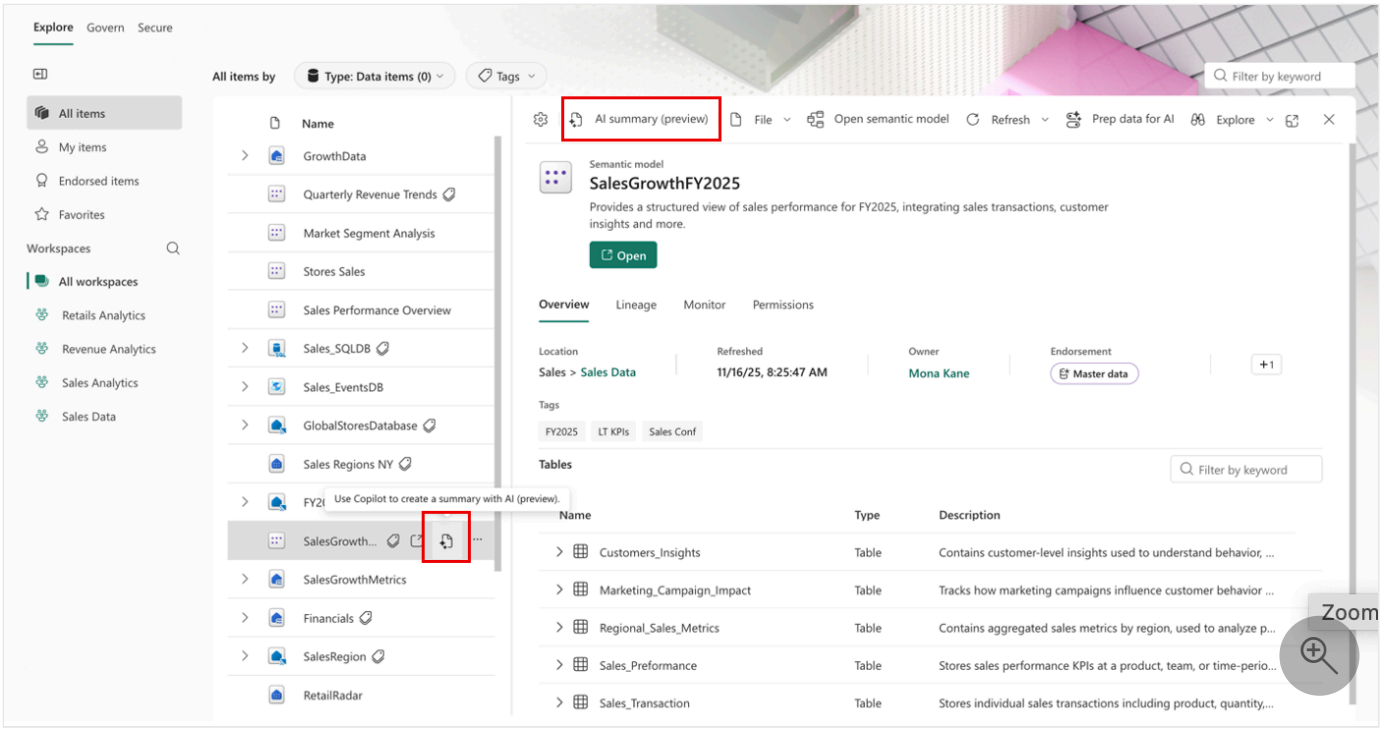

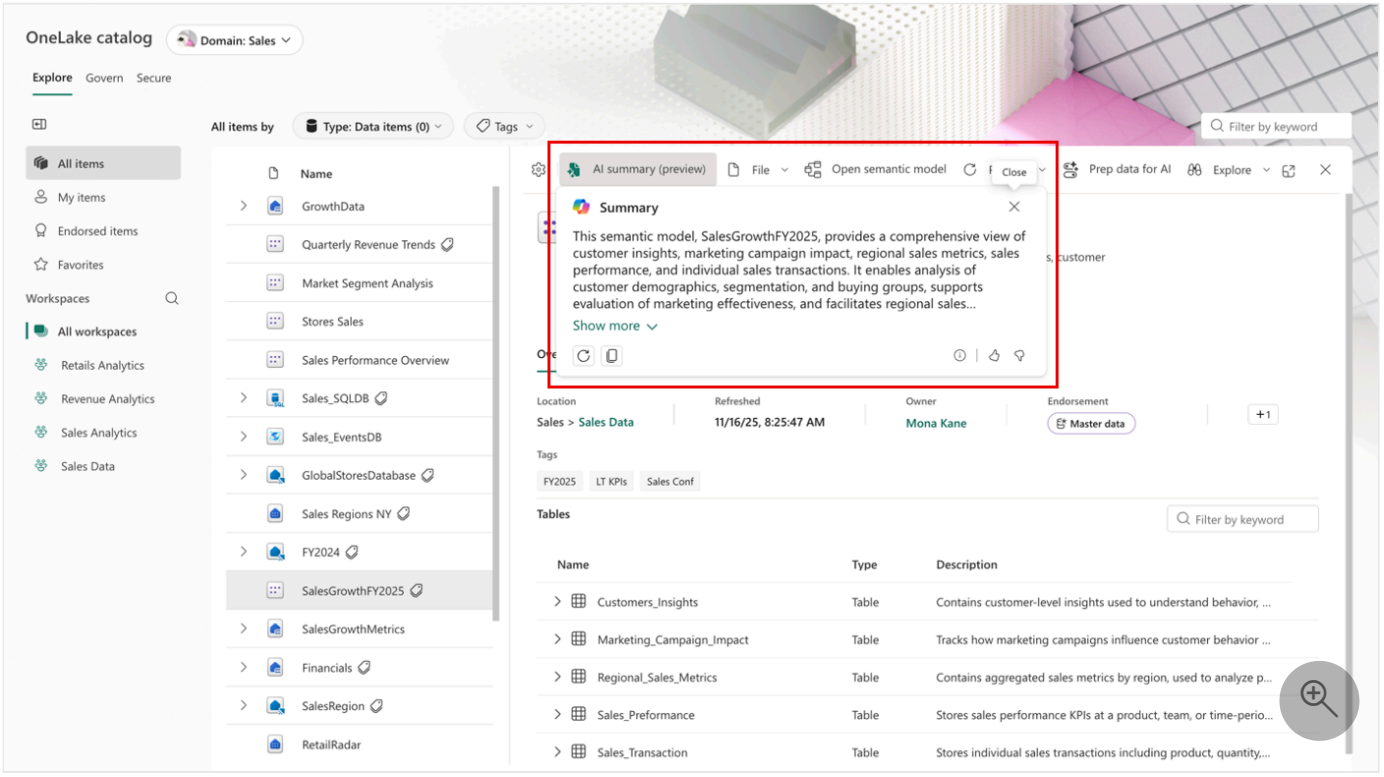

View item details

When you select an item from the items list, its detailed information appears in a dedicated pane alongside your current view. This in-context view allows you to quickly browse through item metadata, lineage, and permissions without losing your place in the item list. You can easily compare and find the most suitable items. For a comprehensive understanding of how item details are presented and the various ways you can access them, see View item details.

Considerations and limitations

- Streaming semantic models are retired and thus aren't shown. For more information, see Announcing the retirement of real-time streaming in Power BI.

Related content

- OneLake catalog overview
- View item details
- Monitor and improve your data's governance posture with the OneLake catalog
- Endorsement
- Fabric domains
- Lineage in Fabric
- Monitor hub

Last updated on 02/18/2026

View item details in the OneLake catalog

The item details page in the **OneLake catalog** explorer provides a comprehensive view of each item, including an item's metadata, lineage, permissions, and more. This page is designed to help users find, explore, and use items effectively across their organization.

The item details view includes several tabs that help you explore and manage the selected item. It's especially useful for:

- **Item owners**: Manage permissions, view refresh history, and track lineage to monitor and manage items.
- **Developers and Analysts**: Access trust signals for an item, identify high-quality data to use as starting points, and avoid creating redundant items.
- **Consumers**: Discover reliable apps, reports, and dashboards to support everyday work.

The item details page adapts to how you access it, offering two convenient ways to explore items:

- **In-context view**: When you click on an item's name in OneLake catalog explorer, its details appear as a pane next to your current filtered list of items. This allows you to seamlessly explore and compare items, helping you find the most suitable one. You can expand this pane for a full-page view, then retract or close it to return to your full list.
- **Full-page view**: If you open an item's details from outside the Catalog, for example by selecting "View Details" from an item's context menu in Workspaces, or by using a direct URL, the page opens in full-screen. This gives you a dedicated, comprehensive view of that single item. Since you're not coming from a filtered list in the Catalog, you won't see other items alongside it. To browse for other items, just select the OneLake Catalog entry point in the left navigation.

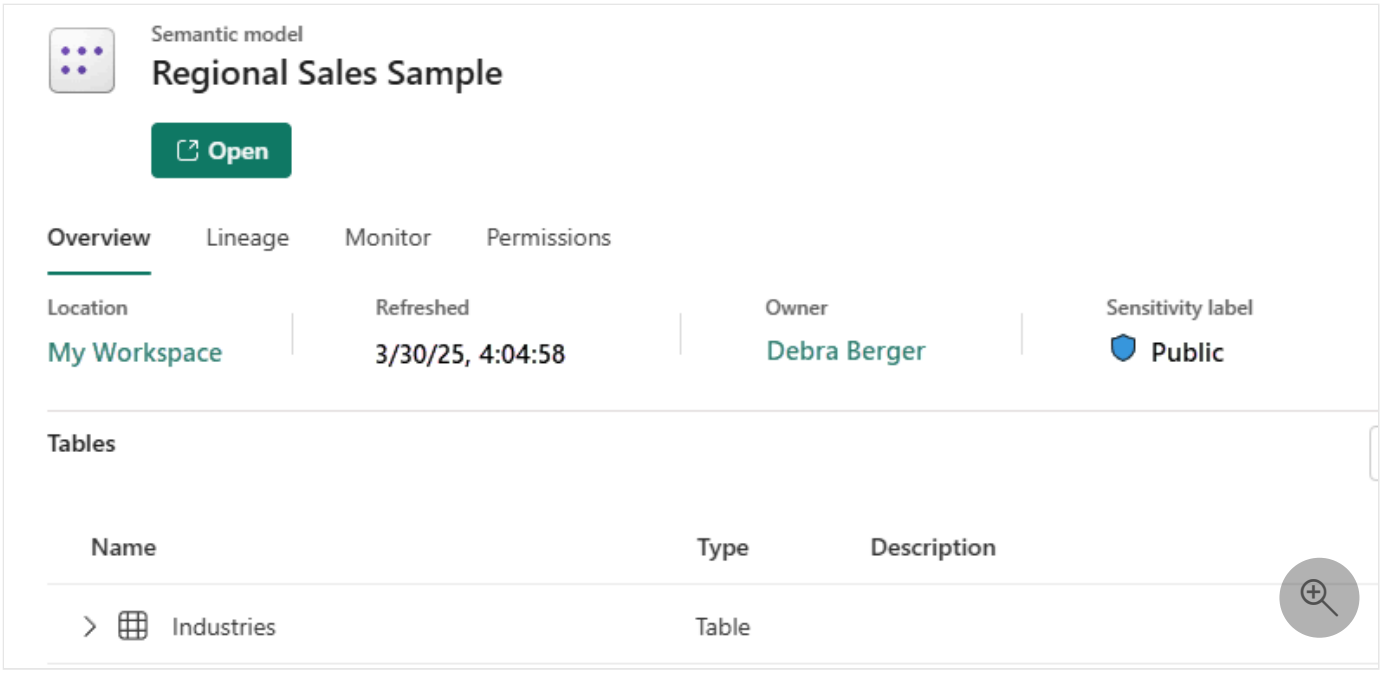

About the item details view

The item details page features multiple tabs that vary by item type (for example, semantic model, dataflow, or pipeline), each providing relevant information and functionality. This article describes each tab, its contents, and the actions you can perform on the item details page.

Overview tab

The **Overview** tab is the default view when you open an item in the OneLake catalog explorer. It shows the following information about the selected item:

- **Description**: The description associated with the item, displayed under the item name. A useful and meaningful description helps users discover the right item for them.
- **Location**: The workspace where the item is located. Selecting the workspace navigates to the workspace.
- **Refreshed**: The last time the item was refreshed or run. A red warning icon appears if the last update was unsuccessful. Selecting the icon displays the error details.
- **Owner**: Displays the item owner or contact. Selecting the name opens an email to them.
- **Sensitivity label**: If applied, the name of the sensitivity label associated with the item.
- **Tags**: If applied, the list of tags associated with the item.
- **Endorsement**: Endorsement status and the details of the endorsing user. Endorsing your item (promoting or certifying it) makes it easier for users to find and signifies it's a trustworthy data source.

Data schema in the Tables section

For data items like semantic models and lakehouses, the **Tables** section displays a structured view of the underlying table and column schema for the item.

- Use the **Filter by keyword** box to search for a specific table or column.
- To explore data in a table or column, point to it and select the binoculars icon. The Explore this data view opens, which displays the data for ad-hoc analysis.

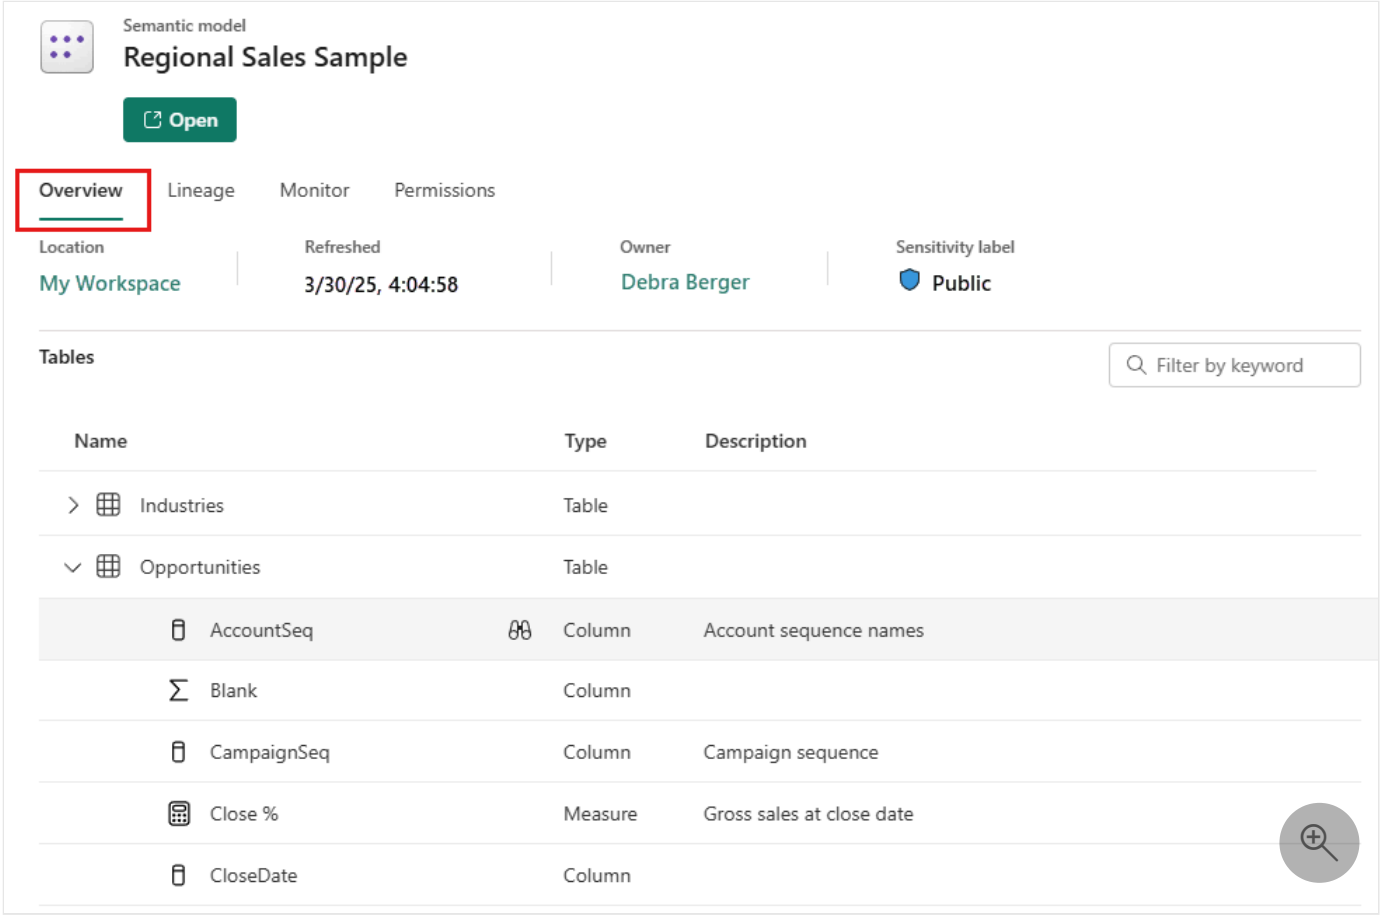

Details about semantic models

In the OneLake catalog explorer, you can view and interact with semantic models based on your permissions and whether the model is set to be discoverable.

- **Discoverable**: These are semantic models that have been made discoverable by their owner. Discoverable semantic models make it easier for users across the organization to find relevant data assets. While you can see the semantic model and its metadata, you can't see the owner details, view underlying data, or build content on top of the semantic model. You need to **request access** (Build permission) to access the semantic model's full information and capabilities.
- **Read-only Permission**: If someone shares a report or semantic model with you but doesn't grant Build permission, you might have read-only access, which limits your ability to view

and interact with the semantic model. To create content or perform more actions, you need at least Build permission.

- **Build Permission**: If you have Build permission on a semantic model, you can create new reports based on it, connect to it from Excel, and export its data.

Export semantic model data

If you have Build permission on a semantic model, you can export its data in two ways:

- **Analyze in Excel**: Select Analyze in Excel in the action bar to quickly generate an Excel workbook with the entire semantic model's content. This is the fastest way to export data for ad-hoc analysis.

- **Create a paginated report:** If you prefer to choose specific tables and columns before exporting, select **Explore this data** > **Create a paginated report**. In the paginated report editor, select the columns you need, then go to **Home** > **Export** and choose your preferred format (PDF, Excel, Word, and more).

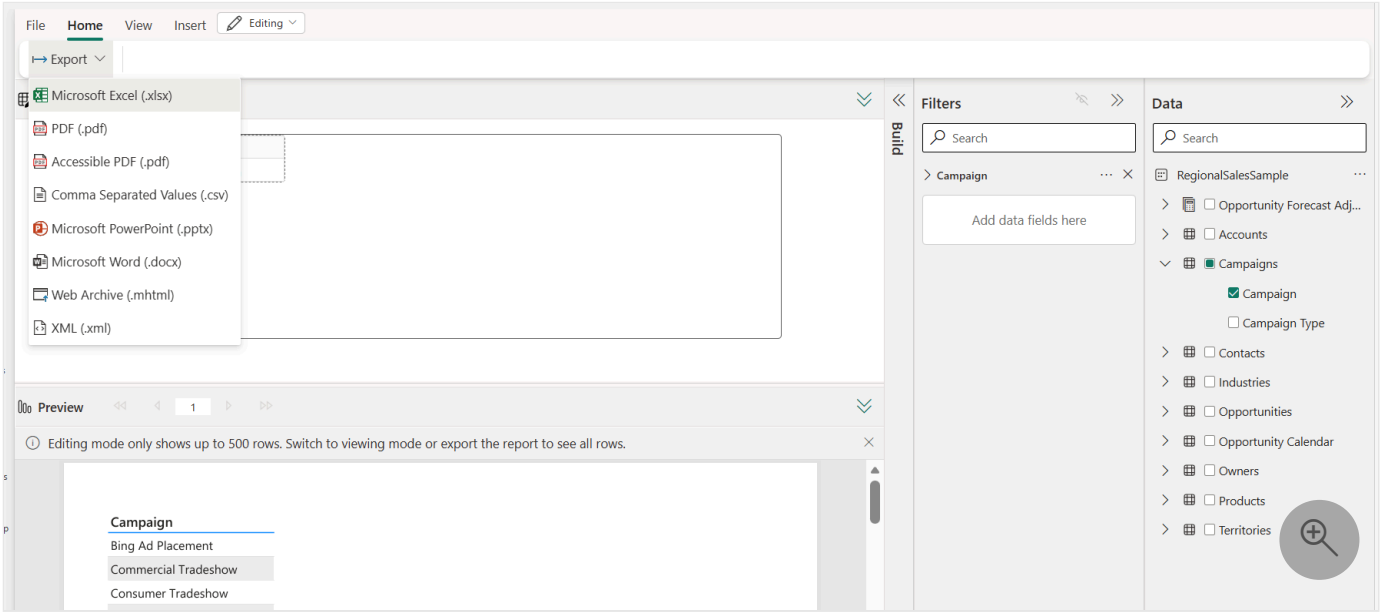

Lineage tab

The **Lineage** tab shows the item's upstream and downstream relationships. It also shows metadata associated with the item, including location, endorsement, and sensitivity information. You can view this lineage in either a list or a graph format.

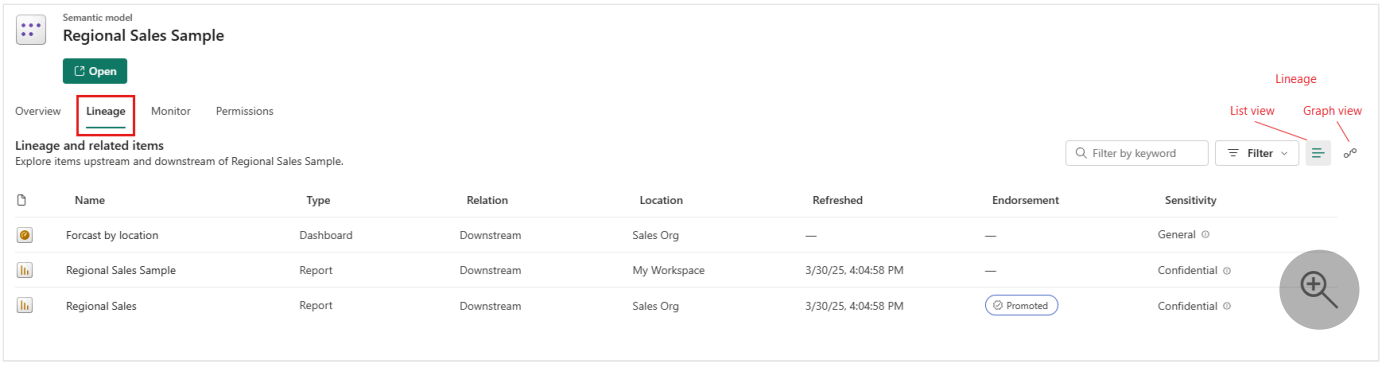

The graph view provides a visual representation of the relationships between the current item and its related items. Each related item is shown as a card containing key metadata, with the current item highlighted for context.

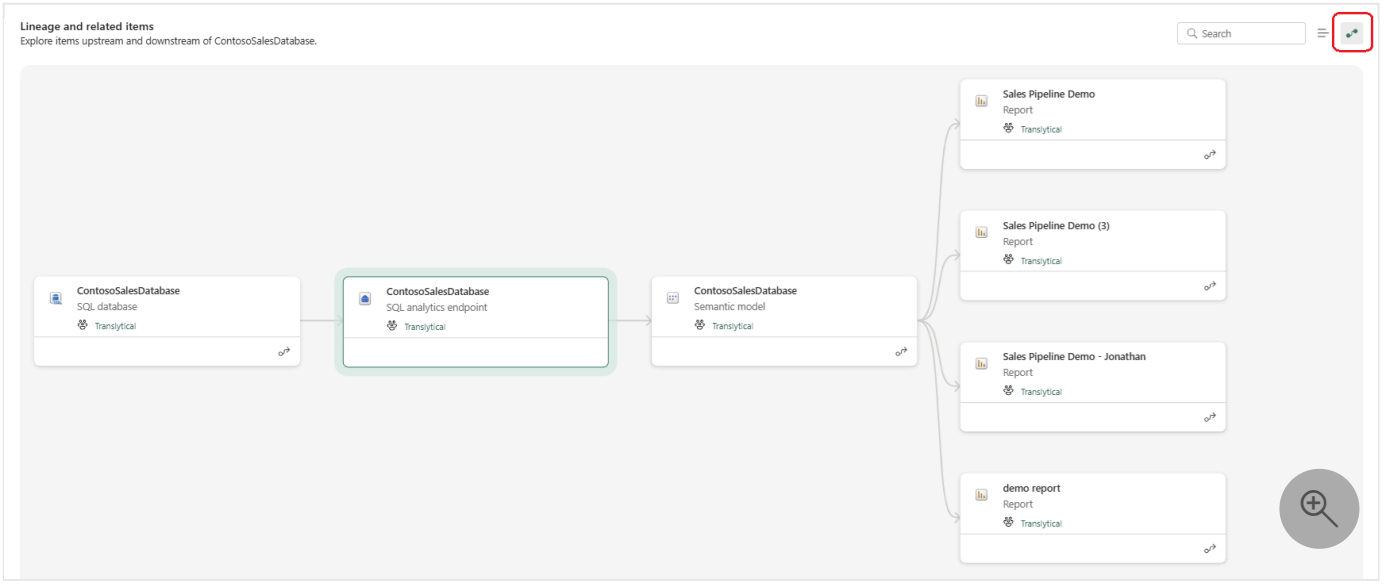

For more information, see Lineage in Fabric.

Monitor tab

The **Monitor** tab displays historical activities for the item. Select **Show** for a record to see the details about that activity. The **Monitor** tab is available for item types supported by the monitor hub.

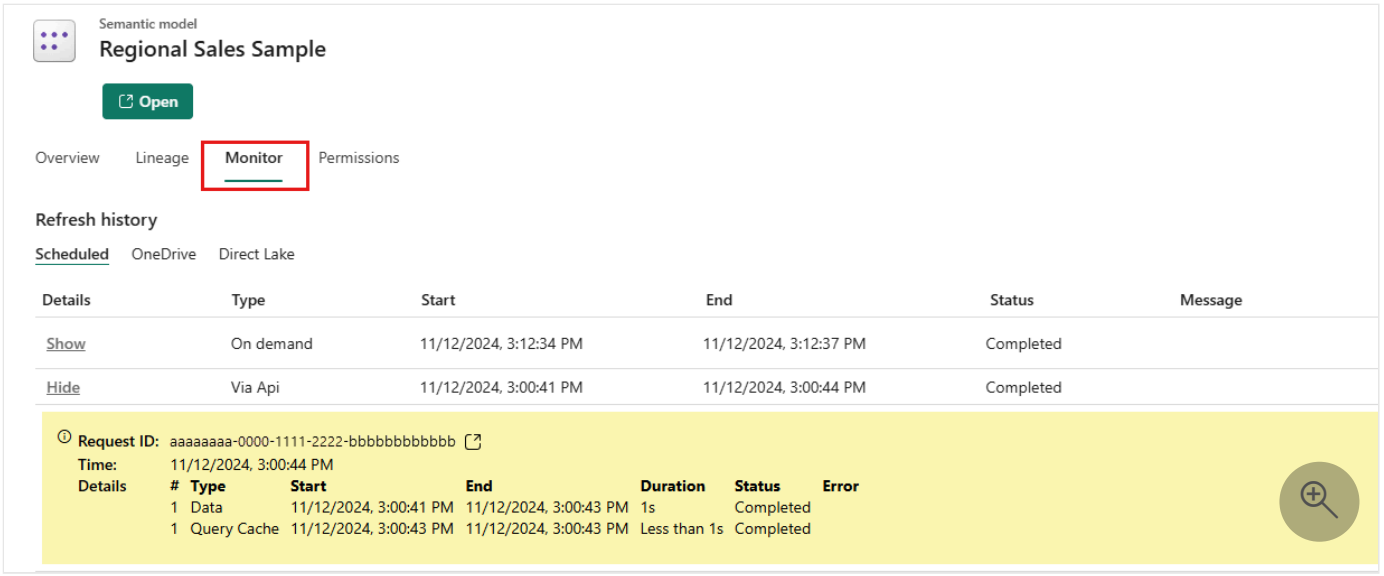

Permissions tab

The **Permissions** tab is available to you if you're a member of the Admin or Member workspace roles. It shows who has direct access to the item, any links that grant access, related items that provide inherited access, and any pending access requests. From this tab, you can manage user permissions, audit or remove access links, and approve or deny access requests.

The subtabs available on the **Permissions** tab vary slightly depending on the item type. To see the available actions for a specific row in any subtab, point to the row and select **More options (...)**.

You can use the filter to refine your search.

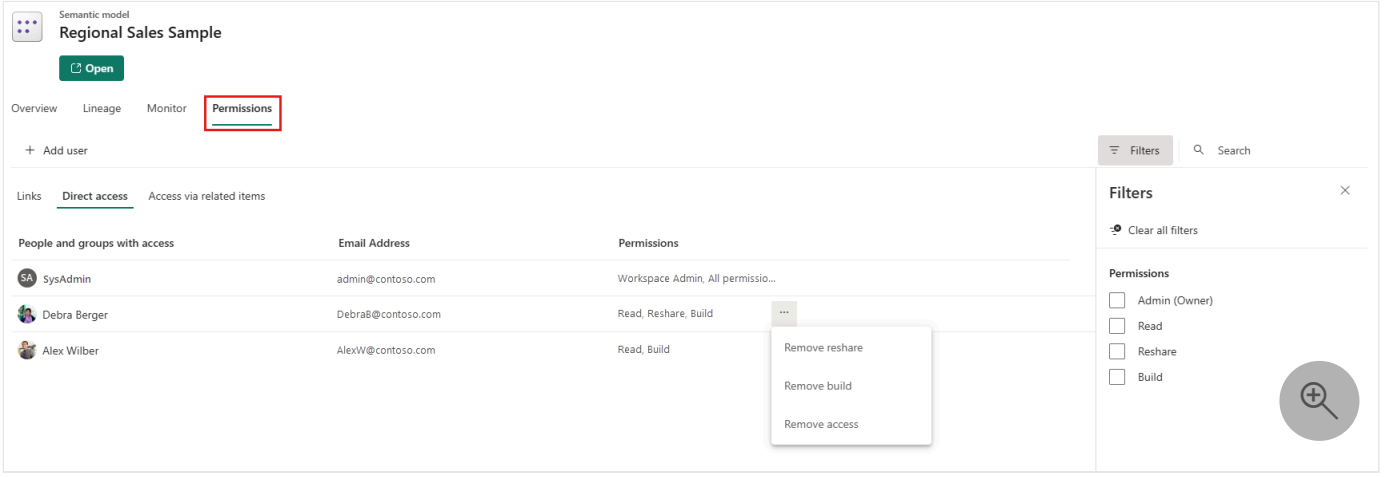

Available actions in the item details view

From within the item details page, you can perform a range of actions on an item based on its type and your permissions. The available actions are grouped into general actions that apply to all item types, and item-specific actions that are unique to data or process items.

General Actions

General actions are universally available across all item types in OneLake catalog explorer. Your access to these actions is based on your permissions.

These actions are available in:

- The item's context menu (three dots icon), which is available when you point to an item in the list.
- The item's details page, which opens when you select an item in the list.

Expand table

Action	Description	Where to find it in the item details page
Settings	Opens the settings page for the item.	**Settings** cog icon
Manage permissions	Opens the manage permissions page for the item.	**Manage permissions** key icon
Share	Allows you to grant other users or group direct access to an item, without giving access to the workspace and the rest of its items. Depending on the item type, more options are available to control user permissions.	**Share** button
Open workspace lineage	Shows the lineage relationships between all the items in the current workspace, and data sources external to the workspace one-step upstream.	**More options (...)** > **Open workspace lineage**
Impact analysis	Impact analysis allows you to assess the potential impact of changes to this item on downstream workspaces and items that depend on it.	**Lineage** > **Impact analysis**
Chat in Teams	Invite people to start chatting in Teams. People you invite receive a Teams chat message with a link to this item's details page. If they have access to the item, the link opens this details page in Teams.	**More options (...)** > **Chat in Teams**

Item-specific actions

Item-specific actions are unique to particular artifact types, providing specialized functionality. They're categorized by whether the item is primarily a data source or a processing artifact.

These actions are available in:

- The item's context menu (three dots icon), which is available when you point to an item in the list.
- The item's details page, which opens when you select an item in the list.

Data Items

This section covers actions unique to Semantic Models, SQL Analytics Endpoints, Warehouses, SQL Databases, and Warehouse Snapshots.

Expand table

Item Type	Action	Description	Where to find it in the item details page
Common Data Item Actions	**Refresh now**	Starts an immediate refresh of the item.	**Refresh** button > **Refresh now**
	Schedule refresh	Opens the settings page where you can configure scheduled refresh.	**Refresh** button > **Schedule refresh**
Semantic Model	**Download this file**	Downloads the .pbix file for this semantic model.	**Download** button
	Version History	With version history, self-service users can recover from the most critical mistakes when editing their semantic models on the web.	**File** > **Version history**
	Explore this data	A lightweight tool for quick, ad-hoc data analysis. Explore the data to create matrix/visual pairs without building a full report. To Explore a specific table or column of a semantic model, point to it in the **Overview** tab and select the binoculars icon.	**Explore this data** button > **Explore this data** (or binoculars icon on Overview tab)
	Auto-create a report	Automatically generates a summary report with key insights and visuals from the selected semantic model, which is useful for quick data checks.	**Explore this data** button > **Auto-create a report**
	Create a blank report	Opens the report editing canvas to a new report built on the semantic model. Saving your new report saves it in the workspace that contains the semantic model if you have write permissions for that workspace. If you	**Explore this data** button > **Create a blank report**

Item Type	Action	Description	Where to find it in the item details page
		don't have write permissions for the workspace, the new report saves in *My workspace*.	
	Create from template	You can create a report template that users can use to get started building their own reports based on your semantic model. This template is a regular report designed to be used as a starting point. When you save it, you should add the suffix "(template)" to the report name, for example, *Monthly Sales (template)*. When a user selects **Create from template** in the details page, a copy of the template is created in the user's *My workspace* and opens in the report editing canvas. Report templates are also easily identifiable in the list of related reports in the semantic model details view.	**Explore this data** > **From a template**.
	Create a paginated report	Opens the paginated report editing canvas. You can then export it using the rich export functionality which retains the applied formatting & styling, or you can save it to a workspace of your choice.	**Explore this data** button > **Create a paginated report**.
	Analyze in Excel	Creates an Excel workbook containing the entire semantic model, allowing you to analyze it using PivotTables, Pivot Charts, and other Excel features.	**Analyze in Excel** button
	Prep data for AI	Prepares the semantic model's data for use with AI capabilities.	**More options (...)** > **Prep data for AI**
	Best practice analyzer	Runs a preconfigured notebook that analyzes the semantic model against a set of best practices to identify potential issues and provide recommendations for optimization.	**More options (...)** > **Best practice analyzer**
	Memory analyzer	Runs a preconfigured notebook that provides insights into the semantic model's memory/storage consumption, helping to identify and resolve performance bottlenecks.	**More options (...)** > **Memory analyzer**
	Community Notebooks	Opens a gallery of community notebooks that can be used to further analyze and extend the semantic model.	**More options (...)** > **Community notebooks**
SQL Analytics Endpoint,	**Open in Visual Studio**	Shows the server connection information to connect the SQL database server externally and offers to	Specific menu option

Item Type	Action	Description	Where to find it in the item details page
Warehouse, Warehouse Snapshot	Code	directly open VS Code.	
	Open in SSMS	Shows the server connection information to connect the SQL database server externally in SQL Server Management Studio.	Specific menu option
	Explore this data	A lightweight tool for quick, ad-hoc data analysis. Explore the data allows you to create matrix/visual pairs without building a full report.	**Explore this data** button
	Copy SQL connection string	Copies the connection string needed to connect from external tools.	Specific menu option
	Analyze in Excel	Creates an Excel workbook containing the item, allowing you to analyze it using PivotTables, Pivot Charts, and other Excel features.	**Analyze in Excel** button
SQL DB	Refresh Git sync status	Refreshes the synchronization status with the connected Git repository.	Specific menu option
	Open Performance summary	Opens the Performance Dashboard to view database performance metrics, identify performance bottlenecks and find solutions to performance issues.	Specific menu option
	Restore database	Recover your database into a specific point in time within the retention period	Specific menu option
Warehouse Snapshot	Capture new state	Capture a read-only representation of a warehouse item at a specific point in time, retained to up to 30 days.	Specific menu option

Process Items

This section covers actions unique to Pipelines, Dataflow Gen2 CI/CD, and Notebooks.

Expand table

Item Type	Action	Description	Where to find it
Common Process Item Actions	Save As	Allows you to save a copy of the item to a workspace of your choice.	Specific menu option

Item Type	Action	Description	Where to find it
	Refresh now	Launches an immediate refresh of the item.	**Refresh** button > **Refresh now**
	Schedule refresh	If your process item needs to be refreshed on a regular interval, you can schedule the refresh using the Fabric scheduler.	**Refresh** button > **Schedule refresh**
Pipeline	**Export**	Enables you to back up your pipelines. Exporting ensures that your data integration processes are portable and can be easily restored or replicated.	Specific menu option
	Download	Downloads the pipeline definition file.	**Download** button
Dataflow Gen2 CI/CD	**Check Validation**	Determine the validity of the dataflow by running a "zero row" evaluation for all the queries in the dataflow.	Specific menu option
Notebook	**Download**	Downloads the notebook file.	**Download** button

Considerations and limitations

- Real-time semantic models don't have a details page, as they're being retired.

Last updated on 06/03/2025

Govern and manage your Fabric data estate with the OneLake catalog

The **Govern** section in the OneLake catalog is your centralized destination for governing and managing your organization's Fabric data estate.

Use Govern to assess and improve the governance status of data across your Fabric estate, act on recommended actions, and access governance and administrative experiences.

> **Note**
>
> OneLake catalog and Govern are rolling out by region. If they're not available in your region, use the **Admin portal** to access the existing administration capabilities during the regional rollout.

All data estate view

The **All data in Fabric** view gives you an organization-wide picture of the governance status of your Fabric estate. It's the default view when you have tenant-wide visibility, and it draws on the entire tenant metadata—from items through workspaces to capacities and domains (see Considerations and limitations for exceptions).

The All data estate view surfaces:

- **Governance Insights** at a glance across the tenant.
- **Recommended Actions** to strengthen governance across the estate.
- Governance health signals across domains and subdomains, capacities, workspaces, and item counts, as currently documented.

The first time you open the **Govern** section, it might take a few moments for the insights and actions to appear. You can switch to the **My items** view to see the governance state of the items you own.

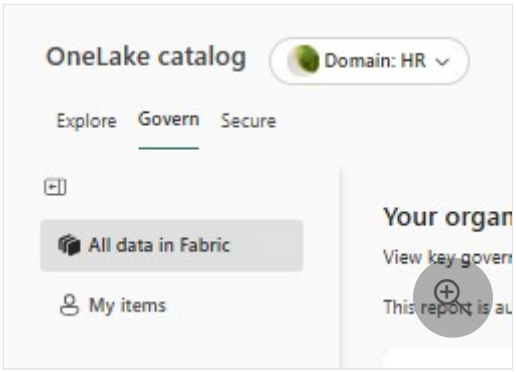

If your organization defines domains, you can use the OneLake catalog's domain selector to choose a specific domain or subdomain. This action scopes the insights and recommended actions to items that reside within the selected domain.

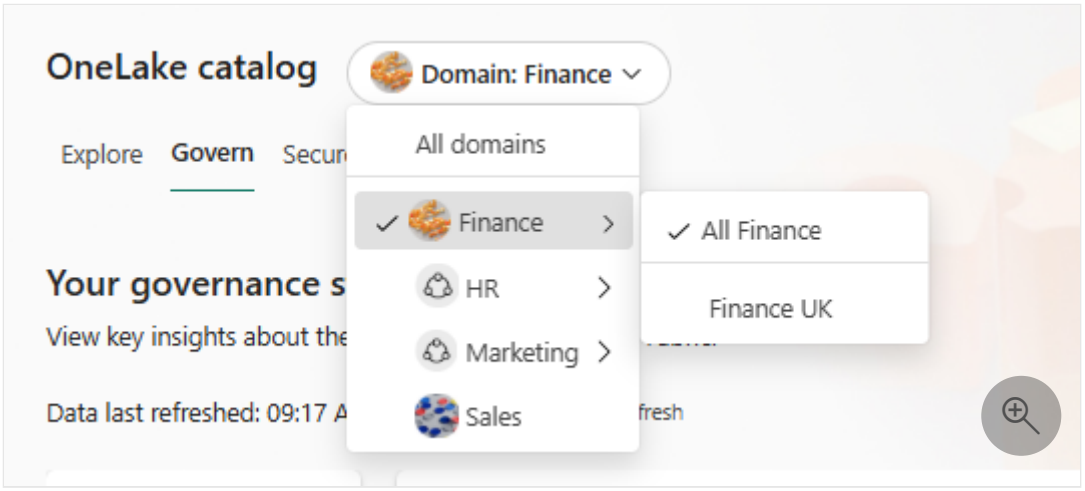

> **Note**
>
> In the **All data estate** view, the domain filter doesn't apply to certain actions and those actions remain unchanged.

Governance insights for the estate

The insights section provides a high-level snapshot of the current governance state of the entire Fabric tenant. Select **View more** to see all available insights.

These insights use data from the last successful refresh of the Admin Monitoring Storage, which is automatically generated in the Admin Monitoring workspace the first time you open the Govern section or access the Admin Monitoring workspace. The data refreshes automatically every day. (See Considerations and limitations for exceptions.)

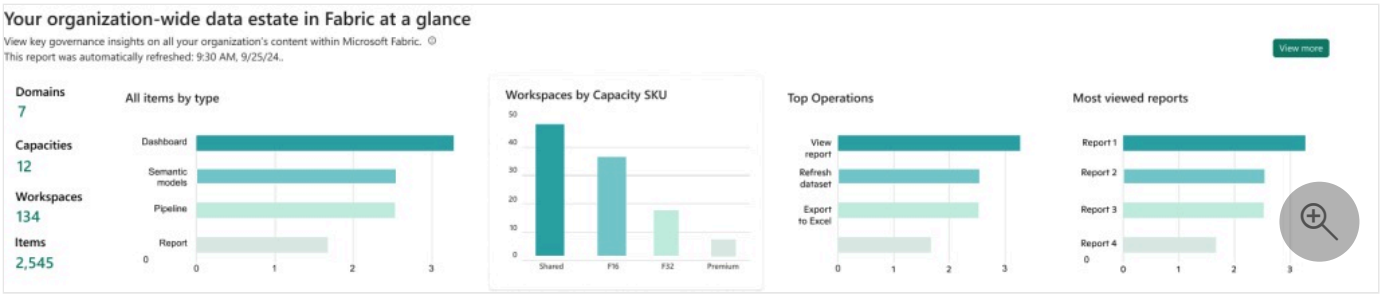

Selecting **View more** gives you expanded governance, administration, and security insights for your Fabric data. The estate-wide report provides expanded insights across three tabs:

- **Manage your data estate** contains inventory overview, capacities and domains information, and details about feature usage across the tenant.

- **Protect, secure & comply** includes information about sensitivity label coverage and data loss prevention policies activated and scanned across the various workspaces in the organization.
  - The **Sensitivity labels** selector shows the most frequently used labels and the percentage of unlabeled items. Drill down by item type and user to identify labeling gaps and policy misalignment. Review a complete inventory of labels, or analyze label distribution by domain or workspace, to understand how labels are applied across the tenant.
  - The **DLP** selector shows the workspaces or data items evaluated by DLP policies, helping you identify policy violations and take action, such as applying a more restrictive label or removing sensitive information. Break down scanned items by type or location, and review the last evaluation time to assess data freshness and trigger a new scan if needed.

- **Discover, trust, and reuse** surfaces insights about data freshness, item curation state, description and endorsement coverage, and content sharing.

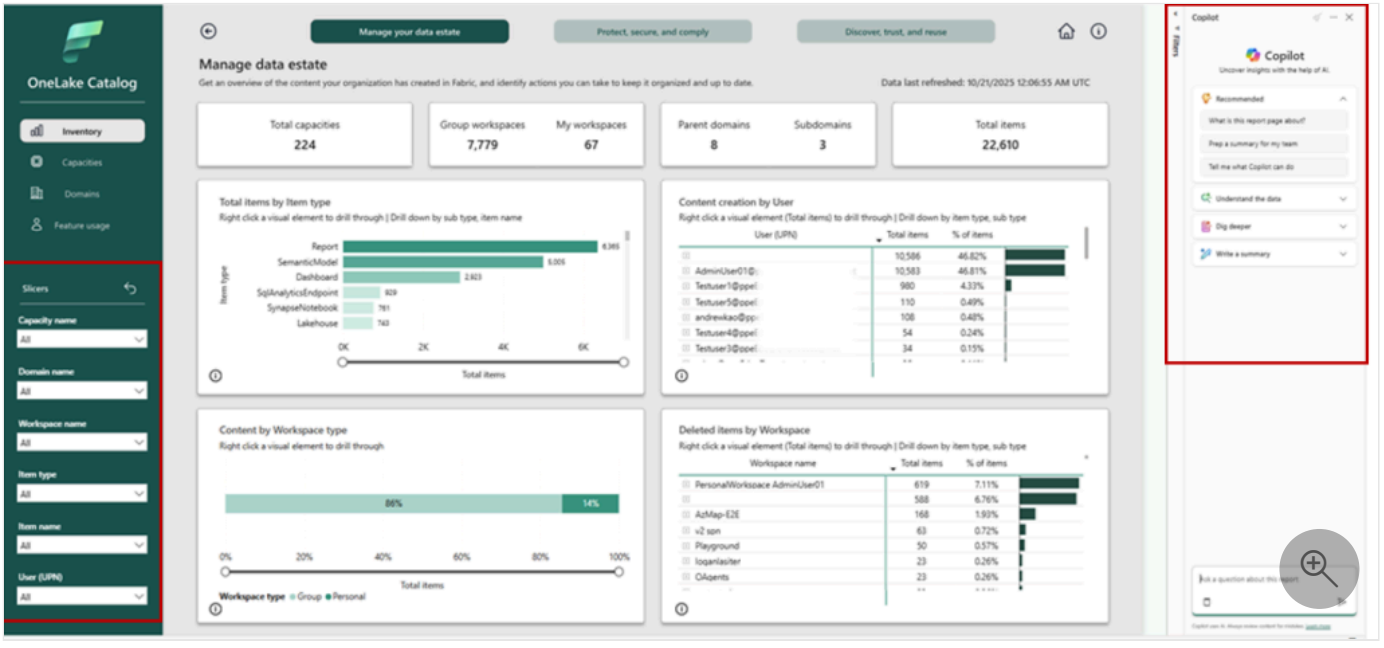

In the report, you can filter, drill through to get more details, and initiate Copilot at any point to gain more insights and support comprehensive, interactive data exploration (see Considerations and limitations for exceptions).

Recommended actions for the estate

The **Recommended Actions** section displays cards suggesting actions you can take to improve the governance posture across the estate. When you select a card, you see an insight highlighting the issue, an explanation of why the issue matters, and a list of steps about how to address the issue.

Example of a recommended action card:

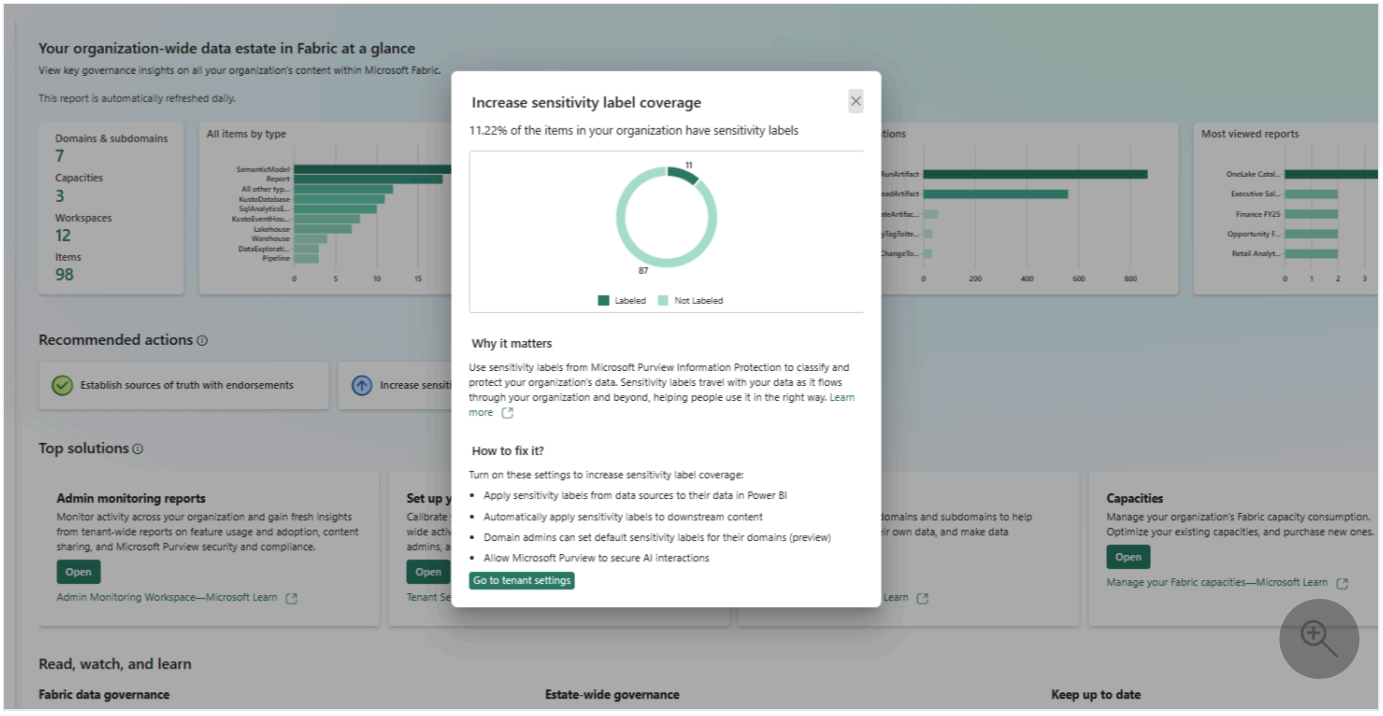

The recommended actions vary depending on what the insights reveal.

My data view

The **My items** view scopes governance to the items you own - the items that appear when you apply the **My items** filter in the Explore section (see Considerations and limitations for exceptions).

The **My data** view surfaces:

- **Governance Insights** scoped to the items you own.
- **Recommended Actions** for a data owner's items.
- Ownership-scoped governance signals, as currently documented.

Governance insights for your items

In the **My items** view, the insights section shows high-level insights about the content you create in Fabric. Select **View more** to see all available insights.

These insights use data from the last successful refresh of your OneLake catalog governance report. The data refreshes automatically every time you open the **Govern** section. Select the **Refresh** button to make sure you have the latest data.

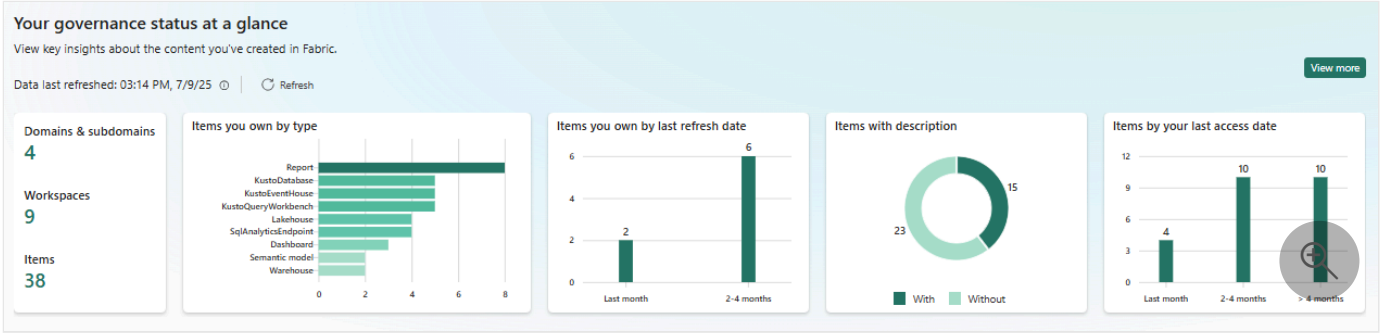

Selecting **View more** opens a simplified report with more insights about your inventory, sensitivity label coverage, and curation state.

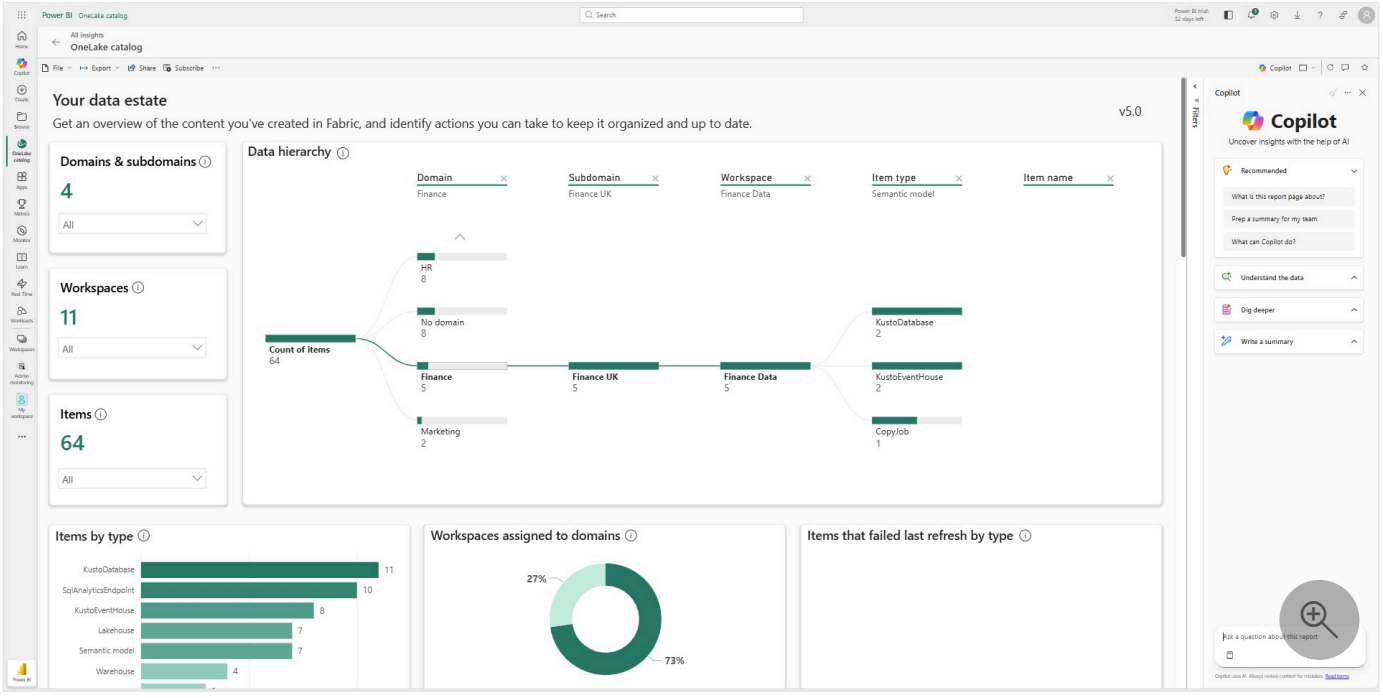

Recommended actions for your items

The **Recommended Actions** section displays cards suggesting actions you can take to improve the governance posture of the items you own. Recommended actions in the **My items** view also include the ability to view all entities associated with a recommended action, including items, workspaces, and more, and open any of them with a single click.

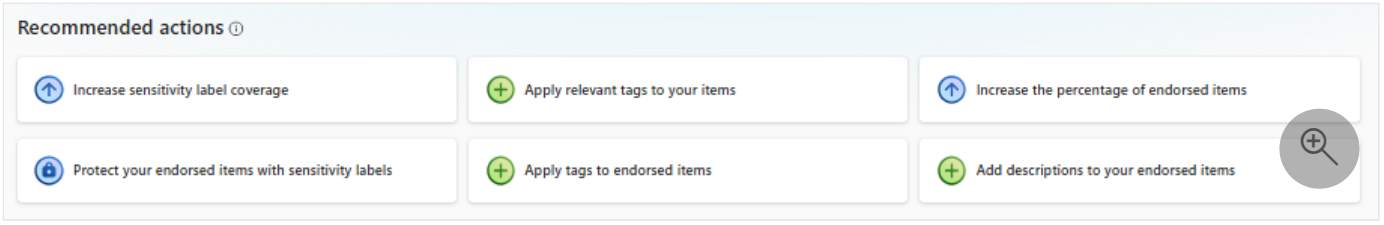

Example of a recommended action page for a data owner:

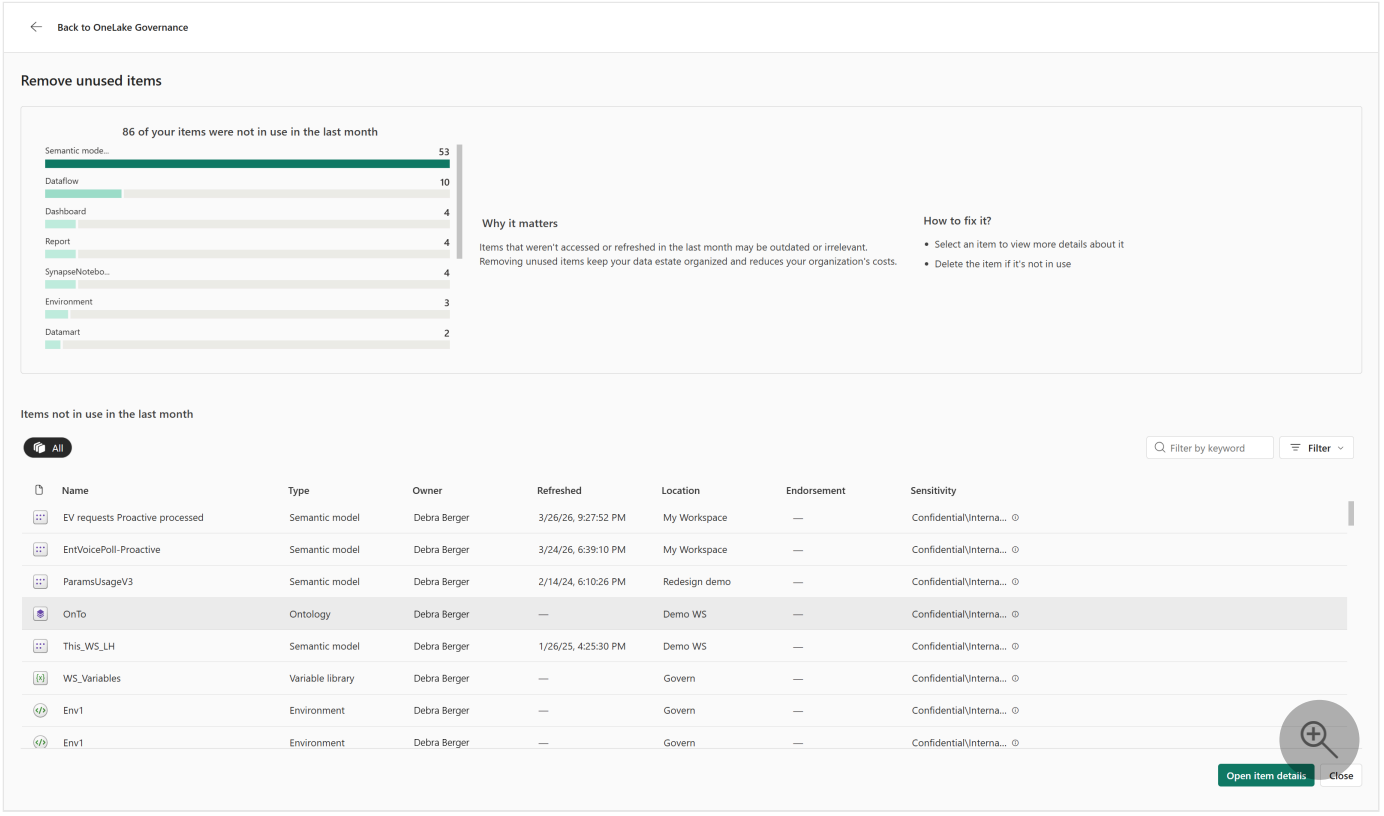

Governance experiences in the Govern section

The Govern section brings Fabric's governance capabilities together in one destination. Each of the following capabilities is publicly documented today:

- Domains — organize your data estate into domains and subdomains.
- Tags — classify and find items with organizational tags.
- Endorsement — signal the trust level of items across the estate.
- Metadata scanning — extract and catalog metadata across the tenant.
- Cross-tenant sharing discovery — understand where data is shared beyond your tenant.
- **Policies** — define and apply governance policies across your Fabric estate.

Administrative experiences in the Govern section

The following administrative experiences are available from Govern:

- Capacities
- Workspaces
- Configurations — access tenant settings and Help and support links.
- Azure connections
- Workloads
- Organizational themes
- Fabric identities
- Branding settings

Who the Govern section is for

A single person often holds multiple responsibilities. The same individual might be a Capacity Admin, a Fabric Admin, a Domain Admin, a Workspace Admin, and a Data Steward. Because of this overlap, the Govern section is organized around what you're trying to do rather than around isolated personas. Whatever combination of roles you hold, you can use the Govern section to:

- **Find** the data and administrative controls relevant to your responsibilities.
- **Understand** the governance status of your estate through Governance Insights.
- **Act** on Recommended Actions to improve governance posture.
- **Secure and manage** your estate through governance capabilities and administrative experiences in one place.

The entries you see in Govern depend on your permissions:

- Users without admin permissions see **Configurations**, with Help and support entries only, when this experience is enabled for their tenant.
- Domain admins see **Configurations**, with Help and support entries, and **Domains**.
- Capacity admins see **Configurations**, with Help and support entries, and **Capacities**.

Open the Govern section

To open the Govern section, select the OneLake catalog icon in the Fabric navigation pane and then select **Govern**.

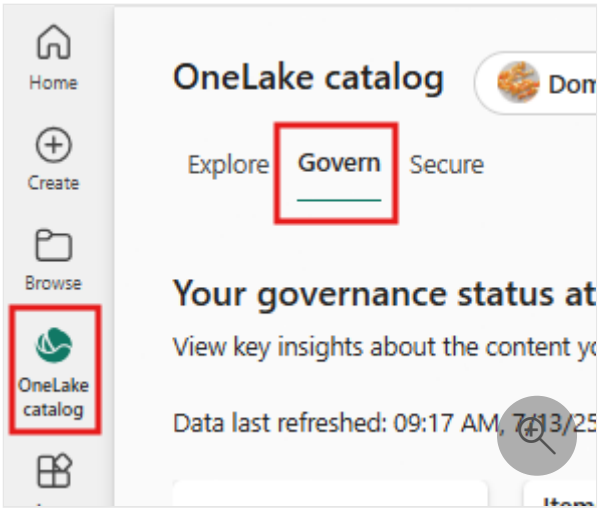

You can also access the Govern section from the settings gear by selecting the **OneLake Catalog | Govern** link.

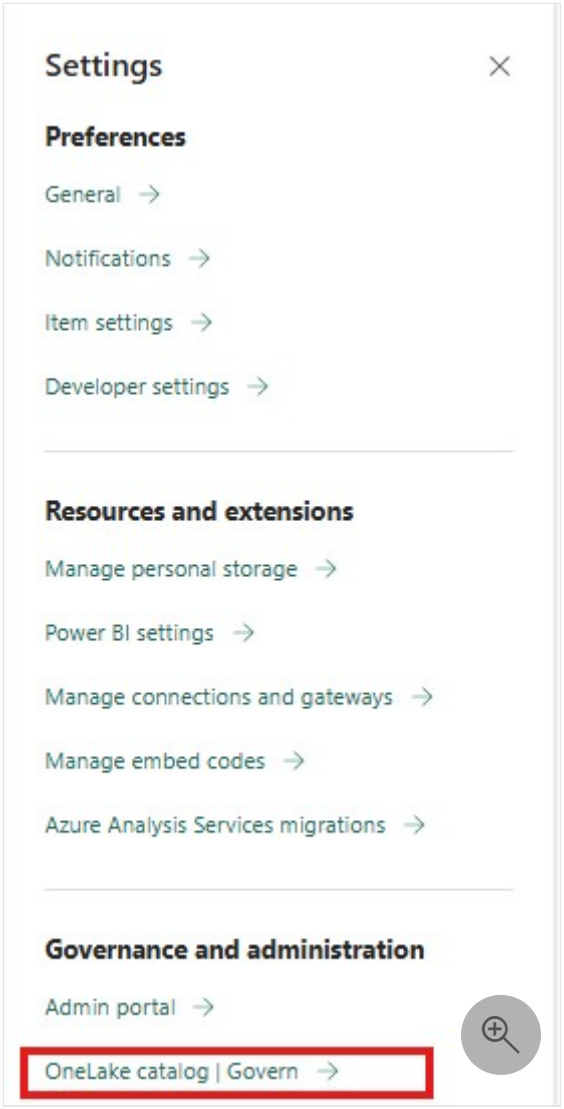

Use Quick access

The **Quick access** section connects you to common administration experiences:

Expand table

Card	Use it to
Admin monitoring workspace	Monitor activity across your organization and gain insights from tenant-wide reports.
Manage users	Manage Fabric users, groups, and administrators in the Microsoft 365 admin center.
Audit logs	View tenant activity and export audit logs in the Microsoft Purview portal.

The Govern section also provides links to tools and resources relevant to your scope of responsibility:

- **Top solutions**: Lists relevant Microsoft Fabric solutions for data governance, compliance, and security, along with links to documentation.
- **Read, watch, and learn**: Provides links to other relevant documentation and resources.

Create custom reports

To customize a report, create a copy of the report and modify the copy, or create a new report based on the autogenerated semantic model.

> **Important**
>
> Don't modify the original autogenerated report or semantic model. Govern requires these artifacts to provide insights, and changes can cause the experience to stop working.

Troubleshooting

- For the *My items* view, if the data isn't refreshing as expected, check the notifications pane and the Monitor page to see if you can identify the cause. If refresh fails repeatedly, or if you can't figure out what's causing it to fail, try regenerating the report. To do so, close the Govern section, and delete the OneLake catalog governance report and its associated semantic model from your *My workspaces*. Then reopen the Govern section.

Considerations and limitations

The following list describes considerations and limitations when using the Govern section:

- **Subitems** - Subitems such as tables aren't supported and don't appear in insights.

- **Cross tenant** - The Govern section doesn't support cross-tenant scenarios or guest users.
- **Private Link** - The Govern section isn't available when Private Link is activated.
- ***View more* reports** -
  - Copilot functionality depends on organizational setup and the capacity the workspace is assigned with. The workspace (*admin monitoring workspace* for the estate view and *My workspaces* for the data owner view) should be allocated to the appropriate capacity in order to activate the Copilot button.
  - Items of third-party workloads aren't included in the charts.
- **Data refresh** - Because estate insights, recommended actions, and *view more* reports are based on admin monitoring storage that refreshes once a day, there could be gaps between the data reflected and the actual state. It takes a day to get an updated view of all the changes made in the organization.
- ***Admin monitoring* workspace** -
  - If you reassign the *Admin monitoring* workspace or *My workspace* to another capacity, you might get an error when accessing the Govern section of the OneLake catalog. In such cases, ensure the newly assigned capacity has enough resources to run the semantic models and reports.
  - All users viewing content in the admin monitoring workspace, including the estate report, must have a Power BI Pro license, unless the workspace is assigned to a capacity.
  - The estate semantic model is read-only and can't be used with Fabric data agents.

Related content

Discover and trust

- OneLake catalog overview
- Discover and explore Fabric items in the OneLake catalog

Govern and manage

- Governance documentation
- Governance and compliance overview
- Fabric administration documentation
- Domains
- Administer Microsoft Fabric

Secure and comply

- Sensitivity labels

- Data protection
- Use Microsoft Purview with Microsoft Fabric

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Manage Microsoft Fabric settings in OneLake catalog

The OneLake catalog in Fabric provides a centralized location for managing and configuring Fabric settings. Administrators can use this interface to efficiently oversee and adjust settings across capacities, workspaces, domains, tags, and other Fabric configurations, including tenant settings.

Manage capacities

In the OneLake catalog, administrators can manage capacities, which are the resources allocated for various workloads in Fabric. This task includes creating new capacities, configuring capacity settings, and monitoring their usage to ensure optimal performance and resource allocation.

To manage capacities:

1. Sign in to Fabric by using your administrator account.
2. Select **OneLake catalog**, and then go to **Govern** > **Capacities**.

    - To view details and usage data for an existing capacity, select a capacity from the list.
    - To take actions on a capacity (such as pause, resize, reassign, configure settings, and more) select the options menu (three dots) next to the capacity name.
    - To create a new capacity, select **+ New capacity**.

For more guidance on managing capacities, refer to the Manage Fabric capacities article.

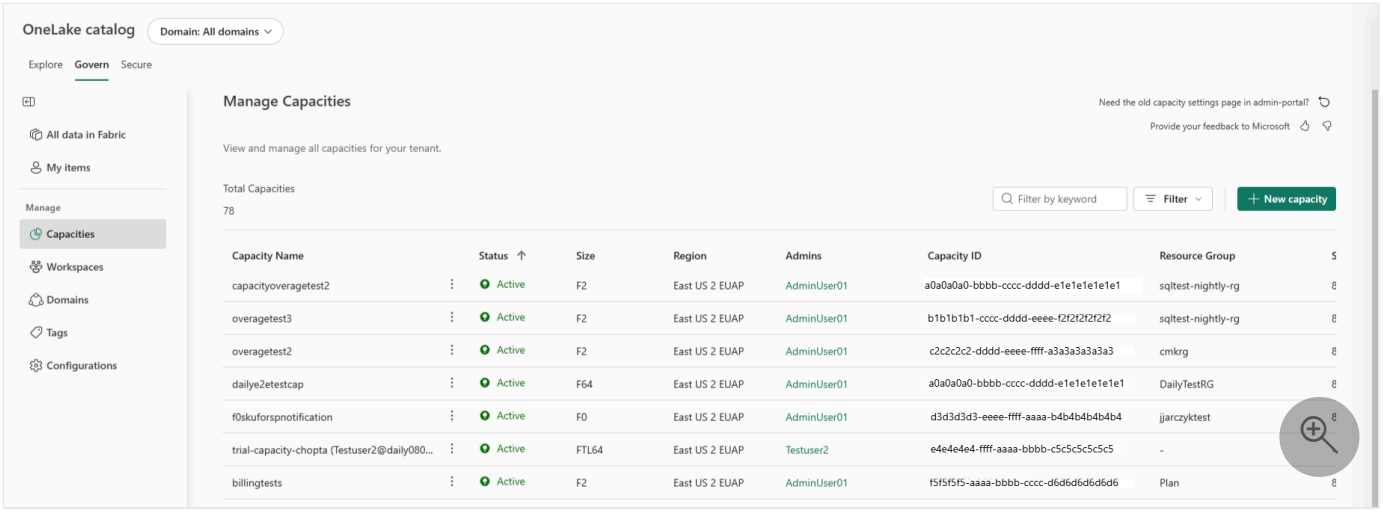

Manage workspaces

Administrators can manage workspaces in the OneLake catalog. Workspaces serve as collaborative environments for organizing and working with data, reports, and other resources. This management includes creating new workspaces, configuring workspace settings, managing access permissions, and monitoring workspace activity to maintain efficient collaboration and security.

To manage workspaces:

1. Sign in to Fabric by using your administrator account.
2. Select **OneLake catalog**, and then go to **Govern** > **Workspaces**.

    - To configure settings for a workspace, select the **More options** menu (three dots) next to a workspace name, then select **Edit**.
    - To manage access to a workspace, select the three dots next to the capacity name to take actions on a capacity such as pause, resize, reassign, and configure settings.
    - To create a new capacity, select **+ New capacity**.

For more guidance on managing workspaces, refer to the Manage Fabric workspaces article.

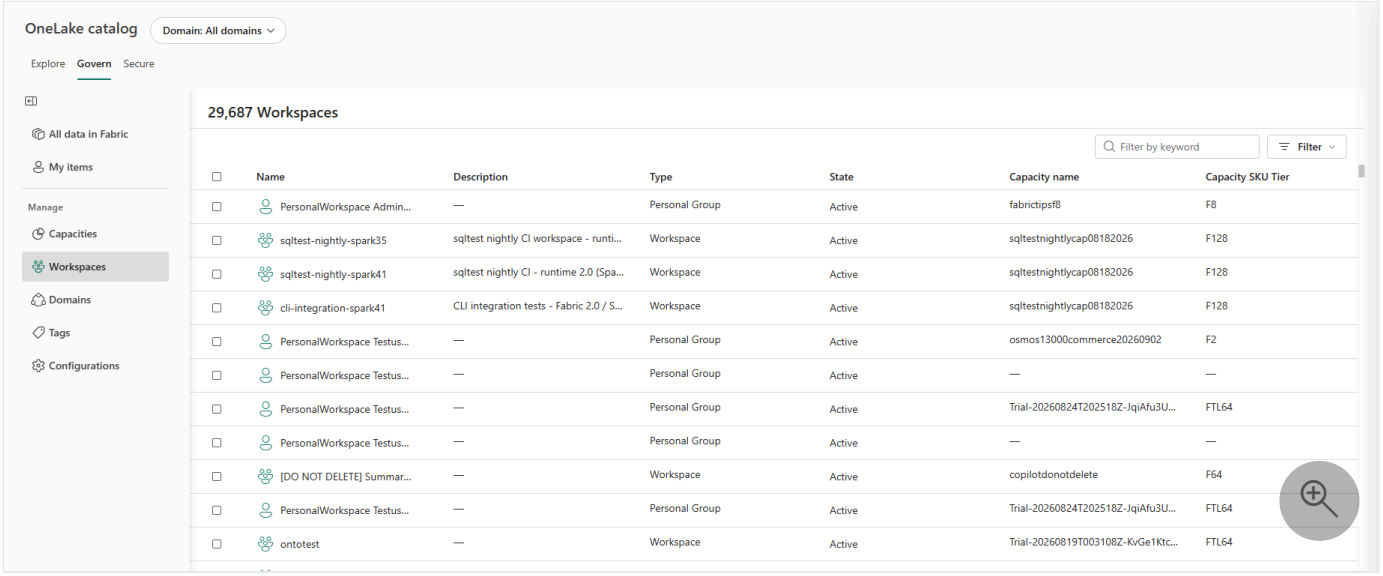

Manage domains

The OneLake catalog allows administrators to manage domains, which organize and control access to data within Fabric. Administrators can create and configure domains, assign ownership, and set access policies to ensure proper data governance and security.

To manage domains:

1. Sign in to Fabric by using your administrator account.
2. Select **OneLake catalog**, and then go to **Govern** > **Domains**.

- To create a new domain, select + **New domain**.
- To view details and manage an existing domain, select a domain from the list.

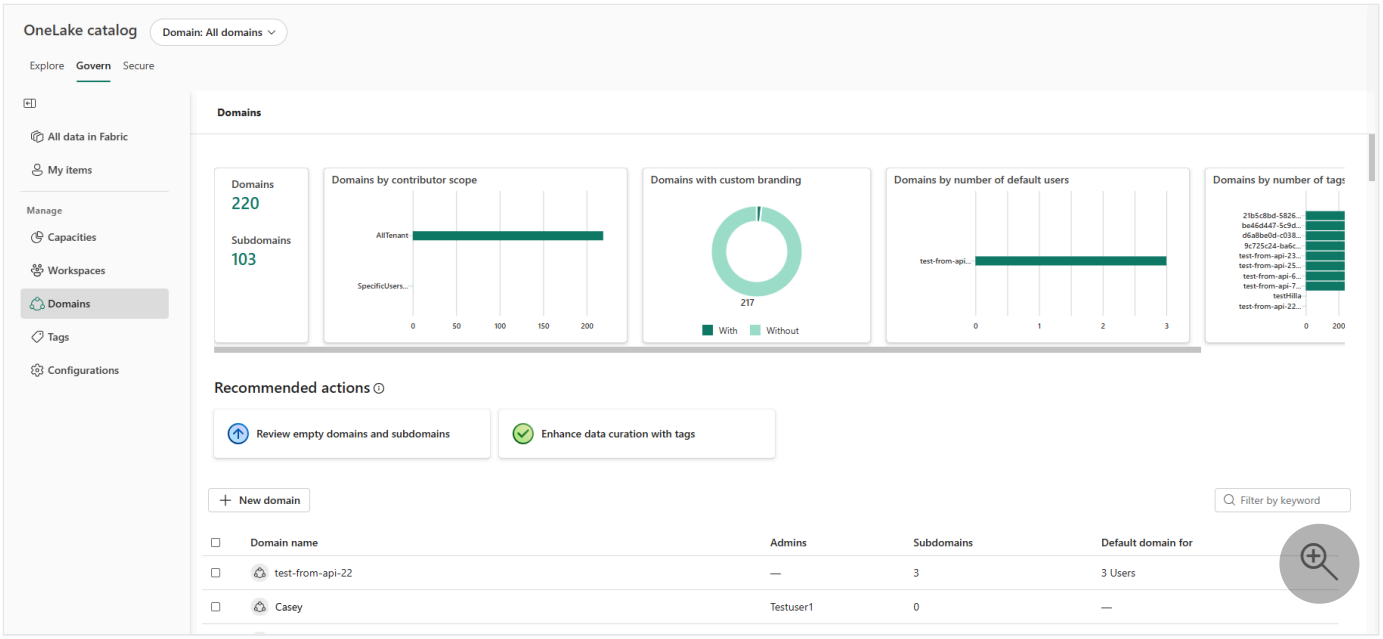

Manage tags

Administrators can use tags in the OneLake catalog to categorize and label various resources for easier management and discovery. This task includes creating new tags, applying them to resources, and managing existing tags to maintain a well-organized and searchable catalog.

To manage tags:

1. Sign in to Fabric by using your administrator account.
2. Select **OneLake catalog**, and then go to **Govern** > **Tags**.

   - To create a new tag, select + **New tag**.
   - To view details and manage an existing tag, select a tag from the list.

For more guidance on managing tags, see Manage Fabric tags.

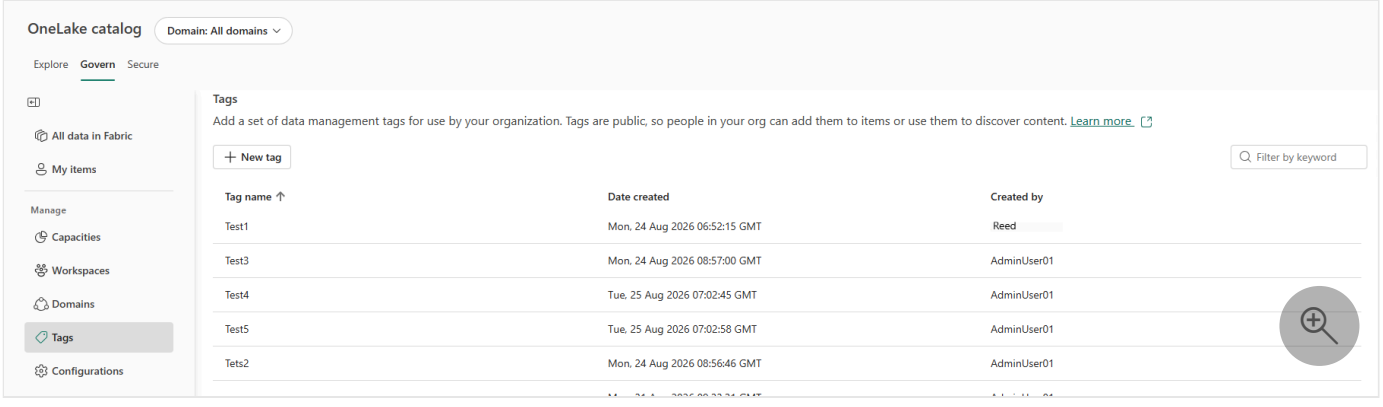

Manage other Fabric settings

In addition to capacities, workspaces, domains, and tags, the OneLake catalog provides options for configuring other Fabric settings, including tenant settings. Administrators can configure these settings to tailor the Fabric environment to their organization's needs, ensuring optimal performance, security, and compliance.

To manage other Fabric settings:

1. Sign in to Fabric by using your administrator account.
2. Select **OneLake catalog**, and then go to **Govern** > **Configurations**.

    - To modify tenant settings, select **Tenant settings** and make the necessary changes.
    - To configure other Fabric options, select the relevant configuration and update the settings as needed.

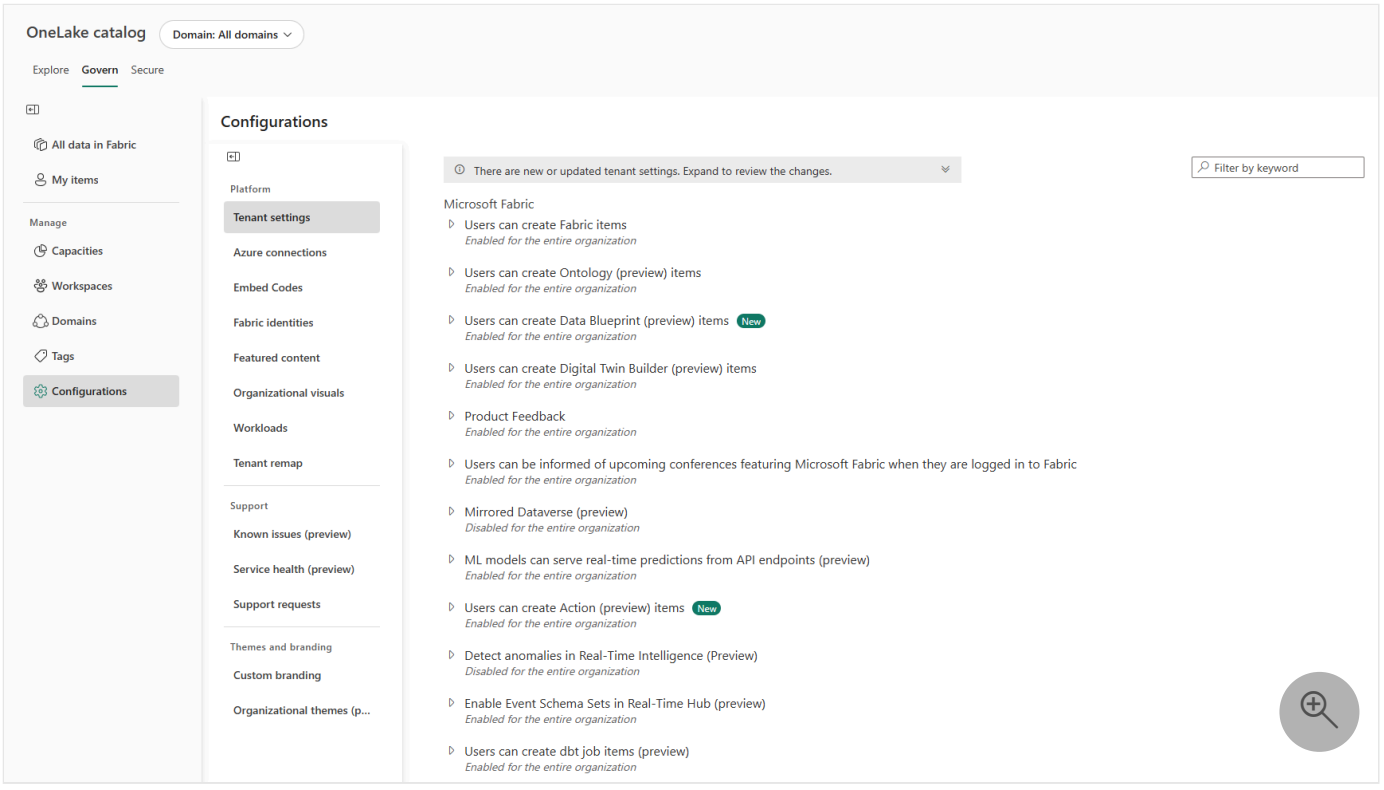

Related content

- Azure connections
- Embed codes
- Fabric identities
- Featured content
- Organizational visuals
- Workloads

- Help and support settings
- Custom branding

Last updated on 09/29/2026

Secure your Fabric data

Use the OneLake catalog's **Secure** tab to view, monitor, and configure security roles across workspaces and items in Microsoft Fabric. It centralizes in one place:

- A view of workspace roles and permissions, for auditing access to workspaces and data.
- A view of OneLake security roles across workspaces and item types. Admins can create, edit, or delete OneLake security roles from a single location.

The **Secure** tab shows you the items that are relevant to you, including which workspaces you have access to. Admin and Member roles are required in a workspace to see data for that workspace. Contributors or Viewers in a workspace only see information about their own access.

View users

On the **View users** page, you can see all of the users with access to selected workspaces, including a summary of their assigned roles.

1. To get started, choose the workspaces from the **Workspaces** control on the left.
2. Each row represents a unique user, group, or application and shows a count of the roles they have in the selected workspaces.
3. Use the **All users by** filter at the top to limit your selection to specific types of user, workspace roles, or workspaces.
4. Use the **search** control to find a specific user or group and verify their permissions.

Manage workspace permissions

Using the ribbon on the **View uers** page, you can manage workspace role assignments across multiple workspaces from a single experience.

- **Add users**: Select this option to add one or more users to the same workspace role across one or more workspaces. For example, you can onboard all new hires to the Viewer role across all relevant workspaces.
- **Manage access**: Select this to choose between editing roles and removing roles. Like with **add users**, you can perform these actions in bulk across multiple users and workspaces. For example, removing all summer interns from their relevant workspaces.

To manage workspace permissions, select the **Add users** or **Manage access** option you want.

- Under **Workspaces**, select one or more workspaces where you want the user to be updated. Note that you must have Admin or Member permissions on a workspace for the update to apply.
- For **Members**, search for the users or groups you want to update.
- Select the **Role** you are granting or updating them to. This option is not present for removing access.
- Select **Update**.
- To cancel the changes, select **Discard**.

View security roles

On the **View security roles** page, you can see all of the OneLake security roles across the selected workspaces. Each item that has a OneLake security role in those workspaces appears in the list. You can view existing roles, edit roles, delete roles, or create new ones from this page.

View and edit an existing role

1. Open **View security roles** in the OneLake catalog.

2. Choose the workspaces from the **Workspaces** control on the left.

3. Each row represents a single OneLake security role that's inside a Fabric item. You can view the following details for a role:

    - **Item** name
    - **Role name**
    - **Role type** - only OneLake security roles are supported
    - **Permission** - the permission granted by the role
    - **Location** - the workspace where the item lives
    - **Data owner** - the owner of the item

4. Select a role to open the OneLake security role management experience.

5. You can view and edit the **Data in role** and **Members in role** in the role details page. For more information about viewing and editing a OneLake security role, see Create and manage OneLake security roles.

Delete a role

1. Open **View security roles** in the OneLake catalog.

2. Choose the workspaces from the **Workspaces** control on the left.
3. Select a role to open the OneLake security role management experience.
4. In the control bar at the top, select **Delete**.
5. Select **Delete** to confirm the deletion, or **Cancel** to return to the previous screen.

Create a new role

1. Open **View security roles** in the OneLake catalog.
2. Choose the workspaces from the **Workspaces** control on the left.
3. Select a role to open the OneLake security role management experience.
4. Use the control bar at the top to create a new role with **New role** or **Duplicate** an existing role. The role is created in the item where the currently selected role lives. To change to a different item, select a OneLake security role for that item.

Related content

- Fabric and OneLake security overview
- OneLake data access roles (preview)
- OneLake catalog overview
- Govern your data

Last updated on 06/05/2026

What are Microsoft Fabric policies (preview)?

Microsoft Fabric policies enable administrators to centrally define and enforce granular governance rules across Fabric. By using Microsoft Fabric policies, you can create rules that control specific operations based on conditions such as workspaces, security groups, sensitivity labels, and more.

This article introduces the key concepts of Microsoft Fabric policies, explains the entry points for creating and activating policies, and walks through the end-to-end policy lifecycle.

Why use Fabric policies?

Organizations using Fabric often face challenges with:

- **Feature enablement.** Organizations need granular control to enable features that align with their governance requirements.
- **Administrative overhead.** Managing resources and ensuring compliant workflows through manual processes or separate systems is time-consuming.
- **Limited visibility.** Administrators struggle to track existing settings, understand default configurations, and manage governance scope across the organization.

Fabric policies address these challenges and provide the following benefits:

- **Granular, condition-based governance.** Define rules that apply only under specific conditions, such as a particular workspace, security group, or sensitivity label.
- **Centralized policy management.** Manage governance from a single place, the OneLake catalog Policies page.
- **Visibility and auditability.** Track configuration and enforcement decisions through Fabric audit logs.
- **Programmatic management.** Automate policy management with REST APIs, Git integration, and service principals.

Together, these capabilities deliver consistent governance across the tenant, increase administrative visibility, and reduce governance friction.

Key concepts in Fabric policies

Fabric policies use a small set of concepts that work together to define and enforce governance. The following diagram shows how those concepts combine to evaluate a user action and allow or block it. The steps after the diagram walk through the same flow.

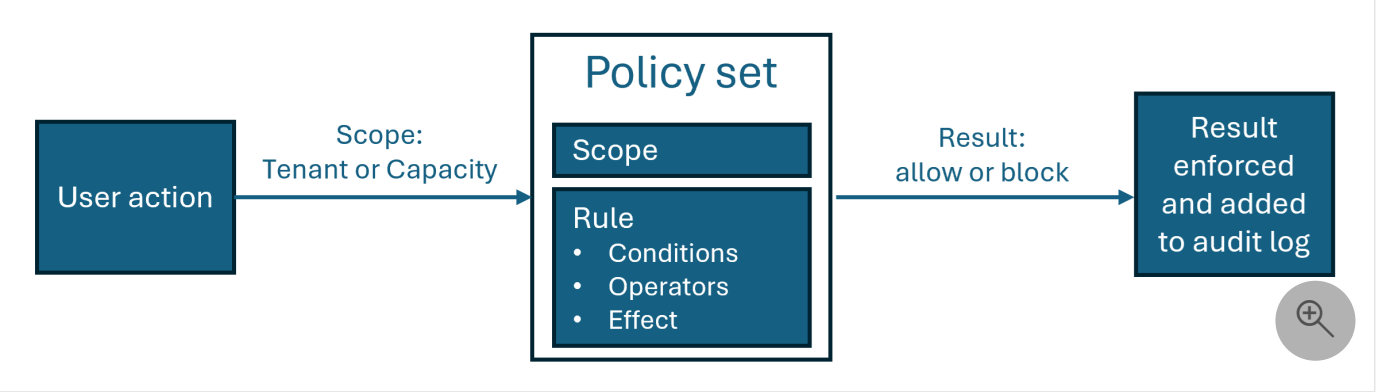

Fabric evaluates a user action in the following order:

1. **Action.** A user attempts a governed action.
2. **Admin scope.** The action is associated with an admin scope: Tenant or Capacity.
3. **Active policy.** Fabric loads the active policy for that scope.
4. **Rule check.** Fabric checks whether the policy has any rules. If it doesn't, Fabric applies the policy default and stops.
5. **Condition match.** If rules exist, Fabric evaluates each rule's conditions with its operators. If no rule matches, Fabric blocks the action. If a rule matches, Fabric applies the Allow effect and permits the action.
6. **Result.** Fabric enforces the result, whether default, block, or allow, and records it in the audit logs.

Fabric Policies tenant setting

The Policies (preview) tenant setting controls the availability of Fabric policies and allows Fabric tenant admins and capacity admins to create policy sets and apply policy rules. You can use the security group selector to limit this access to specific security groups.

Fabric policy items

Fabric stores each policy as an item in a workspace. This item model gives policies standard Fabric item capabilities, including public REST APIs, Git integration, CI/CD, the standard Fabric permission model, and discovery and management in the OneLake catalog.

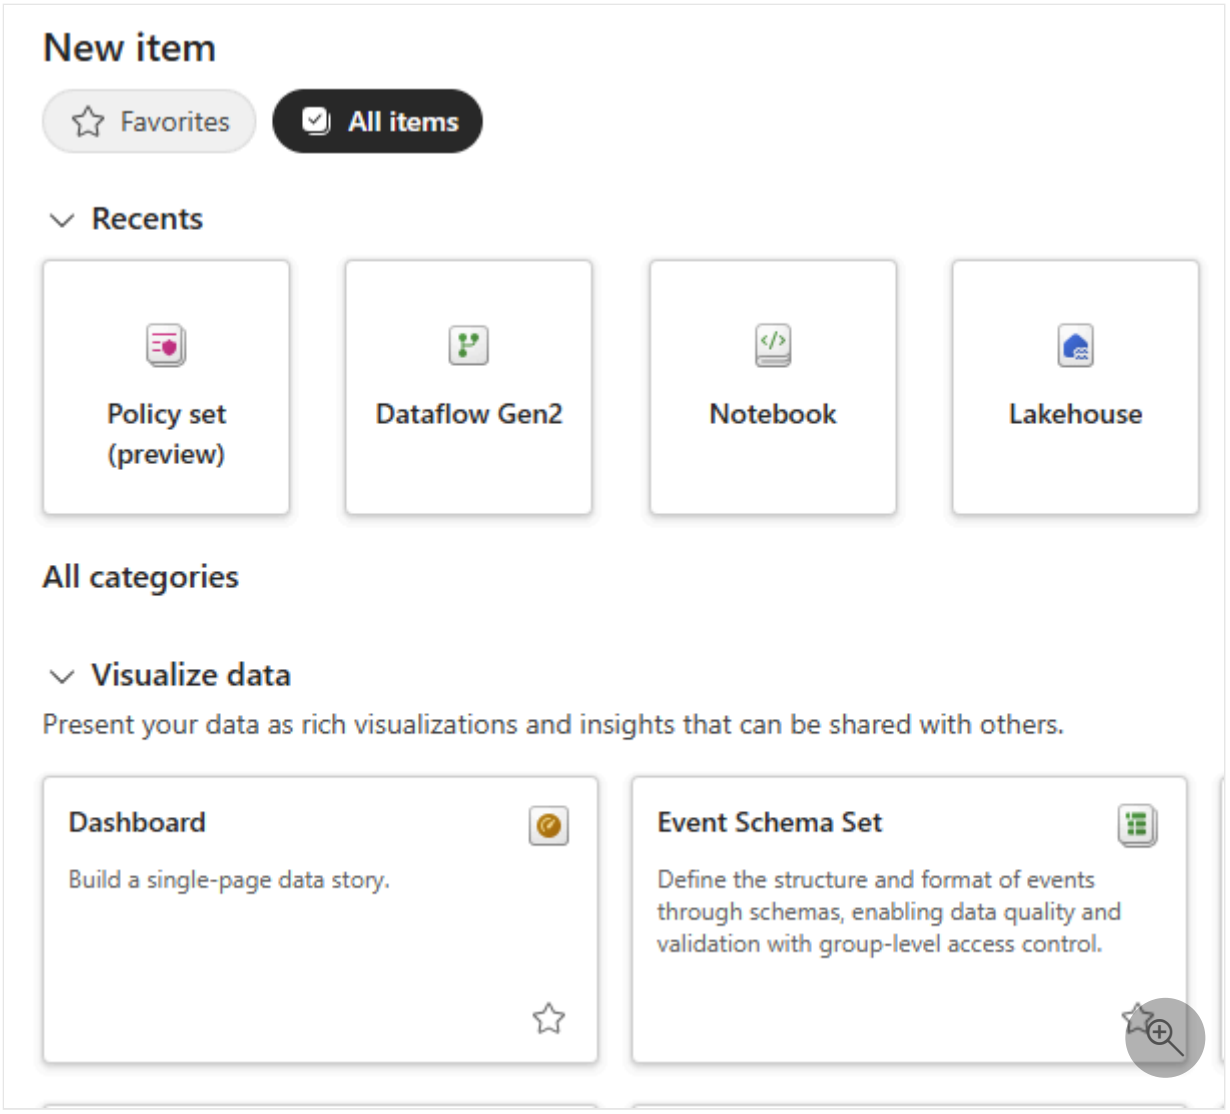

Fabric policies have the following characteristics:

- You assign each policy to an admin scope (Tenant or Capacity) when you create it. You can't change the scope.
- You must *activate* the policy before its rules take effect.
- You can activate only one policy per admin scope at a time.
- You can use Git and CI/CD automation to manage and deploy policies. However, you can't activate a policy through Git.

Anyone with permission can create and manage the policy item or edit its rules, but those rules take effect only after activation. To activate a policy, you need the appropriate admin permissions: a Fabric administrator for Tenant scope or a capacity administrator for Capacity scope.

Fabric policy rules

A policy rule is the core unit of governance. Each rule consists of:

- **Metadata**, including a system-generated rule ID, a display name, a description, and creation and modification dates.
- **Conditions** that determine when the rule applies. Fabric combines multiple conditions with a logical AND.
- **An effect** that determines what happens when the conditions are true.

The number of rules a policy can contain depends on the admin scope: up to 100 rules per policy at Tenant scope, and up to 50 rules per policy at Capacity scope.

Fabric policy conditions and operators

Conditions determine when a rule applies. Each condition pairs a policy attribute, such as a workspace, a security group, or a recipient email domain, with an operator that defines how Fabric tests the action's value for that attribute against the values you configure. The conditions available for a policy depend on what the policy evaluates:

- For an action performed within an item (for example, sharing data externally), you can base conditions on the item, the user, the workspace, and policy-specific properties.
- For an action at the workspace level (for example, changing workspace settings), you can base conditions on the workspace, security groups, and policy-specific properties.

The available condition types are:

Expand table

Condition type	Description	Applies to	Notes
Workspaces	Fabric workspaces where the policy applies.	All policies	Maximum 50 per rule.
Security groups	Microsoft Entra ID security groups the policy applies to.	All policies	Maximum 50 per rule.
Sensitivity labels	Microsoft Purview sensitivity labels that must be present on items.	Selected policies	Includes a **No label** option.
Recipient email domain	Email domains allowed for external data sharing.	External data sharing policy	—
Item type	Fabric item types the policy applies to.	Item creation policy	—
Workspace settings	Specific settings within a workspace that the policy applies to.	Workspace settings policy	—

Condition evaluation logic:

- Multiple conditions within a rule are combined with **AND**.
- Multiple values within a condition are combined with **OR** (for `AnyOf`) or **NOR** (for `NoneOf`).
- Unspecified conditions evaluate to **TRUE** (match all).

The operator available for a condition depends on its format:

- **ID-based conditions** (security groups, workspace IDs, and sensitivity labels) use `AnyOf` (at least one configured value matches) and `NoneOf` (no configured value matches).

Example: Allow selected users to create specific item types

This capacity-scoped example allows approved data team to create DataFlow Gen2 items.

Expand table

Setting	Value
Policy	Allow item creation
Rule name	Only data team can create Dataflow Gen2
Security groups	Any of: Data-team
Item type	Any of: Dataflow Gen2

Fabric policy effects

The effect defines what happens when a policy's conditions evaluate to true. Currently, the only available effect is **Allow**.

Expand table

Effect	Description
Allow	Permits the operation when conditions are met. If no rules match, the operation is denied.

When you configure a policy with rules and conditions, it behaves as an allow list. The action is only permitted if it meets all conditions specified in the policy. Fabric blocks any action that doesn't meet the conditions.

Fabric policy admin scopes

An admin scope is the administrative boundary where Fabric enforces a policy. The scope determines which policies are available, which workspaces a rule can target, and who can activate

the policy.

Expand table

Admin scope	Description	Activated by
Tenant	Policies apply across the entire organization	Fabric administrator
Capacity	Policies apply to the workspaces assigned to a specific capacity	Capacity administrator

Policies center in the OneLake catalog

The Policies center on the OneLake catalog **Govern** tab is the primary place to manage active policies for each admin scope. A policy appears in the center after its activation.

In the Policies center, you can:

- View policies by admin scope.
- See policy status (**Set by admin** or **Default**) and action type (**Allow all**, **Blocked**, or **Advanced**).
- Create, edit, and delete rules.
- Monitor policy activity through audit logs.

Access depends on your role: a tenant administrator sees all admin scopes, while a capacity administrator sees only their assigned capacities.

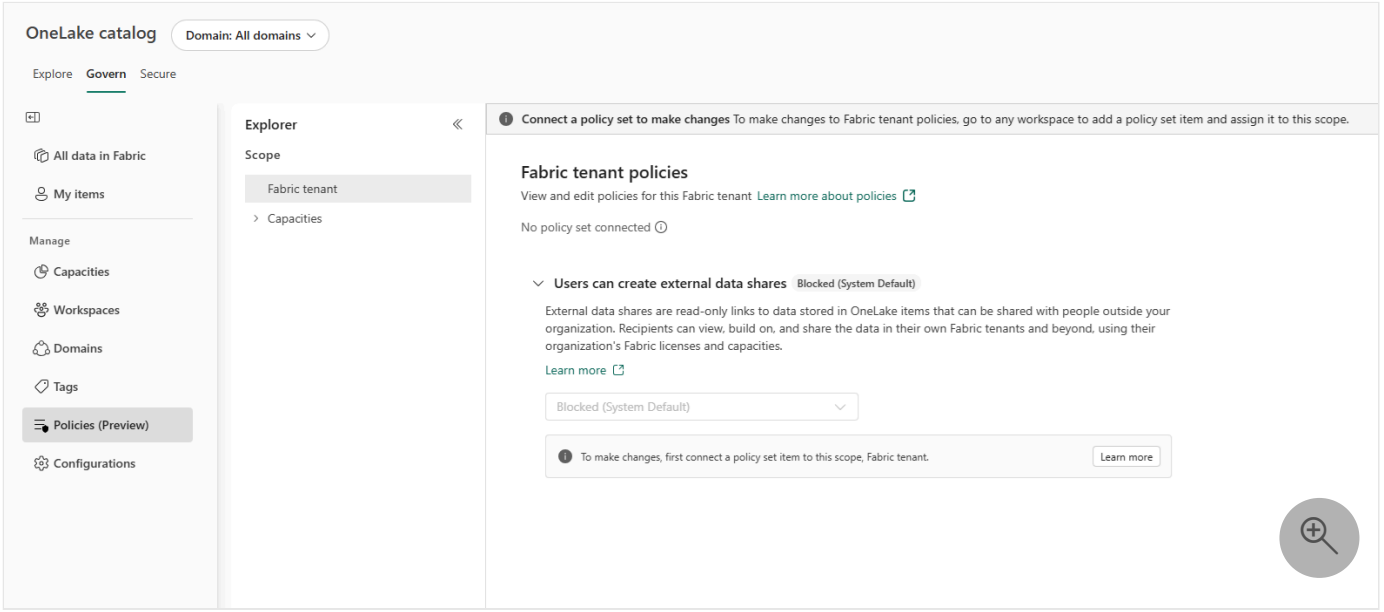

How Fabric policies work

A Fabric policy enforces governance when a user attempts a governed action. The Fabric feature calls the policy engine, which evaluates the active policy rules for the relevant scope and policy and returns an **Allow** or **Deny** result. The feature enforces the result and provides feedback to the user.

Fabric policies move from definition to enforcement in four stages:

1. **Policy definition.** An administrator creates a policy and defines rules by using conditions and effects. Fabric stores the policy in a policy item.
2. **Policy activation.** The administrator activates the policy for an admin scope. Only one policy can be active per scope.
3. **Policy evaluation.** When a user attempts an operation, the owning feature calls the policy engine, which evaluates the operation's attributes against the active rules.
4. **Policy enforcement.** Fabric returns an **Allow** or **Deny** result, the owning feature enforces it, and the user receives feedback about the result.

During evaluation, Fabric checks whether any rules exist for the scope and policy:

- If no rules exist, Fabric returns the policy's default behavior.
- If rules exist, Fabric evaluates each rule's conditions. When the conditions match, Fabric applies the effect.
- If rules exist but no conditions match, Fabric denies the operation for Allow-effect policies.

Fabric audit logs record all evaluation results.

Create and activate policies

You can create a policy in a workspace or programmatically with the REST API. A policy is stored as a policy item in a workspace and associated with an admin scope.

Activation is an explicit step that turns policy intent into enforcement. Until you activate a policy, it has no effect. Only one policy can be active per scope.

For step-by-step instructions, see Get started with Fabric policies.

Create a policy as a workspace item

You can create a policy directly in any workspace. Use a secured, dedicated administrative workspace. Fabric stores each policy in a policy item. A single workspace can contain multiple

policies associated with different capacity scopes.

When creating a policy, choose the admin scope, review the system default for each policy, and add rules. You must activate the policy separately before it takes effect.

Create a policy with the REST API

You can create, configure, and activate policies programmatically for automation and CI/CD scenarios. Git doesn't support activation. For more information, see Automate Fabric policies with the REST APIs.

Example scenarios

The following example scenarios illustrate how Fabric policies can be used to control item creation and external data sharing.

Limit item creation in production workspaces

In this example, a capacity administrator wants to limit the creation of Dataflow Gen2 items in production workspaces to only a specific security group. The administrator creates a policy at the capacity scope with the Allow item creation policy and two rules.

Rule 1 - Allow Dataflow Gen2 only for the approved group

Expand table

Condition	Operator	Values
Security groups	AnyOf	DataflowGen2-Creators
Workspaces	AnyOf	Prod-Analytics, Prod-Reporting
Item type	AnyOf	DataflowGen2

Rule 2 - Allow all other item types for all users

Expand table

Condition	Operator	Values
Workspaces	NoneOf	Prod-Analytics, Prod-Reporting

Result: Because the policy contains rules with conditions, it behaves as an allow list. Rule 1 ensures that only members of DataflowGen2-Creators can create Dataflow Gen2 items in the specified workspaces. On its own, Rule 1 would block all other item creation across the capacity, so Rule 2 ensures that any user can create any item type in other workspaces.

JSON example: Allow policy rule for external data sharing

The following JSON shows the structure of an Allow policy rule for external data sharing:

JSON

```
{
  "policyRuleId": "789e1234-e89b-12d3-a456-426614174999",
  "policy": "ExternalDataShare",
  "displayName": "External Data Sharing Policy for finance and marketing",
  "description": "Controls which users and workspaces are allowed to share data
externally.",
  "conditions": [
    {
      "type": "Principal.UserAndSecurityGroup",
      "predicate": { "operator": "AnyOf", "values": ["DataStewards", "BIAdmins"] }
    },
    {
      "type": "Workspace.Id",
      "predicate": { "operator": "AnyOf", "values": ["MarketingWS", "FinanceWS"] }
    },
    {
      "type": "EDS.RecipientEmailDomain",
      "predicate": { "operator": "AnyOf", "values": ["@contoso.com", "@fabrikam.com"]
}
    }
  ],
  "then": { "effect": "Allow" }
}
```

How this rule evaluates

The rule allows external data sharing when **all** of the following conditions are true:

- The user belongs to the DataStewards or BIAdmins security group.
- The action is performed in the MarketingWS or FinanceWS workspace.
- The recipient email domain is @contoso.com or @fabrikam.com.

Best practices for creating a policy

These recommendations aren't required, but they simplify policy administration and testing.

- Use a secured, admin-managed workspace to store policies.
- Store all policies for the admin scopes managed by the same administrator or admin team in the same workspace. This approach provides one location for tenant- and capacity-scoped policies and simplifies access management, discovery, and lifecycle management.
- Use clear policy set and rule names that identify the admin scope and purpose.
- Use test users who are members and nonmembers of the configured security groups to validate both allow and deny outcomes before broader rollout.

Supported Fabric policies

The following policies are available:

Expand table

Policy	Admin scope	Effect	Default behavior	What it controls
Allow external data sharing	Tenant	Allow	Block all	Controls external data sharing.
Allow item creation	Capacity	Allow	Allow all item types	Controls which item types users can create in workspaces assigned to the scope.
Allow workspace settings editing	Tenant	Allow	Allow editing of all workspace settings	Controls which workspace admins can edit certain workspace settings.

Licensing and billing

Fabric Policies is included with your Fabric license at no extra charge. Policies don't consume capacity units.

General limitations

- Fabric policies aren't currently supported in the following capacity regions: West Europe, North Europe, and West US.
- You can activate only one policy at a time for an admin scope.
- Tenant-scoped policies support up to 100 rules per policy.
- Capacity-scoped policies support up to 50 rules per policy.

- A condition supports up to 50 values.
- Policy changes take effect within 15 minutes of activation.

Related content

- Get started with Fabric policies
- External data sharing policy
- Item creation policy
- Automate Fabric policies with the REST APIs

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Get started with Microsoft Fabric policies (preview)

Microsoft Fabric policies (preview) help administrators govern Fabric by automatically enforcing rules that control user actions across Fabric resources and activities. This article explains how to enable Fabric policies, create a policy set, add policy rules, activate the policy, and manage active policies in the **Policies** center.

To learn more about Fabric policies, see What are Fabric policies?

Prerequisites for creating a Fabric policy

- The **Policies** tenant setting is enabled. See the next section, Enable the Policies feature in the tenant settings.
- A Fabric workspace assigned to a capacity.
- **Contributor**, **Member**, or **Admin** access to the workspace so that you can create a policy item type.
- The appropriate admin role to activate the policy:
  - **Fabric administrator**: Can create and activate policies for tenant-scope policies.
  - **Capacity administrator**: Can create and activate policies for capacity-scope policies.
- Entra security groups for any user-based rule conditions.
- **Microsoft Purview Information Protection** enabled if you plan to use sensitivity-label conditions.
- Understand the best practices for creating a policy.

> **Note**
>
> Creating or saving a policy doesn't start enforcement. The policy must be activated by an admin of its assigned scope.

Enable the Policies feature in the tenant settings

When you enable the **Policies** tenant setting, authorized users can create policies.

1. In the Fabric portal, open the **OneLake catalog**.

2. Select the **Govern** tab.
3. Select **Configurations**.
4. Filter by the keyword **Policies**.
5. Expand the **Policies** section and set the toggle to **Enabled**.
6. Under **Apply to**, select either **The entire organization** or specific security groups. For either option, you can also create exceptions by selecting **Except specific security groups** and choosing the relevant groups.

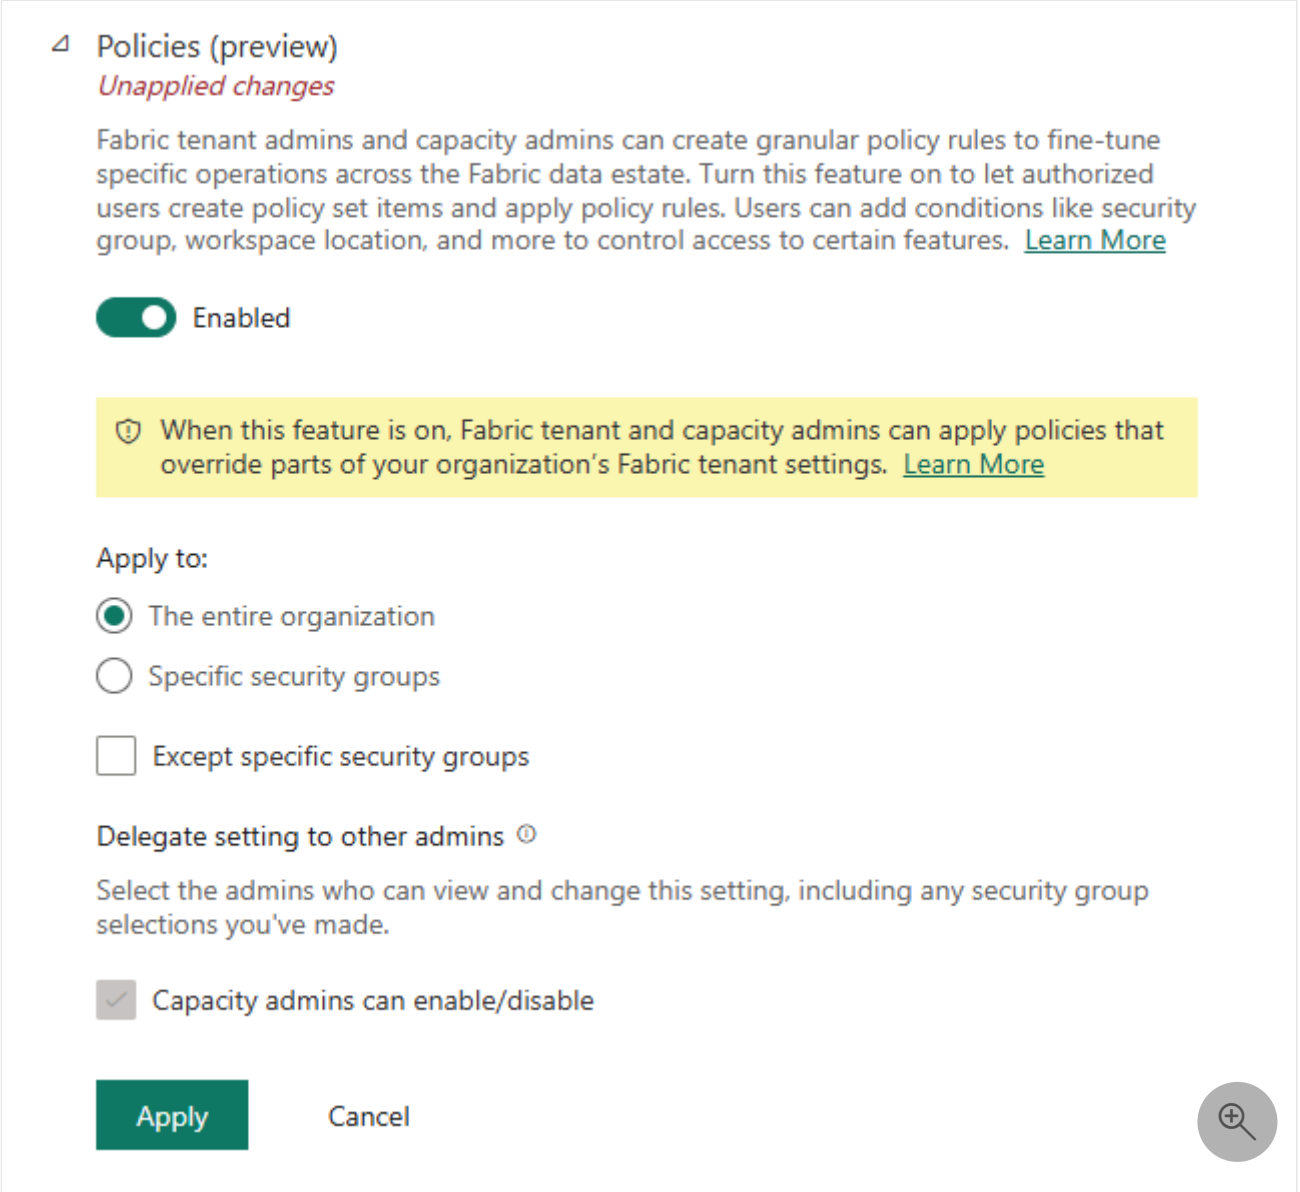

7. Select **Apply** to save the changes.

Create a policy item

In Fabric, a dedicated Fabric item called a *policy set* holds a policy and its rules. When you create the item, choose whether the policy applies at the tenant or capacity scope. You can't change the scope later.

1. In the Fabric portal, select **Workspaces** and open a workspace that's assigned to a Fabric capacity.
2. Select **New item**.
3. In the item gallery, search for and select **Policy set**. If there are no results, select **All items** and search again.
4. In **New policy set** dialog, enter a unique, descriptive **Name** for the policy set. You can't use spaces.
5. Under **Policy set scope**, select the appropriate scope:
   - Select **Fabric tenant** for policies that a Fabric administrator activates and apply at the tenant scope.
   - Select **A capacity** for policies that a capacity administrator activates for a specific capacity.

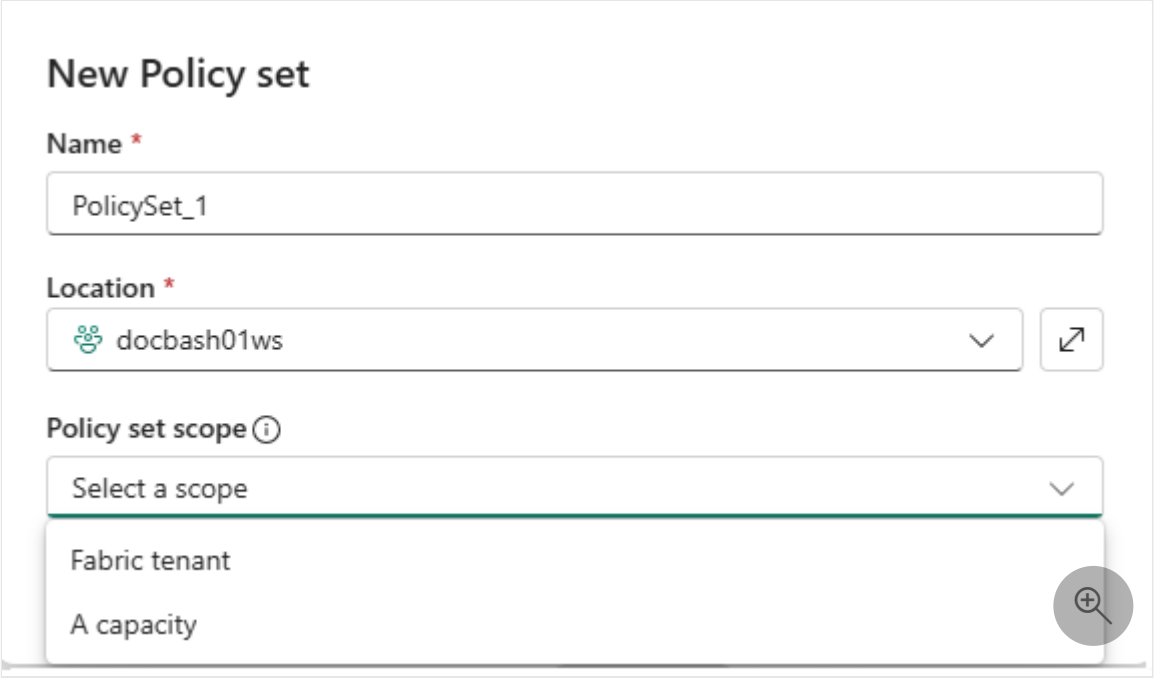

6. Select **Create** to create the policy set item in the workspace.

The new policy set is inactive and initially shows each supported policy in its system-default state.

> **Important**

You can't change the admin scope after the policy set is created. If you select the wrong scope, create a new policy set with the correct scope.

Review the policy defaults

Before you add a rule, review the default behavior shown for the policy.

- **Allowed** means the governed action follows the system default and remains available to users who otherwise have permission.
- **Blocked** means the governed action is denied.
- **Advanced** means Fabric evaluates the configured allow rules.

Policies with **Allow** behavior use the rules configured in **Advanced** settings as an allow list:

- If any rule matches, Fabric allows the operation.
- If no rule matches, Fabric denies the operation.
- Conditions in one rule use **AND** logic.
- Separate rules use **OR** logic.

> ⊗ **Caution**
>
> Adding a narrowly scoped allow rule blocks any operations that don't match it. Add rules for all required allow scenarios before activating the policy set.

Add policy rules

You can apply a single **Allowed** or **Blocked** behavior, or choose the **Advanced** option to use the no-code conditions builder for granular rules.

1. Open the workspace, and then open the policy item.
2. Locate the policy that you want to configure.
3. Select **Allowed**, **Blocked**, or **Advanced (Set by user)** to change the policy behavior. If you select **Advanced**, the rules configuration is enabled.
4. To create an advanced rule, select **Add rule** or **Create rule**.
5. Enter a clear rule **Name** and **Description** that explains who the rule applies to, what it allows, and where it applies:

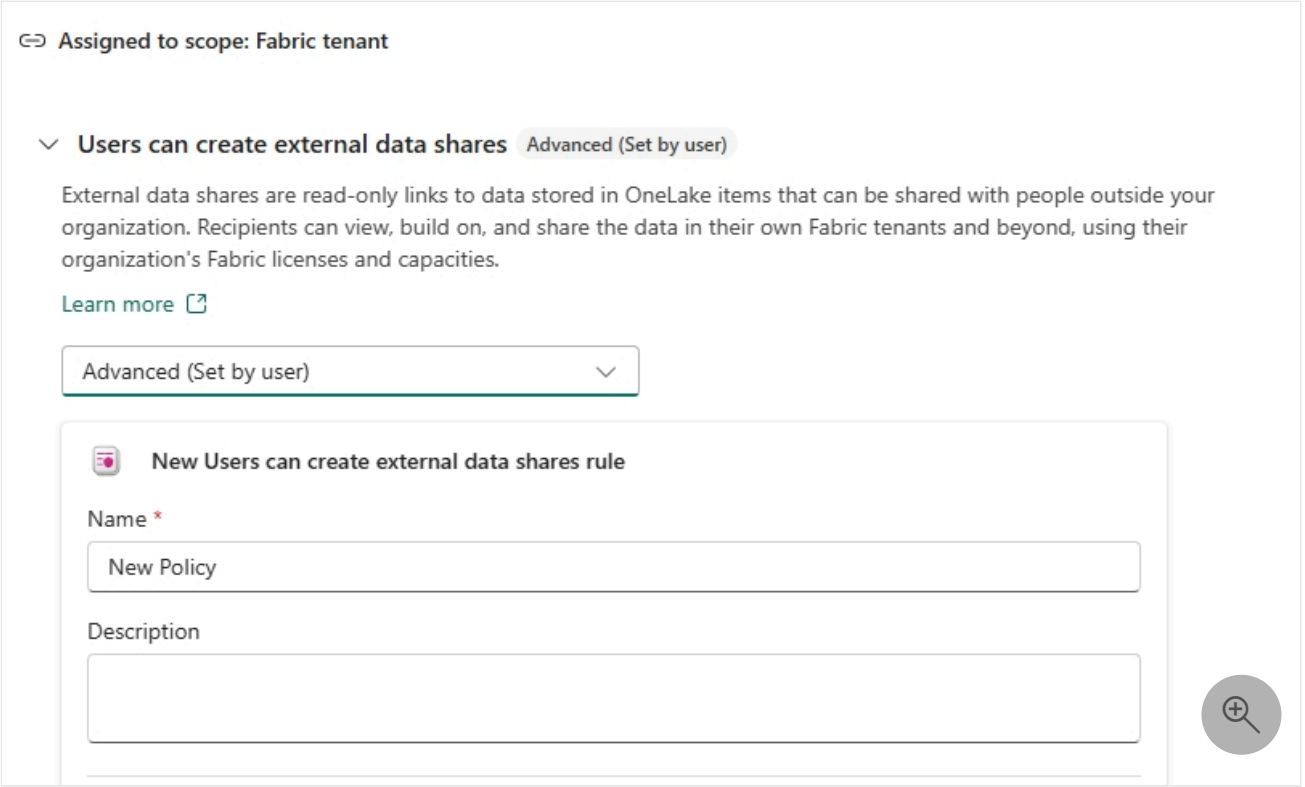

6. Under **Condition type**, choose a condition type supported by the policy.

Expand table

Condition	Purpose
Recipient email domain	Match an external recipient domain for external data sharing.
Security groups	Match the acting user's Entra security-group membership.
Sensitivity labels	Match the Microsoft Purview sensitivity label applied to an item.
Workspaces	Match the workspace where the governed operation occurs.
Item type	Match the Fabric item type being created.

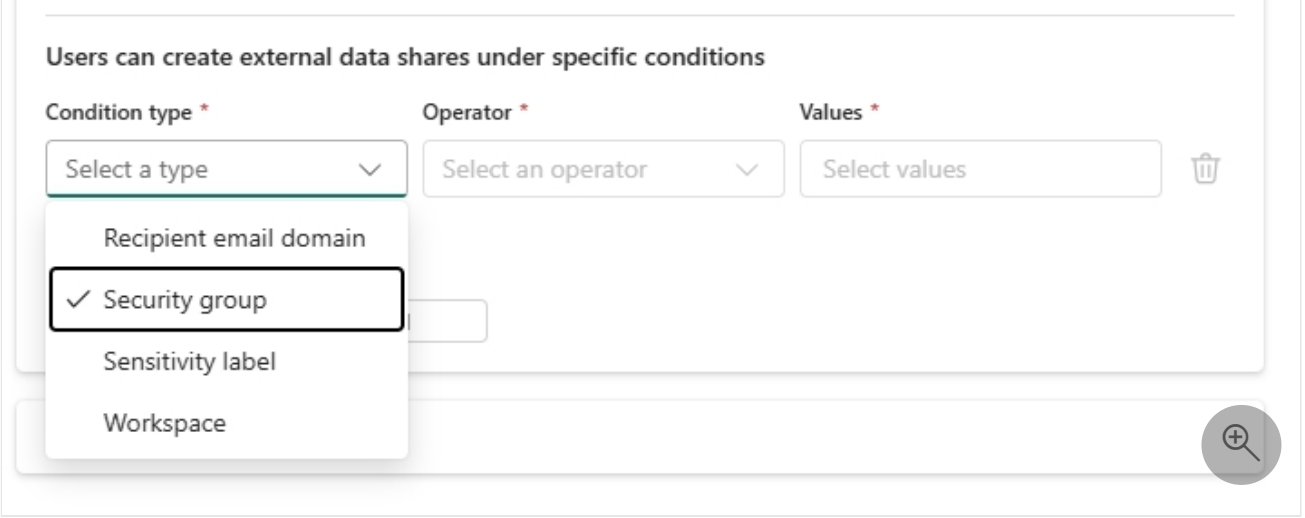

> **Note**
>
> If you omit a condition, the rule isn't restricted by that condition type. For example, a rule without a workspace condition can match operations in any workspace within the policy's admin scope.

7. Select an operator:
    - **Is any of** matches when at least one configured value matches.
    - **Is none of** matches when none of the configured values match.
8. Select or enter one or more **Values** that match the selected condition type.
9. Repeat the condition steps until the rule describes the complete allow scenario.
10. Review the rule. All conditions in the rule must be true for the rule to match.
11. Select **Apply**.

The rule is added to the policy set. If you need to allow a different combination of users, workspaces, item types, labels, or domains, create another rule.

Assign and activate the policy set

Only an administrator of the selected scope can activate a policy set. Only one policy set can be active for an admin scope at a time.

> **Caution**
>
> Activating a policy set can immediately change which operations users are allowed to perform after propagation. Validate the rules and affected user groups before activation.

1. Open the policy set from its workspace. Review each configured policy and rule, and then select **Active**:

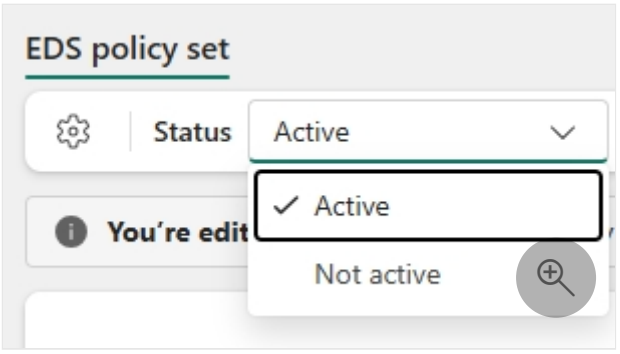

2. For a capacity-scoped policy set, select the target capacity and review the confirmation message.

3. If another policy set is active for the scope, confirm that you want to replace it and confirm the activation.

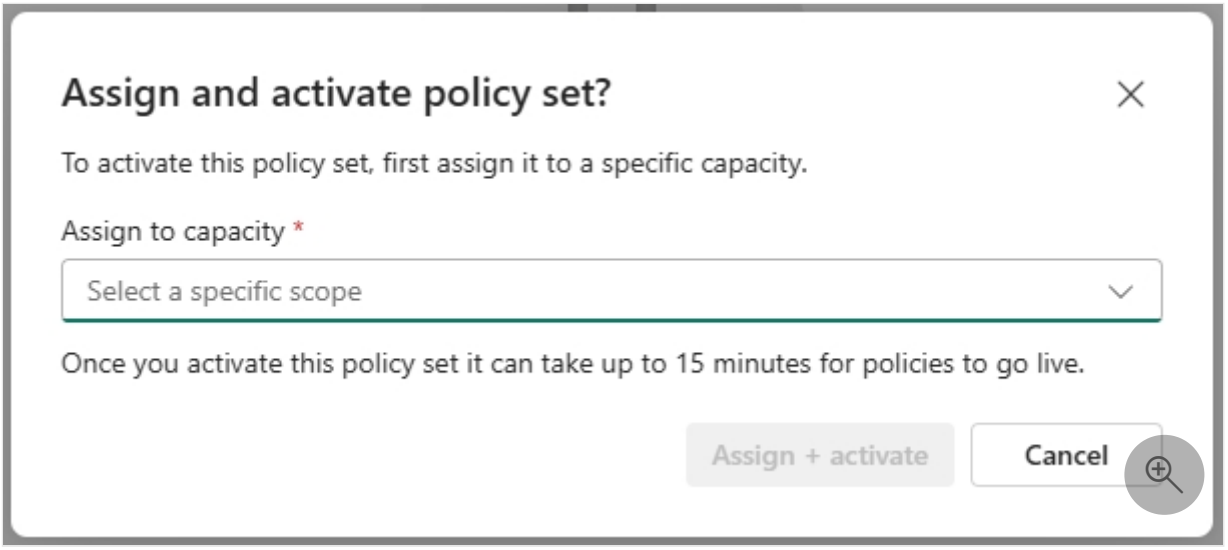

The policy set becomes active. Policy changes can take up to 15 minutes to take effect.

Validate policy enforcement

We recommend validating policy enforcement with test accounts after waiting up to 15 minutes for changes to take effect. Confirm that an eligible user who matches a rule is allowed and that an eligible user who doesn't match any rule is denied. You can also test a user outside the configured security groups and review the Fabric audit logs to verify the enforcement results.

Deactivate or replace a policy set

You can deactivate or replace an active policy set from the workspace where the policy set is stored. Only an administrator of the policy set's admin scope can deactivate or replace it.

Deactivate a policy set

1. Open the active policy set.
2. Select **Deactivate**.
3. Confirm the deactivation.

Deactivation stops the policy set from being evaluated and returns the scope to the applicable system-default behavior. The policy set remains in the workspace.

Replace an active policy set

1. Create and configure the replacement policy set in a workspace.
2. Review all allow scenarios in the replacement.
3. Select **Activate**.
4. Confirm that the existing active policy set will be replaced.

The previous policy set becomes inactive, and the replacement takes effect within 15 minutes.

Monitor enforcement with Fabric audit logs

Fabric records each policy decision so you can verify that enforcement matches your intent. Track these decisions through Fabric audit logs.

1. Open the Fabric audit logs.
2. Review the recorded policy enforcement decisions to confirm that Fabric applies your policies as expected.

Considerations and limitations

- You can currently create policy sets only from the item creation experience in a workspace.
- Use the Policies Center for centralized policy management, not for policy-set creation.
- Only workspaces assigned to a Fabric capacity can store policy sets.
- You can't change a policy set's admin scope after creation.
- Only one policy set can be active for an admin scope at a time.
- Policy and tenant-setting changes can take time to propagate.
- Conditions support **Any of** and **None of** in the current preview.
- A condition supports up to 50 values.
- Tenant-scoped policies support up to 100 rules per policy.
- Capacity-scoped policies support up to 50 rules per policy.
- Domain and workspace admin scopes aren't currently available.
- You can't activate a policy set through Git.

Related content

- What are Fabric policies?
- External data sharing policy
- Item creation policy
- Automate Fabric policies with the REST APIs

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Manage Microsoft Fabric policies in the OneLake catalog Policies center (preview)

The Policies center in the OneLake catalog is a management experience for all active Fabric policies (preview). Use it to change the scope you are viewing, review policy status, and add, update, or delete policy rules.

Prerequisites

To manage policies, you need:

- The Policies tenant setting enabled.
- A Fabric administrator role to manage tenant-scoped policies.
- A capacity administrator role for each capacity whose policies you manage.
- **ReadWrite** access to the active policy set.

Open the Policies center

1. In the OneLake catalog, select **Govern** > **Policies**.
2. Choose an admin scope: **Tenant** or **Capacity**.

The Policies center displays the policies available for the selected administrative scope.

Change the admin scope and view policies

Use the scope explorer to switch between the policies that apply to different administrative scopes.

1. In the Policies center, open the scope explorer.
2. Select Tenant to view tenant-scoped policies, or select a capacity to view the policies for that capacity.
3. Review the active policy set and policy configuration for the selected scope.

The scopes you can view depend on your administrator role:

Expand table

Administrator role	Available scopes
Fabric administrator	Tenant and capacity scopes
Capacity administrator	Capacities for which you are an administrator

> **Note**
>
> Changing the selected scope changes only the Policies center view. It doesn't change the admin-scope type assigned to a policy set. A policy set created for the Tenant scope can't be changed to the Capacity scope, and a capacity-scoped policy set can't be changed to the Tenant scope.

Review active policies

For the selected scope, use the policy list to:

- Review the active policy set.
- Search for a policy by its name or description.
- See whether each policy uses its system default or was configured by an administrator.
- Review whether an allow policy is set to Allow all, Blocked, or Advanced.
- Open a policy to review its rules and conditions.

Only one policy set can be active for an administrative scope at a time.

Manage a policy

1. In the Policies center, select the required administrative scope.
2. Select the policy that you want to manage.
3. Review the policy state and its existing rules.
4. Add, update, or delete rules as required.

Changes made in the Policies center are synchronized with the active policy set item. Changes made directly in the active policy set item are also reflected in the Policies center.

> **Caution**

Advanced allow policies use allow list behavior. If at least one rule matches, the operation is allowed. If no rule matches, the operation is denied. Configure every required allow scenario before you save changes.

Add a policy rule

1. Open the policy that you want to update.
2. Select **Add rule** or **Create rule**.
3. Enter a clear rule name and description.
4. Add the conditions that define the allow scenario.
5. Select an operator and one or more values for each condition.
6. Review the rule, and then select **Save**.

All conditions in one rule use **AND** logic. Separate rules use **OR** logic.

> **Note**
>
> A security-group condition narrows which eligible users can match a rule. It doesn't make a user eligible for enforcement if the user isn't in a group configured in the **Policies** tenant setting.

Update a policy rule

1. Open the policy that contains the rule.
2. Select the rule that you want to update.
3. Change the rule name, description, conditions, operators, or values.
4. Review the effect of the changes.
5. Select **Save**.

Changes to an active policy set can take up to 15 minutes to affect enforcement.

Delete a policy rule

1. Open the policy that contains the rule.
2. Open the rule's context menu, and then select **Delete**.
3. Review the effect of removing the allow scenario.
4. Confirm the deletion.

> **⊗ Caution**
>
> Deleting a rule from an active allow policy can cause previously allowed operations to be denied when no remaining rule matches.

Considerations and limitations

- The Policies center is a management experience and doesn't provide the current policy-set creation flow.
- You can create policy sets only through **New item** > **Policy set** in a workspace assigned to a Fabric capacity.
- You can't activate or deactivate policy sets from the Policies center. Open the policy set item in its workspace to perform these actions.
- Only one policy set can be active for an administrative scope at a time.
- You can't change a policy set's Tenant or Capacity admin-scope type after creation.
- Policy changes can take up to 15 minutes to affect enforcement.
- Policy enforcement is limited to users included through the security groups configured in the Policies (preview) tenant setting.
- Domain and workspace administrative scopes aren't available in the current preview.

Related content

- What are Fabric policies?
- Get started with Fabric policy sets
- Fabric policies REST APIs
- External data sharing policy
- Item creation policy

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Fabric policies REST APIs (preview)

By using the Fabric policies REST APIs (preview), you can programmatically create, configure, and activate policy sets. Automating these operations helps you manage governance as code and integrate policy management into your CI/CD workflows, rather than configuring each policy set manually in the Policies Center.

A policy set contains policies and rule definitions, and you scope and activate it. Fabric stores each policy set in a policy item. Because the policy item is a native Microsoft Fabric item, it supports standard Fabric item capabilities, including public REST APIs, the standard permission model, Git integration, and CI/CD. This article introduces the policy APIs and explains how the item model supports policy-as-code and automated deployment.

Fabric policies REST API authentication and identities

The Fabric policies REST APIs use Microsoft Entra ID for authentication, the same model as the rest of the Fabric REST API. The APIs support:

- User tokens
- Service principals
- Managed identities

By using service principals and managed identities, you can run policy automation without an interactive sign-in, which makes unattended CI/CD pipelines possible.

Git integration and CI/CD for Fabric policy sets

Because Fabric stores a policy set in a policy item, you can use Git integration and CI/CD to version, promote, and deploy governance as code across your environments. By storing policy items in source control, you can review policy set changes, track history, and move a validated configuration from development to production the same way you handle application code.

Git doesn't support activation. To enforce a policy set, use the activation API or the Policies Center. This separation lets you manage the definition of a policy set as code while keeping enforcement an explicit, controlled action.

Fabric policies REST API lifecycle operations

The Fabric policies REST APIs map to the stages of the policy set lifecycle: create, configure, activate, manage, and evaluate. The following table describes the operations available at each step.

Expand table

Step	API	Description
Create	Create Policy Set	Create a policy set with an admin scope and store it in a policy item in a workspace.
Configure	Add Policy Rule	Add a rule definition (policy, conditions, effect) to a policy set with schema validation.
Configure	Update / Delete / List Policy Rule	Manage rules within a policy set.
Activate	Activate Policy Set (Tenant / Capacity)	Enforce the policy set for the scope.
Manage	Deactivate Policy Set, List Policy Sets, Read Active Policy Sets	Manage and discover policy sets and their policy items.
Evaluate	Policy Evaluation API	Runtime check returning Allow, Denied, or System default.

For full request and response details for each operation, see the Fabric REST API reference.

Fabric policy item operations and policy rule operations

The API surface covers two kinds of operations:

- **Policy item operations.** You can create, read, update, and delete the policy item that stores a policy set like any other Fabric item.
- **Policy rule operations.** Within the policy set, you can create, update, delete, and list its individual rules.

Because the policy item APIs extend the Fabric Items REST API, the standard Fabric item operations and permission model apply to the policy item. Activation applies to the policy set and requires the appropriate scope administrator: a Fabric administrator for a Tenant scope or a capacity administrator for a Capacity scope.

Related content

- What are Fabric policies?
- Get started with Fabric policy sets

- External data sharing policy
- Item creation policy

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Microsoft Fabric policy: Allow external data sharing

The **Allow external data sharing** policy controls who can share Microsoft Fabric data externally, from which workspaces, for which sensitivity labels, and with which recipient email domains. Tenant admins configure this policy type.

The policy is evaluated when a user creates or modifies an external data share. The external data sharing feature supplies the recipient domain and evaluation identifiers to the policy evaluation API.

Properties

Expand table

Property	Value
Policy type	ExternalDataSharing
Scope	Tenant
Evaluated when	A user creates or modifies an external data share
Rule logic	All conditions in a rule must match

Effects

Supported effects

Expand table

Effect	Behavior
Allow	Permits the external-sharing operation when the rule matches.

The **Create Policy Rule** API doesn't include an effect property. The **ExternalDataSharing** policy type defines the allow behavior, and an unmatched request is blocked when policy enforcement is active.

Default effect

- When an active policy doesn't manage external data sharing, the system uses the legacy tenant-setting configuration as the default.
- When policy enforcement is active, an operation must match an Allow rule or the operation is denied.
- Omitting a condition means that dimension doesn't restrict the rule. For example, omitting `workspace.id` applies the rule to all workspaces.

Conditions

You define conditions in the conditions array of a policy rule. Each entry is a DynamicCondition that resolves a target property from the external-sharing evaluation context and applies a predicate to produce a true-or-false result.

Key concepts

Expand table

Concept	Description
DynamicCondition	A condition that resolves a property from the evaluation context. Its JSON discriminator is `"type": "Dynamic"`.
targetProperty	A dot-separated property path, such as `item.sensitivityLabel.id` or `externalDataShareRecipient.emailDomain`.
predicate	The comparison definition that contains operator and values.
operator	The comparison to perform. External Data Sharing supports `AnyOf` and `NoneOf`.
values	An array of strings. Each string must be a valid GUID or email-domain value for the selected target property.

Condition structure

JSON

```
{
  "type": "Dynamic",
  "targetProperty": "<propertyPath>",
  "predicate": {
    "operator": "<operator>",
    "values": ["<value>"]
  }
}
```

Expand table

API property	Type	Required	Description
type	string	Yes	Condition discriminator. Use Dynamic.
targetProperty	string	Yes	A dot-separated property path supported by ExternalDataSharing.
predicate	object	Yes	Predicate applied to the resolved target property.
predicate.operator	string	Yes	AnyOf or NoneOf.
predicate.values	array of string	Yes	1-50 values in the format required by targetProperty.

Predicates

Expand table

Operator	Predicate example	Result
AnyOf	{"operator":"AnyOf","values":["@contoso.com"]}	True when the resolved domain is @contoso.com.
NoneOf	{"operator":"NoneOf","values":["@contoso.com"]}	True when the resolved domain isn't @contoso.com.

- AnyOf is true when at least one resolved value equals one configured value.
- NoneOf is true when no resolved value equals a configured value.
- Values within a predicate use OR logic.
- Conditions within a rule use AND logic.
- Separate rules use OR logic.
- AnyOf and NoneOf accept 1-50 values.
- Omit a condition when that dimension shouldn't restrict the rule.

Supported root entities for this policy type

Each policy type exposes a different set of root entities. The ExternalDataSharing policy type supports the root entities listed in the following table. The root entity is the first segment of `targetProperty`. A property that references another entity links to that entity's definition.

Expand table

Root entity	Entity type	Source	Description
principal	Principal	Platform-provided user context	The effective user performing the sharing operation.
workspace	Workspace	Platform-provided workspace context	The workspace from which the item is shared.
item	Item	Platform-provided item context	The Fabric item being shared.
externalDataShareRecipient	ExternalDataShareRecipient	External data sharing evaluation context	The external recipient supplied by the external data sharing enforcement point.

Principal

Root path: principal

The platform resolves the effective human user's transitive Microsoft Entra security-group membership. Display names can appear in the editor, but policy JSON stores stable group object IDs.

Expand table

Property	Type	Supported operators	Values	Description
groups	`list<Group>`	Select the `id` property	None	The effective user's transitive Entra security groups.

Example entity path: `principal.groups.id`.

Group

Referenced from: `principal.groups`

Expand table

Property	Type	Supported operators	Values	Description
id	guid	AnyOf, NoneOf	1-50 security-group object IDs	The stable Entra object ID of a security group.

Workspace

Available paths: `workspace`, `item.workspace`

The API supports the workspace supplied directly in the evaluation context and the workspace referenced by the shared item.

Expand table

Property	Type	Supported operators	Values	Description
id	guid	AnyOf, NoneOf	1-50 Fabric workspace IDs	The source workspace's stable Fabric object ID.

Example entity paths: `workspace.id` and `item.workspace.id`.

Item

Root path: `item`

The platform resolves the Fabric item being shared.

Expand table

Property	Type	Supported operators	Values	Description
sensitivityLabel	SensitivityLabel	Select the `id` property	None	The Microsoft Purview sensitivity label applied to the item.
workspace	Workspace	Select the `id` property	None	The workspace that contains the shared item.

Example entity paths: `item.sensitivityLabel.id` and `item.workspace.id`.

SensitivityLabel

Referenced from: item.sensitivityLabel

The label picker requires the Microsoft Purview Information Protection tenant setting to be enabled. It shows labels available through the acting admin's Purview label policy. JSON can reference other valid label IDs.

Expand table

Property	Type	Supported operators	Values	Description
id	guid	AnyOf, NoneOf	1-50 sensitivity-label IDs	The stable Purview sensitivity-label ID.

For an unlabeled item, use the supported **No label** authoring option rather than a display name.

ExternalDataShareRecipient

Root path: externalDataShareRecipient

The external data sharing feature supplies this custom entity directly in the evaluation request. The policy platform doesn't fetch it.

Expand table

Property	Type	Supported operators	Values	Description
emailDomain	string	AnyOf, NoneOf	1-50 normalized domains in the form @domain.tld	The domain extracted from the external recipient address. Comparisons are case-insensitive.

Example entity path: externalDataShareRecipient.emailDomain.

Common scenarios and examples

Allow approved teams to share with approved domains

This rule allows Sales and Partner Success users to share labeled content from selected workspaces with two approved domains. Every condition must match.

JSON

```
{
  "displayName": "Allow approved partner sharing",
  "description": "Allows approved teams to share labeled data from selected workspaces
```

```
with approved domains.",
  "policy": "ExternalDataSharing",
  "conditions": [
    {
      "type": "Dynamic",
      "targetProperty": "principal.groups.id",
      "predicate": {
        "operator": "AnyOf",
        "values": [
          "8a6e2aa2-6e4d-49ad-a970-2e757ac8a2f1",
          "b2d465a3-734e-47f6-82b7-2fdddb9a91d5"
        ]
      }
    },
    {
      "type": "Dynamic",
      "targetProperty": "workspace.id",
      "predicate": {
        "operator": "AnyOf",
        "values": [
          "c9d4e5f6-a7b8-49c0-9234-56789abcdef0",
          "d0e5f6a7-b8c9-40d1-8345-6789abcdef01"
        ]
      }
    },
    {
      "type": "Dynamic",
      "targetProperty": "item.sensitivityLabel.id",
      "predicate": {
        "operator": "AnyOf",
        "values": [
          "f42a3342-8706-4288-bd31-ebb85995028c",
          "75a9e7c2-7d3f-4566-8f2b-5c123f58a9f8"
        ]
      }
    },
    {
      "type": "Dynamic",
      "targetProperty": "externalDataShareRecipient.emailDomain",
      "predicate": {
        "operator": "AnyOf",
        "values": ["@contoso.com", "@fabrikam.com"]
      }
    }
  ]
}
```

Add a separate finance exception

Rules use OR logic with each other. This second rule allows Finance Leadership to share with an external auditor without changing the first rule.

JSON

```
{
  "displayName": "Allow finance auditor sharing",
  "description": "Allows Finance Leadership to share with the approved external
auditor.",
  "policy": "ExternalDataSharing",
  "conditions": [
    {
      "type": "Dynamic",
      "targetProperty": "principal.groups.id",
      "predicate": {
        "operator": "AnyOf",
        "values": ["db0fc5e1-f9de-47de-909e-417c611bf4ab"]
      }
    },
    {
      "type": "Dynamic",
      "targetProperty": "externalDataShareRecipient.emailDomain",
      "predicate": {
        "operator": "AnyOf",
        "values": ["@finance-auditor.com"]
      }
    }
  ]
}
```

Related content

- What are Fabric policies?
- Get started with Fabric policy sets
- Item creation policy
- Automate Fabric policies with the REST APIs
- External data sharing in Microsoft Fabric

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Microsoft Fabric policy: Allow item creation

The **Allow item creation** policy controls who can create selected Microsoft Fabric item types in workspaces assigned to a capacity. Capacity admins configure this policy type.

The policy is evaluated when a user attempts to create a Fabric item. The creation flow supplies the requested item type, user context, and target workspace to the policy evaluation API.

Expand table

Property	Value
Policy type	`ItemCreation`
Scope	Capacity
Evaluated when	A user attempts to create a Fabric item
Rule logic	All conditions in a rule must match

Effects

Supported effects

Expand table

Effect	Behavior
`Allow`	Permits the item-creation operation when the rule matches.
`Block` (implicit)	Denies the item-creation operation when advanced policy enforcement is active and no `Allow` rule matches. You don't author `Block` in the create request; it is the outcome for unmatched requests.

The **Create Policy Rule** API doesn't include an effect property. The `ItemCreation` policy type defines the allow behavior, and blocking is an implicit policy outcome.

Default effect

The default effect is `Allow`. When no active policy or policy rule applies, users can create items according to their workspace permissions.

Policy action behavior

- **Allow all** permits every item type according to the user's workspace permissions.
- **Blocked** denies all item-creation operations covered by the policy.
- **Advanced** uses allow list behavior. Matching `Allow` rules permit an operation; every operation that doesn't match an `Allow` rule is implicitly blocked.
- Omitting a condition means that dimension doesn't restrict the rule. For example, omitting `workspace.id` applies the rule to all workspaces in the selected capacity.

> **Important**
>
> Switching to **Advanced** and adding an `Allow` rule changes the policy from allow-all behavior to allow list behavior. For example, if a rule allows creation of only `Notebook` and `Lakehouse` items, all other item types are immediately blocked unless another active rule allows them.

Conditions

Define conditions in the `conditions` array of a policy rule. Each entry is a `DynamicCondition` that resolves a target property from the item-creation evaluation context and applies a predicate to produce a true-or-false result.

Key concepts

Expand table

Concept	Description
`DynamicCondition`	A condition that resolves a property from the evaluation context. Its JSON discriminator is `"type": "Dynamic"`.
`targetProperty`	A dot-separated property path, such as `item.type` or `principal.groups.id`.
`predicate`	The comparison definition that contains `operator` and `values`.
`operator`	The comparison to perform. Item Creation supports `AnyOf` and `NoneOf`.
`values`	An array of strings. Each string must be a valid API item-type identifier or GUID for the selected target property.

Condition structure

JSON

```
{
  "type": "Dynamic",
  "targetProperty": "<propertyPath>",
  "predicate": {
    "operator": "<operator>",
    "values": ["<value>"]
  }
}
```

Expand table

API property	Type	Required	Description
`type`	string	Yes	Condition discriminator. Use `Dynamic`.
`targetProperty`	string	Yes	Dot-separated property path supported by `ItemCreation`.
`predicate`	object	Yes	Predicate applied to the resolved target property.
`predicate.operator`	string	Yes	`AnyOf` or `NoneOf`.
`predicate.values`	array of string	Yes	1-50 values in the format required by `targetProperty`.

Predicates

Expand table

Operator	Predicate example	Result
`AnyOf`	`{"operator":"AnyOf","values":["Notebook"]}`	True when the requested item type is `Notebook`.
`NoneOf`	`{"operator":"NoneOf","values":["Notebook"]}`	True when the requested item type isn't `Notebook`.

- `AnyOf` is true when at least one resolved value equals one configured value.
- `NoneOf` is true when no resolved value equals a configured value.
- Values within a predicate use **OR** logic.
- Conditions within a rule use **AND** logic.
- Separate rules use **OR** logic.
- `AnyOf` and `NoneOf` accept 1-50 values.

- Omit a condition when that dimension shouldn't restrict the rule.

Supported root entities for this policy type

Each policy type exposes a different set of root entities. The `ItemCreation` policy type supports the root entities listed in the following table. The root entity is the first segment of `targetProperty`.

Expand table

Root entity	Entity type	Source	Description
`principal`	`Principal`	Platform-provided user context	The effective user performing the create operation.
`item`	`Item`	Item-creation request	The item the user intends to create.
`workspace`	`Workspace`	Platform-provided workspace context	The workspace in which the item will be created.

`Principal`

Root path: `principal`

The platform resolves the effective human user's transitive Microsoft Entra security-group membership. Display names can appear in the editor, but policy JSON stores stable group object IDs.

Expand table

Property	Type	Supported operators	Values	Description
`groups`	`list<Group>`	Select the `id` property	None	The effective user's transitive Entra security groups.

Example entity path: `principal.groups.id`.

`Group`

Referenced from: `principal.groups`

Expand table

Property	Type	Supported operators	Values	Description
`id`	`guid`	`AnyOf`, `NoneOf`	1-50 security-group object IDs	The stable Entra object ID of a security group.

`Item`

Root path: `item`

The item-creation flow supplies this custom entity because the requested item doesn't exist yet.

Expand table

Property	Type	Supported operators	Values	Description
`type`	`enum`	`AnyOf`, `NoneOf`	1-50 supported API item-type identifiers	The type of Fabric item the user intends to create.

Example entity path: `item.type`.

> **Important**
>
> Use API item-type identifiers, not UI display names. For example, use `DataPipeline` for the UI item type **Pipeline**.

`Workspace`

Root path: `workspace`

The platform resolves the target workspace. For this capacity-scoped policy, referenced workspaces must belong to the selected capacity.

Expand table

Property	Type	Supported operators	Values	Description
`id`	`guid`	`AnyOf`, `NoneOf`	1-50 Fabric workspace IDs	The target workspace's stable Fabric object ID.

Example entity path: `workspace.id`.

`workspace.name` isn't a supported substitute. The API rejects it with `UnsupportedPropertyValue`.

Current REST API limitations

- Policy-set activation supports `ScopeType``: "Capacity"` for this policy type. `ScopeType``: "Workspace"` is rejected.
- A policy set can be active on only one scope at a time. Deactivate it before activating it on another capacity.

Common scenarios and examples

Allow data engineers to create engineering items

This rule allows members of the Data Engineers group to create notebooks, data pipelines, and lakehouses in two production workspaces. All three conditions must match.

JSON

```
    {
      "displayName": "Allow engineering item creation",
      "description": "Allows Data Engineers to create selected engineering items in
production workspaces.",
      "policy": "ItemCreation",
      "conditions": [
        {
          "type": "Dynamic",
          "targetProperty": "principal.groups.id",
          "predicate": {
            "operator": "AnyOf",
            "values": ["2f9f8a1e-6b77-4b1a-9af7-9a611f29de18"]
          }
        },
        {
          "type": "Dynamic",
          "targetProperty": "item.type",
          "predicate": {
            "operator": "AnyOf",
            "values": ["Notebook", "DataPipeline", "Lakehouse"]
          }
        },
        {
          "type": "Dynamic",
          "targetProperty": "workspace.id",
          "predicate": {
            "operator": "AnyOf",
            "values": [
```

```
            "a7b2c3d4-e5f6-47a8-9012-3456789abcde",
            "b8c3d4e5-f6a7-48b9-8123-456789abcdef"
          ]
        }
      }
    ]
  }
```

Apply an item-type rule across the capacity

Omit `workspace.id` to apply the rule to every workspace in the selected capacity.

JSON

```
  {
    "displayName": "Allow approved analytical items",
    "description": "Allows the Analytics Builders group to create approved item
types across the capacity.",
    "policy": "ItemCreation",
    "conditions": [
      {
        "type": "Dynamic",
        "targetProperty": "principal.groups.id",
        "predicate": {
          "operator": "AnyOf",
          "values": ["7b785ad9-1a87-4f31-a336-cbb8f0de28c4"]
        }
      },
      {
        "type": "Dynamic",
        "targetProperty": "item.type",
        "predicate": {
          "operator": "AnyOf",
          "values": ["Warehouse", "SemanticModel", "Report"]
        }
      }
    ]
  }
```

Related content

- What are Fabric policies?
- Get started with Fabric policy sets
- External data sharing policy
- Automate Fabric policies with the REST APIs

ⓘ**Note:** The author created this article with assistance from AI. Learn more

Last updated on 09/29/2026

Microsoft Fabric policy: Allow workspace settings editing

The **Allow workspace settings editing** policy is a tenant-scoped policy that controls which workspace admins can modify security-sensitive workspace settings. It lets tenant admins restrict changes to customer-managed keys and workspace networking settings without removing the workspace admin role or limiting unrelated workspace management tasks.

The policy governs **who can edit a setting**. It doesn't configure the setting value. Existing tenant settings continue to determine whether a capability is available.

Expand table

Property	Value
Policy type	WorkspaceSettingsEditing
Scope	Tenant
Evaluated when	A user opens the workspace settings in the Microsoft Fabric portal, or attempts to change a governed workspace setting through API
Rule logic	Conditions in a rule use AND logic. Rules use OR logic.

> **Important**
>
> A tenant setting takes precedence over this policy. If a tenant setting disables a capability, allowing a workspace admin to edit the related workspace setting doesn't enable the capability.

Effects

This policy type supports two effects that determine whether a workspace admin can edit a governed setting.

Expand table

Effect	Behavior
`Allow`	Permits the setting editing by the user when the rule matches.
`Block` (implicit)	Denies the editing of the setting when advanced policy enforcement is active and no `Allow` rule matches. `Block` isn't authored in the create request; it is the outcome for unmatched requests.

Default effect

The default effect is `Allow`. When no active policy or policy rule applies, workspace admins can edit all workspace settings, including those that this policy type supports.

Policy action behavior

- **Allow all** permits workspace admins to edit all supported network-security settings, subject to tenant settings status. Other workspace settings that this policy type doesn't support remain editable.
- **Blocked** denies all workspace admins from editing the network security settings on all workspaces. Other workspace settings that this policy type doesn't support remain editable.
- **Advanced** uses allow list behavior. Matching `Allow` rules permit an operation; every operation that doesn't match an `Allow` rule is implicitly blocked.
- Omitting a condition means that dimension doesn't restrict the rule. For example, omitting `workspace.id` applies the rule to all workspaces.

> **Important**
>
> Switching to **Advanced** and adding an `Allow` rule changes the policy from allow-all behavior to allow list behavior. For example, if a rule allows editing of customer-managed keys, all other security settings (outbound and inbound settings) are immediately blocked unless another active rule allows them.

Conditions

Define conditions in the `conditions` array of a policy rule. Each entry is a `DynamicCondition` that resolves a target property from the edit-workspace-setting evaluation context and applies a predicate to produce a true-or-false result.

Key concepts

Expand table

Concept	Description
`DynamicCondition`	A condition that resolves a property from the evaluation context. Its JSON discriminator is `"type": "Dynamic"`.
`targetProperty`	A dot-separated property path, such as `workspaceSettingsEditing.settingType` or `principal.groups.id`.
`predicate`	The comparison definition that contains `operator` and `values`.
`operator`	The comparison to perform. Workspace settings editing supports `AnyOf` and `NoneOf`.
`values`	An array of strings. Each string must be a valid API item-type identifier or GUID for the selected target property.

Condition structure

JSON

```
{
  "type": "Dynamic",
  "targetProperty": "workspaceSettingsEditing.settingType",
  "predicate": {
    "operator": "AnyOf",
    "values": ["Encryption", "OutboundNetworking"]
  }
}
```

Expand table

API property	Type	Required	Description
`type`	string	Yes	Condition discriminator. Use `Dynamic`.
`targetProperty`	string	Yes	Dot-separated property path supported by WorkspaceSettingsEditing.
`predicate`	object	Yes	Predicate applied to the resolved target property.
`predicate.operator`	string	Yes	`AnyOf` or `NoneOf`.
`predicate.values`	array of string	Yes	1-50 values in the format required by `targetProperty`.

Predicates

Expand table

Operator	Predicate example	Result
`AnyOf`	`{"operator":"AnyOf","values":["Encryption","OutboundNetworking"]}`	True when the setting edit requested field is `customer managed keys`.
`NoneOf`	`{"operator":"NoneOf","values":["Encryption"]}`	True when the setting edit request isn't `customer managed keys`.

- `AnyOf` is true when at least one resolved value equals one configured value.
- `NoneOf` is true when no resolved value equals a configured value.
- Values within a predicate use **OR** logic.
- Conditions within a rule use **AND** logic.
- Separate rules use **OR** logic.
- `AnyOf` and `NoneOf` accept 1-50 values.
- Omit a condition when that dimension shouldn't restrict the rule.

Supported root entities for this policy type

Each policy type exposes a different set of root entities. The WorkspaceSettingsEditing policy type supports the root entities listed in the following table. The root entity is the first segment of `targetProperty`.

Expand table

Root entity	Entity type	Source	Description
principal	`Principal`	Platform-provided user context	The effective user performing the edit operation. Must be a workspace admin on the specified workspace.
workspaceSettingsEditing	`WorkspaceSettingsEditing`	Workspace setting edit request	The workspace setting the user attempts to edit.
workspace	`Workspace`	Platform-provided workspace context	The workspace in which the setting will be edited.

`Principal`

Root path: `principal`

The platform resolves the effective human user's transitive Microsoft Entra security-group membership. Display names can appear in the editor, but policy JSON stores stable group object IDs.

Expand table

Property	Type	Supported operators	Values	Description
groups	`list<Group>`	Select the `id` property	None	The effective user's transitive Microsoft Entra security groups.

Example entity path: `principal.groups.id`.

`Group`

Referenced from: `principal.groups`

Expand table

Property	Type	Supported operators	Values	Description
`id`	`guid`	`AnyOf`, `NoneOf`	1-50 security-group object IDs	The stable Microsoft Entra object ID of a security group.

`WorkspaceSettingsEditing`

Root path: `workspaceSettingsEditing`

Expand table

Property	Type	Supported operators	Values	Description
`settingType`	`string`	`AnyOf`, `NoneOf`	`Encryption` or `OutboundNetworking`	The type of workspace setting the user attempts to edit.

Example entity path: `workspaceSettingsEditing.settingType`.

`Workspace`

Root path: `workspace`

The platform resolves the target workspace.

Expand table

Property	Type	Supported operators	Values	Description
`id`	`guid`	`AnyOf`, `NoneOf`	1-50 Fabric workspace IDs	The target workspace's stable Fabric object ID.

Example entity path: `workspace.id`.

Current REST API limitations

- Policy-set activation supports `ScopeType: "Tenant"` for this policy type. The API rejects `ScopeType: "Workspace"`.

Related content

- What are Microsoft Fabric policies?
- Get started with Microsoft Fabric policies
- Item creation policy

Last updated on 09/29/2026

Use Microsoft Purview to protect Microsoft Fabric

Microsoft Fabric provides built-in governance capabilities for discovering, governing, and managing data. Microsoft Purview works with Fabric to protect sensitive data and support audit, risk, and compliance scenarios. This article explains when to use each experience.

Governance in Microsoft Fabric

Microsoft Fabric includes built-in governance capabilities that help organizations discover, govern, manage, and protect data.

The OneLake catalog is the primary Fabric experience for discovering and governing Fabric data. From the catalog, data consumers and data owners can find trusted data, understand where it comes from, and act on its governance state without leaving Fabric. Built-in governance capabilities include:

- **Discovery** — find and explore data across your Fabric estate from a single catalog experience.
- **Governance Insights** — see the governance state of your data across the organization and for the data you own.
- **Recommended Actions** — get prioritized guidance to improve the governance posture of your data.
- **Domains** — organize data by business area to support federated, distributed ownership.
- **Endorsement** — promote or certify trusted items so consumers can find high-quality, approved data.
- **Lineage** — trace how data flows across Fabric items, from source to report.

To learn more about these capabilities, see Fabric governance documentation, Governance and compliance in Microsoft Fabric, OneLake catalog overview, and Govern your Fabric data with the OneLake catalog.

Fabric brings governance and administrative experiences together in the **Govern** section of the OneLake catalog. Use **Govern** to review governance insights and recommended actions and to access tenant configuration, workspace and capacity management, policies, and other organization-wide controls.

A single person often holds more than one of these responsibilities. Someone who manages capacities might also be a Fabric administrator, a domain administrator, a workspace administrator, or a data steward. The **Govern** section is organized around what you're trying to accomplish rather than around a single role.

When to use Microsoft Purview with Fabric

Microsoft Purview works with Fabric to support information protection, compliance, risk management, and security scenarios.

Microsoft Fabric provides the built-in governance capabilities described previously. Microsoft Purview adds protection and monitoring when your security, risk, and compliance needs extend beyond Fabric. Consider Microsoft Purview when you need to:

- Discover and classify sensitive data, and apply protection with information protection and sensitivity labels.
- Prevent risky sharing or movement of sensitive data with data loss prevention.
- Audit and investigate user activity for compliance and security.
- Address broader compliance and risk-management requirements.

Using Microsoft Purview doesn't change where Fabric governance lives. The OneLake catalog remains the primary experience for governing Fabric data, and Microsoft Purview adds protection, audit, risk, and compliance coverage.

How Fabric and Purview work together

Microsoft Fabric surfaces its governance capabilities natively, and Microsoft Purview integrates with Fabric so you can protect and monitor Fabric data. The following integrations are available today:

- **Microsoft Purview Information Protection** — discover, classify, and protect Fabric data using sensitivity labels. Sensitivity labels can be set on all Fabric items, and data remains protected when it's exported through supported export paths. For more information, see Information Protection in Microsoft Fabric. *Customer outcome:* keep sensitive Fabric data classified and protected wherever it goes.
- **Microsoft Purview Data Loss Prevention (DLP)** — DLP policies support structured data in Fabric, such as lakehouses, warehouses, databases, and semantic models. Policies evaluate sensitivity labels and sensitive info types to limit user actions when data is classified as sensitive, and can generate policy tips and alerts. For more information, see data loss

prevention policies. *Customer outcome:* prevent the risky sharing or movement of sensitive Fabric data.

- **Microsoft Purview Audit** — all Microsoft Fabric user activities are logged and available in the Microsoft Purview audit log. For more information, see track user activities in Microsoft Fabric and track user activities in Power BI. *Customer outcome:* audit and investigate activity across Fabric for compliance and security.
- **Microsoft Purview Insider Risk Management (IRM)** — IRM policies support ready-to-use risk indicators for Microsoft Fabric, such as Power BI and lakehouse activities, to help detect potential data theft or leakage. For more information on supported Fabric workloads, see Configure policy indicators in Insider Risk Management. *Customer outcome:* identify and respond to insider risks involving Fabric data.
- **Microsoft Purview governance for Fabric Copilots and agents** — Purview provides governance and risk controls for Fabric Copilots and agents, including risk discovery in prompts and responses, audit coverage for AI interactions, and retention and eDiscovery applicability to AI-generated content. For more information, see Microsoft Purview data security and compliance protections for generative AI apps. *Customer outcome:* govern and monitor AI usage across supported Fabric workloads.

Security and compliance scenarios

Use the following scenarios to decide which experience to reach for. Fabric's built-in governance is the starting point, and Microsoft Purview extends protection, audit, risk, and compliance.

Discover and govern data

Start with Fabric's built-in governance experiences to find and govern your Fabric data:

- OneLake catalog overview
- Govern your Fabric data with the OneLake catalog
- Fabric governance documentation
- Governance and compliance in Microsoft Fabric

Protect sensitive data

Fabric supports sensitivity labels natively, and Microsoft Purview extends protection with information protection and data loss prevention:

- Information Protection in Microsoft Fabric
- Microsoft Purview data security and compliance protections for generative AI apps

Audit and investigate activity

Fabric user activity is available in the Microsoft Purview audit log for compliance and investigation:

- Track user activities in Microsoft Fabric
- Track user activities in Power BI

Related documentation

Start with Fabric governance

- Fabric governance documentation
- Governance and compliance in Microsoft Fabric
- OneLake catalog overview

Govern data

- Govern your Fabric data with the OneLake catalog

Microsoft Purview integrations

- Microsoft Purview data security and compliance protections for generative AI apps

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Information protection in Microsoft Fabric

Information protection in Fabric and Power BI enables organizations to categorize and protect their sensitive data using sensitivity labels and policies from Microsoft Purview Information Protection. In addition, Fabric and Power BI provide additional capabilities to help organizations achieve maximum sensitivity label coverage, and to manage and govern their sensitive data.

This article describes Fabric and Power BI's information protection capabilities. You can find details about current support in the considerations and limitations section.

Requirements

- An appropriate license for Microsoft Purview Information Protection.

  > **Note**
  >
  > If your organization uses Azure Information Protection sensitivity labels, they need to be migrated to the Microsoft Purview Information Protection unified labeling platform in order for them to be used in Fabric. Learn more about migrating sensitivity labels.

- To be able to apply labels to Power BI items, a user must have a Power BI Pro or Power BI Premium Per-User (PPU) license in addition to the licenses needed for Microsoft Purview Information Protection.

- Office apps have their own licensing requirements for viewing and applying sensitivity labels.

- Before enabling sensitivity labels on your tenant, make sure that sensitivity labels have been defined and published for relevant users and groups. See Create and configure sensitivity labels and their policies for detail.

- Customers in China must enable rights management for the tenant and add the Microsoft Purview Information Protection Sync Service service principal, as described in steps 1 and 2 under Configure Azure Information Protection for customers in China.

- Using sensitivity labels in Power BI Desktop requires the Desktop December 2020 release or later.

> **Note**
>
> If you try to open a protected .pbix file with a Desktop version earlier than December 2020, it will fail, and you'll be prompted to upgrade your Desktop version.

Access control

Sensitivity labels can apply access control to Fabric and Power BI data and content in the following cases:

- In the tenant where the sensitivity labels were applied. This scenario relies on sensitivity labels that are associated with Microsoft Purview *protection* policies. When a user logged in to the Fabric tenant where the labels were applied tries to access an item that has a label associated with a protection policy, their access is controlled by that protection policy. See Protection policies in Microsoft Fabric for more information.
- In Power BI Desktop (*.pbix*) files. This scenario relies on sensitivity labels that are associated with Microsoft Purview *publishing* policies. When a user tries to open a *.pbix* file that has a sensitivity label associated with a publishing policy, their access depends on the permissions they have under that policy. See Restrict access to content by using sensitivity labels to apply encryption.
- In supported export paths. This scenario relies on sensitivity labels that are associated with Microsoft Purview *publishing* policies. When a user tries to open a file generated via one of the supported export paths, their access depends on the permissions they have under that policy. See Restrict access to content by using sensitivity labels to apply encryption.

> **Important**
>
> Access control in all other scenarios is unsupported. This includes cross-tenant scenarios, such as **external data sharing**, where data is accessed from another tenant, or other export paths, such as export to *.csv* files or *.txt* files.

Supported export paths

Sensitivity labels and the access control they apply (if they are associated with a Microsoft Purview publishing policy) stays with data and content that leaves Fabric and Power BI in the following export paths:

- Export to Excel, PDF files, and PowerPoint.
- Analyze in Excel from Fabric, which triggers download of an Excel file with a live connection to a Power BI semantic model.
- PivotTable in Excel with a live connection to a Power BI semantic model, for users with Microsoft 365 E3 and above.
- Download to a Power BI Desktop (*.pbix*) file from Fabric.

Capabilities

The following table summarizes the information protection capabilities in Fabric that help you achieve maximum coverage of the sensitive information in your organization. Fabric support is indicated in the third column. See the sections under Considerations and limitations for details.

Expand table

Capability	Scenario	Support status
Manual labeling	Users can manually apply sensitive labels to Fabric items	Supported for all Fabric items.
Default labeling	When an item is created or edited, it gets a default sensitivity label unless a label is applied through other means.	Supported for all Fabric items, with limitations.
Mandatory labeling	Users can't save items unless a sensitivity label is applied to the item. This means they can't remove a label either.	Currently fully supported for Power BI items only. Supported for some non-Power BI Fabric items, with limitations.
Programmatic labeling	Sensitivity labels can be added, changed, or deleted programmatically via Power BI admin REST APIs.	Supported for all Fabric items.
Downstream inheritance	When a sensitivity label is applied to an item, the label propagates downstream to all dependent items.	Supported for all Fabric items, with limitations.
Inheritance upon creation	When you create a new item from an existing item, the new item inherits the label of the existing item.	Supported for all Power BI Fabric items. Supported for some non-Power BI Fabric items as described in the considerations and limitations.

Capability	Scenario	Support status
Inheritance from data sources	When a Fabric item ingests data from a data source that has a sensitivity label, that label is applied to the Fabric item. The label then propagates downstream to the child items of that Fabric item via downstream inheritance.	Currently supported for Power BI semantic models only.
Inheritance upon update and relationship changes	When you modify an existing item, connect it to a new data source, or link it to other items through new lineage relationships, its sensitivity label is reevaluated and updated if necessary to ensure consistent protection as the item evolves.	Supported for all Fabric items, with limitations.
Export	When a user exports data from an item that has a sensitivity label, the sensitivity label moves with it to the exported format.	Currently supported for Power BI items in supported export paths.

Monitor and govern sensitivity label adoption

Beyond applying sensitivity labels, Fabric provides governance experiences that help administrators and data owners understand how labels are being used across their data estate. Using governance experiences in OneLake catalog, organizations can:

- Identify labeled and unlabeled assets.
- Understand sensitivity label coverage across Fabric items.
- Discover assets associated with specific sensitivity labels.
- Monitor adoption of information protection practices.
- Prioritize remediation of unlabeled or incorrectly classified content.

These insights help organizations improve information protection posture and achieve broader label coverage across their Fabric environment.

Considerations and limitations

Manual labeling

When you enable sensitivity labels on your tenant, you specify which users can apply sensitivity labels. While the other information protection capabilities described in this article can ensure that most items get labeled without someone having to manually apply a label, manual labeling makes it possible for users to change labels on items. For more information about how to manually apply sensitivity labels to Fabric items, see How to apply sensitivity labels.

> **Note**
>
> For a user to be able to apply sensitivity labels to Fabric items, it's not enough just to include the user in the list of specified users. The sensitivity label must also be published to the user as part of the label's policy definitions in the Microsoft Purview compliance portal. For more information, see **Create and configure sensitivity labels and their policies**.

Default labeling

Default labeling is fully supported in Power BI and is described in Default label policy for Power BI. In Fabric, there are some limitations.

- When a non-Power BI Fabric item is created, if there's a clear, substantive create dialog, the default sensitivity label will be applied to the item if the user doesn't choose a label. If the item is created in a process where there's no clear create dialog, the default label **won't** be applied.
- When a Fabric item that has no label is updated, if the item is a Power BI item, a change to any of its attributes causes the default label to be applied if the user doesn't apply a label. If the item is a non-Power BI Fabric item, only changes to certain attributes, such as name and description, cause the default label to be applied. And this is only if the change is made in the item's flyout menu. For changes made in the experience interface, default labeling isn't currently supported.

Mandatory labeling

Mandatory labeling is currently supported for Power BI items only. Mandatory labeling isn’t enforced if changes are made via the flyout menu.

For lakehouses, pipelines, and data warehouses: Assuming that information protection is enabled, if mandatory labeling is on and default labeling is off, it will be possible for the user to select a label. However, mandatory labeling logic isn't enforced. That means that the user can save the item without a label, unless the experience itself requires that a label be set.

For more information about mandatory labeling, see Mandatory label policy for Fabric and Power BI.

Programmatic labeling

Programmatic labeling is supported for all Fabric items. For more information, see Set or remove sensitivity labels using Power BI REST admin APIs.

Downstream inheritance

Downstream inheritance is on by default. It's supported in Fabric as follows:

Supported:

- Power BI item to Power BI item
- Fabric item to Fabric item
- Fabric item to Power BI item

Not supported:

- Power BI item to Fabric item

Autogenerated items from a lakehouse or data warehouse take their sensitivity label from their parent lakehouse or data warehouse. They don't inherit the label from items further upstream.

For more information about downstream inheritance, see Sensitivity label downstream inheritance.

Inheritance upon creation

Inheritance upon creation is supported for Power BI Fabric items and in other scenarios with non-Power BI items where one item is created from another item:

- A Pipeline created from a Lakehouse inherits the sensitivity label of the Lakehouse.
- A Notebook created from a Lakehouse inherits the sensitivity label of the Lakehouse.
- A Lakehouse shortcut created from a Lakehouse inherits the sensitivity label of the Lakehouse.
- A Pipeline created from a Notebook inherits the sensitivity label of the Notebook.
- A KQL Queryset created from a KQL Database inherits the sensitivity label of KQL Database.
- A Pipeline created from a KQL Database inherits the sensitivity label of KQL Database.

For more information about downstream inheritance, see Sensitivity label inheritance upon creation of new content.

Inheritance from data sources

Inheritance from data sources is currently supported for Power BI semantic models only. For more information, see Sensitivity label inheritance from data sources.

Sensitivity label inheritance upon update

When you create a new Fabric item from a labeled item, the new item inherits the parent's sensitivity label at the moment of creation. But what happens when the relationship between items changes *after* creation? Label inheritance upon update ensures that sensitivity labels continue to propagate in the following scenarios:

- **Item modification**: You update an existing item in a way that changes its data content or connections.
- **New data source connections**: You connect an existing item to a new, labeled data source.
- **New lineage relationships**: You establish a new lineage relationship for an existing item, such as when you connect a notebook to a lakehouse through the user interface (UI).

In each of these scenarios, Fabric evaluates the sensitivity labels of the related items and applies the appropriate label to ensure consistent data protection. For more information, see Sensitivity label inheritance upon update and relationship changes.

Export

Sensitivity label inheritance upon export is supported for Power BI items only in supported export paths. Currently no other Fabric experience uses an export method that transfers the sensitivity label to the exported output. However, if they do export an item that has a sensitivity label, a warning is issued.

To see the supported export paths for Power BI items, see Supported export paths in Power BI.

Related content

- Sensitivity labels in Power BI

Last updated on 07/16/2026

Protected sensitivity labels in Fabric and Power BI

Protected labels are sensitivity labels that have file protection policies associated with them and can be used to protect files and data. A file protection policy for a label grants usage rights to users, such as the ability to open a file, edit a file, copy from a file, etc. When a protected label is applied to a file or item, users who are included in the policy can perform the actions they have the usage rights for under the label policy. A file protection policy can grant different sets of users different sets of usage rights. For example, under the policy, one set of users might be granted full control over the file, while another set of users might only be permitted to open and view the file.

Having a set of sensitivity labels and policies in place is a prerequisite for using sensitivity labels in Fabric and Power BI. The labels and their file protection policies are defined by security admins in the Microsoft Purview compliance portal. Typically, the same set of protected labels is used across an organization for Office apps such as Word, Excel, and PowerPoint as well as for Fabric and Power BI.

How protected labels work in Fabric and Power BI

In Fabric and the Power BI service, protected labels only control the ability to change or remove labels on items. They don't control access to content; for access control via sensitivity label, see Protection policies. In order for a user to be able to change or remove a protected label from an item, the user must either be the user who applied the sensitivity label (the RMS owner), or have at least one of the following usage rights requirements for the label:

- OWNER
- EXPORT
- EDIT and EDITRIGHTSDATA

If the label on the item was set via an automated process, such as inheritance from data sources or downstream inheritance, then the third option, EDIT and EDITRIGHTSDATA, is reduced to just EDIT. See Relaxations to accommodate automatic labeling scenarios for details.

In Power BI Desktop, protected labels control not only the ability to change or remove the protected label, but also encrypt the file's content (viewing, editing, exporting, etc.). Having the previously described usage rights is sufficient to decrypt the file and have full control over it.

Users with the following permissions can open, edit, and republish protected PBIX files with restrictions.

Permissions:

- View content (VIEW)
- Edit content (DOCEDIT)
- Save (EDIT)
- Copy and extract content (EXTRACT)
- Allow macros (OBJMODEL)

Restrictions:

- They can't export to formats that don't support sensitivity labels, such as CSV files.
- They can't save as PBIP or PBIT.
- They can't change the label on the PBIX file.
- They can only republish the PBIX file to the original workspace the file came from.

> **Note**
>
> The file must have been published at least once for other users to be able to publish it back to that specific workspace. If the file hasn't yet been published, then the latest label issuer (the one who most recently set the protected label) or a user with sufficient usage rights must publish it and then share the file with the other editors.

These restrictions ensure that the security of the content remains under the control of those who have high enough permissions to set the label.

These usage rights are a subset of the **editor** (formerly known as co-author) permissions preset in the Microsoft Purview compliance portal.

Relaxations to accommodate automatic labeling scenarios

Fabric and Power BI support several capabilities, such as label inheritance from data sources and downstream inheritance, which automatically apply sensitivity labels to content. These automated scenarios can result in situations where no user has been set as the RMS label issuer for a label on an item. This means that there's no user who is guaranteed to be able to change or remove the label.

In such cases, the usage rights requirements for changing or removing the label are relaxed - a user needs just one of the following usage rights to be able to change or remove the label:

- OWNER
- EXPORT
- EDIT

If no user has even these usage rights, nobody can change or remove the label from the item, and access to the item is potentially endangered.

To avoid this situation, the Fabric admin can enable the **Allow workspace admins to override automatically applied sensitivity labels** tenant setting. This makes it possible for workspace admins to override automatically applied sensitivity labels without regard to label change enforcement rules.

To enable this setting, go to **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings** > **Information protection**, and enable the toggle on the **Allow workspace admins to override automatically applied sensitivity labels** setting.

Related content

- Information protection
- Set or remove sensitivity labels using Power BI REST admin APIs

---

Last updated on 06/11/2025

Default label policy for Fabric and Power BI

To help ensure comprehensive protection and governance of sensitive data, organizations can create default label policies for Fabric and Power BI that automatically apply default sensitivity labels to unlabeled items.

This article describes how to enable a default label policy, both in the Microsoft Purview compliance portal and by using the Security & Compliance PowerShell setLabelPolicy API.

> **Note**
>
> The default label policy settings for Fabric and Power BI are independent of the default label policy settings for files and email.

What happens when a default label policy is in effect?

- **In Power BI Desktop**

  When a user to whom the policy applies opens a new .pbix file or an existing unlabeled .pbix file, the default label is applied to the file. If the user is working offline, the label is applied when the user signs in.

- **In Fabric**

  When a user to whom the policy applies creates a new item, the default label is applied to that item.

  When a user to whom the policy updates an unlabeled item:

  - If the item is a Power BI item, a change to any of its attributes causes the default label to be applied if the user doesn't apply a label.
  - If the item is a non-Power BI Fabric item, only changes to certain attributes, such as name and description, cause the default label to be applied. And this is only if the change is made in the item's flyout menu. For changes made in the experience interface, default labeling isn't currently supported.

Enabling a default label policy for Fabric and Power BI

A Microsoft 365 administrator can enable a default label policy for Fabric and Power BI by selecting the desired label in the **Apply this default label to Power BI content** dropdown menu in the policy settings for Power BI in the Microsoft Purview compliance portal. For more information, see What label policies can do.

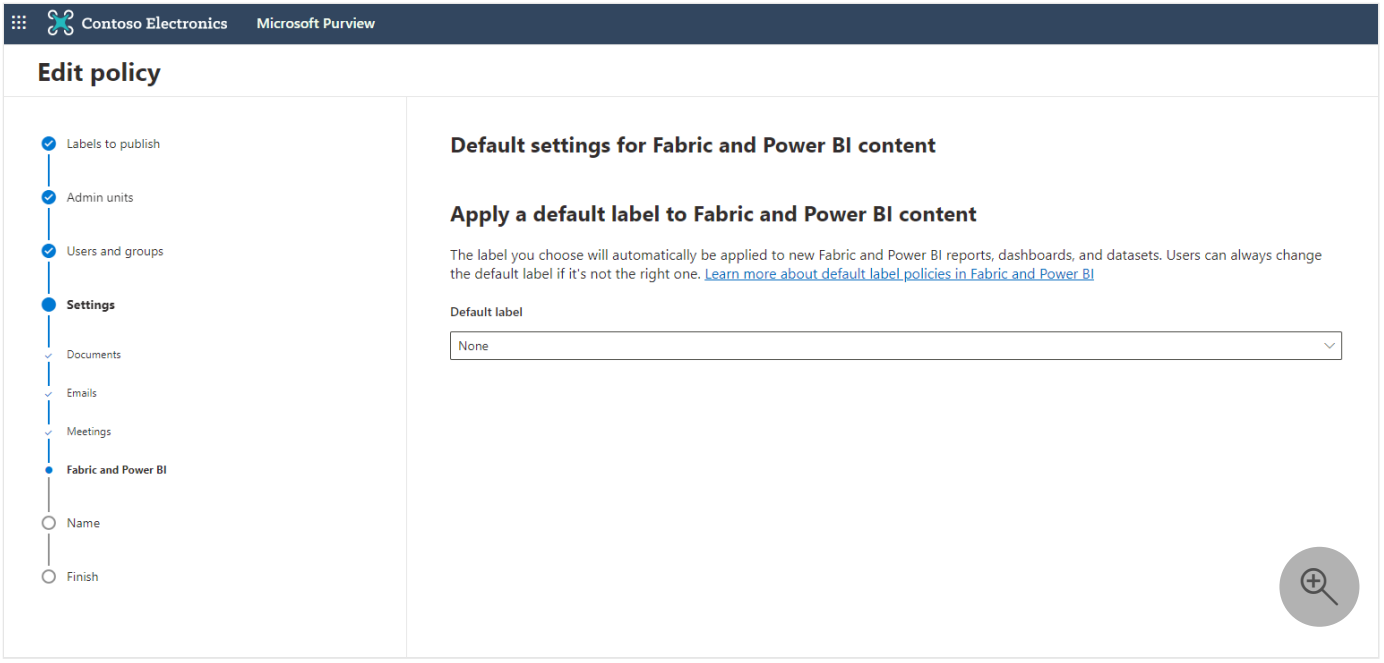

For existing policies, it's also possible to enable default label policies for Fabric and Power BI using the Security & Compliance PowerShell setLabelPolicy API.

PowerShell

```
Set-LabelPolicy -Identity "<default label policy name>" -AdvancedSettings
@{powerbidefaultlabelid="<LabelId>"}
```

Where:

- `<default label policy name>` is the name of the policy whose associated sensitivity label you want to be applied by default to unlabeled items in Fabric and Power BI.

> **Important**
>
> If a user has more than one label policy, the default label setting is always taken from the policy with the highest priority, make sure to configure the default label on that policy.

Requirements for using PowerShell

- The Exchange Online PowerShell V2 (EXO V2) module. For more information, see About the Exchange Online PowerShell V2 module

- A connection to the Microsoft Purview compliance portal. For more information, see Connect to Security & Compliance PowerShell

Documentation

- Admin Guide: Custom configurations for the Azure Information Protection unified labeling client
- Create and configure sensitivity labels and their policies
- Set-LabelPolicy documentation

Considerations and limitations

- Default labeling in Fabric Power BI covers most common scenarios, but there might be some less common flows that still allow users to open or create unlabeled *.pbix* files or Fabric items.
- Default label policy settings for Fabric and Power BI are independent of the default label policy settings for files and email.
- Default labeling in Fabric and Power BI isn't supported for service principals and APIs. Service principals and APIs aren't subject to default label policies.
- Default label policies in Fabric and Power BI aren't supported for external guest users (Microsoft Entra B2B). When a B2B user opens or creates an unlabeled *.pbix* file in Power BI Desktop or Fabric or Power BI item in Fabric, no default label is applied automatically.

Related content

- Mandatory label policy for Power BI

Last updated on 03/03/2025

Domain-level default sensitivity labels in Microsoft Fabric

This article describes the domain-level default sensitivity label feature in Microsoft Fabric and how to enable it. Its target audience is Fabric administrators, security and compliance officers, Fabric domain administrators, Fabric domain contributors.

Domain admins can set a default sensitivity label for their domains. This label will be applied by default to items in workspaces that belong to the domain in the following instances:

- When a new item is created and saved in the workspace.
- When an unlabeled item that already exists in the workspace is updated and saved.

When an existing labeled item in the workspace is updated and saved, the existing label is either retained or overwritten according to the following logic:

Expand table

Existing label	Override with domain-level default label
Manually applied, any priority	No
Automatically applied, lower priority	Yes
Automatically applied, higher priority	No
Default label from policy, lower priority	Yes
Default label from policy, higher priority	No

Requirements

The tenant setting **Domain admins can set default sensitivity labels for their domains (preview)** must be enabled.

Limitations

Deployment pipelines and Git integration are currently not supported.

Related content

- Fabric domains
- Sensitivity labels in Fabric and Power BI

Last updated on 11/10/2025

Mandatory label policy for Fabric and Power BI

To help ensure comprehensive protection and governance of sensitive data, you can require your organization's Fabric and Power BI users to apply sensitivity labels to items they create or edit. You do this by enabling, in their sensitivity label policies, a special setting for mandatory labeling in Fabric and Power BI. This article describes the user actions that are affected by a mandatory labeling policy, and explains how to enable a mandatory labeling policy for Fabric and Power BI.

> **Note**
>
> The mandatory label policy setting for Fabric and Power BI is independent of the mandatory label policy setting for files and email.
>
> Mandatory labeling in Fabric and Power BI isn't supported for service principals and APIs. Service principals and APIs aren't subject to mandatory label policies.

What happens when a mandatory label policy is in effect?

In Fabric and the Power BI service:

- Users must apply a sensitivity label before they can save new items (if the item is a supported item type).
- Users must apply a sensitivity label before they can save changes to the settings or content of existing, unlabeled items (if the item is a supported item type).
- If users try to import data from an unlabeled *.pbix* file, a prompt requires them to select a label before the import can continue. The label they select is applied to the resulting semantic model and report in the service. **It's not applied to the *.pbix* file itself**.

In Power BI Desktop:

- Users must apply sensitivity labels to unlabeled *.pbix* files before they can save or publish to the service.

Supported item types

Mandatory labeling in Fabric and Power BI is supported for all item types except:

- Scorecard
- Dataflow Gen 1
- Dataflow Gen 2
- Streaming semantic model
- Streaming dataflow

Enabling a mandatory label policy for Fabric and Power BI

Users with appropriate permissions can enable a mandatory label policy for Fabric and Power BI by selecting the **Require users to apply a label to their Fabric and Power BI content** checkbox at the **Default settings for Fabric and Power BI content** stage of the policy configuration process in the Microsoft Purview portal. For more information, see What label policies can do.

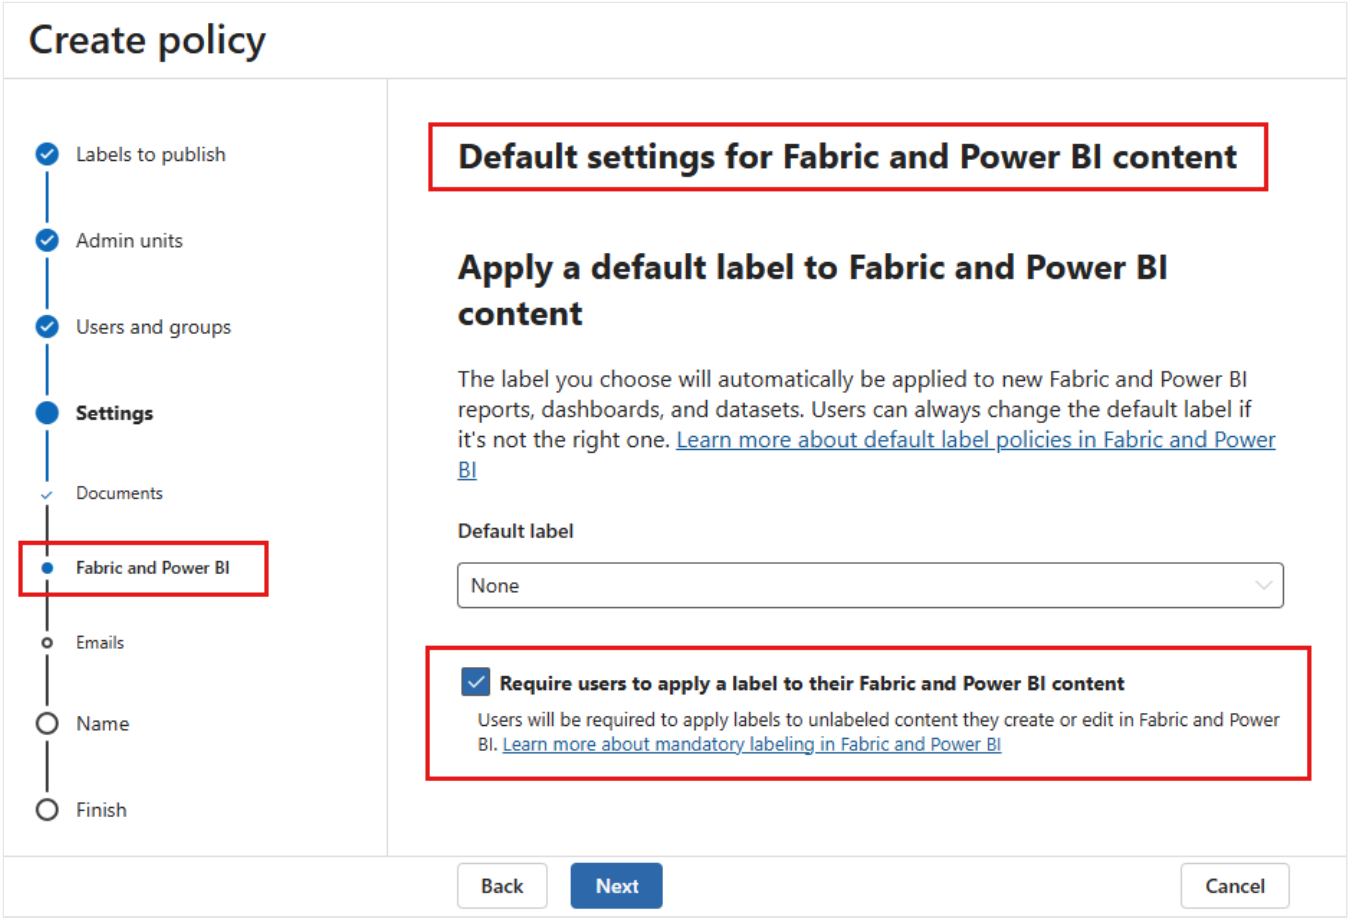

If you have an existing policy, and you want to enable mandatory labeling in Power BI, you can use the Security & Compliance PowerShell setLabelPolicy API.

PowerShell

```
Set-LabelPolicy -Identity "<policy name>" -AdvancedSettings
@{powerbimandatory="true"}
```

Where:

- `<policy name>` is the name of the policy where you want to set labeling in Fabric and Power BI as mandatory.

Requirements for using PowerShell

- You need the Exchange Online PowerShell (EXO) module to run this command. For more information, see About the Exchange Online PowerShell module.
- A connection to the Microsoft Purview portal is also required. For more information, see Connect to Security & Compliance PowerShell using the EXO module.

Documentation

- Admin Guide: Advanced settings for Microsoft Purview Information Protection client
- Use PowerShell for sensitivity labels and their policies
- Set-LabelPolicy documentation

Considerations and limitations

- Mandatory labeling in Fabric and Power BI covers most common scenarios, but there might be some less common flows that still allow a user to create or edit unlabeled content.
- The mandatory label policy setting for Fabric and Power BI is independent of the mandatory label policy setting for files and email.
- Mandatory labeling in Fabric and Power BI isn't supported for service principals and APIs. Service principals and APIs aren't subject to mandatory label policies.
- Mandatory labeling in Fabric and Power BI isn't supported for external guest users (B2B users). B2B users aren't subject to mandatory label policies.

Related content

- Default label policy for Fabric and Power BI

Last updated on 12/02/2025

Sensitivity label support for paginated reports

Sensitivity labels can be applied to paginated reports hosted in the Power BI service. After uploading a paginated report to the service, apply the label to the report just as you would to a regular Power BI report.

When you export data from a labeled paginated report to a supported file type (Excel, PDF, PPTX, and Word), the sensitivity label on the paginated report is applied to the exported file.

Sensitivity labels on paginated reports are included in the OneLake catalog insights, and can be audited (label-change audits only) and modified by public APIs, just like labels on regular Power BI reports.

Considerations and limitations

- Downstream inheritance isn't supported. The label of an upstream model won't propagate down to its downstream paginated reports. Likewise, the label of a paginated report won't propagate down to the report's downstream content.
- Mandatory labeling doesn't apply to paginated reports.

Paginated Report visuals

A Paginated Report visual is a special type of visual that you can include in a regular Power BI report. It renders a selected paginated report inside the regular Power BI report.

When a supported file type is exported from a Paginated Report visual that is included in a Power BI report, and the original paginated report being rendered in the visual has a sensitivity label, the exported file inherits the sensitivity label of the original paginated report. If the original paginated report doesn't have a label, the exported file inherits the label of the Power BI report, if it has one.

Related content

- Apply sensitivity labels in Power BI
- Sensitivity label overview

Last updated on 05/11/2026

Set or remove sensitivity labels using Power BI REST admin APIs

To meet compliance requirements, organizations are often required to classify and label all sensitive data in Power BI. This task can be challenging for tenants that have large volumes of data in Power BI. To make the task easier and more effective, you can use Power BI admin REST APIs to programmatically setLabels and removeLabels on large numbers of Power BI artifacts.

The APIs set or remove labels from artifacts by artifact ID.

Requirements and considerations

- Users must be Fabric administrators to call these APIs.
- The admin user (and the delegated user, if provided) must have sufficient usage rights to set or remove labels.
- To set a sensitivity label using the setLabels API, the admin user (or the delegated user, if provided) must have the label included in their label policy.
- The APIs allow a maximum of 25 requests per hour. Each request can update up to 2000 artifacts.
- **Required scope**: Tenant.ReadWrite.All

API documentation

- setLabels
- removeLabels

Sample

The following sample demonstrates how to set and remove sensitivity labels on Power BI dashboards. Similar code can be used to set and remove labels on semantic models, reports, and dataflows.

rest

```
const string adminBearerToken = "<adminBearerToken>";
const string ApiUrl = "<api url>";
var persistedDashboardId = Guid.Parse("<dashboard object Id>");
var credentials = new TokenCredentials(adminBearerToken, "Bearer");
```

```
            var artifacts = new InformationProtectionArtifactsChangeLabel();
            artifacts.Dashboards = new List<ArtifactId> { new ArtifactId(id:
persistedDashboardId) };

            using (PowerBIClient client = new PowerBIClient(credentials))
            {
                client.BaseUri = new Uri(ApiUrl);

                // Delete labels

               var removeResponse =
client.InformationProtection.RemoveLabelsAsAdmin(artifacts);

                foreach (var updateLabelResult in removeResponse.Dashboards)
                {
                    if (updateLabelResult.Status == Status.Succeeded)
                    {
                        Console.WriteLine($"label has been deleted from artifact
{updateLabelResult.Id}");
                    }
                    else
                    {
                        Console.WriteLine($"label has not been deleted from artifact
{updateLabelResult.Id}");
                    }
                }

                // Set labels

                var setLabelRequest = new InformationProtectionChangeLabelDetails();
                setLabelRequest.Artifacts = artifacts;
                setLabelRequest.LabelId = Guid.Parse("<label Id>");

                // assignmentMethod (optional)
                setLabelRequest.AssignmentMethod = AssignmentMethod.Privileged;

                // delegatedUser (optional)
                var delegatedUser = new DelegatedUser();
                delegatedUser.EmailAddress = "<delegated user email address>";

                setLabelRequest.DelegatedUser = delegatedUser;

                var setResponse =
client.InformationProtection.SetLabelsAsAdmin(setLabelRequest);
                foreach (var updateLabelResult in setResponse.Dashboards)
                {
                    if (updateLabelResult.Status == Status.Succeeded)
                    {
                        Console.WriteLine($"label has been upsert on artifact
{updateLabelResult.Id}");
                    }
                    else
                    {
                        Console.WriteLine($"label has not been upsert on artifact
```

```
{updateLabelResult.Id}");
                    }
                }

            }

        }
```

Related content

- setLabels API
- removeLabels API

Last updated on 08/08/2024

Sensitivity label downstream inheritance

When a sensitivity label is applied to a semantic model or report in the Power BI service, it's possible to have the label trickle down and be applied to content that's built from that semantic model or report. For semantic models, this means other semantic models, reports, and dashboards. For reports, this means dashboards. This capability is called downstream inheritance.

Downstream inheritance is a critical link in Power BI's end-to-end information protection solution. Together with inheritance from data sources, inheritance upon update and relationship changes, inheritance upon creation of new content, inheritance upon export to file, and other capabilities for applying sensitivity labels, downstream inheritance helps ensure that sensitive data remains protected throughout its journey through Power BI, from data source to point of consumption.

Downstream inheritance is illustrated using lineage view. When a label is applied to the semantic model *Customer profitability*, that label filters down and gets applied to the semantic model's downstream content: the reports that are built using that semantic model, and, in this case, a dashboard that's built from visuals from one of those reports.

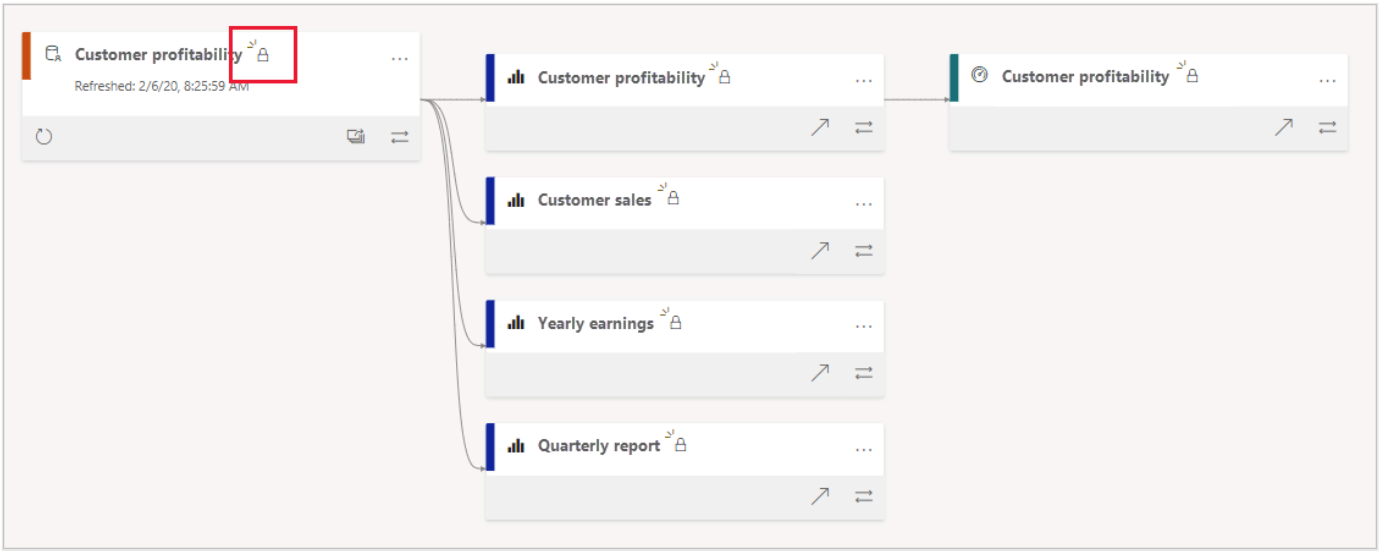

Important

- Downstream inheritance never overwrites labels that were applied manually.
- Downstream inheritance never overwrites a label with a less restrictive label.

Downstream inheritance vs. inheritance upon creation

It's important to understand the difference between downstream inheritance and inheritance upon item creation:

- **Downstream inheritance** applies labels to *existing* downstream items when you change a label on a parent item (semantic model or report). This feature is what this article describes.
- **Inheritance upon creation** automatically applies a label to a *new* item when you create it based on a labeled parent item. This happens regardless of the downstream inheritance settings.

For example, when you create a new report based on a labeled semantic model, the report automatically inherits the semantic model's label to prevent creating unlabeled artifacts that could expose data. This inheritance upon creation occurs even if the **Automatically apply sensitivity labels to downstream content** tenant setting is disabled.

The **Automatically apply sensitivity labels to downstream content** setting only controls whether labels propagate to existing downstream items when you change a parent item's label. It doesn't affect label inheritance when you create new items.

For more information about inheritance upon creation, see Sensitivity label inheritance upon creation of new content.

Downstream inheritance modes

Downstream inheritance operates in one of two modes. The Power BI admin decides via a tenant setting which mode is operable on the tenant.

- **Downstream inheritance with user consent** (default): In this mode, when users apply sensitivity labels on semantic models or reports, they can choose whether to also apply that label downstream. They make their choice by selecting the box that appears with the sensitivity label selector.
- **Fully automated downstream inheritance** (when enabled by a Power BI admin): In this mode, downstream inheritance happens automatically whenever a label is applied to a semantic model or report. There's no checkbox provided for user consent.

The two downstream inheritance modes are explained in more detail in the following sections.

Downstream inheritance with user consent

In user consent mode, when a user applies a sensitivity label to a semantic model or report, they can choose whether to apply the label to its downstream content as well. A checkbox appears

along with the label selector:

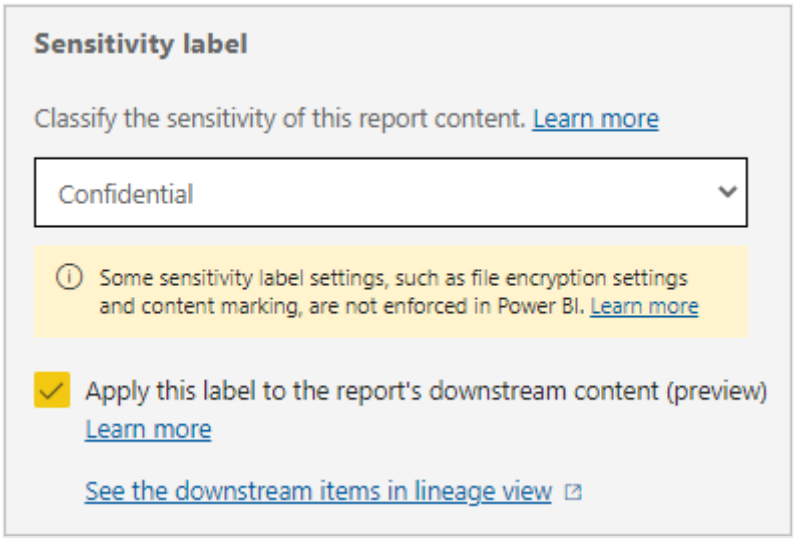

By default, the checkbox is selected. This means that when the user applies a sensitivity label to a semantic model or report, the label propagates down to its downstream content. For each downstream item, the label is applied only if:

- The user who applied or changed the label has Power BI edit permissions on the downstream item (that is, the user is an admin, member, or contributor in the workspace where the downstream item is located).
- The user who applied or changed the label is authorized to change the sensitivity label that already exists on the downstream item.

Clearing the checkbox prevents the label from being inherited downstream.

Fully automated downstream inheritance

In fully automated mode, a label applied to either a semantic model or report is automatically propagated and applied to the semantic model or report's downstream content, without regard to edit permissions on the downstream item and the usage rights on the label.

Relaxed label change enforcement

In certain cases, downstream inheritance (like other automated labeling scenarios) can result in a situation where no user has all the required permissions needed to change a label. For such situations, label change enforcement relaxations are in place to guarantee access to affected items. See Relaxations to accommodate automatic labeling scenarios for detail.

Enabling fully automated downstream inheritance

The tenant setting **Automatically apply sensitivity labels to downstream content** controls fully automated downstream inheritance. This setting is enabled by default when information protection is enabled (that is, when the tenant setting **Allow users to apply sensitivity labels for content** is enabled). To change the downstream inheritance setting, go to the tenant settings and enable or disable the **Automatically apply sensitivity labels to downstream content** setting as desired.

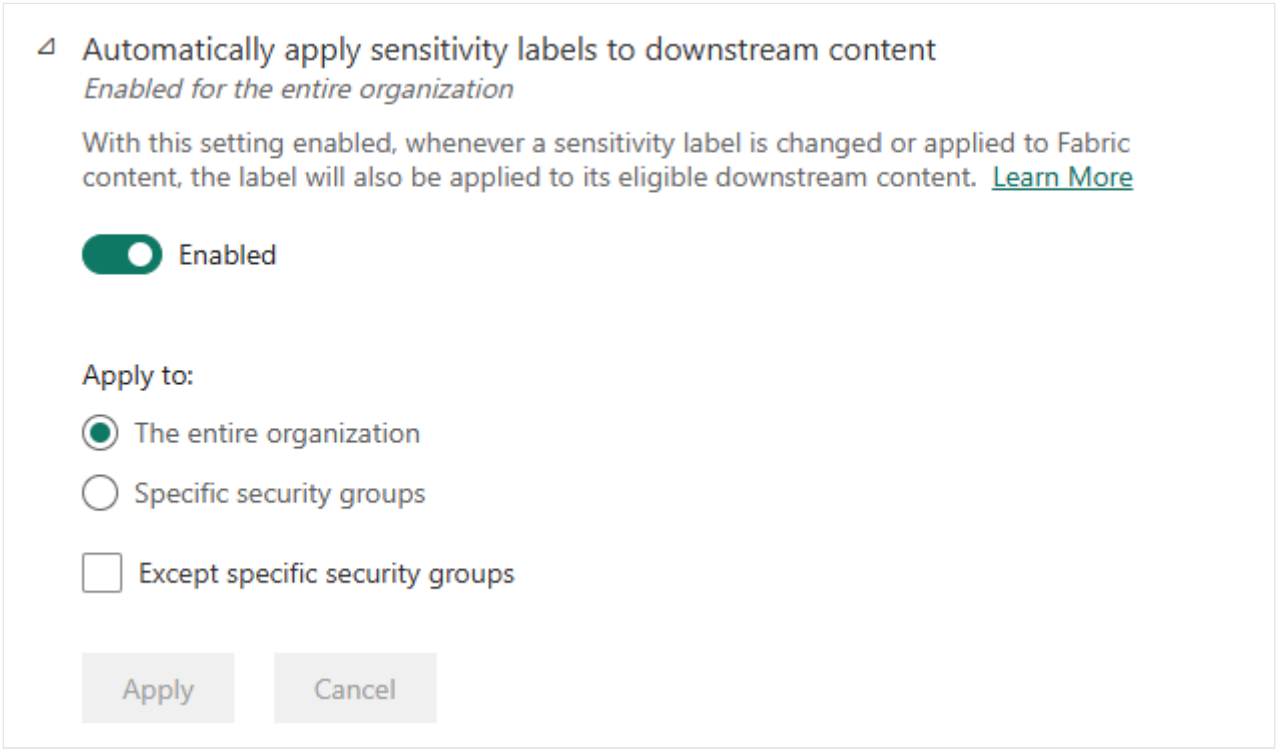

Considerations and limitations

- Downstream inheritance is limited to 80 items. If the number of downstream items exceeds 80, no downstream inheritance takes place. Only the item the label was actually applied to receives the label.
- Downstream inheritance never overwrites manually applied labels. See the following section for a significant consideration.
- Downstream inheritance never replaces a label on downstream content with a label that's less restrictive than the currently applied label.
- Sensitivity labels inherited from data sources are automatically propagated downstream only when fully automated downstream inheritance mode is enabled.

Downstream inheritance between semantic models and reports published from *.pbix* files

When you publish a *.pbix* file that has a sensitivity label, the label that's inherited by the semantic model and report created in the service is considered to be automatically or manually applied depending on whether the label on the *.pbix* file was automatically or manually applied. This has implications for subsequent downstream inheritance from the semantic model to its associated report. If the label on the *.pbix* file was automatically applied, then later, after publishing, if the label on the semantic model is changed, the associated report inherits the change. If, however, the label on the *.pbix* file was manually applied, then if the label on the semantic model is changed, the label on its associated report **is not** overwritten, as it's considered to be manually applied. This is in keeping with the rule that downstream inheritance never overwrites a manually applied label.

Related content

- Sensitivity label inheritance upon update and relationship changes
- Sensitivity label inheritance from data sources
- Sensitivity label overview
- Label change enforcement

Last updated on 07/13/2026

Sensitivity label inheritance from data sources

Power BI semantic models that connect to sensitivity-labeled data in supported data sources can inherit those labels, so that the data remains classified and secure when brought into Power BI.

Currently supported data sources:

- Excel files stored on OneDrive or SharePoint Online*
- Azure Synapse Analytics (formerly SQL Data Warehouse)
- Azure SQL Database

*Inheritance from Excel files requires specific configuration and isn't supported for Excel files stored behind a gateway, such as files stored locally. See Sensitivity label inheritance from Excel files for more detail.

To be operative, sensitivity label inheritance from data sources must be enabled on the tenant.

Requirements

- The data in the data source must be labeled with sensitivity labels from Microsoft Purview Information Protection.

  For Azure Synapse Analytics and Azure SQL Database, this is accomplished using a two-step Purview flow:

  1. Automatically apply sensitivity labels to your data.
  2. Classify your Azure SQL data using Azure Purview labels.

  - The scope of the labels must be **Files and emails** and **Azure Purview assets**. See Extending sensitivity labels to Azure Purview and Creating new sensitivity labels or modifying existing labels.

- Sensitivity labels must be enabled in Power BI.

- A Power BI Pro or Premium Per-User (PPU) license is required for the user whose credentials are used to connect to the data source. In addition, all other conditions for applying a label must be met.

- The **Apply sensitivity labels from data sources to their data in Power BI** tenant admin setting must be enabled. **Note**: This requirement applies to the Power BI service only. In Desktop, a *.pbix* file will inherit the label from the data source even if the tenant admin setting is off. However, after publishing to the service, upon refresh, changes to the label in the data source will only be inherited by the report and semantic model if the setting is enabled.

Inheritance behavior

- In the Power BI service, when the semantic model is connected to the data source, Power BI inherits the label and applies it automatically to the semantic model. Subsequently, inheritance occurs upon semantic model refresh. In Power BI Desktop, when you connect to the data source via **Get data**, Power BI inherits the label and automatically applies it to the *.pbix* file (both the semantic model and report). Subsequently inheritance occurs upon refresh.
- If the data source has sensitivity labels of different degrees, the most restrictive is chosen for inheritance. In order to be applied, that label (the most restrictive) must be published for the semantic model owner.
- Labels from data sources never overwrite manually applied labels.
- Less restrictive labels from the data source never overwrite more restrictive labels on the semantic model.
- In Desktop, if the incoming label is more restrictive than the label that is currently applied in Desktop, a banner will appear that recommends to the user to apply the more restrictive label.
- Semantic model refresh will succeed even if for some reason the label from the data source isn't applied.

> **Note**
>
> No inheritance takes place if the semantic model owner is not authorized to apply sensitivity labels in Power BI, or if the specific label in question has not been published for the semantic model owner.

Sensitivity label inheritance from Excel files

Sensitivity label inheritance from an Excel file is supported for Excel files stored on OneDrive or SharePoint Online.

To make sure sensitivity label inheritance from an Excel file works:

1. Store the Excel file on OneDrive or SharePoint Online.

2. In Power BI Desktop, connect to the Excel file using the web connector, as described in Use OneDrive for work or school links in Power BI Desktop. The process described in that article applies to both OneDrive and SharePoint Online.

3. After publishing the semantic model, to enable refresh, reconfigure the authentication credentials for the semantic model, also as described in the above article. Be sure to select **OAuth2** as the authentication method, otherwise you might encounter an error when you attempt to connect or refresh.

Considerations and limitations

- Inheritance from data sources is supported only for semantic models with enhanced metadata. See Using enhanced semantic model metadata for more information.
- Inheritance from data sources is supported only for semantic models using the Import data connectivity mode. Live connection and DirectQuery connectivity isn't supported.
- Inheritance from data sources isn't supported in connections via gateways or Azure Virtual Network (VNet). This means that inheritance from an Excel file located on a local machine won't work, because this requires a gateway.

Related content

- Sensitivity label inheritance upon update and relationship changes
- Sensitivity label downstream inheritance
- Enable sensitivity label inheritance from data sources
- Sensitivity label overview

Last updated on 07/13/2026

Sensitivity label inheritance upon update and relationship changes in Microsoft Fabric

When you modify existing items, connect them to new data sources, or establish new lineage relationships, Microsoft Fabric automatically reevaluates and updates their sensitivity labels to ensure consistent data protection across your environment.

This article explains when and how sensitivity labels are automatically updated as items change, so that you can ensure consistent protection across your Microsoft Fabric environment.

How sensitivity label inheritance works for existing Microsoft Fabric items

When you create a new Microsoft Fabric item from a labeled item, the new item inherits the parent's sensitivity label at the moment of creation. If the relationship between items changes *after* creation, sensitivity labels continue to propagate in the following scenarios:

- **Item modification**: You update an existing item in a way that changes its data content or connections.
- **New data source connections**: You connect an existing item to a new, labeled data source.
- **New lineage relationships**: You establish a new lineage relationship for an existing item, such as when you connect a notebook to a lakehouse through the user interface (UI).

In each of these scenarios, Microsoft Fabric evaluates the sensitivity labels of the related items and applies the appropriate label to ensure consistent data protection.

> **Important**
>
> Label inheritance upon update follows the same core rules as other automatic labeling capabilities in Microsoft Fabric:
>
> - Automatically applied labels never overwrite labels that you manually applied.
> - A less restrictive label never overwrites a more restrictive label that is already applied.

Sensitivity label behavior when you modify Fabric items

When you modify an existing Fabric item, the system reevaluates the item's sensitivity label based on its current relationships and any applicable label policies.

The following behaviors apply:

- **Default label application**: If you edit an unlabeled item and a default label policy is active, Fabric applies the default label. For Power BI items, any change to the item's attributes triggers default labeling. For non-Power BI Fabric items, only changes to certain attributes, such as name and description, trigger default labeling.
- **Downstream propagation**: When you modify a labeled item and change its label, downstream inheritance ensures the updated label propagates to dependent items, provided fully automated downstream inheritance is enabled or you consent to propagation.
- **Label reevaluation upon refresh**: When you refresh a semantic model and its data source has a sensitivity label, the system reevaluates and updates the semantic model's label if the data source label is more restrictive than the current label.

> ⓘ **Note**
>
> For non-Power BI Fabric items, default label application on update currently works only when you make changes via the item's **flyout menu**. Changes made in the experience interface don't trigger default labeling.

Sensitivity label inheritance when you establish new data source connections

When you connect a Power BI semantic model to sensitivity-labeled data in a supported data source, the item can inherit the data source's label. This behavior ensures that as Fabric items expand their data connections, they reflect the sensitivity classification of the most restrictive data they access.

The following behaviors apply when you establish a new data source connection:

- The item inherits the label of the newly connected data source if the data source label is more restrictive than the item's current label.
- If the item connects to multiple data sources with different sensitivity labels, Fabric applies the most restrictive label.

- If the data source label is less restrictive than the item's current label, the label doesn't change.
- You must publish the label for the item owner for Fabric to apply it.
- Label inheritance from data sources is currently supported for Power BI semantic models only. Connecting a non-Power BI Fabric item to a labeled data source doesn't trigger data source label inheritance.

For detailed requirements for data source label inheritance, including supported data sources and tenant settings, see Sensitivity label inheritance from data sources.

Sensitivity label inheritance when new lineage relationships are created

When an existing Microsoft Fabric item establishes a new lineage relationship, Microsoft Fabric evaluates sensitivity label inheritance based on the labels of the related items. This scenario covers cases such as:

- You connect a notebook to a lakehouse through the Microsoft Fabric UI.
- You configure a pipeline to read from a new data source item.

When you establish a new lineage relationship:

- The newly connected item evaluates the labels of its upstream items.
- If the upstream item has a more restrictive label, the downstream item inherits that label, provided the label wasn't manually applied by a user. For the full set of inheritance conditions, see downstream inheritance.
- If fully automated downstream inheritance is enabled, the label propagates automatically. Otherwise, the user is prompted to consent to label propagation.

Example: Notebook connected to a lakehouse

Consider a notebook with a *General* sensitivity label that you connect through the Microsoft Fabric UI to a lakehouse labeled *Confidential*. Because *Confidential* is more restrictive than *General*, the notebook inherits the *Confidential* label. Any items downstream of the notebook also receive the *Confidential* label through downstream inheritance.

Sensitivity label propagation rules

The following rules govern how sensitivity labels propagate during update and relationship change scenarios in Microsoft Fabric:

Expand table

Rule	Description
More restrictive label wins	A label is applied only if it's more restrictive than the currently applied label.
Manual labels are preserved	Automatically inherited labels never overwrite a label that a user manually applied.
Owner must have label access	The label must be published for the item owner; otherwise, inheritance doesn't take place.
Downstream propagation	When a label is updated through inheritance, the change propagates downstream if fully automated downstream inheritance is enabled.
Refresh triggers reevaluation	For semantic models, a data refresh reevaluates and potentially updates the label based on the current data source labels.

Considerations and limitations

Sensitivity label inheritance upon update has the following considerations and limitations:

- For non-Power BI Fabric items, label inheritance upon update is limited to changes made to attributes like name and description. You can only make these changes through certain UI paths, such as the item's flyout menu. For more information, see Default labeling in Fabric.
- Downstream inheritance is limited to 80 items. If the number of downstream items exceeds 80, downstream inheritance doesn't take place for any of the downstream items.
- Inheritance from data sources currently supports only Power BI semantic models.
- Label inheritance doesn't apply when you connect items through gateways or Azure Virtual Networks (VNets).
- Relaxed label change enforcement applies to accommodate automatic labeling scenarios where no user has all the required permissions to change a label.

Related content

- Information protection in Microsoft Fabric
- Sensitivity label downstream inheritance
- Sensitivity label inheritance from data sources
- Sensitivity labels in Power BI
- Default label policy for Power BI

ⓘ**Note:** The author created this article with assistance from AI. Learn more

Last updated on 07/16/2026

Audit schema for sensitivity labels in Fabric

Whenever a sensitivity label on a Fabric item is applied, changed, or removed, that activity is recorded in the audit log for Fabric, where you can track it. For information about tracking activities in the audit log, see Track user activities in Microsoft Fabric.

This article documents the information in the Fabric auditing schema that's specific to sensitivity labels. It covers the following activity keys:

- SensitivityLabelApplied
- SensitivityLabelChanged
- SensitivityLabelRemoved

SensitivityLabelEventData

Expand table

Field	Type	Must appear in the schema	Description
SensitivityLabelId	Edm.Guid		The guid of the new label. This field is only present when the activity key is SensitivityLabelApplied or SensitivityLabelChanged.
OldSensitivityLabelId	Edm.Guid		The guid of the label on the item before the action. This field is only present when the activity key is SensitivityLabelChanged or SensitivityLabelRemoved.
ActionSource	Edm.Enum	Yes	This field indicates whether the label change is the result of an automatic or manual process.
ActionSourceDetail	Edm.Enum	Yes	This field gives more detail about what caused the action to take place.
LabelEventType	Edm.Enum	Yes	This field indicates whether the action resulted in a more restrictive label, less restrictive label, or a label of the same degree of sensitivity.

ArtifactType

This field indicates the type of item the label change took place on.

Expand table

Value	Field
1	Power BI dashboard
2	Power BI report
3	Power BI semantic model
7	Power BI dataflow
11	Datamart
12	Fabric item

ActionSource

This field indicates whether the label change is the result of an automatic or manual process.

Expand table

Value	Meaning	Description
2	Auto	An automatic process performed the action.
3	Manual	A manual process performed the action.

ActionSourceDetail

This field gives more detail about what caused the action to take place.

Expand table

Value	Meaning	Description
0	None	There are no other details.
3	AutoByInheritance	The label change took place as a result of an automatically triggered inheritance process.
4	AutoByDeploymentPipeline	The label change took place automatically as a result of the deployment pipeline process.
5	PublicAPI	The label change action was performed by one of the following Fabric public admin REST APIs: bulkSetLabels, bulkRemoveLabels.

LabelEventType

This field indicates whether the action resulted in a more restrictive label, less restrictive label, or a label of the same degree of sensitivity.

Expand table

Value	Meaning	Description
1	LabelUpgraded	A more restrictive label was applied to the item.
2	LabelDowngraded	A less restrictive label was applied to the item.
3	LabelRemoved	The label was removed from the item.
4	LabelChangedSameOrder	The label was replaced by another label with the same level of sensitivity.

Related content

- Track user activities in Fabric

Last updated on 08/04/2024

Sensitivity labels troubleshooting

Licensing and requirements

What licenses do I need to be able to view and apply sensitivity labels?

- An Azure Information Protection Premium P1 or Premium P2 license is required to apply or view sensitivity labels from Microsoft Purview Information Protection in Power BI.
- To be able to apply labels to Power BI content and files, a user must have a Power BI Pro or Power BI Premium Per-User (PPU) license, in addition to one of the Azure Information Protection licenses mentioned above.
- Office apps have their own licensing requirements for viewing and applying sensitivity labels.

What requirements and prerequisites are there for enabling sensitivity labels in my tenant?

- If your organization uses Azure Information Protection sensitivity labels, they need to be migrated to the Microsoft Purview Information Protection Unified Labeling platform in order for them to be used in Power BI. Learn more about migrating sensitivity labels.
- Before enabling sensitivity labels on your tenant, make sure that sensitivity labels have been defined and published for relevant users and groups. See Create and configure sensitivity labels and their policies for detail.
- Customers in China must enable rights management for the tenant and add the Microsoft Purview Information Protection Sync Service service principal, as described in steps 1 and 2 under Configure Azure Information Protection for customers in China.
- Using sensitivity labels in Desktop requires the Desktop December 2020 release and later.

General problems with sensitivity labels

I can't enable sensitivity labels on my tenant

Power BI uses sensitivity labels from Microsoft Purview Information Protection. Thus, if you encounter an error message when trying to enable sensitivity labels, it might be due to one of the following:

- You don't have an Azure Information Protection license.
- Sensitivity labels haven't been migrated to the Microsoft Purview Information Protection version supported by Power BI.
- No sensitivity labels from Microsoft Purview Information Protection have been defined in the organization.

I can't apply sensitivity labels

To be able to apply or change a sensitivity label, you must

- Have a Power BI Pro or Power BI Premium Per-User (PPU) license.
- Have create and edit permissions on the item you want to apply the label to.
- Belong to a security group that has permissions to apply sensitivity labels, as described in Enable sensitivity labels in Power BI.

The sensitivity label I want to apply is greyed out

If a particular label you wish to change is greyed out, you may not have the correct usage rights to change that label. If you need to change a sensitivity label and can't, either ask the person who applied the label in the first place to modify it, or contact the Microsoft 365/Office security administrator and request the necessary usage rights for the label.

The sensitivity button in Desktop is greyed out

If the sensitivity button is greyed out, it may indicate that you don't have an appropriate license or that you don't belong to a security group that has permissions to apply sensitivity labels, as described in Enable sensitivity labels in Power BI.

The sensitivity label doesn't protect an exported file

Sensitivity labels and file encryption protect data only when it leaves Power BI via supported export paths. Data that leaves Power BI via unsupported export paths won't inherit the sensitivity label and won't be encrypted.

To prevent leakage of sensitive data, the Power BI admin can block export from non-supported export paths using Power BI's export and sharing settings.

Miscellaneous problems with sensitivity labels

- Don't use parent labels. A parent label is a label that has sublabels. You can't apply parent labels, but a label that is already applied may become a parent label if it acquires sublabels. If you come across an item that has a parent label, apply the appropriate sublabel. To change a parent label, you must have sufficient usage rights on the label.

  If an item has a parent label, note the following behavior:
  - Parent labels won't be inherited.
  - Mandatory label policies won't be applied to items that have a parent label. This means users won't be required to apply a meaningful label in order to save the item, and the item will escape mandatory label policies designed to promote total coverage.
  - If you try to export data from an item that has a parent label, export will fail.
  - It's possible to publish a *.pbix* file that has a parent label, but if the parent label is protected, publish will fail. The solution is to apply a suitable sublabel.

- In the Power BI service, if a semantic model has a label that has been deleted from the label admin center, you won't be able to export or download the data. In Analyze in Excel, a warning will be issued and the data will be exported to an .odc file with no sensitivity label. In Desktop, if a *.pbix* file has such an invalid label, you won't be able to save the file.

- Power BI doesn't support sensitivity labels of the Do Not Forward, user-defined, and HYOK protection types. The Do Not Forward and user-defined protection types refer to labels defined in the Purview compliance portal.

- Get data and refresh scenarios from encrypted Excel (*.xlsx*) files are supported, unless the file is stored behind a gateway, in which case the Get data/refresh action will fail. Get data and refresh actions from an Excel file that is stored behind a gateway and that has an **unprotected** sensitivity label will succeed, but the sensitivity label won't be inherited. See Sensitivity label inheritance from data sources for detail.

Problems with PBIX files

Export from the Power BI service to a .pbix file fails when I try to export a report that has a protected sensitivity label

When you export a report with a protected sensitivity label from the Power BI service to a pbix file, if the pbix file size gets to be greater than 6 GB, the protected label can't be applied (due to a Microsoft Purview Information Protection limitation) and the export fails.

I can see a report and semantic model in the Power BI service, but when I download them to pbix, I get a message that says I don't have sufficient permissions to open the file

In the Power BI service, sensitivity labeling doesn't affect access to content. Access to content in the service is determined solely by the permissions a user has on the content. While the labels are visible in the service, any associated encryption settings (configured in the Microsoft Purview portal) aren't applied. They're applied only to data that leaves the service via supported export paths.

In Power BI Desktop, sensitivity labels with encryption settings affect access to content. If a user doesn't have sufficient permissions according to the encryption settings of the sensitivity label on the *.pbix* file, they won't be able to open the file. In addition, in Desktop, when you save your work, any sensitivity label you've added and its associated encryption settings will be applied to the saved *.pbix* file.

Can't open protected *.pbix* file in Desktop

Using sensitivity labels in Desktop requires the Desktop December 2020 release and later. If you try to open a protected *.pbix* file with a Desktop version earlier than December 2020, it will fail, and you'll be prompted to upgrade your Desktop version.

Users with a free license can't open protected *.pbix* files.

Protected *.pbix* files can be only opened by a user who has an appropriate license and Full control and/or Export usage rights for the relevant label. The user that set the label also has Full control and can never be locked out. See more detail

In rare cases, it may happen that no one has the necessary usage rights for the relevant label except the person that set the label. Then, if that one person leaves the organization or changes aliases within the organization, all access to the .pbix file will be lost. The solution for regaining access to the file in such cases is to either change or remove the sensitivity label on the file using the set/remove sensitivity label Admin APIs. Contact your Power BI admin for assistance (only admins can run the Admin APIs).

Can't save a labeled *.pbix* file in Desktop

Power BI Desktop users may experience problems saving their work when internet connectivity is lost, such as after going offline. With no internet connection, some actions related to sensitivity

labels and rights management might not complete properly. In such cases it's recommended to go back online and try saving again.

In general, when you protect a file with a sensitivity label that applies encryption, it's good practice to use another encryption method as well, such as pagefile encryption, NTFS encryption, BitLocker instances, antimalware, etc.

Can't publish or get data of a protected *.pbix* file

"Publish" or "Get data" of a protected *.pbix* file requires that the label on the *.pbix* file be in the user's label policy. If the label isn't in the user's label policy, the Publish or Get data action will fail.

Can't publish or import a *.pbix* file with a sensitivity label to the service using APIs running under a service principal

Publishing or importing a .pbix file that has a **protected** sensitivity label to the service via APIs running under a service principal **is not** supported and will fail. To mitigate, users can remove the label and then publish using service principals.

Can't upload a protected file to Desktop via Get data

Import of sensitivity-labeled *.pbix* files (both protected and unprotected) stored on OneDrive or SharePoint Online, as well as on-demand and automatic semantic model refresh from such files, is supported, **with the exception of the following scenarios**:

- Protected live-connected *.pbix* files and protected Azure Analysis Services *.pbix* files: Refresh will fail. Neither report content nor label will be updated.
- Labeled unprotected Live Connect *.pbix* files: Report content will be updated but label won't be updated.
- When the *.pbix* file has had a new sensitivity label applied that the semantic model owner doesn't have usage rights to. In this case, refresh will fail. Neither report content nor label will be updated.
- If the semantic model owner's access token for OneDrive/SharePoint has expired. In this case, refresh will fail. Neither report content nor label will be updated.

Can't open protected *.pbix* file in Power BI Desktop for Power BI Report Server

Power BI Desktop for Power BI Report Server doesn't support information protection. If you try to open a protected *.pbix* file, the file won't open, and you'll receive an error message. Sensitivity-

labeled *.pbix* files that aren't encrypted can be opened as normal.

Connecting to data sources

To successfully connect from Fabric or Power BI (including Power BI Desktop) to a data source (such as an Excel file) that has a sensitivity label that applies file encryption, information protection must be enabled in Fabric/Power BI (that is, the tenant setting **Allow users to apply sensitivity labels for content** must be set to *Enabled*).

Sovereign clouds

Sensitivity labels are supported in the following sovereign clouds:

- US Government: GCC, GCC High, DoD
- China: Customers in China must enable rights management for the tenant and add the Microsoft Purview Information Protection Sync Service service principal, as described in steps 1 and 2 under Configure Azure Information Protection for customers in China.

Sensitivity label support in template apps

Data sensitivity labels aren't supported for template apps. Sensitivity labels set by the template app creator are removed when the app is extracted and installed, and sensitivity labels added to artifacts in an installed template app by the app consumer are lost (reset to nothing) when the app is updated.

Default labeling

No default label is applied to new content I create in the Power BI service.

Default labeling in Power BI covers most common scenarios, but there may be some less common flows that still allow users to open or create unlabeled *.pbix* files or Power BI artifacts.

Default labeling in Power BI isn't supported for service principals and APIs. Service principals and APIs aren't subject to default label policies.

Default label policies in Power BI aren't supported for external guest users (Microsoft Entra B2B). When a B2B user opens or creates an unlabeled *.pbix* file in Power BI Desktop or Power BI artifact in the Power BI service, no default label is applied automatically.

The pbix file I created didn't get the default label, even though default labeling is enabled on my tenant

Default labeling in Power BI covers most common scenarios, but there may be some less common flows that still allow users to open or create unlabeled *.pbix* files or Power BI artifacts.

The default label applied to my Power BI content isn't the same as the label applied to my emails and files

Default label policy settings for Power BI are independent of the default label policy settings for files and email.

B2B users manage to open and create unlabeled *.pbix* files, even though default labeling is on.

Default label policies in Power BI aren't supported for external guest users (B2B users). When a B2B user opens or creates an unlabeled file in Power BI Desktop, no default label will be applied to the file automatically.

Mandatory labeling

Mandatory labeling for Power BI is enabled but some artifacts are getting created or saved without a label having to be applied.

Mandatory labeling in Power BI isn't supported for service principals and APIs. Service principals and APIs aren't subject to mandatory label policies.

There may be flows that allow the user to create or edit unlabeled content.

Mandatory labeling in Power BI isn't supported for external guest users (B2B users). B2B users aren't subject to mandatory label policies.

Downstream inheritance

Downstream inheritance is enabled, but some or all downstream items don't inherit the label

Downstream inheritance is limited to 80 items. If the number of downstream items exceeds 80, no downstream inheritance takes place. Only the item the label was actually applied to will receive the label.

Downstream inheritance never overwrites labels that were applied manually.

Downstream inheritance never overwrites a label with a less restrictive label.

Sensitivity labels inherited from data sources are automatically propagated downstream only when fully automated downstream inheritance mode is enabled.

Sensitivity labels are applied to new reports even though "Automatically apply sensitivity labels to downstream content" is disabled

When you create a new Power BI report based on a semantic model that has a sensitivity label, the report automatically inherits the label from the semantic model, even if the **Automatically apply sensitivity labels to downstream content** tenant setting is disabled.

This behavior is by design. Sensitivity label inheritance upon item creation is a separate feature from downstream inheritance. When you create a new report based on a labeled parent item, the parent's label is automatically applied to the new report to prevent the creation of unlabeled reports, which could expose data.

The **Automatically apply sensitivity labels to downstream content** setting controls whether labels are propagated to *existing* downstream items when you change the label on a parent item. It doesn't affect label inheritance when you create *new* items.

Sensitivity label inheritance from data sources

Sensitivity labels from a data source aren't inherited into Power BI.

- The data in the data source must be labeled with sensitivity labels from Microsoft Purview Information Protection.
- The scope of the labels must be Files and emails and Azure Purview assets. See Extending sensitivity labels to Azure Purview and Creating new sensitivity labels or modifying existing labels.
- Sensitivity labels must be enabled in Power BI.

- The **Apply sensitivity labels from data sources to their data in Power BI (preview)** tenant admin setting must be enabled.
- All conditions for applying a label must be met.
- Inheritance from data sources isn't supported for semantic models located in classic workspaces. My Workspace and V2 workspaces are supported.
- Inheritance from data sources is supported only for semantic models with enhanced metadata. See Using enhanced dataset metadata for more information.
- Inheritance from data sources is supported only for semantic models using the Import data connectivity mode. Live connection and DirectQuery connectivity isn't supported.
- Inheritance from data sources isn't supported in connections via gateways or Azure Virtual Network (VNet). This means that inheritance from an Excel file located on a local machine won't work, because this requires a gateway.

Label inheritance upon update and relationship changes

A sensitivity label isn't applied when I modify an existing item

Label inheritance upon update follows the same core rules as other automatic labeling capabilities: automatically applied labels never overwrite manually applied labels, and a less restrictive label never overwrites a more restrictive label. In addition:

- For non-Power BI Fabric items, default label application on update currently supports only when you change attributes like name and description through the item's flyout menu. Changes made in the experience interface don't trigger default labeling.
- You must publish the label for the item owner. If the item owner isn't authorized to apply the label, inheritance doesn't take place.

A label isn't inherited when I connect an item to a new data source or establish a new lineage relationship

- Label inheritance from data sources currently supports only Power BI semantic models. Connecting a non-Power BI Fabric item to a labeled data source doesn't trigger data source label inheritance.
- Inheritance from data sources doesn't support connections via gateways or Azure Virtual Networks (VNets).
- If fully automated downstream inheritance isn't enabled, you must consent to label propagation when a new lineage relationship is established.

For more information, see Sensitivity label inheritance upon update and relationship changes.

Problems setting and removing sensitivity labels using Power BI REST APIs

Can't set or remove sensitivity labels using Power BI REST admin APIs

- Users must be Fabric administrators to call these APIs.
- The admin user (and the delegated user, if provided) must have sufficient usage rights to set or remove labels.
- To set a sensitivity label using the setLabels API, the admin user (or the delegated user, if provided) must have the label included in their label policy.
- The APIs allow a maximum of 25 requests per hour. Each request can update up to 2000 artifacts.
- **Required scope**: Tenant.ReadWrite.All

Defender for Cloud Apps

To use Defender for Cloud Apps with Power BI, you must use and configure relevant Microsoft security services, some of which are set outside Power BI. In order to have Defender for Cloud Apps in your tenant, you must have one of the following licenses:

- Defender for Cloud Apps: Provides Defender for Cloud Apps capabilities for all supported apps, part of the EMS E5 and Microsoft 365 E5 suites.
- Office 365 Cloud App Security (a subset of Defender for Cloud Apps): Provides Cloud App Security capabilities only for Office 365, part of the Office 365 E5 suite.

Using Defender for Cloud Apps with Power BI is designed to help secure your organization's content and data, with detections that monitor user sessions and their activities. When using Defender for Cloud Apps with Power BI, there are a few considerations and limitations you should keep in mind:

- Defender for Cloud Apps can only operate on Excel, PowerPoint, and PDF files.
- If you want to use sensitivity labels capabilities in your session policies for Power BI, you need to have an Azure Information Protection Premium P1 or Premium P2 license. Microsoft Azure Information Protection can be purchased either standalone or through one of the Microsoft licensing suites. See Azure Information Protection pricing for detail. In addition, sensitivity labels must have been applied on your Power BI assets.

- Session control is available for any browser on any major platform on any operating system. We recommend using Internet Explorer 11, Microsoft Edge (latest), Google Chrome (latest), Mozilla Firefox (latest), or Apple Safari (latest). Power BI public API calls and other non-browser-based sessions aren't supported as part of Defender for Cloud Apps session control. See more detail.
- If you experience sign in difficulties, such as having to log in more than once, it could be related to the way some apps handle authentication. See Slow login in the Defender for Cloud Apps documentation for more information and remediation steps.

> **Caution**
>
> In the session policy, in the "Action" part, the "protect" capability works only if no label exists on the item. If a label already exists, the "protect" action won't apply; you can't override an existing label that has already been applied to an item in Power BI.

I can't access the Defender for Cloud Apps information.

In order to access Defender for Cloud Apps information, your organization must have the appropriate Defender for Cloud Apps license.

Paginated reports

My paginated report doesn't inherit the sensitivity label of the model it's based on.

Downstream inheritance isn't supported. The label of an upstream model won't propagate down to its downstream paginated reports.

The sensitivity label on my paginated report doesn't get applied to the report's downstream content.

The label of a paginated report won't propagate down to the report's downstream content.

I can create and save paginated reports with no sensitivity label even though mandatory labeling is enabled.

Mandatory labeling won't apply to paginated reports.

Last updated on 07/13/2026

Protection policies in Microsoft Fabric

Microsoft Purview protection access control policies (protection policies) enable organizations to control access to items in Fabric using sensitivity labels.

The target audience for this article is security and compliance admins, Fabric admins and users, and anyone else who wants to learn about how protection policies control access to items in Fabric. If you want to see how to create a protection policy for Fabric, see Create and manage protection policies for Fabric.

How do protection policies for Fabric work?

Each protection policy for Fabric is associated with a sensitivity label. The policy controls access to an item that has the associated label by allowing users and groups specified in the policy to retain permissions they have on the item, while blocking access for everyone else.

The policy can allow specified users and groups to retain full control on the labeled item if they currently have it. All other users and groups are blocked from accessing the item.

> **Note**
>
> A protection policy doesn't apply to a label issuer. That is, the user that last applied a label associated with a protection policy to an item won't be denied access to that item, even if they aren't specified in the policy. For example, if a protection policy is associated with label A, and a user applies label A to an item, that user will be able to access the item even if they're not specified in the policy.

View an item's restricted users

In the OneLake catalog, you can view an item's permissions and determine which users are restricted from accessing it. The **Permissions** tab in the item's details is visible to you if you have a role of Admin or Member in the workspace containing the item.

To view an item's permissions, open the **OneLake catalog**, locate the item (optionally filter by workspace), and select its name to open the item details. Then, select the **Permissions** tab to see the list of users and groups that have access to the item, including those restricted by a protection policy.

Use cases

Protection policies are useful in scenarios where organizations need to restrict access to sensitive items. For example, an organization might want to ensure that only specific internal users can access items labeled as "Confidential," while blocking access for everyone else.

Who creates protection policies for Fabric?

Protection policies for Fabric are generally configured by an organization's Purview security and compliance teams. The protection policy creator must have role of Information Protection Admin or higher. For more information, see Create and manage protection policies for Fabric.

Requirements

- A Microsoft 365 E3/E5 license is required for sensitivity labels from Microsoft Purview Information Protection. For more information, see Microsoft Purview Information Protection: Sensitivity labeling.
- At least one "appropriately configured" sensitivity label from Microsoft Purview Information Protection must exist in the tenant. "Appropriately configured" in the context of protection policies for Fabric means that when the label was configured, it was scoped to **Files & other data assets**, and its protection settings were set to include **Control access** (for information about sensitivity label configuration, see Create and configure sensitivity labels and their policies). Only such "appropriately configured" sensitivity labels can be used to create the protection policies for Fabric.
- For protection policies to be enforced in Fabric, the Fabric tenant setting **Allow users to apply sensitivity labels for content** must be enabled. This setting is required for all sensitivity label related policy enforcement in Fabric, so if sensitivity labels are already being used in Fabric, this setting will already be on. For more information about enabling sensitivity labels in Fabric, see Enable sensitivity labels.

Supported item types

Protection policies are supported for all native Fabric item types, including lakehouses, notebooks, pipelines, and other core Fabric assets.

Additionally, protection policies are supported for Power BI semantic models. Other Power BI item types, such as reports and dashboards, aren't currently supported.

Considerations and limitations

- Up to 50 protection policies can be created.
- Up to 100 users and groups can be added to a protection policy.
- Protection policies for Fabric don't support guest/external users.
- Protection policies don't integrate with Fabric CI/CD solutions. Deployment pipelines and Git integration use only workspace permissions, not item-level permissions from protection policy labels.
- After a policy has been created, it can take up to 24 hours for it to start detecting and protecting items labeled with the sensitivity label that was associated with the policy.
- For native Fabric item types, if the label is applied automatically by the system—such as in the case of downstream inheritance—and there is no designated label issuer, the user who created the artifact will not be restricted by protection policies.

Related content

- Create and manage protection policies for Fabric
- Authoring and publishing protection policies

Last updated on 12/03/2025

Create and manage protection policies for Fabric

This article describes how to create, edit, and delete Microsoft Purview protection policies for Microsoft Fabric. Its target audience is security and compliance admins who need to create protection policies for Fabric.

For an overview of how protection policies for Fabric work, see Protection policies in Fabric.

Prerequisites

To create a protection policy, the following requirements must be met:

- You must have a Microsoft 365 E3/E5 license as required for sensitivity labels from Microsoft Purview Information Protection. For more information, see Microsoft Purview Information Protection: Sensitivity labeling.
- You must have at least the Information Protection Admin role to create protection policies in the Microsoft Purview portal. The Information Protection Admin role is assigned by default under the Compliance Administrator role group.
- At least one "appropriately configured" sensitivity label from Microsoft Purview Information Protection must exist in the tenant. "Appropriately configured" in the context of protection policies for Fabric means that when the label was configured, it was scoped to **Files & other data assets**, and its protection settings were set to include **Control access** (for information about sensitivity label configuration, see Create and configure sensitivity labels and their policies). Only such "appropriately configured" sensitivity labels can be used to create the protection policies for Fabric.
- If service principals need to access items protected by the protection policy you're configuring, add them to an Azure security group that you'll include in the policy. Service principals can't be added to the policy directly in the Microsoft Purview portal.

  If you don't add service principals to the allowed list of users, service principals that currently have access to data will be denied access, potentially causing your application to break. For example, service principals might be used for application authentication to access semantic models.

Create a protection policy for Fabric

1. Open the **Protection policies (preview)** page in the Microsoft Purview portal.

   https://purview.microsoft.com/informationprotection/protectionpolicy .

2. On the **Protection policies (preview)** page, select **New protection policy**.

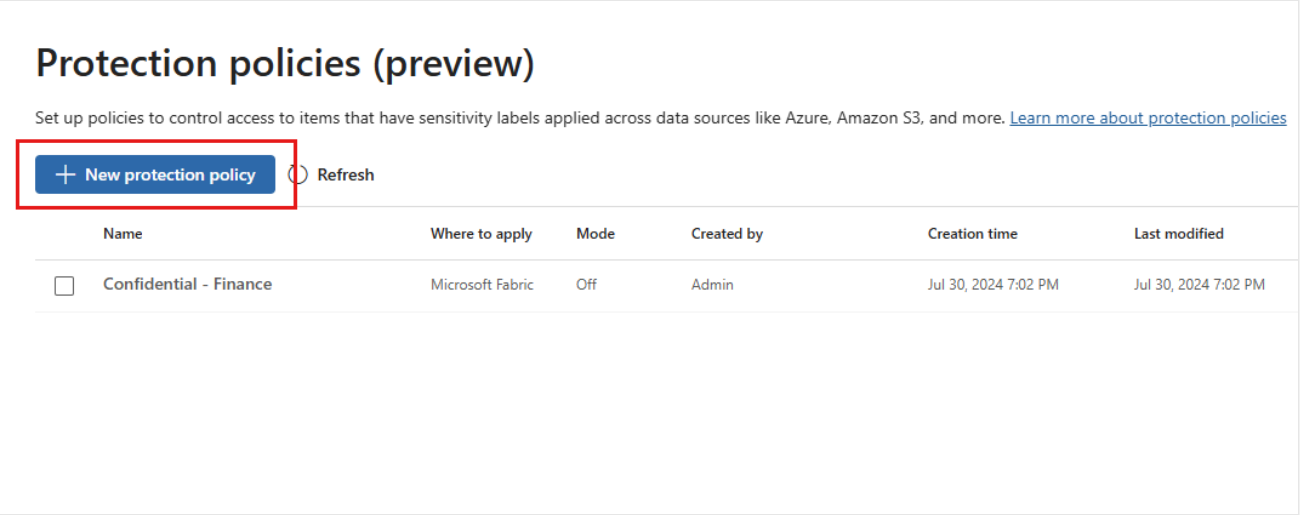

3. On the **Name and describe your protection policy** page, provide a name and description for the policy. When done, select **Next**.

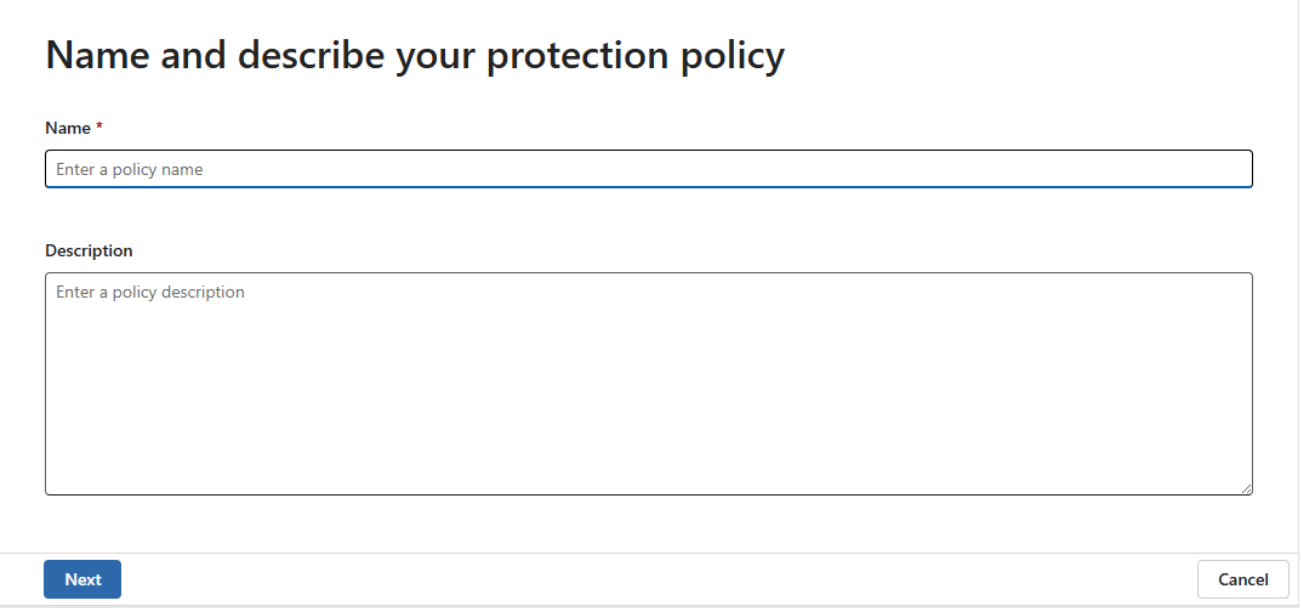

4. On the **Choose the sensitivity label used to detect sensitive items** page, select **+ Add sensitivity label** and choose the sensitivity label you want to associate with the policy. The label must be a label that applies encryption and is correctly configured. See the prerequisites for details. You can select only one label per policy.

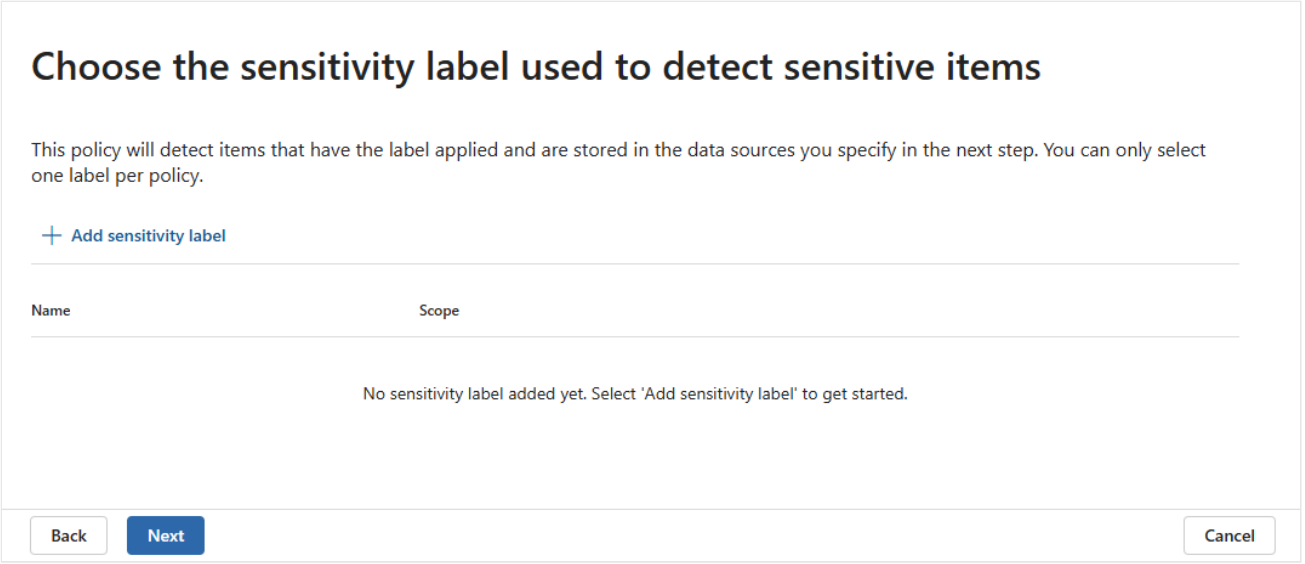

After you've chosen the label, select **Add** and then **Next**.

5. On the **Choose data sources to apply the policy** page, select **Microsoft Fabric**. If you see multiple data sources listed, be sure to select only Fabric. The policy will apply to supported item types in all workspaces. When done, select **Next**.

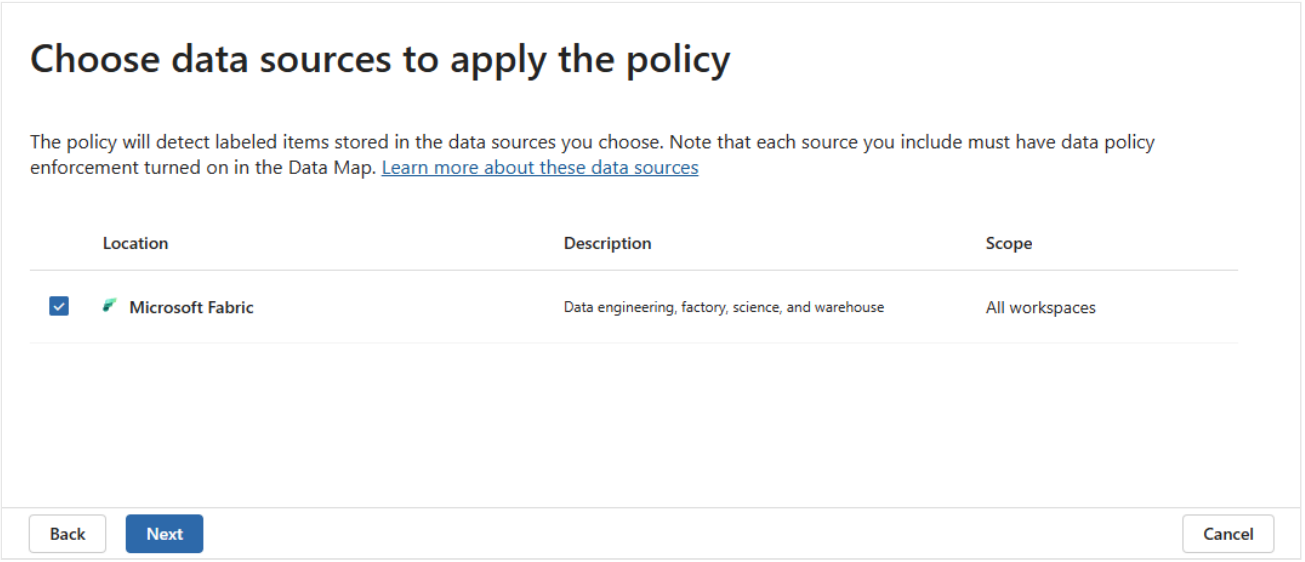

6. On the **Define access control settings** page, specify which users and groups should retain their access to items labeled with the sensitivity label you selected in step 4.

   a. Select the **Allow users to retain full control** checkbox if you want certain users and groups to keep their existing permissions on labeled items. When enabled, this setting allows users with full control of labeled items to retain full control.

   b. Select **Add users and groups** to choose specific users and groups who should retain their access to protected items.

c. Repeat the previous step to add all users and groups that need continued access to the protected items.

> **Note**
>
> Service principals can't be added to protection policies directly via the Microsoft Purview portal. To enable service principals to access items protected by a protection policy, first add them to an Azure security group, and then add this group to the protection policy.

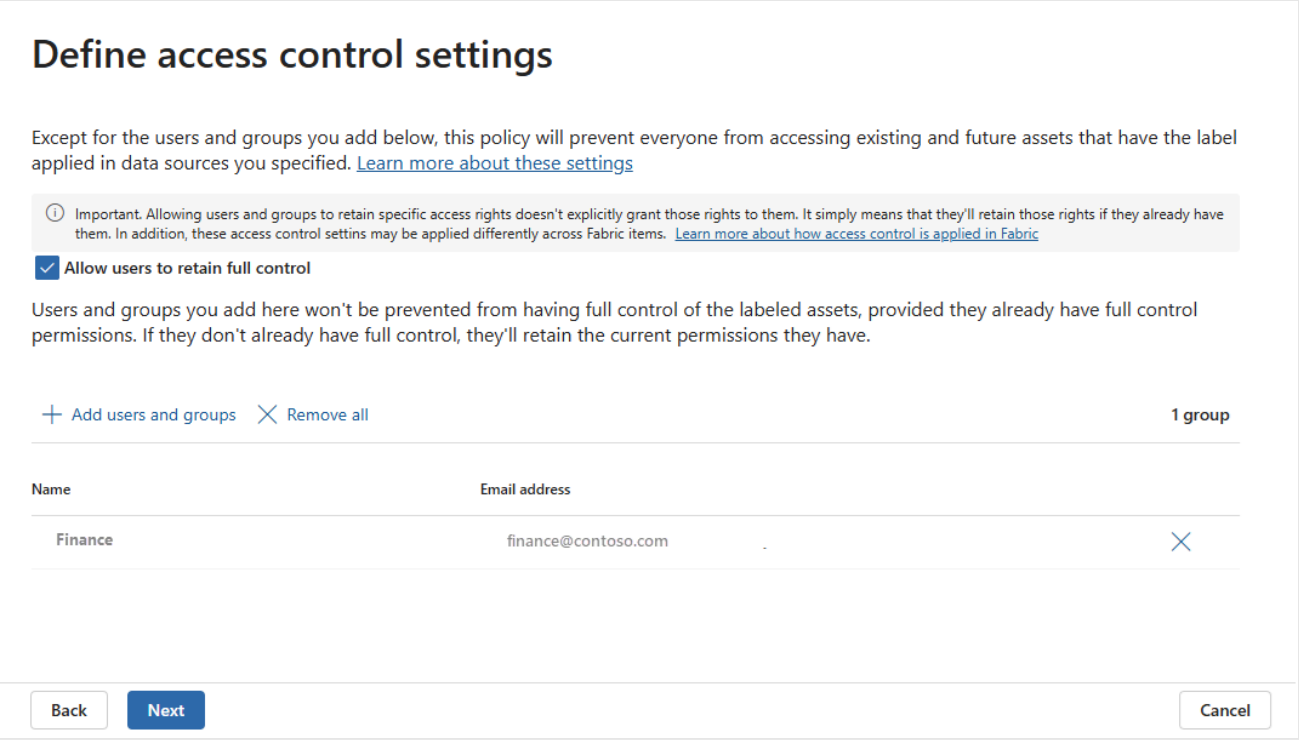

The policy will block access to items labeled with the associated sensitivity label for all users not specified in one of the above control.

When done, select **Next**.

7. On the **Decide whether to turn on the policy on right away or keep it off** page, choose whether turn on the policy right away or not. When done, select **Next**.

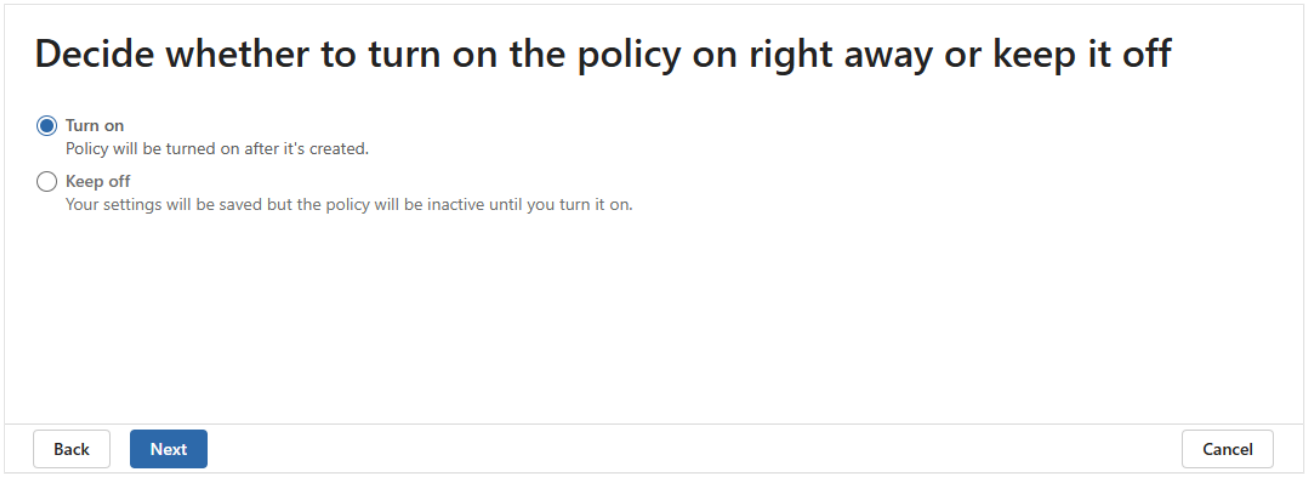

8. On the **Review your policy settings** page, review the policy settings. When you're satisfied, select **Submit**, and then **Done**.

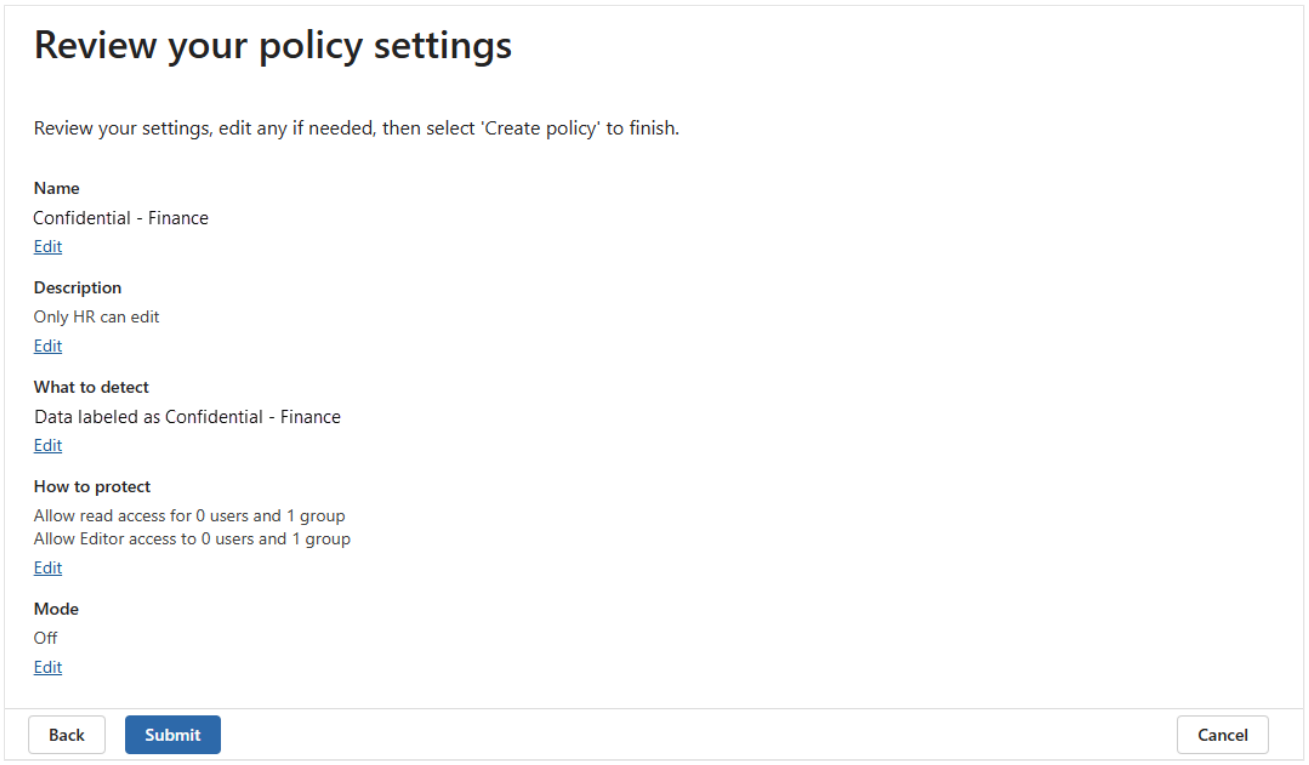

9. You'll be informed that your new protection policy has been created. It can take up to 24 hours for the new policy to start detecting and protecting items labeled with the sensitivity label you chose.

Your new policy now appears in the list of protection policies. Select it to confirm that all the details are correct.

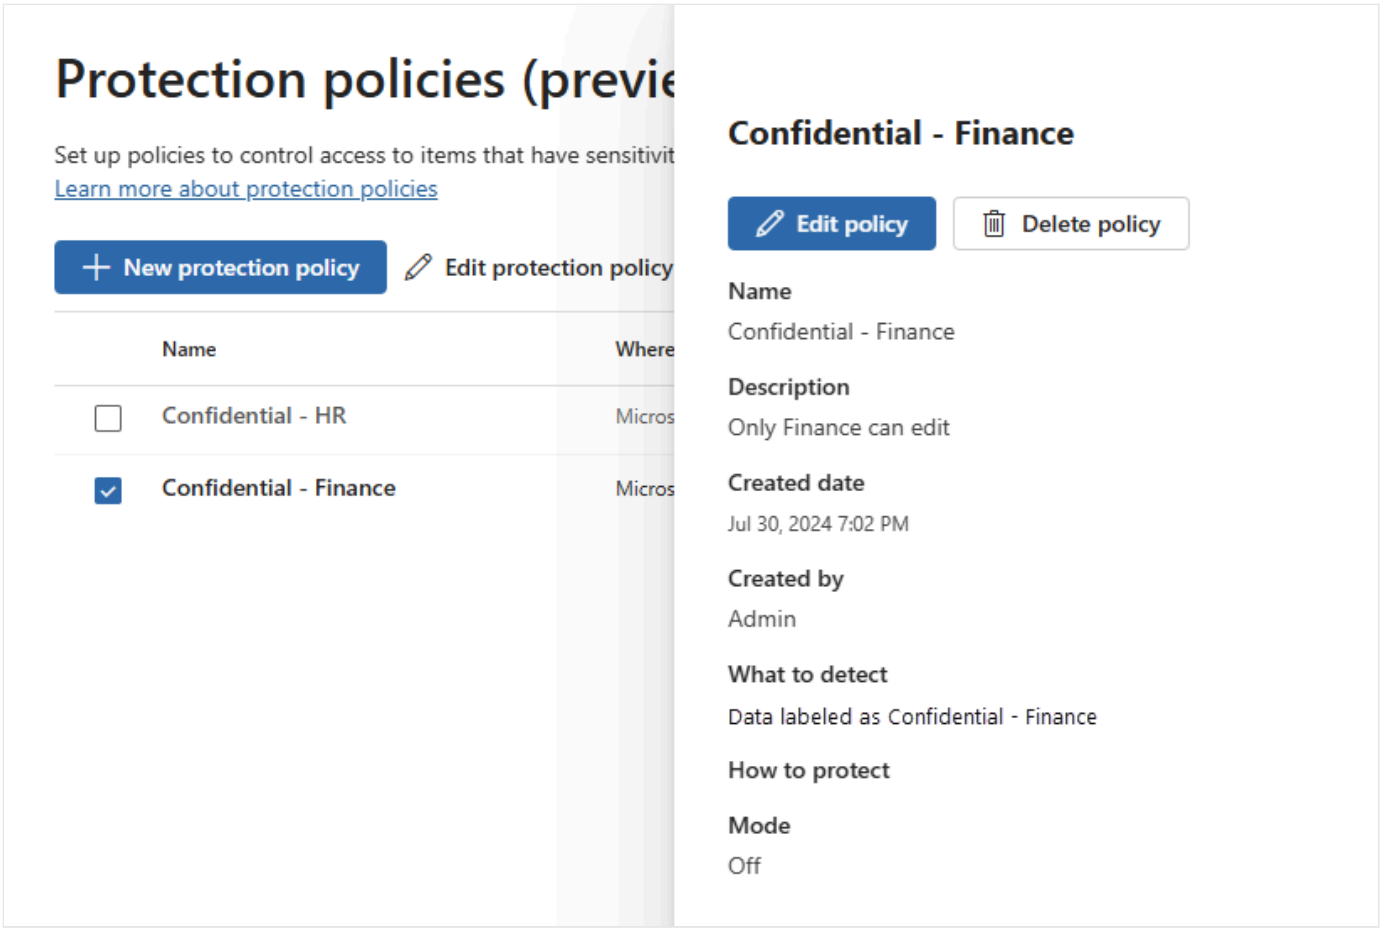

Manage protection policies

To edit or delete a protection policy:

1. Open the **Protection policies (preview)** page in the Microsoft Purview portal.

   https://purview.microsoft.com/informationprotection/protectionpolicy.

2. Select the policy you wish to edit or delete, and then select Edit policy or Delete policy either on the ribbon or in the details pane.

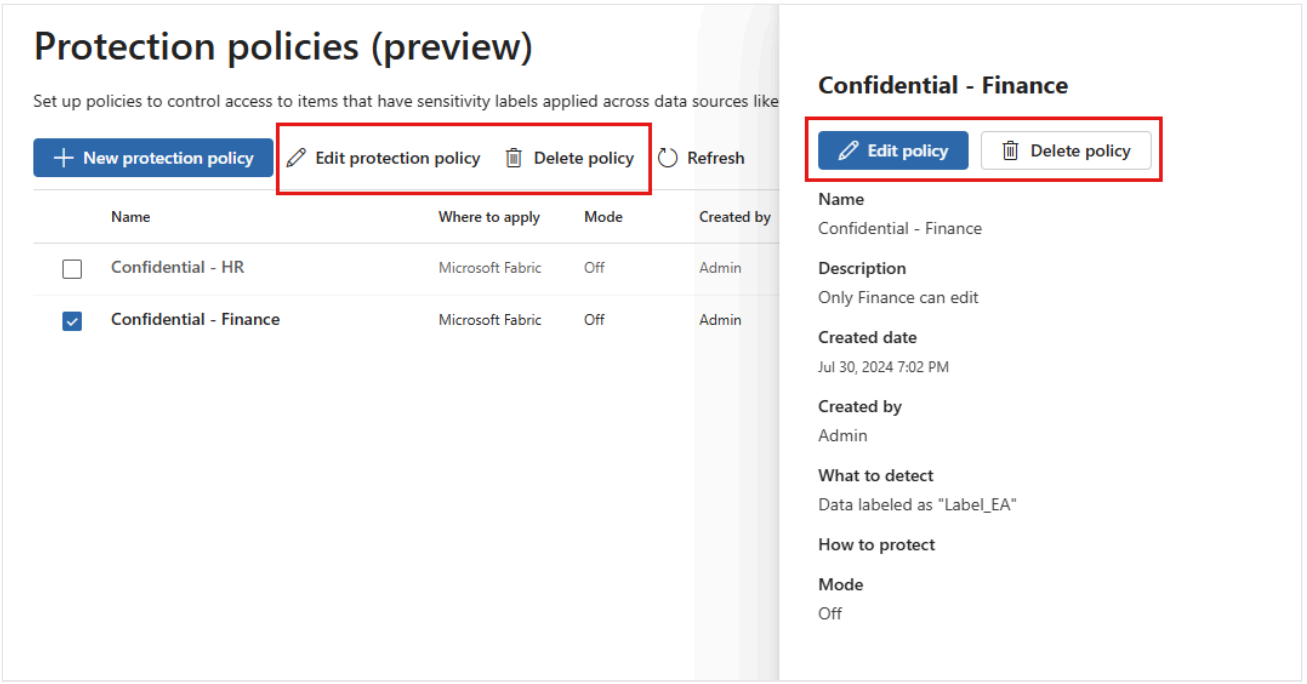

If you're editing the policy, continue cycling through the configuration pages as in Step 3 of the policy creation flow.

Related content

- Protection policies in Fabric
- Authoring and publishing protection policies (preview)

Last updated on 12/03/2025

Troubleshoot access restrictions

The Microsoft Fabric permission model allows users and workspace admins to grant permissions to other users at the workspace level. These permissions define the default access for all items in the workspace. Users can also be given elevated permissions for specific items at the item level, even if they're only viewers at the workspace level. Learn more about the Fabric permission model.

Fabric provides other access controls. For example, OneLake security enables admins to assign roles to items and control which portions of data users can view. More importantly for this troubleshooting article, organization-wide policies can automatically restrict access to items based on their sensitivity or content.

Because there are multiple layers controlling access, you could encounter restrictions from policies you're unaware of. This article helps you identify if a policy is affecting your access and how to restore it if possible.

User's access to an item is blocked

If you previously had access to an item but can no longer see it in your workspace or open it, consider the following common causes.

Access was manually removed

A user with administrative or ownership permissions manually revoked your access to the item. This action is often taken to comply with company regulations or security policies. To resolve this issue:

- **Contact the item owner directly**: If you know who owns or created the item, reach out to them to inquire about the access change.
- **Reach out to your workspace administrator**: If you don't know the item owner, your workspace admin can help you contact them or reinstate your access if they removed it by mistake.

Workspace permissions changed

You were removed from the workspace or your access level was downgraded, which automatically revoked some of your permissions to the items within it.

To understand the reason for the change in permissions, contact your workspace administrator.

Company policies changed

Your organization implemented company-wide policies that restrict access to certain items based on the applied their sensitivity or content. These policies can override individual permissions and workspace settings. There are two main types of policies that can enforce such restrictions:

- **Protection policies**: Microsoft Purview protection policies automatically enforce access restrictions based on the applied *sensitivity label*. For example, the label "Confidential" can automatically limit access to a specific security group, for any Fabric item labeled by it.
- **Data Loss Prevention (DLP) policies with the "restrict access" action**: Microsoft Purview DLP policies with the *restrict access* action allow for access enforcement on items containing sensitive information.

Both policies produce similar user experiences when access is revoked. The following sections provide troubleshooting steps for these scenarios.

How to tell if policies are the cause of access issues

You can tell a policy change is the cause of your restricted access if you experience one of the following issues:

- You no longer see the item in your workspace list view or the OneLake catalog.

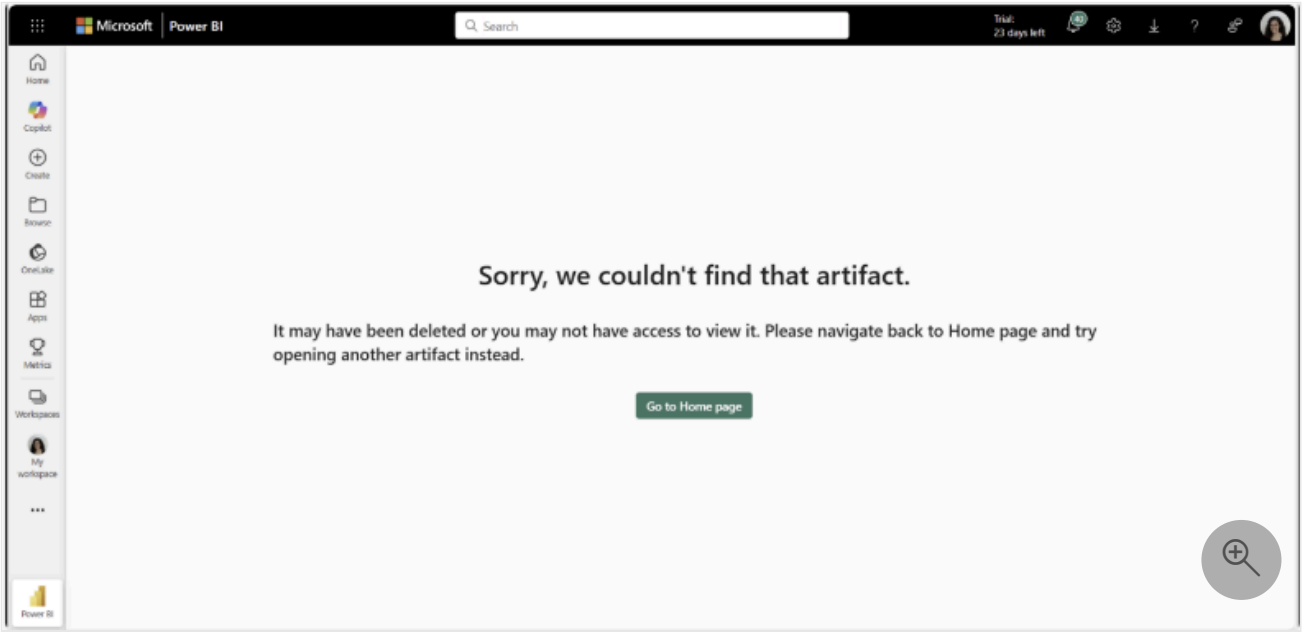

- You see "Permission required" or "Sorry we couldn’t locate that artifact" when trying to open the item using a link.

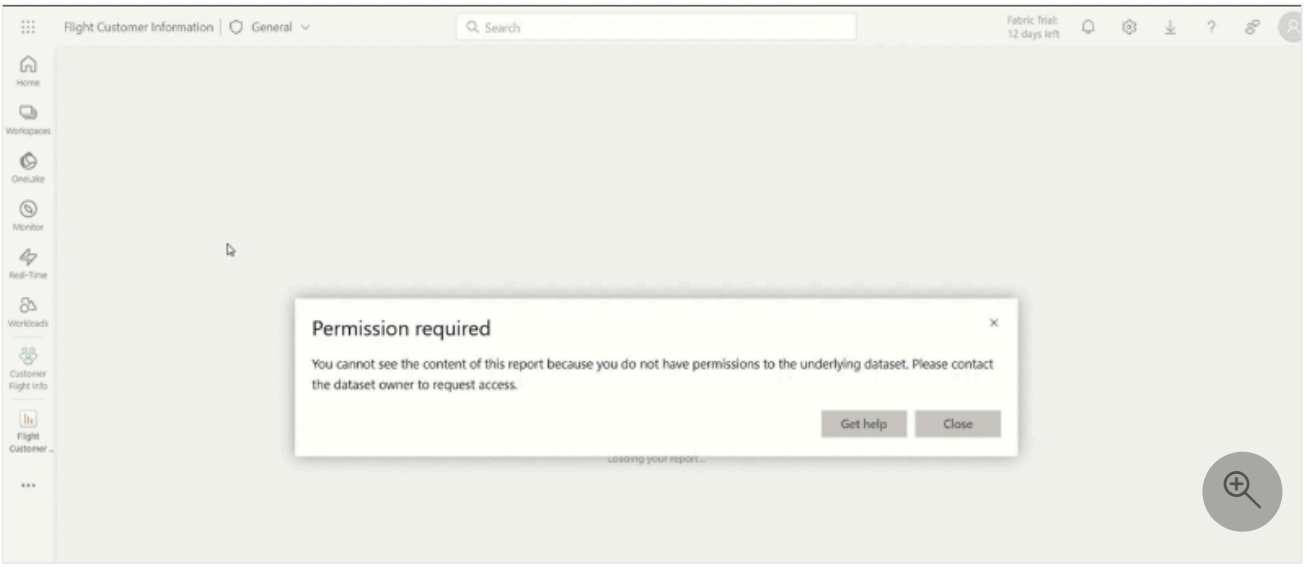

- You see an indicator on the item showing that it contains data from a restricted underlying source.

- You see tiles on a dashboard that display a "Permission required" message. This message indicates that these tiles contain data from a restricted underlying source. Tiles that don't use restricted data render normally.

These symptoms can indicate that either a protection policy or a DLP policy with the "restrict access" action is blocking your access.

What should I do if I think a policy is blocking my access?

Contact your workspace admin to investigate and resolve the issue. If the workspace admin is also blocked by a policy, the label issuer (the person who applied the sensitivity label) or the item creator can access the **Manage Permissions** page to determine if a policy is blocking your access.

The next sections provide guidance for workspace admins and others with appropriate permissions on how to troubleshoot and resolve these issues.

How to investigate a restriction scenario

When a user reports losing access to an item, the workspace admin can investigate the cause along with other users who can assist depending on their permissions:

- **Workspace admins** can access the **Manage Permissions** page to investigate access issues for all users.

- **Label issuers** (users who applied the sensitivity label to the item) retain access even if not included in the protection policy, preventing lockout scenarios.
- **Item creators** also retain access when labels are applied automatically (such as through downstream inheritance), ensuring at least one user can manage the item.

To investigate the issue:

1. Open the item's **Manage Permissions** page and check the user's **Permissions** setting.
2. If the user's setting is **User is blocked**, it indicates that the user's access was revoked by an organizational security policy.

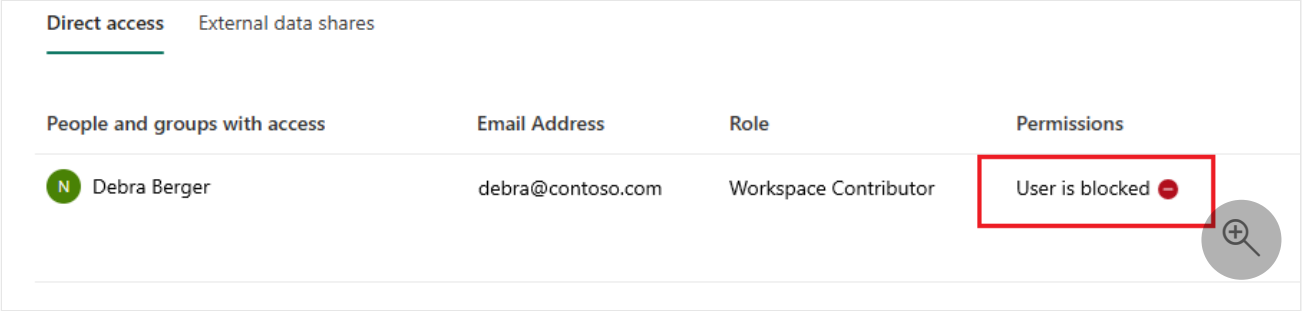

3. Identify the policy type:
    - Recent label changes likely mean an issue with a Microsoft Purview protection policy.
    - Recent item updates or sensitive data likely mean an issue with a Microsoft Purview DLP restriction.
4. If the user's **Permissions** setting doesn't show **User is blocked**, investigate other causes such as manual access removal or workspace permission changes.

How to resolve policy-related issues

To unblock users, the workspace admin or a user with appropriate permissions, such as the label issuer or item creator, can take action. The steps to resolve the issue vary based on whether a Microsoft Purview protection policy or Microsoft Purview DLP policy is causing the restriction.

> ⓘ **Note**
>
> Only take these actions after ensuring they don't risk organizational compliance or data security. These actions are audited and might be subject to investigation by your security administrator.

Resolving protection policy issues

When a protection policy revoked access and the user's permission shows as **No access**, the workspace admin can take the following actions. The label issuer (the person who applied the label to the item) or the item creator can also assist with investigating and resolving the issue.

- If the sensitivity label was applied incorrectly, either apply the correct label yourself or request that the Protection Owner update it.
- If the label is correct but the user needs access to this classified data, contact your organization's security team to request adding the user to the protection policy in Microsoft Purview. After the policy is updated, have the user sign back in to Fabric to receive the updated permissions.

Resolving DLP "restrict access" policy issues

If a DLP policy caused the restriction, the workspace admin can take the following actions:

- If the DLP policy incorrectly flagged the data as sensitive and you're confident it's a false positive, the workspace admin can override the policy. Be prepared to provide justification to your compliance team for this action.
- If the data was mistakenly added to the data item, the workspace admin can edit the item to remove the sensitive data. This action triggers a new evaluation that automatically removes the restriction.

Understanding how DLP policies affect different user types

If a DLP policy with the "restrict access" action is causing the restriction, it's important to understand how it affects different user types:

- **All users in the tenant**: When applied to this group, DLP policies block all users except workspace admins, who can troubleshoot and fix the data, where relevant.
- **Guest users**: When applied to this group, DLP policies block guest users, while retaining access for internal users and workspace admins. Guest users are users who aren't defined as **Member** in Microsoft Entra. If an external user is configured as a tenant **Member**, the policy treats them as an internal user.

The workspace admin can assess how DLP policies affect different user types to decide how to handle the restriction.

Mirrored databases restricted by DLP – special case

Mirrored databases present a slightly different use case. Usually, these types of databases are updated by service principals (SPNs) rather than by human users. Because organizations typically wouldn't grant these SPNs full control of workspaces (the workspace admin role), the updating user might be restricted by the DLP **restrict access** action.

How to resolve a database-restricted user

Workspace admins can always override the policy if the circumstances merit an override.

If sensitive information was justifiably found in the data, workspace admins can:

1. Edit the mirroring configuration to exclude the problematic table.
2. Edit the existing item to remove the sensitive information.

Once these two steps are completed, the next database update triggers an automated DLP evaluation. Because the sensitive data has been removed, the evaluation finds no sensitive information, and the restriction is lifted.

Related articles

- Microsoft Fabric permission model
- Microsoft Purview protection policies
- Data Loss Prevention (DLP) in Microsoft Purview

Last updated on 05/11/2026

Get started with data loss prevention policies for Fabric and Power BI

This article provides a general overview of Microsoft Purview Data Loss Prevention (DLP) policies for Fabric and Power BI. The target audience is Fabric administrators, security and compliance teams, and Fabric data owners. If you're a data owner and want to know how to respond when a policy tip tells you your item has a DLP policy match, see Respond to a DLP policy match in Fabric. If you're a Fabric admin or a security and compliance admin and need to audit alerts on DLP policy matches, see Monitor DLP policy matches in Fabric.

Overview

To help organizations detect and protect their sensitive data, Fabric supports Microsoft Purview Data Loss Prevention (DLP) policies. When a DLP policy for Fabric detects a supported item type containing sensitive information, the policy triggers the configured actions. These actions can include:

- Attaching a policy tip to the item that explains the nature of the sensitive content.
- Registering an alert for administrators on the data loss prevention **Alerts** page in the Microsoft Purview portal.
- Sending email alerts to administrators and specified users.
- Restricting access to the item.

For more detail, see How do DLP policies for Fabric and Power BI work.

For DLP policies, a *member of the organization* means a user with UserType = Member. This includes users who sign in with an external account. Because their UserType is Member, they're treated as internal users. In contrast, users with UserType = Guest are considered external. The **Block only people outside your organization** setting applies to these Guest users.

> **Tip**
>
> Get started with Microsoft Security Copilot to explore new ways to work smarter and faster using the power of AI. Learn more about Microsoft Security Copilot in Microsoft Purview.

Licensing and permissions

Licensing

For information on licensing, see

- Microsoft 365 Enterprise Plans
- Microsoft 365 Service Descriptions

Permissions

You can view data from DLP for Fabric and Power BI in Activity explorer. Four roles grant permission to Activity explorer; the account you use for accessing the data must be a member of any one of them.

To view the Activity explorer, your account must be a member of any of the following roles:

- Compliance administrator
- Security administrator
- Compliance data administrator

Billing

DLP evaluation workloads impact capacity consumption. For information about how this consumption is metered and billed, see Learn about Microsoft Purview billing models.

How DLP policies for Fabric and Power BI work

You define a DLP policy in the data loss prevention section of the Microsoft Purview portal. In the policy, you specify the conditions, such as which sensitivity labels and sensitive info types you want to detect. You also specify the actions that the system takes when it detects a policy match.

When DLP policies evaluate a supported item type, they check if it matches the conditions in a DLP policy. If it does, the policy's specified actions take place. DLP policies start with the following actions:

Semantic models

DLP policies evaluate a semantic model whenever one of the following events occurs:

- Publish
- Republish
- On-demand refresh

- Scheduled refresh

> **Note**
>
> DLP evaluation of the semantic model doesn't occur if either of the following is true:
>
> - An account using service principal authentication initiates the event (publish, republish, on-demand refresh, scheduled refresh).
> - The semantic model owner is a service principal.

Fabric Items

DLP policies evaluate a Fabric item, such as a lakehouse, SQL database, or Mirrored database, when the data within it changes. Changes include getting new data, connecting a new source, adding tables, updating existing tables, and more.

What happens when a Fabric DLP policy flags an item

When a DLP policy detects an issue with an item, it takes the following actions:

If you enable **user notification** in the policy, Fabric marks the item with an icon that indicates a DLP policy detected an issue with the item. When you hover over the icon, a hover card appears that provides an option to see the full details in a side panel. For more information about what you see in the side panel, see Respond to a DLP violation in Fabric.

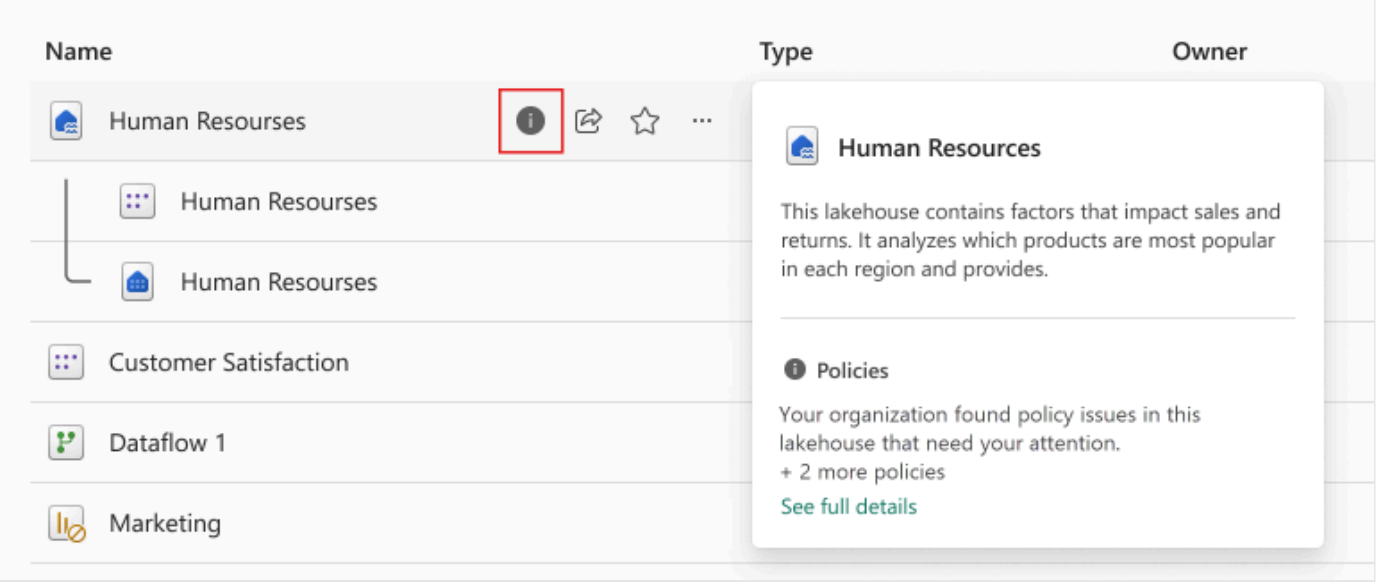

For semantic models, opening the details page shows a policy tip that explains the policy violation and how to handle the type of sensitive information detected. Selecting **View all** opens

a side panel with all the policy details.

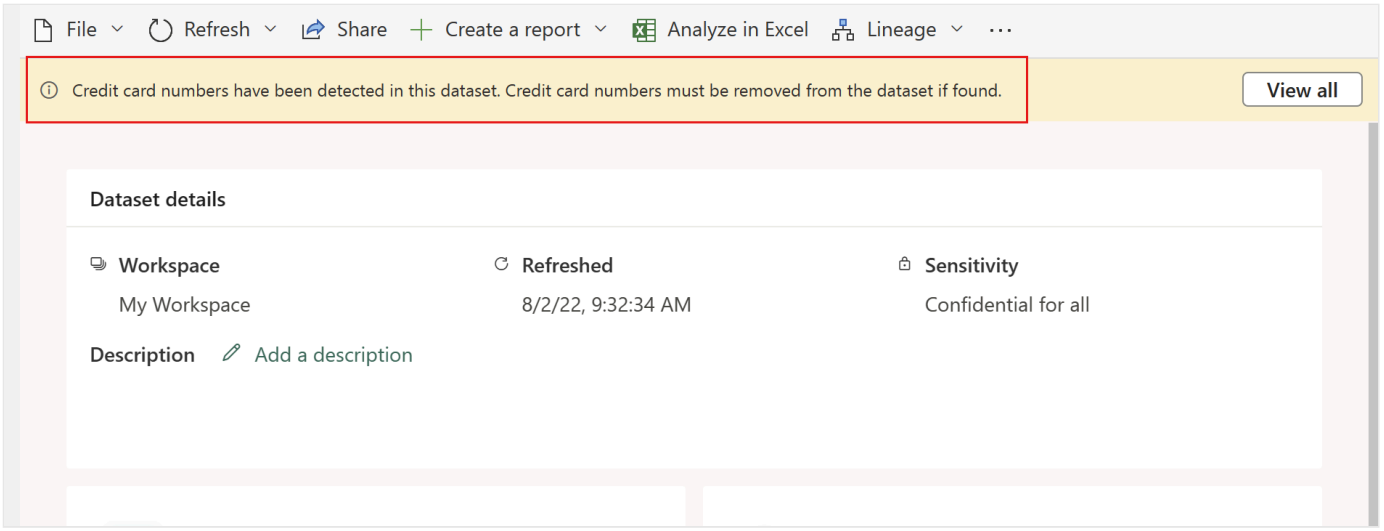

> **Note**
>
> If you hide the policy tip, it doesn't get deleted. It appears the next time you visit the page.

For Fabric items, such as lakehouses, the indication appears in the header in edit mode. Opening the flyout makes it possible to see more details about the policy tips affecting the item. Selecting **View all** opens a side panel with all the policy details.

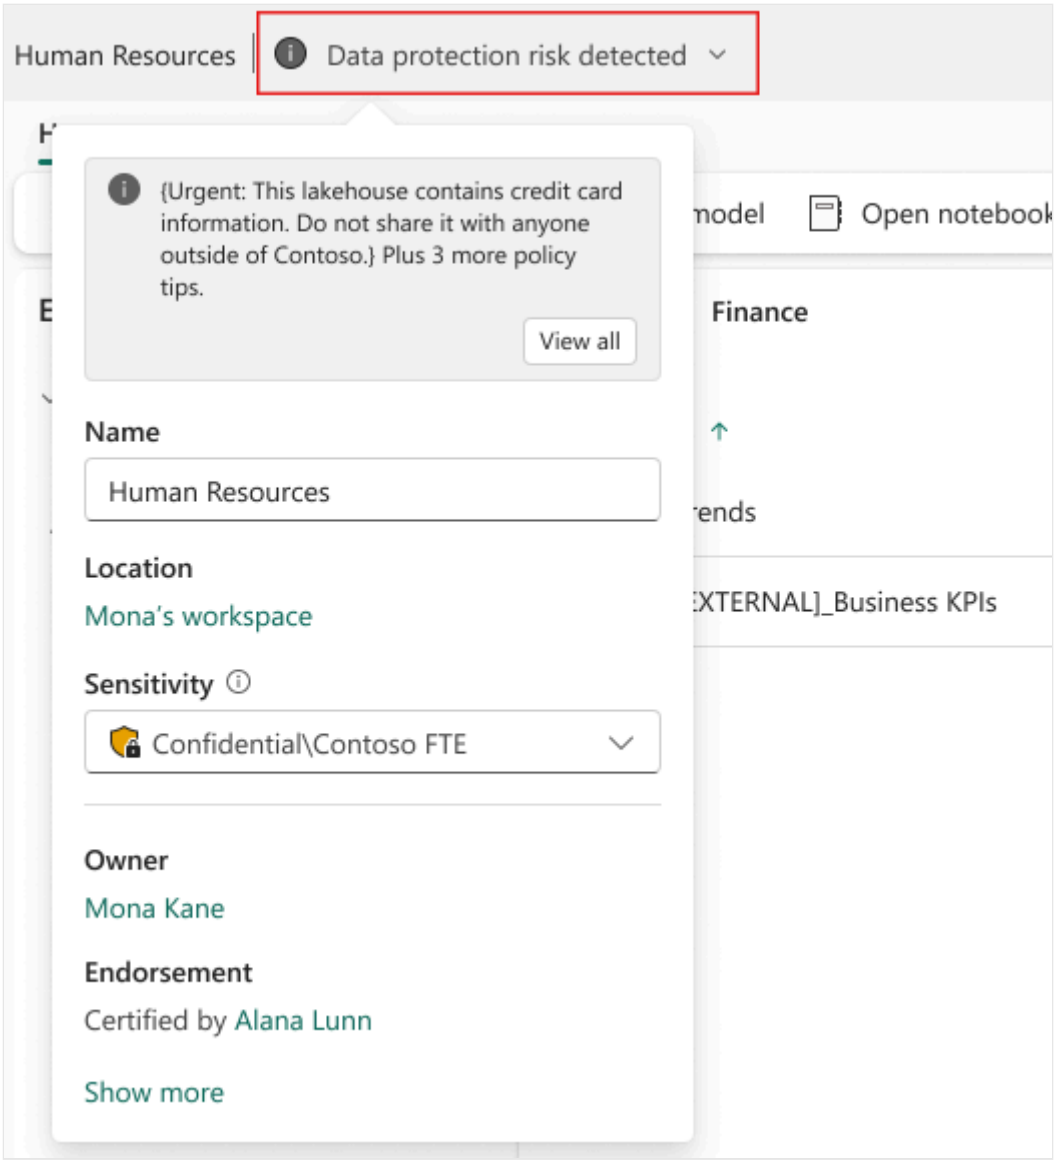

If you enable alerts in the policy, an alert records on the data loss prevention **Alerts** page in the Microsoft Purview portal. If configured, an email is sent to administrators and/or specified users. For more information, see Monitor and manage DLP policy violations.

Supported actions

When Fabric items or semantic models meet the conditions in a DLP policy, the actions defined in the policy are applied. DLP policies for Fabric and Power BI support three actions:

- Notify the user through a policy tip.
- Generate an alert. Administrators and users can receive alert emails. Additionally, administrators can monitor and manage alerts on the **Alerts** tab in the Purview portal.
- Restrict access (preview). When you configure a policy with the restrict access action and a policy match occurs, the policy restricts access to the data owners or to members of the organization, depending on the policy configuration. All other users lose access to the item.

For information about what triggers DLP evaluation, see How DLP policies for Fabric and Power BI work.

Supported item types

DLP policies for Fabric and Power BI currently support the following item types.

- Semantic models
- Lakehouses
- Warehouses
- KQL databases
- Mirrored databases
- SQL databases
- Cosmos databases

See Considerations and limitations for exceptions.

Supported condition types

DLP policy rules for Fabric and Power BI support sensitivity labels and a subset of sensitive info types (see considerations and limitations) as conditions. Learn more about Purview Sensitive Info Types and confidence levels.

Configure a DLP policy for Fabric and Power BI

For information about creating a DLP policy for Fabric or Power BI, see Help prevent sharing Power BI reports with credit card numbers and adapt it to your own scenario.

Gain visibility into DLP activity and adoption

After deploying DLP policies, administrators can use Fabric governance experiences in the OneLake catalog to understand how data loss prevention policies are being used across their organization. These insights can help organizations:

- Identify workspaces and items evaluated by DLP policies.
- Understand where sensitive information is being detected.
- Track policy adoption across business units and workspaces.
- Prioritize remediation of high-risk data assets.
- Measure the effectiveness of DLP investments and governance programs.

Together, DLP insights help security and compliance teams move beyond individual policy violations and understand their overall data protection posture.

Considerations and limitations

- DLP policies apply to workspaces. Only workspaces hosted in Fabric or Premium capacities are supported. For more information, see Microsoft Fabric concepts and licenses.
- DLP policy templates aren't supported for Fabric DLP policies. When creating a DLP policy for Fabric, choose the custom policy option.
- DLP policies for Fabric aren't supported for sample semantic models, streaming datasets, or semantic models that connect to their data source via DirectQuery or live connection. This limitation includes semantic models with mixed storage, where some of the data comes via import-mode and some comes via DirectQuery.
- DLP policies for Fabric apply only on data in tables stored in Delta format.
- Delta tables are treated as text files for the purpose of scanning considerations.
- DLP policies for Fabric support all the primitive Delta types except *timestamp_ntz*.
- DLP policies for Fabric aren't supported for the following Delta Parquet data types:
  - Binary, timestamp_ntz, Struct, Array, List, Map, Json, Enum, Interval, Void.
  - Data with LZ4, Zstd, and Gzip compression codecs.
- Advanced classifiers, such as exact data match (EDM), trainable classifiers, credential classifiers, and named entities aren't supported by DLP for Fabric. If you select an advanced classifier in the condition of a policy, the policy returns no results even if the item contains data that satisfies the advanced classifier. Other classifiers specified in the policy return results if any.
- DLP policies for Fabric aren't supported in the China North region. See How to find the default region for your organization to learn how to find your organization's default data region.
- Onboarding a new tenant to DLP can take a few hours, depending on the number of supported workspaces that are being onboarded.

See also

- Learn about data loss prevention
- Sensitivity labels in Power BI
- Audit schema for sensitivity labels in Power BI
- Configure a DLP policy for Fabric.
- Respond to a DLP policy violation in Fabric.
- Monitor and manage DLP policy violations

- Learn about data loss prevention

Last updated on 07/22/2026

Configure data loss prevention policies for Fabric

Data loss prevention policies for Fabric help organizations protect their sensitive data by detecting upload of sensitive data in supported item types. When a policy violation occurs, data owners can see this indicated, and alerts can be sent to data owners and security admins, and violations can be investigated. For more information, see Get started with Data loss prevention policies for Fabric and Power BI.

This article describes how to configure Purview data loss prevention (DLP) policies for Fabric. The target audience is compliance administrators who are responsible for data loss prevention in their organization.

Prerequisites

The account you use to create DLP policies must be a member of one of these role groups

- Compliance administrator
- Compliance data administrator
- Information Protection
- Information Protection Admin
- Security administrator

SKU/subscriptions licensing

Before you get started with DLP for Fabric and Power BI, you should confirm your Microsoft 365 subscription. The admin account that sets up the DLP rules must be assigned one of the following licenses:

- Microsoft 365 E5
- Microsoft 365 E5 Compliance
- Microsoft 365 E5 Information Protection & Governance
- Purview capacities

Configure a DLP policy for Fabric

1. Open the data loss prevention policies page in the Microsoft Purview portal and select + **Create policy**.

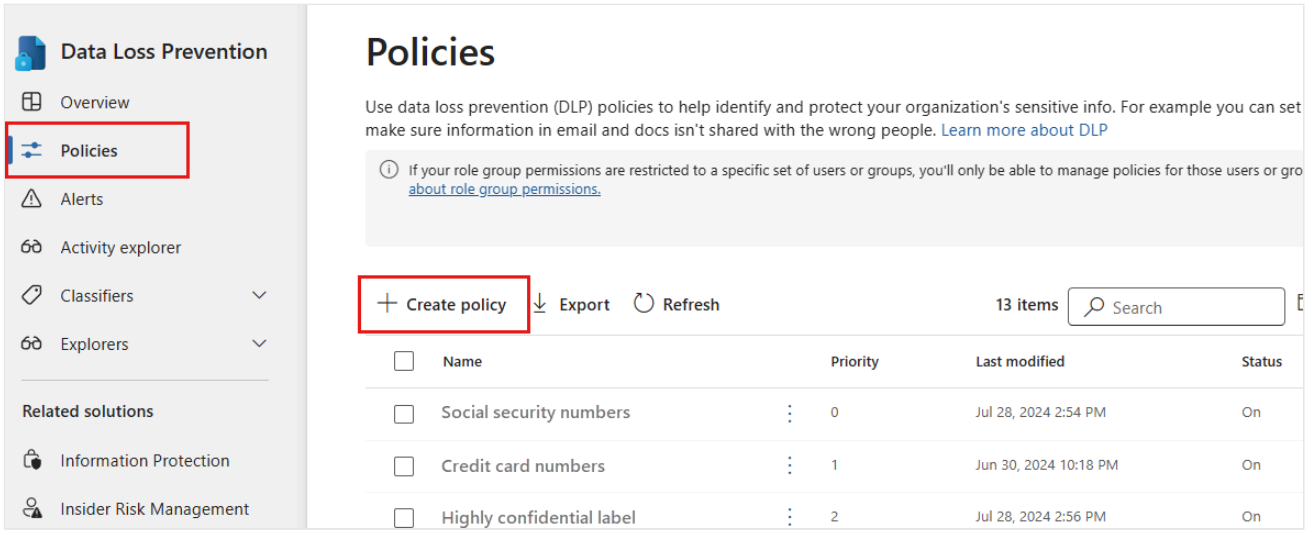

> **Note**
>
> The + **Create policy** option is only available if the prerequisites have been met.

2. Choose the **Custom** category and then the **Custom policy** template. When done, select **Next**.

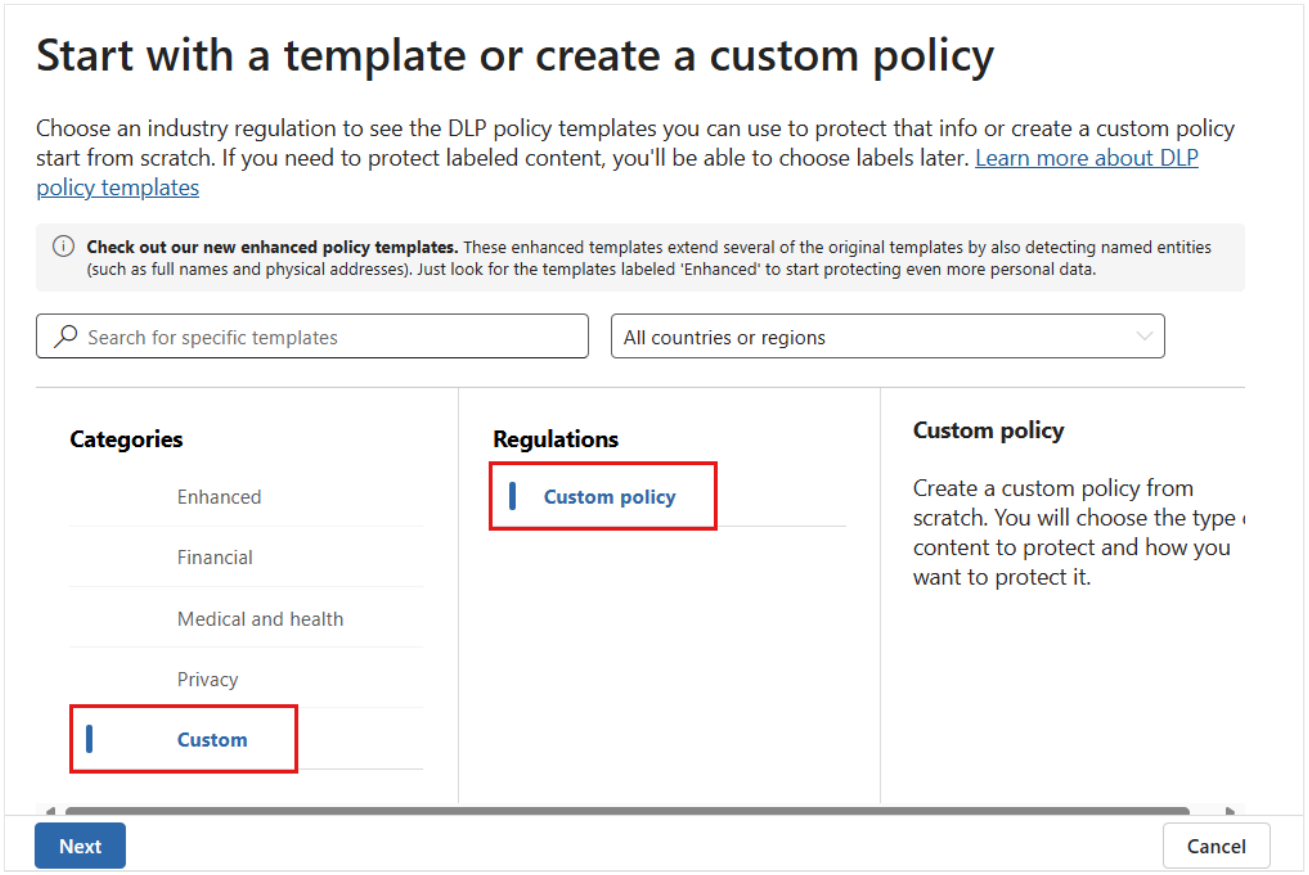

> **Note**
>
> No other categories or templates are currently supported.

3. Name the policy and provide a meaningful description. When done, select **Next**.

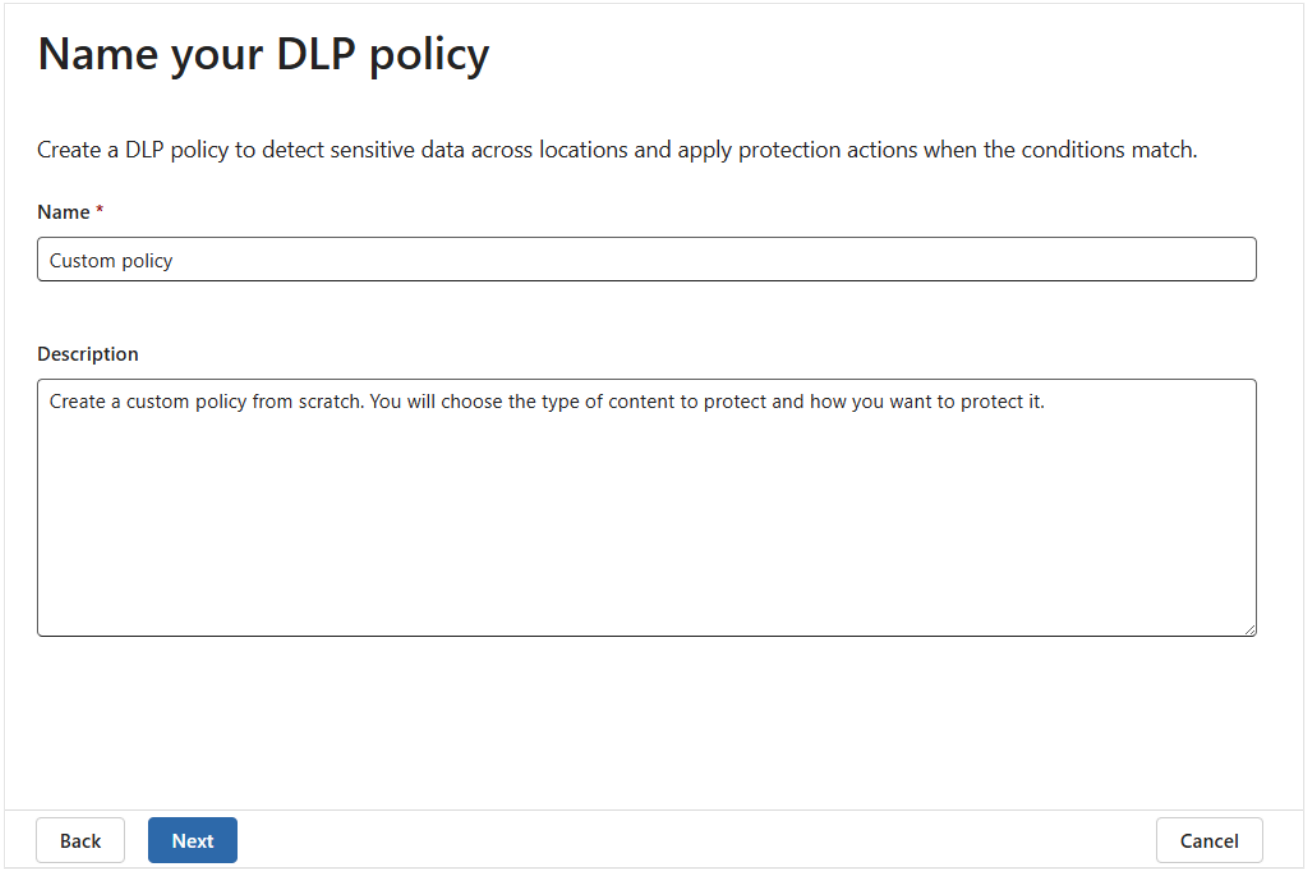

4. Select **Next** when you get to the Assign admin units page. Admin units are not supported for DLP in Fabric and Power BI.

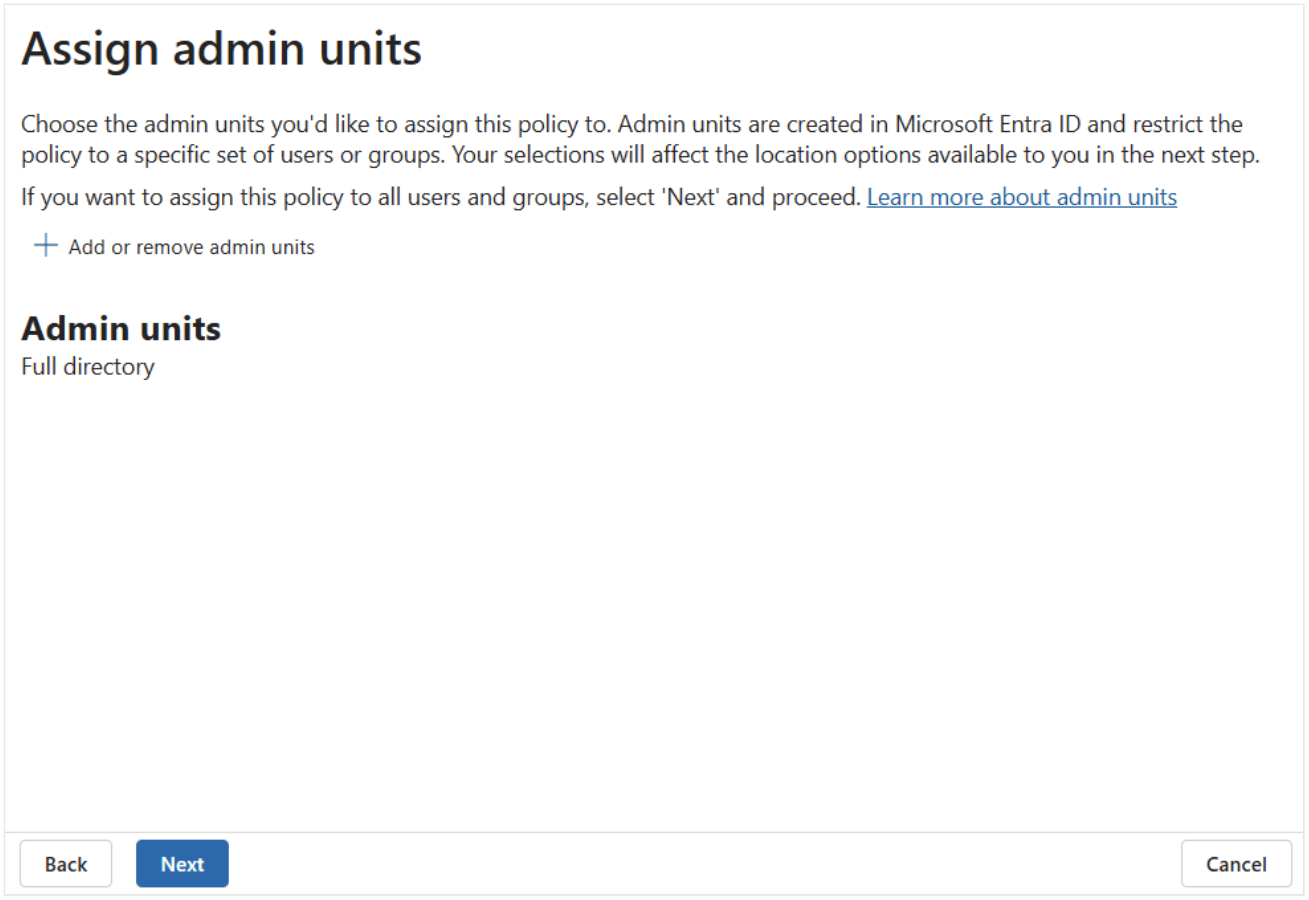

5. Select **Fabric and Power BI workspaces** as the location for the DLP policy. All other locations will be disabled, as DLP policies for Fabric and Power BI only support this location.

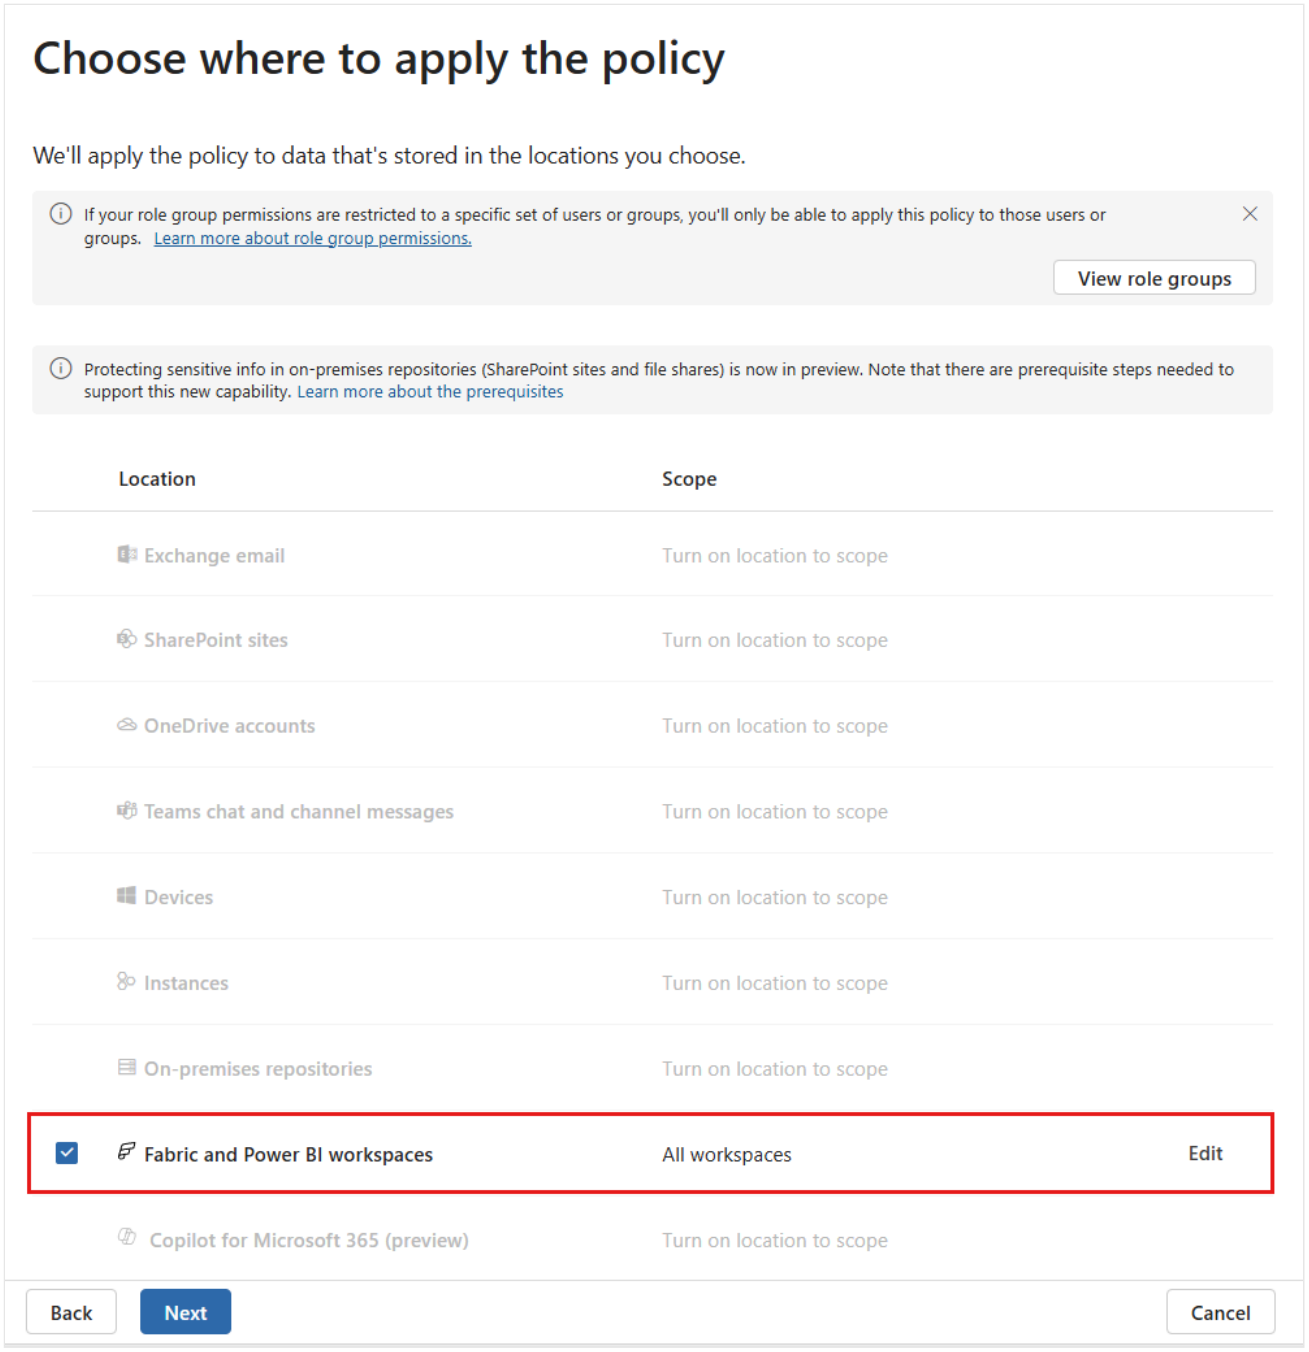

By default the policy will apply to all workspaces. However, you can specify particular workspaces to include in the policy as well as workspaces to exclude from the policy. To specify specific workspaces for inclusion or exclusion, select **Edit**.

> **Note**
>
> DLP actions are supported only for workspaces hosted in Fabric or Premium capacities.

After you've enabled Fabric and Power BI as the DLP location for the policy and chosen which workspaces the policy will apply to, select **Next**.

6. The **Define policy settings** page appears. Choose **Create or customize advanced DLP rules** to begin defining your policy.

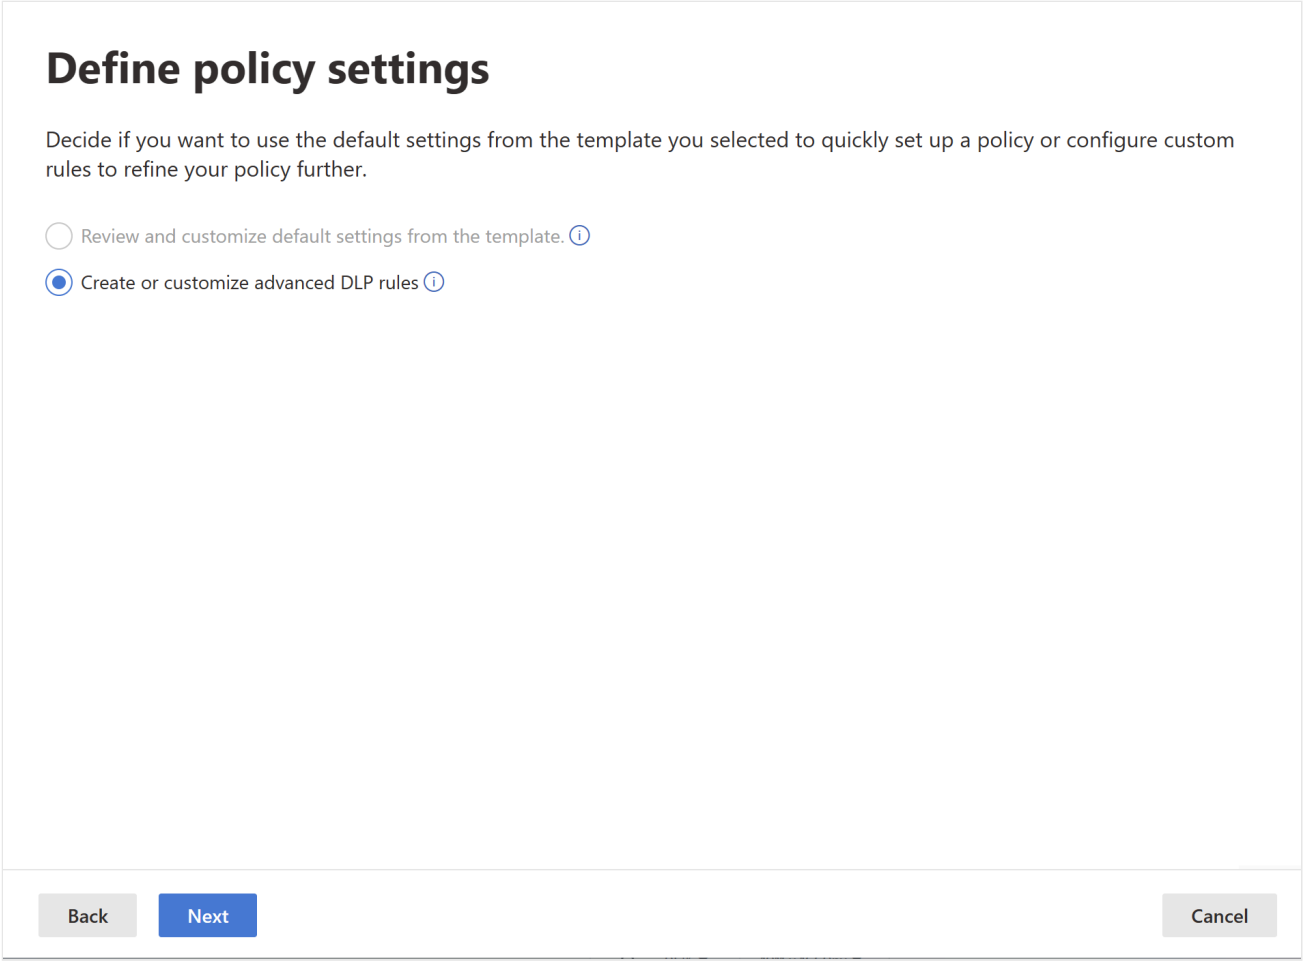

When done, select **Next**.

7. On the **Customize advanced DLP rules** page, you can either start creating a new rule or choose an existing rule to edit. Select **Create rule**.

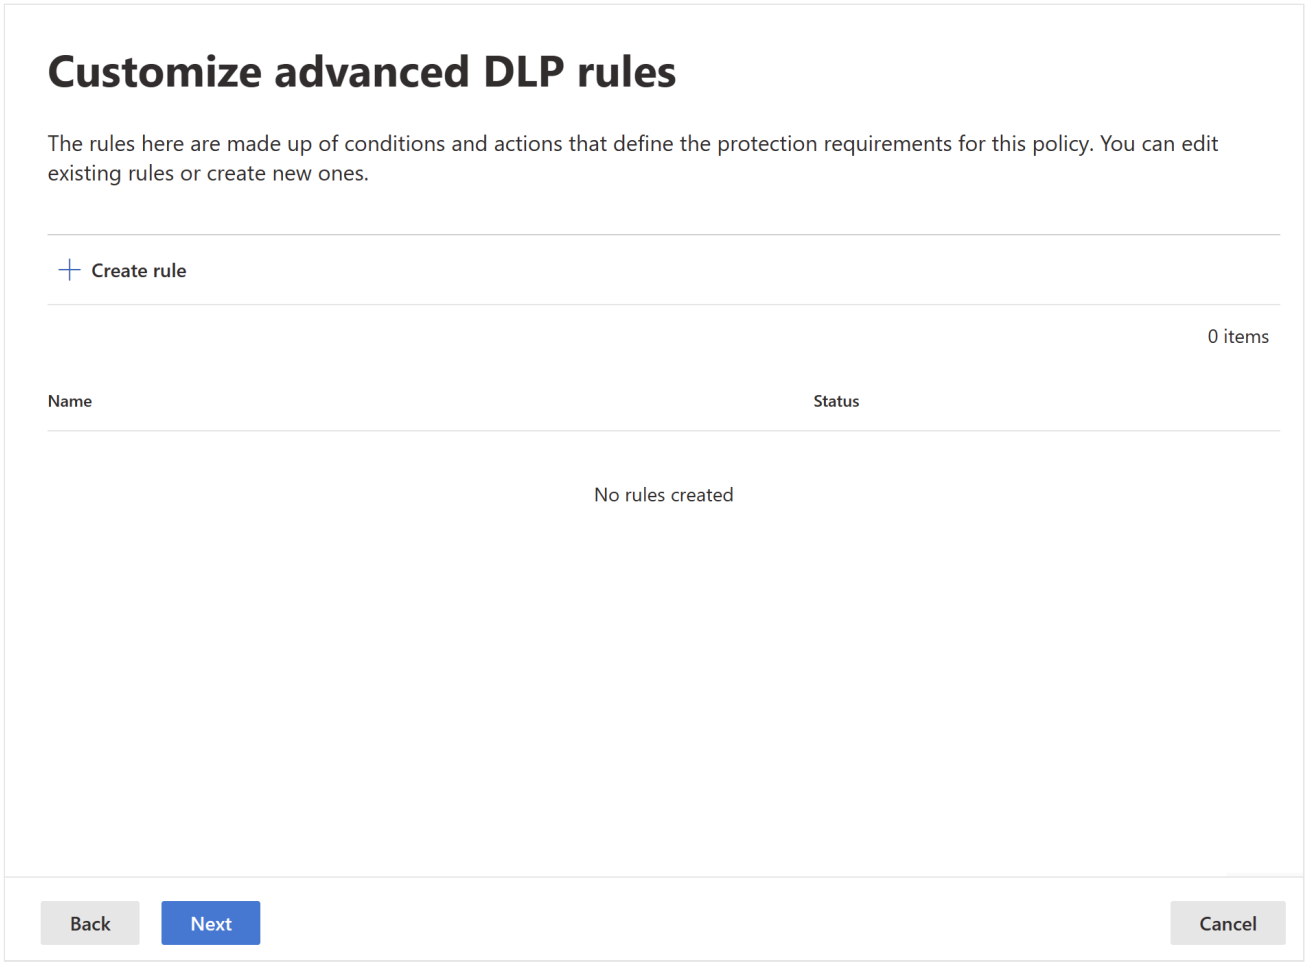

8. The **Create rule** page appears. On the create rule page, provide a name and description for the rule, and then configure the other sections, which are described following the image below.

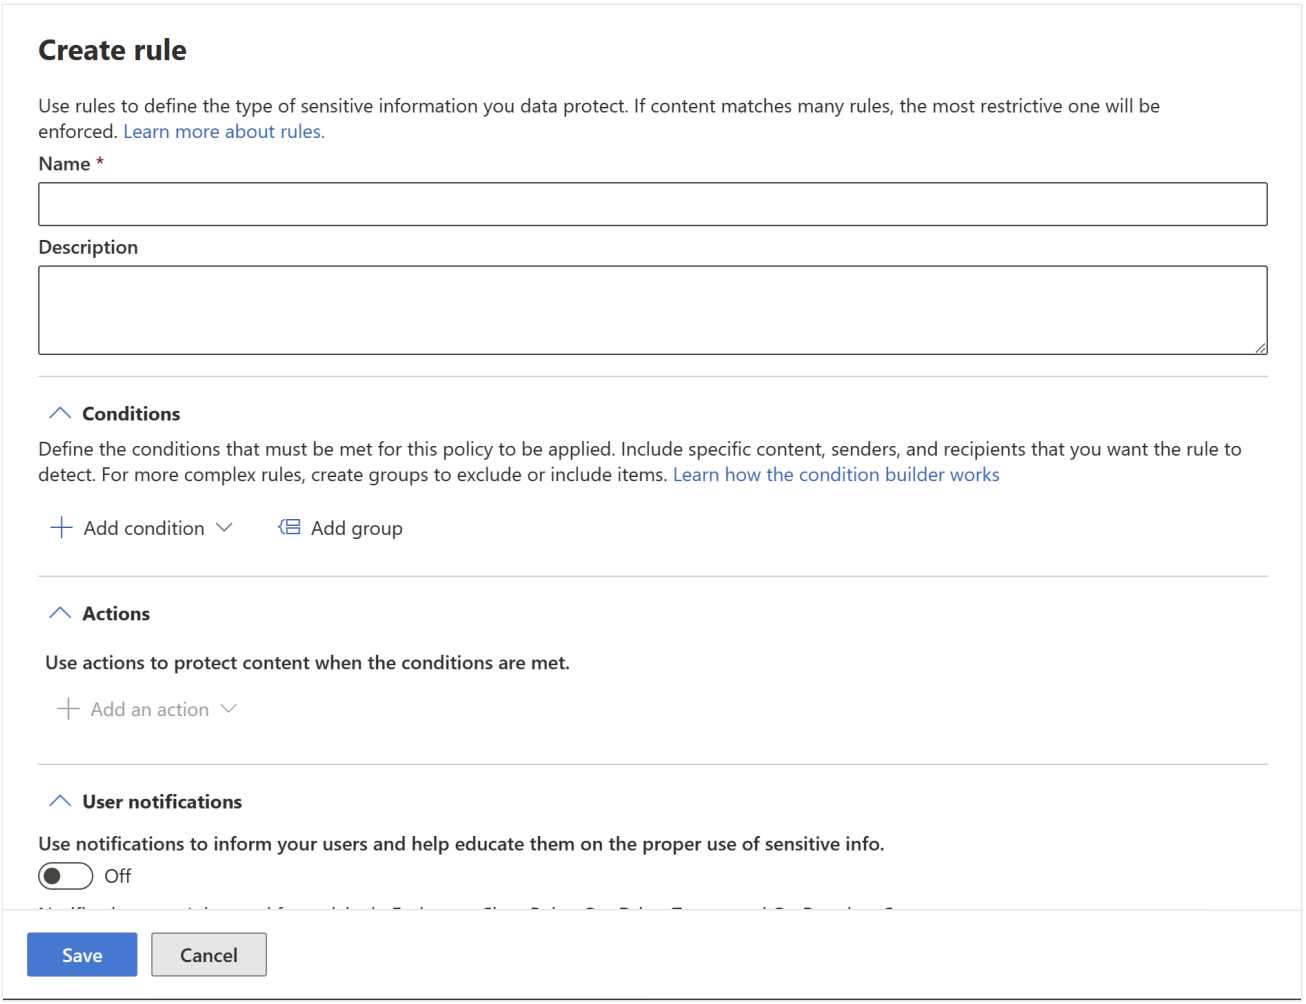

Conditions

In the condition section, you define the conditions under which the policy will apply to supported item types. Conditions are created in groups. Groups make it possible to construct complex conditions.

1. Open the conditions section. Choose **Add condition** if you want to create a simple or complex condition, or **Add group** if you want to start creating a complex condition.

   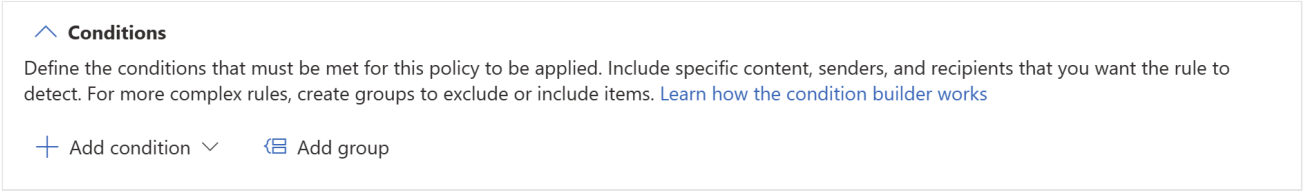

   For more information about using the condition builder, see Complex rule design.

2. If you chose **Add condition**, next choose **Content contains**, then **Add**, and then either **Sensitive info types** or **Sensitivity labels**.

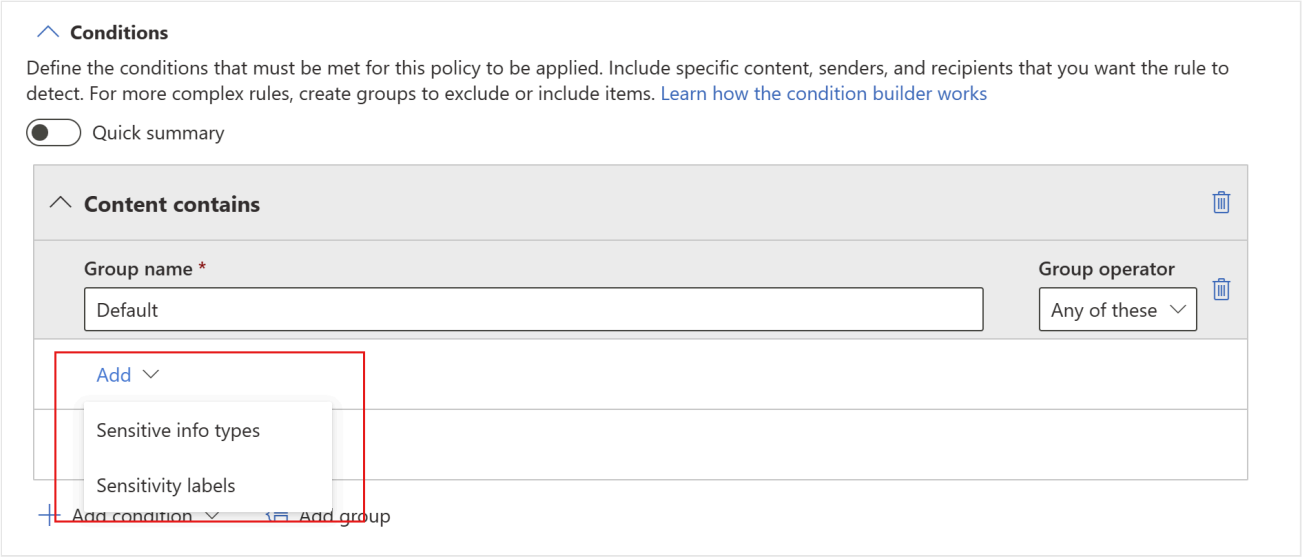

If you started with **Add group**, you'll eventually get to **Add condition**, after which you continue as described above.

When you choose either **Sensitive info types** or **Sensitivity labels**, you'll be able to choose the particular sensitivity labels or sensitive info types you want to detect from a list that will appear in a sidebar.

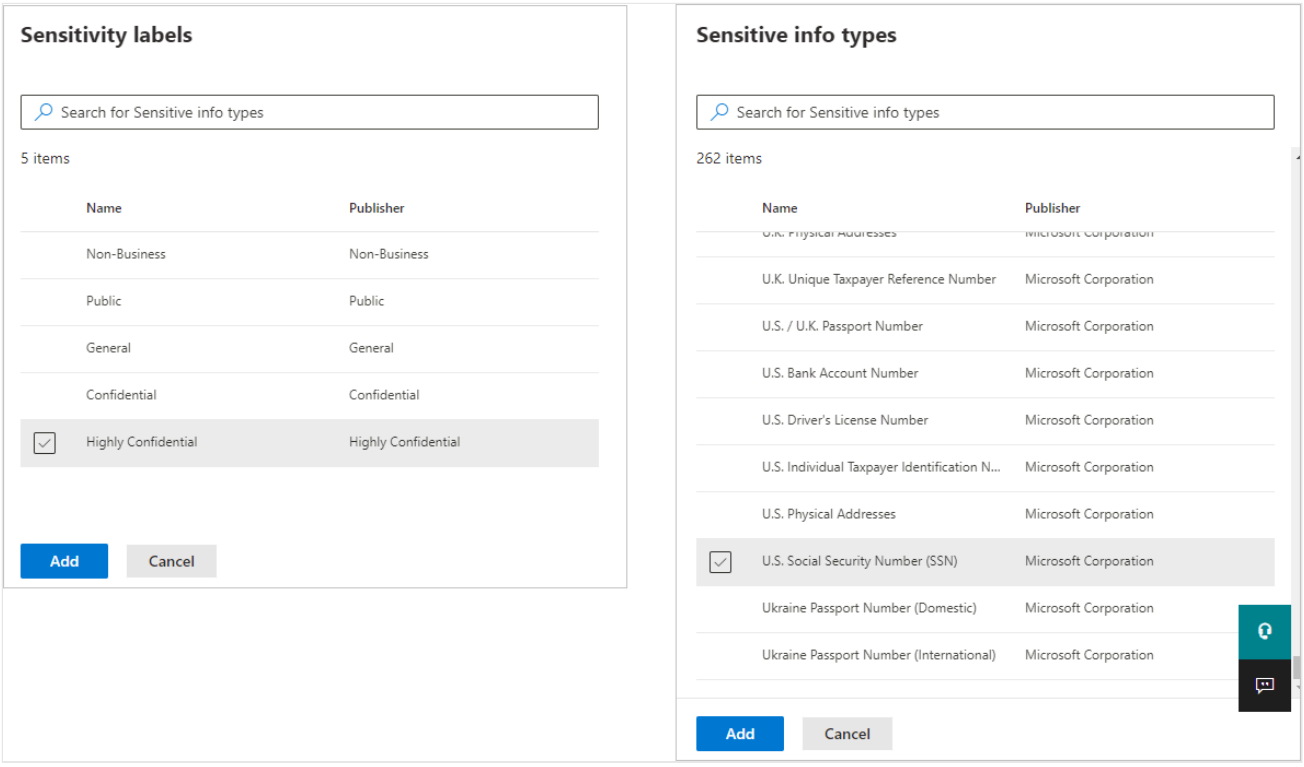

When you select a sensitive info type as a condition, you then need to specify how many instances of that type must be detected in order for the condition to be considered as met. You can specify from 1 to 500 instances. If you want to detect 500 or more unique instances, enter a range of '500' to 'Any'. You also can select the degree of confidence in the matching algorithm. Select the info button next to the confidence level to see the definition of each level.

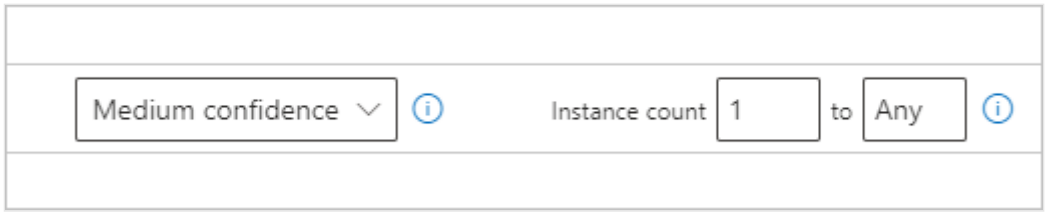

You can add additional sensitivity labels or sensitive info types to the group. To the right of the group name, you can specify **Any of these** or **All of these**. This determines whether matches on all or any of the items in the group is required for the condition to hold. If you specified more than one sensitivity label, you'll only be able to choose **Any of these**, since Fabric and Power BI items can't have more than one label applied.

The image below shows a group (Default) that contains two sensitivity label conditions. The logic Any of these means that a match on any one of the sensitivity labels in the group constitutes *true* for that group.

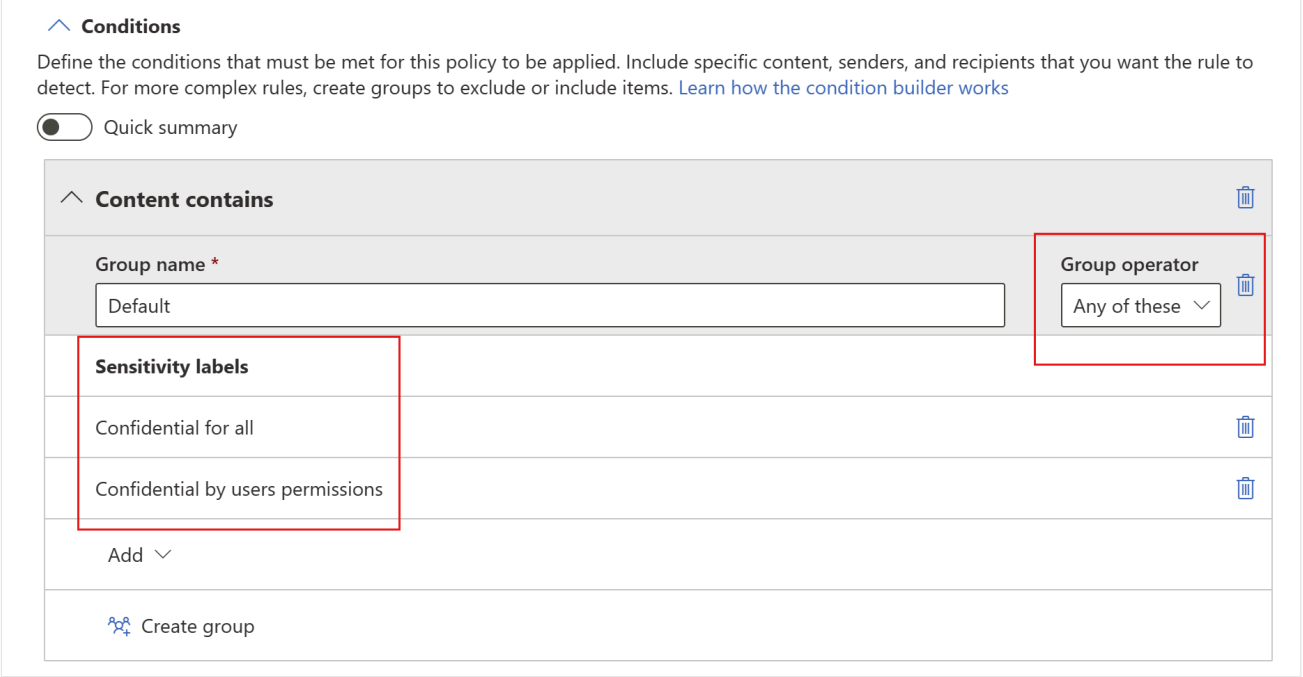

You can use the **Quick summary** toggle to get the logic of the rule summarized in a sentence.

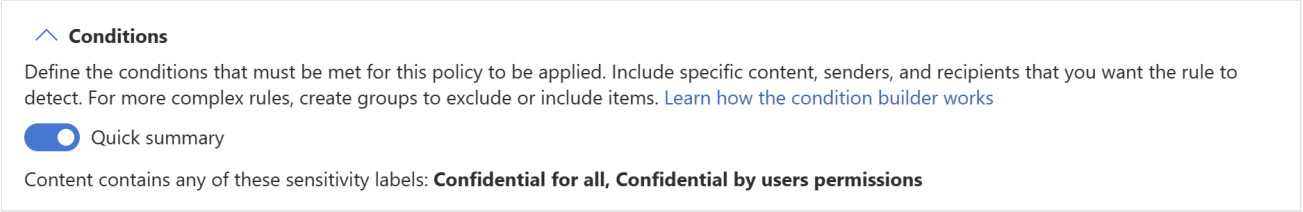

You can create more than one group, and you can control the logic between the groups with **AND** or **OR** logic.

The image below shows a rule containing two groups, joined by **OR** logic.

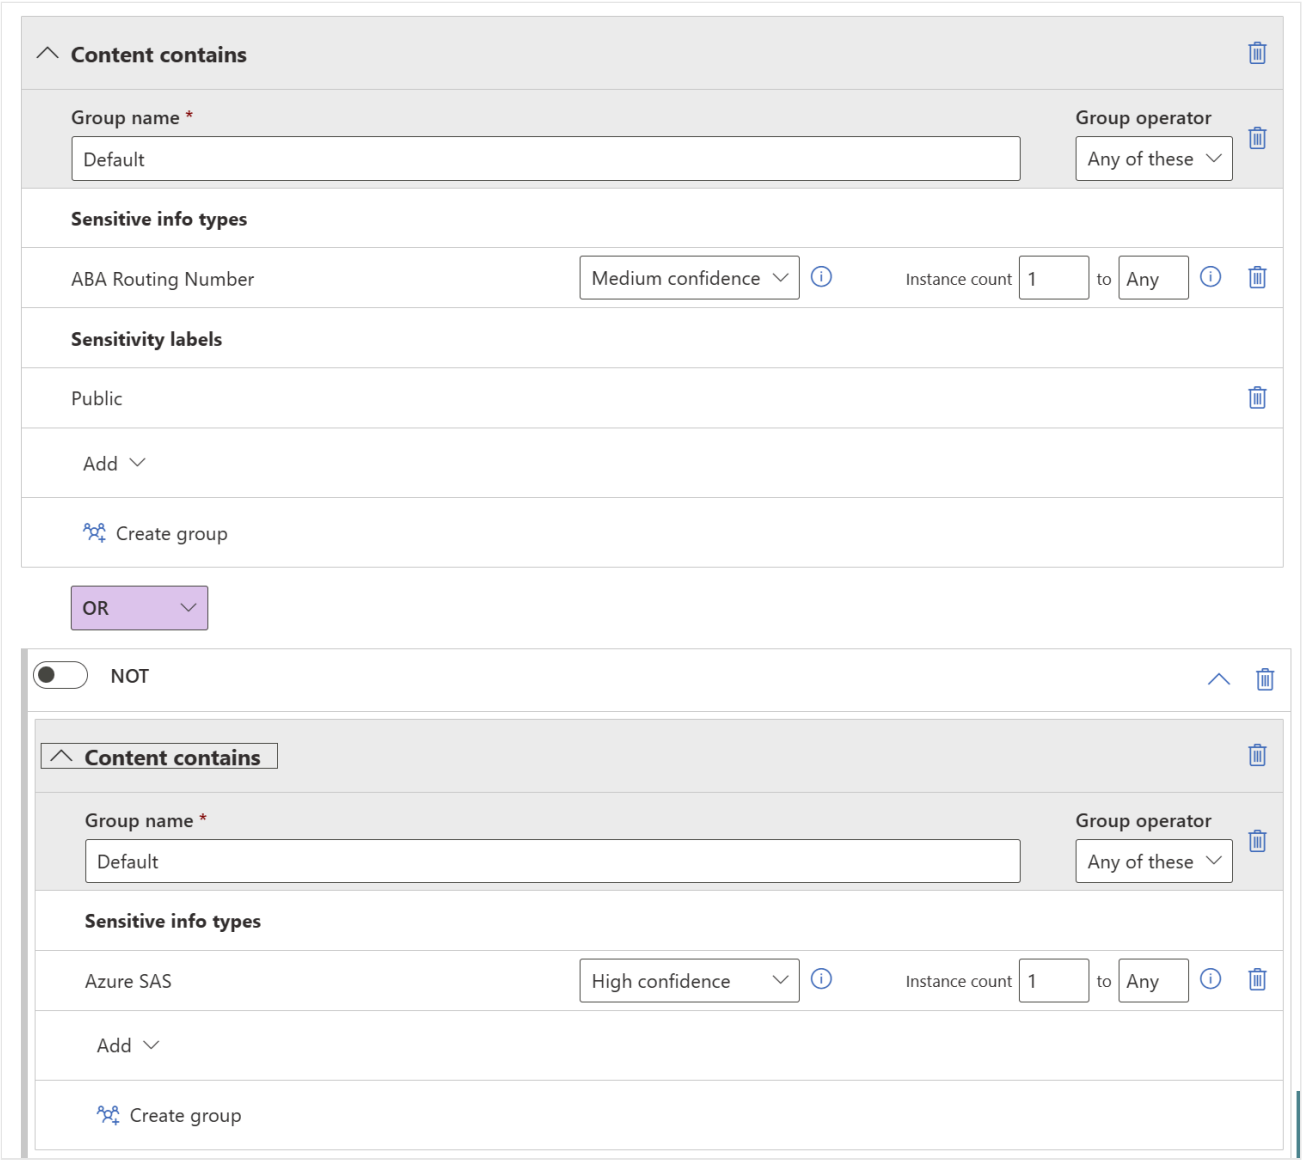

Here's the same rule shown as a quick summary.

Actions

If you want the policy to restrict access to items that trigger the policy, expand the **Restrict access or encrypt the content in Microsoft 365 locations** section and select **Block users from receiving email, or accessing shared SharePoint, OneDrive, and Teams files, and Power BI items.** Then choose whether to block everyone or only people outside your organization.

When you enable the restrict access action, user overrides are automatically allowed.

User notifications

The user notifications section is where you configure your policy tip. Turn on the toggle, select the **Notify users in Office 365 service with a policy tip or email notifications** check box, and then select the **Policy tips** checkbox. Write your policy tip in the text box that appears.

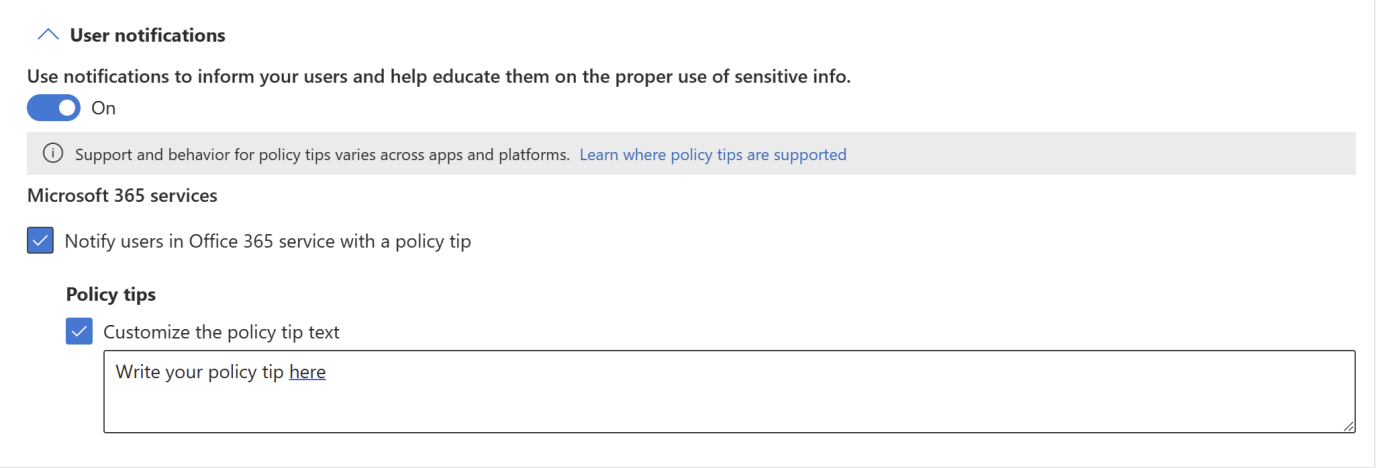

User overrides

If you enabled user notifications and selected the **Notify users in Office 365 service with a policy tip** checkbox, owners of items that have DLP policy violations (owners - that is, users with an Admin or Member role in the workspace where the item is located) will be able to respond to violations on the Data loss prevention policies side pane, which they can display from a button or a link on the policy tip. The selection of response options they have depends on the choices you make in the **User overrides** section.

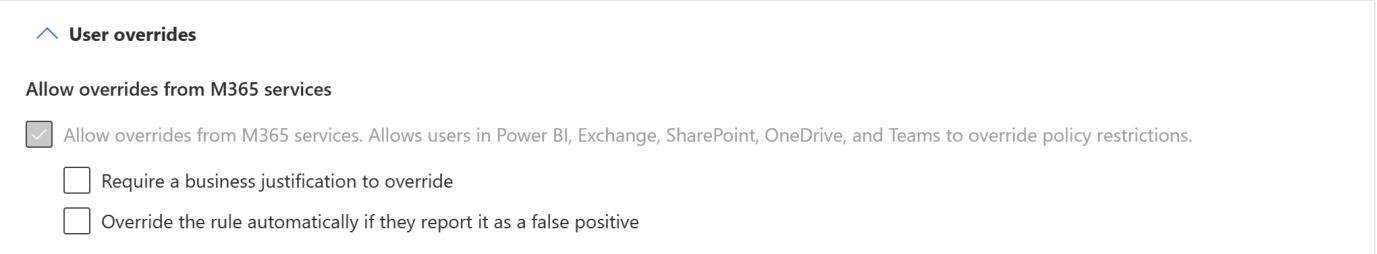

The options are described below.

- **Allow overrides from M365 services. Allows users in Power BI, Exchange, SharePoint, OneDrive, and Teams to override policy restrictions** (automatically selected when you've enabled user notifications and selected the **Notify users in Office 365 service with a policy tip** checkbox): Users will be able to either report the issue as a false positive or override the policy.

- **Require a business justification to override**: Users will be able to either report the issue as a false positive or override the policy. If they choose to override, they'll need to provide a business justification.
- **Override the rule automatically if they report it as a false positive**: Users will be able to report the issue as a false positive and automatically override the policy, or they can just override the policy without reporting it as a false positive.
- If you select both **Override the rule automatically if they report it as a false positive** and **Require a business justification to override**, users will be able to report the issue as a false positive and automatically override the policy, or they can just override the policy without reporting it as a false positive, but they'll have to provide a business justification.

Overriding a policy means that from now on the policy will no longer check the item for sensitive data.

Reporting an issue as a false positive means that the data owner believes that the policy has mistakenly identified non-sensitive data as sensitive. You can use false positives to fine tune your rules.

Any action the user takes is logged for reporting.

Incident reports

Assign a severity level that will be shown in alerts generated from this policy. Enable (default) or disable email notification to admins, specify users or groups for email notification, and configure the details about when notification will occur.

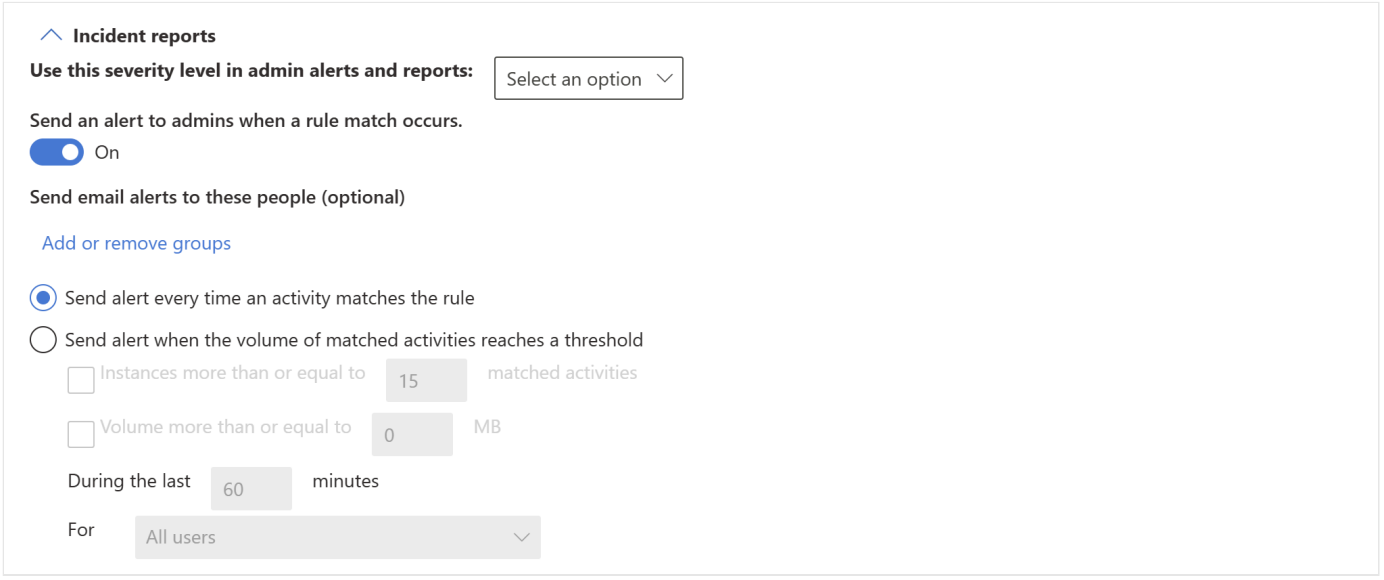

Additional options

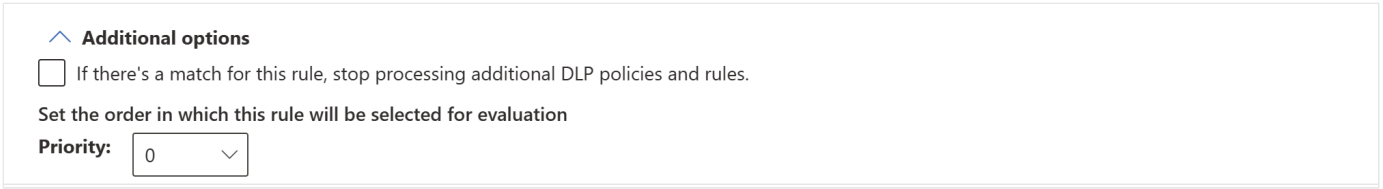

Considerations and limitations

- DLP policy templates aren't yet supported for Fabric DLP policies. When creating a DLP policy for Fabric, choose the *custom policy* option.
- Fabric DLP policy rules currently support sensitivity labels and sensitive info types as conditions.

Related content

- Create and Deploy data loss prevention policies
- Get started with Data loss prevention policies for Fabric and Power BI
- Respond to a DLP policy violation
- Monitor DLP policy violations

Last updated on 10/21/2025

Respond to a DLP policy violation in Fabric and Power BI

When a supported item type you own violates a DLP policy for Fabric and Power BI, a violation icon on the item in the OneLake catalog or workspace list view appears. Hover over the icon and select **Show full details** to display a side panel that displays the details of your item's DLP policy violations and provides options for responding to them. Alternatively:

- From a semantic model details page, select the **View all** button on the policy tip.
- In a Fabric item's edit mode, select the policy violation notice in the header and then select **View all** in the policy violation section of the flyout that appears.

This article describes the information you see on the side pane and the actions you can take regarding the violations.

View an item's DLP violations

The data loss prevention side pane lists the name of the item and all the DLP policy issues detected by a content scan of that item. You can select the item's name if you want to view the item's details.

Each DLP policy issue is shown on a card. The card shows you the policy tip, indicates what kind of sensitive data was detected, and offers actions you can take if you believe the data was falsely identified.

Violations of policies configured with the restrict access action are indicated in red and indicate who access is restricted to.

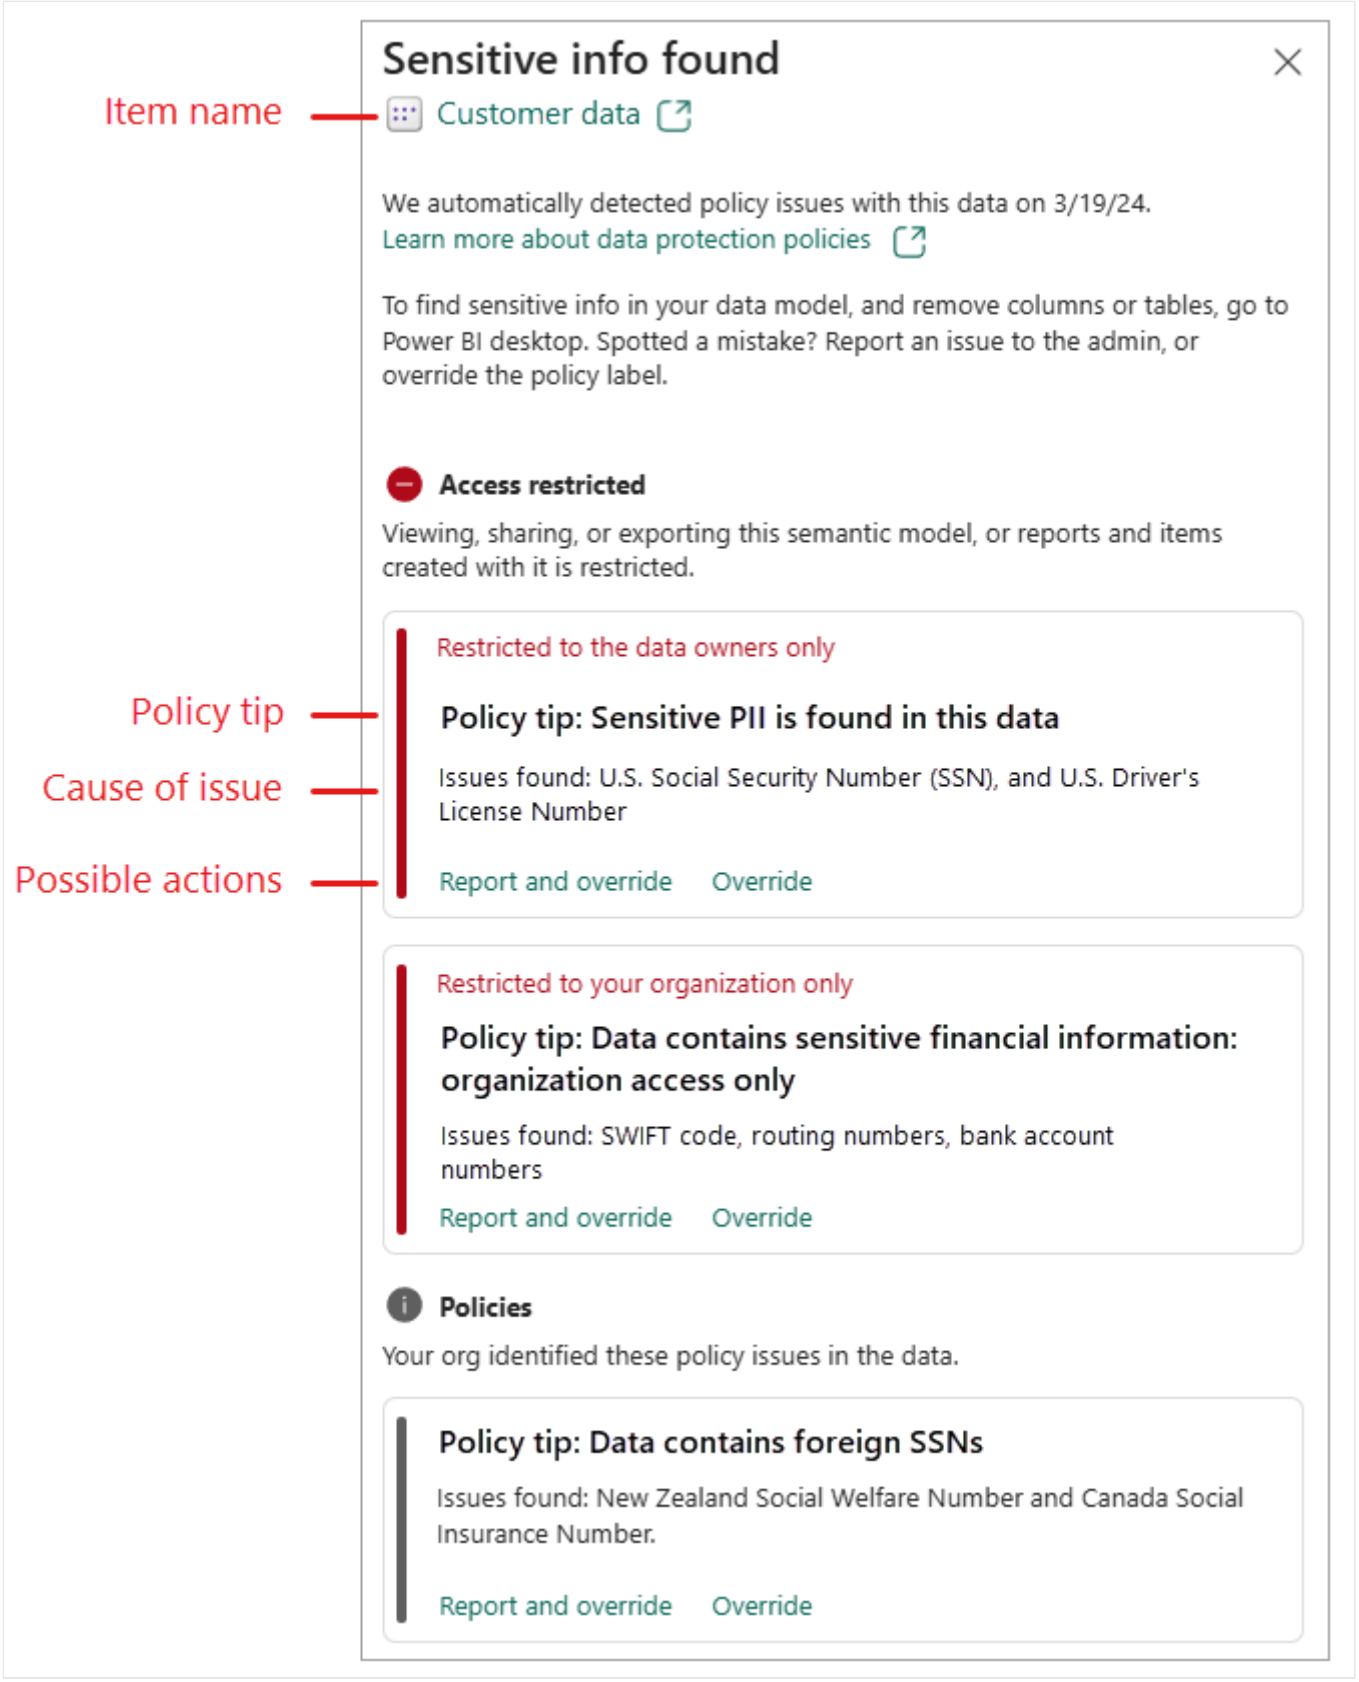

Take action on the violation

The action or combination of actions you see might vary depending on the policy configuration. The possible actions are:

- **Report an issue**: Report the issue as a false positive (that is, report that the policy has mistakenly identified nonsensitive data as sensitive).
- **Override**: Override the policy. Overriding a policy means that this policy will no longer check this item for sensitive data. Depending on the policy configuration, you might be required to provide a justification for the override.
- **Report and override**: Report the issue as a false positive and override the policy.

> **Note**
>
> Any action you take will be recorded in the audit log where it can be reviewed by security admins.
>
> The purpose of the policy tip is to flag sensitive information found in your item. If you own this data and determine that it shouldn't be in the item, when you remove the data from the item, the next time the item is evaluated, the policy tip will be removed, since the sensitive data will no longer be found.

Related content

- Get started with Data loss prevention policies for Fabric and Power BI
- Configure a DLP policy for Fabric and Power BI
- Monitor and manage DLP policy violations
- Learn about data loss prevention

Last updated on 11/19/2024

Monitor and manage DLP policy violations

This article describes how to monitor and manage DLP policy violations in Microsoft Fabric. The target audience is an organization's security and compliance teams who want to detect leakage of sensitive data and information.

Monitor and manage policy alerts

Log into the Microsoft Purview compliance portal, expand the **Data loss prevention** solution, and choose **Alerts**.

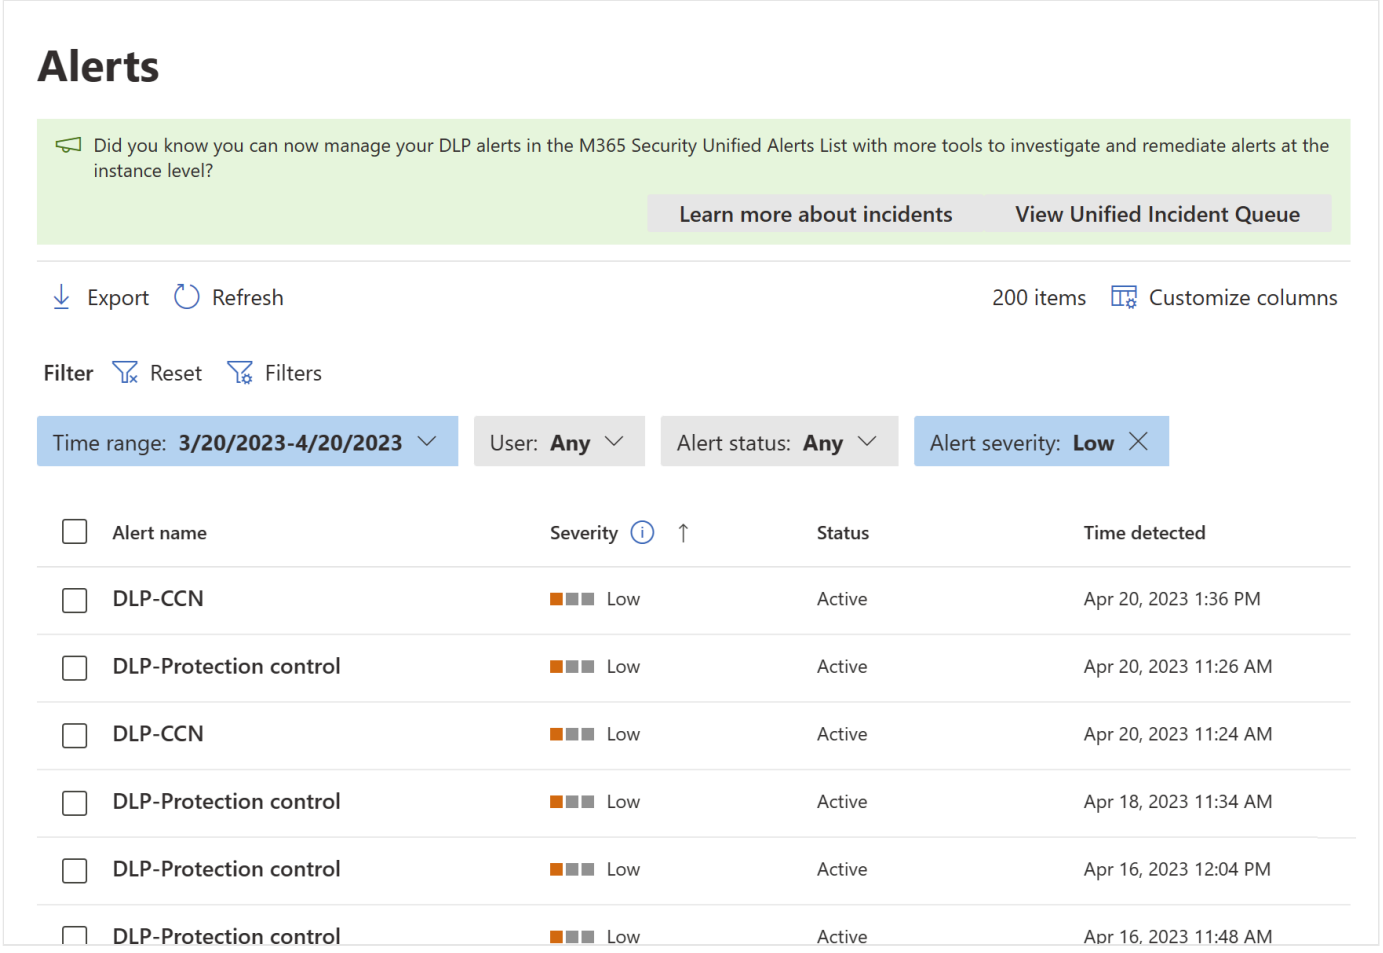

Select an alert to start drilling down to its details and to see management options.

Related content

- Get started with Data loss prevention policies for Fabric and Power BI
- Configure a DLP policy for Fabric.
- Respond to DLP policy violation in Fabric.
- Learn about data loss prevention

Last updated on 11/19/2024

Using Microsoft Defender for Cloud Apps controls in Power BI

By using Defender for Cloud Apps with Power BI, you can help protect your Power BI reports, data, and services from unintended leaks or breaches. With Defender for Cloud Apps, you can create conditional access policies for your organization's data, by using real-time session controls in Microsoft Entra ID, that help to ensure your Power BI analytics are secure. Once these policies are set, administrators can monitor user access and activity, perform real-time risk analysis, and set label-specific controls.

> **Note**
>
> Microsoft Defender for Cloud Apps is now part of Microsoft Defender XDR. For more information, see Microsoft Defender for Cloud Apps in Microsoft Defender XDR.

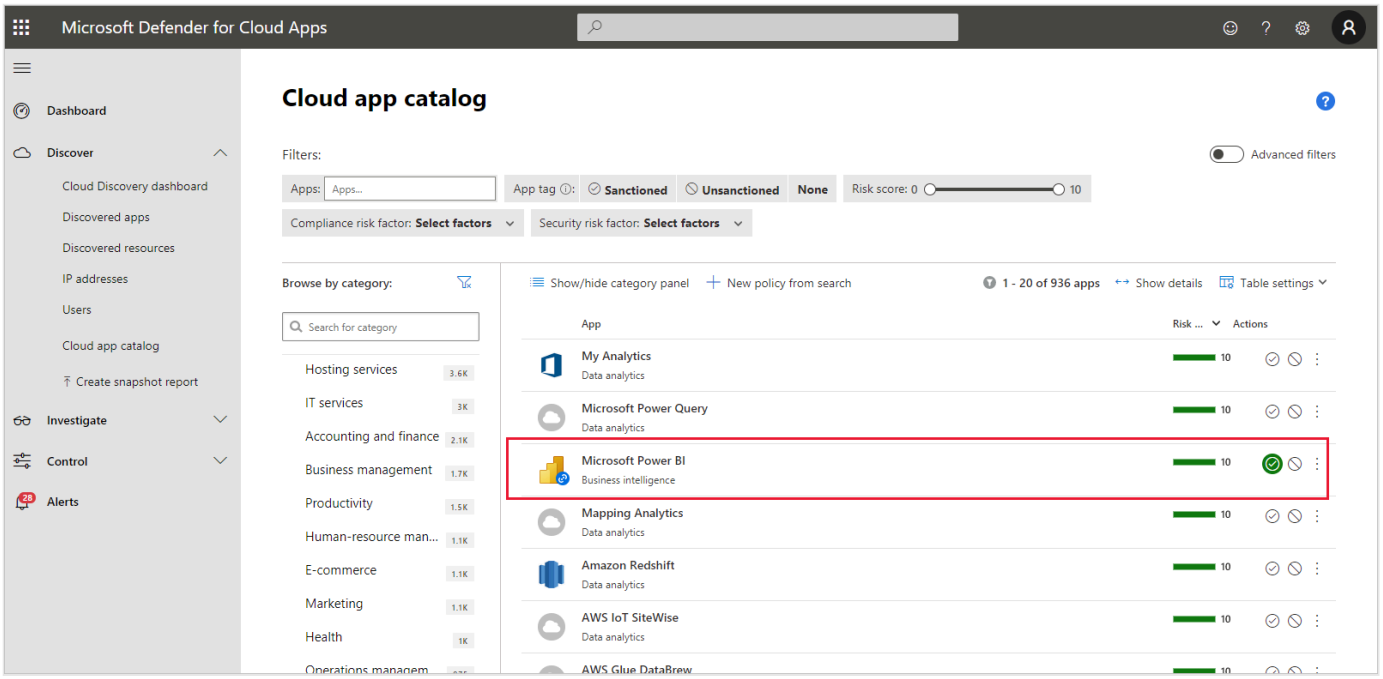

You can configure Defender for Cloud Apps for all sorts of apps and services, not only Power BI. You'll need to configure Defender for Cloud Apps to work with Power BI to benefit from Defender for Cloud Apps protections for your Power BI data and analytics. For more information about Defender for Cloud Apps, including an overview of how it works, the dashboard, and app risk scores, see the Defender for Cloud Apps documentation.

Defender for Cloud Apps licensing

To use Defender for Cloud Apps with Power BI, you must use and configure relevant Microsoft security services, some of which are set outside Power BI. In order to have Defender for Cloud Apps in your tenant, you must have one of the following licenses:

- Microsoft Defender for Cloud Apps: Provides Defender for Cloud Apps capabilities for all supported apps, part of the EMS E5 and Microsoft 365 E5 suites.
- Office 365 Cloud App Security: Provides Defender for Cloud Apps capabilities only for Office 365, part of the Office 365 E5 suite.

Configure real-time controls for Power BI with Defender for Cloud Apps

> **Note**
>
> A Microsoft Entra ID P1 license is required in order to benefit from Defender for Cloud Apps real-time controls.

The following sections describe the steps for configuring real-time controls for Power BI with Defender for Cloud Apps.

Set session policies in Microsoft Entra ID (required)

The steps necessary to set session controls are completed in the Microsoft Entra ID and Defender for Cloud Apps portals. In the Microsoft Entra admin center, you create a conditional access policy for Power BI, and route sessions used in Power BI through the Defender for Cloud Apps service.

Defender for Cloud Apps operates in a reverse-proxy architecture, and is integrated with Microsoft Entra Conditional Access to monitor Power BI user activity in real-time. The following steps are provided to help you understand the process, and detailed step-by-step instructions are provided in the linked content in each of the following steps. For a description of the whole process, see Defender for Cloud Apps.

1. Create a Microsoft Entra Conditional Access test policy
2. Sign into each app using a user scoped to the policy
3. Verify the apps are configured to use access and session controls
4. Enable the app for use in your organization
5. Test the deployment

The process for setting session policies is described in detail in Session policies.

Set anomaly detection policies to monitor Power BI activities (recommended)

You can define anomaly Power BI detection policies that can be independently scoped, so that they apply to only the users and groups you want to include and exclude in the policy. For more information, see Anomaly detection policies.

Defender for Cloud Apps has two dedicated, built-in detections for Power BI. See Built-in Defender for Cloud Apps detections for Power BI.

Use sensitivity labels from Microsoft Purview Information Protection (recommended)

Sensitivity labels enable you to classify and help protect sensitive content, so that people in your organization can collaborate with partners outside your organization, yet still be careful and aware of sensitive content and data.

For information about the process of using sensitivity labels for Power BI, see Sensitivity labels in Power BI. See the example later in this article of a Power BI policy based on sensitivity labels.

Custom policies to alert on suspicious user activity in Power BI

Defender for Cloud Apps activity policies enable administrators to define their own custom rules to help detect user behavior that deviates from the norm, and even possibly act upon it automatically, if it seems too dangerous. For example:

- **Massive sensitivity label removal.** For example, alert me when sensitivity labels are removed by a single user from 20 different reports in a time window shorter than 5 minutes.
- **Encrypting sensitivity label downgrade.** For example, alert me when a report that had a **Highly confidential** sensitivity label is now classified as **Public**.

> ⓘ **Note**
>
> The unique identifiers (IDs) of Power BI artifacts and sensitivity labels can be found using Power BI REST APIs. See Get semantic models or Get reports.

Custom activity policies are configured in the Defender for Cloud Apps portal. For more information, see Activity policies.

Built-in Defender for Cloud Apps detections for Power BI

Defender for Cloud Apps detections enable administrators to monitor specific activities of a monitored app. For Power BI, there are currently two dedicated, built-in Defender for Cloud Apps detections:

- **Suspicious share** – detects when a user shares a sensitive report with an unfamiliar (external to the organization) email. A sensitive report is a report whose sensitivity label is set to **INTERNAL-ONLY** or higher.
- **Mass share of reports** – detects when a user shares a massive number of reports in a single session.

Settings for these detections are configured in the Defender for Cloud Apps portal. For more information, see Unusual activities (by user).

Power BI admin role in Defender for Cloud Apps

A new role is created for Power BI admins when using Defender for Cloud Apps with Power BI. When you sign in as a Power BI admin to the Defender for Cloud Apps portal, you have limited access to data, alerts, users at risk, activity logs, and other information relevant to Power BI.

Considerations and limitations

Using Defender for Cloud Apps with Power BI is designed to help secure your organization's content and data, with detections that monitor user sessions and their activities. When you use Defender for Cloud Apps with Power BI, there are a few considerations and limitations you should keep in mind:

- Defender for Cloud Apps can only operate on Excel, PowerPoint, and PDF files.
- If you want to use sensitivity label capabilities in your session policies for Power BI, you need an Azure Information Protection Premium P1 or Premium P2 license. Microsoft Azure Information Protection can be purchased either standalone or through one of the Microsoft licensing suites. For more information, see Azure Information Protection pricing. In addition, sensitivity labels must have been applied on your Power BI assets.

- Session control is available for any browser on any major platform on any operating system. We recommend using the latest version of Microsoft Edge, Google Chrome, Mozilla Firefox, or Apple Safari. Power BI public API calls and other non-browser-based sessions aren't supported as part of Defender for Cloud Apps session control. For more information, see Supported apps and clients.
- If you experience sign in difficulties, such as having to sign in more than once, it could be related to the way some apps handle authentication. For more information, see the troubleshooting article Slow login.

> ⊗ **Caution**
>
> In the session policy, in the "Action" part, the "protect" capability works only if no label exists on the item. If a label already exists, the "protect" action won't apply; you can't override an existing label that has already been applied to an item in Power BI.

Example

The following example shows you how to create a new session policy using Defender for Cloud Apps with Power BI.

First, create a new session policy. In the **Defender for Cloud Apps** portal, select **Policies** on the navigation pane. Then on the policies page, select **Create policy** and choose **Session policy**.

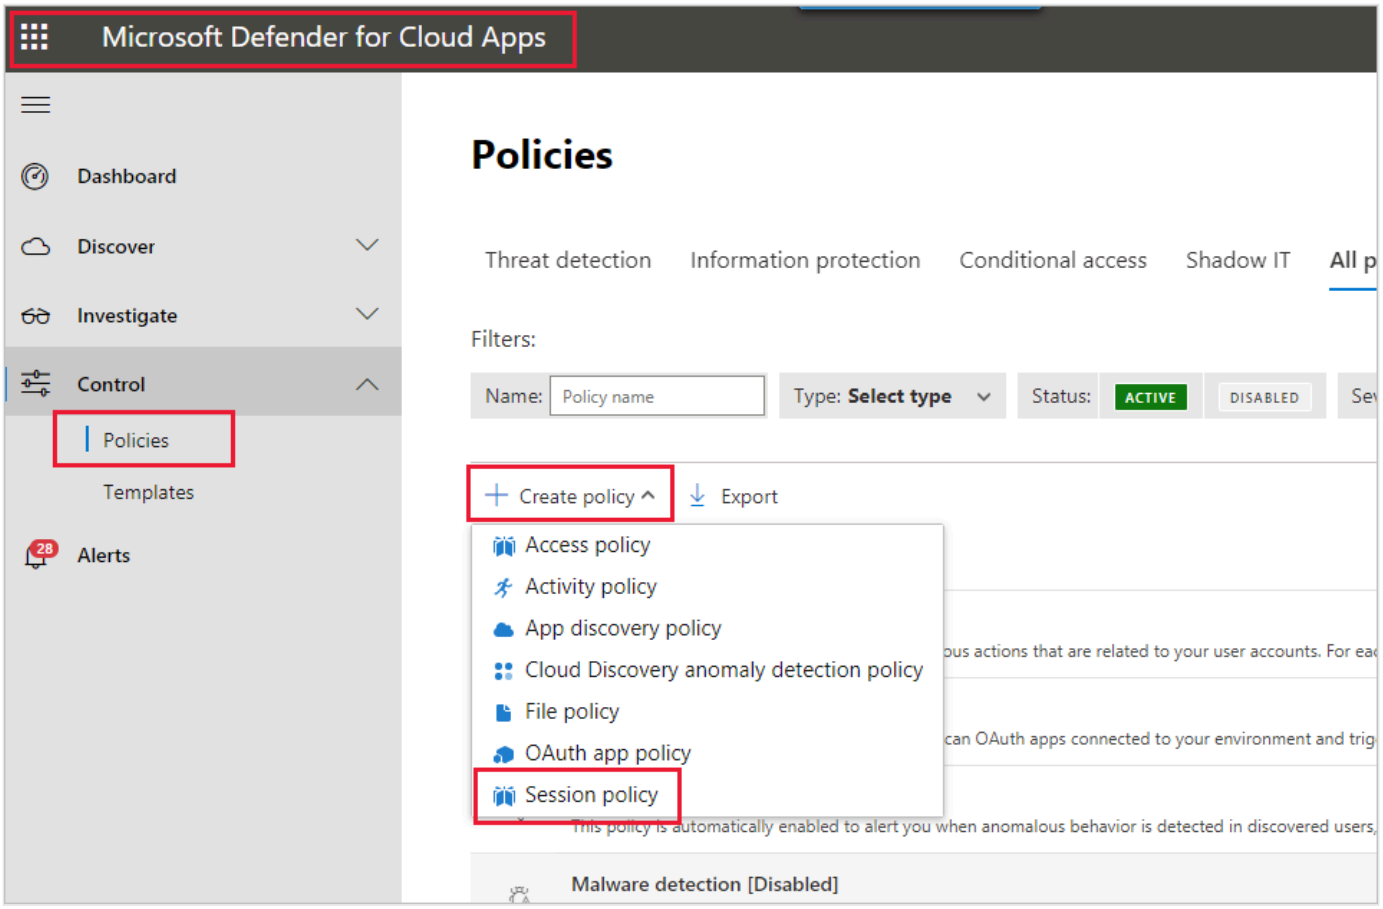

In the window that appears, create the session policy. The numbered steps describe settings for the following image.

1. In the **Policy template** dropdown, choose *No template*.
2. For **Policy name**, provide a relevant name for your session policy.
3. For **Session control type**, select *Control file download (with inspection)* (for DLP).

   For the **Activity source** section, choose relevant blocking policies. We recommend blocking un-managed and non-compliant devices. Choose to block downloads when the session is in Power BI.

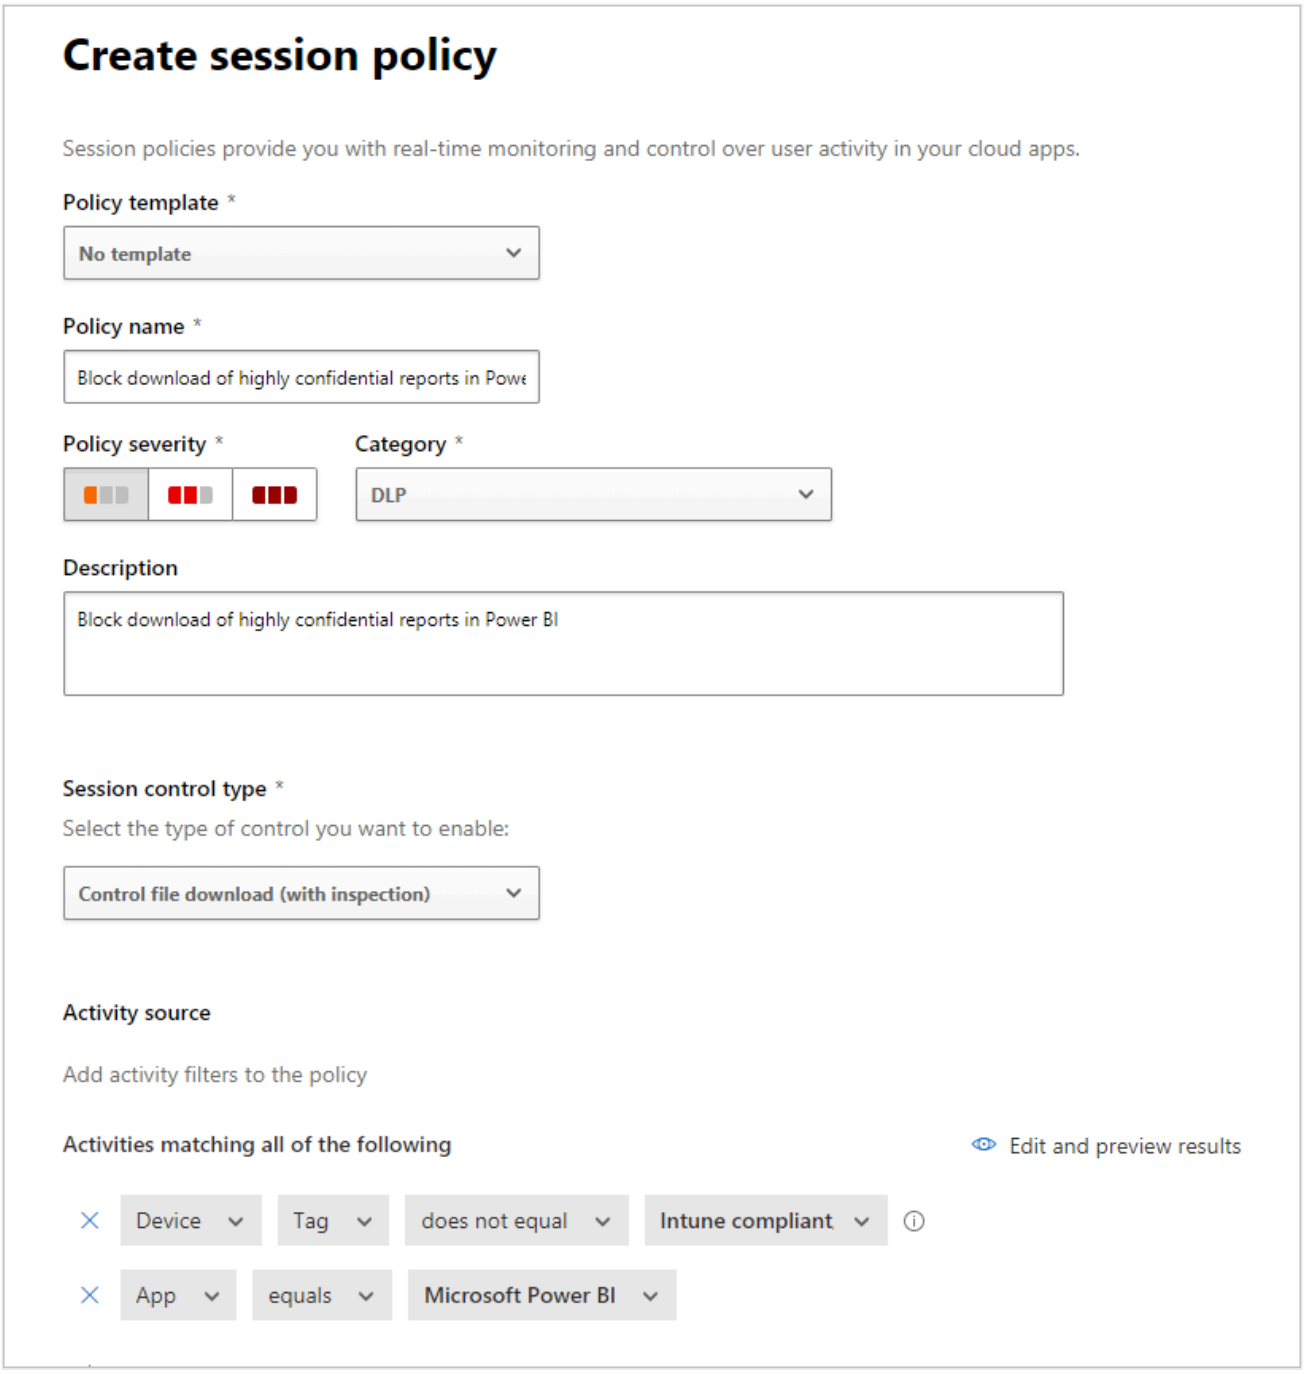

More options appear when you scroll down. The following image shows those options, with other examples.

4. Create a filter on **Sensitivity label** and choose *Highly confidential* or whatever best fits your organization.

5. Change the **Inspection method** to *none*.

6. Choose the **Block** option that fits your needs.

7. Create an alert for such an action.

Add file filters to the policy

files matching all of the following
Filters:

Sensitivity label | equals | **Highly confidential**

+ Add a filter

Inspection method

None

Actions

Select an action to be applied when user activity matches the policy.

- **Test**
  Monitor login activities
- **Block**
  A default block message is displayed when possible
  - [x] Also notify user by email
  - [ ] Customize block message
- **Protect**
  Apply sensitivity label to downloads & monitor all activities
- **Require step-up authentication** PREVIEW
  Re-evaluate Microsoft Entra Conditional Access polices based on the authentication context.
  Unpublished authentication context will not be enforced
  Configure authentication context
  No authentication context configured

- [ ] **Always apply the selected action even if data cannot be scanned**

Alerts

- [x] Create an alert for each matching event with the policy's severity

Save as default settings | Restore default settings

- [ ] Send alert as email
- [ ] Send alert as text message

Daily alert limit per policy: 5

- [ ] Send alerts to Power Automate

Create a playbook in Power Automate

Session control applies to browser-based apps.
To block access from mobile and desktop apps, create an Access policy

8. Select **Create** to complete the session policy.

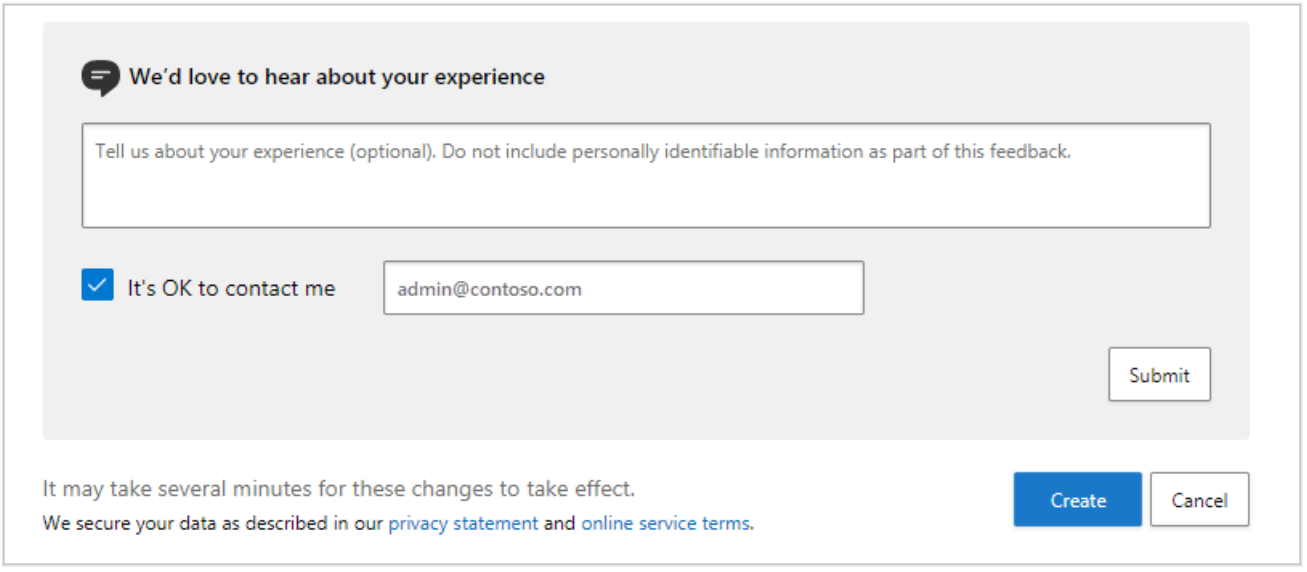

Related content

This article described how Defender for Cloud Apps can provide data and content protections for Power BI. For more information about Data Protection for Power BI and supporting content for the Azure services that enable it, see:

- Sensitivity labels in Power BI
- Enable sensitivity labels in Power BI
- How to apply sensitivity labels in Power BI

For information about Azure and security articles, see:

- Protect apps with Microsoft Defender for Cloud Apps Conditional Access App Control
- Deploy Conditional Access App Control for catalog apps with Microsoft Entra ID
- Session policies
- Learn about sensitivity labels
- Power BI implementation planning: Defender for Cloud Apps for Power BI

Last updated on 04/03/2024

External data sharing in Microsoft Fabric

Fabric external data sharing is a feature that enables Fabric users to share data from their tenant with users in another Fabric tenant. The data is shared *in-place* from OneLake storage locations in the sharer's tenant, meaning that no data is actually copied to the other tenant. Rather, this cross-tenant sharing creates a OneLake shortcut in the other tenant that points back to the original data in the sharer's tenant. Because the shortcut provides live access, changes in the source are immediately reflected in the consumer tenant; no ETL or duplication is required. Data that is shared across tenant boundaries is exposed to users in the other tenant as read-only, and can be consumed by any OneLake-compatible Fabric workload in that tenant.

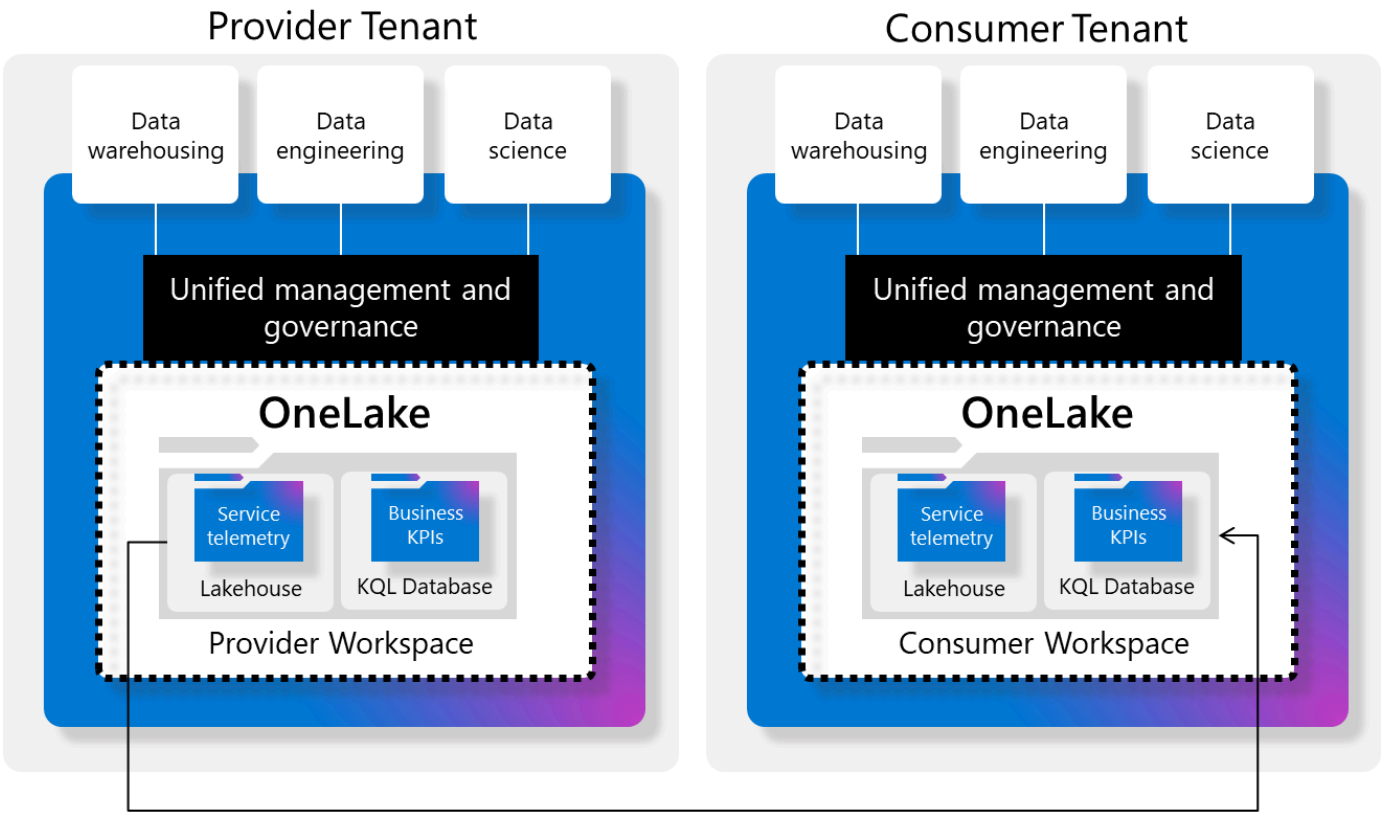

Cross-tenant OneLake data share

This external data sharing feature for Fabric OneLake data isn't related to the mechanism that exists for sharing Power BI semantic models with Microsoft Entra B2B guest users.

How does external data sharing work

As a prerequisite to external data sharing, Fabric admins need to turn on external data sharing both in the sharer's tenant and in the external tenant. Enabling external data sharing includes

specifying who can create and accept external data shares. For more information, see Enable external data sharing.

Users who are allowed to create external data shares can share data residing in tables or files within supported Fabric items, as long as they have the standard Fabric Read and Reshare permissions for the item being shared. The user creating the share invites a user from another tenant to accept the external data share. This user receives a link that they use to accept the share. Upon accepting the share, the recipient chooses a lakehouse where a shortcut to the shared data will be created. The resulting shortcut provides read-only access to the live source data; the consumer can query it directly without staging pipelines or intermediate copies.

External data share links don't work for users who are in the tenant where the external data share was created. They work only for users in external tenants. To share data from OneLake storage accounts with users in the same tenant, use OneLake shortcuts.

> ⓘ **Note**
>
> Cross-tenant data access is enabled via a dedicated Fabric-to-Fabric authentication mechanism and does not require Entra B2B guest user access.

Supported Fabric item types

The Fabric item types that can be used in external data sharing are listed below.

- **Creating an external data share (provider tenant)**: External data shares can be created for tables or files in lakehouses and warehouses, and in KQL, SQL, and mirrored databases.
- **Accepting an external data share (consuming tenant)**: Only lakehouses can be chosen as the location of the external data share shortcut when accepting an external data share. After the shortcut is created, consumer-side workloads such as Power BI reports, Spark notebooks, data science and ML experiments, and SQL analytics endpoint queries can all access the shared data.

Revoking external data shares

Any user in the sharing tenant who has Read and Reshare permissions on an externally shared item can revoke the external data share at any time using the **External data shares** tab on the manage permissions page. Revoking external data shares can have serious implications for the

consuming tenants and should be considered carefully. For more information, see Revoke external data shares.

Security considerations

Sharing data with users outside your home tenant has implications for data security and privacy that you should consider. It's important to understand the underlying flows of data sharing to better evaluate these implications. While external data sharing preserves a single version of truth (both tenants read the same live data), governance controls from the provider tenant don't flow across tenant boundaries. The consumer tenant must apply its own governance, access control, and compliance policies to the incoming shared data.

Fabric-internal security mechanisms share data across tenants. The share security mechanism grants read-only access to **any user** within the home tenant of the user that you invited to accept the share. Data is shared "in-place": no data is copied, and the data isn't even accessed until the user in the receiving tenant executes a Fabric workload over the shared data. Fabric evaluates and enforces Entra ID-based roles and permissions within the tenant where they're defined. Governance controls from the provider tenant don't automatically transfer to or execute within the consumer tenant. As a result, capabilities such as label visibility, downstream label inheritance, export encryption, and similar Microsoft Purview Information Protection and Data Loss Prevention behaviors aren't enforced in the consumer tenant. However, provider-tenant policies that restrict access to shared data continue to be enforced at the source and can prevent external consumption. After access is granted, governance, access control, and compliance policies that you define in the consumer tenant apply to the incoming share.

With this understanding in mind, be aware of the following:

- The sharer can't control who has access to the data in the consumer's tenant. For more information about how access control applied in the consumer tenant is enforced, see Information protection in Microsoft Fabric and Power BI: Access control.
- The consumer can grant access to the data to anyone, even guest users from outside the consumer's organization.
- Data might be transferred across geographic boundaries when it's accessed within the consumer's tenant.

Considerations and limitations

- **Shortcuts:** Shortcuts contained in folders that are shared via external data sharing won't resolve in the consumer tenant.

Related content

- Create an external data share
- Accept an external data share
- Manage external data shares
- Set up external data sharing on your tenant (Fabric admins)

Last updated on 08/27/2026

Create an external data share

This article shows how you can share data from your Fabric tenant's OneLake storage with users in other Fabric tenants.

Before sharing data with a user in another tenant via external data sharing, be sure to review the external data sharing security considerations.

Prerequisites

- A supported Fabric item with at least one folder or table.
- Standard Fabric read and reshare permissions on the item you're sharing.

Create an external data share

1. In a workspace or in the OneLake catalog, find the Fabric item that contains the data you want to share. See the list of supported Fabric item types.

2. Open the context menu of the selected item and choose **External data share**.

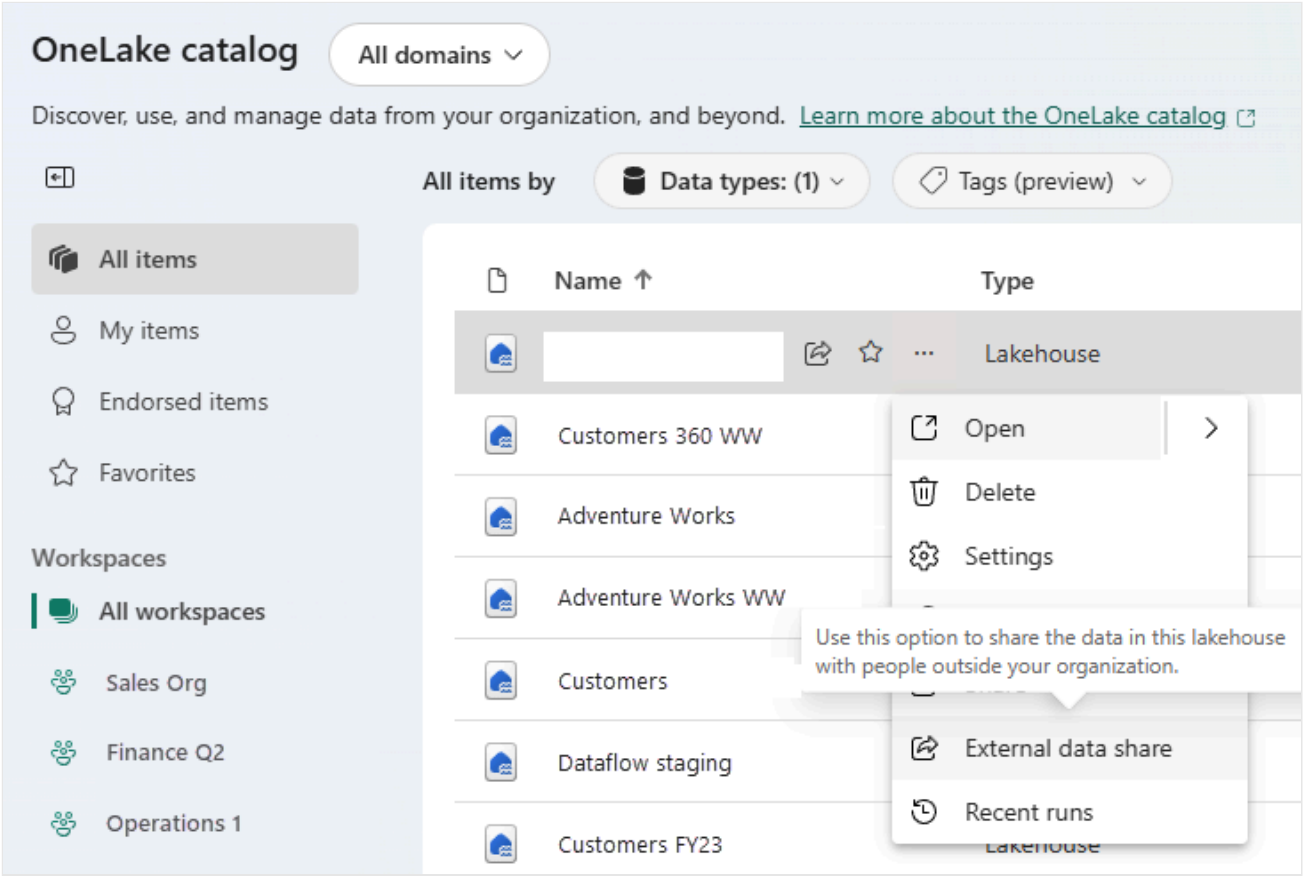

3. In the **New external data share** dialog that opens, choose the folders and/or tables to be shared and select **Save and continue**.

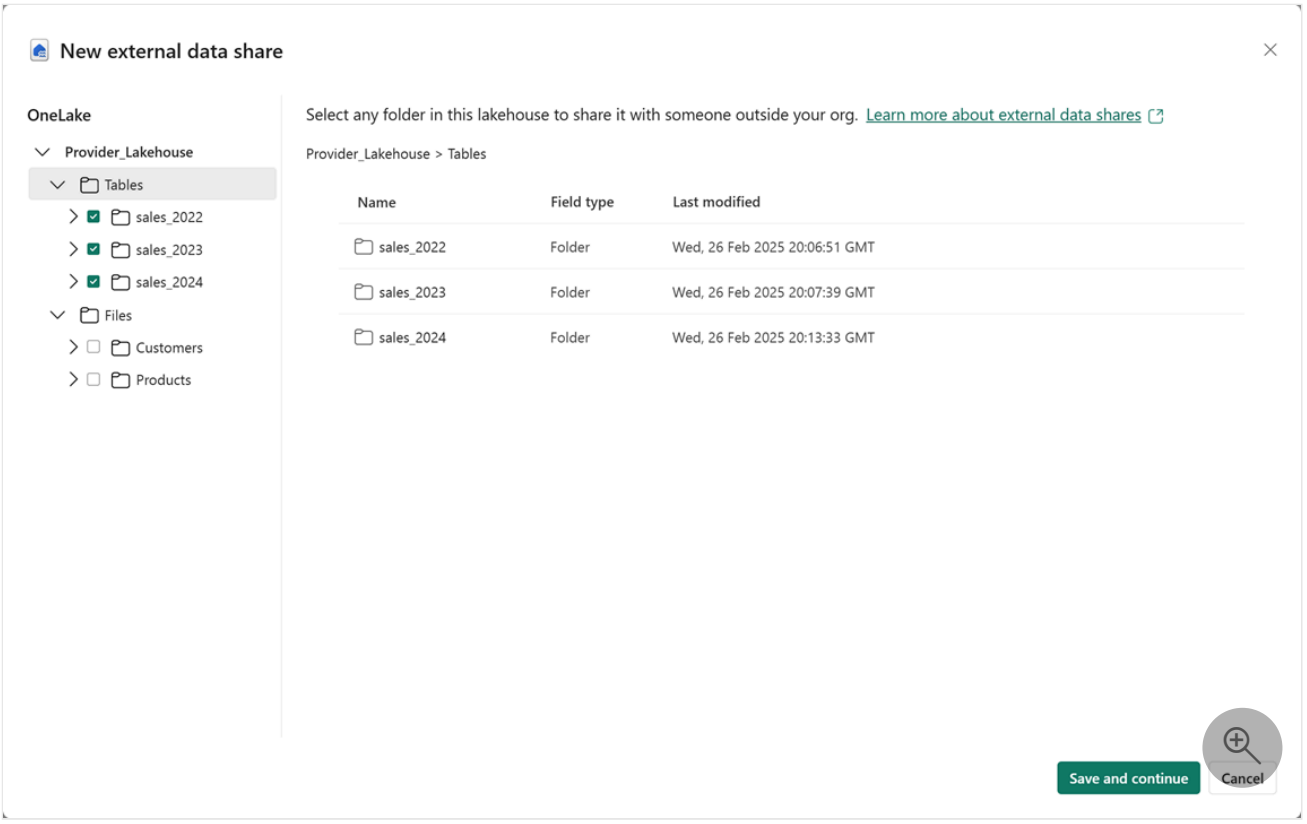

4. Enter the email address of the user you want to share the data with and select **Send**. An email will be sent to the user inviting them to accept the invitation to share. Alternatively, you can select **Copy link** and then paste it into an email that you yourself compose.

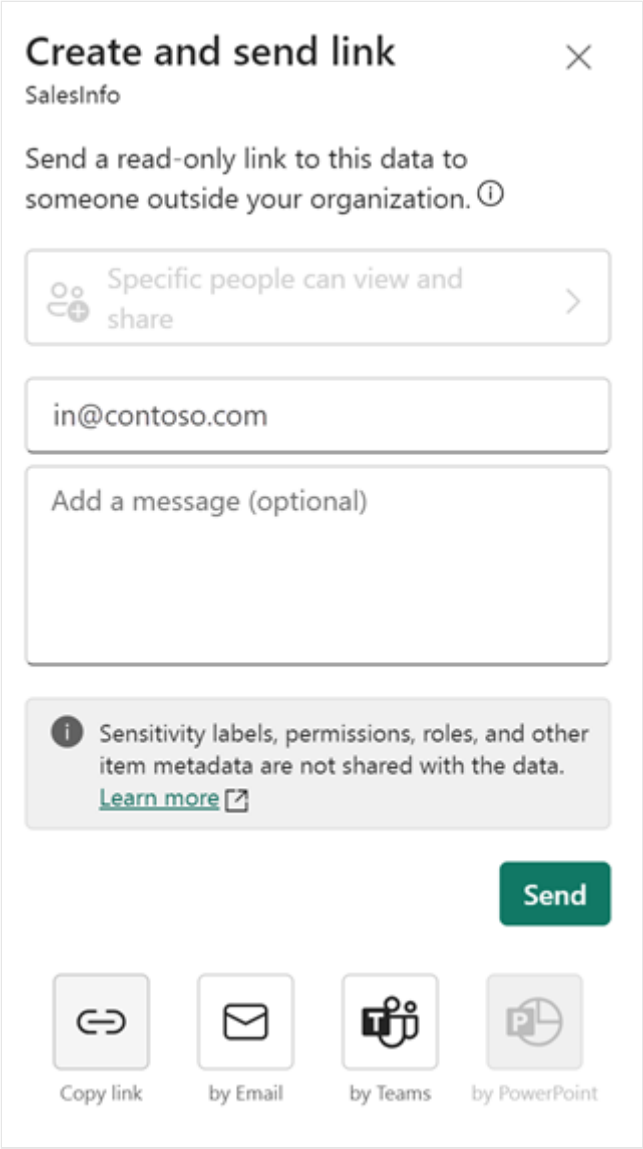

The person you invited will now be able to accept the share and access the data from within their tenant. They have 90 days to accept the invitation, after which the invitation expires.

You can revoke the external share at any time, although doing so can have serious implications for the consuming tenant. For information see Manage external data shares.

Related content

- External data sharing overview
- Accept an external data share
- Manage external data shares
- Fabric admins: Enable external data sharing

Last updated on 11/19/2024

Accept an external data share

When someone invites you to share OneLake data from their organization, you receive a link that you can use to accept the invitation and to create a shortcut to the shared data in a Fabric item of your choosing. You have 90 days to accept the invitation.

Prerequisites

To accept and use an external data share, you need standard Fabric write permissions on the lakehouse item where you want to create the shortcut to the shared data.

Accept a share

To accept a share:

1. Select the share link or paste the URL in a browser. This navigates you to a dialog that displays the name of the share and the details of the data provider's tenant.

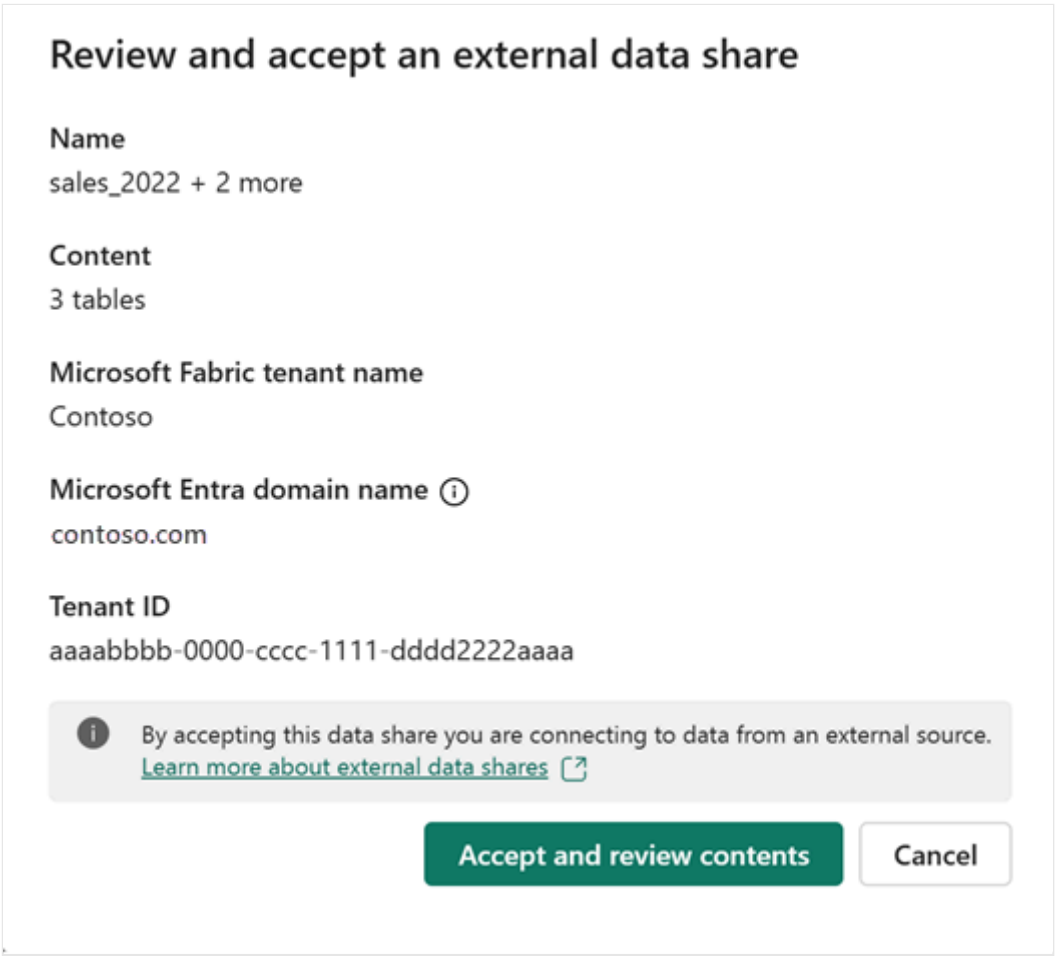

> **Note**
>
> The link only works for users who aren't in the tenant where the shared data is stored. If you're in the sharing tenant, contact the person who sent you the link and ask them to share with you using a OneLake shortcut instead. For more information, see How does external data sharing work.

2. Select **Accept and select a location**. This opens the OneLake catalog. Select a lakehouse, select **Next**, select the table or folder location in which to create the incoming share shortcuts, and then select **Apply**.

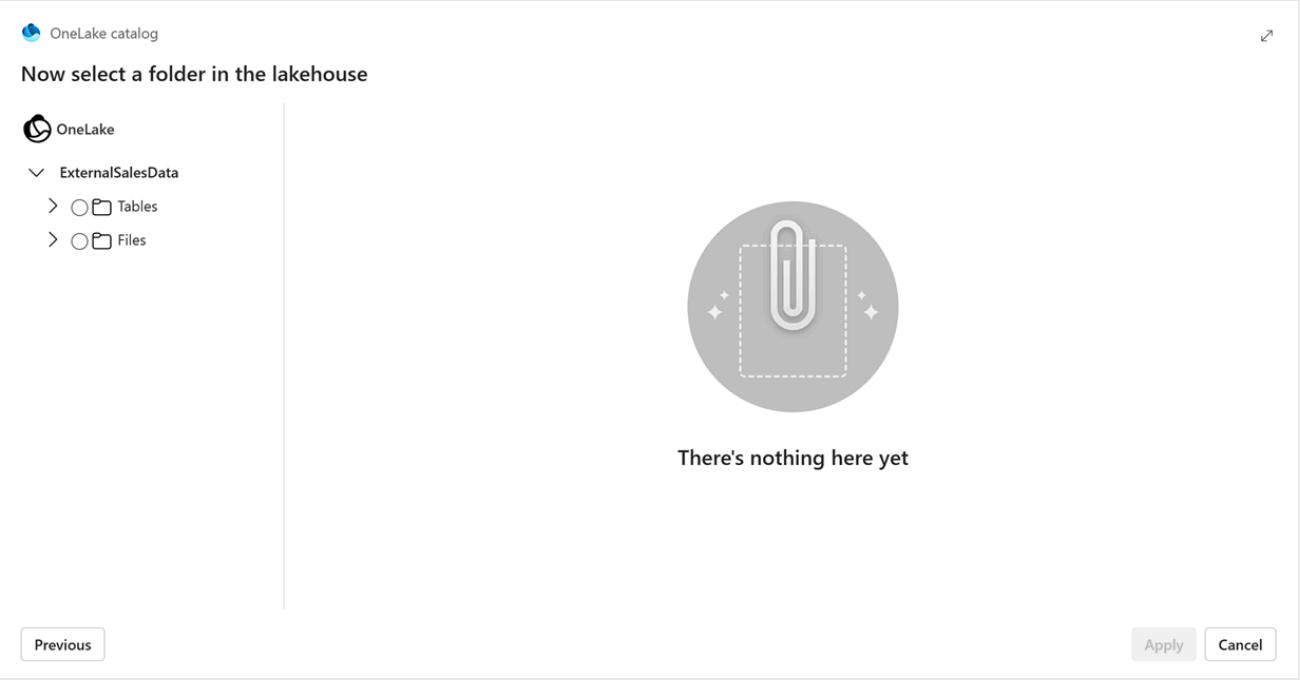

The share has now been created in your OneLake location. The data within this share location can be consumed using any Fabric workload in your tenant.

Security considerations

Using data shared from another tenant has important implications for data security and privacy. See Security considerations for detail.

Related content

- External data sharing overview
- Create an external data share
- Manage external data shares
- Fabric admins: Enable external data sharing

Last updated on 03/16/2025

Manage external data shares

This article explains what you see on the external data share manage permissions tab, and shows how to revoke external data shares.

Prerequisites

- Fabric read and reshare permissions on the item whose external data shares you want to manage.

View an item's external data shares

To view and manage existing external data shares, navigate to the shared item in the workspace or in the OneLake catalog, open the options menu, and select **Manage permissions**.

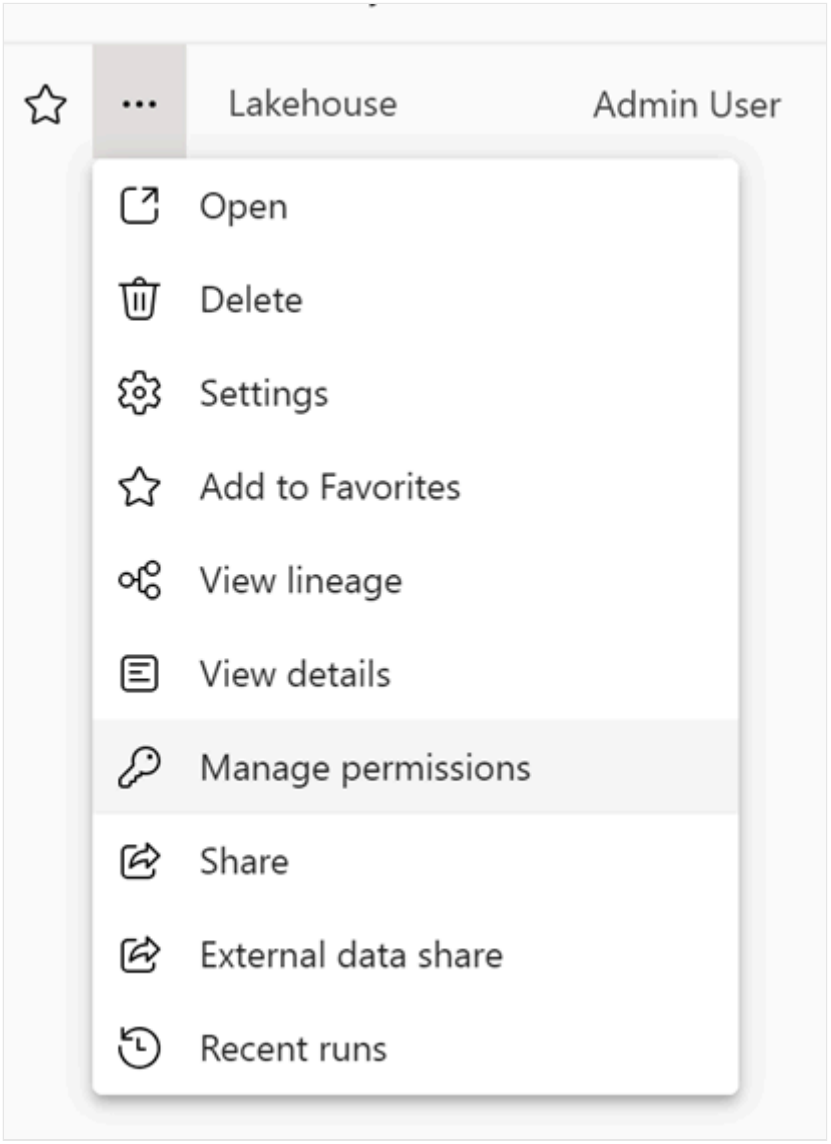

You'll see a list of all the external shares for this item.

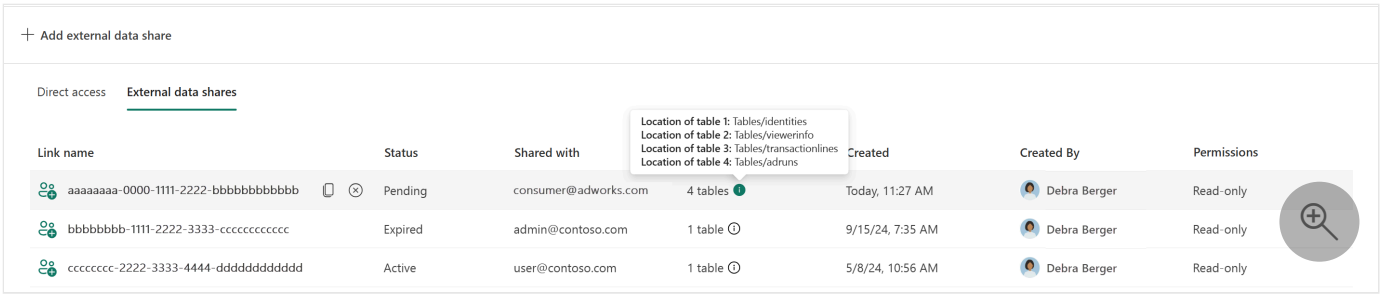

The following table describes the columns.

Expand table

Column	Description
Link name	A guid that is assigned to the external data share when it's created.
Status	Share status. - Pending: Share invitation sent but not yet accepted. - Active: Share accepted. - Expired: Share sent but not accepted within 90 days of it being created. - Revoked: Share revoked. Data no longer accessible via this share.
Shared with	The user to whom the external data invitation was sent.
Contents	A summary of the tables or files that were shared. The locations can be viewed in a tooltip by hovering over the information icon
Created	Creation date of the external data share.
Created by	The user who created the external data share.
Permissions	Permissions granted by the external share.

Revoke external data shares

External data shares can be revoked at any time. You can revoke the external data shares of any item you have read and reshare permissions on.

To revoke a share, hover over its Link name and select the **Revoke share** icon that appears.

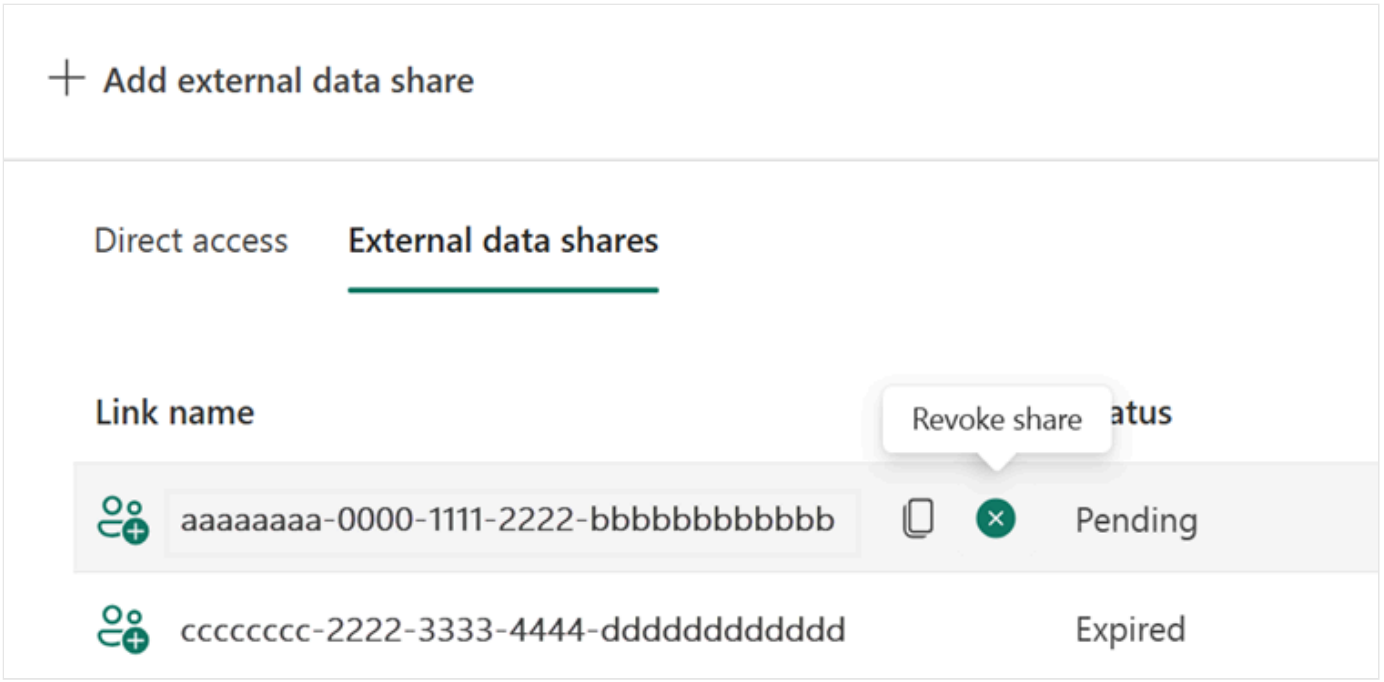

> **⚠ Warning**
>
> Revoking an external data share is a serious matter that should be considered carefully and in consultation with the receiving tenant. It completely and irreversibly severs all access from the receiving tenant to the shared data. This means that any and all data artifacts built in the receiving tenant on the basis of the shared data will cease to function. A revoked external data share can't be restored. A new external data share can be created, but all work done in the receiving tenant based on the revoked share will have to be rebuilt from scratch.

Related content

- External data sharing overview
- Create an external data share
- Accept an external data share
- Fabric admins: Enable external data sharing

Last updated on 03/16/2025

Enable external data sharing in a Fabric tenant

This article is intended for Fabric administrators who want to enable external data sharing in their Fabric tenant.

External data sharing involves two tenants - a providing tenant where the data to be shared is located, and the consuming tenant, where the shared data is consumed. For external data sharing to work, it must be enabled on both tenants as described in this article.

Sharing data via external data sharing has important security considerations. For more information, see Security considerations.

For more information about external data sharing, see the External data sharing overview.

Enable external data sharing in the providing tenant

You must have the Fabric administrator role in the providing tenant to perform these steps.

1. Go to the tenant settings in the providing tenant.
2. Find the **External data sharing** tenant setting (under export and sharing settings), enable the toggle, and specify who in the tenant can create external data shares.

Enable external data sharing in the consuming tenant

You must have the Fabric administrator role in the consuming tenant to perform these steps.

1. Go to the tenant settings in the consuming tenant.
2. Find the **Users can accept external data shares** tenant setting (under export and sharing settings), enable the toggle, and specify who in the tenant can create external data shares.

Related content

- External data sharing overview
- Create an external data share
- Accept an external data share
- Manage external data shares

Last updated on 11/19/2024

Support for private link enabled workspaces

By default, External Data Share doesn't work on workspaces that have inbound network restrictions (workspace-level private links) turned on. A workspace administrator can enable the **Allow external data shares to access this workspace** setting for data consumers to access data protected workspaces.

Prerequisites

- External data sharing is enabled in both the data producer and data consumer tenants.
- You have the **Admin** role on the workspace you want to configure.
- Inbound network restrictions (workspace private links) is enabled on the workspace.

Enable the setting in the workspace

You must have the workspace **Admin** role to perform these steps.

1. Go to the providing workspace and open **Workspace settings**.
2. Select **Network security** > **Inbound networking**.
3. Under the workspace connection settings, turn on **Allow external data shares to access this workspace**.
4. Save your changes.

After you enable the setting, it can take up to 30 minutes for the setting to take effect.

Verify access

1. Ask the consumer to refresh or retry reading the shared data.
2. Confirm the consumer can access the shared data.

What happens when the setting is off

If **Allow external data shares to access this workspace** is off and inbound network restrictions are enabled, consumers see a message that access is blocked by network restrictions. They must

ask the workspace administrator to enable the setting.

Related content

- External data sharing overview
- Create an external data share
- Accept an external data share
- Manage external data shares

Last updated on 07/09/2026

Manage embed codes

As a Fabric administrator, you can view the embed codes that are generated for sharing reports publicly, using the Publish to web from Power BI feature. You can also disable or delete embed codes.

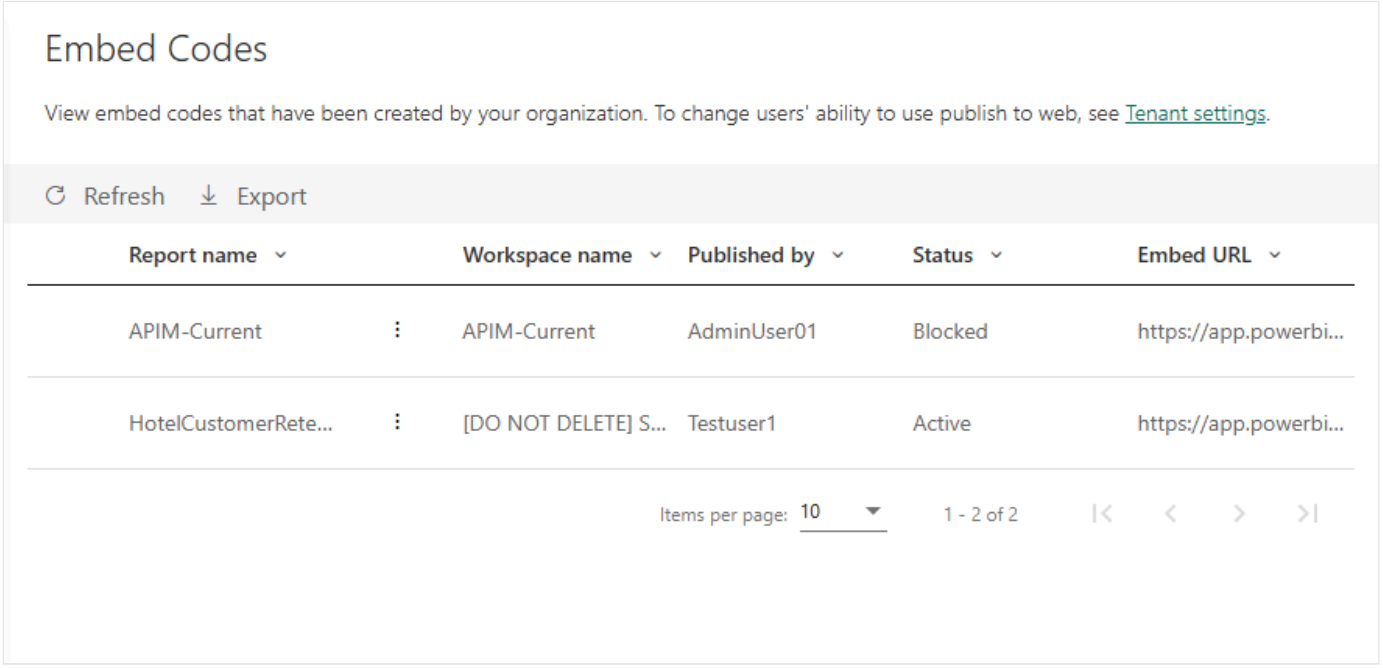

To access the **Embed codes** page in the Admin portal:

1. Sign in to Fabric using your admin account credentials.
2. Open the **OneLake catalog** and go to **Govern** > **Configurations** > **Embed codes**.

Disable embed codes

You can disable the *Publish to web* feature, or allow embed codes to work only in your organization. If you disable *Publish to web*, the existing embed codes aren't deleted. When you re-enable *Publish to web*, the existing embed codes become active again.

Disabling the embed codes is described in Publish to web.

Delete embed codes

To delete embed codes, select the codes you want to delete and then select **Delete**.

Transfer embed code ownership

Embed codes are linked directly to the publisher who creates them. This means that if the publisher loses access to the workspace where a report is published, users can no longer view the embedded report. When a publisher leaves a workspace or an organization, Tenant admins can reassign ownership through the Admin portal, thereby restoring user access.

To change ownership from the Admin portal, follow these steps:

1. Sign in to Fabric using your admin account credentials.
2. Open the **OneLake catalog** and go to **Govern** > **Configurations** > **Embed codes**.
3. On the **Embed codes** page, select **Change ownership**.

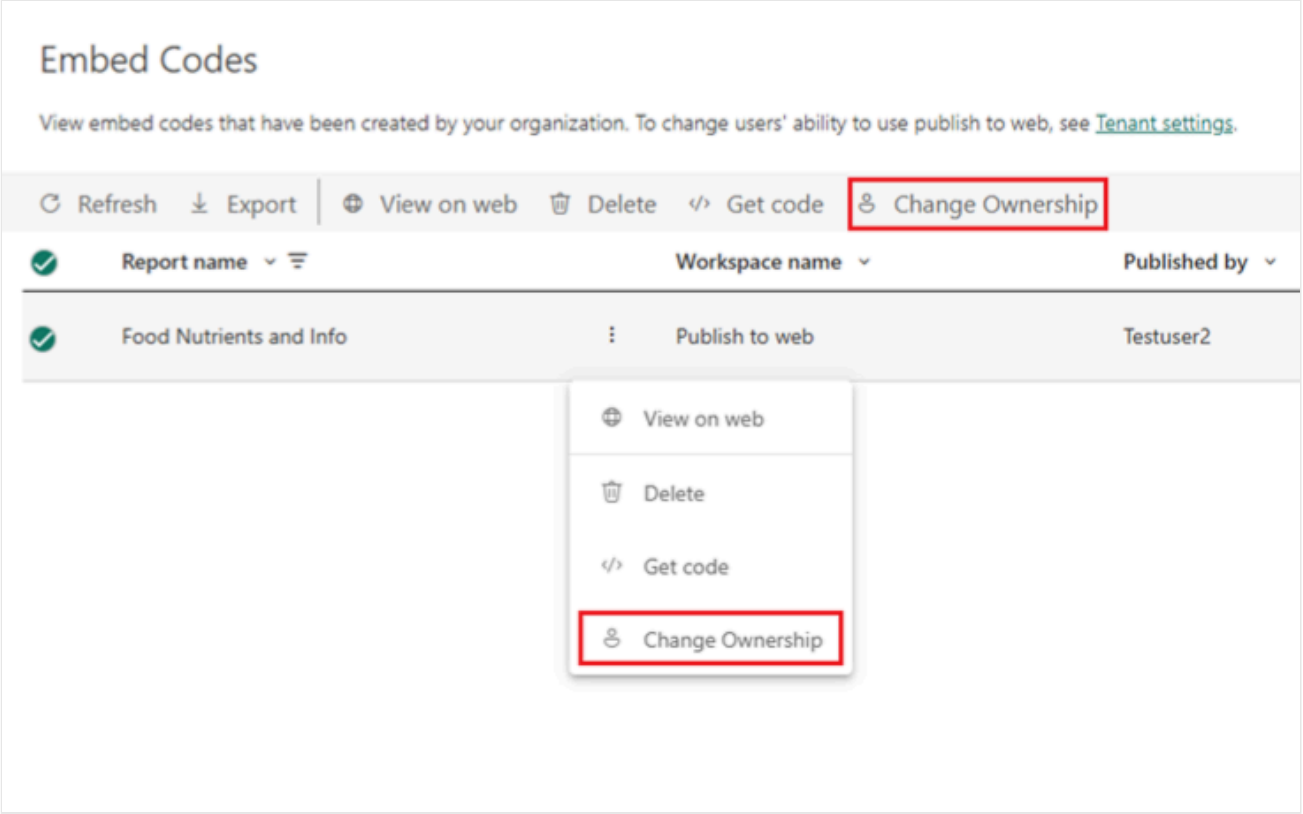

4. Choose the new embed code owner from the dropdown menu, then select **OK**.

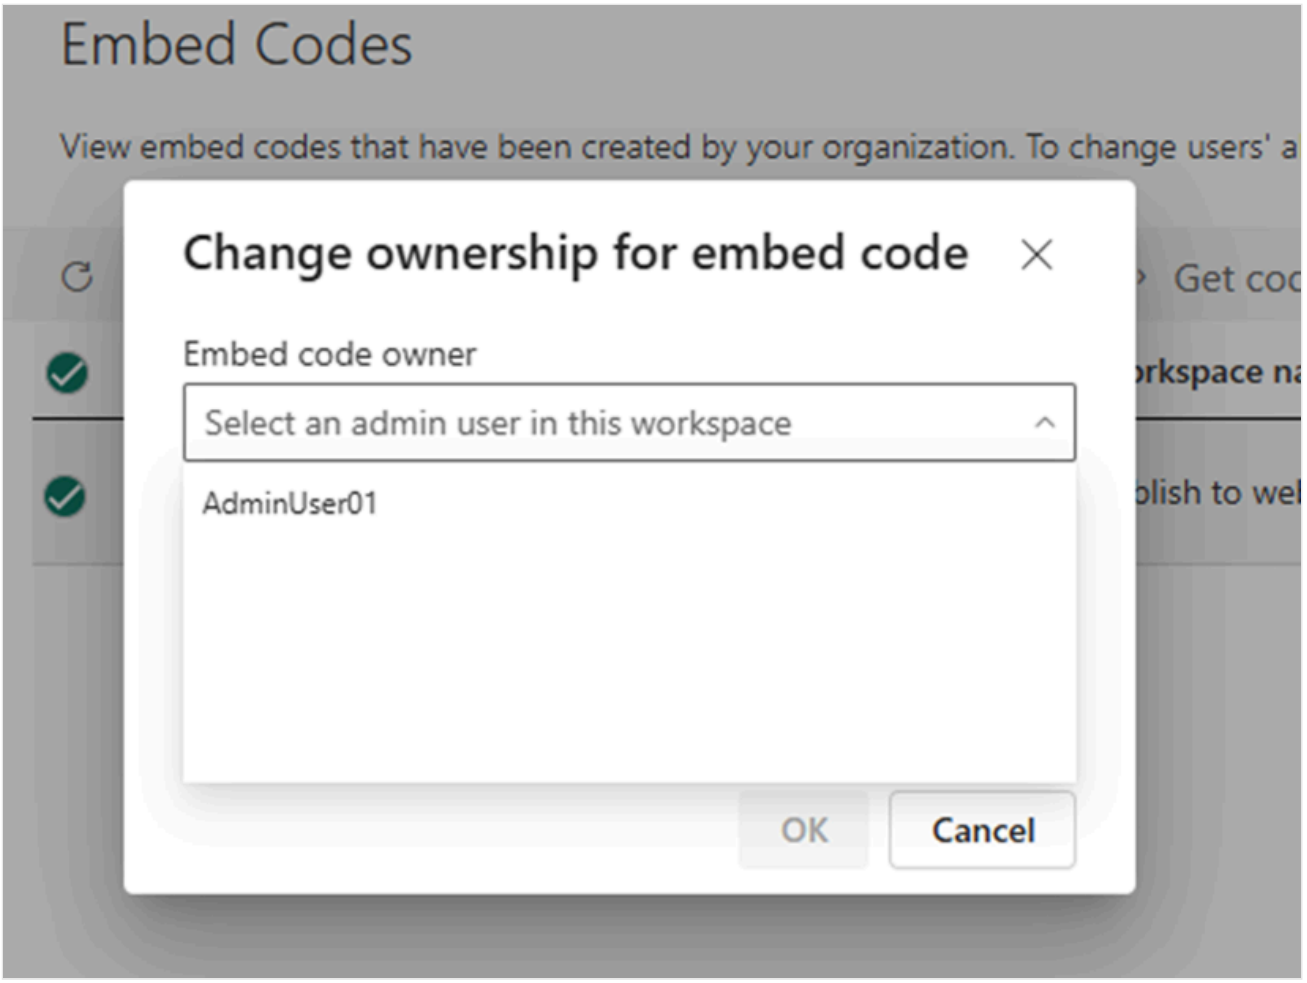

Related content

- Publish to web
- What is the admin portal?

ⓘ**Note:** The author created this article with assistance from AI. Learn more

Last updated on 09/29/2026

Endorsement overview

Endorsement is a Fabric and Power BI feature that makes it easier for users in an organization to find high-quality, trust-worthy content and data. Endorsed items are clearly labeled in the UI with a badge, and in some lists they are given precedence and are listed first.

Items can receive one of three endorsement badges - **Promoted**, **Certified**, or **Master data**:

- **Promoted**: When an item has the **Promoted** badge, it means that the item creators think the item is ready for sharing and reuse. Any Fabric or Power BI item except Power BI dashboards can be promoted. Any user with write permissions on an item can promote it.

  Learn how to promote items.

- **Certified**: When an item has the **Certified** badge, it means that an organization-authorized reviewer has certified that the item meets the organization's quality standards, can be regarded as reliable and authoritative, and is ready for use across the organization. Any Fabric or Power BI item except Power BI dashboards can be certified. Any user can request certification for an item, but only users specified by a Fabric administrator can actually certify items.

  Learn how to certify items or request certification.

- **Master data**: When an item has the **Master data** badge, it means the data in the item is a core source of organizational data. The master data designation is often used to indicate that a data item is to be regarded as the authoritative, single source of truth for certain kinds of organizational or business data, such as product codes or customer lists. The master data label can only be applied to items that contain data, such as lakehouses, and semantic models. Only users specified by the Fabric administrator can label data items as master data.

  Learn how to label data items as master data or request master data designation.

> **Note**
>
> Certification and master data endorsement are available only if a Fabric administrator has enabled them for your organization. Fabric admins should see Enable item certification and Enable master data endorsement.

Certification enablement can be delegated to domain administrators, making it possible to specify a different set of reviewers for each domain. For more information, see Enable item certification.

How endorsed items appear in Fabric and Power BI

The following image illustrates how promoted, certified, and master data items are clearly identified when you're searching for an item.

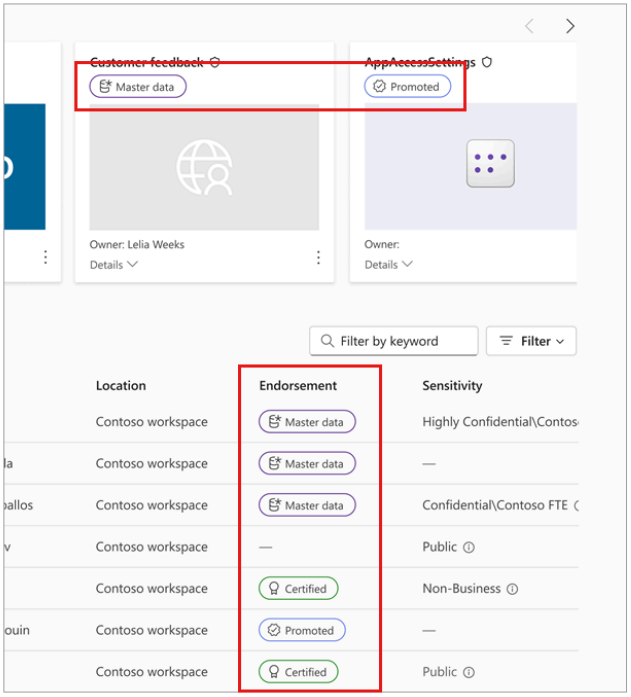

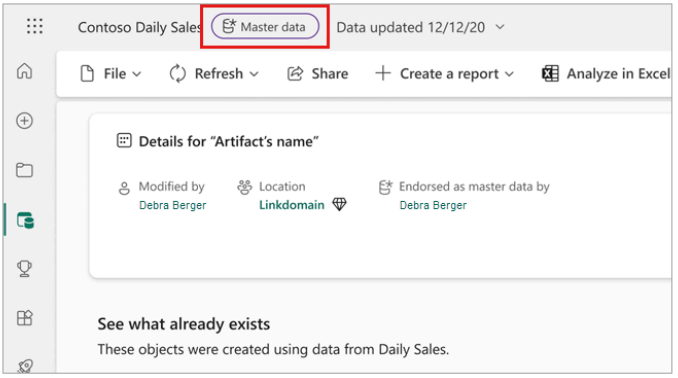

Types of items that can be endorsed

- All Fabric items and Power BI items except Power BI dashboards can be promoted or certified.
- All Fabric and Power BI items that contain data can be labeled as master data.

Related content

- Promote or certify Fabric content
- Enable item certification (Fabric admins)
- Enable master data endorsement (Fabric admins)

Last updated on 07/11/2024

Enable item certification

Your organization can certify selected items to identify them an as authoritative sources for critical information. Currently, all Fabric items except Power BI dashboards can be certified.

As a Fabric admin, you're responsible for enabling and setting up the certification process for your organization. This means:

- Enabling certification on your tenant.
- Defining a list of security groups whose members will be authorized to certify items.
- Providing a URL that points to the documentation for the organization's item certification process, if such documentation exists.
- Deciding whether to delegate certification setup to domain administrators, so that they can set up certification specifically for their domain. When you delegate certification setup to domain administrators, the administrators of each domain can override any or all tenant-level certification settings, including enable/disable, for their domain.

Certification is part of Power BI's *endorsement* feature. For more information, see the endorsement overview.

Enable item certification

1. Go to **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**.

2. Under the Export and sharing settings section, expand the Certification section.

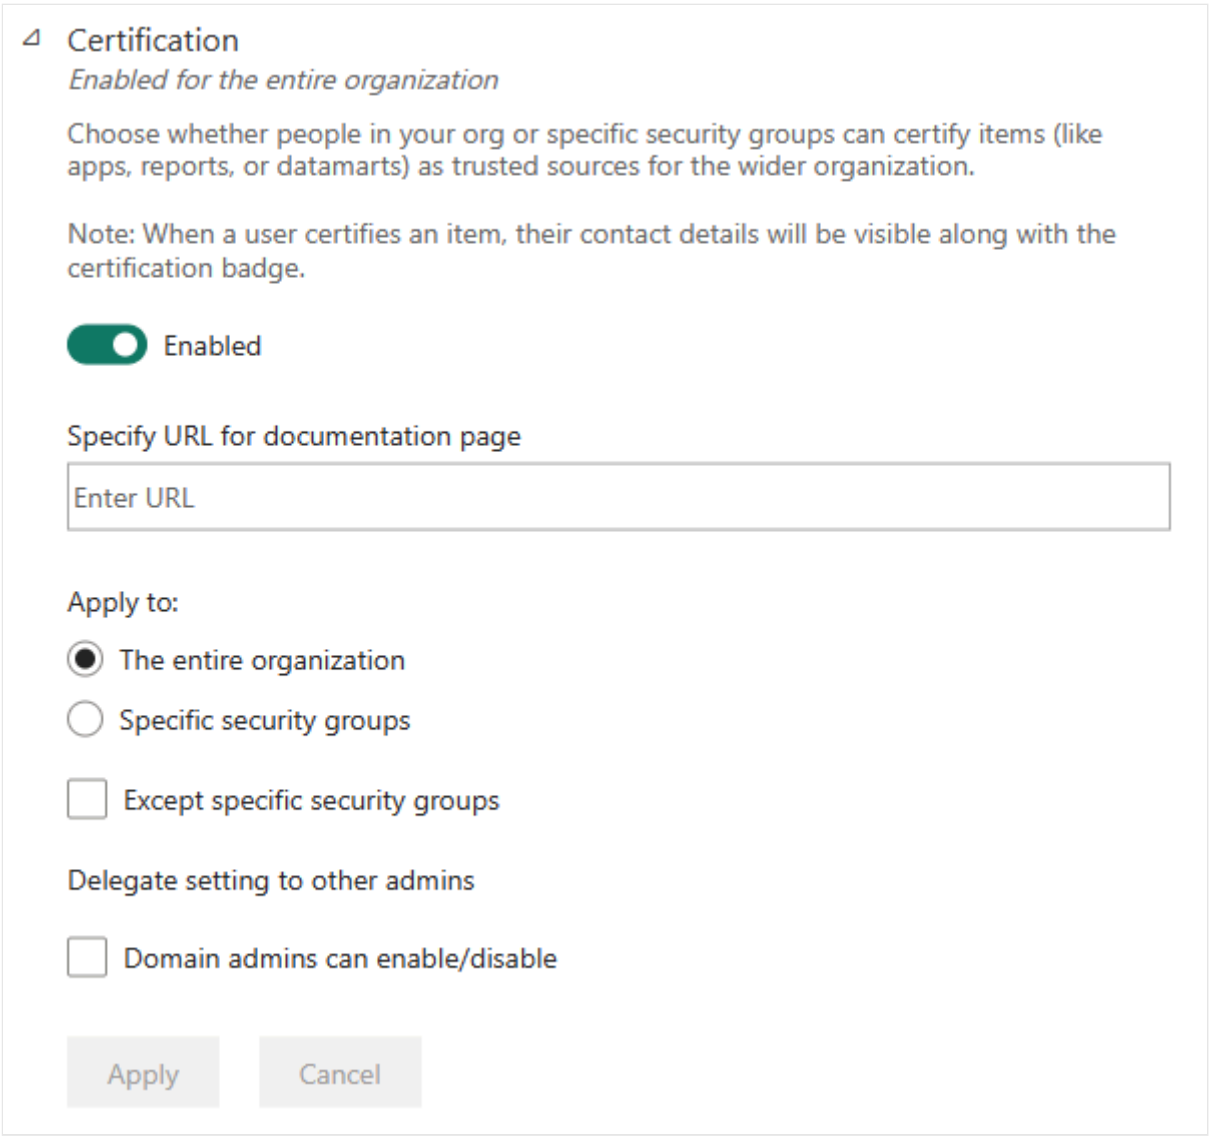

3. Set the toggle to **Enabled**.

4. If your organization has a published certification policy, provide its URL here. This becomes the **Learn more** link in the certification section of the endorsement settings dialog. If you don't supply a link, users who want to request certification of their item will be advised to contact their Fabric administrator.

5. Specify one or more security groups whose members will be authorized to certify items. These authorized certifiers will able to use the Certification button in the certification section of the endorsement settings dialog. This field accepts security groups only. You can't enter named users.

   If a security group contains subsecurity groups that you don't want to give certification rights to, you can check the **Except specific security groups** box and enter the names of those groups in a text box that will appear.

6. Check the **Domain admins can enable/disable** checkbox if you want domain administrators to be able to override any or all tenant-level certification settings.

> **Note**
>
> Selecting the checkbox enables domain admins to override any or all tenant-level certification settings, including enable/disable, even though the checkbox description only mentions enable/disable.

7. Select **Apply**.

Related content

- Read about endorsement in Fabric
- Enable master data endorsement
- Promote Fabric items
- Certify Fabric items

Last updated on 09/29/2026

Enable master data endorsement

This article describes how to enable master data endorsement. Its target audience is Fabric administrators.

Master data endorsement is a way organizations can help its Fabric users find and use data that the organization regards as core, single-source-of-truth data, such as customer lists, product codes, etc. This is accomplished by qualified users applying a **Master data** badge to data items that they judge to contain that core, single-source-of-truth data.

As a Fabric administrator, you're responsible for enabling master data endorsement in your organization and specifying who in the organization is qualified to designate data items as master data.

For more information about endorsement and master data, see Endorsement overview.

Prerequisites

You must be a Fabric administrator to enable master data endorsement.

Enable master data endorsement

To enable master data endorsement, turn on the **Endorse master data (preview)** tenant setting and specify who is authorized to apply the **Master data** badge to data items.

1. Sign in to Fabric using your admin account credentials. Open the **OneLake catalog**, and then select the **Govern** tab > **Configurations** > **Tenant settings**.

2. Find and expand the **Endorse master data (preview)** tenant setting.

3. Switch the toggle to **Enabled**.

4. Specify who can apply the **Master data** badge to data items by choosing the appropriate options:

    - **The entire organization**: Everyone in the organization is authorized to apply the master data badge to data items.
    - **Specific security groups**: Only members of the specified security groups are authorized to apply the **Master data** badge to data items.

- **Exclude specific security groups**: If you want to exclude particular users who are included in the option you chose, select the **Exclude specific security groups** checkbox and provide the names of a security group or groups that include the users you want to exclude.

> **Note**
>
> Users who are authorized to apply the **Master data** badge must also have write permission on data items they wish to apply the badge to.

Related content

- Endorsement overview
- Enable item certification
- Endorse items
- Governance documentation

Last updated on 09/29/2026

Fabric domains

This article introduces key concepts of domains in Fabric, and shows how to set up and manage them. To get started planning domains for your organization, see Best practices for planning and creating domains in Microsoft Fabric.

Introduction

Today, organizations are facing massive growth in data, and there's an increasing need to be able to organize and manage that data in a logical way that facilitates more targeted and efficient use and governance.

To meet this challenge, organizations are shifting from traditional IT centric data architectures, where the data is governed and managed centrally, to more federated models organized according to business needs. This federated data architecture is called *data mesh*. A data mesh is a decentralized data architecture that organizes data by specific business domains, such as marketing, sales, human resources, etc.

Currently, Microsoft Fabric's data mesh architecture primarily supports organizing data into domains and enabling data consumers to be able to filter and find content by domain. It also enables federated governance, which means that some governance currently controlled at the tenant level can be delegated to domain-level control, enabling each business unit/department to define its own rules and restrictions according to its specific business needs.

Key concepts

Domains

In Fabric, a domain is a way of logically grouping together all the data in an organization that is relevant to a particular area or field. One of the most common uses for domains is to group data by business department, making it possible for departments to manage their data according to their specific regulations, restrictions, and needs.

To group data into domains, workspaces are associated with domains. When a workspace is associated with a domain, all the items in the workspace are also associated with the domain, and they receive a domain attribute as part of their metadata. Currently, the association of workspaces and the items in them with domains primarily enables a better consumption experience. For instance, in the OneLake catalog, users can filter content by domain in order find

content that is relevant to them. In addition, some tenant-level settings for managing and governing data can be delegated to the domain level, thus allowing domain-specific configuration of those settings.

> **Note**
>
> Domain assignment doesn't affect item visibility or accessibility for tenant users. Item discovery, visibility, and access depend on such things as workspace role and item permissions, but not domain assignment.
>
> Likewise, all users within a tenant can see all the domains defined in the tenant, regardless of their specific domain roles. For example, users who are neither contributors nor admins of a domain called "Finance" in their tenant can still see this domain in the **domains filter** of the Onelake Catalog.

Subdomains

A subdomain is a way for fine tuning the logical grouping of your data. You can create subdomains under domains. For information about how to create subdomains, see Create subdomains.

Domain roles

There are three roles involved in the creation and management of domains:

- **Fabric admin** (or higher): Fabric admins can create and edit domains, specify domain admins and domain contributors, and associate workspaces with domains. Fabric admins see all the defined domains on the Domains tab in the admin portal, and they can edit and delete domains.

- **Domain admin**: Ideally, the domain admins of a domain are the business owners or designated experts. They should be familiar with the data in their area and the regulations and restrictions that are relevant to it.

  Domain admins can access to the **Domains** tab in the admin portal, but they can only see and edit the domains they're admins of. Domain admins can update the domain description, define/update domain contributors, and associate workspaces with the domain. They also can define and update the domain image and override tenant settings for any

specific settings the tenant admin has delegated to the domain level. They can't delete the domain, change the domain name, or add/delete other domain admins.

- **Domain contributor**: Domain contributors are workspace admins whom a domain or Fabric admin has authorized to assign the workspaces they're the admins of to a domain, or to change the current domain assignment.

  Domain contributors assign the workspaces they're an admin of in the settings of the workspace itself. They don't have access to the **Domains** tab in the admin portal.

  > **Note**
  >
  > Remember, to be able to assign a workspace to a domain, a domain contributor must be a workspace admin (that is, have the Admin role in the workspace).

Domain settings delegation

To allow domain-specific configuration, some tenant-level settings for managing and governing data can be delegated to the domain level. Domain settings delegation enables each business unit/department to define its own rules and restrictions according to its specific business needs.

Domain image

When users look for data items in the OneLake catalog, they might want to see only the data items that belong to a particular domain. To do this, they can select the domain in the domain selector in the OneLake catalog to display only items belonging to that domain. To remind them which domain's data items they're seeing, you can choose an image to represent your domain. Then, when your domain is selected in the domain selector, the image becomes part of OneLake catalog's theme, as illustrated in the following image.

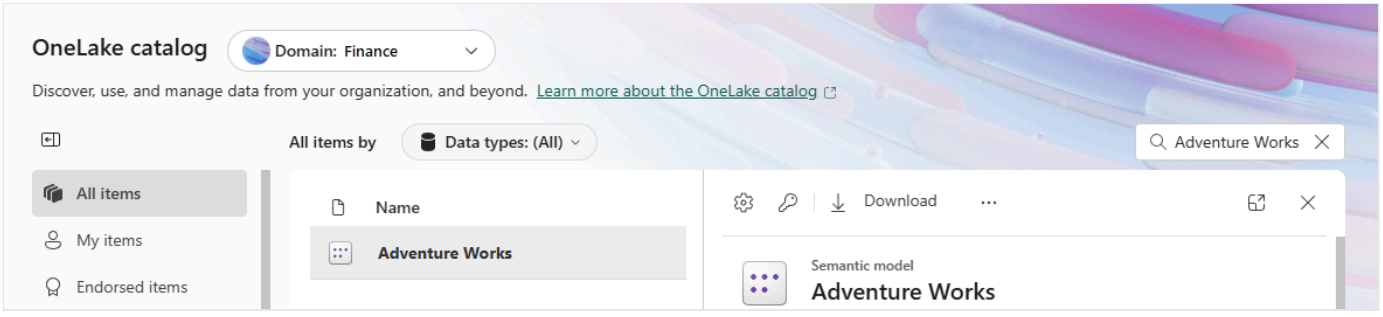

For information about how to specify an image for a domain, see Specify a domain image.

Default domain

A default domain is a domain that has been defined as the default domain for specified users and/or security groups. When you define a domain as the default domain for specified users and/or security groups, the following happens:

1. The system scans the organization's workspaces. When it finds a workspace whose admin is a specified user or member of a specified security group:

    - If the workspace already has a domain assignment, it is preserved. The default domain doesn't override the current assignment.
    - If the workspace is unassigned, it is assigned to the default domain.

2. After this, whenever a specified user or member of a specified security group creates a new workspace, it is assigned to the default domain.

The specified users and/or members of the specified security groups generally automatically become domain contributors of workspaces that are assigned in this manner.

For information about defining a domain as a default domain, see Define the domain as a default domain.

Create a domain

Before you start creating domains for your organization, it is recommended to review Best practices for planning and creating domains in Microsoft Fabric.

To create domain, you must be a Fabric admin.

1. Sign in to Fabric by using your admin account credentials. Open the **OneLake catalog**, and then select **Govern** > **Domains**.

2. On the **Domains** tab, select **Create new domain**.

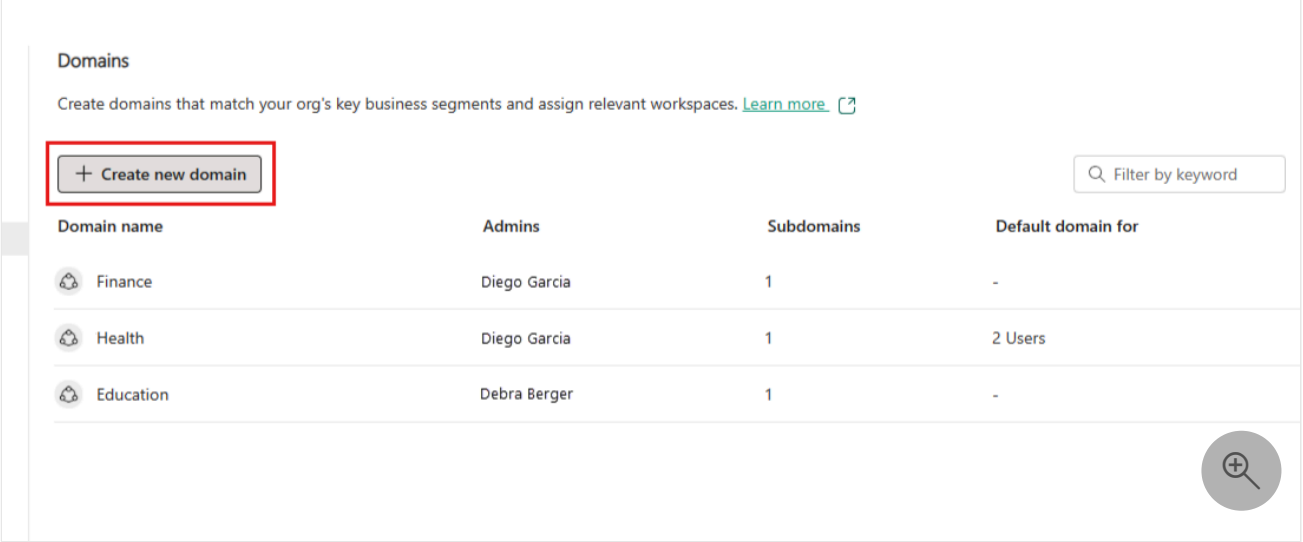

3. In the **New domain** dialog that appears, provide a name (mandatory) and specify domain admins (optional). If you don't specify domain admins, you can do this later in the domain settings.

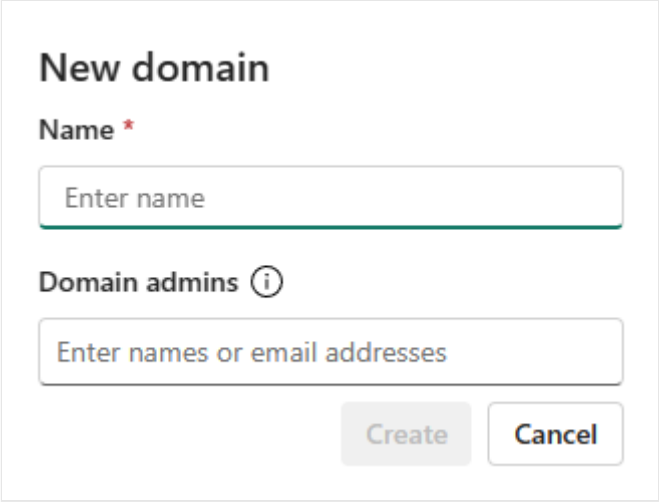

4. Select **Create**. The domain is created, and you can continue configuring the domain as described in the following sections.

Structure your data in the domain

Once you create some domains, you can refine the logic of the way you're structuring your data by creating subdomains for the domains.

You organize your data into the appropriate domains and subdomains by assigning the workspaces the data is located in to the relevant domain or subdomain. When a workspace is assigned to a domain, all the items in the workspace are associated with the domain.

Create subdomains

To create subdomains for a domain, you must be Fabric admin or domain admin.

1. Open the domain you want to create a subdomain for and select **New subdomain**.

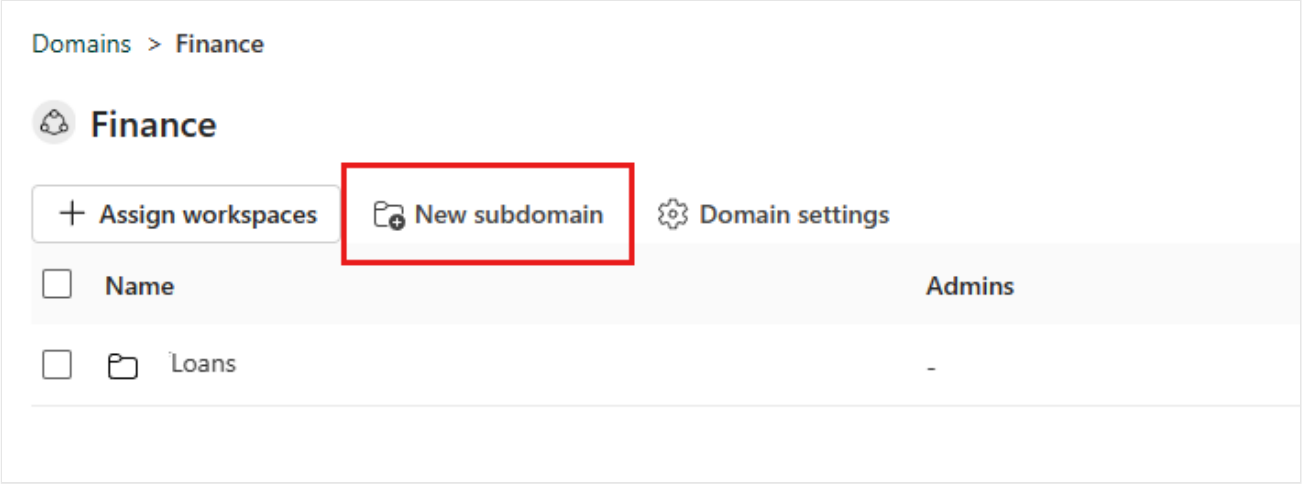

2. Provide a name for the subdomain in the **New subdomain** dialog that appears. When done, select **Create**.

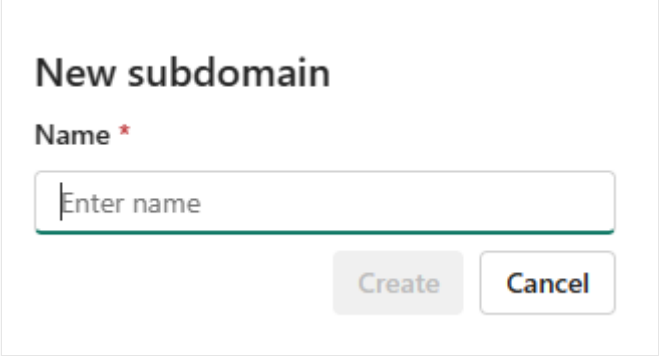

> ⓘ **Note**
>
> Subdomains don't have their own domain admins. A subdomain's domain admins are the domain admins of its parent domain.

Assign workspaces to domains and subdomains

To assign workspaces to a domain or subdomain in the admin portal, you must be a Fabric admin or a domain admin.

1. Go to the domain or subdomain's page and select **Assign workspaces**.

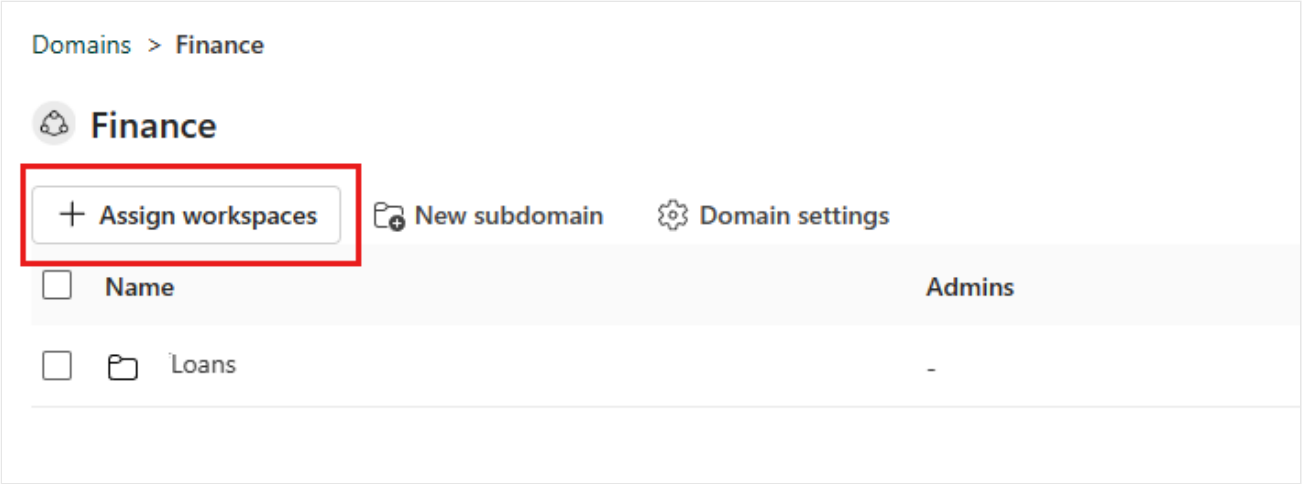

2. In the **Assign workspaces to this domain** side pane, select how to assign the workspaces.

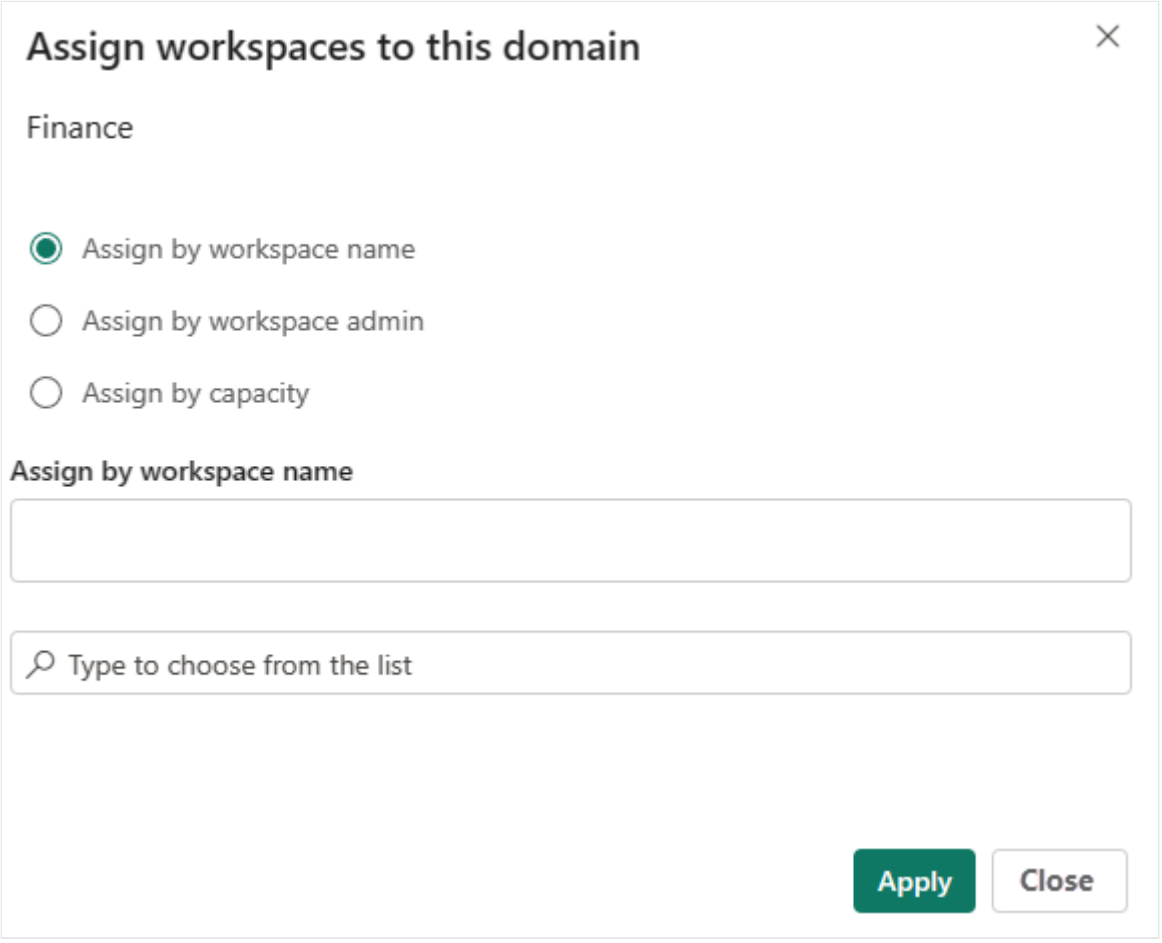

- **Assign by workspace name**
  - Some organizations have naming conventions for workspaces that make it easy to identify the data's business context.
  - You can search for and select multiple workspaces at once
  - If a workspace is already associated with another domain, you'll see an icon next to the specific name. If you chose to continue the action, a warning message pops up, but you'll be able to continue and override the previous association.
- **Assign by workspace admin**

- You can select specific users or security groups as per your business structure. When you confirm the selection, all the workspaces the users and security groups are admins of will be associated to the domain.
- This action excludes "My workspaces".
- If some of the workspaces are already associated with another domain, a warning message will pop up, but you'll be able to continue and override the previous association.
- This action affects existing workspaces only. It won't affect workspaces the selected users create after the action has been performed.

- **Assign by capacity**
  - Some organizations have dedicated capacities per department/business unit.
  - You can search for and select multiple capacities at once. When you confirm your selection, all the workspaces associated to the selected capacities will be assigned to the domain.
  - If some of the workspaces are already associated with another domain, a warning message will pop up, but you'll be able to continue and override the previous association.
  - This action excludes "My workspaces".
  - This action affects existing workspaces only. It won't affect workspaces that are assigned to the specified capacities after the action has been performed.

> ⓘ **Note**
>
> Workspace domain assignments by Fabric and domain admins will override existing assignments only if the **Allow tenant and domain admins to override workspace assignments (preview)** tenant setting is enabled. For more information, see Allow tenant and domain admins to override workspace assignments (preview).

To unassign a workspace from a domain or subdomain, select the checkbox next to the workspace name and then select the **Unassign** button above the list. You can select several checkboxes to unassign more than one workspace at a time.

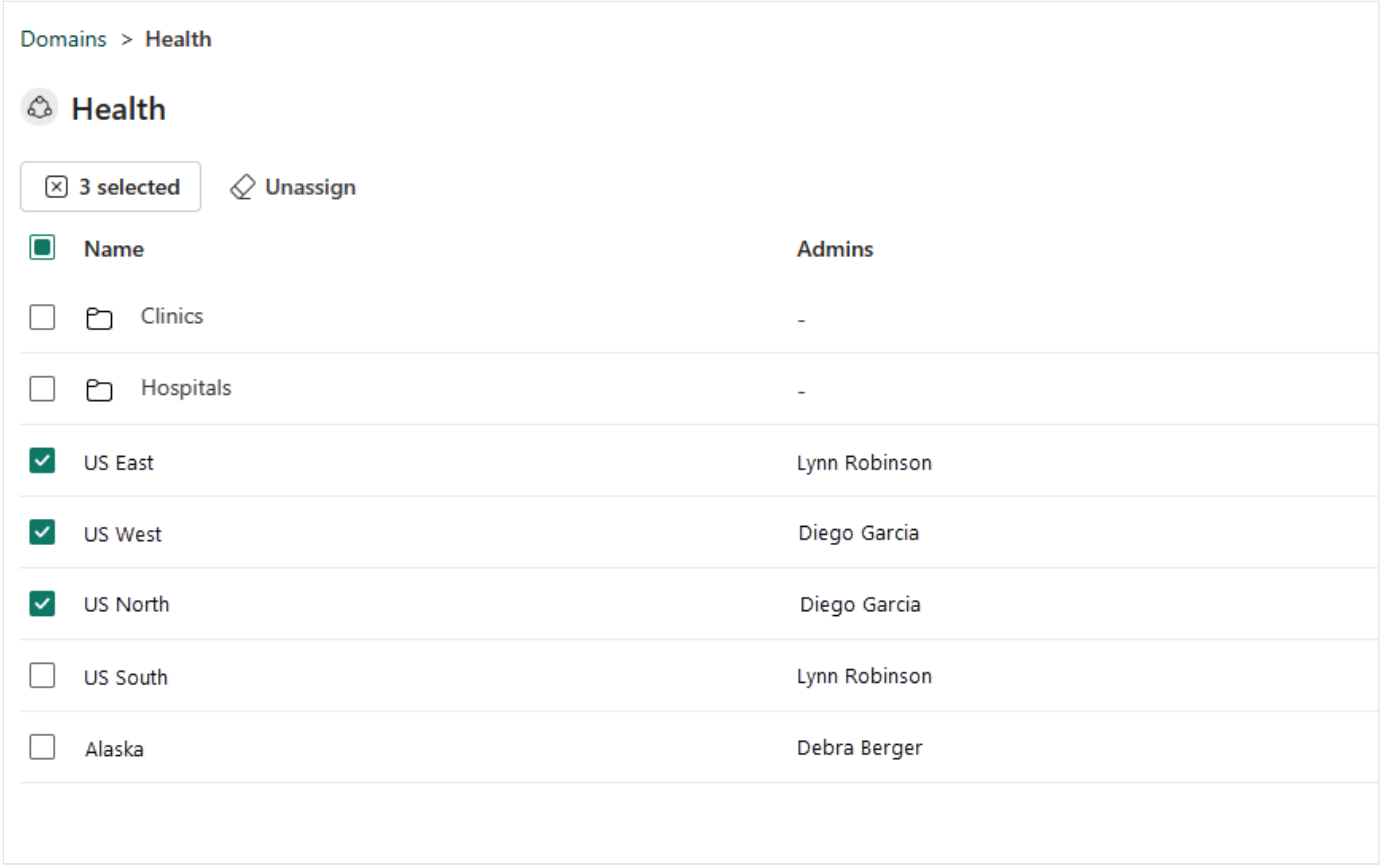

Configure domain settings

You configure domain and subdomain settings on the domain or subdomain's **Domain settings** side pane.

The domain settings side pane has the following tabs:

- General settings: Edit domain name and description
- Image: Specify domain image
- Admins: Specify domain admins
- Contributors: Specify domain contributors
- Default domain: Set up domain as a default domain
- Delegated settings: Override tenant-level settings

> ⓘ **Note**
>
> Subdomains currently have general settings only.

To open the **Domain settings** side pane, open the domain or subdomain and select **Domain settings** (for subdomains, **Subdomain settings**).

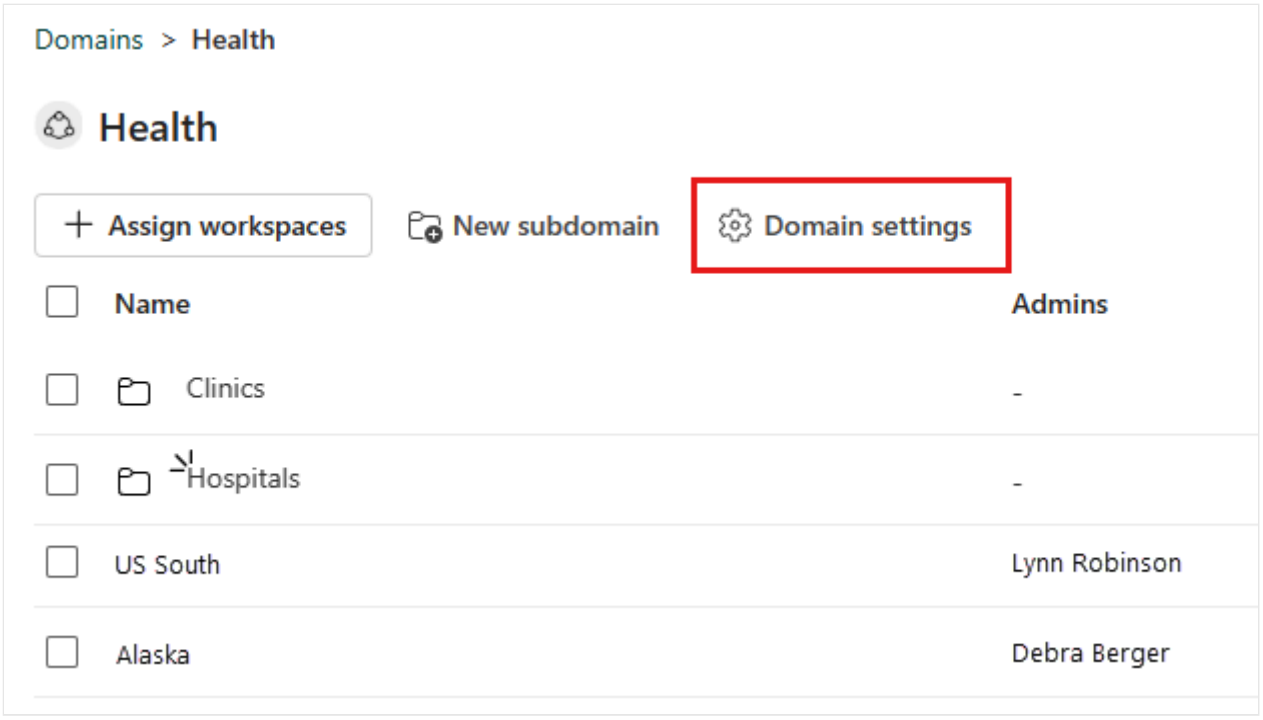

Alternatively, for domains, you can hover over the domain on the Domain tab, select **More options (...)**, and choose **Settings**.

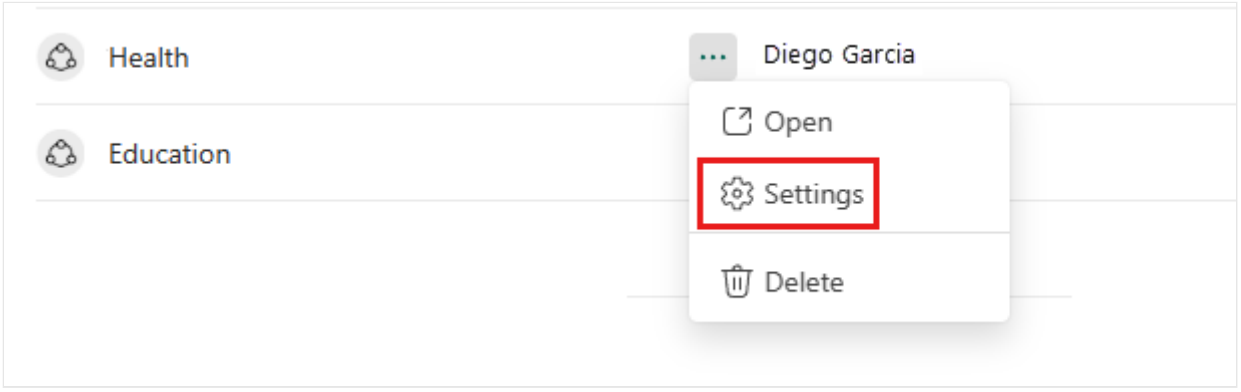

Edit name and description

1. Select **General settings** and then edit the name and description fields as desired.

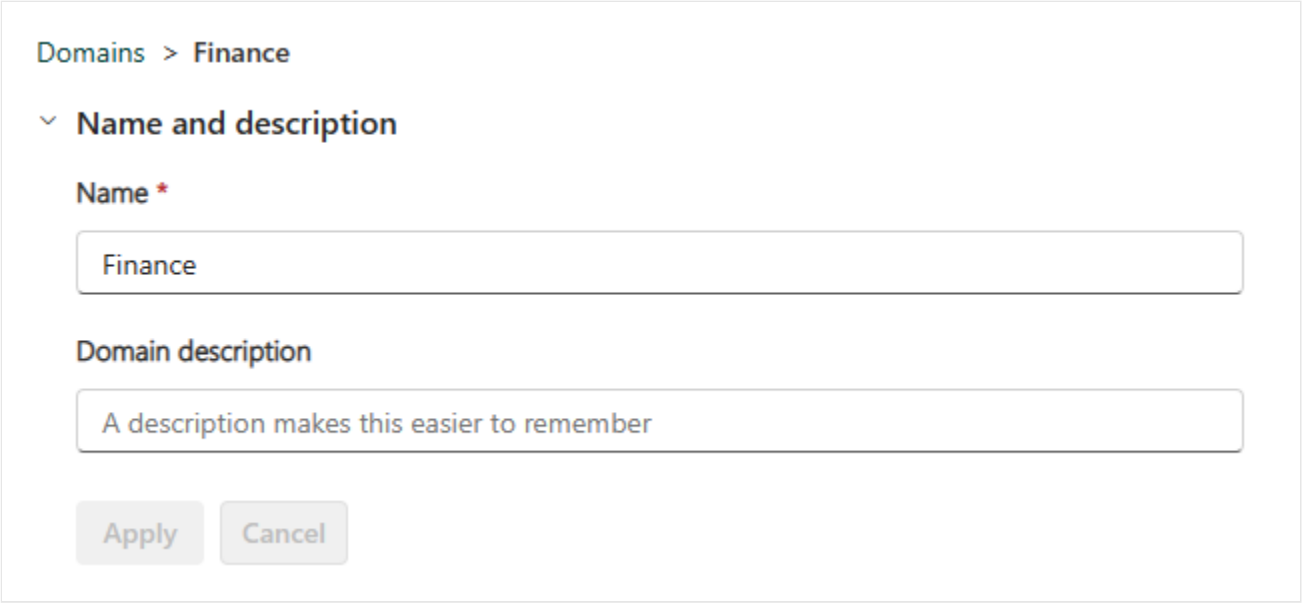

> **Note**
>
> Domain admins can only edit the description field.

2. When done, select **Apply**.

Specify a domain image

Select **Image** and then select **Select an image**.

In the photo gallery that pops up you can choose an image or color to represent your domain in the OneLake catalog when your domain is selected.

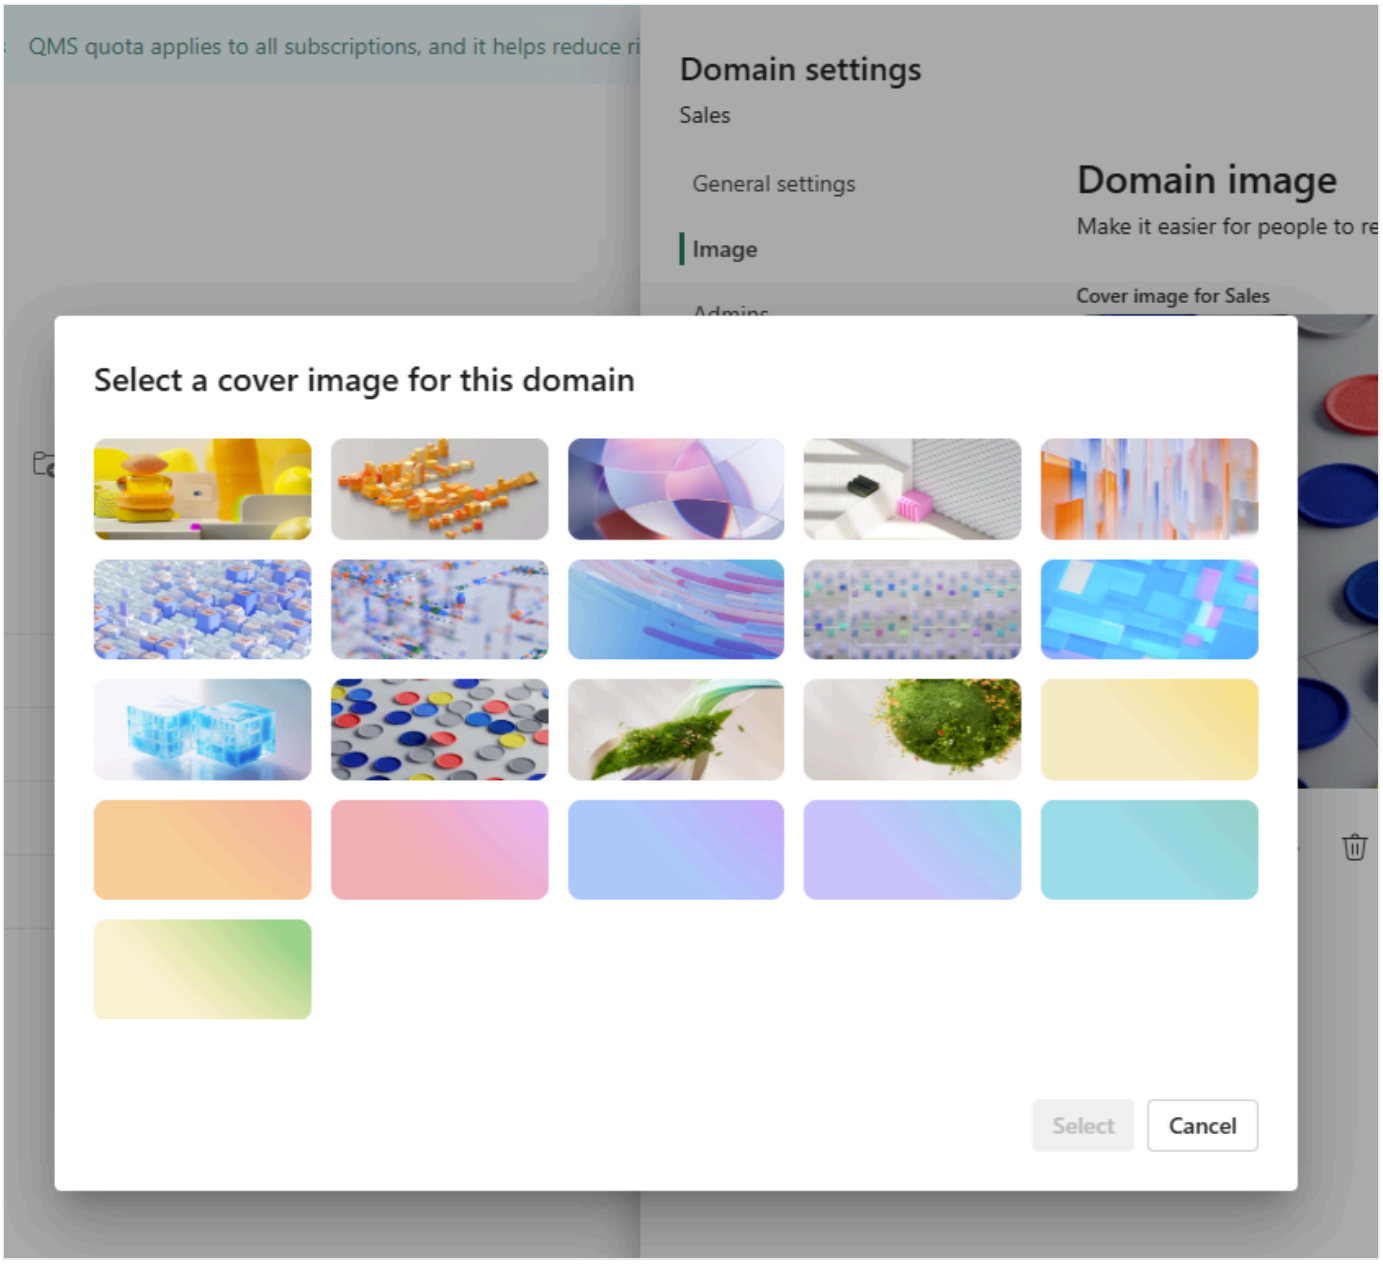

Specify domain admins

You must be a Fabric admin to specify domain admins.

Select **Admins** and then specify who can change domain settings and add or remove workspaces. When done, select **Apply**.

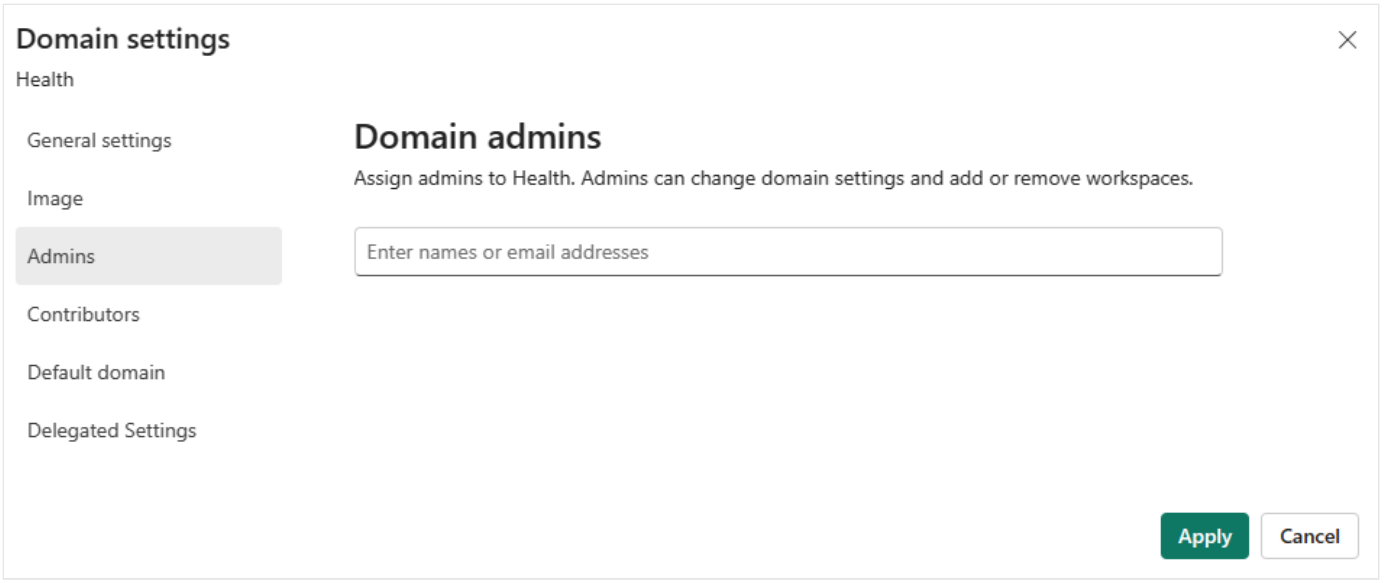

Specify domain contributors

You must be a domain admin of the domain or a Fabric admin to specify domain contributors.

Select **Contributors** and then specify who can assign workspaces to the domain. You can specify everyone in the organization (default), specific users/groups only, or you can allow only tenant admins and the specific domain admins to assign workspaces to the domain. When done, select **Apply**.

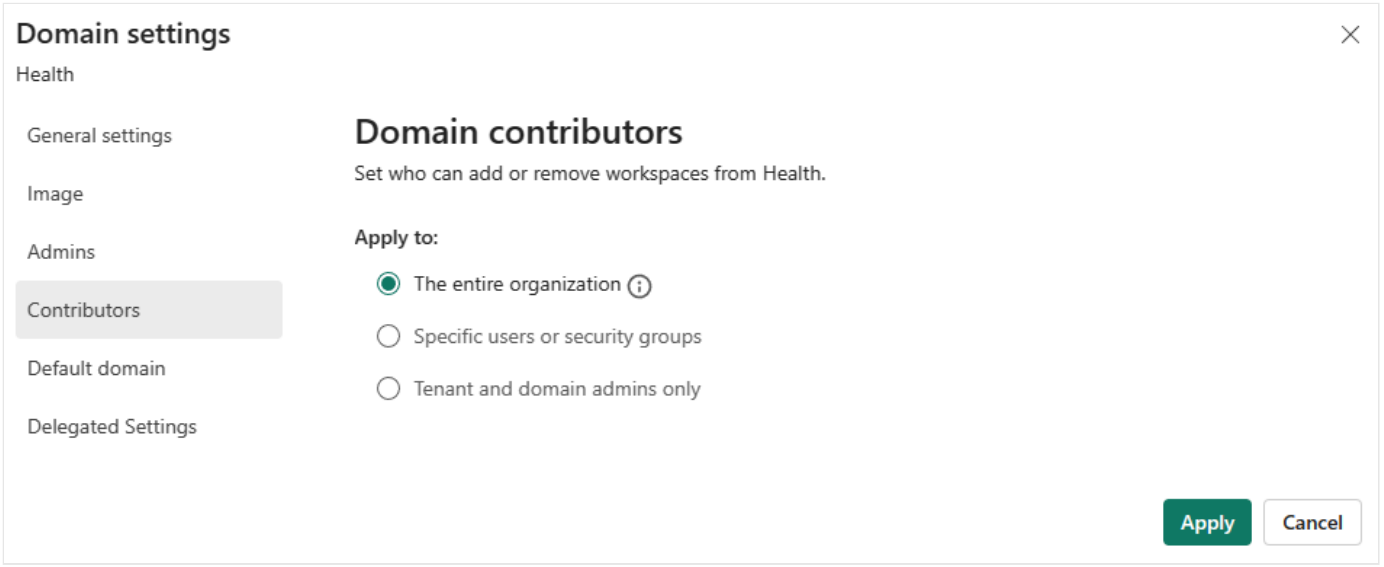

> **Note**
>
> For domain contributors to be able to associate their workspaces with their domains, they must have an admin role in the workspaces they are trying to associate with the domain.

Define the domain as a default domain

To define a domain as a default domain, you must be a Fabric admin or a domain admin of the domain.

Select **Default domain** and specify users and/or security groups. When you add people to the default domain list, unassigned workspaces they're admins of, and new workspaces they create, will automatically be assigned to the domain. For a detailed description of the process, see Default domain.

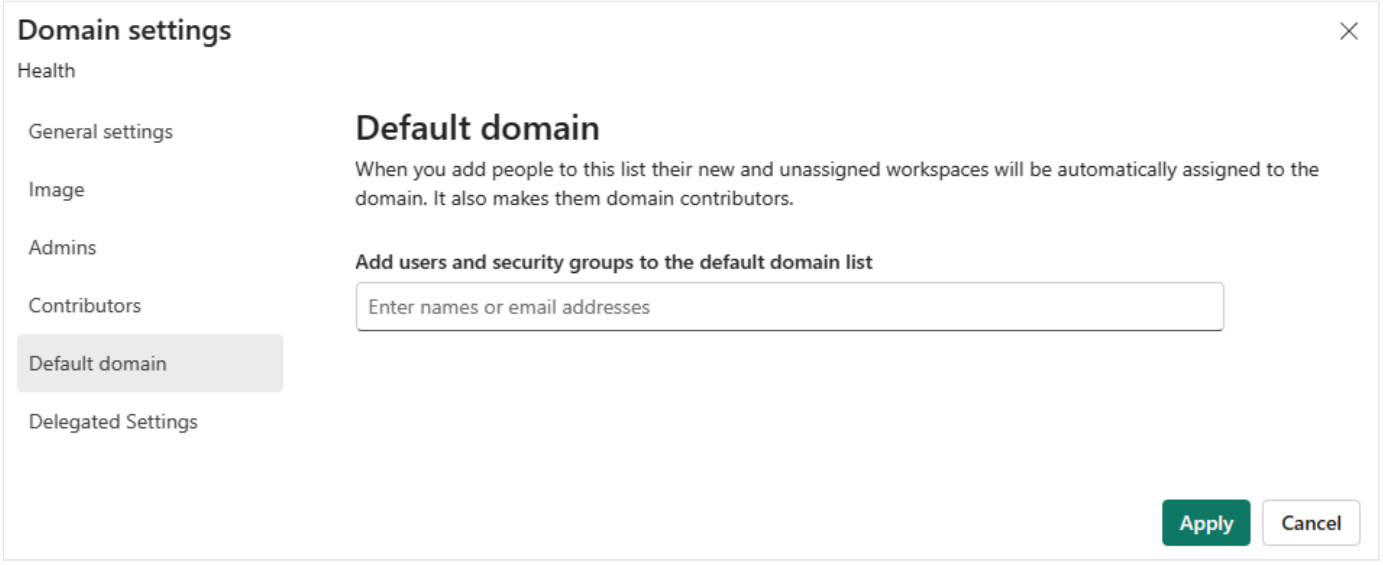

> **Note**
>
> The users and/or members of the security groups specified in the default domain definition generally automatically become domain contributors of the workspaces that get assigned to the domain via the default domain mechanism.

Delegate settings to the domain level

Some tenant-level settings can potentially be overridden at the domain level. To see these settings, select **Delegated Settings**. The following admin settings can potentially be overridden.

Domain-level default sensitivity label

If the domain-level default sensitivity label feature is enabled in your organization, you can specify a sensitivity label that will be applied by default to items in workspaces that are assigned to the domain.

To specify a default sensitivity label for your domain, you must be a Fabric admin or a domain admin of the domain.

Expand **Delegated Settings** and choose **Information protection**. You'll see the option **Set a default label for this domain**. Select the drop down menu and select the desired sensitivity label. The label will be applied to items in workspaces associated with the domain according to the logic described in Domain-level default sensitivity labels in Microsoft Fabric.

Certification settings

Certification is a way for organizations to label items that it considers to be quality items. For more information about certification, see Endorsement.

Certification settings at the domain level mean you can:

- Enable or disable certification of items that belong to the domain.
- Specify certifiers who are experts in the domain.
- Provide a URL to documentation that is relevant to certification in the domain.

To override the tenant-level certification settings, expand the certification section, select the **Override tenant admin selection** checkbox, and configure the settings as desired.

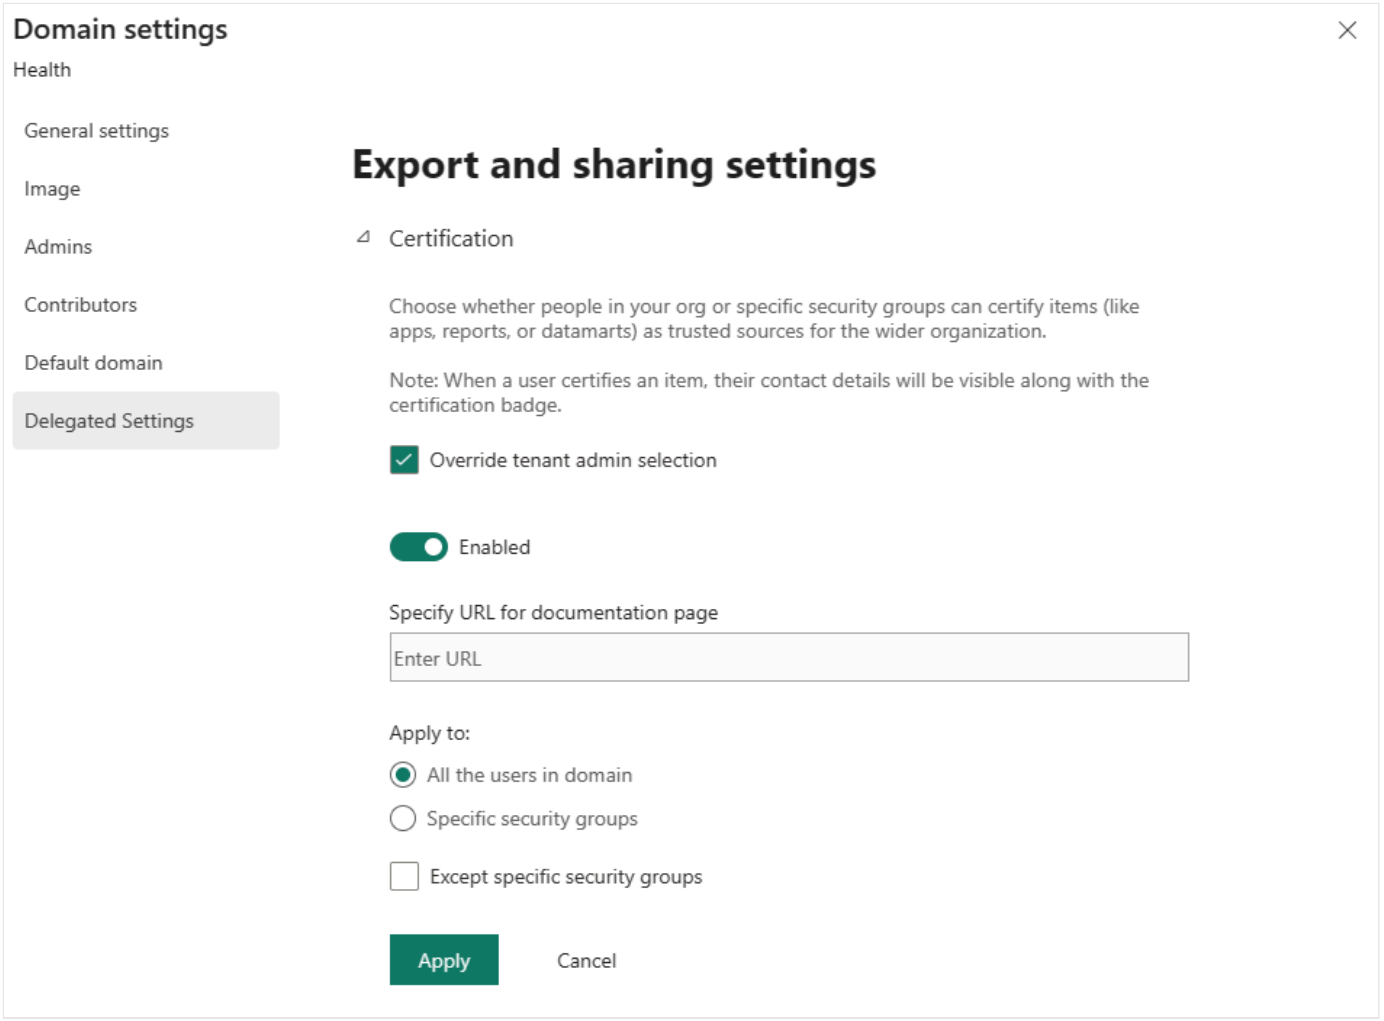

For descriptions of the things you need to set, see Set up certification.

Microsoft Fabric REST Admin APIs for domains

Most of the actions available from the UI are available through the Fabric REST Admin APIs for domains. For more information, see Domains API reference.

Track user activity on domains

Whenever a domain is created, edited, or deleted, that activity is recorded in the audit log for Fabric. You can track these activities in the unified audit log or in the Fabric activity log. For information about the information in the Fabric auditing schema that's specific to domains, see Audit schema for domains.

Related content

- Domain management tenant settings
- Domain design - best practices
- Domain-level default sensitivity labels
- Microsoft Fabric REST Admin APIs for domains
- Audit schema for domains
- Admin role in workspaces
- Blog: Easily implement data mesh architecture with domains in Fabric

---

Last updated on 05/01/2025

Best practices for planning and creating domains in Microsoft Fabric

Domains in Fabric help to organize business data according to the needs and goals of the organization and to facilitate consumption.

Domains also promote distributed governance by making possible for tenant settings to be delegated to domain admins who can set up appropriate controls for each domain.

Planning

When implementing domains, there are a few things to think about:

First, consider involving the following roles in the domain design phase:

- Center of excellence, business, and technical architects.
- Center of excellence, business, and technical leads and owners.
- Security and compliance officers.

Next map out answers to the following questions:

- Who is responsible for the data?
- What is the best structure for my organization's data?
- Do I need an extra level of hierarchy? (this consideration helps you decide whether and how to create subdomains within each domain).

Common organizational structures in industry

The following section describes several of the most common organizational structures in industry.

- **Functional structure**

  In a functional structure, the organization is divided into units (such as Finance, HR, Sales, etc.) based on the roles and business functions they perform. This structure has a clear hierarchy, centralized leadership, and well-defined responsibilities and authorities.

A functional structure enables specialization, scalability, and accountability. It also sets clear expectations and provides a direct chain of command.

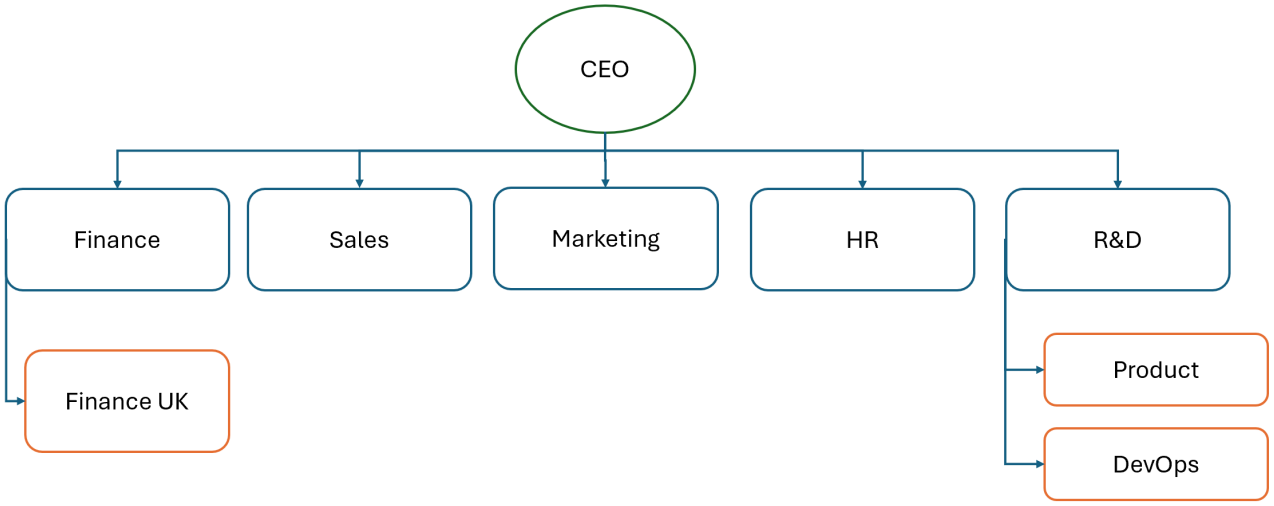

- **Product/project structure**

  A product/project based structure is suitable for companies that have multiple product lines or projects that require different teams and resources.

  A product/project structure allows the company to assign dedicated teams to each product or project, and fosters innovation and collaboration.

  A drawback of a product/project based structure is that it can create duplication of functions, competition for resources, and lack of coordination among teams.

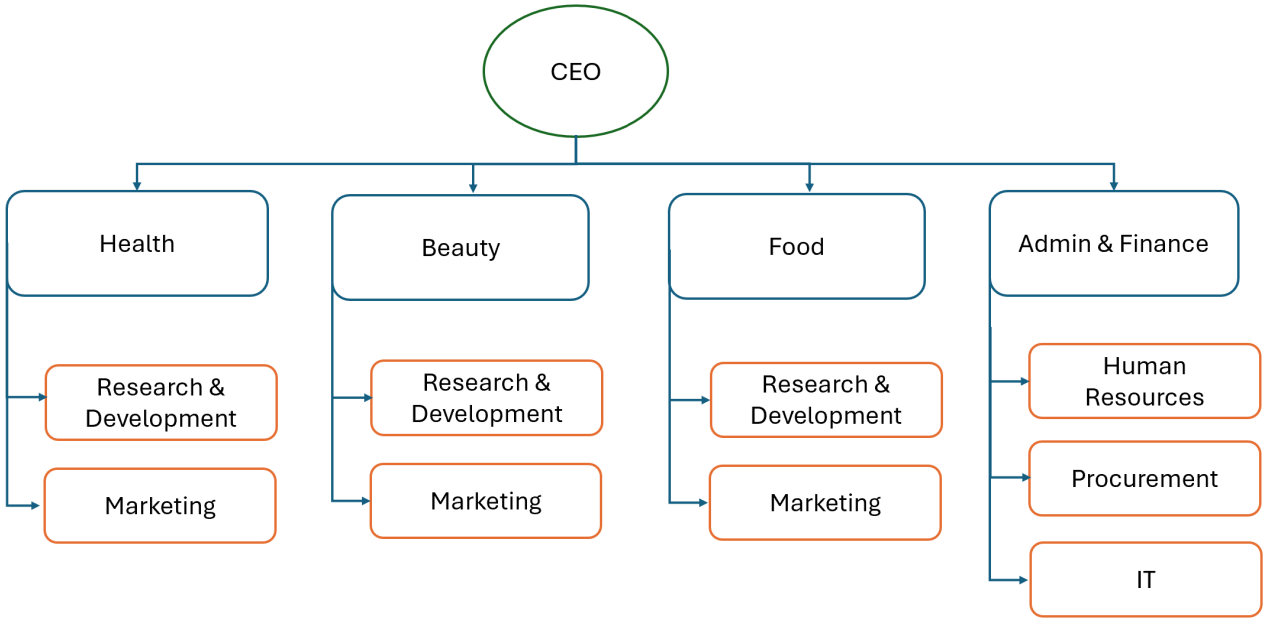

- **Process-based structure**

A process-based structure is suitable for companies that have standardized or repetitive processes that span different products or markets. A process-based structure allows the company to optimize the efficiency, quality, and consistency of each process and use the expertise and skills of the process teams. However, a process-based structure can also create silos, communication gaps, and reduced customer focus among teams.

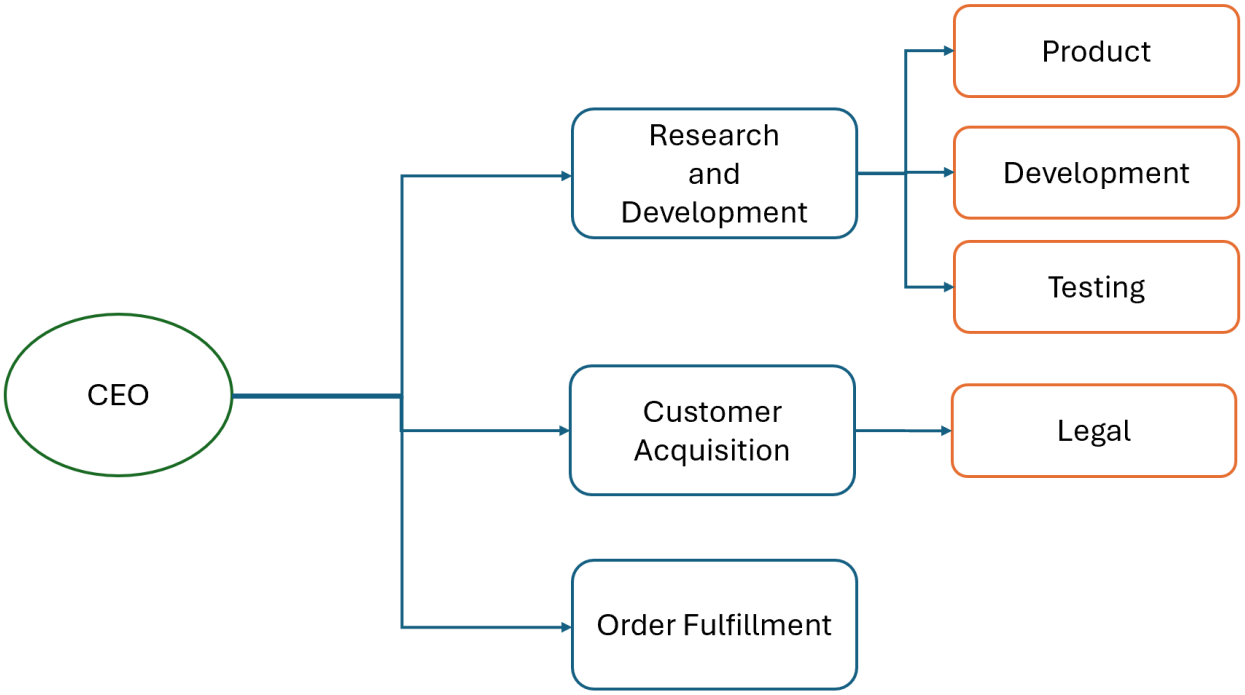

- **Region-based structure**

  A region-based structure is suitable for companies that operate in different regions or countries and need to adapt to the local environment, culture, and regulations. A region-based structure allows the company to delegate decision-making authority to the regional managers who can tailor their products and services to regional customer needs and references.

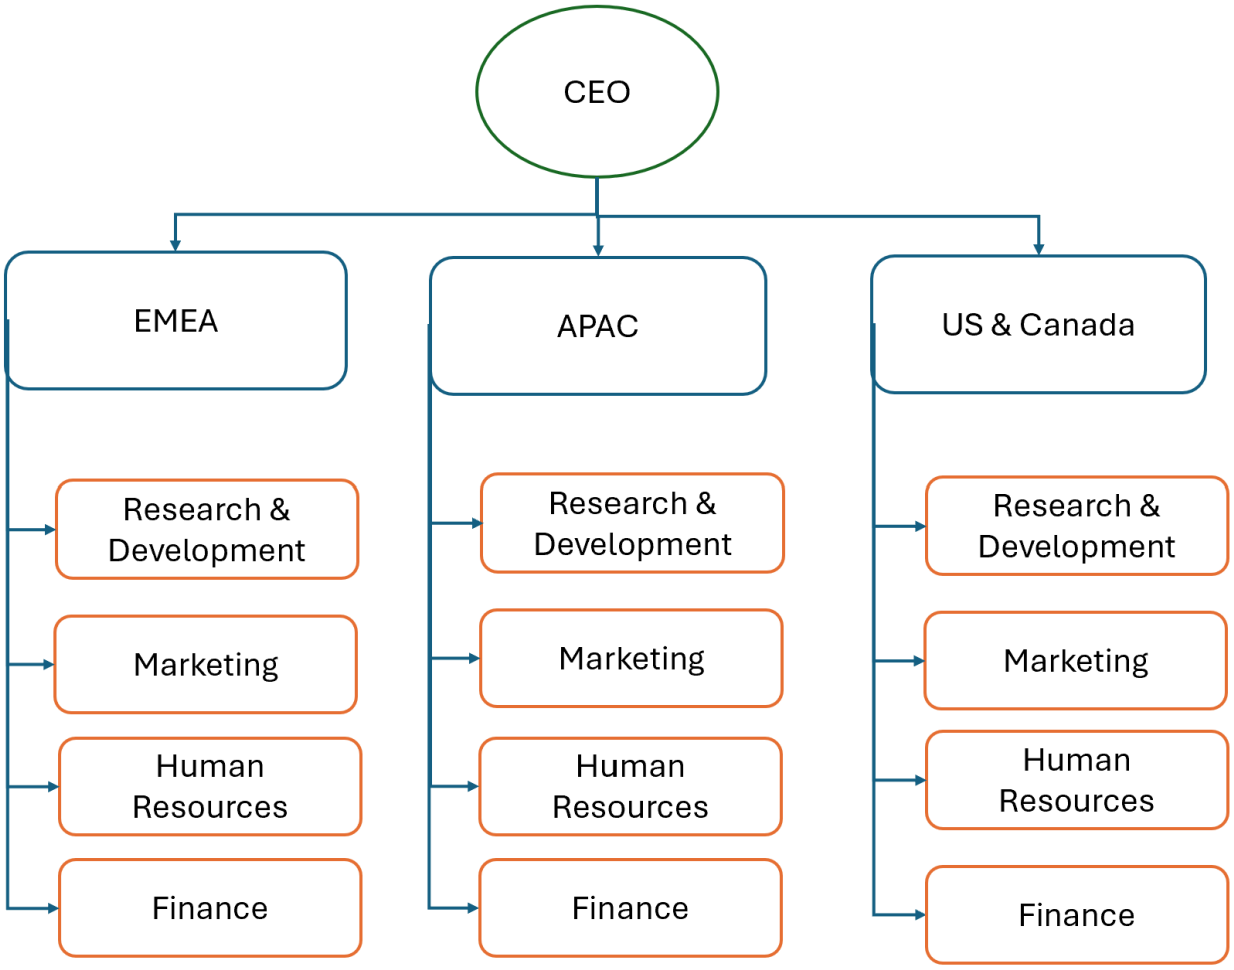

- **Mixed structure**

  A mixed structure is a combination of two or more organizational structures, such as functional, product, or market. A mixed structure can help the company balance the advantages and disadvantages of each structure and achieve greater efficiency and flexibility. For example, a company might have global functional departments for finance, human resources, and research and development, and product or market departments for specific business units or segments.

Subdomain structure

Subdomains can follow the same logic of the parent domain, or they can have their own structure based on organization needs. For example, an organization's domains might be built according to a function-based structure, while the subdomains might be built with a region-based structure.

Workspace assignment

After creating the structures of the domains and subdomains, the next step is to assign workspaces to each domain or subdomain. There are different ways to assign workspaces to

domains, depending on the naming convention and criteria used to create the domains and subdomains. Some possible methods are:

- **By workspace name**: This method is useful for organizations that follow a consistent and clear naming pattern for their workspaces and have names that reflect or relate to the relevant domain or subdomain. For example, if the domain is *Finance* and the subdomain is *Accounting*, then a workspace named *Finance-Accounting-Report* would be assigned to that *Accounting* subdomain.

- **By workspace owner**: This method is applicable for organizations that have a clear and stable ownership structure for their workspaces and have owners that correspond to the domains or subdomains. For example, if the domain is *Product* and the subdomain is *Fabric*, then a workspace owned by *Product Manager-Fabric* would be assigned to that *Fabric* subdomain.

- **By capacity**: This method is suitable for organizations that have adopted the data mesh architecture with capacities, and have capacities that match the domains or subdomains. For example, if the domain is *Marketing* and the subdomain is *Analytics*, then a workspace assigned to the capacity *Marketing-Analytics* would be assigned to that *Analytics* subdomain.

Related content

- Domains in Microsoft Fabric

Last updated on 10/16/2024

Audit schema for domains in Fabric

Whenever a domain is created, edited, or deleted, that activity is recorded in the audit log for Fabric. You can track these activities using Microsoft Purview Audit.

This article explains the information in the Fabric auditing schema that's specific to domains. This information is recorded in the OperationProperties section of the details side pane that opens when you select a domain-related activity on the Audit search page.

On the Audit search page:

1. Search for activities by their operation names.
2. When the search finishes, select the search to open the search results page.
3. On the search results page, select one of the search results to open the side pane that displays the record details.

For domains, the domain-specific details are found under the **OperationProperties** section of the side pane, in json format.

Expand table

Field	Type	Must appear in the schema	Value
OperationName	Edm.Enum	Yes	Activity name as described in the following table.
OperationProperties	Edm.Enum	Yes	Per the properties described in the following table.

Expand table

Activity flow	Activity operation name	Properties
Create domain/sub-domain	InsertDataDomainAsAdmin	**operationName**: - InsertDataDomainAsAdmin **operationProperties**: - DataDomainObjectId: <guid> - DataDomainDisplayName: <string> - ParentObjectId?: <guid>
Delete domain/sub-domain	DeleteDataDomainAsAdmin	**operationName**: - DeleteDataDomainAsAdmin **operationProperties**: - DataDomainObjectId: <guid> - DataDomainDisplayName: <string> - ParentObjectId?: <guid>

Activity flow	Activity operation name	Properties
Update domain/sub-domain	UpdateDataDomainAsAdmin	**operationName**: - UpdateDataDomainAsAdmin **operationProperties**: - DataDomainObjectId: <guid> - DataDomainDisplayName: <DataDomainObjectId> - ParentObjectId?: <guid>
Assign/Unassign workspace to the domain	UpdateDataDomainFoldersRelationsAsAdmin	**operationName**: - UpdateDataDomainFoldersRelationsAsAdm **operationProperties**: - DataDomainObjectId: <guid> - DataDomainDisplayName: <string> - ParentObjectId?: <guid> - FoldersToSetCounter?: <long> - FoldersToUnsetCount?: <long>
Unassign all workspaces to the domain	DeleteAllDataDomainFoldersRelationsAsAdmin	**operationName**: - DeleteAllDataDomainFoldersRelationsAsAd **operationProperties**: - DataDomainObjectId: <guid> - DataDomainDisplayName: <string> - ParentObjectId?: <guid>
Assign/Unassign workspaces to the domain as contributor	UpdateDataDomainFoldersRelationsAsContributor	**operationName**: - UpdateDataDomainFoldersRelationsAsContri **operationProperties**: - DataDomainObjectId: <guid> - DataDomainDisplayName: <string> - ParentObjectId?: <guid> - FoldersToSetCounter?: <long> - FoldersToUnsetCount?: <long>
Remove domain from workspace settings as workspace owner	DeleteDataDomainFolderRelationsAsFolderOwner	**operationName**: - DeleteDataDomainFoldersRelationsAsFolderO **operationProperties**: - DataDomainObjectId: <guid> - DataDomainDisplayName: <string> - ParentObjectId?: <guid> - FolderId?: <long>
Initiate/Process bulk assign domain by workspace owners	BulkAssignDataDomainByWsOwnersAsAdmin?	
Initiate/Process bulk assign domain by capacities	BulkAssignDataDomainByCapacitiesAsAdmin?	
Add/Delete/Update domain access	UpdateDataDomainAccessAsAdmin	**operationName**: - UpdateDataDomainAccessAsAdmin **operationProperties**:

Activity flow	Activity operation name	Properties
		- DataDomainObjectId: <guid> - DataDomainDisplayName: <string> - ParentObjectId?: <guid> - Value: <long>: • 0 - None • 7 - Contributor • 15 - Admin - UsersToSetCounter?: <long> - UsersToUnsetCounter?: <long> - GroupsToSetCounter?: <long> - GroupsToUnsetCounter?: <long>
Add/Delete/Update default domain	UpdateDefaultDataDomainAsAdmin	**operationName**: - UpdateDefaultDataDomainAsAdmin **operationProperties**: - DataDomainObjectId: <guid> - DataDomainDisplayName: <string> - ParentObjectId?: <guid> - UsersToSetCounter?: <long> - UsersToUnsetCounter?: <long> - GroupsToSetCounter?: <long> - GroupsToUnsetCounter?: <long>
Add/Delete/Update contributors	UpdateDataDomainContributorsScopeAsAdmin	**operationName**: - UpdateDataDomainContributorsScopeAsAc **operationProperties**: - DataDomainObjectId: <guid> - DataDomainDisplayName: <string> - ParentObjectId?: <guid> - Value: <long>: • 0 - AllTenant • 1 - SpecificUsersAndGroups • 2 - AdminsOnly
Set/Remove domain branding	UpdateDataDomainBrandingAsAdmin	**operationName**: - UpdateDataDomainBrandingAsAdmin **operationProperties**: - DataDomainObjectId: <guid> - DataDomainDisplayName: <string> - ParentObjectId?: <guid> - Value: <long> // Branding ID
Updated delegation at domain level	UpdateDomainTenantSettingDelegation	
Deleted override at domain level		

Related content

- Domains
- Domains management tenant settings
- Microsoft Fabric REST Admin APIs for domains
- Track user activities in Fabric

Last updated on 03/24/2024

Tags in Microsoft Fabric

Tags in Microsoft Fabric let you apply additional metadata to workspaces and items, making it easier to categorize, organize, and discover data.

Tags are configurable text labels, such as *Sales – FR 2025*, *HR – Summer Event*, or *FY 2025*, that admins define at the tenant or domain level.

How tags work

Here's how tags flow from definition to discovery:

1. **Admins define tags.** Fabric admins create tenant-level tags; Fabric or domain admins create domain-level tags.

    - **Tenant-level tags** are available across all items and workspaces in the tenant. They're suitable for broad classifications, compliance, or security labels that apply universally across your organization.
    - **Domain-level tags** are available only for items and workspaces assigned to that domain. If a workspace isn't assigned to a domain, only tenant-level tags are available. Domain-level tags let domain owners implement more targeted governance policies that reflect the needs of their area. A domain-level tag can't duplicate a tenant-level tag name, but the same name can exist across different domains.

2. **Workspace admins apply tags to workspaces.** Workspace tags let you tag at a broad level without tagging each item individually. They're useful for cost attribution, governance reporting, and policy enforcement. A workspace can have up to 10 tags, counted independently from item tags. Non-admin members (Viewer, Member, Contributor) can view workspace tags but can't modify them.

3. **Data owners apply tags to items.** Users with Write or Contributor permissions can apply up to 10 tags per item. Available tags include tenant-level tags and, if the item's workspace is assigned to a domain, that domain's tags.

4. **Users discover content by tag.** Once tags are applied, users can filter or search for relevant content across the catalog, workspace list, and other surfaces.

5. **Admins govern at scale.** Use the metadata scanning (scanner) APIs to programmatically fetch tag associations for items and workspaces and feed them into downstream

governance solutions.

How tags enhance data discoverability

Once tags are applied, they enhance visibility across multiple surfaces:

Workspace tags

- **Workspace list:** A tag icon appears next to the workspace name in the workspaces list panel. Hover to view applied workspace tags.

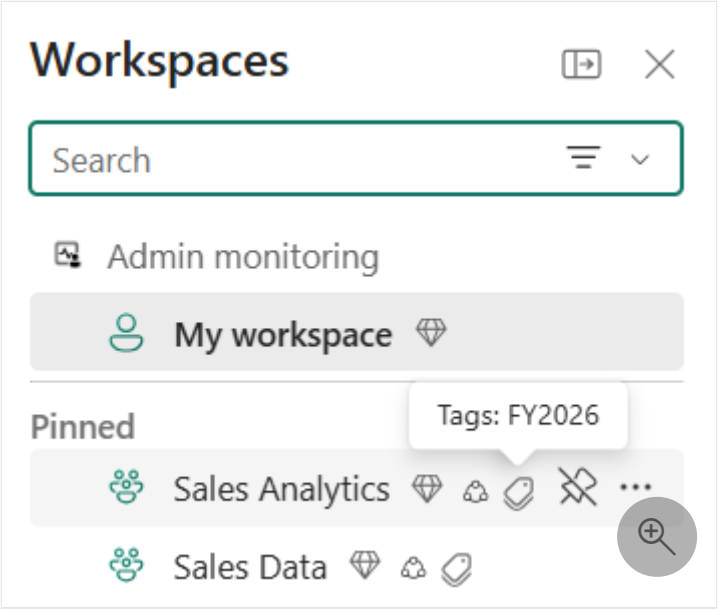

- **Workspace list filtering:** Filter workspaces by applied tags in the workspaces list panel and OneLake Catalog Explorer.

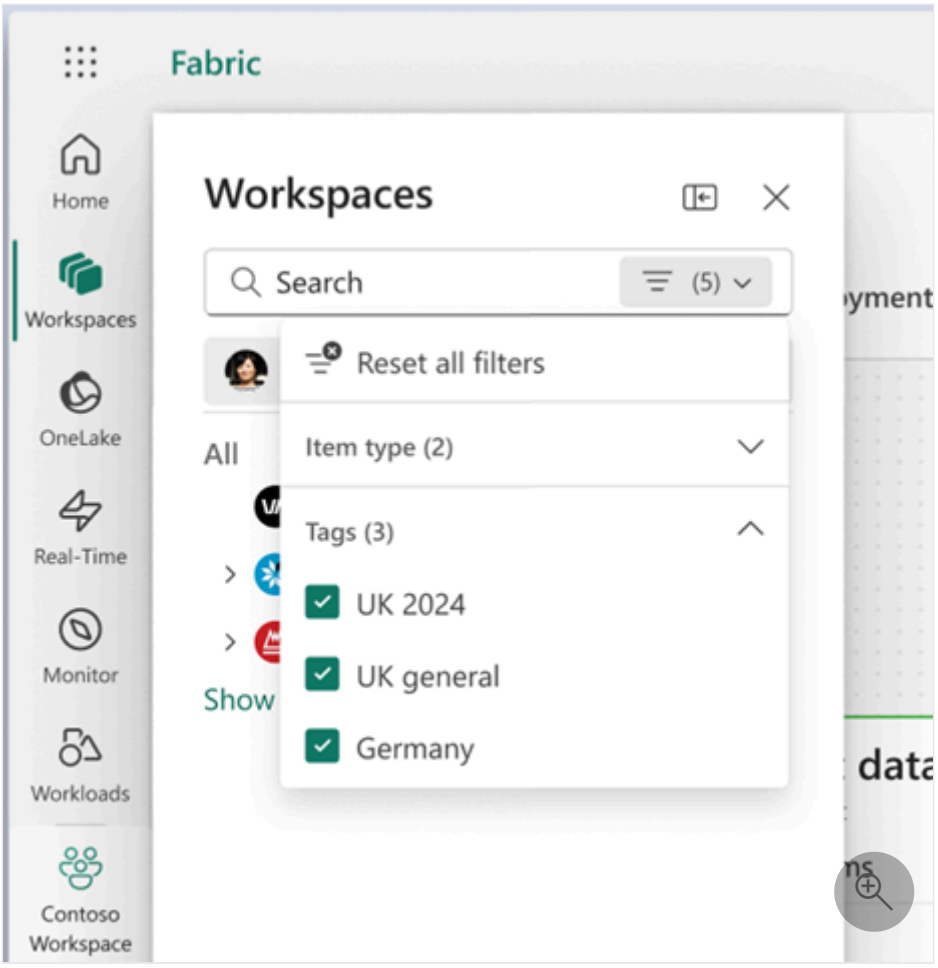

- **Workspace view:** Workspace tag names appear directly in the workspace screen.

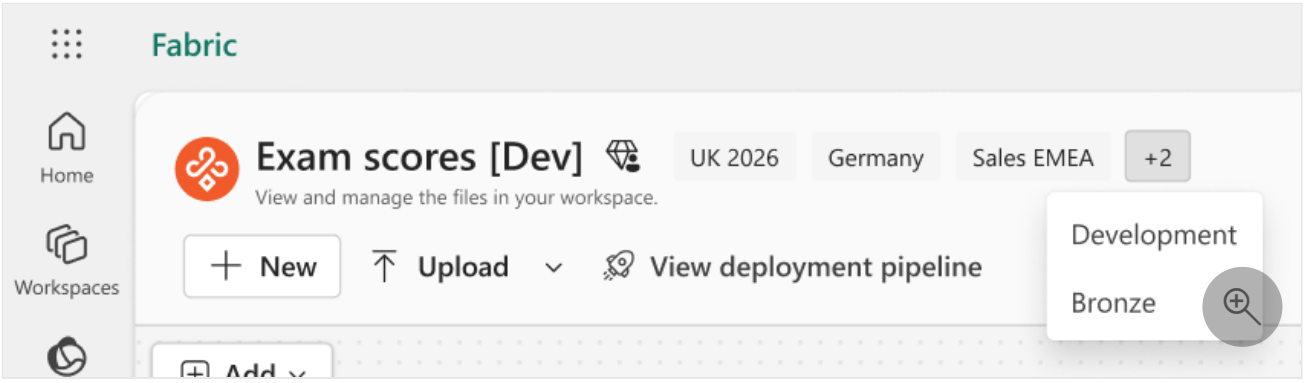

Item tags

- **Item list views:** A tag icon appears next to the item name. Hover to view applied tags.

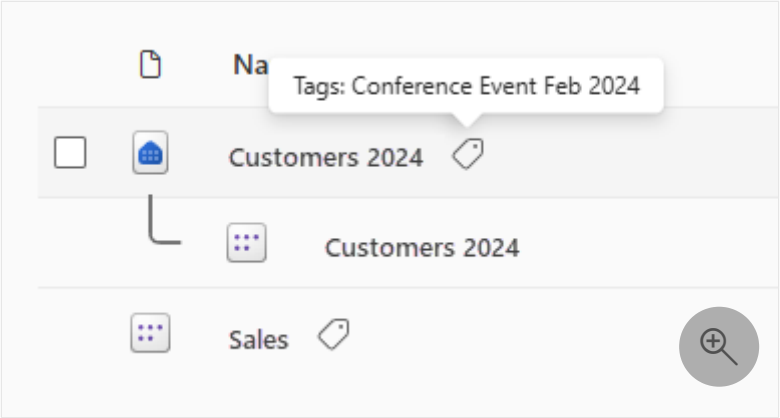

- **Workspace item list:** Filter items list by assigned tag.

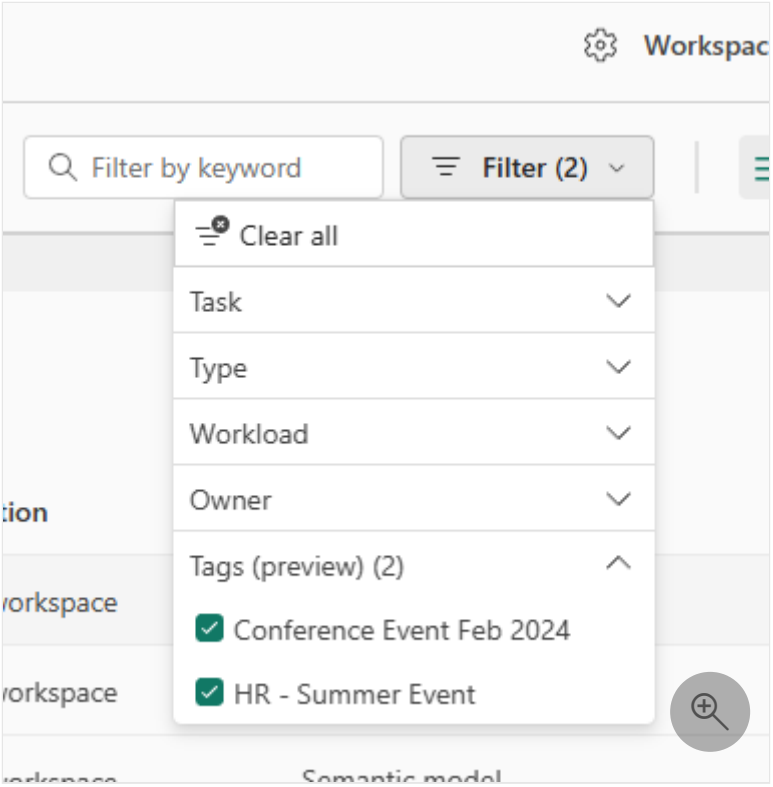

- **Item details:** Tags are shown in the OneLake Catalog item details pane of each item.

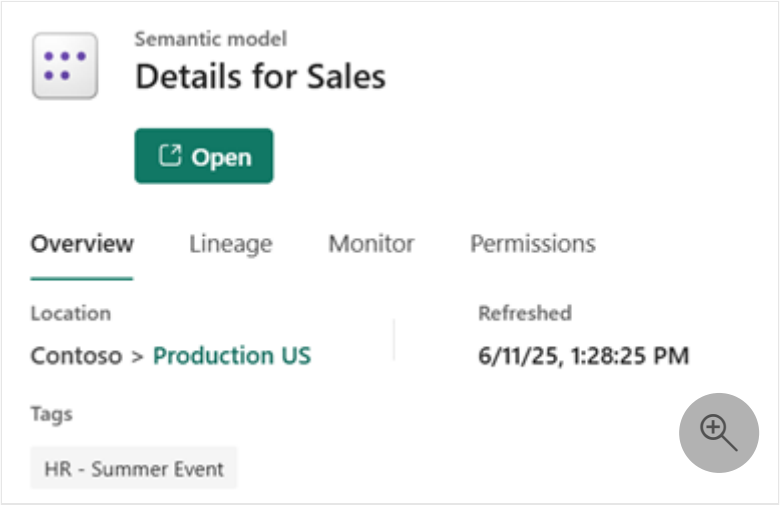

- **Flyout card:** When editing an item, select the item name or sensitivity label to view tags.

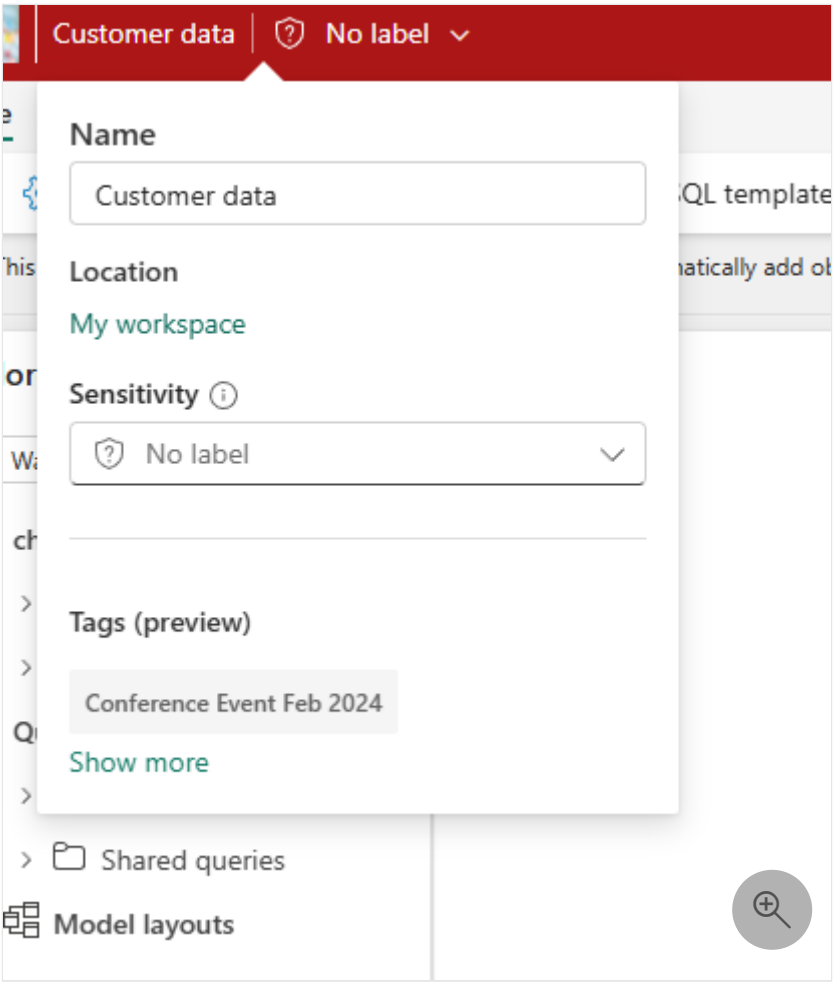

- **Lineage view:** Tags appear in workspace lineage and item-level lineage views.

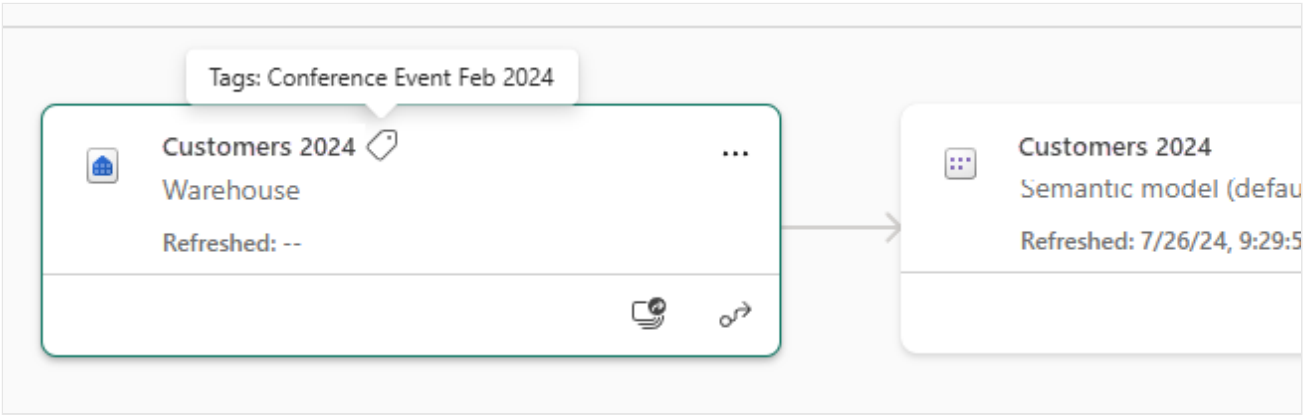

- **Global search:** Search by tag name in global search to find all relevant results, along with other metadata like item owner and location.

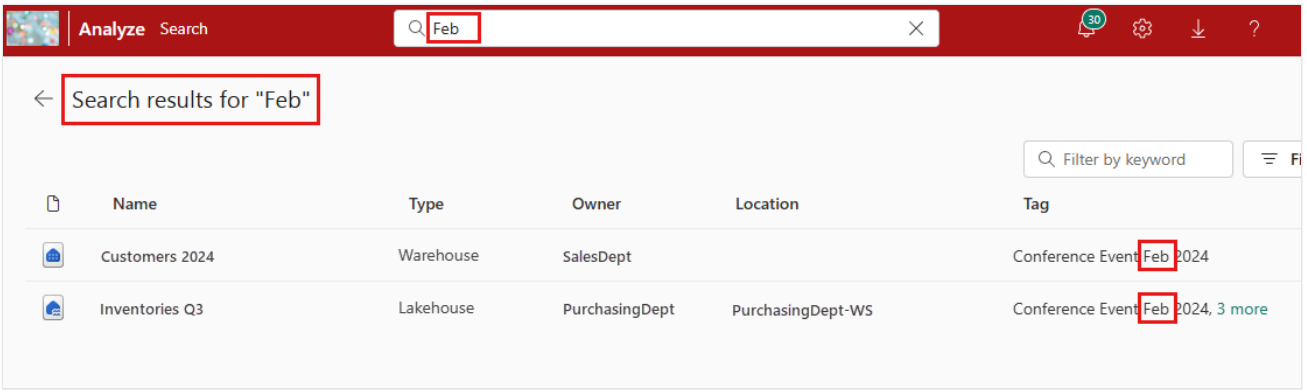

Tags in APIs

- **Tag management APIs**: Use the Fabric REST Admin APIs for tags to create, update, delete, and list tags at the tenant and domain levels.
- **Item APIs**: Use the Apply Tags and Unapply Tags APIs to add or remove tags from individual items. Use the List Tags API to retrieve applied tag IDs and names for an item.
- **Scanner API**: The metadata scanning (scanner) APIs include tag data, so governance and discovery solutions can retrieve tag assignments at scale for both items and workspaces.

  For every applicable item and workspace returned in a scan, the payload includes a `tags` field containing a list of applied tag IDs. Use the List Tags API to resolve the IDs to tag names.

Considerations and limitations

Limits

- A maximum of 10,000 unique tags can be created in a tenant.
- An item or workspace can have a maximum of 10 tags applied to it at a time. Workspace and item tag limits are counted independently.
- There's no limit on the number of tagged items or workspaces.

Propagation and visibility

- After you apply a tag to an item, the icon might take several hours to appear next to the item name. It might also take time before the item appears in global search results when you use the tag name as a search term.
- Workspace tags are visible only where other workspace metadata is visible to you.
- When a workspace is moved to a different domain, existing domain-level tags remain applied. However, those tags might not be available in the new domain, so you can't reapply them if removed.

Related content

- Create and manage a set of tags
- Apply tags
- Microsoft Fabric REST Admin APIs for tags

Last updated on 03/18/2026

Create and manage a set of tags

This article describes how to create and manage a set of tags in Microsoft Fabric. The target audience is Fabric and domain administrators who want to define and manage tags within their organization or specific domains.

For more information about tags, see Tags in Microsoft Fabric.

As a Fabric or domain administrator, you can:

- **Create a set of tags** that users in your organization or domain can use to categorize their Fabric items.
  - **Rename tags.** If desired or necessary, you can rename a tag. When you rename a tag, its name changes wherever the tag is applied.
  - **Delete a tag.** If you decide that a tag isn't needed, you can delete it from the set of tags you defined. When you delete a tag from the set of defined tags, the tag is removed from all items where it was previously applied.

Tags you define can also be applied to workspaces by workspace admins via workspace settings. Applied workspace tags are visible as indicators on workspace pages and can be used for discovery and filtering in the workspaces list and in OneLake Catalog Explorer.

Prerequisites

You must have a **Fabric administrator** role to create and manage tenant-level tags.

You must have a **Fabric administrator** or **Domain administrator** role to create and manage domain-level tags within your assigned domains.

Limits

- Up to 10,000 tags of all types can be created in a tenant.
- Workspace admins can apply up to 10 tags per workspace.
- Only tags defined at the tenant or domain level are available for selection in workspace settings.

Create tenant-level tags

This section shows how Fabric administrators create a set of tenant-level tags. These tags are available across the entire Fabric tenant.

1. Open the admin portal and select **Tags**. The **Tags** tab opens. All currently defined tenant-level tags are listed on the tab.

2. Select + **New tag**. If no tags are defined yet, select **Add your first tags**.

3. In the dialog that appears, provide names for your new tags. Use comma-separated names to create more than one tag at a time. Tag names can be up to 40 characters long and can contain letters, numbers, spaces (not at the beginning), and special characters (must also contain letters and/or numbers).

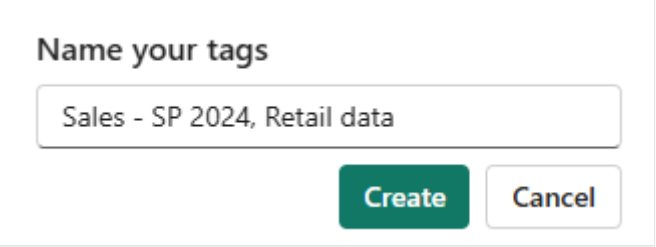

4. When done, select **Create**. The new tags are added to the set of tags.

Create domain-level tags

This section shows how domain administrators create and manage tags specific to their assigned domains. These tags are only available for items within workspaces assigned to that domain and subdomains under it.

Tag uniqueness rules

When creating domain-level tags, consider the following uniqueness rules:

- A tag created at the tenant level can't be duplicated at the domain level and vice-versa. If you attempt to create a tag that duplicates an existing one at a different scope (tenant or domain), an error message indicates the existing tag and its scope.

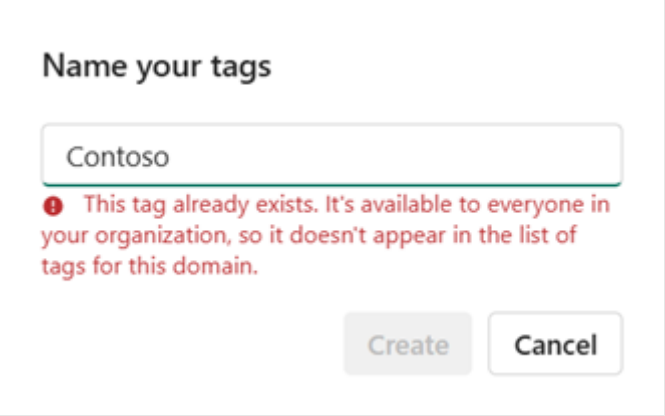

- You can create the same tag name in different domains. For example, you can add a tag named "FY2025" separately in both the "Sales" and "Marketing" domains.
- If you reassign a workspace from one domain to another, any domain-specific tags already applied to items in that workspace remain on those items. This approach maintains continuity and helps prevent confusion about previous tagging. However, you can't reapply tags from a previous domain to items if the workspace is no longer associated with that domain.

Steps to create domain-level tags

1. Open the admin portal.
2. In the admin portal, select **Domains**.
3. Choose the specific domain for which you want to create or manage tags.
4. In the domain settings, select **Tags** to display all currently defined domain-level tags for that domain.
5. To create a new domain tag, select **+ New tag** or **Add your first tags** if none exists.
6. In the dialog, provide names for your new tags. You can use comma-separated names to create multiple tags simultaneously.
7. Tag names can contain letters, numbers, spaces (not at the beginning), and special characters (must also contain letters and/or numbers).
8. Select **Create** to add the new tags to the domain's set of tags.

Rename a tag

You can rename both tenant-level and domain-level tags.

1. Open the admin portal and select **Tags** (for tenant-level tags) or navigate to **Domains** and then **Tags** within your specific domain (for domain-level tags).
2. Open the options menu of the tag you want to rename and select **Rename**.

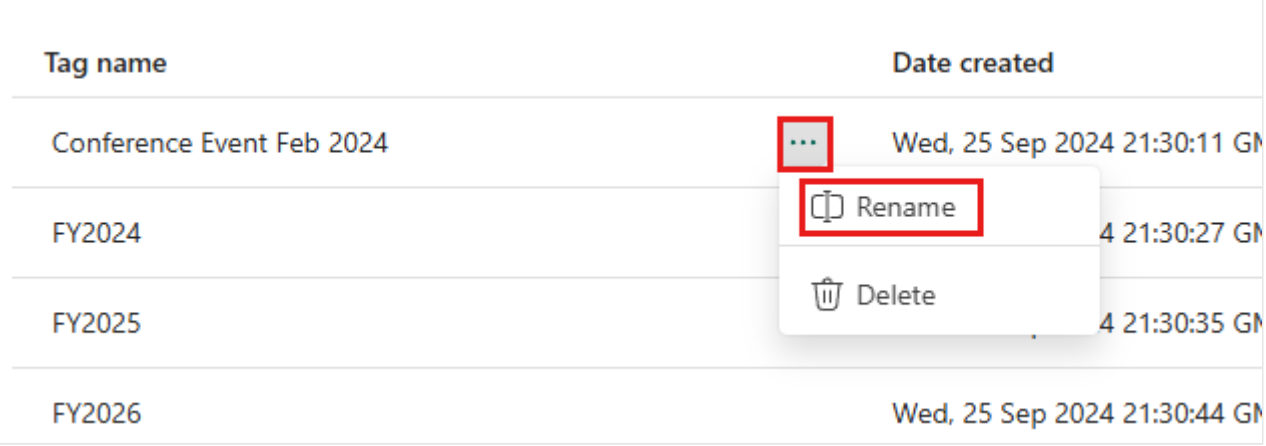

3. Change the name as desired and then select **Rename**. The new name is reflected wherever the tag is applied.

Delete a tag from the set of tags

You can delete both tenant-level and domain-level tags.

1. Open the admin portal and select **Tags** (for tenant-level tags) or navigate to **Domains** and then **Tags** within your specific domain (for domain-level tags).

2. Open the options menu of the tag you want to delete and select **Delete**.

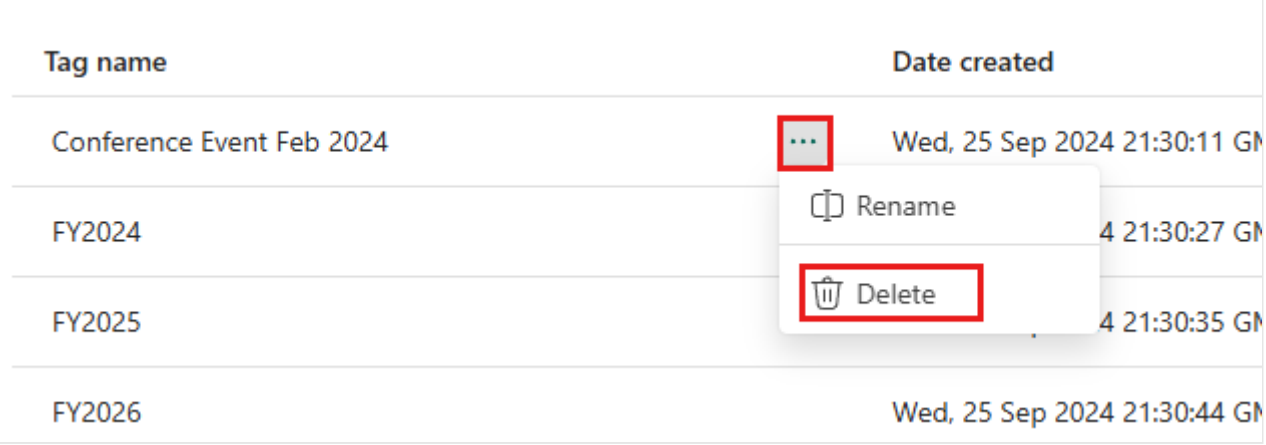

   The tag is deleted from the set of defined tags and is removed from all items where it was previously applied.

Create and manage tags programmatically

All actions related to creating and managing tags can be performed programmatically via APIs. In addition, tags are included in Metadata Scanning (Scanner) APIs so that governance and discovery solutions can harvest tag assignments at scale.

Tags REST APIs

Use the Fabric REST Admin APIs for tags to programmatically create, list, rename and delete tenant and domain-level tags, allowing you to automate tag lifecycle and governance workflows. For more information, see Fabric REST Admin APIs for tags.

Tags in metadata scanning (scanner) APIs

The Scanner APIs include applied tags for each scanned item. For every applicable item returned in a scan, the payload includes a `tags` field containing a list of applied tag UUIDs. To resolve tag IDs to tag names, use the List Tags Admin REST API.

For more information, see Metadata scanning overview.

Related content

- Tags overview
- Apply tags to items
- Fabric REST Admin APIs for tags

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 04/23/2026

Apply tags to items and workspaces in Microsoft Fabric

This article describes how to apply tags to items and workspaces in Microsoft Fabric.

For more information about tags, see Tags in Microsoft Fabric.

Prerequisites

- You must be a workspace admin to apply or remove tags on a workspace. All workspace roles (Viewer, Member, Contributor, and Admin) can view workspace tags, but only workspace admins can modify them.
- You must have Write or Contributor permissions on an item to apply or remove tags from it.
- The tags available in the dropdown are defined and managed by your tenant or domain admins. If you don't see the expected tags, contact your admin. For more information, see Create and manage a set of tags.

Apply tags to a workspace

1. Open the workspace settings and go to the **Tags** section.

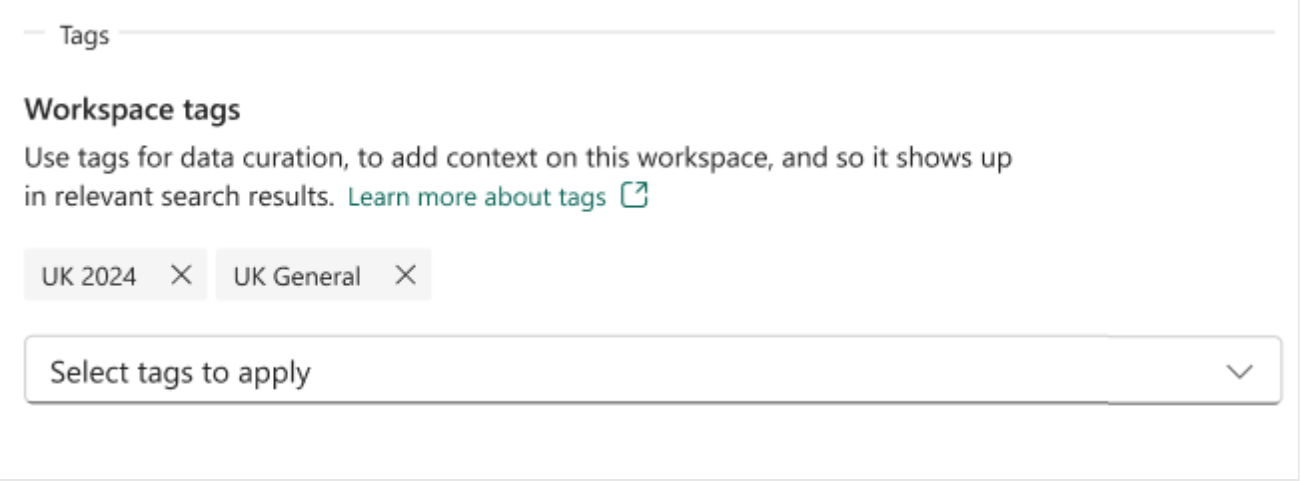

2. Use the **Select tags to apply** dropdown list to select one or more tags for your workspace. A workspace can have up to 10 tags.

> **Note**

If the drop-down is disabled, you might not have workspace admin permissions.

3. When done, close the settings pane.

Remove tags from a workspace

1. Open the workspace settings and go to the **Tags** section.

   You'll see all tags currently applied to the workspace.

2. Select the **X** next to the names of the tags you want to remove from the workspace.

3. When done, close the settings pane.

Apply tags to an item

1. Open the item's settings and go to the **Tags** tab.

2. Use the **Select tags to apply** dropdown list to select one or more tags for your item. An item can have up to 10 tags.

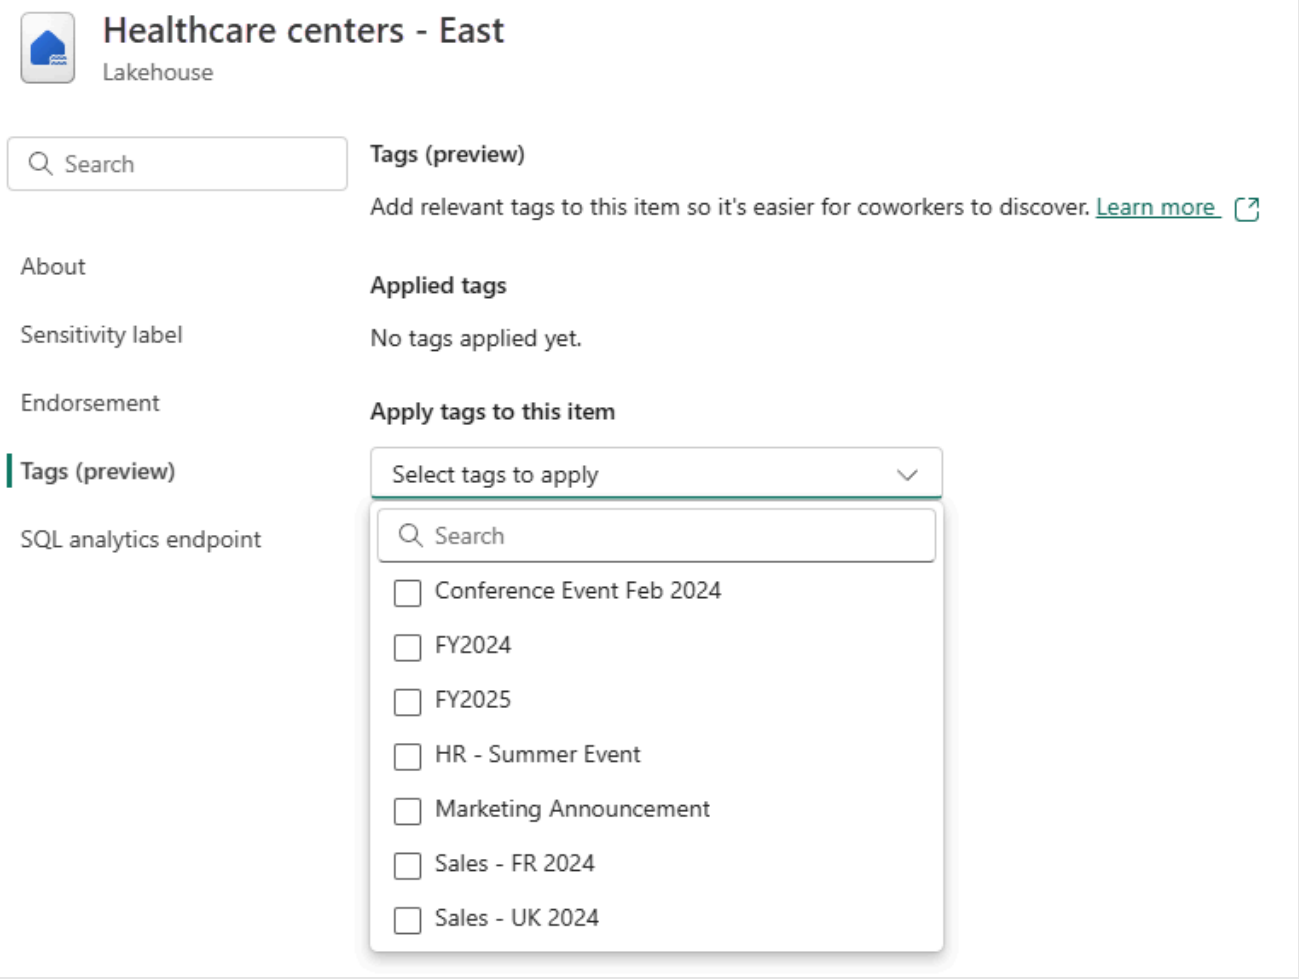

> **Note**
>
> If the **Select tags to apply** drop-down is disabled, you might not have Write or Contributor permissions on the item.

3. When done, close the settings pane. For Power BI items, select **Save** or **Apply**.

Remove tags from an item

1. Open the item's settings and go to the **Tags** tab.

   All the tags applied to the item appear under **Applied tags**.

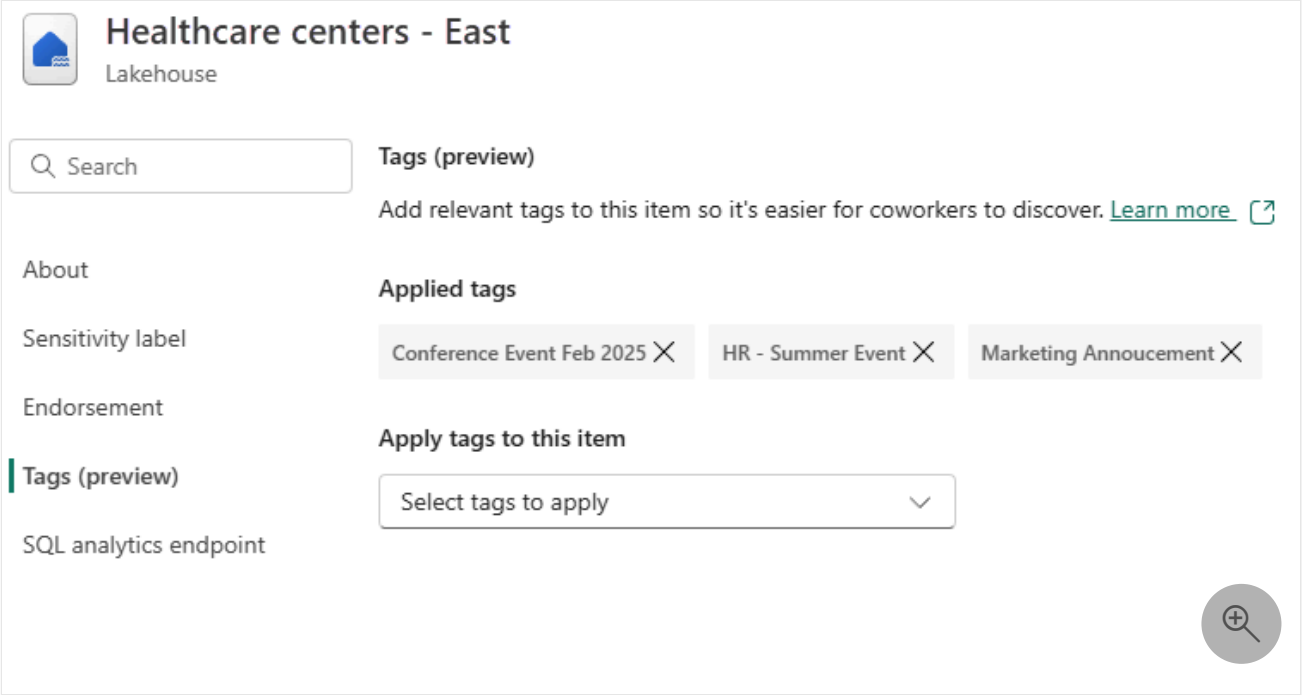

2. Select the **X** next to the names of the tags you want to remove from the item.

3. When done, close the settings pane.

Apply or remove tags using APIs

You can apply or remove tags programmatically using REST APIs:

- For workspaces, use the Apply Workspace Tags and Unapply Workspace Tags APIs to add or remove tags from a workspace.
- For items, use the Apply Tags and Unapply Tags APIs to add or remove tags from an item.

Related content

- Tags overview
- Create and manage a set of tags
- Fabric REST Admin APIs for tags

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 04/23/2026

Lineage in Fabric

In modern business intelligence (BI) projects, understanding the flow of data from the data source to its destination can be a challenge. The challenge is even bigger if you've built advanced analytical projects spanning multiple data sources, data items, and dependencies. Questions like "What happens if I change this data?" or "Why isn't this report up to date?" can be hard to answer. They might require a team of experts or deep investigation to understand. Fabric's lineage view helps you answer these questions.

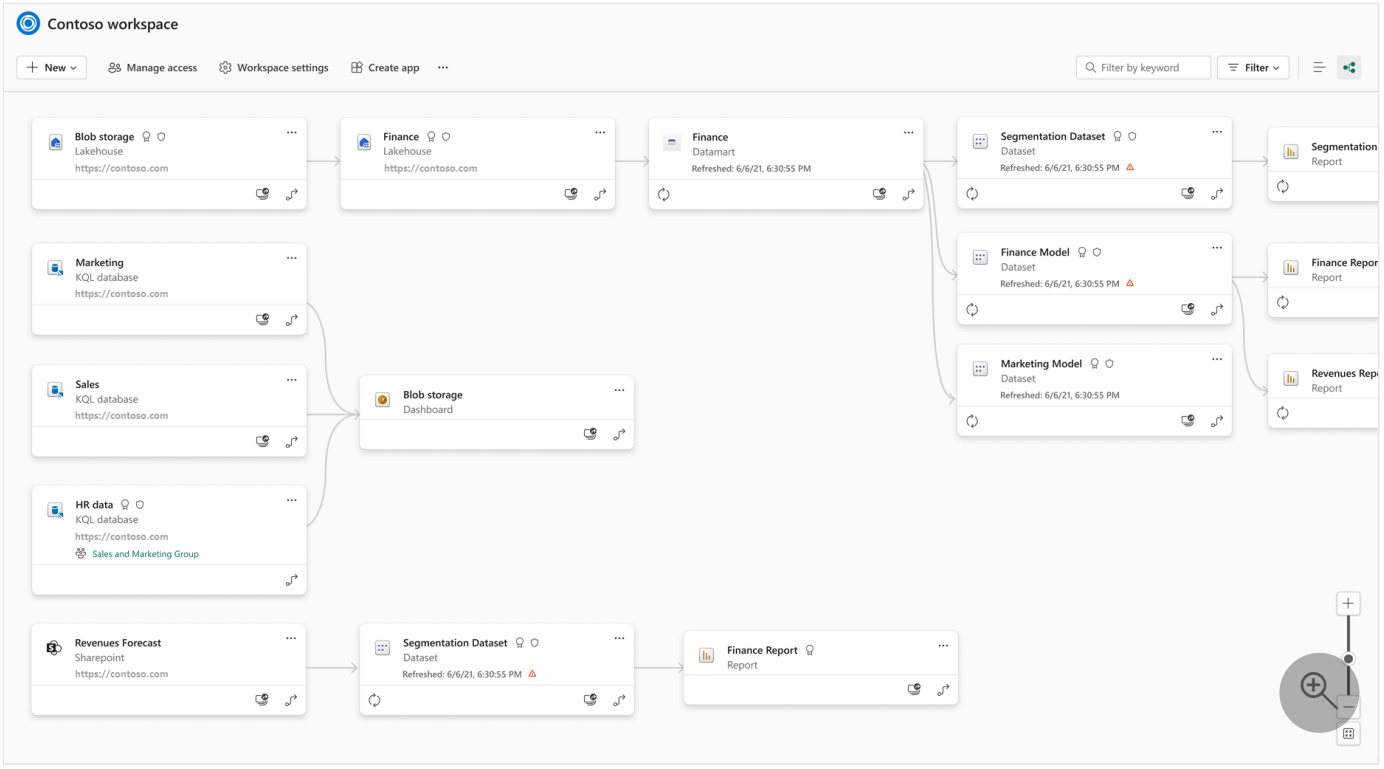

In lineage view, you see the lineage relationships between all the items in a workspace, as well as data sources external to the workspace one-step upstream.

Every workspace automatically has a lineage view.

Permissions

Any user with a role in a workspace can access that workspace's lineage view. However, users with the *Viewer* role won't see data sources.

Open lineage view

Lineage is accessible from multiple locations. Typically, you can get to it

- From the workspace toolbar

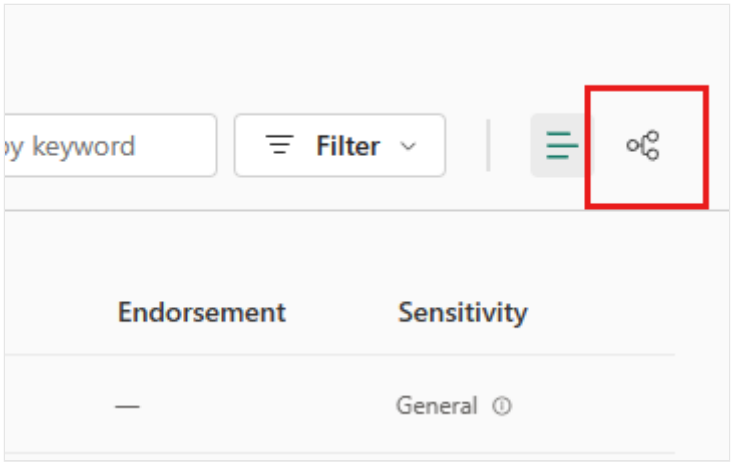

- From an item's option menu (for instance, in the OneLake catalog)

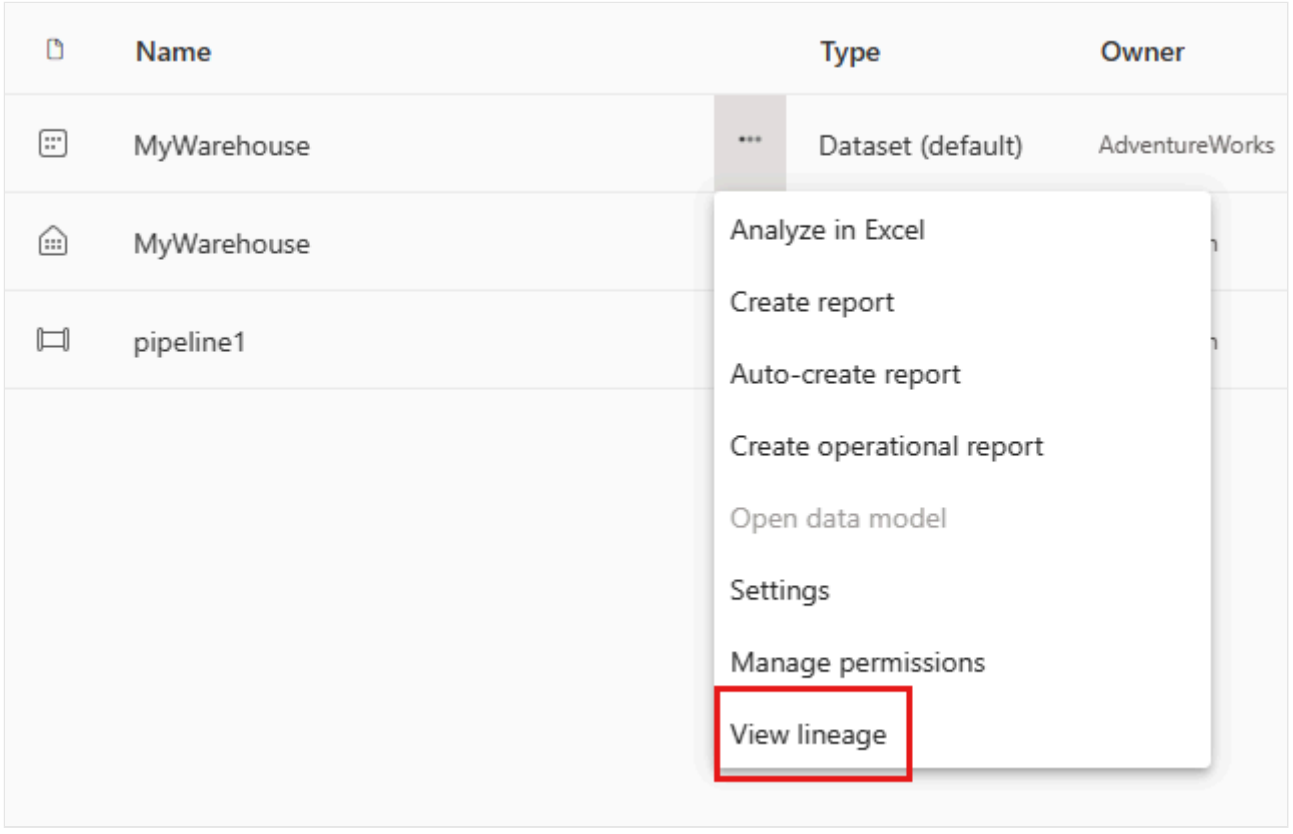

- From the menu items at the top of the item's details page

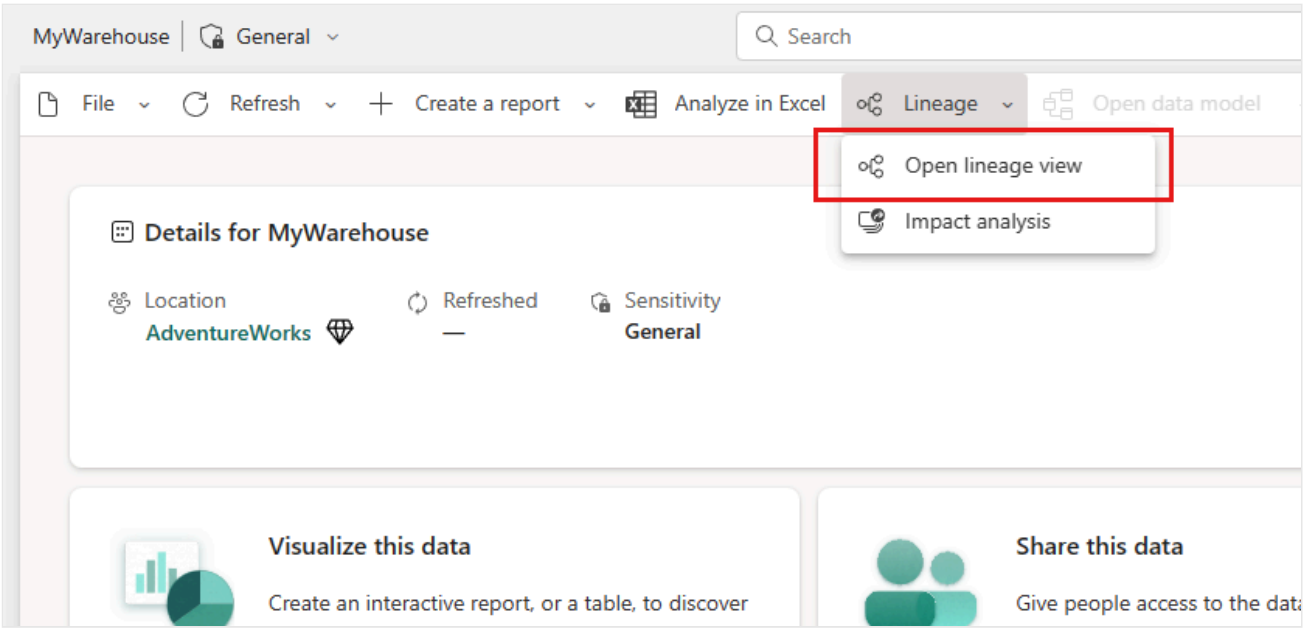

What do you see in lineage view

When you open lineage view on an item, you'll see the connections between all the items in the workspace the item is located in.

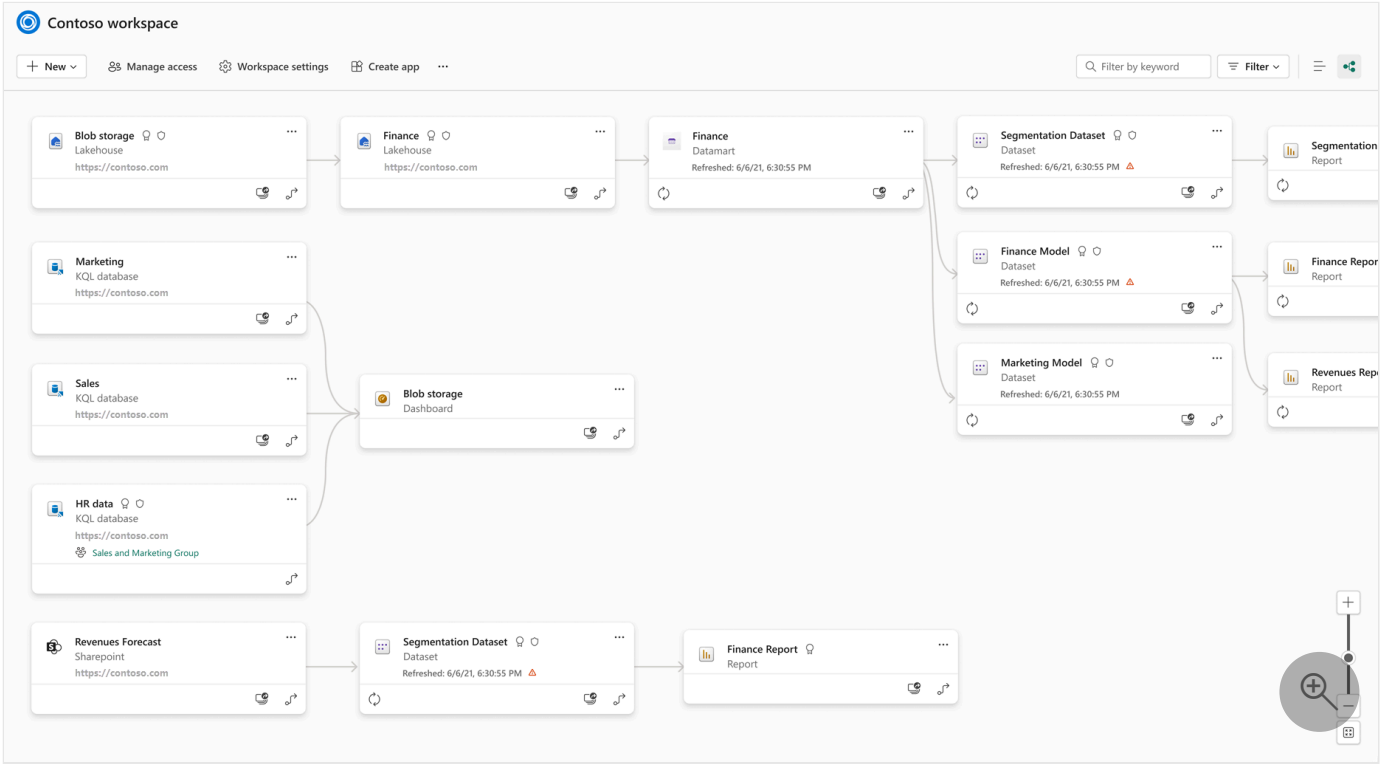

Lineage view displays:

- All the items in the workspace and how they are connected to each other.
- Upstream connections outside the workspace, one level up. You can tell if an item is outside the workspace because the name of the external workspace is indicated on the card, as illustrated by the HR Data KQL Database in the image above.

Downstream items in different workspaces aren't shown. To explore an item's downstream connections outside the workspace, open the item's impact analysis.

Items are represented by cards that provide some information about the item.

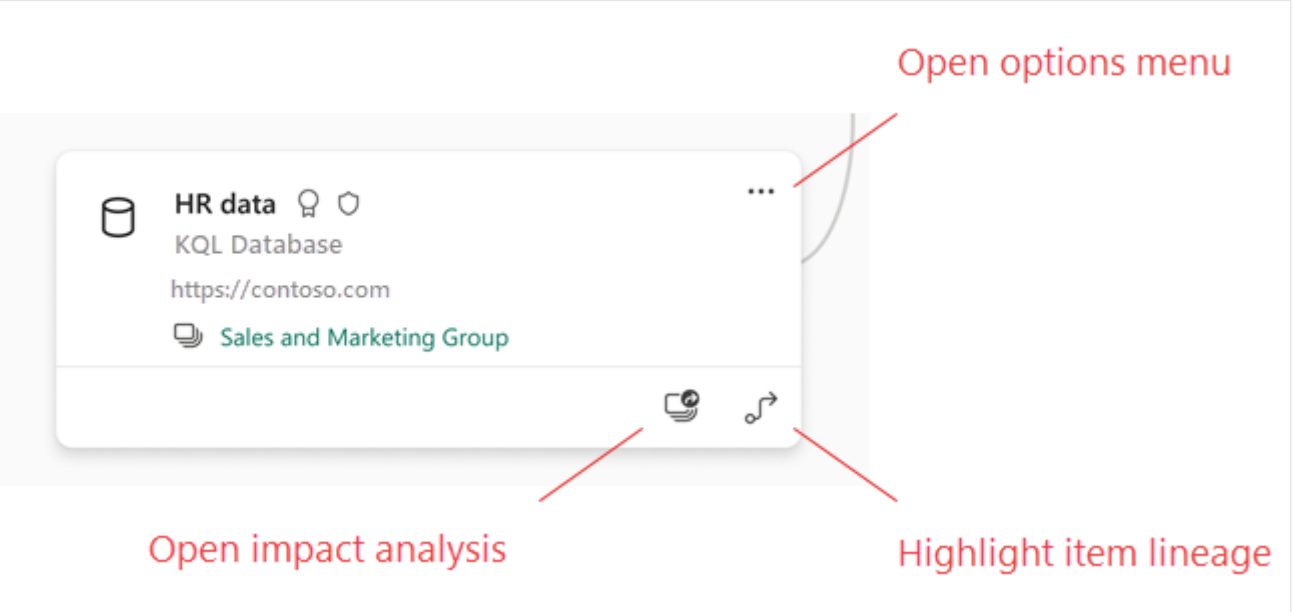

Data sources

You see the data sources from which the semantic models and dataflows get their data. On the data source cards, you see information that can help identify the source. For example, for Azure SQL server, you also see the database name.

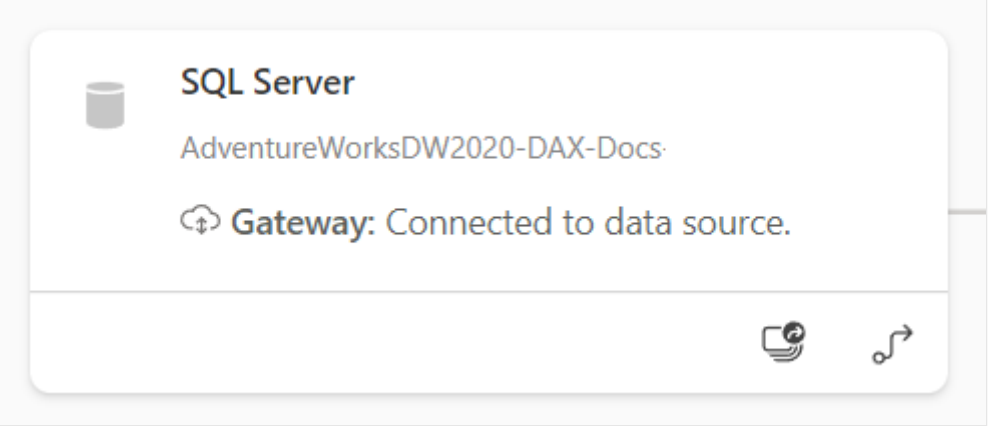

Highlight an item's lineage

To highlight the lineage for a specific item, select the arrow at the bottom right corner of the card.

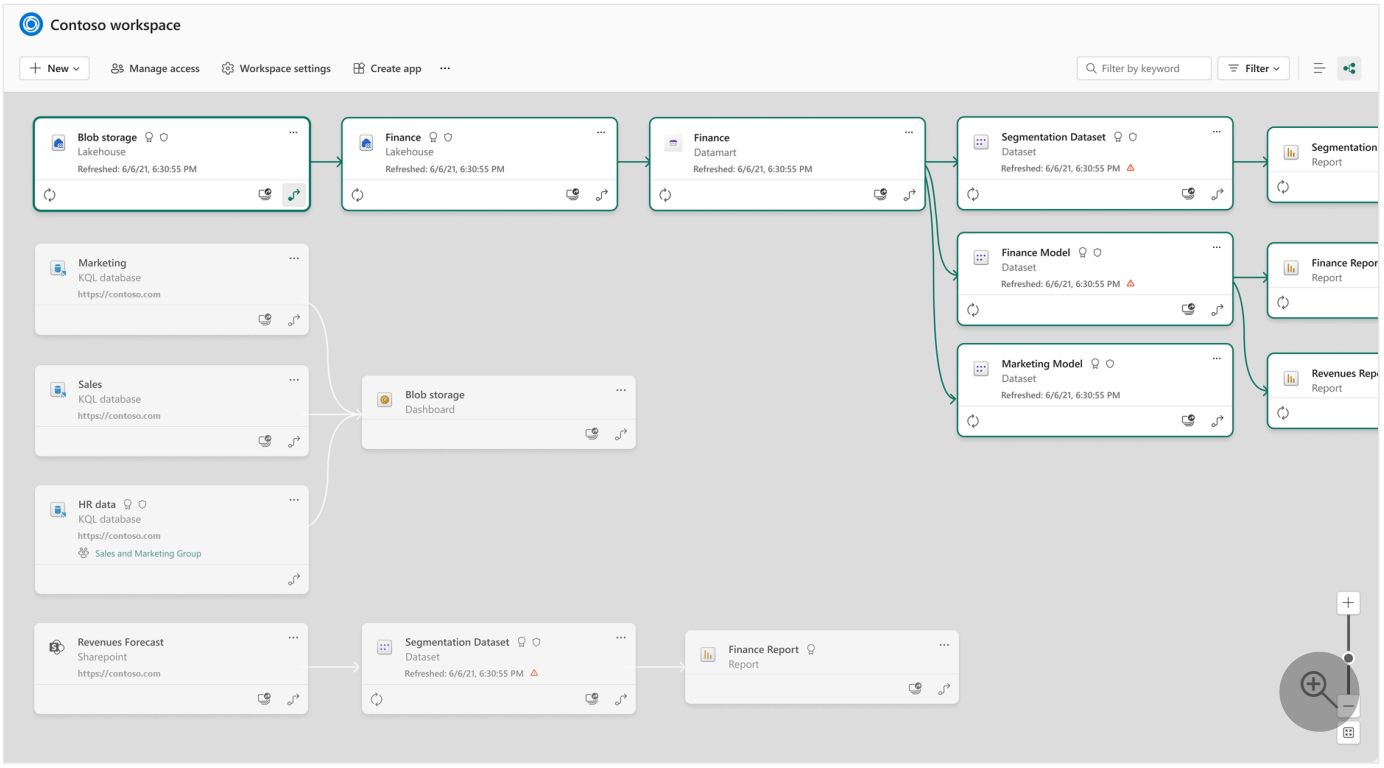

Fabric highlights all the items related to that item, and dims the rest.

Zoom and full screen

Lineage view is an interactive canvas. You can use the mouse and touchpad to navigate in the canvas, as well as to zoom in or out.

- To zoom in and out, use either the menu in the bottom-right corner or your mouse or touchpad.
- To have more room for the graph itself, use the full screen option at the bottom-right corner.

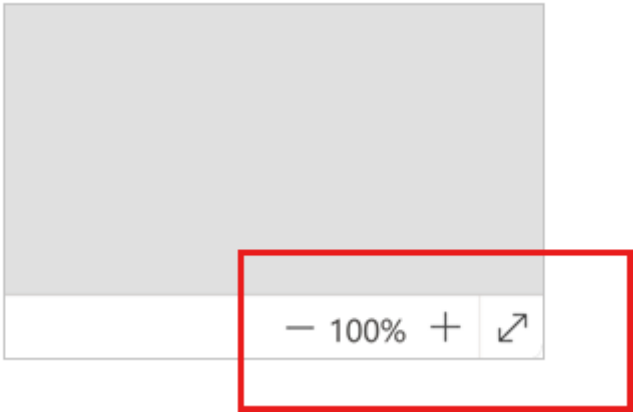

Accessibility

Lineage supports keyboard navigation and screen readers.

Tab navigation follows a depth-first search (DFS) traversal of the graph, which may differ from the visual layout. Screen readers announce node relationships to help understand lineage dependencies.

Considerations and limitations

- Lineage view isn't available in Internet Explorer. For more information, see Supported browsers for Power BI.
- Correct display of the lineage between semantic models and dataflows is guaranteed only if the **Get Data** UI is used to set up the connection to the dataflow, and the **Dataflows** connector is used. Correct display of the lineage between semantic models and dataflows isn't guaranteed if a manually created Mashup query is used to connect to the dataflow.

Related content

- Impact analysis
- Entity diagram view in KQL database

Last updated on 04/17/2026

Impact analysis

When you make changes to a Fabric item, or are considering making changes, it's important to be able to assess the potential impact of those changes on downstream items that depend on that item. **Impact analysis** provides you with information that can help you make this assessment.

- It shows you the workspaces and Fabric items that might be affected by your change. You can choose to see only direct children, or all affected downstream items and workspaces. For data sources, you can only see direct children. To see further downstream items for data sources, use impact analysis on the direct children of the data source.
- It allows you to view potentially affected items either by type or by workspace.
- It provides an easy way to notify the relevant people about a change you made or are thinking about making.
- For data sources, it shows you the connection string used to connect to the data source.

Open impact analysis for an item

You can open an item's impact analysis from its card in lineage view, or from the **Lineage** option on the items detail's page.

- From a card in lineage view:

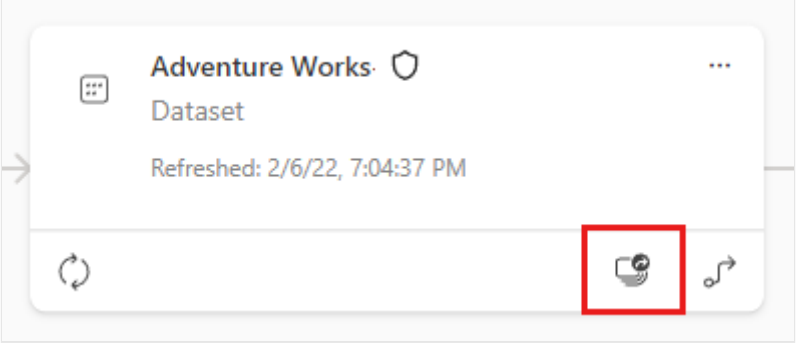

- From an item's details page

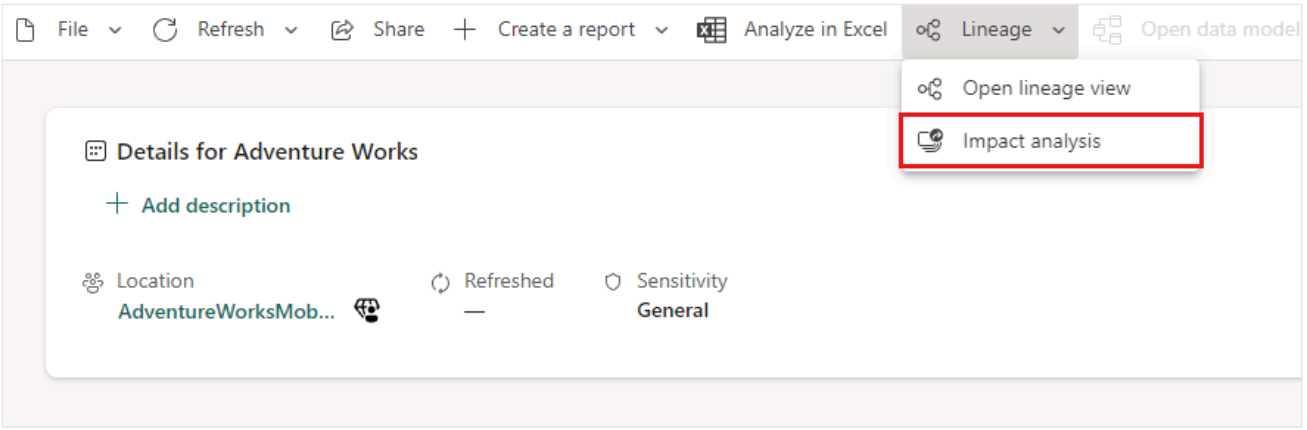

Impact analysis pane

The following images show an example of the impact analysis pane. The image on the left shows the pane when viewing impacted items by type. The image on the right shows the pane when viewing impacted items by workspace.

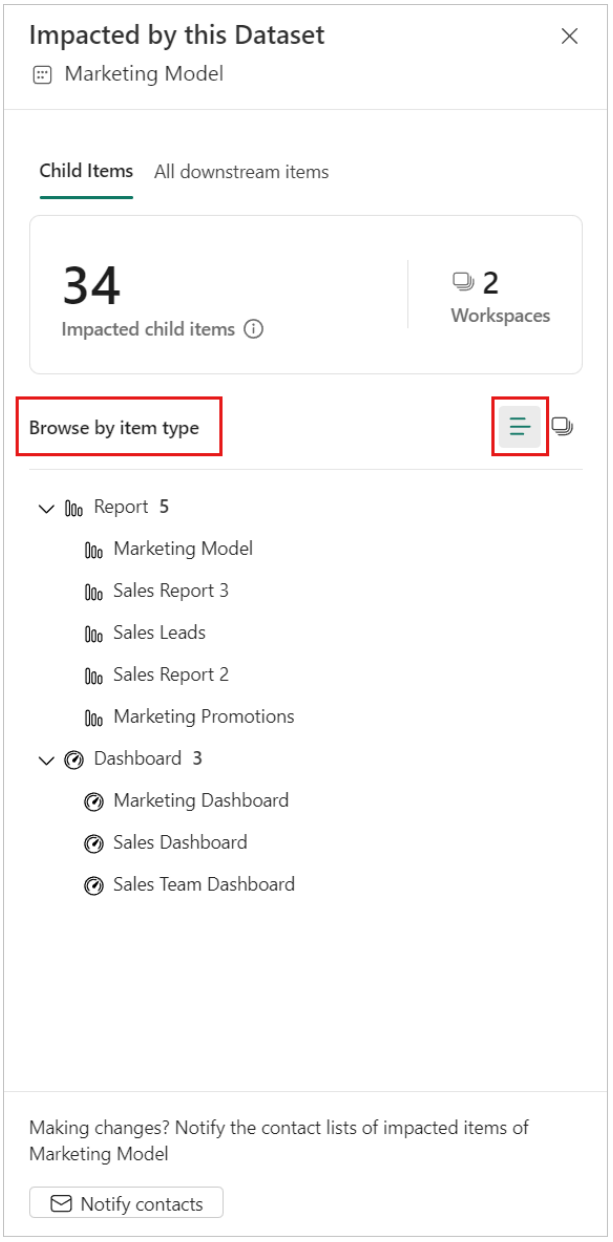

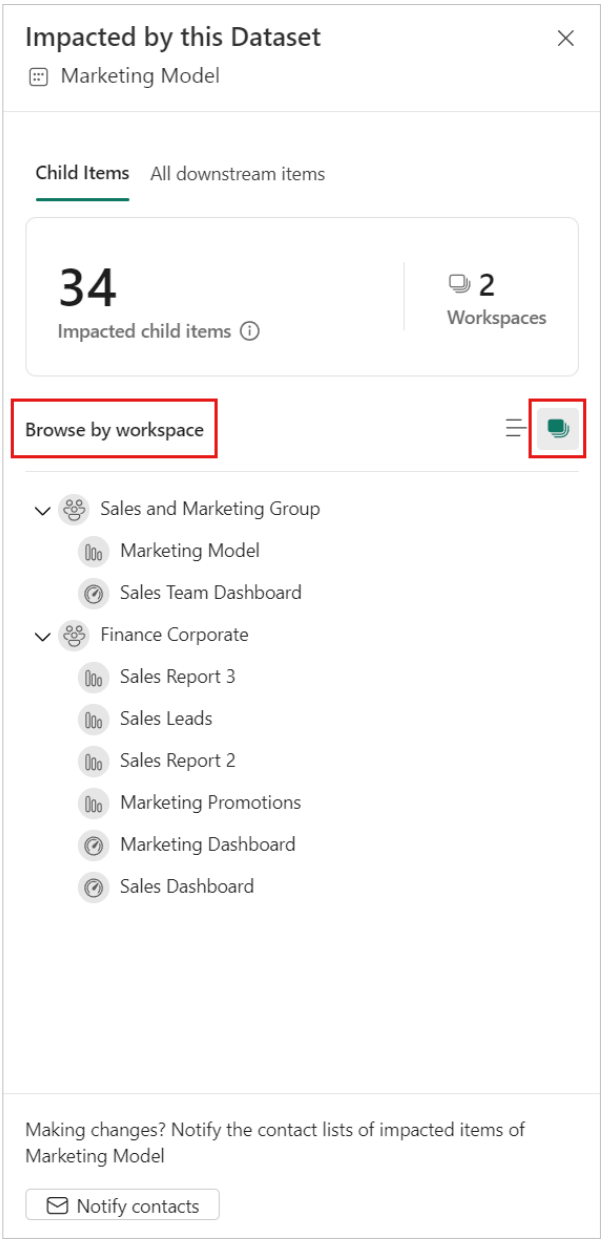

- The **Child items/All downstream items** tabs enable you to switch between seeing only direct children of the item, one-step downstream, or all the affected dependent items downstream. When you select a tab, you see the number of items and workspaces affected in each case. Data sources show direct child items only.
- The **Browse by item type/Browse by workspace** icons enable you to see affected items grouped either by type or by the workspace they're in.
- The **Notify contacts** button opens a dialog where you can create a message about any changes you make to an item, and send it to the contact lists of the affected workspaces.

The pane may say that there are more items with limited access. This occurs when there are affected items that you don't have access permissions for. Such items aren't listed by name because of privacy concerns. See Privacy for detail.

Notify contacts

If you've made a change to an item or are thinking about making a change, you might want to contact the relevant users to tell them about it. When you notify contacts, an email is sent to the contact lists of all the impacted workspaces. Your name appears on the email so the contacts can find you and reply back in a new email thread.

1. Select **Notify contacts** in the impact analysis side pane. The notify contacts dialog appears.

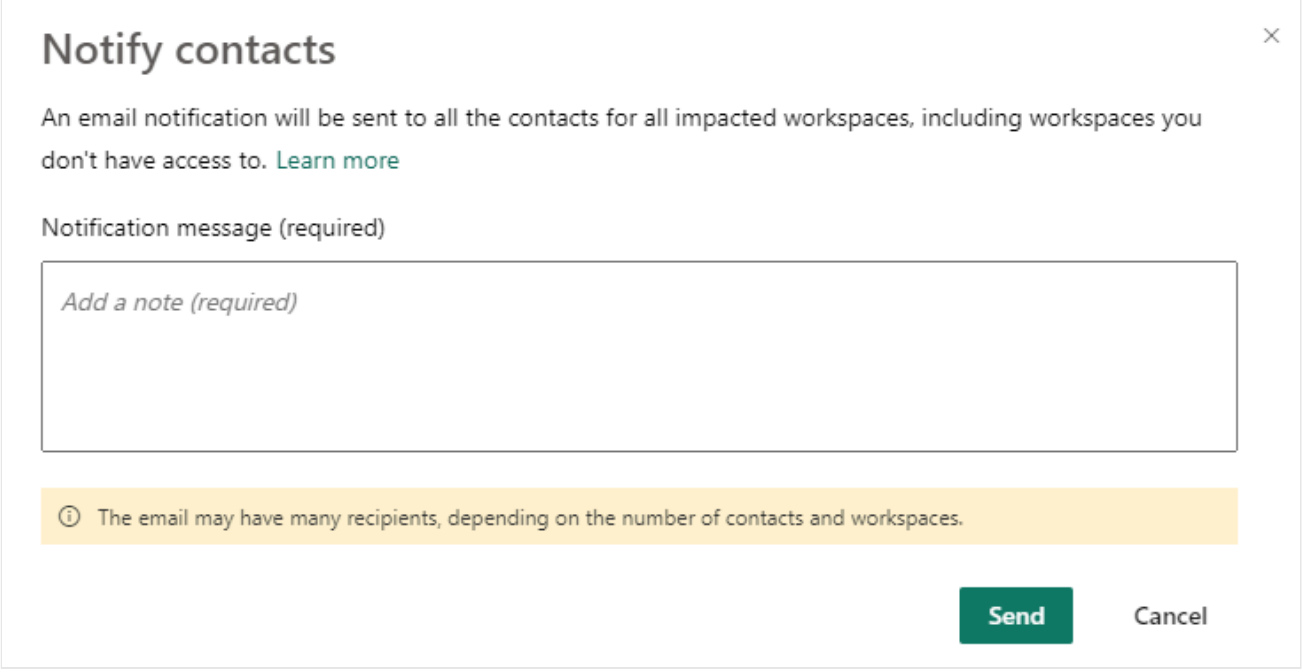

2. In the text box, provide some detail about the change.

3. When the message is ready, select **Send**.

Privacy

In order to perform impact analysis on an item, you must have write permissions to it. In the impact analysis side pane, you only see real names for workspaces and items that you have access to. Items that you don't have access to are listed as **Limited access**. This is because some item names may contain personal information.

Even if you don't have access to some workspaces, your notify contacts messages will still reach the contact lists of those workspaces.

Related content

- Lineage

Last updated on 05/23/2023

Metadata scanning overview

Metadata scanning facilitates governance of your organization's Microsoft Fabric data by making it possible to catalog and report on all the metadata of your organization's Fabric items. It accomplishes this using a set of Admin REST APIs that are collectively known as the *scanner APIs*.

With the scanner APIs, you can extract information such as item name, owner, sensitivity label, applied tags and endorsement status. For Power BI semantic models, you can also extract the metadata of some of the objects they contain, such as table and column names, measures, DAX expressions, mashup queries, and so forth. The metadata of these semantic model internal objects is referred to as subartifact metadata.

For a more extensive list of the artifact and subartifact metadata that metadata scanning returns, see the documentation for the Admin - WorkspaceInfo GetScanResult API.

The following are the scanner APIs. They support both public and sovereign clouds.

- GetModifiedWorkspaces
- PostWorkspaceInfo
- GetScanStatus
- GetScanResult

> **Important**
>
> The app you develop for scanning can authenticate by using either a standard delegated admin access token or a service principal. The two authentication paths are mutually exclusive. **When running under a service principal, there must be no Power BI admin-consent-required permissions set on your app**. For more information, see Enable service principal authentication for admin APIs.

Basic flow

- **Set up metadata scanning in the organization**: Before metadata scanning can be run, a Fabric admin needs to set it up in your organization. Fabric admins should see Set up metadata scanning.
- **Enable service principal authentication for admin read-only APIs**. Service principal is an authentication method that can be used to let a Microsoft Entra application access

Microsoft Fabric content and APIs. See Enable service principal authentication for admin APIs.

- **Run metadata scanning**: See Run metadata scanning for a walk-through that demonstrates how to run a scan.

Licensing

Metadata scanning requires no special license. It works for all of your tenant's metadata, including that of items located in non-Premium workspaces.

Related content

- Set up metadata scanning
- Enable service principal authentication for admin APIs
- Run metadata scanning
- Learn about Power BI REST Admin APIs

Last updated on 06/11/2025

Set up metadata scanning in your organization

Before metadata scanning can be run over an organization's Fabric workspaces, it must be set up by a Fabric administrator. Setting up metadata scanning involves two steps:

1. Enabling service principal authentication for read-only admin APIs.
2. Enabling tenant settings for detailed semantic model metadata scanning.

Enable service principal authentication for read-only admin APIs

Service principal is an authentication method that can be used to let a Microsoft Entra application access Power BI APIs. With this authentication method, you don't have to maintain a service account with an admin role. Rather, to allow your app to use the Admin APIs, you just have to give your approval once as part of the tenant settings configuration.

To see how to enable service principal access to read-only Admin APIs, see Enable service principal authentication for read-only admin APIs.

If you don't want to enable service principal authentication, metadata scanning can be performed with standard delegated admin access token authentication.

Enable tenant settings for metadata scanning

Two tenant settings control metadata scanning:

- **Enhance admin APIs responses with detailed metadata**: This setting turns on Model caching and enhances API responses with low-level semantic model metadata (for example, name and description) for tables, columns, and measures.
- **Enhance admin APIs responses with DAX and mashup expressions**: This setting allows the API response to include DAX expressions and Mashup queries. This setting can only be enabled if the first setting is also enabled.

To enable these settings, go to **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings** > **Admin API settings**.

Related content

- Metadata scanning overview
- Enable service principal authentication for admin APIs
- Run metadata scanning
- Power BI REST Admin APIs

Last updated on 09/29/2026

Run metadata scanning

The following short walkthrough shows how to use the scanner APIs to retrieve metadata from your organization's Fabric items. It assumes that a Fabric admin has set up metadata scanning in your organization.

For the list of the artifact and subartifact metadata that metadata scanning returns, see the documentation for the Admin - WorkspaceInfo GetScanResult API.

The following are the scanner APIs. They support both public and sovereign clouds.

- GetModifiedWorkspaces
- WorkspaceGetInfo
- WorkspaceScanStatus
- WorkspaceScanResult

> **Important**
>
> The app you develop for scanning can authenticate by using either a standard delegated admin access token or a service principal. The two authentication paths are mutually exclusive. **When running under a service principal, there must be no Power BI admin-consent-required permissions set on your app**. For more information, see **Enable service principal authentication for admin APIs**.

Step 1: Perform a full scan

Call workspaces/modified without the **modifiedSince** parameter to get the complete list of workspace IDs in the tenant. This scan retrieves all the workspaces in the tenant, including personal workspaces and shared workspaces. If you wish to exclude personal workspaces from the scan, use the *workspaces/modified* **excludePersonalWorkspaces** parameter.

Divide the list into chunks of 100 workspaces at most.

For each chunk of 100 workspaces:

Call workspaces/getInfo to trigger a scan call for these 100 workspaces. You'll receive the scanId in the response to use in the next steps. In the location header, you'll also receive the Uniform Resource Identifier (URI) to call for the next step.

> **Note**
>
> Not more than 16 calls can be made simultaneously. The caller should wait for a scan succeed/failed response from the *scanStatus* API before invoking another call.
>
> If some metadata you expected to receive is not returned, check with your Fabric admin to make sure they have enabled all relevant admin switches.

Use the URI from the location header you received from calling *workspaces/getInfo* and poll on workspaces/scanStatus/{scan_id} until the status returned is "Succeeded". This status means the scan result is ready. It's recommended to use a polling interval of 30-60 seconds. In the location header, you also receive the URI to call in the next step. Use it only after the status is "Succeeded".

Use the URI from the location header you received from calling *workspaces/scanStatus/{scan-id}* and read the data using workspaces/scanResult/{scan_id}. The data contains the list of workspaces, item info, and other metadata based on the parameters passed in the *workspaces/getInfo* call.

Step 2: Perform an incremental scan

Now that you have all the workspaces and the metadata and lineage of their assets, it's recommended that you perform only incremental scans that reference the previous scan that you did.

Call workspaces/modified with the **modifiedSince** parameter set to the start time of the last scan in order to get the workspaces that have changed, and which therefore require another scan. The **modifiedSince** parameter should be set for a date within the last 30 days.

Divide this list into chunks of up to 100 workspaces, and get the data for these changed workspaces by using the three API calls, workspaces/getInfo, workspaces/scanStatus/{scan_id}, and workspaces/scanResult/{scan_id}, as described in Step 1.

Considerations and limitations

- semantic models that haven't been refreshed or republished will be returned in API responses but without their subartifact information and expressions. For example, semantic model name and lineage are included in the response, but not the semantic model's table and column names.

- semantic models containing only **DirectQuery** tables will return subartifact metadata only if some sort of action has been taken on the semantic model, such as someone building a report on top of it, someone viewing a report based on it, etc.
- Real-time datasets, semantic models with object-level security, semantic models with a live connection to *AS-Azure* and *AS on-premises*, and Excel full fidelity datasets aren't supported for subartifact metadata. For unsupported datasets, the response returns the reason for not getting the subartifact metadata from the dataset. It's found in a field named **schemaRetrievalError**, for example, **schemaRetrievalError: Unsupported request. RealTime dataset are not supported**.
- The API doesn't return subartifact metadata for semantic models that are larger than 1 GB in shared workspaces. In Premium workspaces, there's no size limitation on semantic models.

Licensing

Metadata scanning requires no special license. It works for all of your tenant's metadata, including that of items located in non-Premium workspaces.

Related content

- Metadata scanning overview
- Set up metadata scanning.
- Enable service principal authentication for admin APIs.
- Learn about Power BI REST Admin APIs.
- More questions? Ask the Power BI Community.

Last updated on 05/01/2024

Troubleshoot access restrictions

The Microsoft Fabric permission model allows users and workspace admins to grant permissions to other users at the workspace level. These permissions define the default access for all items in the workspace. Users can also be given elevated permissions for specific items at the item level, even if they're only viewers at the workspace level. Learn more about the Fabric permission model.

Fabric provides other access controls. For example, OneLake security enables admins to assign roles to items and control which portions of data users can view. More importantly for this troubleshooting article, organization-wide policies can automatically restrict access to items based on their sensitivity or content.

Because there are multiple layers controlling access, you could encounter restrictions from policies you're unaware of. This article helps you identify if a policy is affecting your access and how to restore it if possible.

User's access to an item is blocked

If you previously had access to an item but can no longer see it in your workspace or open it, consider the following common causes.

Access was manually removed

A user with administrative or ownership permissions manually revoked your access to the item. This action is often taken to comply with company regulations or security policies. To resolve this issue:

- **Contact the item owner directly**: If you know who owns or created the item, reach out to them to inquire about the access change.
- **Reach out to your workspace administrator**: If you don't know the item owner, your workspace admin can help you contact them or reinstate your access if they removed it by mistake.

Workspace permissions changed

You were removed from the workspace or your access level was downgraded, which automatically revoked some of your permissions to the items within it.

To understand the reason for the change in permissions, contact your workspace administrator.

Company policies changed

Your organization implemented company-wide policies that restrict access to certain items based on the applied their sensitivity or content. These policies can override individual permissions and workspace settings. There are two main types of policies that can enforce such restrictions:

- **Protection policies**: Microsoft Purview protection policies automatically enforce access restrictions based on the applied *sensitivity label*. For example, the label "Confidential" can automatically limit access to a specific security group, for any Fabric item labeled by it.
- **Data Loss Prevention (DLP) policies with the "restrict access" action**: Microsoft Purview DLP policies with the *restrict access* action allow for access enforcement on items containing sensitive information.

Both policies produce similar user experiences when access is revoked. The following sections provide troubleshooting steps for these scenarios.

How to tell if policies are the cause of access issues

You can tell a policy change is the cause of your restricted access if you experience one of the following issues:

- You no longer see the item in your workspace list view or the OneLake catalog.

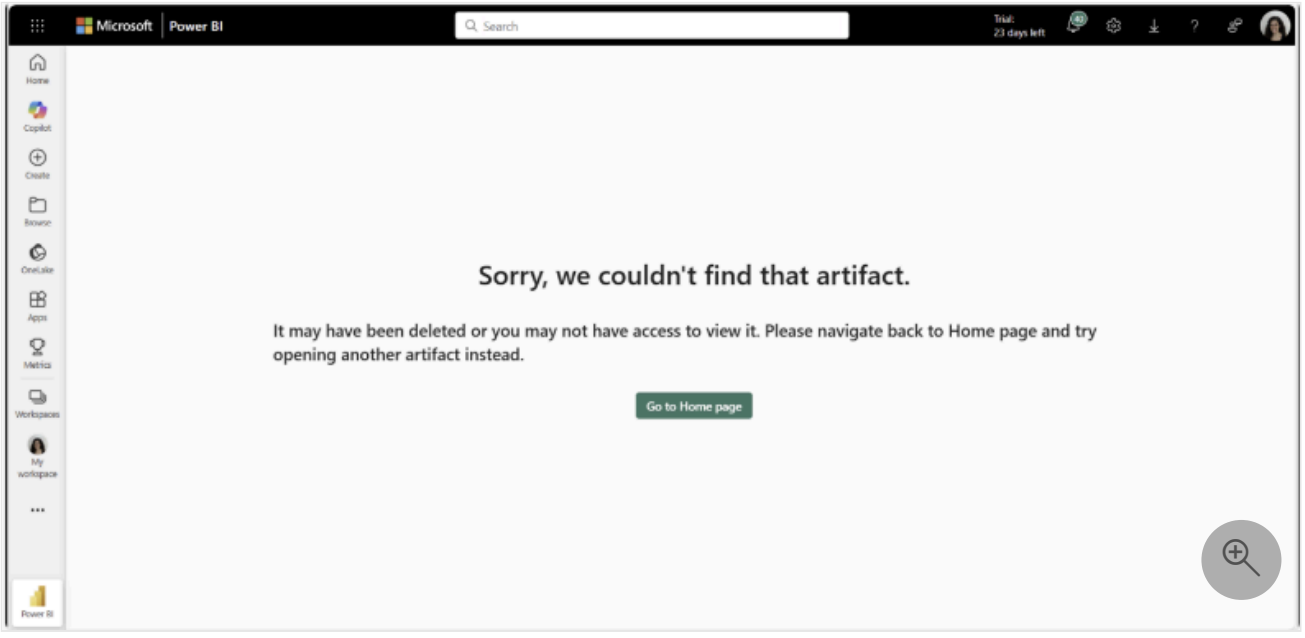

- You see "Permission required" or "Sorry we couldn’t locate that artifact" when trying to open the item using a link.

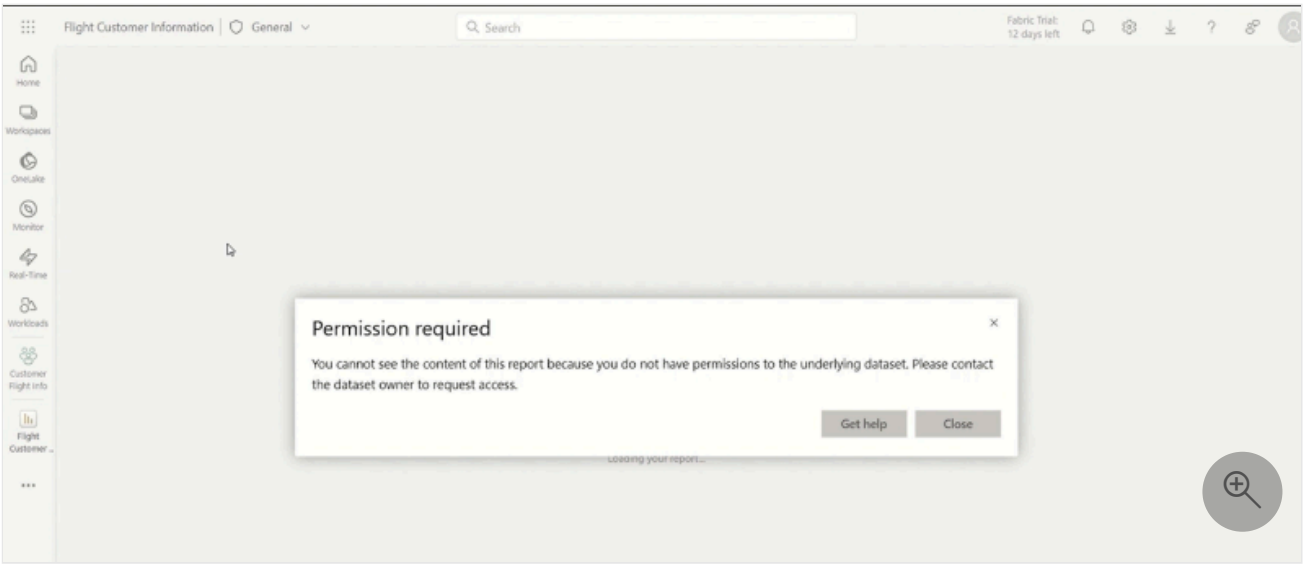

- You see an indicator on the item showing that it contains data from a restricted underlying source.

- You see tiles on a dashboard that display a "Permission required" message. This message indicates that these tiles contain data from a restricted underlying source. Tiles that don't use restricted data render normally.

These symptoms can indicate that either a protection policy or a DLP policy with the "restrict access" action is blocking your access.

What should I do if I think a policy is blocking my access?

Contact your workspace admin to investigate and resolve the issue. If the workspace admin is also blocked by a policy, the label issuer (the person who applied the sensitivity label) or the item creator can access the **Manage Permissions** page to determine if a policy is blocking your access.

The next sections provide guidance for workspace admins and others with appropriate permissions on how to troubleshoot and resolve these issues.

How to investigate a restriction scenario

When a user reports losing access to an item, the workspace admin can investigate the cause along with other users who can assist depending on their permissions:

- **Workspace admins** can access the **Manage Permissions** page to investigate access issues for all users.

- **Label issuers** (users who applied the sensitivity label to the item) retain access even if not included in the protection policy, preventing lockout scenarios.
- **Item creators** also retain access when labels are applied automatically (such as through downstream inheritance), ensuring at least one user can manage the item.

To investigate the issue:

1. Open the item's **Manage Permissions** page and check the user's **Permissions** setting.
2. If the user's setting is **User is blocked**, it indicates that the user's access was revoked by an organizational security policy.

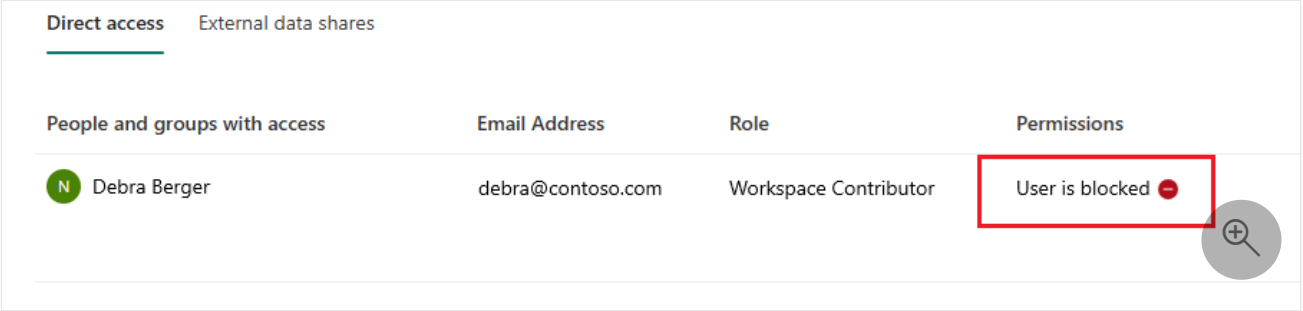

3. Identify the policy type:
    - Recent label changes likely mean an issue with a Microsoft Purview protection policy.
    - Recent item updates or sensitive data likely mean an issue with a Microsoft Purview DLP restriction.
4. If the user's **Permissions** setting doesn't show **User is blocked**, investigate other causes such as manual access removal or workspace permission changes.

How to resolve policy-related issues

To unblock users, the workspace admin or a user with appropriate permissions, such as the label issuer or item creator, can take action. The steps to resolve the issue vary based on whether a Microsoft Purview protection policy or Microsoft Purview DLP policy is causing the restriction.

> ⓘ **Note**
>
> Only take these actions after ensuring they don't risk organizational compliance or data security. These actions are audited and might be subject to investigation by your security administrator.

Resolving protection policy issues

When a protection policy revoked access and the user's permission shows as **No access**, the workspace admin can take the following actions. The label issuer (the person who applied the label to the item) or the item creator can also assist with investigating and resolving the issue.

- If the sensitivity label was applied incorrectly, either apply the correct label yourself or request that the Protection Owner update it.
- If the label is correct but the user needs access to this classified data, contact your organization's security team to request adding the user to the protection policy in Microsoft Purview. After the policy is updated, have the user sign back in to Fabric to receive the updated permissions.

Resolving DLP "restrict access" policy issues

If a DLP policy caused the restriction, the workspace admin can take the following actions:

- If the DLP policy incorrectly flagged the data as sensitive and you're confident it's a false positive, the workspace admin can override the policy. Be prepared to provide justification to your compliance team for this action.
- If the data was mistakenly added to the data item, the workspace admin can edit the item to remove the sensitive data. This action triggers a new evaluation that automatically removes the restriction.

Understanding how DLP policies affect different user types

If a DLP policy with the "restrict access" action is causing the restriction, it's important to understand how it affects different user types:

- **All users in the tenant**: When applied to this group, DLP policies block all users except workspace admins, who can troubleshoot and fix the data, where relevant.
- **Guest users**: When applied to this group, DLP policies block guest users, while retaining access for internal users and workspace admins. Guest users are users who aren't defined as **Member** in Microsoft Entra. If an external user is configured as a tenant **Member**, the policy treats them as an internal user.

The workspace admin can assess how DLP policies affect different user types to decide how to handle the restriction.

Mirrored databases restricted by DLP – special case

Mirrored databases present a slightly different use case. Usually, these types of databases are updated by service principals (SPNs) rather than by human users. Because organizations typically wouldn't grant these SPNs full control of workspaces (the workspace admin role), the updating user might be restricted by the DLP **restrict access** action.

How to resolve a database-restricted user

Workspace admins can always override the policy if the circumstances merit an override.

If sensitive information was justifiably found in the data, workspace admins can:

1. Edit the mirroring configuration to exclude the problematic table.
2. Edit the existing item to remove the sensitive information.

Once these two steps are completed, the next database update triggers an automated DLP evaluation. Because the sensitive data has been removed, the evaluation finds no sensitive information, and the restriction is lifted.

Related articles

- Microsoft Fabric permission model
- Microsoft Purview protection policies
- Data Loss Prevention (DLP) in Microsoft Purview

Last updated on 05/11/2026

Find your Fabric home region

A Fabric home region is the Azure datacenter region linked to your Fabric tenant, chosen when your tenant was created. Knowing your home region is important for understanding workload and feature availability, as well as data residency, performance, and compliance.

To find your Fabric home region, follow these steps:

1. Sign in to Fabric.
2. Open the Help pane and choose the **About** link (the text of the link varies depending on which workload is selected).

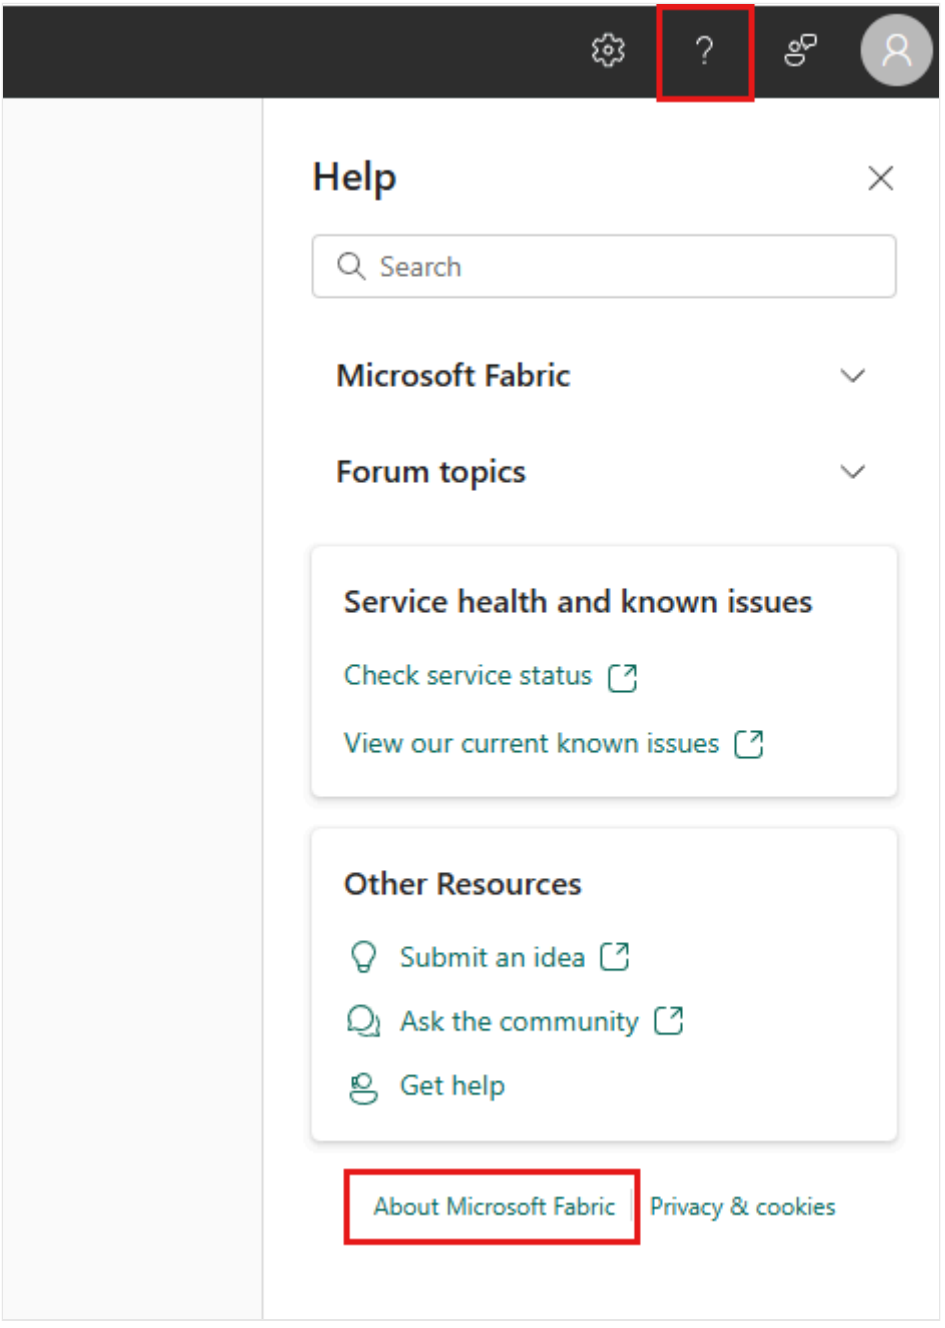

3. Look for the value next to **Your data is stored in**. The location shown is the default region where your data is stored. You might also be using capacities in different regions for your workspaces.

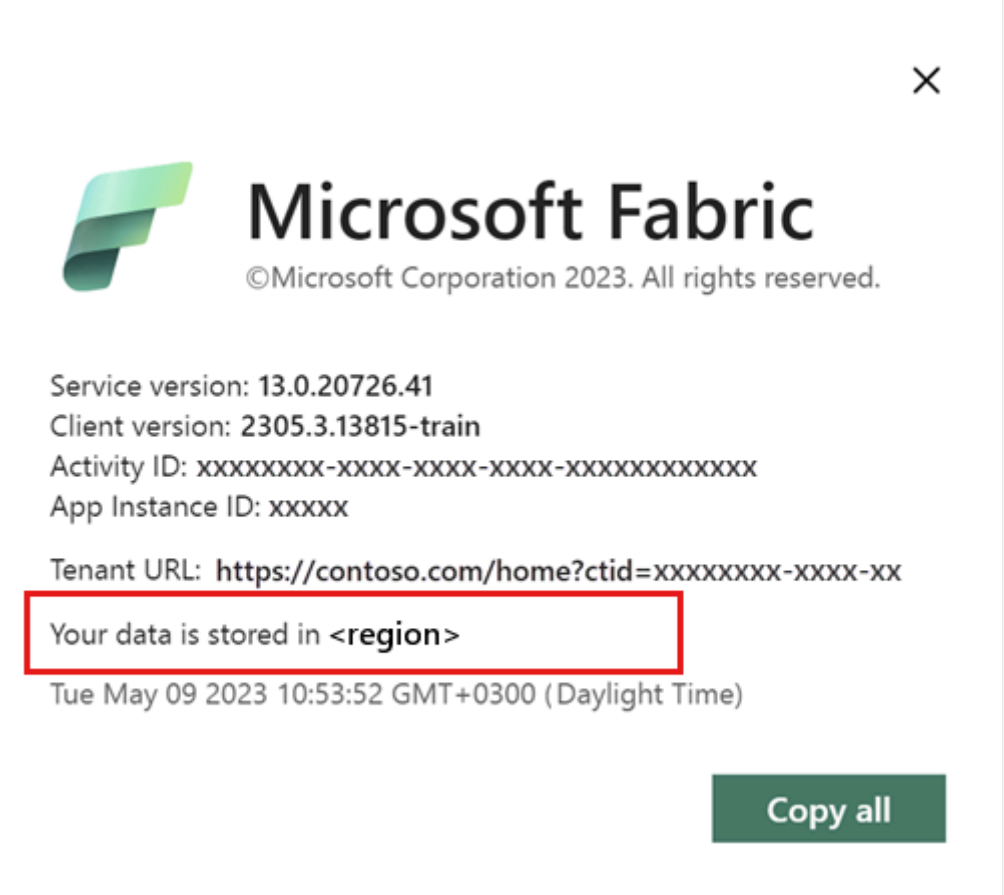

Related content

- Buy Fabric capacity in Azure
- Region availability
- Azure public cloud geographies and regions

Last updated on 08/05/2026

Manage your capacities in the OneLake catalog (preview)

A Fabric capacity is a dedicated pool of compute resources that powers the items and workloads across your organization. As a Fabric administrator or capacity administrator, you need a single place to see the capacities you're responsible for and keep them sized and configured the way your organization requires.

This article shows you how to manage your capacities from the **Capacities** section in the Govern section of the OneLake catalog. From this centralized location, you can view the capacities you administer, create and delete capacities, change capacity settings, and configure delegated tenant settings.

> **Note**
>
> OneLake catalog and Govern are rolling out by region. If they're not available in your region, select **Settings** (gear) icon > **Admin portal** > **Capacity settings** during the regional rollout.

Prerequisites

To access the **Manage capacities** page, you must be a Fabric administrator or a capacity administrator. Fabric administrators can access the page even if they aren't an administrator of any capacity. Capacity administrators can access the capacities they administer.

Fabric roles and Azure role-based access control (Azure RBAC) permissions are independent. To perform an operation on an Azure-managed capacity, you also need the applicable Azure RBAC permissions on the Azure subscription or capacity resource. For available resource provider operations, see the Azure permissions for Microsoft Fabric and Power BI Embedded.

View all capacities

To get to your capacities in the OneLake catalog, follow these steps:

1. Sign in to Fabric by using your admin account credentials.
2. Select **OneLake catalog**, and then select the **Govern** tab.
3. Select **Capacities**.

The **Manage capacities** page shows a list of all the capacities you have permission to manage in your tenant.

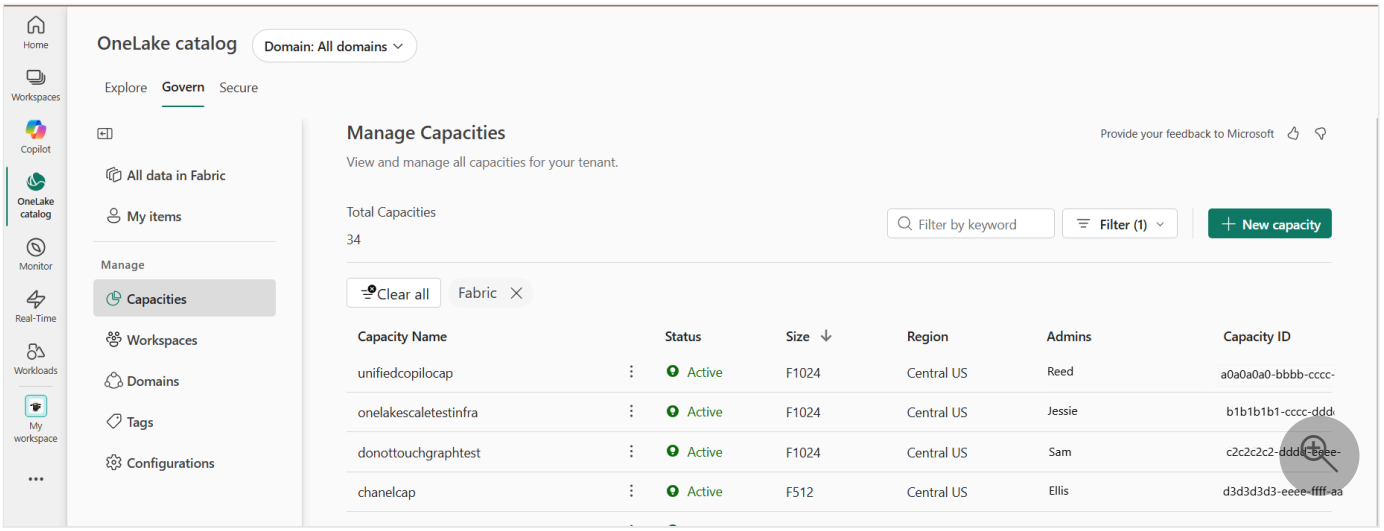

From this page, you can search for a capacity by name, filter the list by capacity type, size, region, and subscription ID, sort by any available column, and open the pane to create a new capacity.

The list includes all capacity types:

- **Power BI Premium** - A capacity you buy as part of a Power BI Premium subscription. These capacities use P SKUs.

  > **Note**
  >
  > Power BI Premium per-capacity (P SKU) subscriptions are retiring. To keep your Power BI workloads running, migrate to Fabric capacity (F SKUs). For an end-to-end view of the migration, see Power BI Premium to Microsoft Fabric migration overview. For answers to common questions, see the Power BI Premium to Microsoft Fabric migration FAQ.

- **Power BI Embedded** - A capacity you buy as part of a Power BI Embedded subscription. These capacities use A or EM SKUs.
- **Trial** - A Fabric trial capacity. These capacities use Trial SKUs.
- **Fabric capacity** - A Fabric capacity. These capacities use F SKUs.

View details for a capacity

To view details for a capacity, follow these steps:

1. Go to **OneLake catalog** > **Govern** > **Capacities**.

2. On the **Manage capacities** page, filter or search the list to find the capacity.

3. Select the capacity to open the details page.

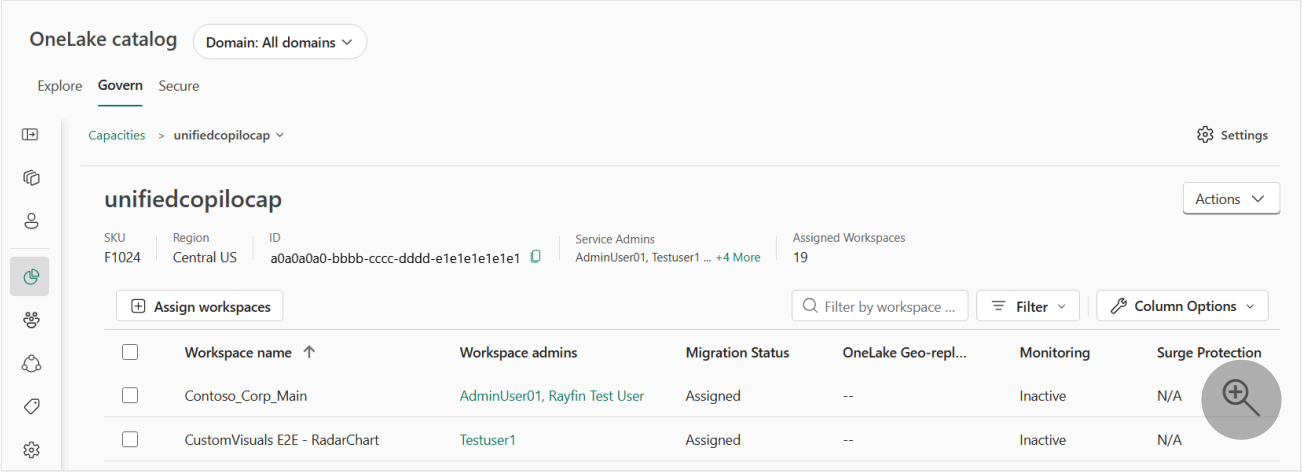

The capacity details view includes:

- Basic capacity information, such as name, SKU, Region, and ID.
- The service admin assigned to the capacity.
- The list of workspaces assigned to the capacity.
- Workspace details, including workspace admins and migration, OneLake geo-replication, monitoring, and surge protection status.

Create a new capacity

The **Manage capacities** page contains a **New capacity** option, where you can create a new Fabric capacity or Azure Embedded capacity. You must be a Fabric administrator. To create an Azure-managed capacity, you also need Azure RBAC permission to create capacity resources in the selected subscription. For F SKU requirements, see Buy Fabric capacity in Azure. For a Power BI Embedded capacity, the applicable permissions include `Microsoft.PowerBIDedicated/capacities/read` and `Microsoft.PowerBIDedicated/capacities/write`.

To create a new capacity, follow these steps:

1. In Fabric, go to **OneLake catalog** > **Govern** > **Capacities**.

2. On the **Manage capacities** page, select **New capacity**.

3. On the **New capacity** page, enter the following information:

    - **Capacity Type** - Select the type of capacity.

- **Capacity name** - Give your capacity a name.
- **Capacity admins** - Add capacity admins.
- **Subscription** - Select the Azure subscription you want to use for the capacity.
- **Resource group** - Select the Azure resource group you want to use for the capacity.
- **Region** - Select the region where you want to create the capacity.
- **Capacity size** - Select the size of the capacity.
- **Capacity admins** - Select the capacity admins.

  > ⓘ **Note**
  > - You can create a new Power BI Embedded capacity with an A SKU or an EM SKU. If you select an EM size, you create a Power BI Embedded capacity.
  > - To create a new Trial capacity, see Fabric trial.

4. To add tags, expand **Additional settings**, select **Add tag**, and then enter the tag name and value.

5. Select **Apply**.

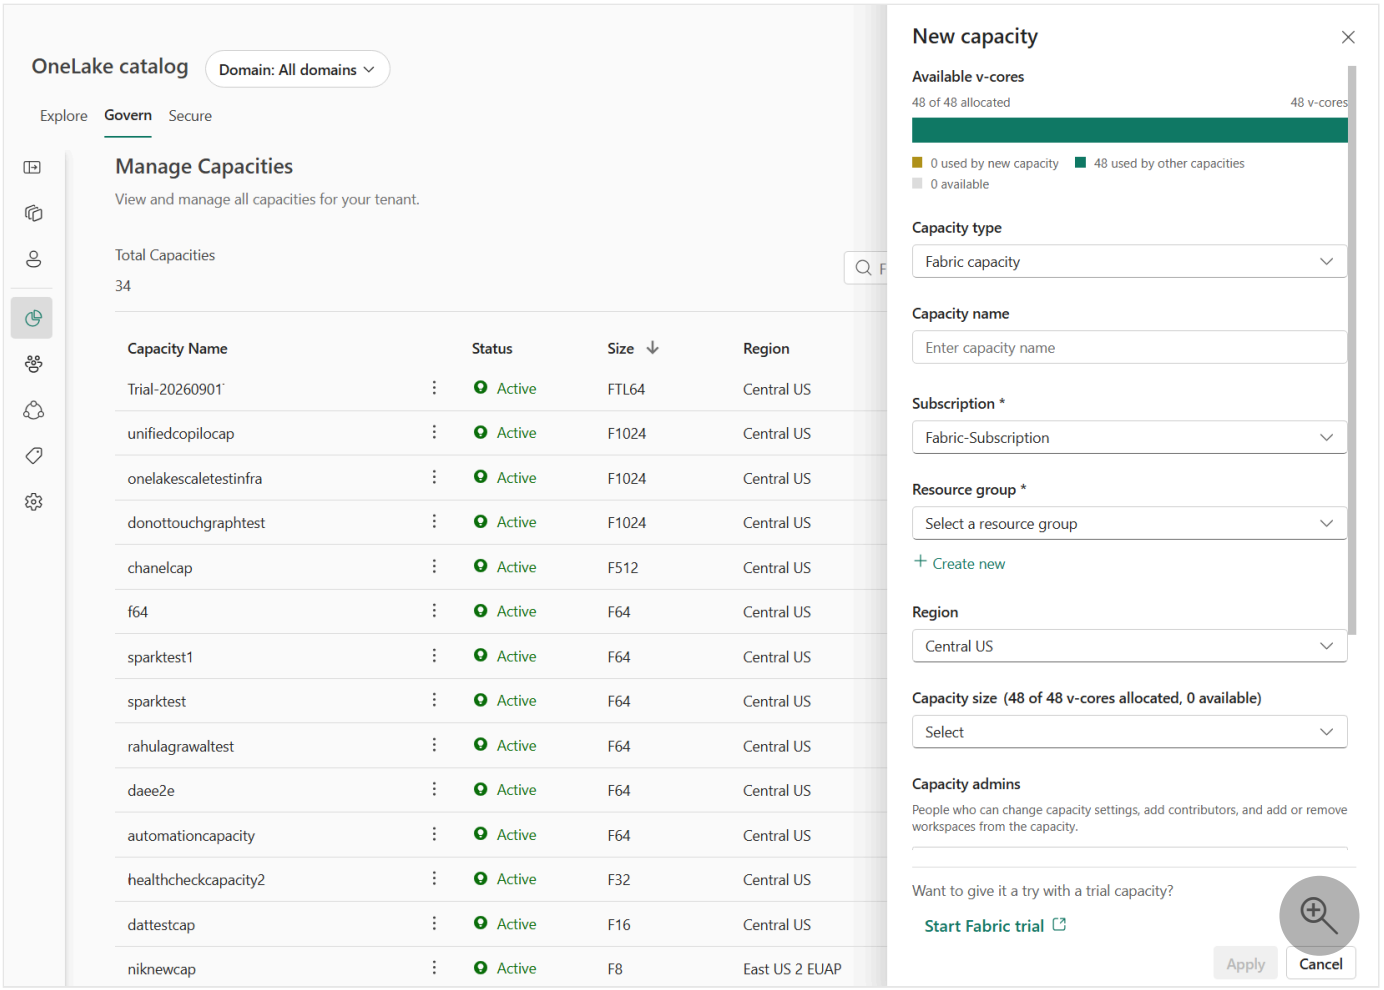

Manage capacities with the actions menu

From either the **Manage capacities** page or a capacity's details page, use the actions menu to perform basic capacity management tasks. For example, pause or resize a capacity, reassign workspaces, configure surge protection, or set capacity overage limits. The actions available depend on the capacity type and your permissions.

To select an action, follow these steps:

1. Go to **OneLake catalog** > **Govern** > **Capacities**.
2. On the **Manage capacities** page, filter or search the list to find the capacity.
3. Do one of the following:
    - Next to the capacity name, select **More options** (three dots), and then select the action you want to take.

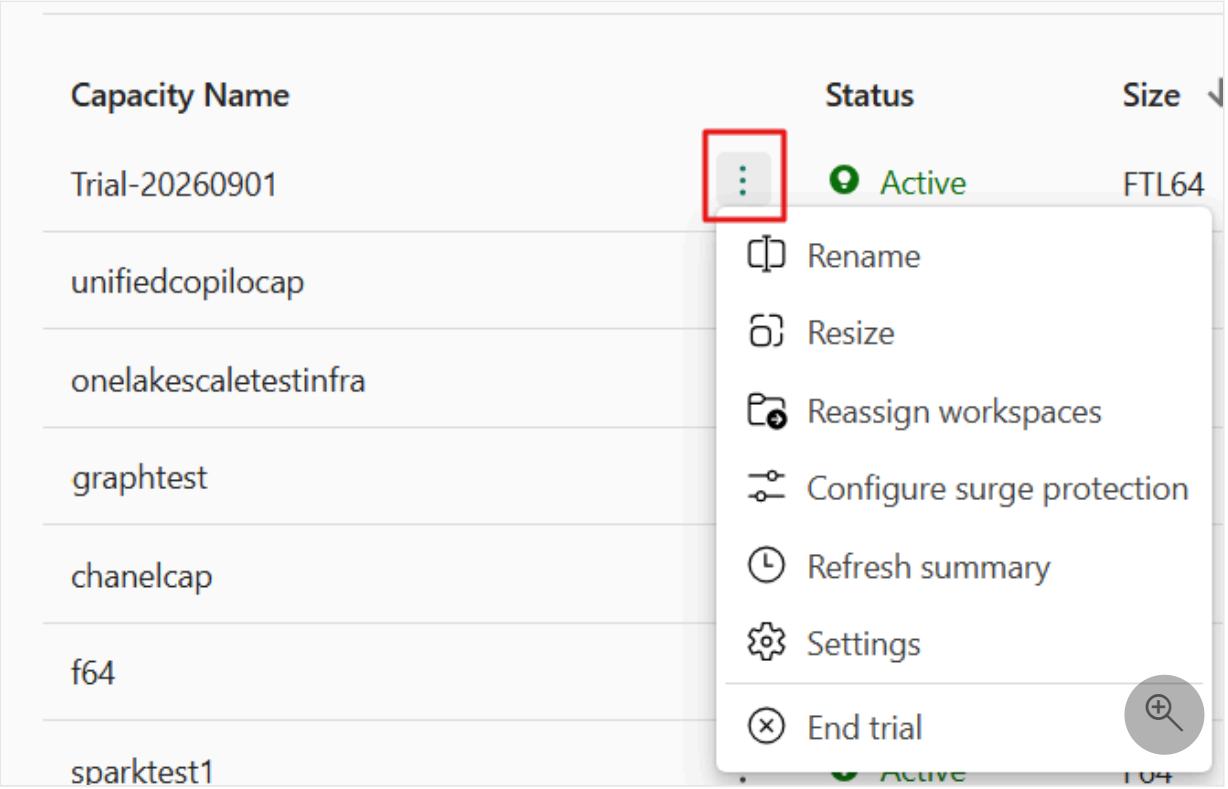

- Open the details page by selecting the capacity name. Use the **Actions** menu to take the action you want.

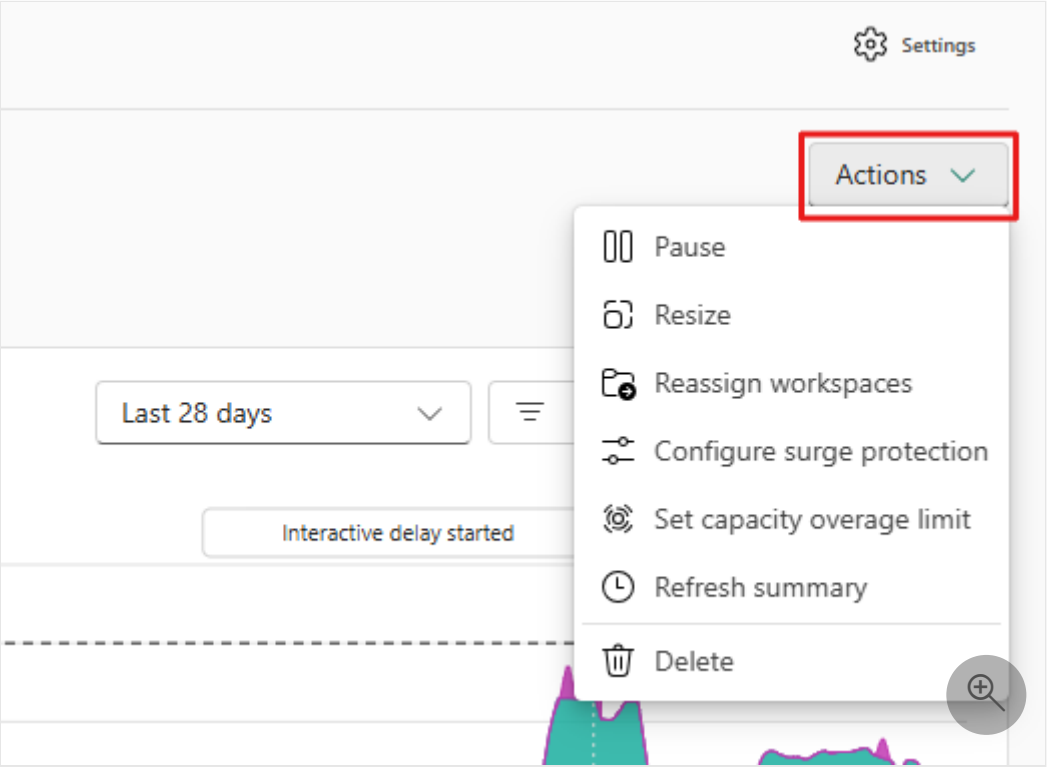

Rename your capacity

To change the name of your Power BI Premium, Power BI Embedded EM SKU, or Trial capacity, follow these steps. You must be a capacity administrator. To rename a Power BI Embedded EM SKU, you also need the `Microsoft.PowerBIDedicated/capacities/read` and

`Microsoft.PowerBIDedicated/capacities/write` Azure RBAC permissions on the capacity resource. For information about assigning capacity administrators, see Add and remove admins.

> **Note**
>
> You can't rename:
>
> - A SKUs
> - Fabric capacities (F SKUs)

1. Go to **OneLake catalog** > **Govern** > **Capacities**.
2. On the **Manage capacities** page, filter or search the list to find the capacity.
3. Next to the capacity name, select **More options** (three dots) > **Rename**.
4. Type the new name in the **Capacity name** field.
5. Select (**Apply**).

Pause and resume a capacity

You can pause a Fabric capacity to stop billing when the capacity isn't in use, and resume it when you need it again. For F SKU prerequisites and steps, see Pause and resume your capacity. For an A SKU, you need the `Microsoft.PowerBIDedicated/capacities/read`, `Microsoft.PowerBIDedicated/capacities/write`, `Microsoft.PowerBIDedicated/capacities/suspend/action`, and `Microsoft.PowerBIDedicated/capacities/resume/action` Azure RBAC permissions on the capacity resource.

> **Note**
>
> You can only pause and resume A SKU and F SKU capacities.

Resize a capacity

The size of a capacity determines the amount of resources allocated to it, such as memory and processing power. To resize a capacity, you must be a capacity administrator. To resize an Azure-managed capacity, you also need Azure RBAC permission to read and update the capacity resource. For F SKU requirements, see Scale your Fabric capacity. For a Power BI Embedded capacity, the applicable permissions include `Microsoft.PowerBIDedicated/capacities/read` and

`Microsoft.PowerBIDedicated/capacities/write`. For information about assigning capacity administrators, see Add and remove admins.

> **Note**
>
> - You can't resize a trial capacity.

To resize a capacity, follow these steps:

1. Go to **OneLake catalog** > **Govern** > **Capacities**.
2. On the **Manage capacities** page, filter or search the list to find the capacity.
3. Next to the capacity name, select **More options** (three dots) > **Resize**.
4. In the **Resize capacity** pane, specify the new **Target size** for the capacity.
5. Select **Apply** to apply the changes.

Autoscale a capacity

You can enable autoscale for a Power BI Premium or Power BI Embedded capacity to automatically adjust its size based on workload demand. For steps, see Using Autoscale with Power BI Premium.

> **Note**
>
> Autoscale isn't available for Trial and Fabric F SKU capacities.

To enable autoscale for a capacity, follow these steps:

1. Go to **OneLake catalog** > **Govern** > **Capacities**.
2. On the **Manage capacities** page, filter or search the list to find the capacity.
3. Next to the capacity name, select **More options** (three dots) > **Manage Autoscale**.
4. In the **Autoscale settings** pane, toggle the switch to **Enable Autoscale**.
5. Select a **Subscription** and **Resource group** for billing.
6. Select **Save**.

Reassign a workspace to a different capacity

You can move a workspace from one capacity to another from the OneLake catalog **Govern** section. For steps, see Reassign a workspace to a different capacity. For information about limitations you might encounter, see Capacity reassignment restrictions and common issues.

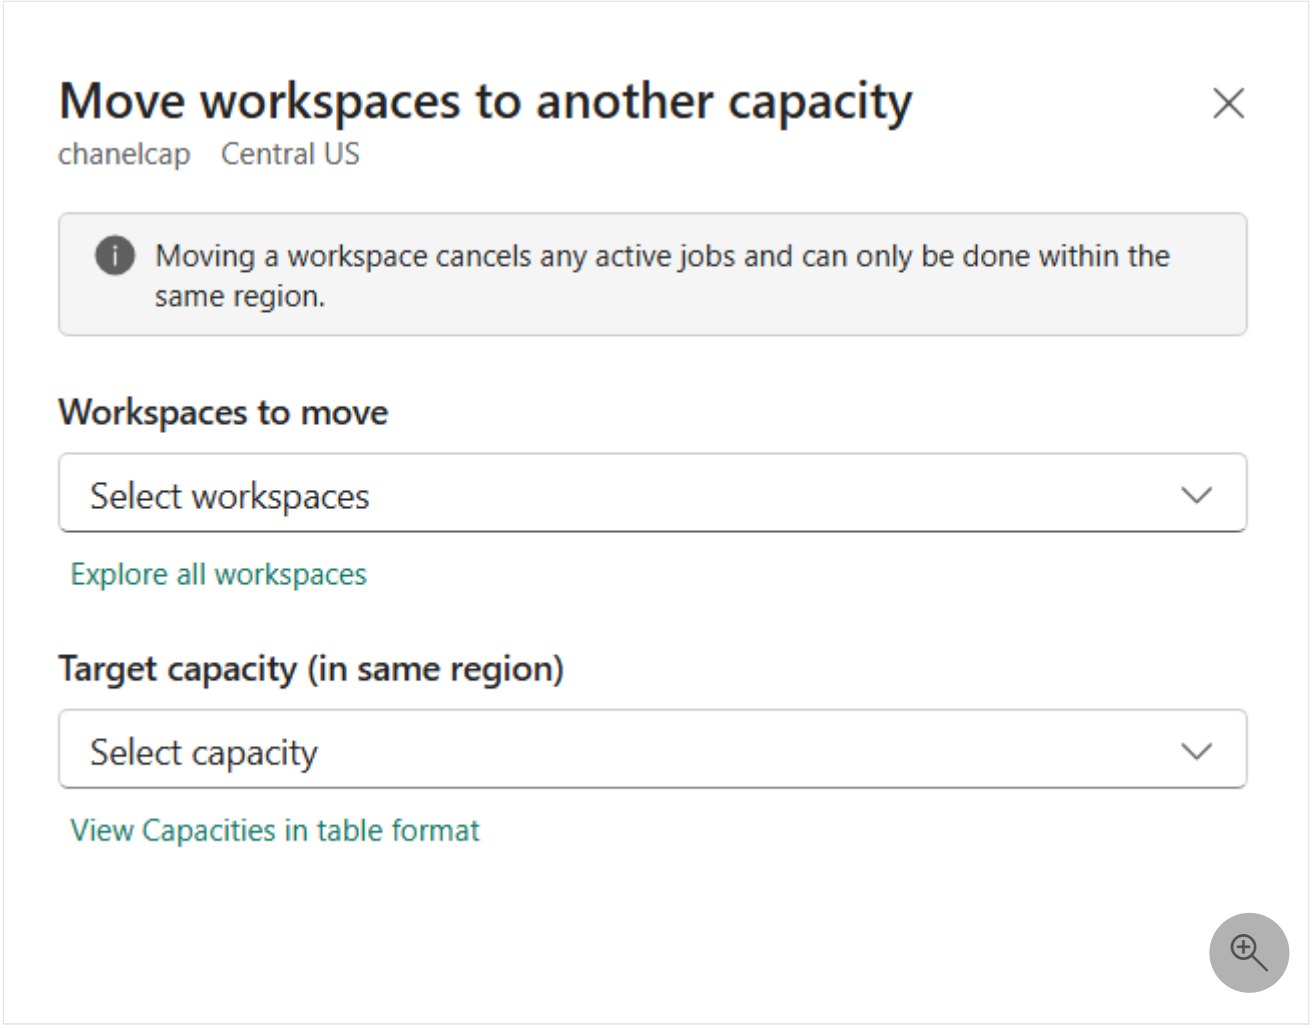

Configure surge protection

Surge protection helps protect your capacities from unexpected spikes in background workload demand. For steps, see Surge protection.

Set a capacity overage limit

Set a capacity overage limit to control how much overage a capacity can accrue. For steps, see Enable capacity overage in Microsoft Fabric.

Delete a capacity

When you delete a Power BI Premium, Trial, or Fabric capacity, the process soft deletes non-Power BI Fabric items in workspaces assigned to the capacity. You can still see these Fabric items in the OneLake catalog and the workspace list, but you can't open or use them. If you associate the workspace that holds these items to a capacity (other than Power BI Embedded) from the same region as the deleted capacity within seven days, the deleted items are restored. This seven-day period is separate from the workspace retention setting.

You must be a Fabric administrator to delete a capacity. To delete an Azure-managed capacity, you also need the applicable Azure RBAC delete permission on the capacity resource, such as `Microsoft.Fabric/capacities/delete` for an F SKU or `Microsoft.PowerBIDedicated/capacities/delete` for a Power BI Embedded capacity.

> **Note**
>
> To delete a trial capacity, you need to **end the trial**.

To delete a capacity, follow these steps:

1. Go to **OneLake catalog** > **Govern** > **Capacities**.
2. On the **Manage capacities** page, filter or search the list to find the capacity.
3. Next to the capacity name, select **More options** (three dots) > **Delete**.
4. In the Delete capacity pane, confirm the deletion by following the prompts.

End a trial

To end a trial capacity, follow these steps:

1. Go to **OneLake catalog** > **Govern** > **Capacities**.
2. On the **Manage capacities** page, filter or search the list to find the trial capacity.
3. Next to the capacity name, select **More options** (three dots) > **End trial**.
4. In the End trial pane, confirm the action by following the prompts.

Manage capacity settings

In addition to the basic actions available in the actions menu, you can access all settings for a capacity.

To open the full capacity settings experience, follow these steps:

1. Go to **OneLake catalog** > **Govern** > **Capacities**.
2. On the **Manage capacities** page, filter or search the list to find the capacity.
3. Next to the capacity name, select **More options** (three dots) > **Settings**.

The **Capacity settings** section opens and displays all settings specific to the capacity. The **Delegated tenant settings** section allows Fabric admins to delegate tenant settings for capacity admins to manage. Changes to these settings affect only the capacity where you make them.

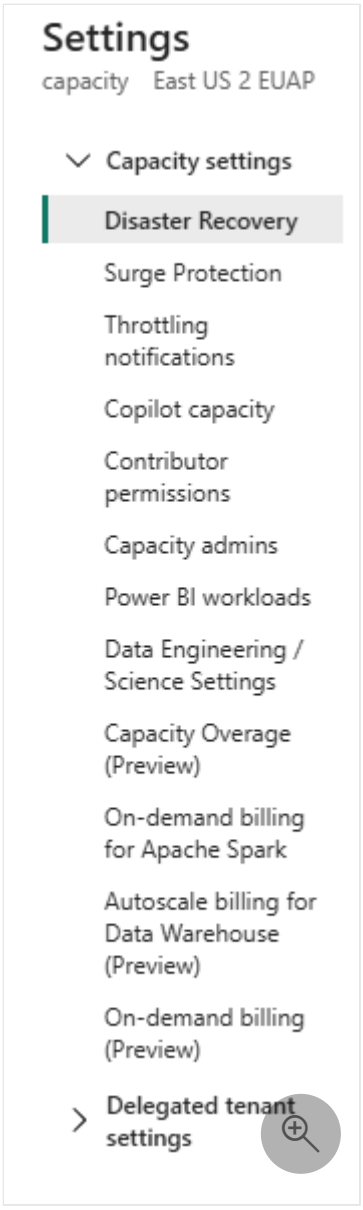

> **Note**
>
> Delegated tenant settings are available for Power BI Premium and Fabric capacities.

Settings details

This table summarizes the settings available for a capacity.

> **Note**
>
> - Some features in the table are available only if they're enabled in the tenant.
> - Trial capacities only have some of the settings listed in the table.

Expand table

Details setting name	Description
Disaster Recovery	Enable disaster recovery for the capacity
Surge protection	Enable surge protection for the capacity
Throttling notifications	Enable notification for your capacity
Copilot capacity	Designate this capacity as a Copilot in Fabric capacity
Contributor permissions	Set up the ability to add workspaces to the capacity. Select one of these two options: • The entire organization • Specific users or security groups
Capacity admins	Give specific users the ability to do the following: • Change capacity settings • Add contributors to the capacity • Add or remove workspaces from the capacity
Power BI workloads	Configure Power BI workloads for: • Semantic models • Paginated reports • AI
Data Engineering/Science Settings	Allow workspace admins to set the size of their Spark pools
Capacity overage (preview)	Enable capacity overage to allow the capacity to temporarily exceed its allocated resources.
On-demand billing for Apache Spark	Enable on-demand billing for Apache Spark to allow the capacity to be billed based on actual usage rather than allocated resources.

* To assign a workspace to a Fabric capacity or a capacity with an A SKU, you need a capacity **contributor** role and a workspace admin role. A **contributor** on a capacity can assign workspaces to that capacity but can't modify capacity settings or delete the capacity. By using this role, workspace admins can move their workspaces into a managed capacity without full administrative control.

Delegated tenant settings

Use delegating admin settings to grant granular access to features in the capacity. The delegated tenant settings section lists these tenant settings:

- Workload management tenant settings that Fabric automatically delegates to the capacity.
- Tenant settings that the Fabric admin delegates.

By default, delegated tenant settings inherit their configuration from the tenant. To override this configuration, follow these steps. When you enable tenant setting delegation, you can disable it by clearing the **Override tenant admin selection** checkbox.

1. From the **Delegate tenant setting** list, open the setting you want to delegate permissions for.

2. Select the **Override tenant admin selection** checkbox.

3. Select **Enabled**.

4. In the **Apply to** section, select one of the following options:

    - **All the users in capacity** - Delegate the setting to all the users in the capacity.
    - **Specific security groups** - Apply the setting to specific security groups. Enter the security groups you want to apply the setting to.

    To exclude specific security groups from the setting, select **Except specific security groups** and enter the security groups you want to exclude. This setting is optional, and you can use it together with the **Apply to** setting.

5. Select **Apply**.

Add and remove admins

To add and remove admins from your capacity, you need to be a capacity admin. Only users that belong to the tenant the capacity is part of can be admins of the capacity.

1. Go to **OneLake catalog** > **Govern** > **Capacities**.
2. On the **Manage capacities** page, filter or search the list to find the capacity.
3. Next to the capacity name, select **More options** (three dots) > **Settings**.
4. In **Settings**, select **Capacity admins**, and then add or remove admins as needed.
5. Select **Apply** to save your changes.

Related content

- Fabric licenses
- About tenant settings

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Manage your Fabric capacity

This article describes the Microsoft Fabric capacity settings. The article is aimed at admins who want to understand how to manage their Microsoft Fabric capacities.

Get to the capacity settings

To get to the capacity settings, follow these steps:

1. In Microsoft Fabric, select the gear icon (⚙), and then select **Admin portal**.
2. In the Admin portal, select **Capacity settings**.

View your capacity

The capacity settings page shows a list of all the capacities in your tenant. At the top of the page you can see a list of the different Fabric capacity types. Select a capacity type to view all the capacities of that type in your tenant.

- **Power BI Premium** - A capacity that was bought as part of a Power BI Premium subscription. These capacities use P SKUs.

  > ⓘ **Note**
  >
  > Power BI capacities are transitioning to Fabric. For more information, see Power BI Premium transition to Microsoft Fabric.

- **Power BI Embedded** - A capacity that was bought as part of a Power BI Embedded subscription. These capacities use A or EM SKUs.
- **Trial** - A Microsoft Fabric trial capacity. These capacities use Trial SKUs.
- **Fabric capacity** - A Microsoft Fabric capacity. These capacities use F SKUs.

The rest of this article is divided to sections based on the different capacity types. To view the settings of your capacity, select the tab that matches your capacity type. If there's no tab to select, the section applies to all capacity types.

Manage your capacity

This section lists basic capacity management tasks, such as creating a new capacity, changing a capacity's name and deleting a capacity.

Create a new capacity

To create a new capacity you need to be a Microsoft Fabric admin.

Power BI Premium

To create a new Power BI Premium capacity, follow these steps:

1. On the **Capacity settings** page, select **Power BI Premium**.
2. Select **Set up new capacity**.
3. On the **Set up a new capacity** page, enter the following information:
    - **Capacity name** - Give your capacity a name.
    - **Capacity admins** - Add capacity admins.
    - **Region** - Select the region you want to create the capacity in.
    - **Available v-cores** - Select the number of v-cores you want to use for the capacity.
    - **Capacity size** - Select the size of the capacity.

      > ⓘ **Note**
      >
      > If you select an EM size, you'll create a Power BI Embedded capacity.
4. Select **Create**.

Change the name of your capacity

To change the name of your capacity, you need to be a capacity admin. To become a capacity admin, you need to be assigned the capacity admin role in the capacity settings. For more information, see Add and remove admins.

Power BI Premium

To change the name of your Power BI Premium capacity, follow these steps:

1. On the **Capacity settings** page, select **Power BI Premium**.

2. From the list of capacities, select the gear icon (⚙) next to the capacity you want to change.
3. Select the pencil icon next to the **Capacity name** field.
4. Enter the new name for the capacity, and then select the checkmark icon (✓).

Add and remove admins

To add and remove admins from your capacity, you need to be a capacity admin. Only users that belong to the tenant the capacity is part of, can be added as admins to the capacity.

Power BI Premium

To add or remove admins in a Power BI Premium capacity, follow these steps:

1. On the **Capacity settings** page, select **Power BI Premium**.
2. From the list of capacities, select the capacity you want to make changes to.
3. From the *Details* tab, select expand **Admin permissions**.
4. Add or remove admins from the text box.
5. Select **Apply**.

Resize a capacity

To resize your capacity, you need to be a capacity admin. To become a capacity admin, you need to be assigned the capacity admin role in the capacity settings. For more information, see Add and remove admins.

Power BI Premium

To resize a Power BI Premium capacity, follow these steps:

1. On the **Capacity settings** page, select **Power BI Premium**.
2. Select the capacity you want to resize.
3. Select **Change size**.
4. In the *Change size* window, from the **Capacity size** dropdown, select the new size for the capacity.
5. Select **Apply**.

Delete a capacity

To delete a capacity you need to be a Microsoft Fabric admin.

When you delete a Power BI Premium, Trial or Fabric Capacity, non-Power BI Fabric items in workspaces assigned to the capacity are soft deleted. These Fabric items can still be seen in OneLake Data Hub and in the workspace list, but can't be opened or used. If the workspace that holds these items is associated to a capacity (other than Power BI Embedded) from the same region as the deleted capacity within seven days, the deleted items are restored. This seven-day period is separate from the workspace retention setting.

Power BI Premium

To delete a Power BI Premium capacity, follow these steps:

1. On the **Capacity settings** page, select **Power BI Premium**.
2. From the list of *Power BI Premium* capacities, select the gear icon (⚙) next to the capacity you want to delete.
3. Select **Delete capacity**.
4. From the confirmation dialog, select **Delete**.

Autoscale

Power BI Premium

To enable autoscale on a Power BI Premium capacity, see Using Autoscale with Power BI Premium.

Capacity settings

After selecting a capacity, you can control its settings from these two tabs:

- **Details** - Capacity details are settings that are specific to the capacity.
- **Delegated tenant settings** - Tenant settings are delegated by Fabric admins to be managed by capacity admins. Changes to these settings only affect the capacity the changes are made in.

> ⓘ **Note**

Delegated tenant settings are available for Power BI Premium and Fabric capacities.

To view the settings of a specific capacity, follow these steps:

1. Go to the **Capacity settings** page.
2. Select the capacity type your capacity belongs to.
3. From the list, select the capacity you want to view.

Details

This table summarizes the actions you can take in the details section.

> **Note**
>
> - Some of the features in the table are only available if they are enabled in the tenant.
> - Trail capacities only have some of the settings listed in the table

Expand table

Details setting name	Description
Disaster Recovery	Enable disaster recovery for the capacity
Capacity usage report	The usage report is replaced with the capacity metrics app
Notifications	Enable notification for your capacity
Copilot capacity	Designate this capacity as a Fabric Copilot capacity
Contributor permissions	Set up the ability to add workspaces to the capacity. Select one of these two options: • The entire organization • Specific users or security groups
Admin permissions	Give specific users the ability to do the following: • Change capacity settings • Add contributors to the capacity • Add or remove workspaces from the capacity
Power BI workloads	Configure Power BI workloads for: • Semantic models • Paginated reports • AI
Preferred capacity for My workspace	Designate the capacity as the default capacity for My workspaces

Details setting name	Description
Data Engineering/Science Settings	Allow workspace admins to set the size of their spark pools
Workspaces assigned to this capacity	*Add or remove workspaces assigned to the capacity

* To assign a workspace to a Fabric capacity or a capacity with an A SKU, you need to have a capacity **contributor** role, and a workspace admin role. A **contributor** on a capacity can assign workspaces to that capacity but can't modify capacity settings or delete the capacity. This role is typically used to allow workspace admins to move their workspaces into a managed capacity without giving them full administrative control.

Delegated tenant settings

Delegating admin settings can be used to grant granular access to features in the capacity. The delegated tenant settings section lists these tenant settings:

- Workload management tenant settings that are automatically delegated to the capacity.
- Tenant settings delegated by the Fabric Admin.

By default, delegated tenant settings inherit their configuration from the tenant. To override this configuration, follow the steps below. When the tenant setting delegation is enabled, you can disable delegation by clearing the *Override tenant admin selection* checkbox.

1. From the **Delegate tenant setting** list, open the setting you want to delegate permissions for.

2. Select the **Override tenant admin selection** checkbox.

3. Select **Enabled**

4. In the *Apply to* section, select one of the following options:

    - **All the users in capacity** - Delegate the setting to all the users in the capacity.
    - **Specific security groups** - Apply the setting to specific security groups. Enter the security groups you want to apply the setting to.

    To exclude specific security groups from the setting, select **Except specific security groups** and enter the security groups you want to exclude. This setting is optional and can be used with together with the *Apply to* setting.

5. Select **Apply**.

Manage max memory for semantic models

To optimize performance and prevent memory-related errors, administrators can adjust the **Max memory (%)** setting for semantic models within the Power BI Premium and Power BI Embedded capacity settings.

1. Go to the **Admin portal**.
2. Select **Capacity settings**.
3. Choose the relevant capacity and click on the **Workloads** tab.
4. Locate the **Semantic models** section.
5. Adjust the **Max memory (%)** slider to allocate more memory as needed.

> [NOTE] Increasing memory allocation may help resolve issues such as model loading failures or performance bottlenecks. For more context on SKU limitations, see Semantic model SKU limitation.

Related content

- Microsoft Fabric licenses
- About tenant settings

ⓘ **Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 03/06/2026

About tenant settings

Tenant settings enable fine-grained control over the features that are made available to your organization. If you have concerns around sensitive data, some of our features might not be right for your organization, or you might only want a particular feature to be available to a specific group.

Tenant settings that control the availability of features in the Power BI user interface can help to establish governance policies, but they're not a security measure. For example, the **Export data** setting doesn't restrict the permissions of a Power BI user on a semantic model. Power BI users with read access to a semantic model have the permission to query this semantic model and might be able to persist the results without using the **Export data** feature in the Power BI user interface.

For a list and brief description of all the tenant settings, see the tenant settings index.

> **Note**
>
> It can take up to 15 minutes for a setting change to take effect for everyone in your organization.

New tenant settings

To help you quickly identify changes and respond, a message at the top of the tenant settings page appears when there's a change. The message lists new tenant settings and changes to existing ones.

You can identify new settings according to their *new* icon.

How to get to the tenant settings

To access the tenant settings:

1. Sign in to Fabric using your admin account credentials.
2. Select **OneLake catalog**, and then select the **Govern** tab.
3. Select **Configurations** > **Tenant settings**.

> **Note**
>
> OneLake catalog and Govern are rolling out by region. If they're not available in your region, select the **Settings** (gear) icon > **Admin portal** > **Tenant settings** during the regional rollout.

How to use the tenant settings

Many of the settings can have one of three states:

- **Disabled for the entire organization**: No one in your organization can use this feature.

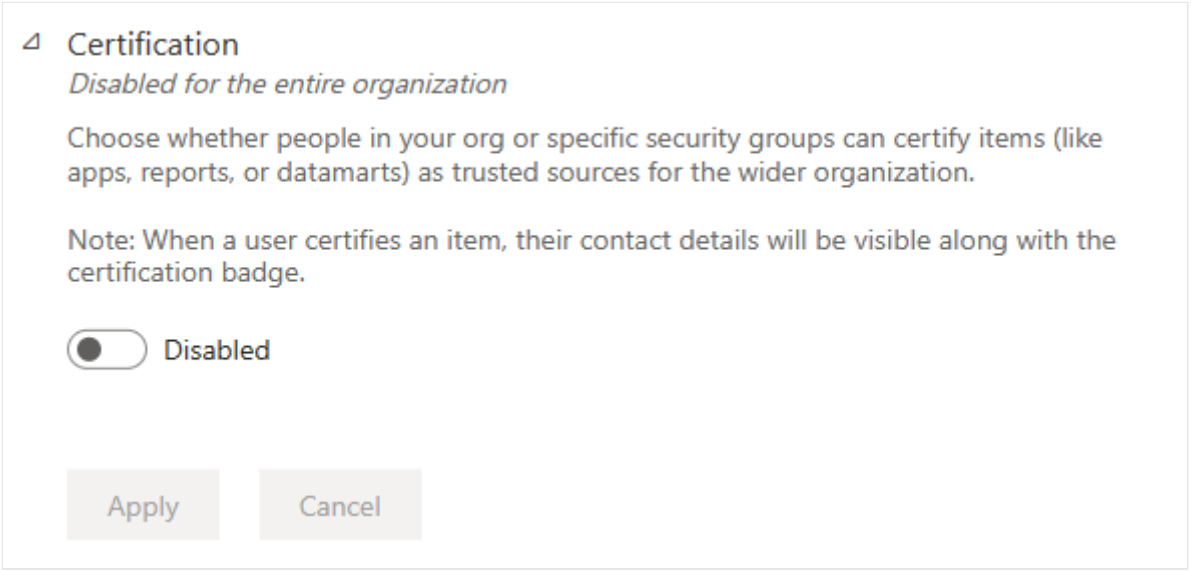

- **Enabled for the entire organization**: Everyone in your organization can use this feature.

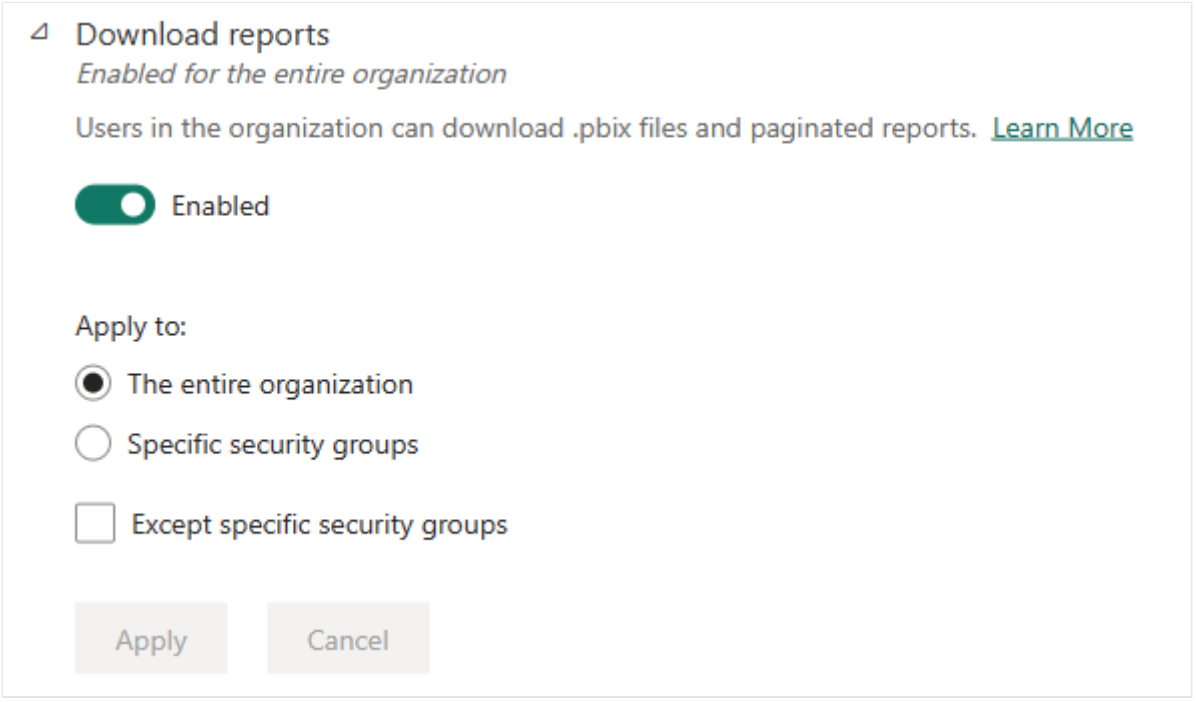

- **Enabled for the entire organization except for certain groups**: Everyone in your organization can use this feature except for users who belong to the specified groups.

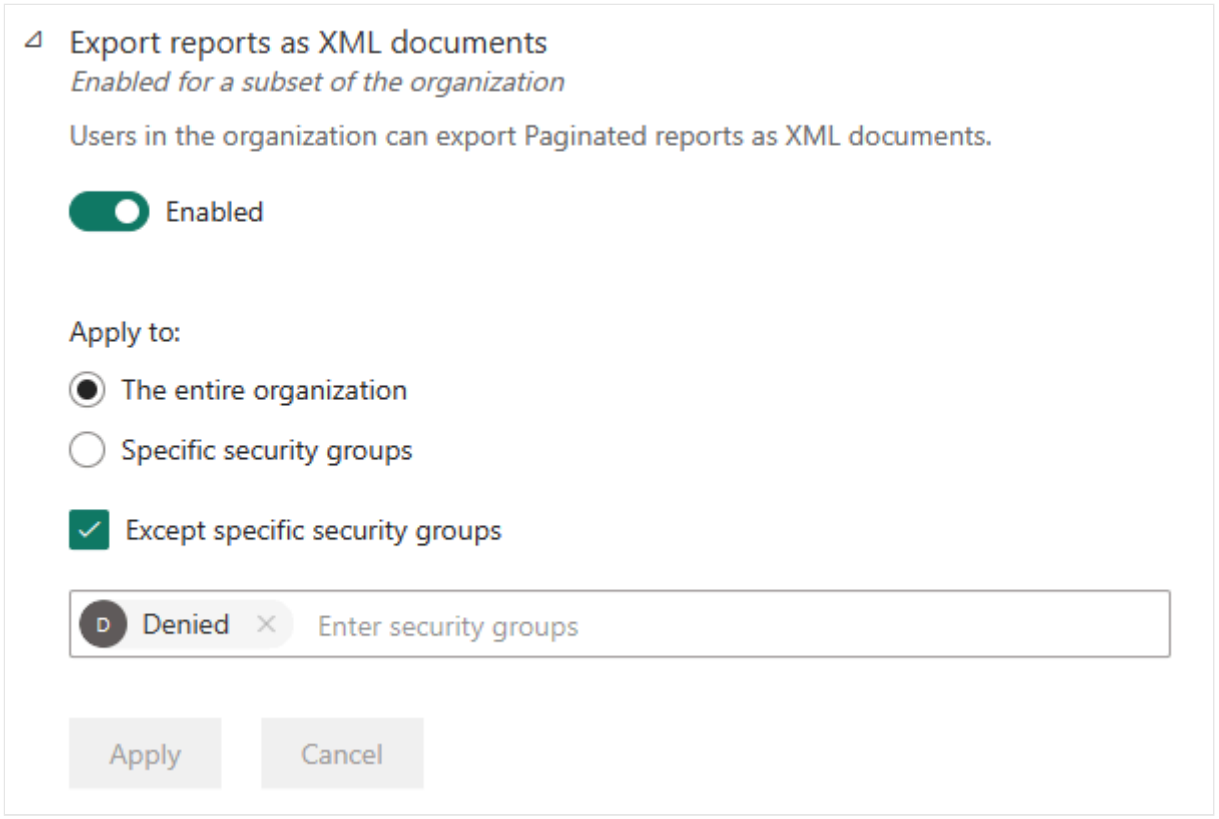

- **Enabled for a subset of the organization**: Specific security groups in your organization are allowed to use this feature.

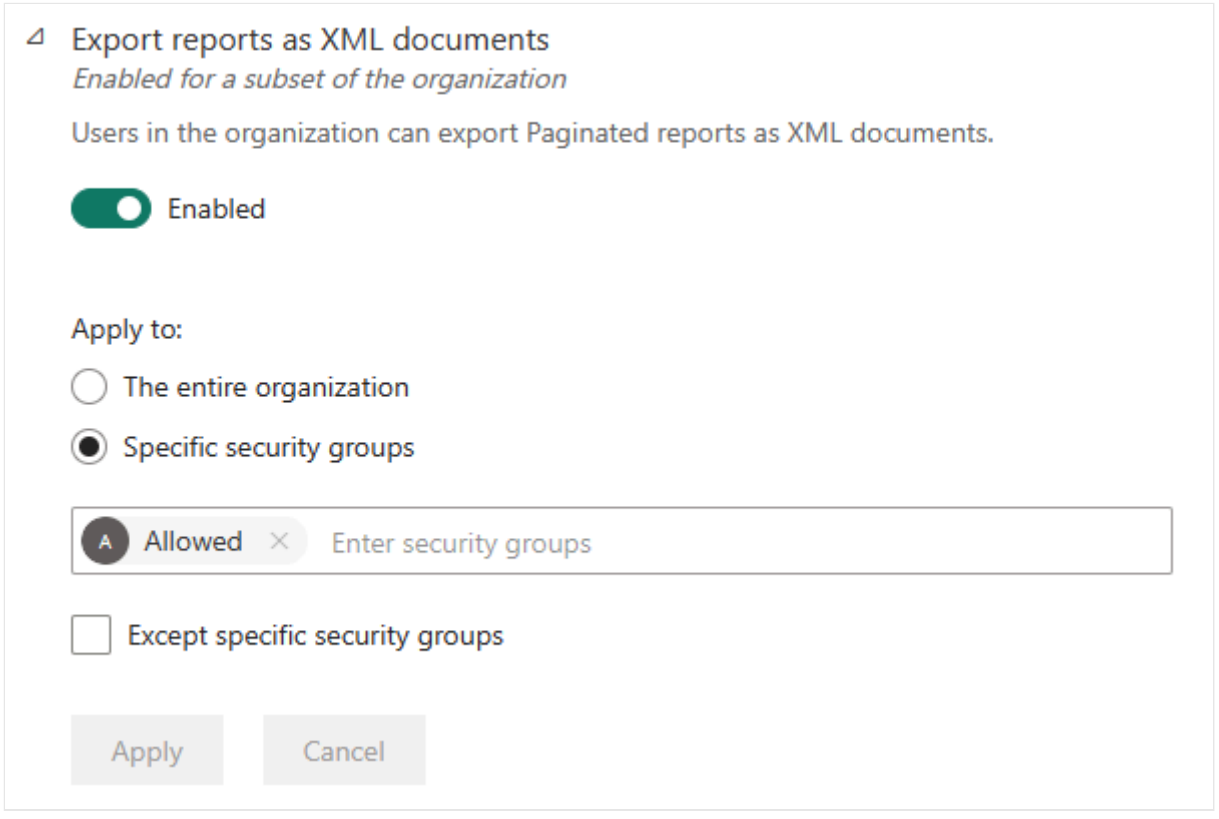

- **Enabled for specific groups except for certain groups**: Members of the specified security groups are allowed to use this feature, unless they also belong to an excluded group. This approach ensures that certain users don't have access to the feature even if they're in the allowed group. The most restrictive setting for a user applies.

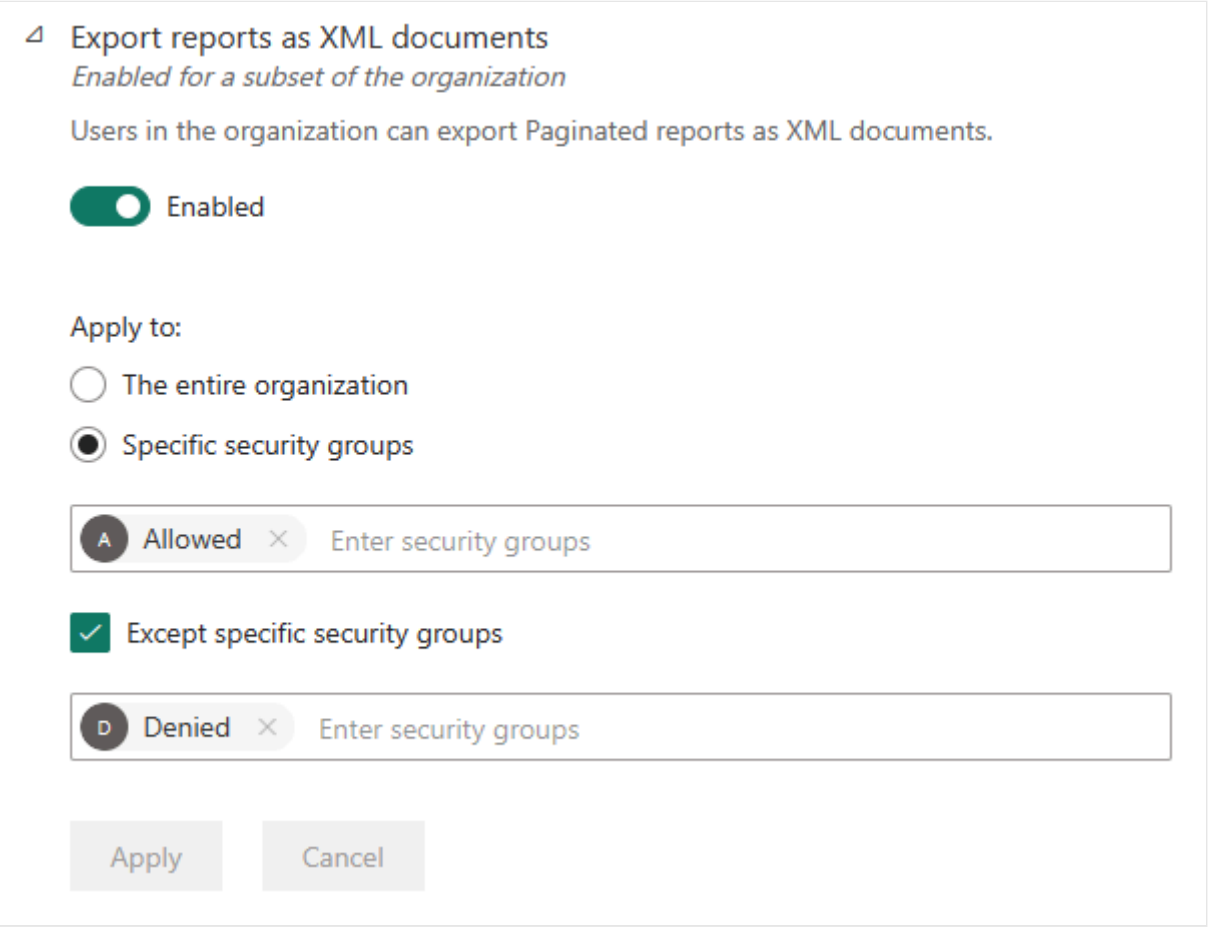

Related content

- Use the Fabric REST API to list tenant settings
- Administration overview
- Tenant settings index

Last updated on 09/29/2026

Delegate tenant settings

Article • 04/21/2024

Microsoft Fabric allows organizations to delegate settings from the tenant to the capacity, and from the capacity to workspaces. Delegation allows the organization to give admins control over specific settings relevant to their area of responsibility. Delegation prevents centralized admins from becoming a bottle neck for teams across the organization that require control over specific settings.

Here are some key concepts related to delegating settings in Fabric:

- Tenant Settings - Global settings controlled by tenant administrators that impact the entire tenant.
- Domain - A logical grouping of workspaces aimed at facilitating data mesh architecture.
- Capacity - A dedicated compute resource. Capacity admins can delegate settings concerning the performance and consumption of compute resources.
- Workspace - Collaborative environments where Fabric items are stored and shared. Certain tenant settings can be delegated to workspaces through a domain or a capacity.

Domain, capacity and workspace admins can override tenant settings. Overriding tenant settings allows admins to modify their environment to meet specific requirements. When a setting is adjusted in a domain or capacity, it affects only the workspaces linked to those administrative units. Similarly, changes to a workspace impact only the items stored within that workspace.

Delegation design

A setting can be delegated either through domains or capacities, but not both. This ensures clarity in governance and prevents conflicts in the management of settings.

After delegating to a domain or capacity, certain settings can be further delegated to workspaces. This allows for finer granularity in control, empowering workspace owners to customize the settings to their requirements. Tenant admins can bypass domain and capacity admins and delegate directly to workspaces.

Delegate settings to workspaces

Follow these steps to delegate settings to workspaces:

1. Delegate tenant settings to a capacity or a domain.
2. Locate the setting you want to delegate to a workspace in your domain or capacity.
3. Select the option to delegate to a workspaces. Some settings can't be delegated to workspaces. In such cases, there isn't an option to delegate to a workspace.
4. Select **Apply**

Audit your delegated settings

The following activity events represent tenant setting changes:

- **UpdatedAdminFeatureSwitch** - Tenant delegation changes.
- **UpdateCapacityTenantSettingDelegation** - Capacity delegation changes.
- **UpdateDomainTenantSettingDelegation** - Domain delegation changes.
- **UpdateWorkspaceTenantSettingDelegation** - Workspace delegation changes.

Related content

- What is the admin monitoring workspace?
- Workspace tenant settings
- Manage workspaces

Feedback

Was this page helpful? Yes No

Provide product feedback | Ask the community

Tenant settings index

This article lists all Fabric tenant settings, along with a brief description of each, and links to relevant documentation, if available. For more information about tenant settings in general, see About tenant settings.

If you want to get to the tenant settings in the Fabric portal, see How to get to the tenant settings.

Microsoft Fabric

Expand table

Setting name	Description
Users can create Fabric items	Users can use production-ready features to create Fabric items. Turning off this setting doesn't impact users' ability to create Power BI items. This setting can be managed at both the tenant and the capacity levels. Learn More
Users can create Ontology (preview) items	Users can create ontologies to unify enterprise semantics across data, models and logic to operationalize decision intelligence with context-aware AI agents. Learn More
Users can create Digital Twin Builder (preview) items	Users can create digital twin builder items to build comprehensive digital twins of real world environments and processes, to enable big-picture data analysis and drive operational efficiency.
Product Feedback	This setting allows Microsoft to prompt users for feedback through in-product surveys within Microsoft Fabric and Power BI. Microsoft will use this feedback to help improve product features and services. User participation is voluntary. Learn More
Users can be informed of upcoming conferences featuring Microsoft Fabric when they are logged in to Fabric	Attending conferences can help your data teams learn in-depth about how to best use the Fabric platform for your business needs and build professional relationships with community members, Microsoft engineering and product teams, and the Fabric Customer Advisory Team (CAT). These conferences may be organized and hosted by third parties. Learn More
ML models can serve real-time predictions from API endpoints (preview)	With this setting on, users can create real-time predictions from model and version endpoints. Even with real-time endpoints turned off, batch predictions can still be generated. Learn More

Setting name	Description
Detect anomalies in Real-Time Intelligence (Preview)	This setting allows users to use statistical detection algorithms to detect anomalies in real-time data. Learn More
Enable Event Schema Sets in Real-Time Hub (preview)	This setting allows users to define the structure and format of events through schemas, enabling data quality and validation with group-level access control. Learn More
Users can create dbt job items (preview)	Users can import, author and execute dbt (data build tool) projects directly within Fabric. This allows users to access a powerful transformation engine that connects seamlessly with SQL-based workflows without having to setup a CLI environment. Learn More
All Power BI users can see "Set alert" button to create Fabric Activator alerts	When enabled, all Power BI users will see the "Set alert" button in reports. However, only users with permission to create Fabric items can actually set up Fabric Activator alerts, which send real-time notifications based on predefined data conditions. Learn More
Enable new mirrored catalog items (preview)	Turn on this setting to allow users to create mirrored catalog items in Fabric, to mirror tables from external catalogs into OneLake through shortcuts. This only affects mirrored catalog item types that are in a Preview state. Learn More
Enable Fabric App Items (preview)	Users can create Fabric apps which is a Fabric-native suite of services for building intelligent applications for front end and backend of your application. It allows developers and AI agents to create and operate complete app backends directly on Microsoft Fabric. Fabric Apps unifies operational data, analytics, and AI while meeting enterprise security and governance requirements. Learn More
Enable anonymous data access for Fabric Apps (Preview)	This setting allows users to enable anonymous data access for Fabric Apps that are intended for public consumption. Admins can limit this capability to specific users or security groups.
Users can create Azure Databricks Storage items (preview)	Users can create Azure Databricks storage item that uses OneLake for storage.
Enable sending Operations Agents observability data to Agent 365 (preview)	Allow Microsoft Fabric Operation Agent data to be processed and stored by Agent 365, such as information about the Fabric item for this agent, the user ID of the creator of the agent, the Entra Agent ID that the agent is operating with, and the prompts and tools that it uses. This will help Agent 365 administrators monitor what Fabric agents are being created, where they are being used and more. When this setting is turned on, data sent to Agent 365 can be processed and/or stored outside your capacity's geographic region, compliance boundary, or national cloud instance. Learn More
Insights for operations agents (preview)	Users can view operational insights for events generated by operations agents. These insights help users understand what happened and why by showing summarized findings and key observations based on agent activity and underlying

Setting name	Description
	data. Learn More Turning on this setting can help your organization improve issue triage by identifying possible causes behind events, understand the context behind agent recommendations, and increase transparency into how agents analyze and act on data.
Fabric data agents can send operational metadata for observability in Microsoft Foundry	Agent creators can use operational metadata to monitor agent health, detect issues, and understand how their published agents are used across applications and services. These insights help agent creators improve agent reliability and performance. When this setting is turned on, operational metadata from Fabric data agents is processed and stored in a configured Application Insights resource. This data includes agent display names, data source names, reasoning step IDs, and conversation IDs. The Application Insights resource that processes and stores this data may be located outside your tenant's geographic region, compliance boundary, or national cloud instance. Accessing the operational metadata requires a role with read permissions (such as Monitoring Reader or Log Analytics Reader) on the configured Application Insights resource. Users with this access can view the data in Application Insights. Users can also view the data through a Foundry project that connects to the configured Application Insights resource. Access through Foundry requires role-based access to the Foundry project (Foundry User or higher) in addition to read permissions on the Application Insights resource. Access to a Foundry project alone doesn't include access to the telemetry data. Learn More

Help and support settings

Expand table

Setting name	Description
Publish "Get Help" information	Users in the organization can go to internal help and support resources from the Power BI help menu.
Receive email and Teams notifications for service outages or incidents	Mail-enabled security groups will receive email and Teams notifications if this tenant is impacted by a service outage or incident.
Receive In-product notifications for service outages or incidents	Mail-enabled security groups will receive In-product notifications if this tenant is impacted by a service outage or incident. Learn More
Users can try Microsoft Fabric paid features	When users sign up for a Microsoft Fabric trial, they can try Fabric paid features for free for 60 days from the day they signed up. Learn More

Setting name	Description
Show a custom message before publishing reports	When people attempt to publish a report, they'll see a custom message before it gets published.

Domain management settings

Expand table

Setting name	Description
Allow tenant and domain admins to override workspace assignments (preview)	Tenant and domain admins can reassign workspaces that were previously assigned to one domain to another domain.

Workspace settings

Expand table

Setting name	Description
Create workspaces	Users in the organization can create app workspaces to collaborate on dashboards, reports, and other content. Even if this setting is disabled, a workspace will be created when a template app is installed.
Use semantic models across workspaces	Users in the organization can use semantic models across workspaces if they have the required Build permission.
Block users from reassigning personal workspaces (My Workspace)	Turn on this setting to prevent users from reassigning their personal workspaces (My Workspace) from Premium capacities to shared capacities. Learn More
Define workspace retention period	Turn on this setting to define a retention period during which you can restore a deleted workspace and recover items in it. At the end of the retention period, the workspace is permanently deleted. By default, workspaces are always retained for a minimum of 7 days before they're permanently deleted. Turn off this setting to accept the minimum retention period of 7 days. After 7 days the workspace and items in it will be permanently deleted. Enter the number of days to retain a workspace before it's permanently deleted. My Workspace workspaces will be retained for 30 days automatically. Other workspaces can be retained for up to 90 days.
Automatically convert and store reports using Power BI	Enable this setting to automatically convert reports to PBIR format after editing and save them in the workspace.

Setting name	Description
enhanced metadata format (PBIR) (preview)	When this is activated, new reports will also be created in PBIR format. PBIR format provides source control-friendly file structures, enhancing co-development and boosting development efficiency for Power BI reports. Learn More
Fabric item recovery	Item recovery is turned on by default, and deleted items are retained for a minimum of 3 days. You can specify a longer retention period for a maximum of 90 days. At the end of the retention period, the items are permanently deleted. Note: Some Fabric items do not support item recovery. Learn more If you turn this setting off, deleting items will be a permanent action. Enter the number of days to retain items before they're permanently deleted.

Information protection

Expand table

Setting name	Description
Allow users to apply sensitivity labels for content	With this setting enabled, Microsoft Purview Information Protection sensitivity labels published to users by your organization can be applied. All prerequisite steps must be completed before enabling this setting. Important: Sensitivity-label-based access control for Fabric and Power BI data and content is only enforced in the tenant where the labels were applied, in Power BI Desktop (.pbix) files, and in Excel, PowerPoint, and PDF files generated via supported export paths. Sensitivity-label-based access control is not supported in cross-tenant scenarios, such as external data sharing, or in any other export scenario, such as export to .csv or .txt formats. For more information, see Information protection in Microsoft Fabric: Access control. Note: Sensitivity label settings, such as encryption and content marking for files and emails, are not applied to content in Fabric. Learn More. Encryption is applied to content in supported export paths.
Apply sensitivity labels from data sources to their data in Power BI	Only sensitivity labels from supported data sources will be applied. Please see the documentation for details about supported data sources and how their sensitivity labels are applied in Power BI. Learn about supported data sources
Automatically apply sensitivity labels to downstream content	With this setting enabled, whenever a sensitivity label is changed or applied to Fabric content, the label will also be applied to its eligible downstream content. Learn More

Setting name	Description
Allow workspace admins to override automatically applied sensitivity labels	With this setting enabled, workspace admins can change or remove sensitivity labels that were applied automatically by Fabric, for example, as a result of label inheritance. Learn More
Restrict content with protected labels from being shared via link with everyone in your organization	This setting will prevent content with protection settings in the sensitivity label from being shared via link with everyone in your organization. Learn More
Domain admins can set default sensitivity labels for their domains (preview)	Domain admins can set a default sensitivity label for their domains. The label they set will override your organization's default labels in Microsoft Purview, as long as it has a higher priority than the existing default labels set for your tenant. A domain's default label will automatically apply to new Fabric items created within the domain. Reports, semantic models, dataflows, dashboards, scorecards, and some additional item types aren't currently supported. Learn More
Allow Microsoft Purview to secure AI interactions	Allow Microsoft Purview to access, process, and store prompts and responses-including metadata-for data security and compliance outcomes such as sensitive info type (SIT) classification, reporting in Microsoft Purview Data Security Posture Management for AI, Audit, Insider Risk Management, Communication Compliance, and eDiscovery. Note: This is a Microsoft Purview paid capability and is not included in the Copilot in Fabric pricing. Learn More

Export and sharing settings

Expand table

Setting name	Description
External data sharing	Users can share a read-only link to data stored in OneLake with collaborators outside your organization. When you grant them permission to do so, users can share a link to data in lakehouses and additional Fabric items. Collaborators who receive the link can view, build on, and share the data both within and beyond their own Fabric tenants, using their organization's licenses and capacities. Learn More
Users can accept external data shares	Users can accept a read-only link to data from another organization's Fabric tenant. Users who accept an external share link can view, build on, and share the data, both inside and outside of your organization's tenant. For more information about the security limitations of this preview feature, view the feature documentation. Learn More
Guest users can access Microsoft Fabric	Guest users who've been added to your Microsoft Entra directory can access Microsoft Fabric and any Fabric items they have permissions to. Learn More

Setting name	Description
Users can invite guest users to collaborate through item sharing and permissions	Users can collaborate with people outside the organization by sharing Fabric items with them and granting them permission to access those items. After external users accept an invitation, they're added to your Microsoft Entra directory as guest users. Learn More
Guest users can browse and access Fabric content	Users can invite guest users to browse and request access to Fabric content. Learn More
Users can see guest users in lists of suggested people	With this setting on, users will see both users in your organization and guest users who've been added to your Microsoft Entra directory in lists of suggested people. With this setting off, users will see only users in your organization. Users can still share items with guests by providing their full email address. Learn More
Publish to web	People in your org can publish public reports on the web. Publicly published reports don't require authentication to view them. Go to Embed codes in the admin portal to review and manage public embed codes. If any of the codes contain private or confidential content remove them. Review embed codes regularly to make sure no confidential information is live on the web. Learn more about Publish to web
Copy and paste visuals	Users in the organization can copy visuals from a tile or report visual and paste them as static images into external applications.
Export to Excel	Users in the organization can export the data from a visualization or paginated report to an Excel file. Learn More
Export to .csv	Users in the organization can export data from a tile, visualization, or paginated report to a .csv file. Learn More
Download reports	Users in the organization can download .pbix files and paginated reports. Learn More
Users can work with semantic models in Excel using a live connection	Users can export data to Excel from a report visual or semantic model, or export a semantic model to an Excel workbook with Analyze in Excel, both options with a live connection to the XMLA endpoint. Learn More
Export reports as PowerPoint presentations or PDF documents	Users in the organization can export reports as PowerPoint files or PDF documents.
Export reports as MHTML documents	Users in the organization can export Paginated reports as MHTML documents.

Setting name	Description
Export reports as Word documents	Users in the organization can export Paginated reports as Word documents.
Export reports as XML documents	Users in the organization can export Paginated reports as XML documents.
Export reports as image files	Users in the organization can use the export report to file API to export reports as image files.
Print dashboards and reports	Users in the organization can print dashboards and reports.
Certification	Choose whether people in your org or specific security groups can certify items (like apps, reports, or datamarts) as trusted sources for the wider organization. Note: When a user certifies an item, their contact details will be visible along with the certification badge.
Endorse master data	Choose whether people in your org or specific security groups can endorse items (like lakehouses, warehouses, or datamarts) as one of the core sources for your organization's data records. Learn More Note: When someone endorses an item as master data, their name and email will show with the endorsement badge.
Users can set up email subscriptions	Users can create email subscriptions to reports and dashboards.
B2B guest users can set up and be subscribed to email subscriptions	B2B guest users can set up and be subscribed to email subscriptions. B2B guest users are external users that have been added to your Microsoft Entra ID. Turn this setting off to prevent B2B guest users from setting up or being subscribed to email subscriptions. Learn More
Users can send email subscriptions to external users	Users can send email subscriptions to external users. External users are users you've not added to your Microsoft Entra ID. Turn this setting off to prevent users from subscribing external users to email subscriptions. Learn More
Featured content	Users in the organization can promote their published content to the Featured section of Power BI Home.
Allow shareable links to grant access to everyone in your organization	This setting will grant access to anyone in your organization with the link. It won't work for external users. Learn More
Enable Microsoft Teams integration	This setting allows people in the organization to access features associated with the Microsoft Teams and Power BI integration. This includes launching Teams experiences from the Power BI service like chats, the Power BI app for Teams, and

Setting name	Description
	receiving Power BI notifications in Teams. To completely enable or disable Teams integration, work with your Teams admin.
Install Power BI app for Microsoft Teams automatically	The Power BI app for Microsoft Teams is installed automatically for users when they use Microsoft Fabric. The app is installed for users if they have Microsoft Teams and the Power BI app is allowed in the Teams Admin Portal. When the app is installed, users receive notifications in Teams and can more easily discover and collaborate with colleagues. The Power BI app for Teams provides users with the ability to open all Fabric content. Learn More.
Enable Power BI add-in for PowerPoint	Let people in your org embed Power BI data into their PowerPoint presentations. This integration requires that your organization's Microsoft Office admin has enabled support for add-ins.
Allow DirectQuery connections to Power BI semantic models	DirectQuery connections allow users to make changes to existing semantic models or use them to build new ones. Learn More
Guest users can work with shared semantic models in their own tenants	Authorized guest users can discover semantic models shared with them in the OneLake data hub (in Power BI Desktop), and then work with these semantic models in their own Power BI tenants.
Allow specific users to turn on external data sharing	Turn off this setting to prevent all users from turning on external data sharing. If this setting is on, all or specific users can turn on the external data sharing option, allowing them to share data with authorized guest users. Authorized guest users can then discover, connect to, and work with these shared semantic models in their own Power BI tenants.
Users with read or write permission can download data from notebooks	With this setting on, users with read or write permission can download or export data from tables in notebook outputs. Learn More

Discovery settings

Expand table

Setting name	Description
Make promoted content discoverable	Allow users in this org who can promote content to make content they promote discoverable by users who don't have access to it. Learn More
Make certified content discoverable	Allow users in the org who can certify content to make content they certify discoverable by users who don't have access to it. Learn More

Setting name	Description
Discover content	Allow users to find and request access to content they don't have access to if it was made discoverable by its owners. Learn More
Users can find objects in search	Include objects such as tables and columns in the OneLake catalog and global search. Users can find objects only if they have access to items that contain them. Disable to limit discovery to top-level items, like lakehouses and reports. Learn More

App settings

Expand table

Setting name	Description
Create template organizational apps	Users in the organization can create template apps that use semantic models built on one data source in Power BI Desktop.
Push apps to end users	Users can share apps directly with end users without requiring installation from AppSource.
Publish apps to the entire organization	Users in the organization can publish apps to the entire organization.

Integration settings

Expand table

Setting name	Description
Allow XMLA endpoints and Analyze in Excel with on-premises semantic models	Users in the organization can use Excel to view and interact with on-premises Power BI semantic models. This also allows connections to XMLA endpoints.
Semantic Model Execute Queries REST API	Users in the organization can query semantic models by using Data Analysis Expressions (DAX) through Power BI REST APIs.
Users can use the Power BI Model Context Protocol server endpoint (preview)	When enabled, the Power BI Model Context Protocol (MCP) server endpoint allows users to connect MCP clients to Power BI to use MCP tools to interact with their permissioned Power BI artifacts. GitHub Copilot in Visual Studio will operate by default; configure other MCP clients with service principals. To allow service principals to connect to the MCP Server, turn on "Service principals can call Fabric public APIs." Learn More
Use ArcGIS Maps for Power BI	Users in the organization can use the ArcGIS Maps for Power BI visualization provided by Esri.

Setting name	Description
Use global search for Power BI	Turn on this setting to let users use the global search bar at the top of the page.
Users can use the Azure Maps visual	With this setting on, users can create and view the Azure Maps visual. Your data may be temporarily stored and processed by Microsoft for essential services, including translating location names into latitudes and longitudes. Use of Azure Maps is subject to the following Terms of use .
Data sent to Azure Maps can be processed outside your tenant's geographic region, compliance boundary, or national cloud instance	Azure Maps services are currently not available in all regions and geographies. With this setting on, data sent to Azure Maps can be processed in a region where the service is available, which might be outside your tenant's geographic region, compliance boundary, or national cloud instance. Learn More
Data sent to Azure Maps can be processed by Microsoft Online Services Subprocessors	Some Azure Maps visual services, including the selection tool and the processing of location names within some regions, may require mapping capabilities provided in part by Microsoft Online Services subprocessors. Microsoft shares only necessary data with these Microsoft Online Services subprocessors, who may access data only to deliver the functions in support of online services that Microsoft has engaged them to provide and are prohibited from using data for any other purpose. Microsoft does not share the name of the customer or end user who submits the query. This feature is non-regional and the queries you provide may be stored and processed in the United States or any other country in which Microsoft or its subprocessors operate. Learn More
Map and filled map visuals	Allow people in your org to use the map and filled map visualizations in their reports.
Integration with SharePoint and Microsoft Lists	Users in the organization can launch Power BI from SharePoint lists and Microsoft Lists. Then they can build Power BI reports on the data in those lists and publish them back to the lists.
Dremio SSO	Enable SSO capability for Dremio. By enabling, user access token information, including name and email, will be sent to Dremio for authentication.
Snowflake SSO	Enable SSO capability for Snowflake. By enabling, user access token information, including name and email, will be sent to Snowflake for authentication. Learn More
Redshift SSO	Enable SSO capability for Redshift. By enabling, user access token information, including name and email, will be sent to Redshift for authentication.
Google BigQuery SSO	Enable SSO capability for Google BigQuery. By enabling, user access token information, including name and email, will be sent to Google BigQuery for authentication.

Setting name	Description
Enable Denodo SSO	Enable SSO capability for Denodo. By enabling, user access token information, including name and email, will be sent to Denodo for authentication.
Enable Starburst SSO	Enable SSO capability for Starburst. By enabling, user access token information, including name and email, will be sent to Starburst for authentication.
Microsoft Entra single sign-on for data gateway	Users can use Microsoft Entra single sign-on (SSO) to authenticate to on-premises / virtual network data gateways and access data sources. With this setting on, user access token information, including names and emails, is sent to data sources to authenticate to the on-premises / virtual network data gateway service. Learn More
Microsoft Entra single sign-on for virtual network (VNet) data gateway	Users can use Microsoft Entra single sign-on (SSO) to authenticate to virtual network (VNet) data gateways and access data sources. With this setting on, user access token information, including names and emails, is sent to data sources to authenticate to the virtual network data gateway service. Learn More
Users can view Power BI files saved in OneDrive and SharePoint (preview)	Users in the organization can view Power BI files saved in OneDrive for Business or SharePoint document libraries. The permissions to save and share Power BI files in OneDrive and SharePoint document libraries are controlled by permissions managed in OneDrive and SharePoint. Learn More
Users can share links to Power BI files stored in OneDrive and SharePoint through Power BI Desktop (preview)	Users who have saved Power BI files (.pbix) to OneDrive and SharePoint can share links to those files using Power BI Desktop. Learn More
Enable granular access control for all data connections	Enforce strict access control for all data connection types. When this is turned on, shared items will be disconnected from data sources if they're edited by users who don't have permission to use the data connections. Learn More
Semantic models can export data to OneLake	Semantic models configured for OneLake integration can send import tables to OneLake. Once the data is in OneLake, users can include the exported tables in Fabric items, including lakehouses and warehouses. Learn More
Semantic model owners can choose to automatically update semantic models from files imported from OneDrive or SharePoint	Semantic model owners can choose to allow semantic models to be automatically updated with changes made to the corresponding Power BI files (.pbix) stored in OneDrive or SharePoint. File changes can include new and modified data connections. Turn off this setting to prevent automatic updates to semantic models. Learn More
ArcGIS GeoAnalytics for Fabric Runtime	Users in your organization can use Esri's ArcGIS GeoAnalytics for Fabric Runtime in Microsoft's Fabric Spark Runtime. ArcGIS GeoAnalytics delivers

Setting name	Description
	spatial analysis to your big data by extending Apache Spark with ready-to-use spatial SQL functions and analysis tools. Learn More
Allow non-Entra ID auth in Eventstream	Users can enhance the security of data streaming by disabling key-based authentication in Eventstream's Custom Endpoint, ensuring that only Microsoft Entra ID (formerly Azure Active Directory) authentication is allowed. This reduces the risk of unauthorized access to Fabric Eventstream through non-Entra ID authentication methods.
Users can connect to data sources by using Apache Arrow database connectivity (ADBC)	ADBC offers enhanced performance, security, and reliability compared to legacy connectivity options. When enabled, data connections to supported data sources use ADBC by default. Learn More

Power BI visuals

Expand table

Setting name	Description
Allow visuals created using the Power BI SDK	Users in the organization can add, view, share, and interact with visuals imported from AppSource or from a file. Visuals allowed in the "Organizational visuals" page are not affected by this setting. Learn More
Add and use certified visuals only (block uncertified)	Users in the organization with permissions to add and use visuals can add and use certified visuals only. Visuals allowed in the "Organizational visuals" page are not affected by this setting, regardless of certification. Learn More
Allow downloads from custom visuals	Enabling this setting will let custom visuals download any information available to the visual (such as summarized data and visual configuration) upon user consent. It is not affected by download restrictions applied in your organization's Export and sharing settings. Learn More
AppSource Custom Visuals SSO	Enable SSO capability for AppSource custom visuals. This feature allows custom visuals from AppSource to get Microsoft Entra ID access tokens for signed-in users through the Authentication API. Microsoft Entra ID access tokens include personal information, including users' names and email addresses, and may be sent across regions and compliance boundaries. Learn More
Allow access to the browser's local storage	When this setting is on, custom visuals can store information on the user's browser's local storage. Learn More

R and Python visuals settings

Expand table

Setting name	Description
Interact with and share R and Python visuals	Users in the organization can interact with and share visuals created with R or Python scripts.

Audit and usage settings

Expand table

Setting name	Description
Usage metrics for content creators	Users in the organization can see usage metrics for dashboards, reports and semantic models that they have appropriate permissions to. Learn More
Per-user data in usage metrics for content creators	Usage metrics for content creators will expose display names and email addresses of users who are accessing content.
Show user data in the Fabric Capacity Metrics app and reports	With this setting on, active user data, including names and email addresses, are displayed in the Capacity Metrics app and reports. Learn More
Azure Log Analytics connections for workspace administrators	NO DESCRIPTION IN UI
Workspace admins can turn on monitoring for their workspaces (preview)	Workspace admins can turn on monitoring for their workspaces. When a workspace admin turns on monitoring, a read-only Eventhouse that includes a KQL database is created. After the Eventhouse and KQL database are added to the workspace, logging is turned on and data is sent to the database. Learn More
Microsoft can store query text to aid in support investigations	Query text for some items, including semantic models, is securely stored for usage during support investigations. Turn off this setting to stop the service from storing query text. Turning off this setting might negatively impact Microsoft's ability to provide support for the Fabric service. Learn More

Dashboard settings

Expand table

Setting name	Description
Web content on dashboard tiles	Users in the organization can add and view web content tiles on Power BI dashboards. Note: This may expose your org to security risks via malicious web content.

Developer settings

Expand table

Setting name	Description
Embed content in apps	Users in the organization can embed Power BI dashboards and reports in Web applications using "Embed for your customers" method. Learn More
Service principals can create workspaces, connections, and deployment pipelines	This setting allows service principals to create workspaces, connections, and deployment pipelines. To allow service principals to call the rest of Fabric public APIs, turn on the setting titled "Service principals can call Fabric public APIs". Learn More
Service principals can call Fabric public APIs	This setting allows service principals with the appropriate roles and item permissions to call Fabric public APIs. To allow service principals to create workspaces, connections, and deployment pipelines turn on the setting titled "Service principals can create workspaces, connections, and deployment pipelines". Learn More
Allow service principals to create and use profiles	Allow service principals in your organization to create and use profiles.
Block ResourceKey Authentication	For extra security, block using resource key based authentication. This means users not allowed to use streaming semantic models API using resource key.
Define maximum number of Fabric identities in a tenant	Allow admins to specify the maximum number of Fabric identities that can be created in a tenant. If this setting is disabled, up to 10000 Fabric identities can be created in a tenant. Learn more .

Admin API settings

Expand table

Setting name	Description
Service principals can access read-only admin APIs	Web apps registered in Microsoft Entra ID can use service principals, rather than user credentials, to authenticate to read-only admin APIs. To allow an app to use a service principal as an authentication method, the service principal must be added to an allowed security group. Service principals included in allowed security groups will have read-only access to all the information available

Setting name	Description
	through admin APIs, which can include users' names and emails, and detailed metadata about semantic models and reports. Learn More
Service principals can access admin APIs used for updates	Web apps registered in Microsoft Entra ID can use service principals, rather than user credentials, to authenticate to admin APIs used for updates. To allow an app to use a service principal as an authentication method, add the service principal to an allowed security group. Service principals in allowed security groups have full access to the information available through admin APIs, including users' names and emails, and detailed metadata about items. Learn More
Enhance admin APIs responses with detailed metadata	Users and service principals allowed to call Power BI admin APIs may get detailed metadata about Power BI items. For example, responses from GetScanResult APIs will contain the names of semantic model tables and columns. Learn More Note: For this setting to apply to service principals, make sure the tenant setting allowing service principals to use read-only admin APIs is enabled. Learn More
Enhance admin APIs responses with DAX and mashup expressions	Users and service principals eligible to call Power BI admin APIs will get detailed metadata about queries and expressions comprising Power BI items. For example, responses from GetScanResult API will contain DAX and mashup expressions. Learn More Note: For this setting to apply to service principals, make sure the tenant setting allowing service principals to use read-only admin APIs is enabled. Learn More

Gen1 dataflow settings

Setting name	Description
Create and use Gen1 dataflows	Users in the organization can create and use Gen1 dataflows. Learn More

Template app settings

Setting name	Description
Publish template apps	Users in the organization can publish template apps for distribution to clients outside of the organization. Learn More.
Install template apps	Users in the organization can install template apps created outside the organization. When a template app is installed, an upgraded workspace is created. Learn More
Install template apps not listed in AppSource	Users in the organization who have been granted permission to install template apps which were not published to Microsoft AppSource. Learn More.

Q&A settings

Setting name	Description
Review questions	Allow semantic model owners to review questions people asked about their data.
Synonym sharing	Allow people to share Q&A synonyms with your organization. Learn More

Explore settings (preview)

Setting name	Description
Users with view permission can launch Explore	Explore is a light-weight visual data exploration experience that enables people to quickly and easily do ad hoc analysis. This setting allows people with view permission on a semantic model to launch Explore from that model and from items connected to it. Learn More

Semantic Model Security

Setting name	Description
Block republish and disable package refresh	Disable package refresh, and only allow the semantic model owner to publish updates.

Advanced networking

Setting name	Description
Tenant-level Private Link	Increase security by allowing people to use a Private Link to access your Fabric tenant. Someone will need to finish the set-up process in Azure. If that's not you, grant permission to the right person or group by entering their email. Learn More | Set-up instructions Review the considerations and limitations section before enabling private endpoints.
Block Public Internet Access	For extra security, block access to your Fabric tenant via the public internet. This means people who don't have access to the Private Link won't be able to get in. Keep in mind, turning this on could take 10 to 20 minutes to take effect. Learn More | Set-up instructions
Configure workspace-level inbound network rules	With this setting on, workspace admins can configure inbound private link access protection in workspace settings. When a workspace is configured to restrict inbound network access, existing tenant-level private links can no longer connect

Setting name	Description
	to these workspaces. Turning off this setting reverts all workspaces to their previous configuration. Learn More
Configure workspace-level outbound network rules	With this setting on, workspace admins can configure outbound access protection in workspace settings. Turning off this tenant setting also turns off outbound access protection in all the workspaces in the tenant. Learn More
Configure workspace-level IP firewall rules and trusted resource instances	With this setting enabled, workspace admins can configure IP firewall rules and trusted resource instances in Workspace settings. Even when the tenant blocks public access, this setting still allows workspace admins to grant access from specific IP firewall rules and designated trusted resource instances. Learn More

Encryption

Setting name	Description
Apply customer-managed keys	With this setting turned on, users can configure workspace level encryption using customer-managed keys to protect their data. When turned off, the default is to use Microsoft managed keys. Learn More

Scorecards settings

Setting name	Description
Create and use Scorecards	Users in the organization can create and use Scorecards

User experience experiments

Setting name	Description
Help Power BI optimize your experience	Users in this organization will get minor user experience variations that the Power BI team is experimenting with, including content, layout, and design, before they go live for all users.

Share data with your Microsoft 365 services

Setting name	Description
Share Fabric data with your Microsoft 365 services	When this setting is enabled, Microsoft Fabric data can be stored and displayed in Microsoft 365 services. Fabric data (including Power BI report titles, chart axis labels, Fabric data agent instructions, or open and sharing history) may be used to improve Microsoft 365 services like search results and recommended content lists. Learn More

Setting name	Description
	Users can browse or get recommendations only for content they have access to. Users will see metadata about Fabric items (including refresh dates and workspace names in search listings) and see item content (like chart axis labels or titles reflected in Copilot summarizations) to enhance Microsoft 365 services. This setting is automatically enabled only if your Microsoft Fabric and Microsoft 365 tenants are in the same geographical region. You may disable this setting. Where is my Microsoft Fabric tenant located?

Insights settings

Setting name	Description
Receive notifications for top insights (preview)	Users in the organization can enable notifications for top insights in report settings
Show entry points for insights (preview)	Users in the organization can use entry points for requesting insights inside reports

Datamart settings

Setting name	Description
Create Datamarts (preview)	Users in the organization can create Datamarts

Semantic model settings

Setting name	Description
Users can edit semantic models in the Power BI service	Turn on this setting to allow users to edit semantic models in the service. This setting doesn't apply to DirectLake semantic models or editing a semantic model through an API or XMLA endpoint. Learn More

Kusto settings

Setting name	Description
Allow Real-Time Dashboards to include embedded app visuals	Enabling this setting allows dashboards to accept embedded app visuals, which are external apps embedded as iframes within Real-Time Dashboards

Scale-out settings

Setting name	Description
Scale out queries for large semantic models	For semantic models that use the large semantic model storage format, Power BI Premium can automatically distribute queries across additional semantic model replicas when query volume is high.

OneLake settings

Setting name	Description
Users can access data stored in OneLake with apps external to Fabric	Users can access data stored in OneLake with apps external to the Fabric environment, such as custom applications created with Azure Data Lake Storage (ADLS) APIs, OneLake File Explorer, and Databricks. Users can already access data stored in OneLake with apps internal to the Fabric environment, such as Spark, Data Engineering, and Data Warehouse. Learn More
Use short-lived user-delegated SAS tokens	OneLake SAS tokens enable applications to access data in OneLake through short-lived SAS tokens, based on a Microsoft Fabric user's Entra identity. These token's permissions can be further limited to provide least privileged access and cannot exceed a lifetime of one hour. Learn More
Authenticate with OneLake user-delegated SAS tokens	Allow applications to authenticate using a OneLake SAS token. Fabric users can create OneLake SAS by requesting a user delegation key. The tenant setting, Use short-lived user delegated SAS tokens, must be turned on to generate user delegation keys. The lifetimes of the user delegation keys and SAS tokens cannot exceed one hour. Learn More
Users can sync data in OneLake with the OneLake File Explorer app	Turn on this setting to allow users to use OneLake File Explorer. This app will sync OneLake items to Windows File Explorer, similar to OneDrive. Learn More
Include end-user identifiers in OneLake diagnostic logs	Control whether OneLake diagnostic logs capture end user identifiable information (EUII), such as email addresses and IP addresses. When enabled, these fields are recorded to support diagnostics, investigations, and usage analysis across your tenant. When disabled, these fields are redacted from new events. Learn More

Git integration

Setting name	Description
Users can synchronize workspace items with their Git repositories	Users can import and export workspace items to Git repositories for collaboration and version control. Turn off this setting to prevent users from syncing workspace items with their Git repositories. Learn More
Users can export items to Git repositories in other	The workspace and the Git repository may reside in different geographies. Turn on this setting to allow users to export items to Git

Setting name	Description
geographical locations	repositories in other geographies.
Users can export workspace items with applied sensitivity labels to Git repositories	Turn on this setting to allow users to export items with applied sensitivity labels to their Git repositories.
Users can sync workspace items with GitHub repositories	Users can select GitHub as their Git provider and sync items in their workspaces with GitHub repositories.

Copilot and Azure OpenAI Service

Setting name	Description
Users can use Copilot and other features powered by Azure OpenAI	When this setting is on, users can access Fabric features powered by Azure OpenAI, including Copilot and Fabric AI agents. Check documentation for the most recent list of these features. This setting can be managed at both the tenant and the capacity levels. Learn More For customers in the EU Data Boundary, this setting adheres to Microsoft Fabric's EU Data Boundary commitments. Learn More By turning on this setting, you agree to the Preview Terms for any AI features in preview.
Users can access a standalone, cross-item Power BI Copilot experience (preview)	When this setting is turned on, users will be able to access a Copilot experience that allows them to find, analyze, and discuss different Fabric items in a dedicated tab available via the Power BI navigation pane. This setting requires the following tenant setting to be enabled: "Users can use Copilot and other features powered by Azure OpenAI." Learn More
Data sent to Azure OpenAI can be processed outside your capacity's geographic region, compliance boundary, or national cloud instance	This setting is only applicable for customers who want to use Copilot and AI features in Fabric powered by Azure OpenAI, and whose capacity's geographic region is outside of the EU Data Boundary or the United States. Learn More When this setting is on, data sent to Copilot and other generative AI features can be processed outside your capacity's geographic region, compliance boundary, or national cloud instance. Check documentation for the types of data this might include. This setting can be managed at both the tenant and the capacity levels. Learn More By turning on this setting, you agree to the Preview Terms for any AI features in preview.
Capacities can be designated as Fabric Copilot capacities	With this setting on, capacity admins can designate capacities as Fabric Copilot capacities. Copilot capacities are special capacity types that allow your organization to consolidate users' Copilot usage and billing on a

Setting name	Description
	single capacity. Learn More When users use Copilot features, capacity admins can see the names of the items associated with users' Copilot activity. Learn More
Data sent to Azure OpenAI can be stored outside your capacity's geographic region, compliance boundary, or national cloud instance	This setting is only applicable for customers who want to use Copilot and AI features in Fabric powered by Azure OpenAI, and whose capacity's geographic region is outside of the EU Data Boundary or the United States. Learn More When this setting is turned on, data sent to Azure OpenAI can be stored outside your capacity's geographic region, compliance boundary, or national cloud instance. Check documentation for the types of experiences and data this might include. Learn More By turning on this setting, you agree to the Preview Terms for any AI features in preview.
Only show approved items in the standalone Copilot in Power BI experience (preview)	When this is turned on, only apps, data agents, and items marked as "approved for Copilot" will be shown in standalone Copilot. Users will still be able to manually attach items to ask questions. Copilot item usage is always subject to user permissions. Learn More
Users can use Fabric Copilot (preview)	When this setting is on, users can create new Fabric Copilot sessions and send prompts in existing sessions. When this setting is off, users can't create new sessions or send new prompts and existing session history remains read-only. By turning on this setting, you agree to the Preview Terms for any AI features in preview.

Azure Maps services

Setting name	Description
Users can use Azure Maps services	When this setting is enabled, users can access the features powered by Azure Maps services. Learn More For customers in the EU Data Boundary, this setting adheres to Microsoft Fabric's EU Data Boundary commitments. Learn More Use of Azure Maps is subject to the following Terms of use.
Data sent to Azure Maps can be processed outside your capacity's geographic region, compliance boundary or national cloud instance	Azure Maps services are currently not available in all regions and geographies. With this setting on, data sent to Azure Maps can be processed in a region where the service is available, which might be outside your capacity's geographic region, compliance boundary, or national cloud instance. Learn More

Setting name	Description
Users can use Azure Maps Weather Services (Preview)	When this setting is enabled, users can access weather data from Azure Maps Weather, sourced from AccuWeather Learn More

Additional workloads

Setting name	Description
Workspace admins can add and remove additional workloads (preview)	Workspace admins can add and remove workloads in their workspaces. If this setting is turned off, any existing workloads will stay added and items created with those workloads continue to work normally. When users interact with a workload, their data and access tokens, including name and email, are sent to the publisher. Sensitivity labels and protection settings including encryption aren't applied to items created with workloads. Learn More
Capacity admins and contributors can add and remove additional workloads	Capacity admins or individuals granted Contributor permission in Capacity settings can add and remove additional workloads in capacities. If this setting is turned off, any existing workloads will stay added and items created with those workloads continue to work normally. When users interact with a workload, their data and access tokens, including name and email, are sent to the publisher. Sensitivity labels and protection settings including encryption aren't applied to items created with workloads. Learn More
Workspace admins can develop partner workloads	Workspace admins can develop partner workloads with a local machine development environment. Turning off this feature will prevent developers from uploading to this workspace. Learn More
Users can see and work with additional workloads not validated by Microsoft	Turn on this setting to allow users to see and work with additional workloads not validated by Microsoft. Make sure that you only add workloads from publishers that you trust to meet your organization's policies. Learn More

Azure AI Service

Setting name	Description
Allow data processing outside geographic region for Azure AI services	This setting applies only to customers using prebuilt AI services in Fabric whose capacity is in a region not covered by Azure AI services. Learn More When enabled, data sent to Azure AI services may be processed outside your capacity's geographic region, compliance boundary, or national cloud instance.

Setting name	Description
	Check the documentation for details about what data types are affected. You can manage this setting at both tenant and capacity levels. Learn More

Related content

- What is the admin portal?
- About tenant settings
- Use the Fabric REST API to list tenant settings

Last updated on 08/28/2026

Service interruption notifications

It's important to have insight into the availability of your mission-critical business applications. Fabric incident notification so you can optionally receive emails if there's a service disruption or degradation.

At this time, emails are sent for the following *reliability scenarios*:

- Open report reliability
- Model refresh reliability
- Query refresh reliability

Notifications are sent when there's an *extended delay* in operations like opening reports, semantic model refresh, or query executions. After an incident is resolved, you receive a follow-up email.

Enable notifications for service outages or incidents

A Fabric admin can enable notifications for service outages or incidents in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**:

1. Identify or create an email-enabled security group that should receive notifications.
2. Go to **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. Under **Help and support settings**, expand **Receive email notifications for service outages or incidents**.
3. Enable notifications, enter a security group, and select **Apply**.

Service health in Microsoft 365

This article describes how to receive service notifications through Fabric. You can also monitor Power BI service health through Microsoft 365. Opt in to receive email notifications about service health from Microsoft 365. Learn more in How to check Microsoft 365 service health.

Last updated on 09/29/2026

Diagnostic query text storage

Article • 05/29/2024

Diagnostic query text storage is a Fabric feature that stores query text content for a limited period to help Fabric support teams to resolve issues that might arise from the use of some Fabric items. The feature is on by default. This article describes feature and how to disable it, if necessary or desired. Its target audience is Fabric administrators and others (for example, security and compliance teams) who want to understand the feature and its implications for data privacy and security.

Overview

To improve support and provide more effective troubleshooting, Microsoft might store the query text generated when users use Fabric items such as reports and dashboards. This data is sometimes necessary for debugging and resolving complex issues related to the performance and functionality of Fabric items such as semantic models.

What is stored query text

Query text refers to the text of the queries/commands (for example, DAX, MDX, TMSL, XMLA, etc.) that Fabric executes when users use Fabric items such as reports and dashboards, as well as external applications such as Excel, SQL Server Management Studio, etc. This information helps Fabric support teams understand the context and specific details of issues that arise, and facilitates quicker, more precise resolution.

Privacy and security

In compliance with Microsoft's stringent data protection standards, stored query text is securely handled within Fabric and retained for a limited period of time (less than 30 days). The data is used for approved investigations only. To prevent unauthorized use, access is strictly controlled and monitored. Access can also be protected by Customer Lockbox for Microsoft Azure if that feature is enabled.

Disabling diagnostic query text storage

Diagnostic query text storage is controlled by the tenant setting **Microsoft can store query text to aid in support investigation**. The feature is on by default. This means that

unless a Fabric admin changes the setting, Microsoft stores the query text associated with the use of some Fabric items in the organization.

If there are organizational requirements that don't permit the storage of query text, or if you wish to opt out of this feature for any other reason, you can turn the feature off. Go to the tenant settings, find the setting, and set the toggle to **Disabled**. See Microsoft can store query text to aid in support investigation for more information.

> **Note**
>
> The availability of stored query text enables high-quality support without any additional configuration on the part of Fabric administrators. It is thus not recommended to disable diagnostic query text storage unless there is a specific reason for doing so.

Implications of turning off diagnostic query text storage

Turning off diagnostic query text storage stops the storage of new query text data. This might adversely affect the ability of Fabric support teams to provide swift, efficient support should issues arise with your organization's reports or dashboards. Without access to the recent history of query text content, diagnosing and resolving problems might take longer, and it might not be possible to identify a root cause for some complex issues.

Further support

Contact Fabric support for further assistance regarding the query text storage feature.

Feedback

Was this page helpful? Yes No

Provide product feedback | Ask the community

Configure notifications

Power BI Premium allows you to configure email notifications for your capacity. The emails are sent to the people you specify in the notifications settings.

To calculate when to send emails, Power BI checks the capacity every 15 minutes. During the check, the last 15 to 30 minutes of capacity activity are examined.

> **Note**
>
> To configure email notifications, you must be a capacity administrator.

Configure capacity notifications

To configure the capacity notification emails, follow these steps:

1. In Fabric, select **OneLake catalog** > **Govern** > **Capacities**.

2. Next to the capacity name, select the **More options** menu > **Settings**.

3. Expand the **Throttling notifications** section.

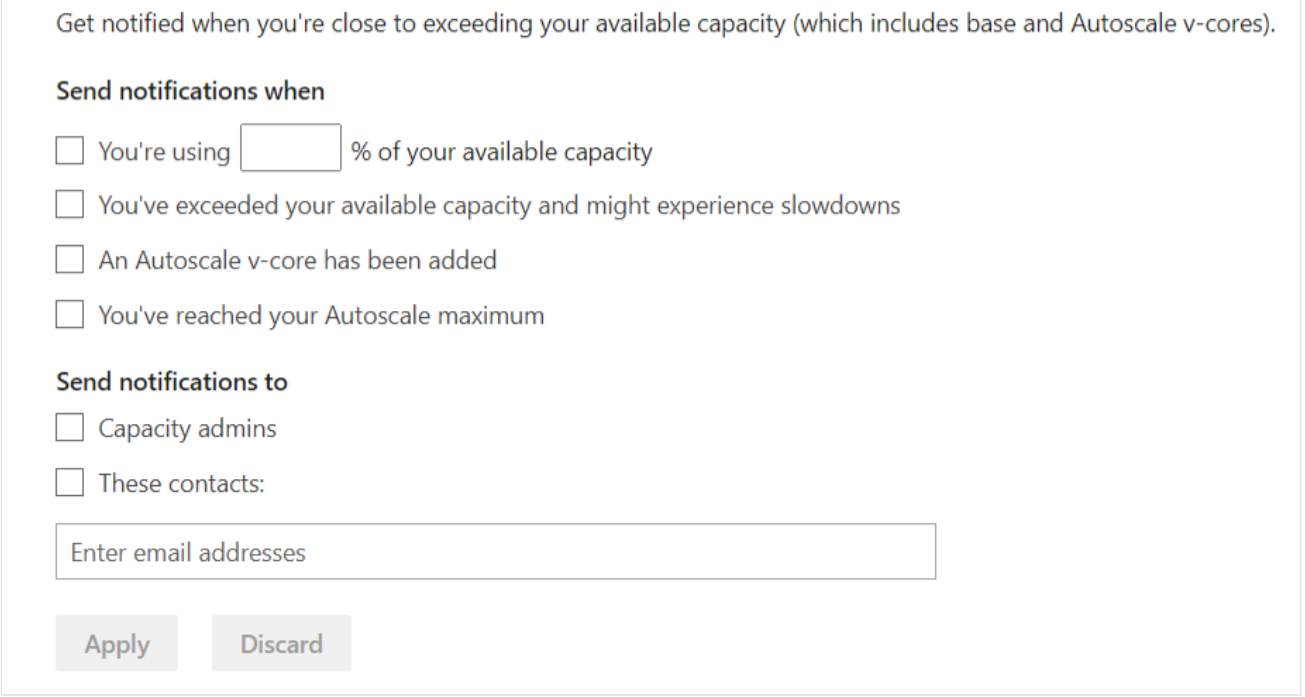

4. In the section **Send notifications when**, configure your required notifications as follows:

- **You're using ___% of your available capacity** - A notification is sent after the capacity reaches the threshold you enter.
- **You've exceeded your available capacity and might experience slowdowns** - A notification is sent after you reach your capacity limit. After the limit is reached, if you have autoscale enabled, autoscale starts. If you don't have autoscale enabled, throttling is applied to your capacity.
- **An Autoscale v-core has been added** - A notification is sent after autoscale starts and every time a v-core is added.
- **You've reached your Autoscale maximum** - A notification is sent when all the autoscale v-cores are fully utilized. Throttling is applied to your capacity if it continues to be overloaded.

5. In the section **Send notifications to**, select who you want the notifications to be emailed to:
   - **Capacity admins** - Email notifications are sent to all the admins of this capacity.
   - **These contacts** - Enter the emails of the contacts you want to receive notifications.
6. Select **Apply**.

Considerations and limitations

- Timestamps aren't included in notification emails.
- Notification emails don't list by how much a threshold was crossed.
- After a notification is sent, there's a three hour period in which new notifications won't be sent, even if your capacity crosses thresholds that are set to trigger these notifications. For example, if you configure your capacity to send a notification after you cross the 75% usage threshold, after that threshold is met you'll receive a notification. If the capacity goes below this threshold to 60%, and then right back over it in the next hour, you won't get another notification for crossing the 75% mark. If you have the autoscale notification turned on, and your capacity crosses the 100% threshold during these three hours, you get a notification that autoscale started.
- A 30-seconds window is applied to calculate your capacity usage. Due to a less granular calculation, capacity usage might appear differently in the Power BI Premium utilization and metrics app. As a result, you might not see the event your notification points to in the app.

For example, a short spike in capacity activity that triggers a notification might not be seen at all in the Power BI Premium utilization and metrics app.

Related content

- What is Power BI Premium?
- What is Fabric admin?

Last updated on 09/29/2026

Add custom branding to the Power BI service

As a Fabric admin, you can change the look and feel of the Power BI service to match your organization's own branding. With custom branding, you can change the theme color that appears in the top navigation bar, add your company logo, and bring your default landing page to life by adding a cover image.

Custom branding changes the look of Power BI for your whole organization. Users can't override your custom branding with their own theme. Custom branding also appears to any external users who have access to your reports in B2B scenarios, helping to easily distinguish your organization.

Before you begin

- Make sure you're a Fabric administrator.
- Prepare your images for upload. You need these files:
  - A logo file that's saved in *.png* format, is 10 KB or smaller, and is at least 200 x 30 pixels. Choosing a PNG file makes sure your logo has a high-resolution appearance on all screens and at all zoom levels. The logo appears on every page.
  - A cover image that's saved in *.jpg* or *.png* format, is 1 MB or smaller, and is at least 1920 x 160 pixels. Get creative with your choice with an image that complements your theme color and feels welcoming. The cover image appears only at the top of Home.
- Identify the hex or decimal code for your theme color. Your theme color appears on every page and provides the background for your logo. Choose a color that complements your logo and cover image or that matches other custom branding in your organization.

The following image indicates where each of these elements appears in the Power BI service:

1. Logo
2. Cover image
3. Theme color

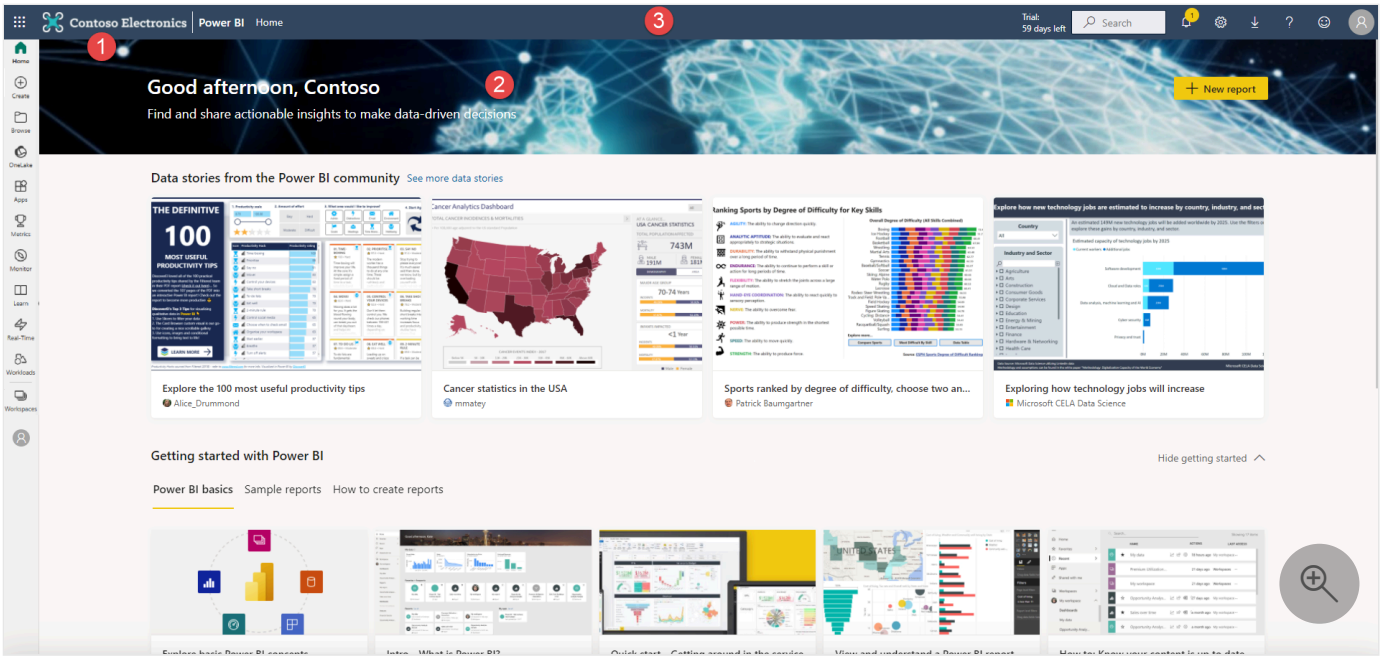

Add custom branding

Follow these steps to customize the look of Power BI for your whole organization:

1. Sign in to the Power BI service as a Fabric admin.

2. Select **OneLake catalog** > **Govern** > **Configurations** > **Custom branding**.

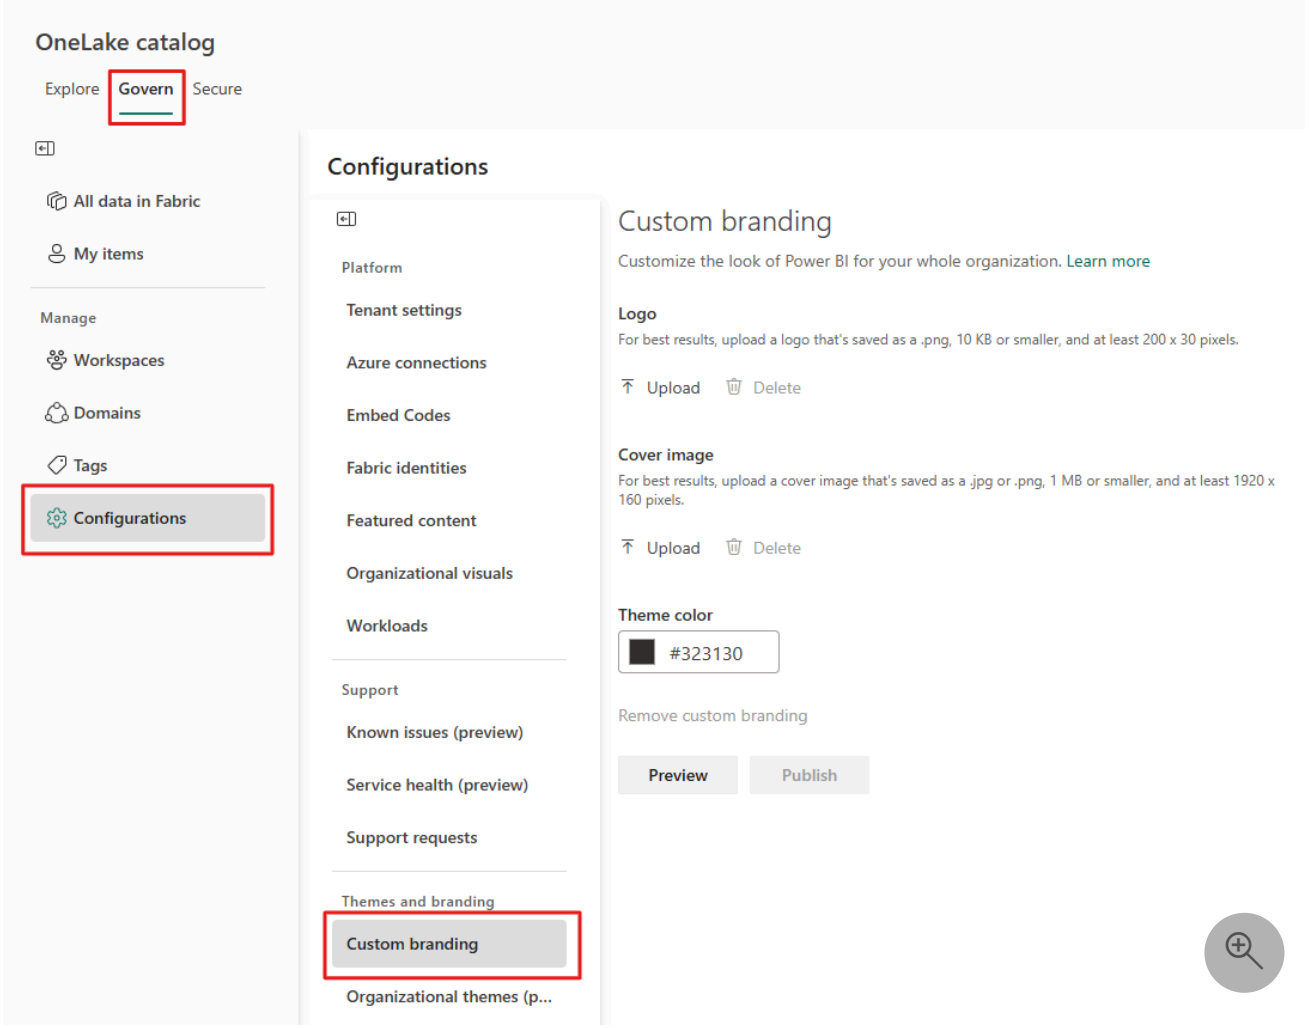

3. Upload a logo file.

4. Upload a cover image file, then crop as needed to adjust how the image appears on the page.

5. Select your theme color by using the color picker or by typing the hex or decimal code.

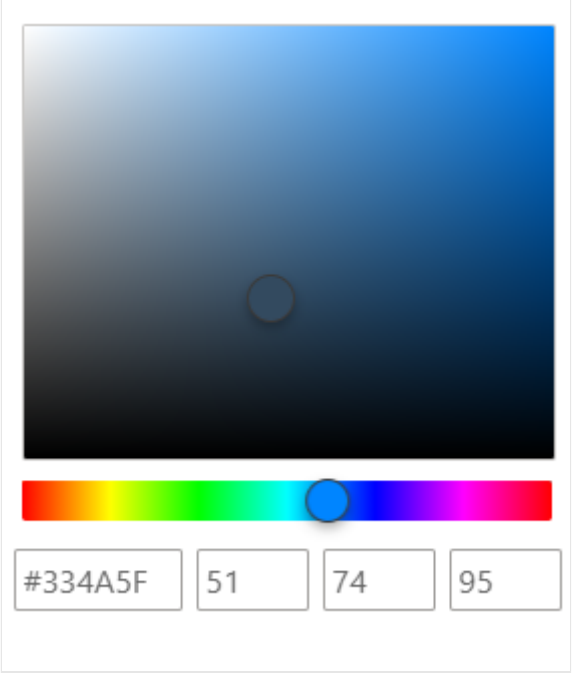

6. Select **Preview** to see how your custom branding looks before you publish.

7. When you're happy with your settings, select **Publish** to make the custom branding the default appearance for all users in your organization. The custom branding appears when you refresh your browser window.

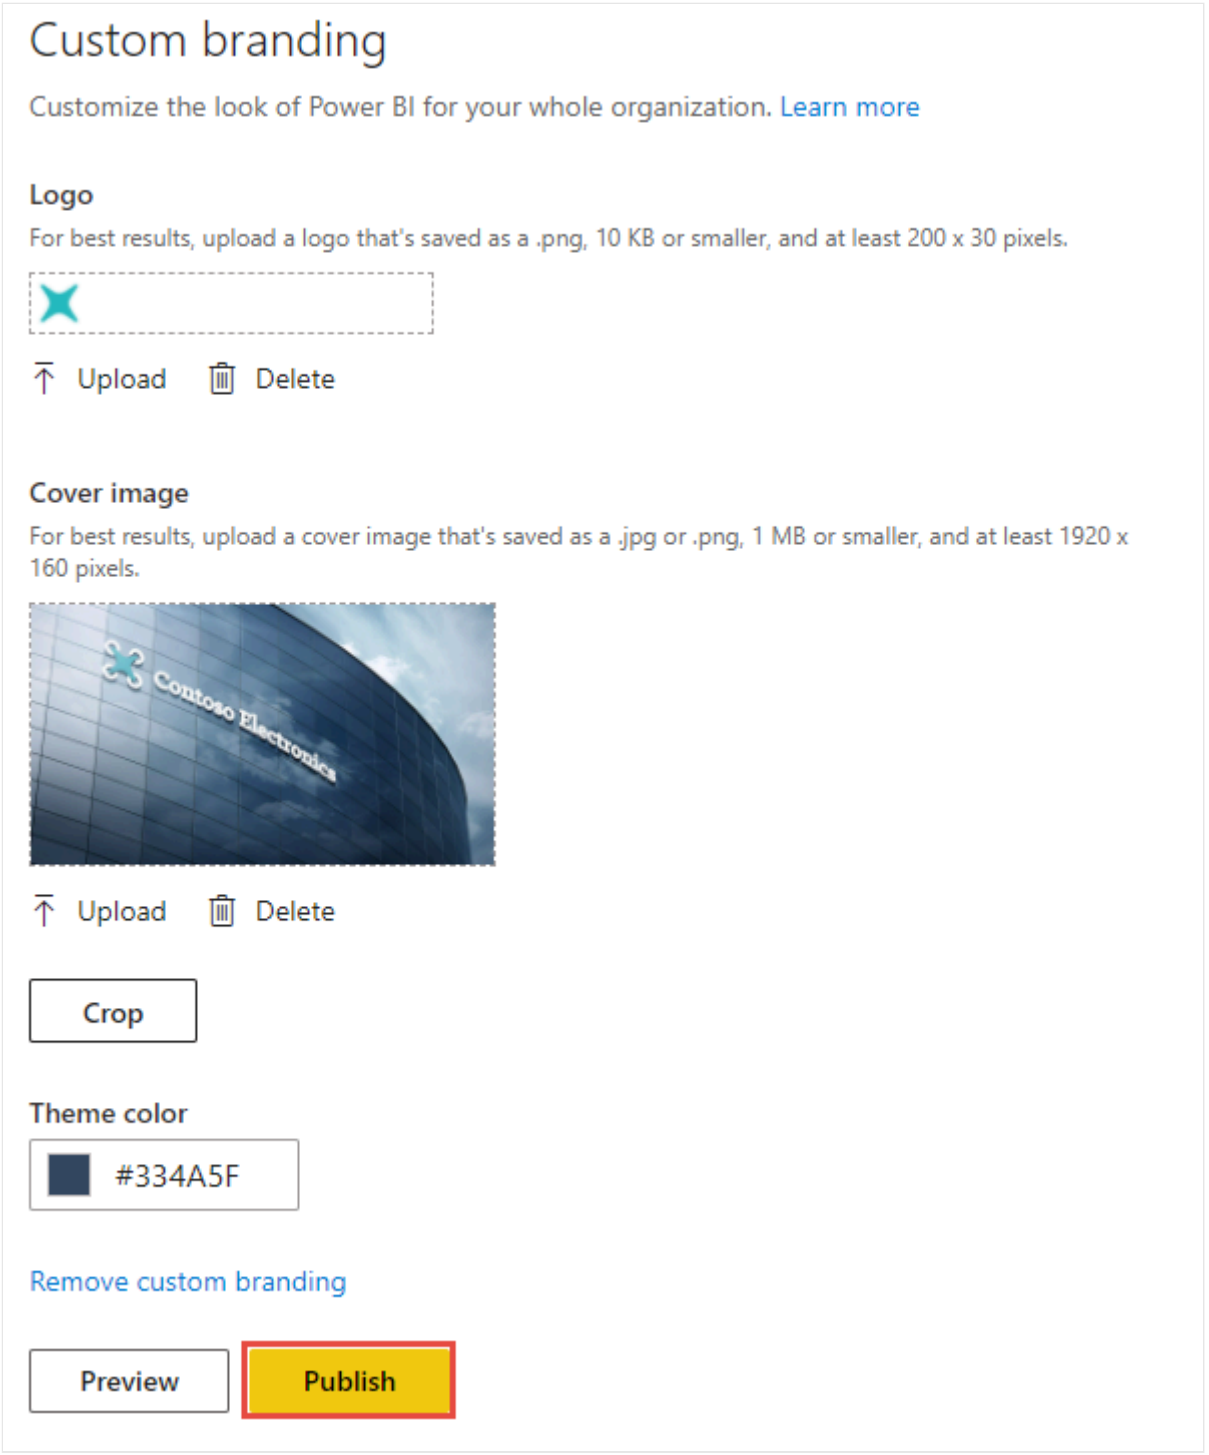

Remove custom branding

Follow these steps to return the look of Power BI to the default settings:

1. Sign in to the Power BI service as a Fabric administrator.
2. Select **OneLake catalog** > **Govern** > **Branding settings**.

3. Select **Remove custom branding**, then select **Publish** to go back to the Power BI default look.

Related content

Give your users a consistent online experience by applying custom branding to other services. Custom branding settings aren't shared between Microsoft 365 and Power BI, but your users will see branding that you apply to your organization's Microsoft Entra sign-in page.

- Add branding to your organization's Microsoft Entra sign-in page
- Customize the Microsoft 365 theme for your organization
- Add featured content to Power BI Home

Last updated on 09/29/2026

Configure Multi-Geo support for Fabric

Multi-Geo is a Fabric feature that helps multinational customers address regional, industry-specific, or organizational data residency requirements. As a Fabric customer, you can deploy content to data centers in regions other than the home region of the Fabric tenant. A geo (geography) can contain more than one region. For example, the United States is a geo, and West Central US and South Central US are regions in the United States. You might choose to deploy content to any of the following geographies (geos) defined in the Azure geography map.

- Sovereign clouds support Multi-Geo across regions within that cloud.
- China North doesn't support Multi-Geo.
- Power BI Embedded supports Multi-Geo.
- Power BI Premium Per-User (PPU) isn't supported for Multi-Geo.

Enable and configure

Enable Multi-Geo by selecting a region other than the default region when you're creating a new capacity. Once a capacity's created, it shows the region where it's currently located.

After you create a capacity, it remains in that region, and any workspaces created under it will have their content stored in that region.

Follow these steps to change the default capacity region when you're creating a new capacity.

Power BI Premium

1. In Fabric, select **settings** (⚙) and from the menu select **Admin portal**.
2. In the *Admin portal*, select **Capacity settings**.
3. Select **Set up new capacity**.
4. From the **Region** dropdown menu, select the region you want to use for this capacity.

Move workspaces between capacities

Follow the steps below to move workspaces from one capacity to another in the same region. During migration, certain operations might fail, such as publishing new semantic models or scheduled data refresh.

> **Important**
>
> Before you start a migration, both the source and destination capacities must be active. Migration isn't supported if either capacity is paused or deleted when the migration starts. Don't pause or delete either capacity before or during a migration. If a capacity is paused or deleted at any point before the migration finishes, the migrated workspace can have missing items or fail to migrate completely. If this condition occurs, resume the capacities as needed and run the migration again.

1. Open the workspace settings.
2. From the side bar, select **Workspace type**.
3. From the **License capacity** dropdown menu, select the capacity you want to move the workspace to.

Change the region of your existing content

To change the region for existing content, do one of the following:

- Create a new capacity and move your workspaces to the new capacities. Free users won't experience any downtime as long as the tenant has spare v-cores.
- Temporarily move your content to a shared capacity. You don't need extra v-cores, but free users will experience some downtime. After you create a new capacity in the desired region, move your workspaces to the new capacity.

Move content to your home region

To move workspaces to your home region, do one of the following:

- Delete the current capacity where the workspace is located. Workspaces in the deleted capacity are moved to a shared capacity in the home region.
- Move individual workspaces to a capacity located in the home tenant.

Large-storage format semantic models shouldn't be moved from the region where they were created. Reports based on a large-format semantic model won't be able to load the semantic model and will return a *Cannot load model* error. Move the large-storage format semantic model back to its original region to make it available again.

Considerations and limitations

- Confirm that any movement you initiate between regions follows all corporate and government compliance requirements prior to initiating data transfer.
- When you're using Multi-Geo, compute and storage (including OneLake and experience-specific storage) is located in the multi-geo region, yet some tenant metadata remains in the home region, including the following:
  - Push datasets
  - Dashboard/report metadata: tile names, tile queries, and any other data
  - Service buses for gateway queries or scheduled refresh jobs
  - Permissions
  - Semantic model credentials
  - Power BI Embedded Analytics Playground saved state
  - Metadata linked to Purview Data Map
- Cached data and queries stored in a remote region stays in that region at rest. Additionally, the data at rest is replicated to another region in the same Azure geography for disaster recovery if the Azure geography contains more than one region. Data in transit might go back and forth between multiple geographies.
- When moving data from one region to another, the source data might remain in the region from which the data was moved for up to 30 days. During that time end users don't have access to it. The data is removed from this region and destroyed during the 30-day period.
- Query text and query result traffic for imported and DirectQuery data models doesn't transit through the home region. However, the report metadata comes from the home region, and certain DNS routing states might take such traffic out of the region.
- Certain features such as screenshots, data alerts and others process data in the home region.
- The detailed semantic model metadata that is cached as part of enhanced metadata scanning is always stored in the home region, even if the scanned semantic model is located in a remote region.

- Detailed semantic model metadata lives in the home tenant.
- To use dataflows gen1 on Multi-Geo you must configure dataflow storage to use Azure Data Lake Storage (ADLS) Gen2.
- It's possible to create and maintain large-storage format semantic models in remote regions to meet data residency requirements. However, you can't move storage format semantic models to another region. Moving large-storage format semantic models from the region where they were created results in reports failing to load the semantic model. Move the large-storage semantic model back to its original region to make it available. If you must move such a model, deploy it as if it was a new model, and then delete the old model from the undesired region.
- Multi-Geo doesn't support Metrics in Power BI.
- Multi-Geo doesn't support Power BI Q&A Settings as detailed semantic model metadata lives in the home tenant.
- Workspaces with non-Power BI Fabric items can't be moved between regions. You must delete all the non-Power BI Fabric items before moving a workspace to a different region. Once the workspace is moved, it can take up to 30 minutes before non-Power BI items can be created.

For more details, see Moving data around.

Related content

- What is Power BI Premium?
- Multi-Geo support for Power BI Embedded
- Moving data around

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 08/14/2026

Manage Azure connections

The Azure connections admin settings connect Azure services to Fabric. Using these settings, you can store your dataflows in your organization's Azure Data Lake Storage Gen2 (ADLS Gen2) account. You can review the benefits of this approach in Reasons to use the ADLS Gen 2 workspace or tenant connection. Workspace-level Log Analytics enables administrators and workspace owners to configure a Log Analytics connection to collect usage and performance logs for the workspace. You can review how the integration between Log Analytics and Power BI works.

The Azure connections admin settings have the following options:

- Tenant-level storage - Use to store dataflows in your organizations tenant settings. This setting can be configured if you want a central data lake storage place, or as a default storage place in addition to workspace level storage.
- Workspace-level storage permissions - Use to store dataflows in specific ADLS Gen 2 accounts, organized per workspace.
- Workspace-level log analytics permissions - Use to configure activity logging in Log Analytics.

To access these settings, open the Govern section of the OneLake catalog, and then select **Azure connections**.

Tenant-level storage

By default, data used with Power BI is stored in internal storage provided by Power BI. With the integration of dataflows and Azure Data Lake Storage Gen2 (ADLS Gen2), you can store your dataflows in your organization's Azure Data Lake Storage Gen2 account. Storing dataflows in Azure Data Lake allows you to access them using the Azure portal, Azure Storage Explorer, and Azure APIs. For more information, see Configuring dataflow storage to use Azure Data Lake Gen 2.

Workspace-level storage permissions

By default, workspace admins can't connect their own storage account. This feature lets Fabric administrators turn on a setting that allows workspace admins to connect their own storage

account.

To activate this feature, go to **OneLake catalog** > **Govern** > **Azure connections** > **Connect to Azure resources** > **Workspace-level storage permissions**, and check the **Allow workspace admins to connect their own storage account** checkbox.

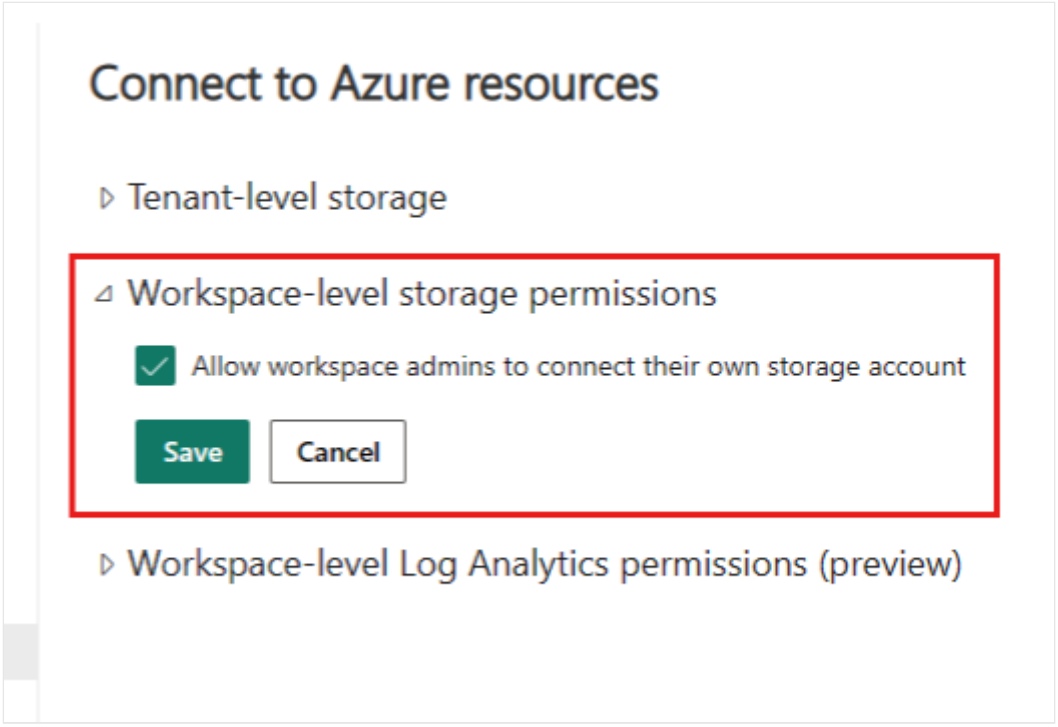

Workspace-level log analytics permissions

Fabric administrators can find and configure activity logging in **Tenant settings** > **Audit and usage settings** > **Azure Log Analytics connections for workspace administrators**. For more information, see Configure workspace-level logging.

Related content

- Govern and manage your Fabric data with the OneLake catalog
- Configuring dataflow storage to use Azure Data Lake Gen 2

Last updated on 09/29/2026

Manage Fabric identities

As a Fabric administrator, you can manage your organization's Fabric identities from **Fabric identities** in the Govern section of the OneLake catalog.

In the **OneLake catalog**, select **Govern** > **Configurations** > **Fabric identities**. The **Fabric identities** page lists all the Fabric identities in your tenant.

The columns of the list of identities are described in following table.

Expand table

Column	Description
Name	The name of the identity.
Service principal ID	The object ID of the Enterprise application that is associated with the identity in Microsoft Entra.
State	The state of the identity. See workspace identity state values.
Workspace	The workspace ID.

View identity details

1. Select the identity that you want to view.
2. Select **Details** on the command bar. The **Details** side pane opens and displays the identity's details.

Expand table

Field	Description
Workspace name	The name of the workspace the identity is associated with.
State	The state of the identity.
State changed date	The date of the last change of state of the identity.
Service principal ID	The object ID of the Enterprise application that is associated with the identity in Microsoft Entra.

Field	Description
Application ID	The application ID of the Enterprise application that is associated with the identity in Microsoft Entra.
Tenant ID	The ID of the tenant the identity is defined in.
Role	The workspace role the identity has been assigned.

Delete an identity

> ⊗ **Caution**
>
> Deleting a workspace identity breaks any Fabric item relying on that identity for trusted workspace access or authentication. Deleted identities can't be restored.

To delete an identity:

1. Select the identity you want to delete.
2. Select **Delete** on the command bar.

Refresh the identities list

Select **Refresh** on the command bar to refresh the list of identities.

Export the identities list as a .csv file

Select **Export** on the command bar to download the list of identities as a *.csv* file.

Related content

- Workspace identity

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Manage Power BI Premium Per-User (PPU)

Summarize this article for me

Power BI Premium Per-User (PPU) is a way to license Power BI Premium features on a per user basis. After the first user is assigned a Power BI Premium Per-User (PPU) license, associated features can be turned on in any workspace. Admins can manage the auto refresh and semantic model workload settings that are shown to users and their default values. For example, access to the XMLA endpoint can be turned off, set to read only, or set to read and write.

PPU settings

You can configure the following PPU settings in the admin portal on the **Power BI Premium Per-User (PPU)** tab. To learn how to access the Fabric admin portal settings, see What is the admin portal?

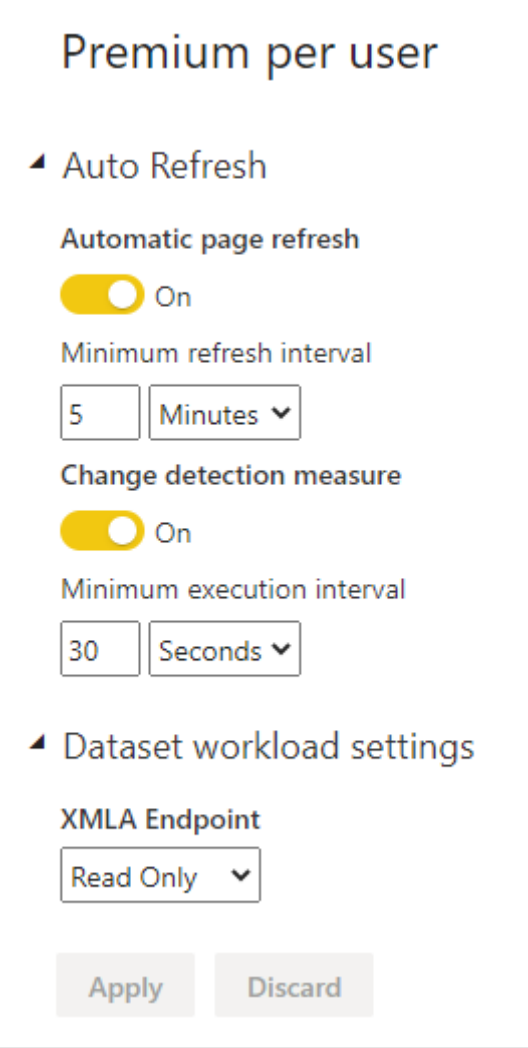

Auto refresh

Automatic refresh enables your active report page to query for new data, during predefined intervals. By default, these settings are turned on. If you turn them off, PPU reports that use automatic refresh and change detection don't get updated automatically.

Use the following settings to override the *automatic refresh* settings in individual reports that reside on the PPU capacity. For example, when the *minimum refresh interval* setting is configured to refresh every 30 minutes, if you have a report that's set to refresh every five minutes, its setting will be overridden and the report is refreshed every 30 minutes instead.

- **Minimum refresh interval** - Use to specify a minimum value for the automatic refresh for all the reports in the PPU capacity. The Power BI service overrides any automatic refresh settings that are higher than this setting.
- **Change detection measure** - Use to specify a minimum value for all the reports in the PPU capacity that use change detection. The Power BI service overrides any change detection settings that are higher than this setting.

Semantic model workload settings

XMLA endpoints allow Microsoft and third-party apps and tools to connect to Power BI semantic models. Use this setting to determine if in the PPU capacity XMLA endpoints are turned off, or configured for read only or read and write.

Related content

- What is the admin portal?
- Power BI Premium Per-User (PPU) FAQ
- Automatic page refresh in Power BI

Last updated on 02/12/2026

Manage users

Article • 11/20/2024

You manage Power BI users, groups, and admins in the Microsoft 365 admin center. The **Users** tab in the Fabric admin portal provides a link to the admin center.

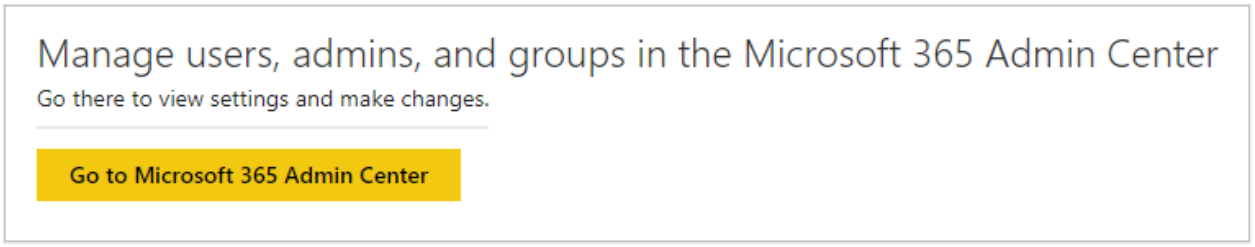

Related content

- Administration overview
- Microsoft 365 admin center help
- Give users access to workspaces

Feedback

Was this page helpful? Yes No

Provide product feedback | Ask the community

Enable service principal authentication for admin APIs

This article shows how to enable service principal authentication for Power BI *read-only* admin APIs and Fabric *update* admin APIs.

Service principal is an authentication method that you can use to let a Microsoft Entra application access Fabric content and APIs. When you create a Microsoft Entra app, the process creates a service principal object. The service principal object, also known simply as the service principal, enables Microsoft Entra ID to authenticate your app. Once authenticated, the app can access Microsoft Entra tenant resources.

An example is the app that runs metadata scanning and Fabric data risk assessments in Microsoft Purview Data Security Posture Management (DSPM) for AI. This app uses service principal authentication to access Fabric admin APIs. You configure this app to authenticate with federated credentials (recommended) or a client secret.

Enable service principal authentication

To enable service principal authentication for Fabric APIs, follow these steps:

1. Create a Microsoft Entra app. You can skip this step if you already have a Microsoft Entra app you want to use. Take note of the app ID, which you need in later steps.

   > **Important**
   >
   > Make sure the app you use doesn't have any admin-consent required permissions for Fabric set on it in the Azure portal. **See how to check whether your app has any such permissions**.

2. Create a new Microsoft Entra Security Group and make sure to select **Security** as the Group type. You can skip this step if you already have a Microsoft Entra security group you'd like to use.

3. Add your app ID as a member of the security group you created. To do so:

4. Navigate to **Azure portal** > **Microsoft Entra ID** > **Groups**, and choose the security group you created in *Step 2*.
   a. Select **Add Members**.

5. Enable the Fabric admin settings:

   a. Sign in to Fabric using your admin account credentials. Open the **OneLake catalog**, and then select the **Govern** tab > **Configurations** > **Tenant settings**.

   b. Under **Admin API settings**, select the switch for the type of admin APIs you want to enable:

      - **Service principals can access read-only admin APIs** (see supported Power BI admin APIs)
      - **Service principals can access admin APIs used for updates** (see supported Fabric admin APIs)

   > **Note**
   >
   > These settings are required to run Fabric data risk assessments in Microsoft Purview DSPM for AI. For more information, see DSPM for AI considerations.

6. Set the toggle to **Enabled**.

7. Select the **Specific security groups** radio button. In the text field that appears below it, add the security group you created in *Step 2*.

8. Select **Apply**.

Supported Power BI admin APIs for read-only

The following read-only admin APIs support service principal authentication. This list may not be exhaustive; for the latest information about APIs not listed here, refer to the Power BI REST API documentation.

- Apps GetAppsAsAdmin
- Apps GetAppUsersAsAdmin
- Dashboards GetDashboardsAsAdmin
- Dashboards GetDashboardsInGroupAsAdmin
- Dashboards GetDashboardSubscriptionsAsAdmin

- Dashboards GetDashboardUsersAsAdmin
- Dashboards GetTilesAsAdmin
- Dataflows ExportDataflowAsAdmin
- Dataflows GetDataflowDatasourcesAsAdmin
- Dataflows GetDataflowsAsAdmin
- Dataflows GetDataflowsInGroupAsAdmin
- Dataflows GetDataflowUsersAsAdmin
- Dataflows GetUpstreamDataflowsInGroupAsAdmin
- Datasets GetDatasetsAsAdmin
- Datasets GetDatasetsInGroupAsAdmin
- Datasets GetDatasetToDataflowsLinksInGroupAsAdmin
- Datasets GetDatasetUsersAsAdmin
- Datasets GetDatasourcesAsAdmin
- Get Activity Events
- Get Capacities As Admin
- Get Power BI Encryption Keys
- Get Refreshable For Capacity
- Get Refreshables For Capacity
- Get Refreshables
- Groups GetGroupAsAdmin
- Groups GetGroupsAsAdmin
- Groups GetGroupUsersAsAdmin
- Groups GetUnusedArtifactsAsAdmin
- Imports GetImportsAsAdmin
- Pipelines GetPipelinesAsAdmin
- Pipelines GetPipelineUsersAsAdmin
- Profiles GetProfilesAsAdmin
- Reports GetReportsAsAdmin
- Reports GetReportsInGroupAsAdmin
- Reports GetReportSubscriptionsAsAdmin
- Reports GetReportUsersAsAdmin
- Users GetUserArtifactAccessAsAdmin
- Users GetUserSubscriptionsAsAdmin
- WidelySharedArtifacts LinksSharedToWholeOrganization
- WidelySharedArtifacts PublishedToWeb
- WorkspaceInfo GetModifiedWorkspaces
- WorkspaceInfo GetScanResult

- WorkspaceInfo GetScanStatus
- WorkspaceInfo PostWorkspaceInfo

How to check if your app has admin-consent required permissions

An app using service principal authentication that calls read-only admin APIs **must not** have any admin-consent required permissions for Power BI set on it in the Azure portal. To check the assigned permissions:

1. Sign in to the **Azure portal**.
2. Select **Microsoft Entra ID**, then **Enterprise applications**.
3. Select the application you want to grant access to Power BI.
4. Select **Permissions**. There must be no admin-consent required permissions of type Application registered for the app.

Supported Fabric admin APIs for updates

The **Service principals can access admin APIs used for updates** setting applies to Fabric admin APIs, such as the Workspaces - Restore Workspace API.

To find out if a specific Fabric admin API supports service principal authentication, check the API's documentation in the Fabric REST API reference. Look for the "Microsoft Entra supported identities" section, which indicates whether service principal authentication is supported.

Considerations and limitations

- The service principal can make REST API calls, but you can't open Fabric with service principal credentials.
- Fabric admin rights are required to enable service principal in the Admin API settings in the Fabric.

Related content

- Fabric REST API reference
- DSPM for AI considerations
- Metadata scanning overview

- Set up metadata scanning
- Run metadata scanning

Last updated on 09/29/2026

Manage featured content

Article • 11/21/2023

If the *featured content* feature is enabled in your organization, users can feature content in the **Featured** section of the Power BI Home page. See Feature content on colleagues' Power BI Home page for details.

As a Fabric admin, you can monitor this featured content and remove it from the **Featured** section, if necessary. You can also disable the featured content feature entirely, in which case users will no longer be able to feature content. See the Enable/disable featured content section.

Monitor and manage featured content

In the admin portal, select **Featured content**.

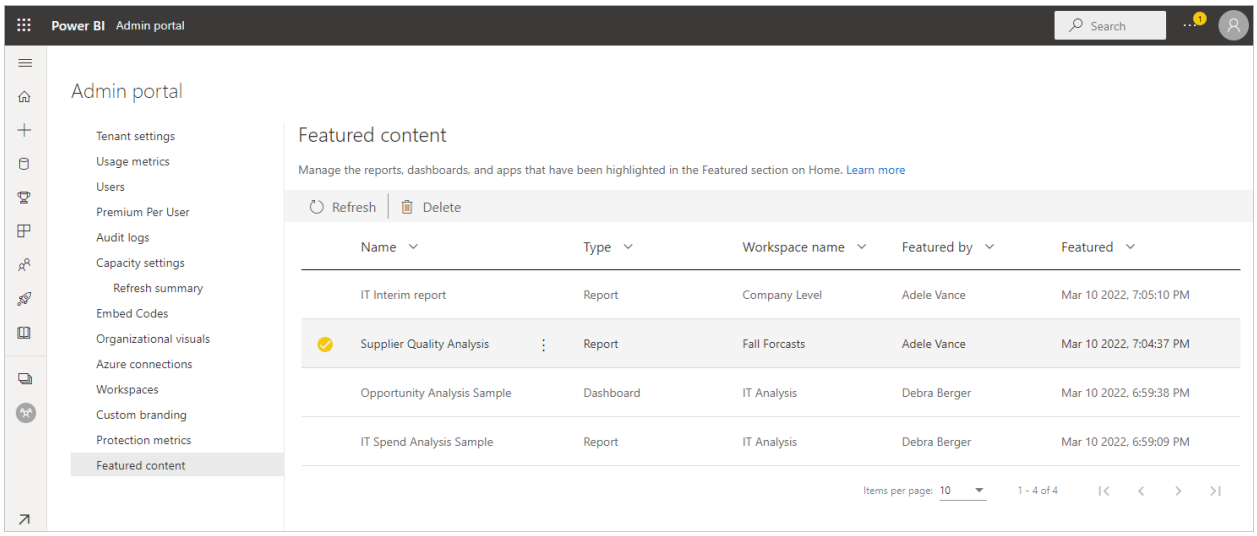

Here you see a list of all featured items along with their relevant metadata. If something looks suspicious, or you want to clean up the **Featured** section, you can delete featured items as needed.

To delete an item, mouse over and select the item, and then click the trash can that appears in the top ribbon, or choose **More options (...)** > **Delete**. It's possible to select multiple items and then delete.

Enable/disable featured content

The featured content feature is enabled, disabled, and configured (for example, specifying who can feature content) via an admin setting. To learn more, see Featured content.

Related content

- Feature content on colleagues' Power BI Home page
- What is the admin portal?

Feedback

Was this page helpful? Yes No

Provide product feedback | Ask the community

Manage organizational visuals

The *organizational visuals* admin setting allows you to manage the list of Power BI visuals available in your organization. For more information and detailed instructions, see Organizational visuals.

Other Power BI visuals admin settings

All the Power BI visuals admin settings, including Power BI visuals tenant settings, are described in Manage Power BI visuals admin settings.

Related content

- Govern and manage your Fabric data with the OneLake catalog
- Manage Power BI visuals admin settings

Last updated on 09/29/2026

Manage Power BI visuals admin settings

As a Fabric administrator for your organization, you can control the type of Power BI visuals that users can access across the organization and limit the actions users can perform.

To manage Power BI visuals, you must be a Fabric administrator. For more information about the Fabric administrator role, see Understand Fabric admin roles.

Power BI visuals tenant settings

To manage the tenant settings for Power BI visuals, sign in to Fabric using your admin account credentials. Open the **OneLake catalog** and go to the **Govern** tab > **Configurations** > **Tenant settings**. Then scroll to **Power BI visuals**.

> **Note**
>
> Manage these settings in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. These settings aren't available in the Power Platform admin center.

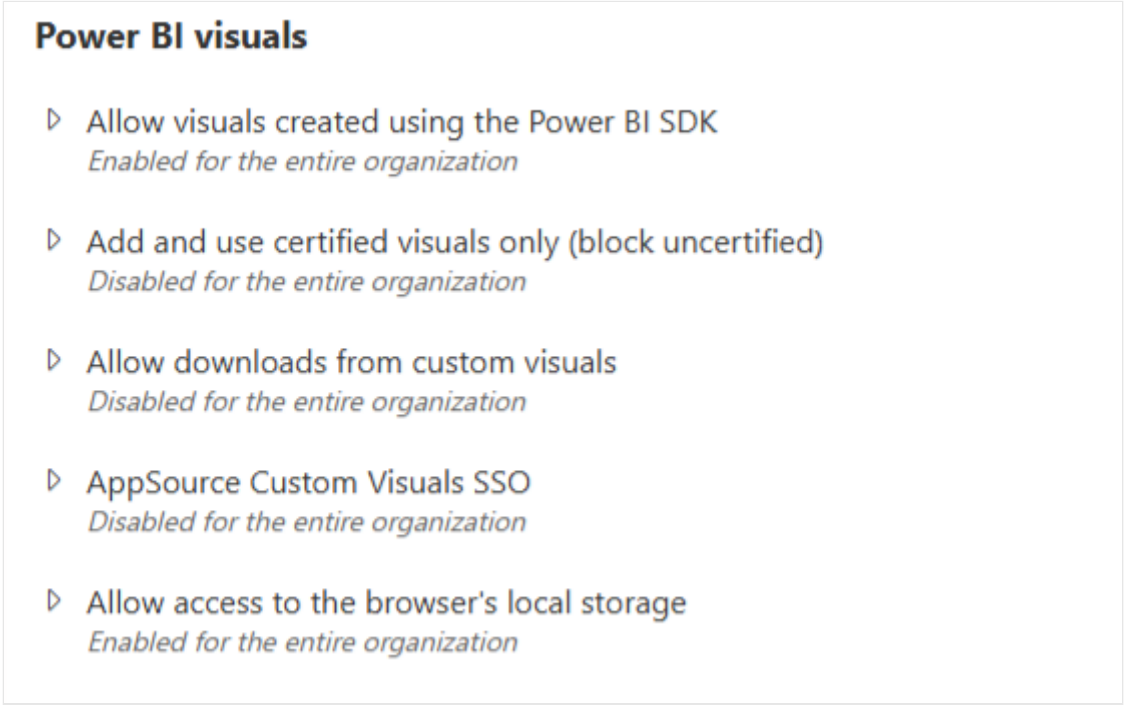

The UI tenant settings only affect the Power BI service. If you want these settings to take effect in Power BI Desktop, use group policies. A table at the end of each section provides details for enabling the setting in Power BI Desktop.

These settings allow you to control the following actions for Power BI visuals in your organization:

- Allow visuals created using the Power BI SDK
- Add and use certified Power BI visuals only
- Allow downloads from custom visuals
- AppSource Custom Visuals SSO
- Allow access to the browser's local storage

Visuals from AppSource or a file

Manage organizational access for the following types of Power BI visuals:

- Custom visuals that developers create by using the Power BI SDK and save as a *.pbiviz* file.
- Visuals that you download from AppSource.

This setting is disabled by default and doesn't affect visuals in the organizational store.

Use the following instructions to enable users in your organization to upload *.pbiviz* files, and add visuals from AppSource to their reports and dashboards:

1. Expand the **Allow visuals created using the Power BI SDK** settings.
2. Select **Enabled**.
3. Choose who can upload *.pbiviz* and AppSource visuals:
    - Select **The entire organization** option to allow everyone in your organization to upload *.pbiviz* files, and add visuals from AppSource.
    - Select the **Specific security groups** option to manage uploading *.pbiviz* files, and adding visuals from AppSource by using security groups. Add the security groups you want to manage to the *Enter security groups* text bar. The security groups you specify are excluded by default. If you want to include these security groups and exclude everyone else in the organization, select the **Except specific security groups** option.
4. Select **Apply**.

UI changes to tenant settings apply only to the Power BI service. To enable users in your organization to upload *.pbiviz* files, and add visuals from AppSource to their visualization pane in Power BI Desktop, use AD Group Policy.

Expand table

Key	Value name	Value
Software\Policies\Microsoft\Power BI Desktop\	EnableCustomVisuals	0 - Disable 1 - Enable (default)

Certified Power BI visuals

Certified Power BI visuals are visuals that meet the Microsoft Power BI team code requirements and testing. The tests check that the visual doesn't access external services or resources. However, Microsoft isn't the author of non-Microsoft visuals, so contact the author directly to verify the functionality of these visuals.

When you enable this setting, only certified Power BI visuals render in your organization's reports and dashboards. Power BI visuals from AppSource or files that aren't certified return an error message. This setting is disabled by default and doesn't apply to visuals in your organizational store.

1. Expand the **Add and use certified visuals only** setting.
2. Select **Enabled**.
3. Select **Apply**.

UI changes to tenant settings apply only to the Power BI service. To manage the certified visuals tenant setting in Power BI Desktop, use AD Group Policy.

Expand table

Key	Value name	Value
Software\Policies\Microsoft\Power BI Desktop\	EnableUncertifiedVisuals	0 - Disable 1 - Enable (default)

Export data to file

When you enable this setting, users can download data from a custom visual into a file on their storage device. This setting is separate from and not affected by download restrictions applied in your organization's export and sharing tenant settings. This setting is disabled by default and applies to all visuals including those visuals managed by your organizational store, Desktop, and web.

> **Note**
>
> When you enable this setting, a custom visual can export to files of the following types:
>
> - .txt
> - .csv
> - .json
> - .tmplt
> - .xml
> - .pdf
> - .xlsx

1. Expand the **Allow downloads from custom visuals** settings.
2. Select **Enabled**.
3. Choose who can download files:
    - Select **The entire organization** option to allow everyone in your organization to download data from a visual into a file.
    - Select the **Specific security groups** option to limit downloading files to specific security groups. Enter the security groups you want in the *Enter security groups* text bar. The security groups you specify are included by default. If you want to exclude these security groups and include everyone else in the organization, select the **Except specific security groups** option.
4. Select **Apply**.

UI changes to tenant settings apply both for Power BI Service and Power BI Desktop. To enable users in your organization to download data from custom visuals in Power BI Desktop, use AD Group Policy.

Expand table

Key	Value name	Value
Software\Policies\Microsoft\Power BI Desktop\	AllowCVToExportDataToFile	0 - Disable 1 - Enable (default)

When `AllowCVToExportDataToFile` is set to *1*, the custom visual can export data to a file only if:

- The feature switch in the tenant settings is enabled.
- The user is signed in.

Local storage

This setting enables visuals to store data on the browser's local storage which helps improve performance. This setting is separate from and not affected by download restrictions applied in your organization's export and sharing tenant settings. The setting is enabled by default and applies to all visuals, including those visuals managed by your organizational store, Desktop, and web.

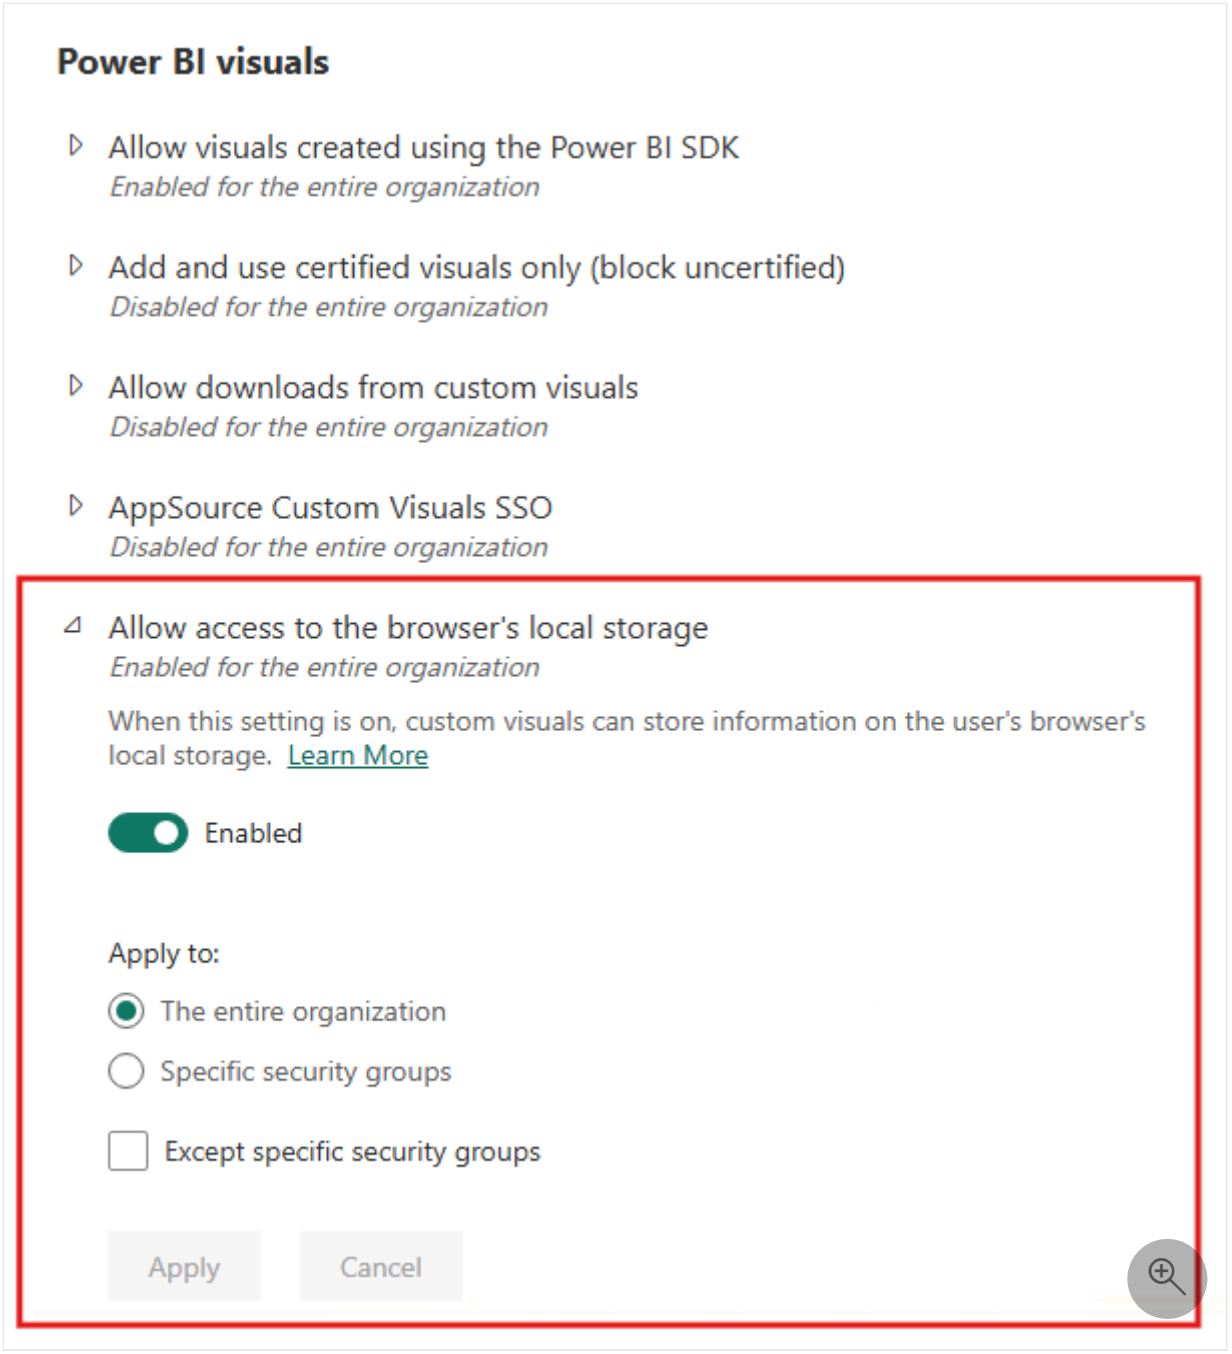

To enable the local storage setting, follow these steps:

1. Expand the **Local storage** settings.

2. Select **Enabled**.

3. Choose who can render this API:

    - Select **The entire organization** option to allow visuals to store data on the local machine for every user in your organization.
    - Select the **Specific security groups** option to limit this privilege to specific security groups. Enter the security groups you want in the *Enter security groups* text bar. The security groups you specify are included by default. If you want to exclude these security groups and include everyone else in the organization, select the **Except specific security groups** option. Only a user listed in the permitted security group can render the API.

4. Select **Apply**.

AppSource Custom Visuals SSO

When you enable this setting, AppSource Custom Visuals can get Microsoft Entra ID (formerly known as Azure Active Directory) access tokens with restricted audiences for signed-in users by using the Authentication API. These tokens include personal information such as the user's name and email address. Custom Visuals can send these tokens across different regions and compliance boundaries, and they're fully responsible for handling the tokens they possess. The setting is disabled by default and applies to all AppSource Custom Visuals, including those visuals managed by your organizational store.

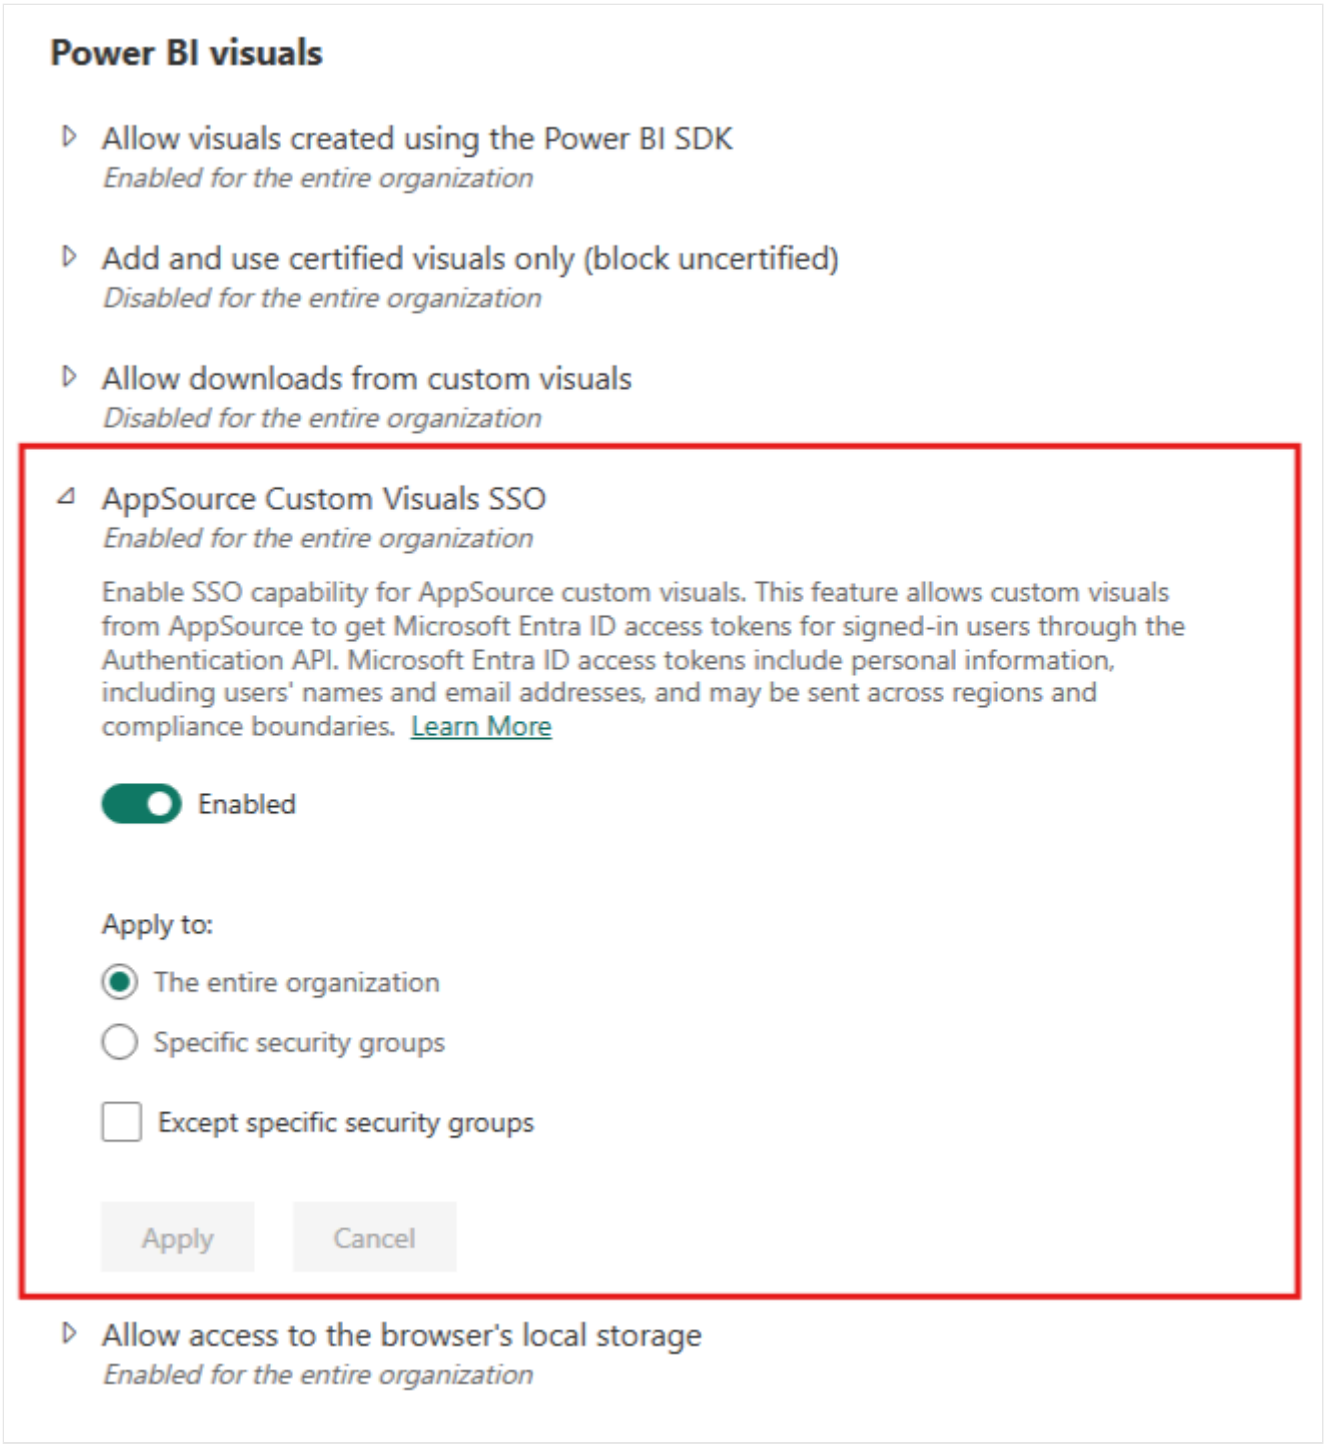

1. Expand the **Allow custom visuals to get user Microsoft Entra access tokens** settings.

2. Select **Enabled**.

3. Choose who can render this API:

    - Select **The entire organization** option to allow visuals to get Microsoft Entra access tokens for every user in your organization.
    - Select the **Specific security groups** option to limit getting access tokens to specific security groups. Enter the security groups you want in the *Enter security groups* text bar. The security groups you specify are included by default. If you want to exclude these security groups and include everyone else in the organization, select the **Except**

specific security groups option. Only a user listed in the permitted security group can render the API.

4. Select **Apply**.

Organizational visuals

As a Fabric admin, you can manage the list of Power BI visuals available in your organization's organizational store. The **Organizational visuals** tab, in the *Admin portal*, allows you to add and remove visuals and decide which visuals automatically display in the visualizations pane of your organization's users. You can add to the list any type of visual including uncertified visuals and *.pbiviz* visuals, even if they contradict the tenant settings of your organization.

Organizational visuals settings automatically deploy to Power BI Desktop.

> ⓘ **Note**
>
> Power BI Report Server doesn't support organizational visuals.

Add a visual from a file

Use this method to add a new Power BI visual from a *.pbiviz* file.

> ⚠ **Warning**
>
> A Power BI visual uploaded from a file could contain code with security or privacy risks. Make sure you trust the author and the source of the visual before deploying to the organization's repository.

1. Select **Add visual** > **From a file**.

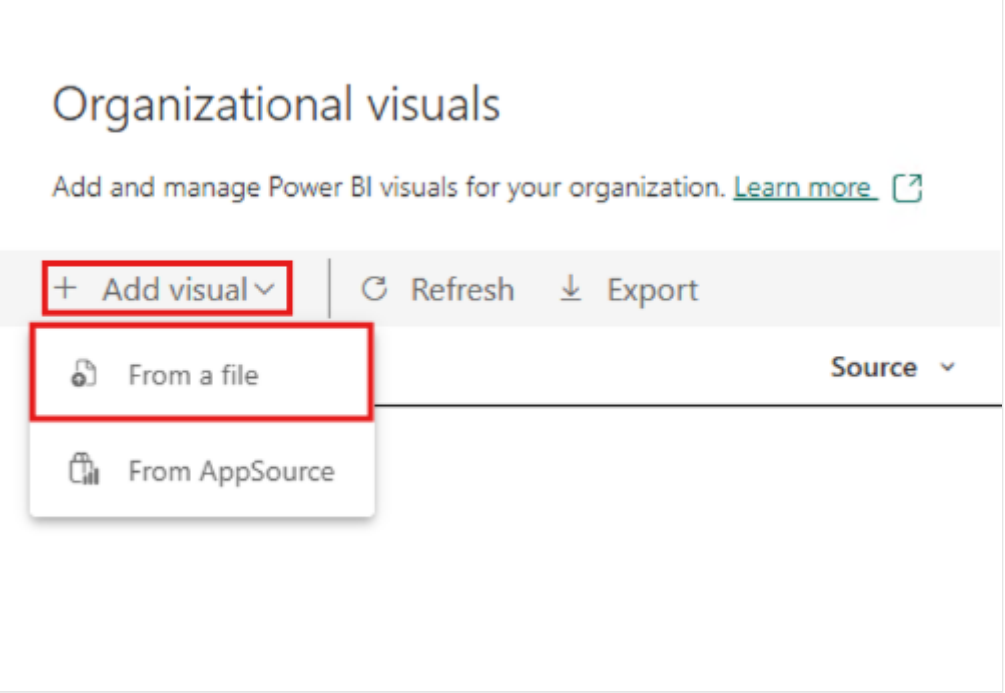

2. Fill in the following fields:

- **Choose a .pbiviz file** - Select a visual file to upload.
- **Name your visual** - Give a short title to the visual, so that report authors can easily understand what it does.
- **Icon** - Upload an icon file to display in the visualization pane.
- **Description** - Provide a short description of the visual to give more context for the user.
- **Access** - Select whether users in your organization can access this visual. This setting is enabled by default.

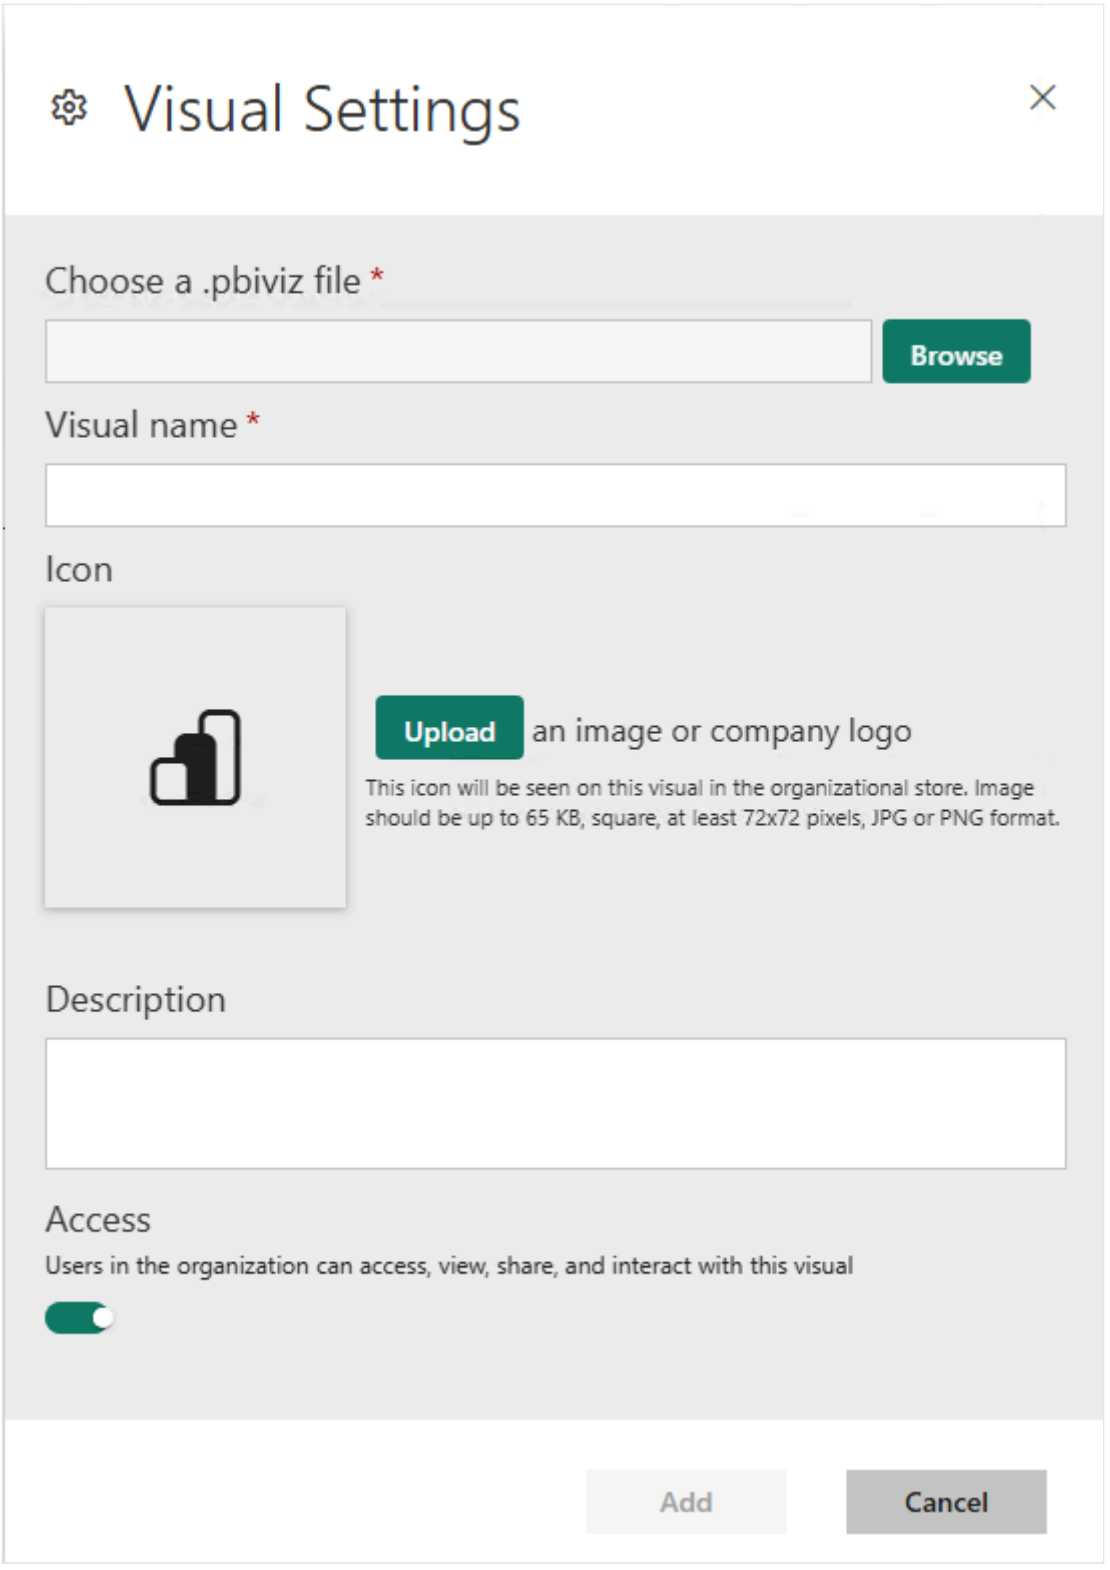

3. Select **Add** to start the upload request. After it uploads, the visual appears in the organizational visuals list.

Add a visual from AppSource

Use this method to add a new Power BI visual from AppSource.

AppSource Power BI visuals automatically update. Users in your organization always have the latest version of the visual.

1. Select **Add visual** > **From AppSource**.

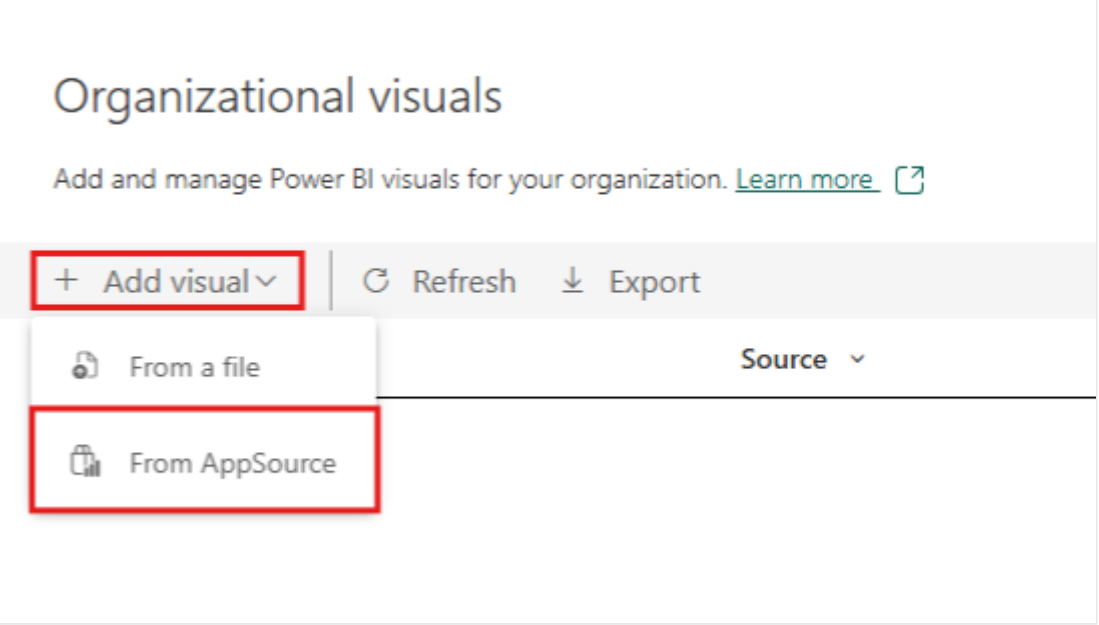

2. In the **Power BI visuals** window, find the AppSource visual you want to add, and select **Add**. After it uploads, the visual displays in the organizational visuals list.

Add a visual to the visualizations pane

You can pick visuals from the organizational visuals page to automatically show on the visualizations pane of all the users in your organization.

1. Select the visual you want to add.

2. Select **Enable for Visualization Pane**.

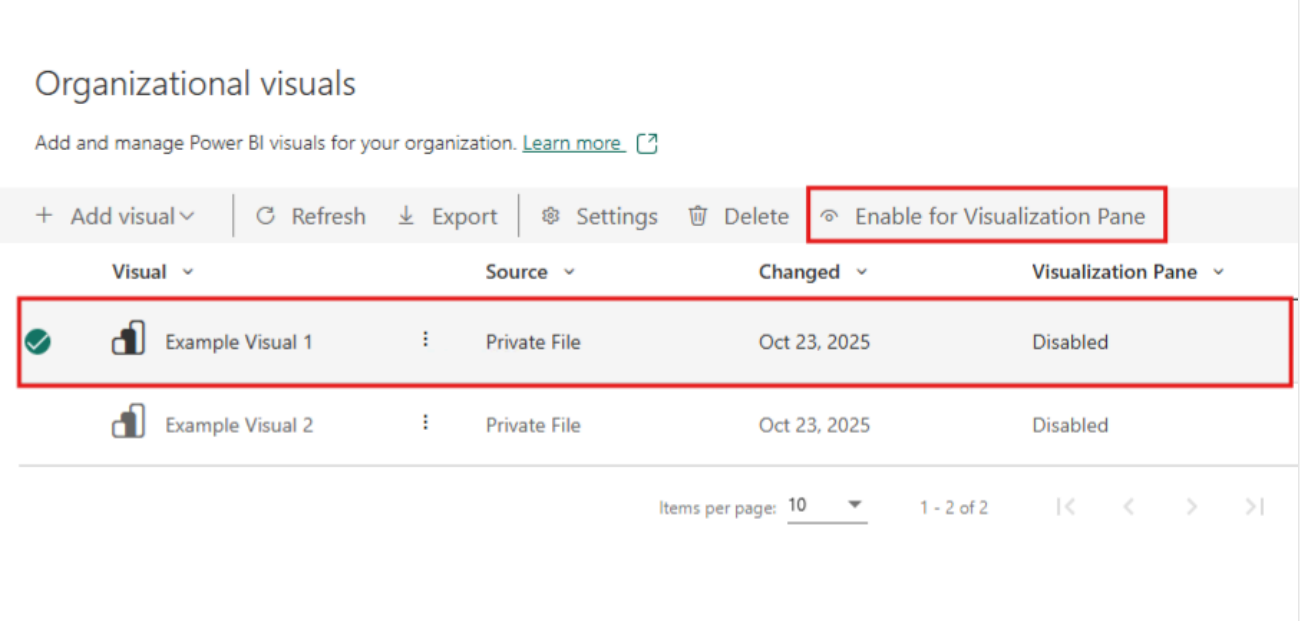

Delete a visual uploaded from a file

To permanently delete a visual, select the visual then select **Delete**.

> **Important**
>
> Deletion is irreversible. After the visual is deleted, it immediately stops rendering in existing reports. Even if you upload the same visual again, it doesn't replace the one that was deleted. However, users can import the new visual again and replace the instance they have in their reports.

Disable a *.pbiviz* visual

You can disable a *.pbiviz* visual from being available through the organizational store, while keeping it on the organizational visuals list.

1. Select the *.pbiviz* visual you want to disable, then select **Settings**.
2. In the **Access** section, disable the setting: **Users in the organization can access, view, share, and interact with this visual**.

After you disable the *.pbiviz* visual, the visual doesn't render in existing reports, and it displays the following error message:

```
This custom visual is no longer available. Contact your administrator for details.
```

> **Note**
>
> *.pbiviz* visuals that are bookmarked continue working even after they're disabled.

Update a visual

AppSource visuals update automatically. After a new version is available from AppSource, it replaces an older version deployed via the organizational visuals list.

To update a *.pbiviz* visual, follow these steps to replace the visual.

1. Select the visual you want to add, then select **Settings**.
2. Select **Browse**, then select the *.pbiviz* file you want to use to replace the current visual.
3. Select **Update**.

Replace a visual from a file with a visual from AppSource

Sometimes an organization develops its own Power BI visual and distributes it internally. After some time, the organization might decide to make this visual public by uploading it to AppSource. To replace the visual uploaded from a file with the one from AppSource, use the following steps:

1. Add the visual from AppSource into the organizational store.
2. Open the report that contains this visual. Both the visual uploaded from a file and the AppSource visual are visible in the visualizations pane.
3. In the report, highlight the visual uploaded from a file and in the visualizations pane, select the AppSource visual to replace it. The visuals are swapped automatically. To verify that you're using the AppSource visual, in the visualizations pane right-click the visual and select *about*.
4. Complete **step 3** for all the reports that contain the visual in your organization.
5. Delete the visual that you uploaded from a file.

Related content

- Administration overview
- Visuals in Power BI
- Organizational visuals in Power BI

Last updated on 09/01/2026

What is the admin monitoring workspace? (preview)

The admin monitoring workspace is a specialized environment designed for Fabric administrators to monitor and manage workloads, usage, and governance within their tenant. Using the resources available within the workspace, admins can perform tasks such as security audits, performance monitoring, capacity management, and more.

Prerequisites

To set up the admin monitoring workspace, you need a Fabric administrator role.

Installing the admin monitoring workspace

The admin monitoring workspace is automatically installed the first time an admin accesses it. Reports in the workspace appear as blank until the first data refresh. The first data refresh begins around five minutes after the workspace is installed, and usually completes within a few minutes.

To trigger the installation of the admin monitoring workspace, follow these steps:

1. Sign in to Fabric as an admin.
2. From the navigation menu, select **Workspaces**.
3. Select **Admin monitoring**. When selected for the first time, the workspace installation begins automatically and usually completes within a few minutes.

Sharing the admin monitoring workspace

Once set up, admins can share all reports in the workspace with users that aren't admins through a workspace viewer role. Admins can also share individual reports or semantic models with users that aren't admins through links or direct access.

Only admins can see the admin monitoring workspace at the top of the workspaces menu. Users that aren't admins can access the workspace's contents indirectly by using the *Browse* or *OneLake catalog* pages, or by bookmarking the workspace URL.

Managing the admin monitoring workspace

By default, the admin monitoring workspace is a Pro-licensed workspace. To take advantage of capacity benefits such as unlimited content sharing for the admin monitoring workspace, follow these steps:

1. Open the **OneLake catalog**, select the **Govern** section, and then select **Workspaces**.
2. Using the **Name** column filter, search for **Admin monitoring**.
3. Select the **Actions** button, and then select **Reassign workspace**.
4. Select the desired **Workspace type**, and then select **Save**.

Reports and semantic models

The Govern report in the OneLake catalog is the primary experience for tenant-wide governance and administration insights. It consolidates and expands on information previously divided among separate administration and security reports. The report covers inventory, capacities, domains, workspaces, feature usage, activities, sharing, endorsement, curation, sensitivity-label coverage, and data loss prevention policy activity.

The administrator report and semantic model are stored in the Admin monitoring workspace. Open the Govern report to explore the data. You can connect to semantic models in this workspace to create reporting solutions optimized for your organization's needs.

Considerations and limitations

- Private links aren't supported.
- Only users whose admin roles are assigned directly can set up the admin monitoring workspace. If the workspace creator's admin role is assigned through a group, data refreshes in the workspace fail.
- The admin monitoring workspace is a read-only workspace. Workspace roles don't have the same capabilities as they do in other workspaces. Workspace users, including admins, aren't able to edit or view properties of items such as semantic models and reports in the workspace.
- Users with *build* permissions for a semantic model in the admin monitoring workspace are shown as having *read* permissions.
- Granular delegated admin privileges (GDAP) aren't supported.

- Once access is provided to the admin monitoring workspace or its underlying content, access can't be removed without reinitializing the workspace. However, sharing links can be modified as with a typical workspace.
- Semantic models in the admin monitoring workspace are read-only and can't be used with Fabric data agents.

Refreshes

The semantic models in the workspace are automatically refreshed once per day, around the same time that the workspace was installed for the first time.

To maintain the scheduled refresh process, consider the following limitations:

- If the user who first accessed the workspace is no longer an admin, scheduled refreshes in the workspace fail. This issue can be mitigated by having any other admin log into Fabric, as their credentials will automatically be assigned to all semantic models in the workspace to support any future data refreshes.
- If the admin who first accessed the workspace uses Privileged Identity Management (PIM), their PIM access must be active during the time of scheduled data refresh, otherwise the refresh fails.

Reinitializing the workspace

Occasionally, administrators may need to reinitialize the workspace, including to reset access to the workspace or its underlying content.

Admins can execute an API to reinitialize the workspace using the following steps:

1. Retrieve the ID of the admin monitoring workspace from the URL when viewing the workspace.
2. Execute the semantic model deletion API, first replacing the 'xxxxxxxx-xxxx-xxxx-xxxx-xxxxxxxxxxxx' with the ID of your admin monitoring workspace.

   ```
   api.powerbi.com/v1/admin/workspaces/xxxxxxxx-xxxx-xxxx-xxxx-xxxxxxxxxxxx -Method Delete
   ```

3. Click the **Workspaces** menu and select **Admin monitoring** to trigger the reinitialization of the workspace, similar to the process of the first installation. On occasion, refreshing the page is also required.

Related content

- Admin overview
- Govern report in the OneLake catalog

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Feature usage and adoption report (preview)

> **Important**
>
> The Govern report in the OneLake catalog is generally available (GA). Use the Govern report for tenant-wide governance and administration insights. The Feature usage and adoption report remains in preview, won't become GA, and won't be expanded with additional data.

The Feature Usage and Adoption Report is aimed at admins who want to understand how Fabric features are utilized across the organization. As an admin, the report insights can help you govern your Fabric tenant and take action when needed.

You can access the report from the Admin monitoring workspace. To access the workspace, you must be a Fabric administrator.

You can also have an admin share the report or semantic model directly with you. With build permissions to the semantic model, users can design a custom report that relies on the same underlying data.

Navigation

The report is designed for admins to analyze Fabric activity in various ways. You can filter activity data across all pages using the date range slicer to focus on a specific time period within the last 30 days. Additionally, the filter pane lets you refine your analysis based on capacity, user, item-related information, and other activity characteristics.

The following image shows an example layout. Your view is different based on activity in your tenant.

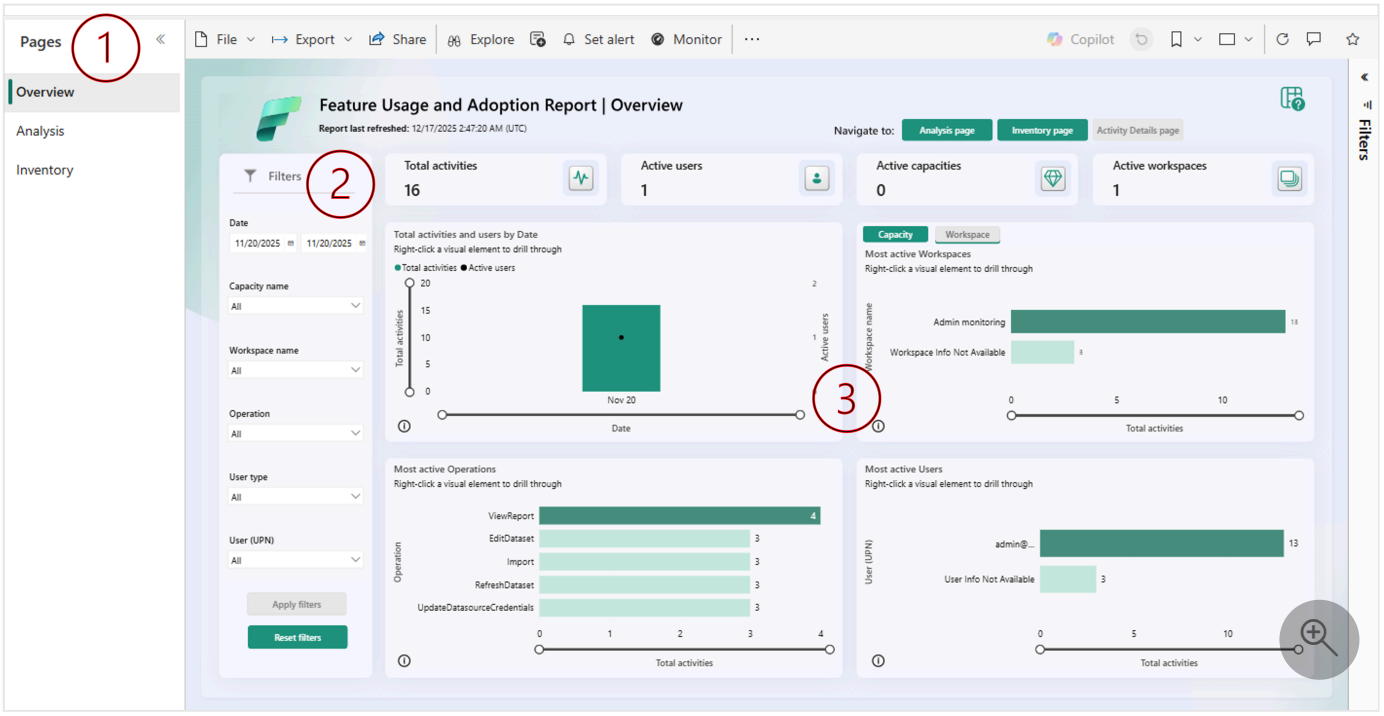

The report consists of these main sections:

1. **Pages** - Navigate between report pages using this section. You can also use the tabs in the upper right of the report to move between report pages and details pages.
2. **Filters** - Specify date ranges and apply filters to analyze Fabric activity based on different characteristics.
3. **Report data and visuals** - Analyze Fabric activity across your organization. Right-click visual elements and select **Drill through** to navigate to details pages for deeper insights.

For detailed information and examples of how to use each report page, see the following sections.

Report pages

The report is composed of five pages:

- **Overview** - Provides a high-level overview of Fabric activity across the organization. This page is one of the main entry points to the report.
- **Analysis** - Visualizes activity across different activity dimensions.
- **Activity Details** - Display this page by drilling through from the Overview or Analysis pages. It shows detailed information on specific activity scenarios.
- **Inventory** - Lists all Fabric items in your tenant.

- **Item Details page** - Display this page by drilling through from the Inventory page. It shows detailed information on specific inventory usage scenarios.

Overview page

The Overview page helps you identify:

- Daily activities and user trends
- The most active capacities and workspaces
- Activities in your organization by your most or least active users

Example

In a large retail organization, you might use the Overview page to check which capacities were most utilized in a given month. You use the date range slicer to filter to the month of December. You notice the *Sales and Marketing* capacity had nearly 1,000 activities while other capacities had under 200. To understand why, you then go to the Analysis page.

Analysis page

On the Analysis page, you can view:

- A daily count of activity and users by date
- A decomposition tree to drill down into activity using dimensions such as operation, capacity, user, and more

Example

Continuing the example from the Overview page, you use the Analysis page to investigate why the *Sales and Marketing* capacity had more activity in December than all other capacities. The decomposition tree reveals the most popular activity on the *Sales and Marketing* capacity was *ViewReport*, which signifies the viewing of a Power BI report. You then drill through to the Activity Details page to identify which reports were most frequently viewed that month on the *Sales and Marketing* capacity.

To drill through to the Activity Details page:

1. Right-click the visual element (such as Operation name) you want to drill through from.
2. Select *Drill through*.

3. Select *Activity Details*.

Activity Details page

The Activity Details page shows detailed information on specific activity scenarios. Users can access this page by drilling through from the Overview or Analysis pages to display the following activity details:

- **Creation time** - The time the activity was registered
- **Capacity name** - The name of the capacity that the activity took place on
- **Capacity ID** - The ID of the capacity that the activity took place on
- **Workspace name** - The name of the workspace that the activity took place in
- **Workspace ID** - The ID of the workspace that the activity took place in
- **User (UPN)** - The user principal name of the user who conducted the activity
- **Operation** - The formal name of the operation
- **Total of activities** - The number of times the activity took place

Example

From the Analysis page, you drill through on frequently conducted *ViewReport* actions on the *Sales and Marketing* capacity in December. Using info from the Activity Details page, you discover that a new report titled "Unclosed Deals" was heavily viewed, prompting further investigation to understand the report's relevance to your organization's sales strategy.

Inventory page

The Inventory page displays all items in your Fabric tenant and how they're utilized. You can filter the Inventory page by:

- **Item type** - Including reports, dashboards, lakehouses, notebooks, and more
- **Workspace name** - The name of the workspace where the items are located
- **Activity status** - Indicates whether the item was recently utilized
  - *Active* - At least one audit log activity was generated related to the item over the last 30 days

- *Inactive* - No audit log activity was generated related to the item over the last 30 days

Example

The Inventory page also includes a decomposition tree visual to breakdown inventory by different factors such as capacity, user, workspace, and more. You can use the decomposition tree to decompose items by activity status; for example, displaying all inactive items by item name so that you can decide whether any of these items can be deleted.

Item Details page

The Item Details page shows information related to specific inventory usage scenarios.

Users can navigate to the Item Details page by drilling through from the Inventory page. To drill through, right-click a visual element (such as Item type) and then select the Item Details page from the *Drill through* menu.

After drilling through, you see the following information for the selected item types:

- **Capacity ID** - The ID of the capacity that the item is hosted on
- **Workspace ID** - The ID of the workspace that the item is located in
- **Workspace name** - The name of the workspace that the item is located in
- **Item ID** - The unique ID of the item
- **Item name** - The display name of the item
- **Item type** - The type of item such as report, semantic model, app, and so on
- **Modified by** - The ID of the user that last modified the item
- **Activity status** - The status of an item whether it's active or inactive based on recent activity
- **Items** - The total number of items

Measures

The following measures are used in visuals throughout the report and are also available in the semantic model.

Measure calculations consider filter context, so measure values change as you apply filters or interact with other visuals.

Expand table

Measure name	Description
Active capacities	The number of capacities with audit activity.
Active users	The number of users who generated audit activity.
Active workspaces	The number of workspaces with audit activity.
Activities	The number of audit activities generated.
Items	The count of items displayed.
Total activities	The number of audit activities generated. Reflected as 0 when no audit data is returned; used exclusively in card visuals.
Total items	The count of items displayed. Reflected as 0 when no items are returned; used exclusively in card visuals.

Considerations and limitations

This section lists the report's considerations and limitations.

Display

- Condensing the zoom slider on a date trend visual to a single day displays a misleading time range, as activities aggregate by day and not by time.
- Using the *next level in the hierarchy* option on the *Most active Capacities* visual doesn't update the dynamic visual title.
- Items with the same name, or items deleted and recreated with the same name, might reflect as one item in certain visuals. To count the total number of unique items, use item identifiers or the *Total items* measure.
- *Info Not Available* represents data that isn't available. Data might not be available when an audit event doesn't have complete information, or when that information isn't applicable for the event.
- The report reflects your tenant's audit activity on a rolling 30-day basis, along with the latest snapshot of your tenant's inventory.

- Deleted inventory (including workspaces with extended retention) doesn't appear in the report after 28 days post deletion. They can be seen in Fabric until they're permanently deleted.
- Items created and deleted within a 24 hour period could have incomplete information.

Power BI Pro and Power BI Premium Per-User (PPU)

Semantic models in *Power BI Pro* and *Power BI Premium Per-User* (PPU) workspaces are hosted on internal logical capacities. The usage of these capacities can be seen in this report.

- **Pro** - Appear as *Reserved Capacity for Pro Workspaces* with the capacity SKU value *Pro*.
- **PPU** - Appear as *Reserved Capacity for Power BI Premium Per-User (PPU) Workspaces* with the capacity SKU value *PPU*.

Counting logic

- All *My workspaces* are counted as separate records as part of the *Active workspaces* measure.

Model limitations

- The semantic model is read-only and can't be used with Fabric data agents.

Related content

- What is the Admin monitoring workspace?
- Admin overview

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Content sharing report (preview)

07/02/2025

The *content sharing* report is aimed at admins who want to understand how Fabric items are distributed and shared across the organization. As an admin, the report insights can help you govern your Fabric tenant and take action if needed.

You can access the report from the Admin monitoring workspace. To access the workspace, you must be a Fabric administrator or a Microsoft 365 global administrator.

You can also have an admin share the report or semantic model directly with you. With build permissions to the semantic model, users can design a custom report that relies on the same underlying data.

Navigation

The report is designed for admins to analyze their Fabric inventory in various ways.

Use the slicers pane on the left side of the report to filter inventory on key attributes such as workspace, item type, and capacity. The filters pane on the right side of the report provides more options to further manipulate the data.

To switch pages, use the navigation buttons at the top right of the report.

Report pages

The report is composed of four pages:

- **Inventory Overview** - Provides a high-level overview of Fabric items across your organization
- **Analysis** - Visualizes inventory across different dimensions
- **Workspace** - Shows detailed inventory information for one or multiple workspaces
- **Details** - Lists all Fabric items in a tabular view

Inventory Overview page

The Inventory Overview page provides a high-level summary of how items are distributed in your organization.

Use the Inventory Overview page to quickly answer questions such as:

- Which workspaces in my tenant have the most items?
- Which Fabric item types are most commonly used in my organization?
- How are my items distributed by endorsement and sensitivity label?
- Is my organization effectively using domains to organize its Fabric inventory?

Analysis page

On the Analysis page, you're provided with a decomposition tree visual to aggregate and drill into your inventory.

Using the toggle at the bottom left of the page, you can decide whether to break down your inventory by total items or by total access count. Definitions for these measures are provided in the Measures section.

You can then break down either measure across different item-related dimensions, including item type, workspace name, endorsement, and more.

The Analysis page can be accessed directly using the page navigation buttons, or by drilling through from other pages in the report.

Example

In a large organization, the Analysis page can be helpful for analyzing item distribution. You can also drill into specific data points for a closer look at item distribution.

To drill into the details of a specific scenario from the decomposition tree visual:

1. With the *Analyze by* toggle set to *Total items*, you select the following dimensions: *Capacity name*, *Workspace name*, and *Item type*.
2. Right-click the data point on the decomposition tree visual that you want to drill into.
3. From the *Drill through* menu, select the Workspace or Details page.

After you drill through, the target page and its visuals are prefiltered to the subset of data selected from the Analysis page.

Workspace page

On the Workspace page, you're provided with details on how items within a specific workspace are shared. The Workspace page includes visuals that highlight certain scenarios such as:

- Most shared items by total access count
- Item distribution by endorsement and sensitivity label
- Items shared with the entire organization using links
- Item deleted within the last 28 days

Example

To drill into the details of a specific workspace:

1. Right-click any visual that uses the *Workspace name* field, then select the *Workspace* page from the *Drill through* menu.
2. After drilling through, you can see detailed metrics for the selected workspace, with the ability to drill through even further to the *Details* page.
3. Once on the Workspace page, users can select attributes from the available slicers or the filters pane to further manipulate the data.

In addition to drilling through, users can also access the *Workspace* page directly using the page navigation buttons.

Details page

The Details page highlights item distribution in a tabular format.

You can also navigate to the *Details* page by drilling through from other pages in the report. To drill through, right-click a value in any visual, then select the *Details* page from the *Drill through* menu. After drilling through to the *Details* page, you can see information for the specific subset of items.

You can export data from the Details page, or any other visual, by clicking *More options* in the visual header and selecting *Export data*.

You can also navigate directly to an item or its workspace using the hyperlinks in the data table.

> **Note**

Workspaces can only be accessed if you have a valid workspace role, else you are redirected to your *My workspace*. Item urls currently only support legacy Power BI items and some Fabric items.

Measures

The following measures are used in visuals throughout the *content sharing* report and are also available in the semantic model.

Expand table

Measure name	Description
Total items	The number of Fabric items across the entire tenant.
Total domains	The number of domains with items.
Total capacities	The number of capacities with items.
Total workspaces	The number of workspaces with items.
User access count	The number of individual users and service principals with access to an item.
Group access count	The number of group members and service principals with access to an item. Group owners aren't included in *group access counts*. Group access counts are calculated by flattening membership of all nested groups, so users aren't double counted if they're members of multiple groups in a nest. *Group access counts* also include +1 for each nested group in a nest.
Total access count	The number of individual users, service principals, and group members with access to an item. *Total access counts* for workspaces, capacities, and domains are a sum of access counts for all underlying items, not the container itself. *Total access counts* include both individual access to an item and access through a group, so users are double counted if they have access to an item in both scenarios.

Note

Access counts include access to an item through *Manage permissions*, or access inherited through a workspace role. Access counts also include service principals and sharing links for a specific persons or group.

Considerations and limitations

This section lists the report's considerations and limitations.

Display

- Items with the same name, or items deleted and recreated with the same name, might reflect as one item in certain visuals. To count the total number of unique items, use item IDs or the *Total items* measure.
- The report retains information for 28 days, including deleted capacities, workspaces, and other items.
- Deleted workspaces with extended retention don't appear in the report after 28 days. They can be seen in the admin portal until they're permanently deleted.
- Items created and deleted within a 24 hour period may have incomplete information.
- Reports and dashboards embedded in apps appear twice. Use the *Item ID* value to differentiate.

Pro and Premium Per User (PPU)

Semantic models in *Pro* and *Premium Per User* (PPU) workspaces are hosted on internal logical capacities.

- **Pro** - Appear as *Reserved Capacity for Pro Workspaces* with the capacity SKU value *Pro*.
- **PPU** - Appear as *Reserved Capacity for Premium Per User Workspaces* with the capacity SKU value *PPU*.

Counting logic

- All *My workspaces* are counted as separate records as part of the *Total workspaces* measure.
- Trial Fabric capacities are counted as separate records as part of the *Total capacities* measure. Trial capacities can be filtered out using the Capacity SKU filter with the value *FT1*.

Related content

- What is the Admin monitoring workspace?

- Admin overview

What is the Monitor hub?

The Monitor hub is the unified observability pane in Microsoft Fabric. It brings job execution health, alerts, capacity, agents, and application telemetry together in one place. Instead of checking each item type on its own, you can spot problems across your Fabric estate and act on them from a single surface.

This article explains what you can monitor in the Monitor hub and how its dedicated pages for job runs, job alerts, applications, capacity, and agents map to different parts of your environment.

Any Fabric user can open the Monitor hub, but you only see activities for Fabric items you have permission to view.

> **Important**
>
> The Monitor hub is currently in preview.
>
> See the Supplemental Terms of Use for Microsoft Azure Previews for legal terms that apply to Azure features that are in beta, preview, or otherwise not yet released into general availability.

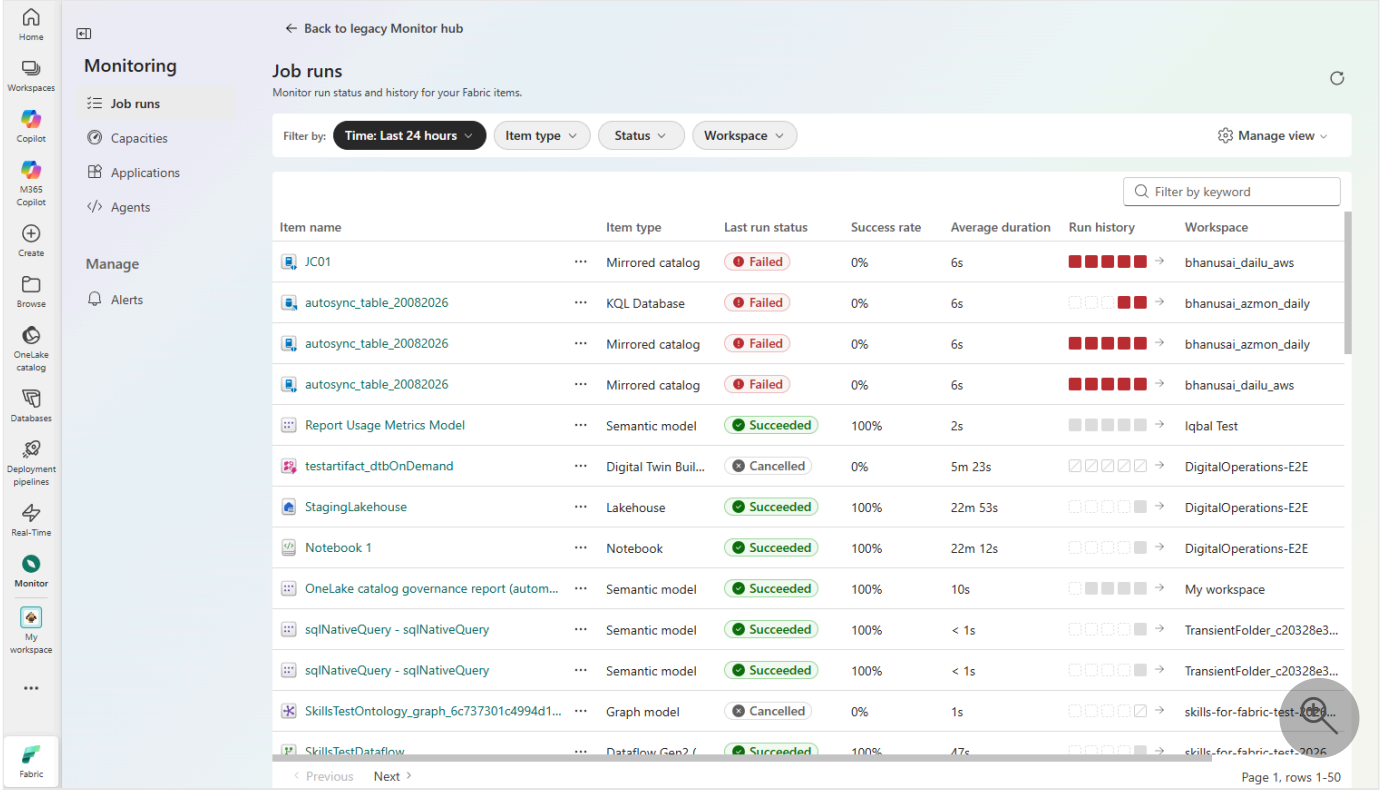

What can you monitor in the Monitor hub?

The Monitor hub organizes observability into dedicated pages, each focused on a different part of your Fabric estate. Select a page to learn how to use it:

- **Job runs** — Track job runs across the items in your workspace.
- **Applications** — Monitor application health and usage.
- **Capacity and cost** — Monitor capacity health and consumption.
- **Agents** — Monitor agent activity and health.
- **Alerts** — Manage who receives notifications about job failures and other job events.

Investigate failed pipeline runs with the Operations agent (preview)

For failed pipeline runs, the Monitor hub gives you an entry point to a read-only **Operations agent** (preview) investigation. From the run history of a failed pipeline job, you can start an investigation that reviews the run and surfaces likely root causes without changing your pipeline or its configuration. This entry point connects job-run monitoring in the Monitor hub to the **Operations agent** (preview) in Real-Time Intelligence.

To start an investigation from a failed run, see Monitor job runs in the Monitor hub. For **Operations agent** (preview) setup, prerequisites, governance, responsible AI guidance, and limitations, see Create and configure operations agents.

Related content

- Feature usage and adoption report
- Job scheduler in Fabric
- Admin overview

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Monitor job runs in the Monitor hub

The **Job runs** page in the Monitor hub gives you a centralized view of job execution health, progress, and outcomes across the items in Fabric. Each row represents one job run, so a single scan of the list tells you what executed, how it finished, and how long it took.

Because every item type that produces jobs lands in one list, the **Job runs** page is usually the fastest way to find out if a scheduled job ran successfully. This article shows you how to open the **Job runs** page, filter the job runs, and open run details to investigate failures.

Any Fabric user can open the Monitor hub, but you only see activities for Fabric items you have permission to view.

> **Important**
>
> The **Job runs** page in the Monitor hub is currently in preview. See the Supplemental Terms of Use for Microsoft Azure Previews for legal terms that apply to Azure features that are in beta, preview, or otherwise not yet released into general availability.

Prerequisites

Before you begin, ensure that you have the following prerequisites:

- Access to at least one item capable of running jobs, such as a notebook, a pipeline, or a dataflow.
- Sufficient permissions to view run history. Some columns and actions require a higher level of access than read-only.
- At least one job run inside the selected time range. The list appears empty when no runs fall within the specified window.

View the Job runs page

To view the job runs in your workspace:

1. Sign in to Microsoft Fabric.
2. Open **Monitor** from the navigation pane. Switch to the new Monitor hub experience if you're still in the classic view.

3. Select **Job runs**. The page shows a list of job runs for your workspace, with a summarized history of the last five runs for each item. The **Last run status** column uses a color-coded indicator for each status (In progress, Succeeded, Failed, and Canceled), making failures easy to see.

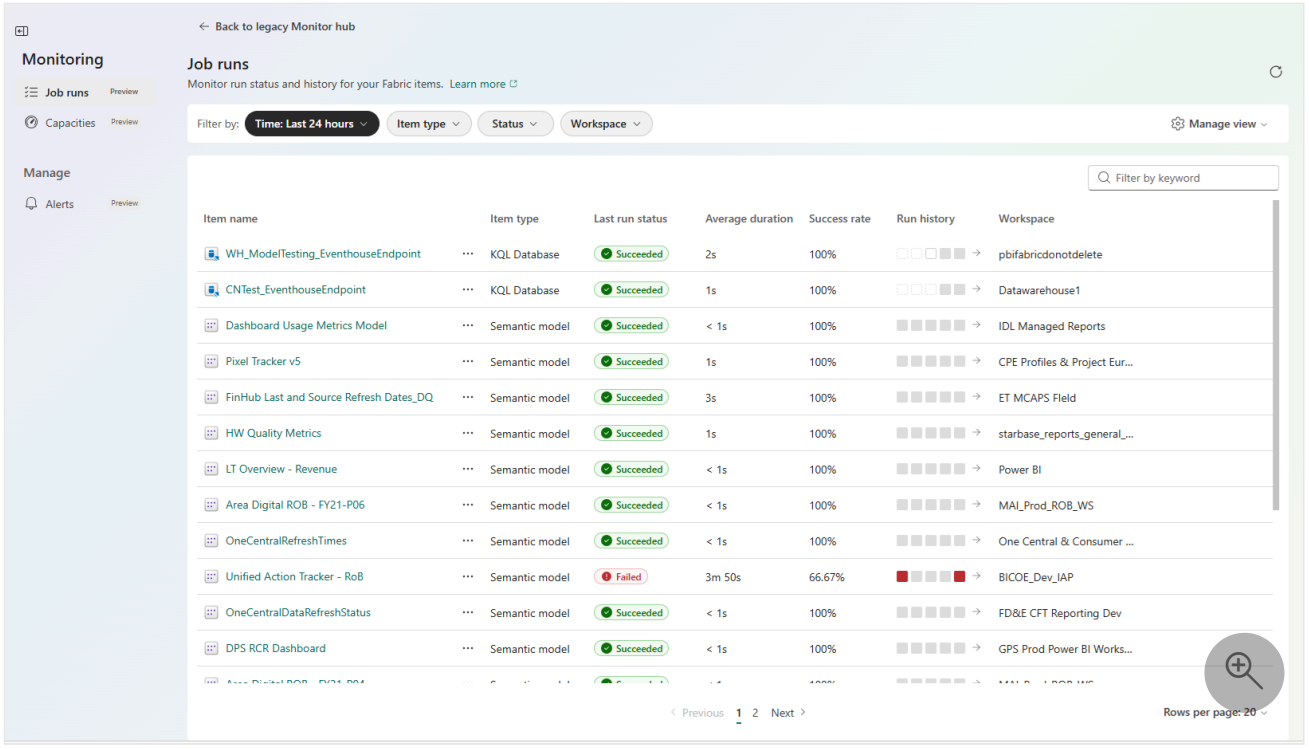

Sort and filter the job runs view

Use the following options to find the job runs you're interested in.

- **Set the time range:** Use the **Time** range selector to scope the list to a time period, such as the last hour, 6 hours, 24 hours, or the number of days. If a run you expect is missing, check the time range and adjust it as needed.
- **Sort and customize columns:** Select a column header to sort the list. For example, sort by **Item name** for an alphabetical view or by **Average duration** to find the longest runs. Use **Manage view** > **Column Options** to add, remove, or reorder columns. Use **Manage view** > **Save view** to save your custom view.
- **Use the filter drop-down menus:** To find a group of runs, use the drop-down menus to filter by **Item type**, job **Status**, or **Workspace**. For example, filter by failed status to create a focused worklist.
- **Filter by keyword:** Type text in the search box to list only those items that contain the text.

> **Tip**

> Keyword search checks only items loaded on the current page. The filter drop-down menus don't have this limitation. If your items span multiple pages, use the drop-down menus to narrow the list, and then select **Manage view** > **Save view**. Search the saved view to find specific items within that focused list.

- **Refresh the list:** The page retrieves job data when it loads and doesn't update continuously. Select the **Reload** icon in the upper right to see updates for active or recently triggered runs.

Save views

After you filter and customize your view of the jobs run list, you can save it for later use:

1. Select **Manage view** > **Save view**.
2. Enter a name for your view and select **Save**.
3. Your saved view is now available under **Manage view** > **Saved views**.

Get job run details

To see details about a job run, select it in the list. A pane opens to the side showing details about the job and its run history. These views help you drill down to the cause of failures without leaving the Monitor hub.

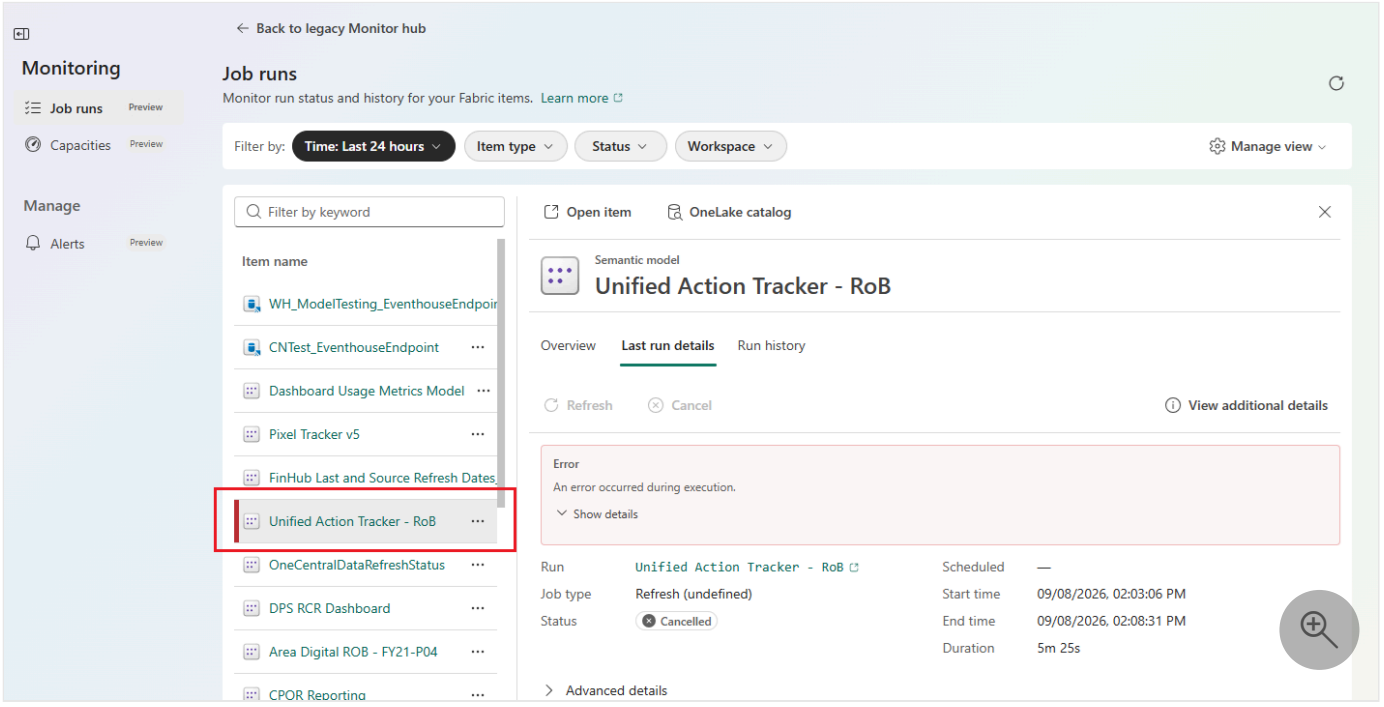

View more details about job runs

- Under **Last run details**, expand **Advanced details** to find out more about the workspace, who submitted the job, and what invoked the job run.
- Select **Run history** and review past runs to understand patterns and recurring issues. Select a run to open a pane that shows advanced **Run details**, including options to view the activities as a table or a Gantt chart.

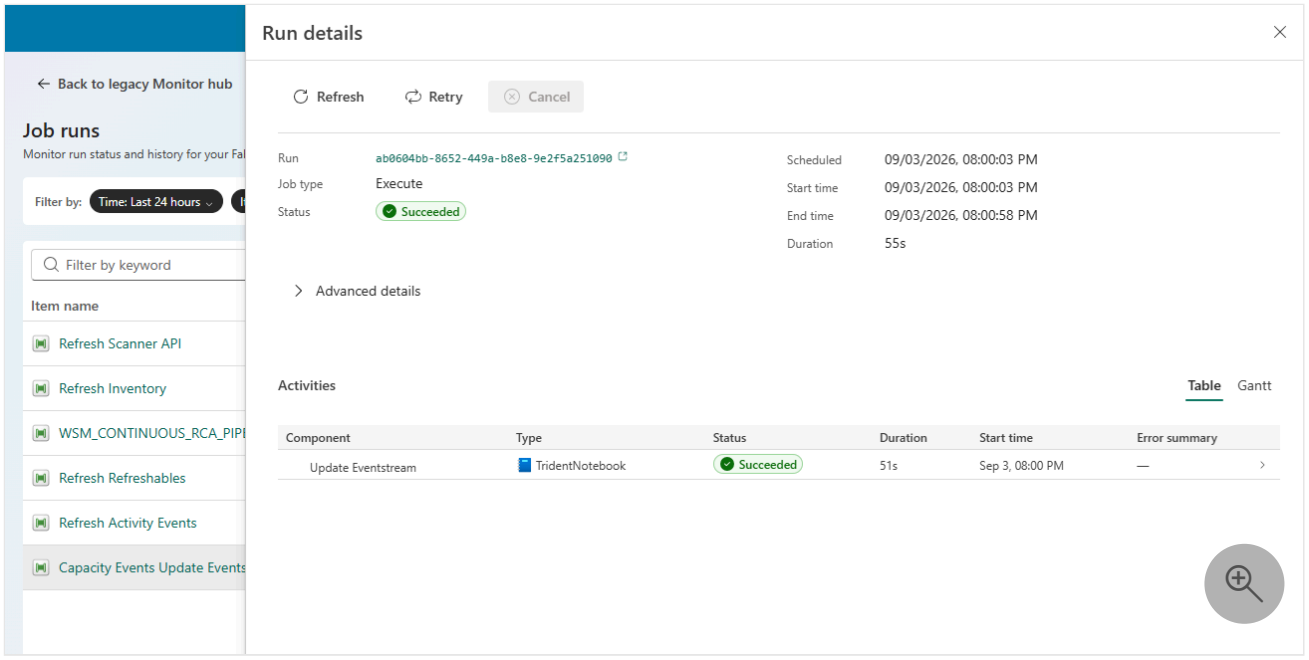

- Select the **Open detail page** icon next to a run in the item's list to open the full detail page. For complex items such as pipelines and dataflows, the details pane includes the execution flow, activity tables, and Gantt charts.

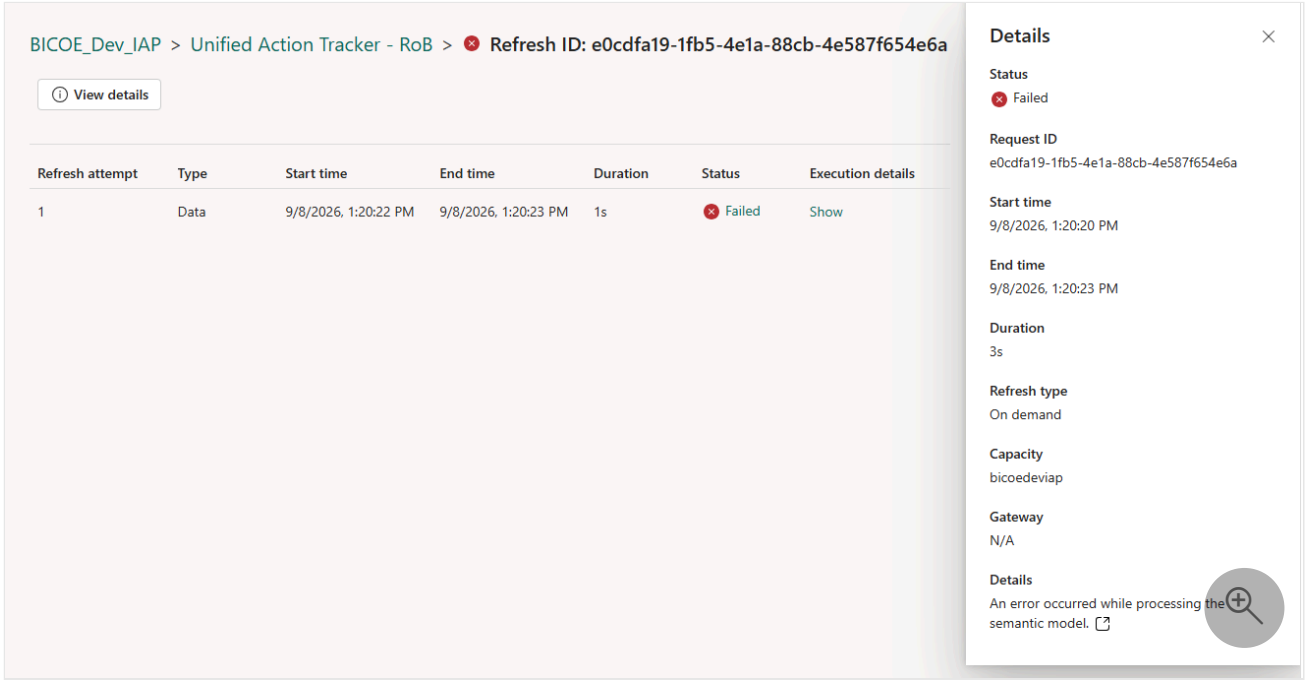

Open an item to view its authoring page

You can open the authoring page for an item directly from the job runs list by using either of these methods:

- In the job runs list, select the **Open item** icon next to an item.
- Select the item to open the job run details pane, and then select **Open item**.

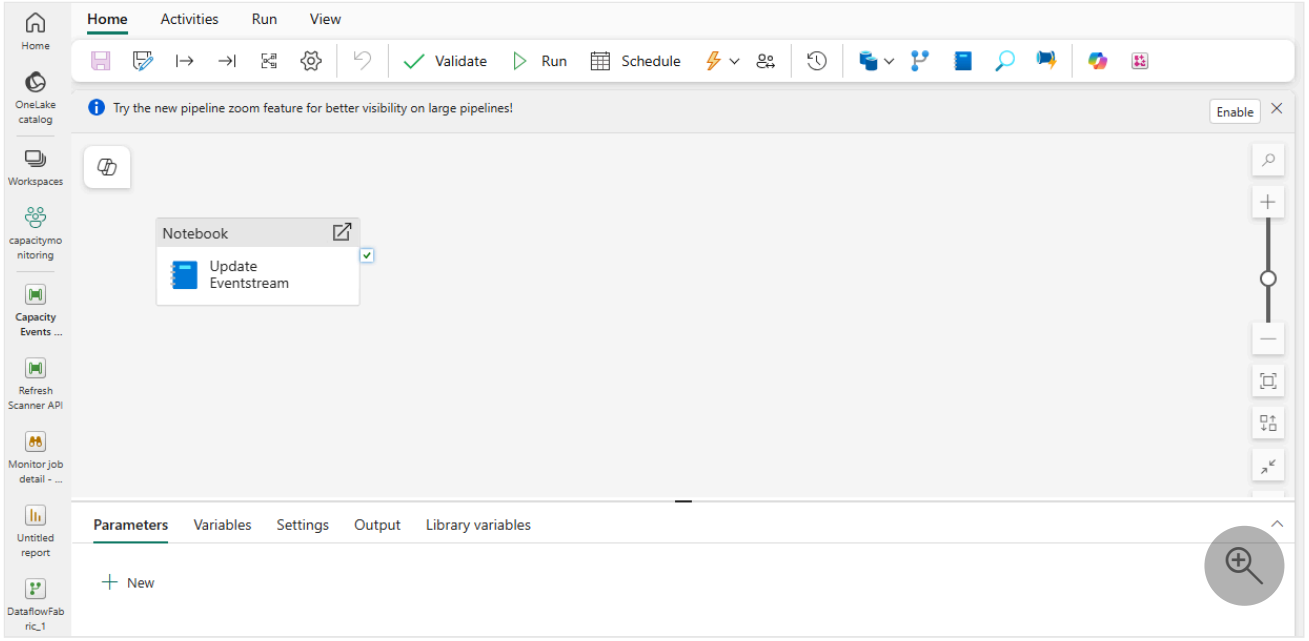

Investigate job runs

On the **Job runs** page, the **Last run status** column provides key information about each run. Use the status to determine what action to take next.

Expand table

Status	What it means	What to do next
Queued or not started	The run is waiting for resources.	Wait. If it stays queued, check capacity availability.
In progress	The run is executing now.	Monitor its duration and refresh for updates.
Succeeded	The run finished without error.	No action needed.
Failed	The run ended with an error.	Open the run details to read the error and the logs.
Canceled	Someone stopped the run before it finished.	Rerun it if the work still needs to happen.

Investigate a failed run with the Operations agent (preview)

For failed pipeline runs, you can start a read-only investigation directly from the Monitor hub. The Operations agent, a Real-Time Intelligence item, reviews the failed run and surfaces likely root causes without changing your pipeline or its configuration.

In the Monitor hub, the **Job runs** list displays one row per job, and the row-level **Investigate** action applies to the latest run. To investigate a specific failed run, follow these steps:

1. In the **Job runs** list, select a failed pipeline job to open its details.
2. Under **Run history**, find the failed run.
3. Select **Investigate** on that run. The Operations agent pane opens and starts a read-only investigation of the run.

For prerequisites, governance, responsible AI guidance, limitations, and troubleshooting, see Run a Fabric observability investigation.

Common tasks

The following table summarizes the most common day-to-day tasks on the page.

Expand table

To do this	Do this
Find a specific run	Type its name in search.
See only failures	Filter the list by failed status.
Look further back in time	Widen the time range.
Change what the list shows	Sort by a column header, or use column options to add, remove, and reorder columns.
Get the latest state	Select **Refresh**.
Investigate a run	Select the row to open its details. For a failed pipeline run, select **Investigate** on the run to launch the Operations agent.

Supported item types in the Job runs page

The **Job runs** page shows activities for these Fabric items:

- Copy Job
- Dataflow Gen2
- Dataflow Gen2 CI/CD
- Datamart
- Data Build Tool (dbt) Job
- Digital Twin Builder Flow
- Experiment
- Graph model
- Lakehouse
- Map
- Notebook
- Pipeline
- Semantic model
- Snowflake database
- Spark job definition
- User data function

The **Job runs** page doesn't support or show Dataflow Gen1.

Related content

- What is the Monitor hub?
- Set up job alerts in the Monitor hub
- Monitor Rayfin applications in the Monitor hub
- Browse the Apache Spark applications in the Fabric Monitor hub
- View refresh history and monitor your dataflows

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 10/05/2026

Monitor job runs in the Monitor hub

The **Job runs** page in the Monitor hub gives you a centralized view of job execution health, progress, and outcomes across the items in Fabric. Each row represents one job run, so a single scan of the list tells you what executed, how it finished, and how long it took.

Because every item type that produces jobs lands in one list, the **Job runs** page is usually the fastest way to find out if a scheduled job ran successfully. This article shows you how to open the **Job runs** page, filter the job runs, and open run details to investigate failures.

Any Fabric user can open the Monitor hub, but you only see activities for Fabric items you have permission to view.

> **Important**
>
> The **Job runs** page in the Monitor hub is currently in preview. See the Supplemental Terms of Use for Microsoft Azure Previews for legal terms that apply to Azure features that are in beta, preview, or otherwise not yet released into general availability.

Prerequisites

Before you begin, ensure that you have the following prerequisites:

- Access to at least one item capable of running jobs, such as a notebook, a pipeline, or a dataflow.
- Sufficient permissions to view run history. Some columns and actions require a higher level of access than read-only.
- At least one job run inside the selected time range. The list appears empty when no runs fall within the specified window.

View the Job runs page

To view the job runs in your workspace:

1. Sign in to Microsoft Fabric.
2. Open **Monitor** from the navigation pane. Switch to the new Monitor hub experience if you're still in the classic view.

3. Select **Job runs**. The page shows a list of job runs for your workspace, with a summarized history of the last five runs for each item. The **Last run status** column uses a color-coded indicator for each status (In progress, Succeeded, Failed, and Canceled), making failures easy to see.

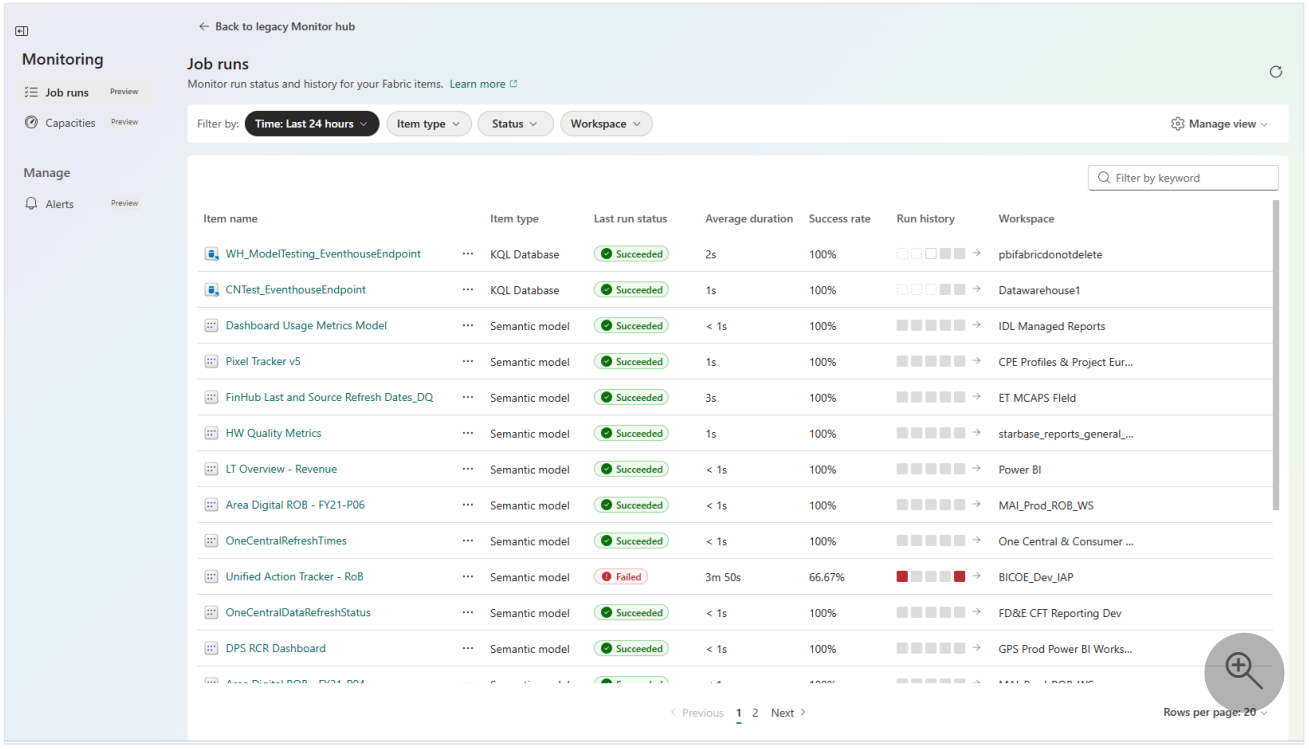

Sort and filter the job runs view

Use the following options to find the job runs you're interested in.

- **Set the time range:** Use the **Time** range selector to scope the list to a time period, such as the last hour, 6 hours, 24 hours, or the number of days. If a run you expect is missing, check the time range and adjust it as needed.
- **Sort and customize columns:** Select a column header to sort the list. For example, sort by **Item name** for an alphabetical view or by **Average duration** to find the longest runs. Use **Manage view** > **Column Options** to add, remove, or reorder columns. Use **Manage view** > **Save view** to save your custom view.
- **Use the filter drop-down menus:** To find a group of runs, use the drop-down menus to filter by **Item type**, job **Status**, or **Workspace**. For example, filter by failed status to create a focused worklist.
- **Filter by keyword:** Type text in the search box to list only those items that contain the text.

> **Tip**

> Keyword search checks only items loaded on the current page. The filter drop-down menus don't have this limitation. If your items span multiple pages, use the drop-down menus to narrow the list, and then select **Manage view** > **Save view**. Search the saved view to find specific items within that focused list.

- **Refresh the list:** The page retrieves job data when it loads and doesn't update continuously. Select the **Reload** icon in the upper right to see updates for active or recently triggered runs.

Save views

After you filter and customize your view of the jobs run list, you can save it for later use:

1. Select **Manage view** > **Save view**.
2. Enter a name for your view and select **Save**.
3. Your saved view is now available under **Manage view** > **Saved views**.

Get job run details

To see details about a job run, select it in the list. A pane opens to the side showing details about the job and its run history. These views help you drill down to the cause of failures without leaving the Monitor hub.

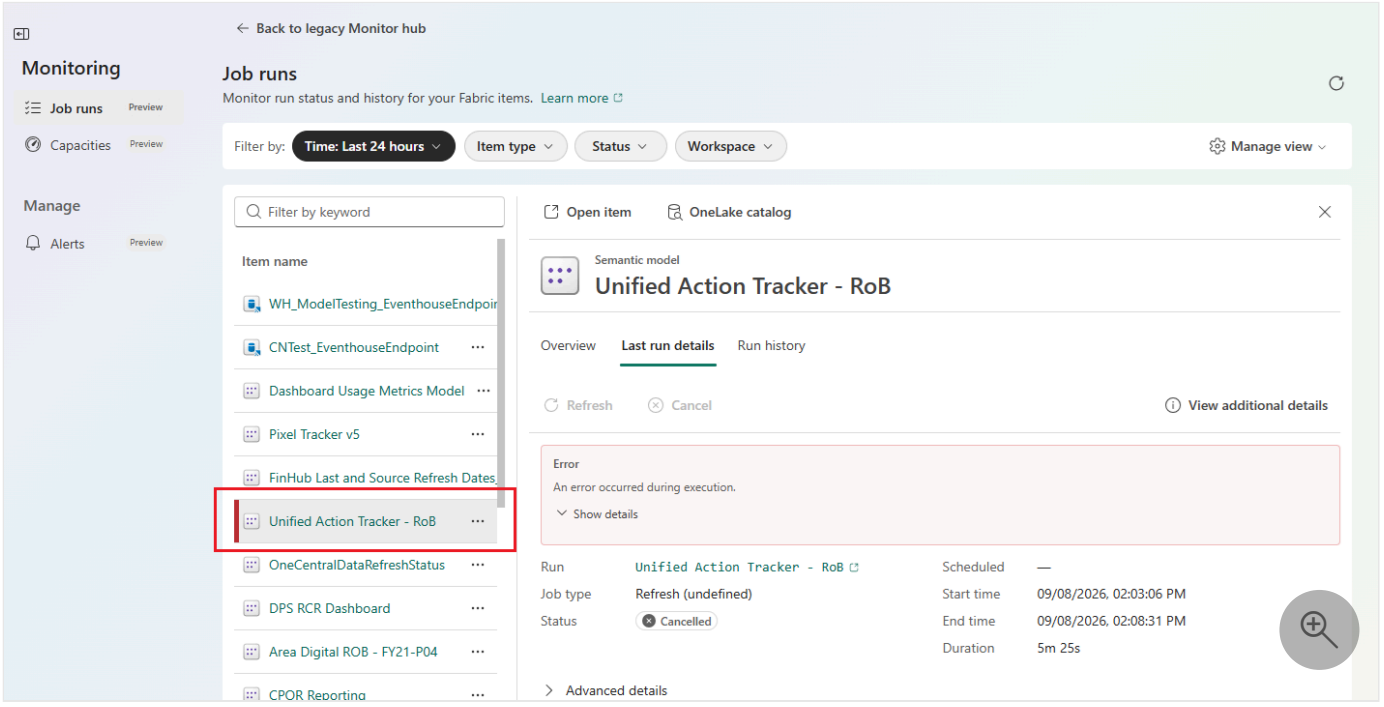

View more details about job runs

- Under **Last run details**, expand **Advanced details** to find out more about the workspace, who submitted the job, and what invoked the job run.
- Select **Run history** and review past runs to understand patterns and recurring issues. Select a run to open a pane that shows advanced **Run details**, including options to view the activities as a table or a Gantt chart.

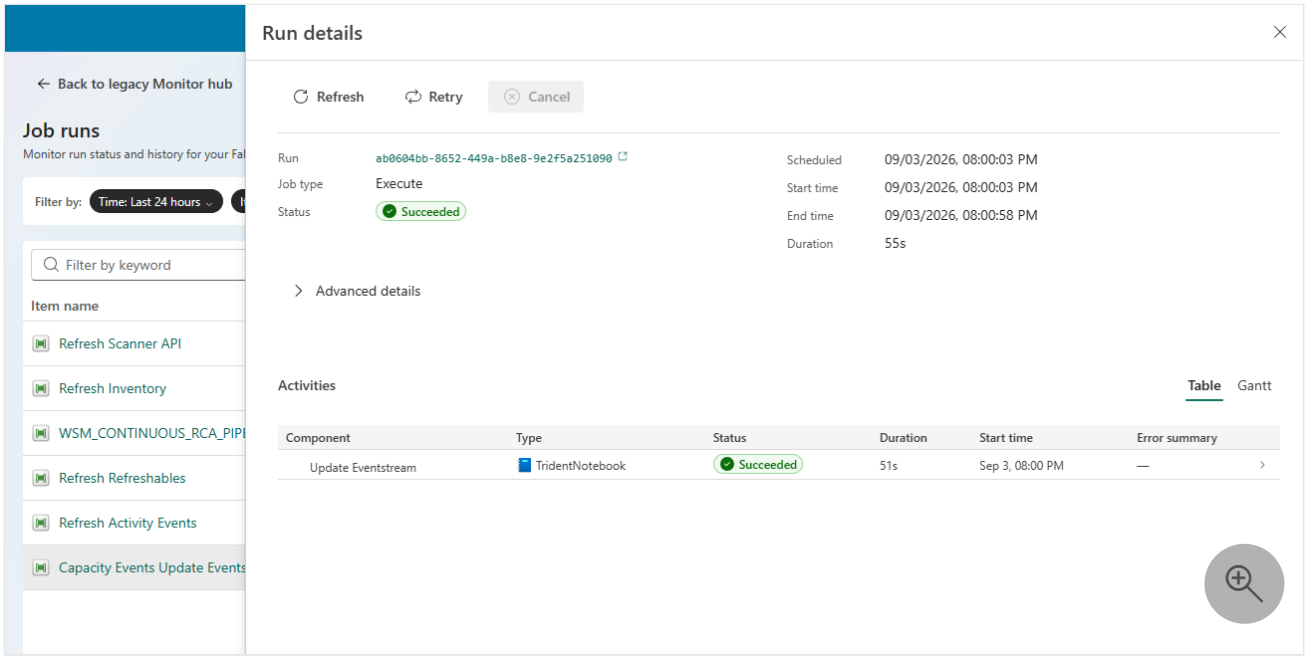

- Select the **Open detail page** icon next to a run in the item's list to open the full detail page. For complex items such as pipelines and dataflows, the details pane includes the execution flow, activity tables, and Gantt charts.

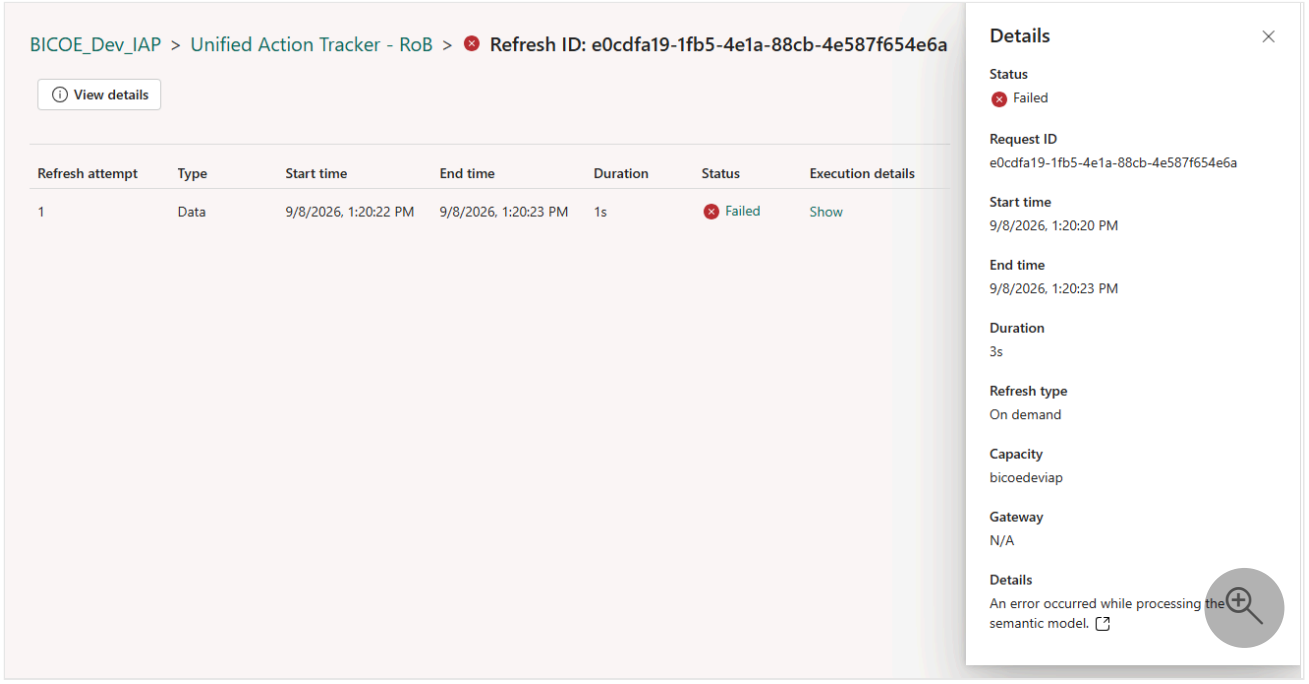

Open an item to view its authoring page

You can open the authoring page for an item directly from the job runs list by using either of these methods:

- In the job runs list, select the **Open item** icon next to an item.
- Select the item to open the job run details pane, and then select **Open item**.

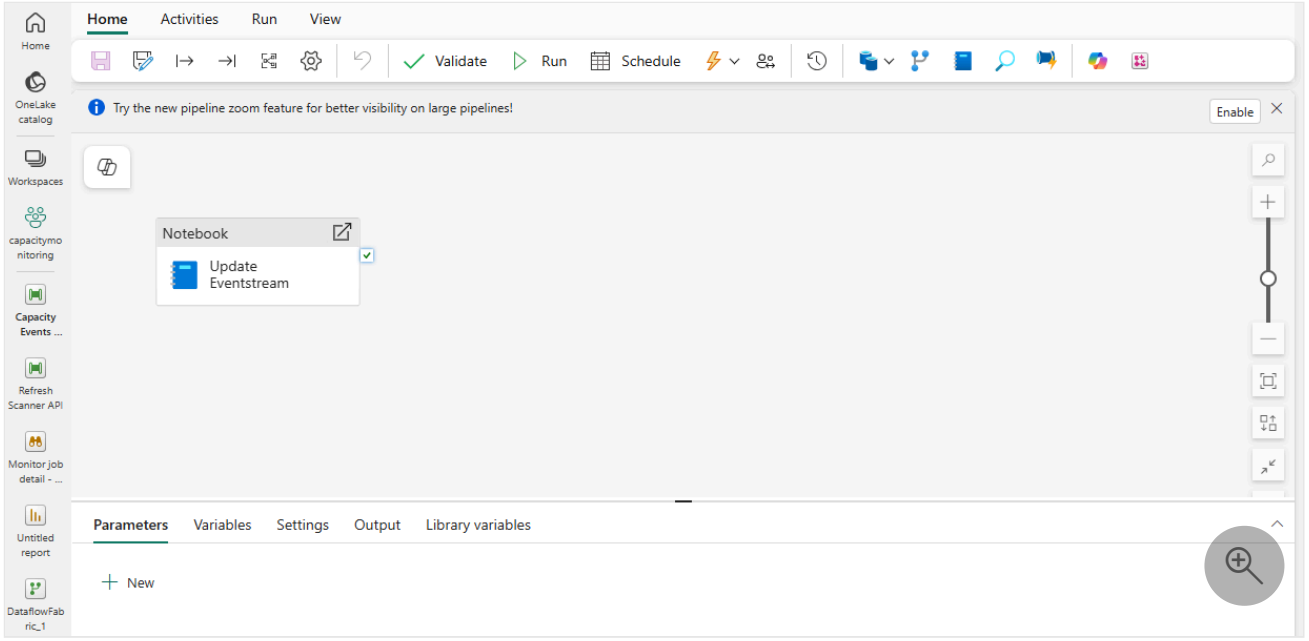

Investigate job runs

On the **Job runs** page, the **Last run status** column provides key information about each run. Use the status to determine what action to take next.

Expand table

Status	What it means	What to do next
Queued or not started	The run is waiting for resources.	Wait. If it stays queued, check capacity availability.
In progress	The run is executing now.	Monitor its duration and refresh for updates.
Succeeded	The run finished without error.	No action needed.
Failed	The run ended with an error.	Open the run details to read the error and the logs.
Canceled	Someone stopped the run before it finished.	Rerun it if the work still needs to happen.

Investigate a failed run with the Operations agent (preview)

For failed pipeline runs, you can start a read-only investigation directly from the Monitor hub. The Operations agent, a Real-Time Intelligence item, reviews the failed run and surfaces likely root causes without changing your pipeline or its configuration.

In the Monitor hub, the **Job runs** list displays one row per job, and the row-level **Investigate** action applies to the latest run. To investigate a specific failed run, follow these steps:

1. In the **Job runs** list, select a failed pipeline job to open its details.
2. Under **Run history**, find the failed run.
3. Select **Investigate** on that run. The Operations agent pane opens and starts a read-only investigation of the run.

For prerequisites, governance, responsible AI guidance, limitations, and troubleshooting, see Run a Fabric observability investigation.

Common tasks

The following table summarizes the most common day-to-day tasks on the page.

Expand table

To do this	Do this
Find a specific run	Type its name in search.
See only failures	Filter the list by failed status.
Look further back in time	Widen the time range.
Change what the list shows	Sort by a column header, or use column options to add, remove, and reorder columns.
Get the latest state	Select **Refresh**.
Investigate a run	Select the row to open its details. For a failed pipeline run, select **Investigate** on the run to launch the Operations agent.

Supported item types in the Job runs page

The **Job runs** page shows activities for these Fabric items:

- Copy Job
- Dataflow Gen2
- Dataflow Gen2 CI/CD
- Datamart
- Data Build Tool (dbt) Job
- Digital Twin Builder Flow
- Experiment
- Graph model
- Lakehouse
- Map
- Notebook
- Pipeline
- Semantic model
- Snowflake database
- Spark job definition
- User data function

The **Job runs** page doesn't support or show Dataflow Gen1.

Related content

- What is the Monitor hub?
- Set up job alerts in the Monitor hub
- Monitor Rayfin applications in the Monitor hub
- Browse the Apache Spark applications in the Fabric Monitor hub
- View refresh history and monitor your dataflows

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 10/05/2026

Set up job alerts in the Monitor hub (preview)

The Monitor hub provides a centralized place to set up and manage alerts for the jobs that run across Fabric. Instead of checking each item individually to find out whether it ran, you can get automatic notifications when a job fails or reaches another state you want to monitor.

This article shows you how to set up and manage both types of job alerts from one place:

- **Notifications for failed scheduled jobs** are free email notifications that tell you when a scheduled job fails.
- **Activator-based alert rules** support advanced job alerting, including more job events (such as **Started** and **Succeeded**), richer notifications, and automated actions.

> **Important**
>
> Job alerting in the Monitor hub is currently in preview. See the Supplemental Terms of Use for Microsoft Azure Previews for legal terms that apply to Azure features that are in beta, preview, or otherwise not yet released into general availability.

Prerequisites

Before you begin, ensure that you have the following prerequisites:

- Access to at least one item capable of running jobs, such as a notebook, a pipeline, or a dataflow.
- To configure, edit, or remove notifications for failed scheduled jobs, at least the **Contributor** role in the workspace or **Write** permission on the item. For more information, see Roles in workspaces in Microsoft Fabric.
- To create Activator-based alert rules, workspace **Owner** or **Contributor** permission, which you need to create the underlying Workspace Monitoring and Activator items.

Choose between notifications for failed scheduled jobs and Activator-based alerts

The Monitor hub manages notifications for failed scheduled jobs and Activator-based alerts together, but they remain separate alert types. You can use one or both, depending on the level of alerting each job needs. Use the following table to decide which alert type fits your needs.

Expand table

Alert type	Best for	Cost
Notifications for failed scheduled jobs	Getting a free email when a scheduled job fails.	Free
Activator-based alert rules	Advanced alerting on more job events (**Started**, **Succeeded**, and others), richer notifications, and automated actions.	Activator capacity usage and billing

Set up notifications for failed scheduled jobs

The **Schedule failure emails** experience in the Monitor hub **Alerts** section provides a centralized place to view and manage failure notifications across items. You can also manage these notifications for individual items in the job scheduler. Both experiences use the same notification settings, so changes made in one appear in the other. For details about per-item configuration and notification content, see Job Scheduler in Microsoft Fabric.

Configure notifications for failed scheduled jobs in the Monitor hub

To set up notifications for failed scheduled jobs for a scheduled item:

1. Sign in to Microsoft Fabric.
2. Open **Monitor** from the navigation pane. Switch to the new the Monitor experience if you're still in the classic view.
3. Under **Manage**, select **Alerts**.
4. Select the **Schedule failure emails** tab, and then select + **Configure notifications**.
5. Only those items that support schedules are available for you to select. Choose an item, and then choose **Select a scheduled item**.

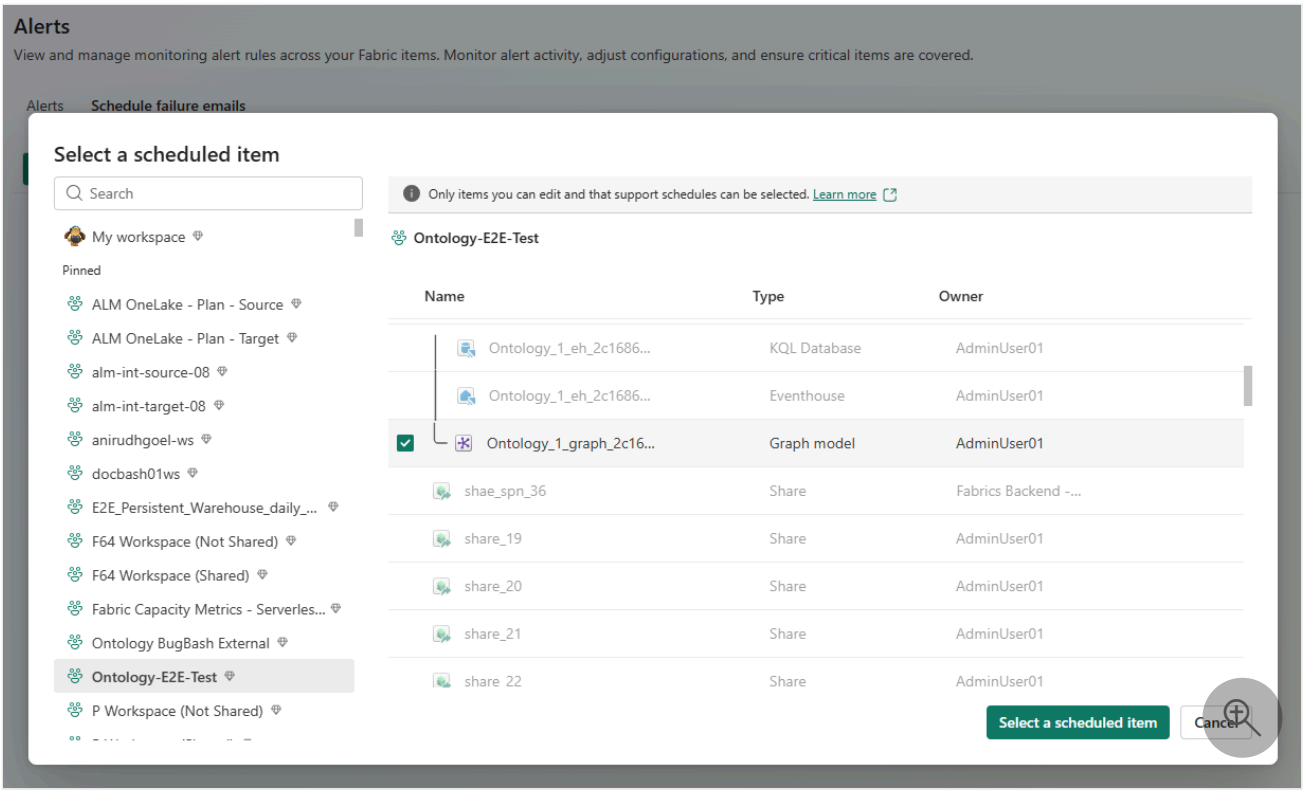

6. In **Select recipients**, enter the names or email addresses of the recipients. Recipients can be users or groups in your Microsoft Entra tenant, including B2B guest users. Direct external email addresses aren't supported.

7. Select **Save**.

> **Tip**
>
> You can also configure notifications for failed scheduled jobs from the **Job runs** page. Select the ellipses (...) next to the item, and then select **Create and manage alerts** > **Schedule failure emails (free)**.

Edit or remove notifications for failed scheduled jobs

To change or stop notifications for failed scheduled jobs for an item:

1. Sign in to Microsoft Fabric.
2. Open **Monitor** from the navigation pane. Switch to the new the Monitor experience if you're still in the classic view.
3. Under **Manage**, select **Alerts**.
4. On **Schedule failure emails**, select the item. The **Notification details** pane opens and shows the current recipients.
5. To change recipients, select **Edit recipients**, update the list, and save your changes.

6. To stop notifications, select **Remove notifications**.

Permissions and limitations

You can see failure notifications for any scheduled item that you have permission to view. To configure, edit, or remove notifications, you need at least the **Contributor** role in the workspace or **Write** permission on the item.

Keep the following limitations in mind:

- Semantic models aren't yet supported on the **Schedule failures** page.
- Failure notifications apply only to scheduled runs, not to manual runs.
- Notifications are sent in the recipient's Fabric display language, with English as the fallback.

Create Activator-based alert rules

Activator-based alert rules extend job alerting beyond notifications for failed scheduled jobs. Use them when you need to alert on more job events (such as **Started** and **Succeeded**), send richer notifications, or trigger automated actions.

> **Note**
>
> To use Activator-based alerts, you must enable workspace monitoring so that Fabric can provision a monitoring item for storing Activator-based alerts. However, you're not required to provision an eventhouse or enable any additional capabilities such as diagnostic data collection or AI powered investigation. Collecting and storing this data might incur additional costs.

Requirements for Activator-based alerts

Before you create Activator-based alerts, review the following requirements:

- Enable workspace monitoring.
- You must have workspace **Owner** or **Contributor** permission to create the Workspace Monitoring and Activator items in the workspace.
- Activator-based alerts support the following job types: pipeline, Spark job, notebook, user data function, warehouse, lakehouse, mirrored database, SQL database, and KQL database.

Enable workspace monitoring

1. Sign in to Microsoft Fabric.
2. Open the workspace where you want to change monitoring settings.
3. Select the **Workspace settings** (gear) icon. Make sure you're using the new workspace experience.
4. In the **Workspace settings** page, select **Monitoring**.
5. Select **Enable**.

> **Note**
>
> By default, all additional monitoring capabilities are enabled, including diagnostic data, telemetry data, and AI powered investigations. These capabilities aren't required for Activator-based alerts, and you can deselect them if you prefer, as shown in the following image.

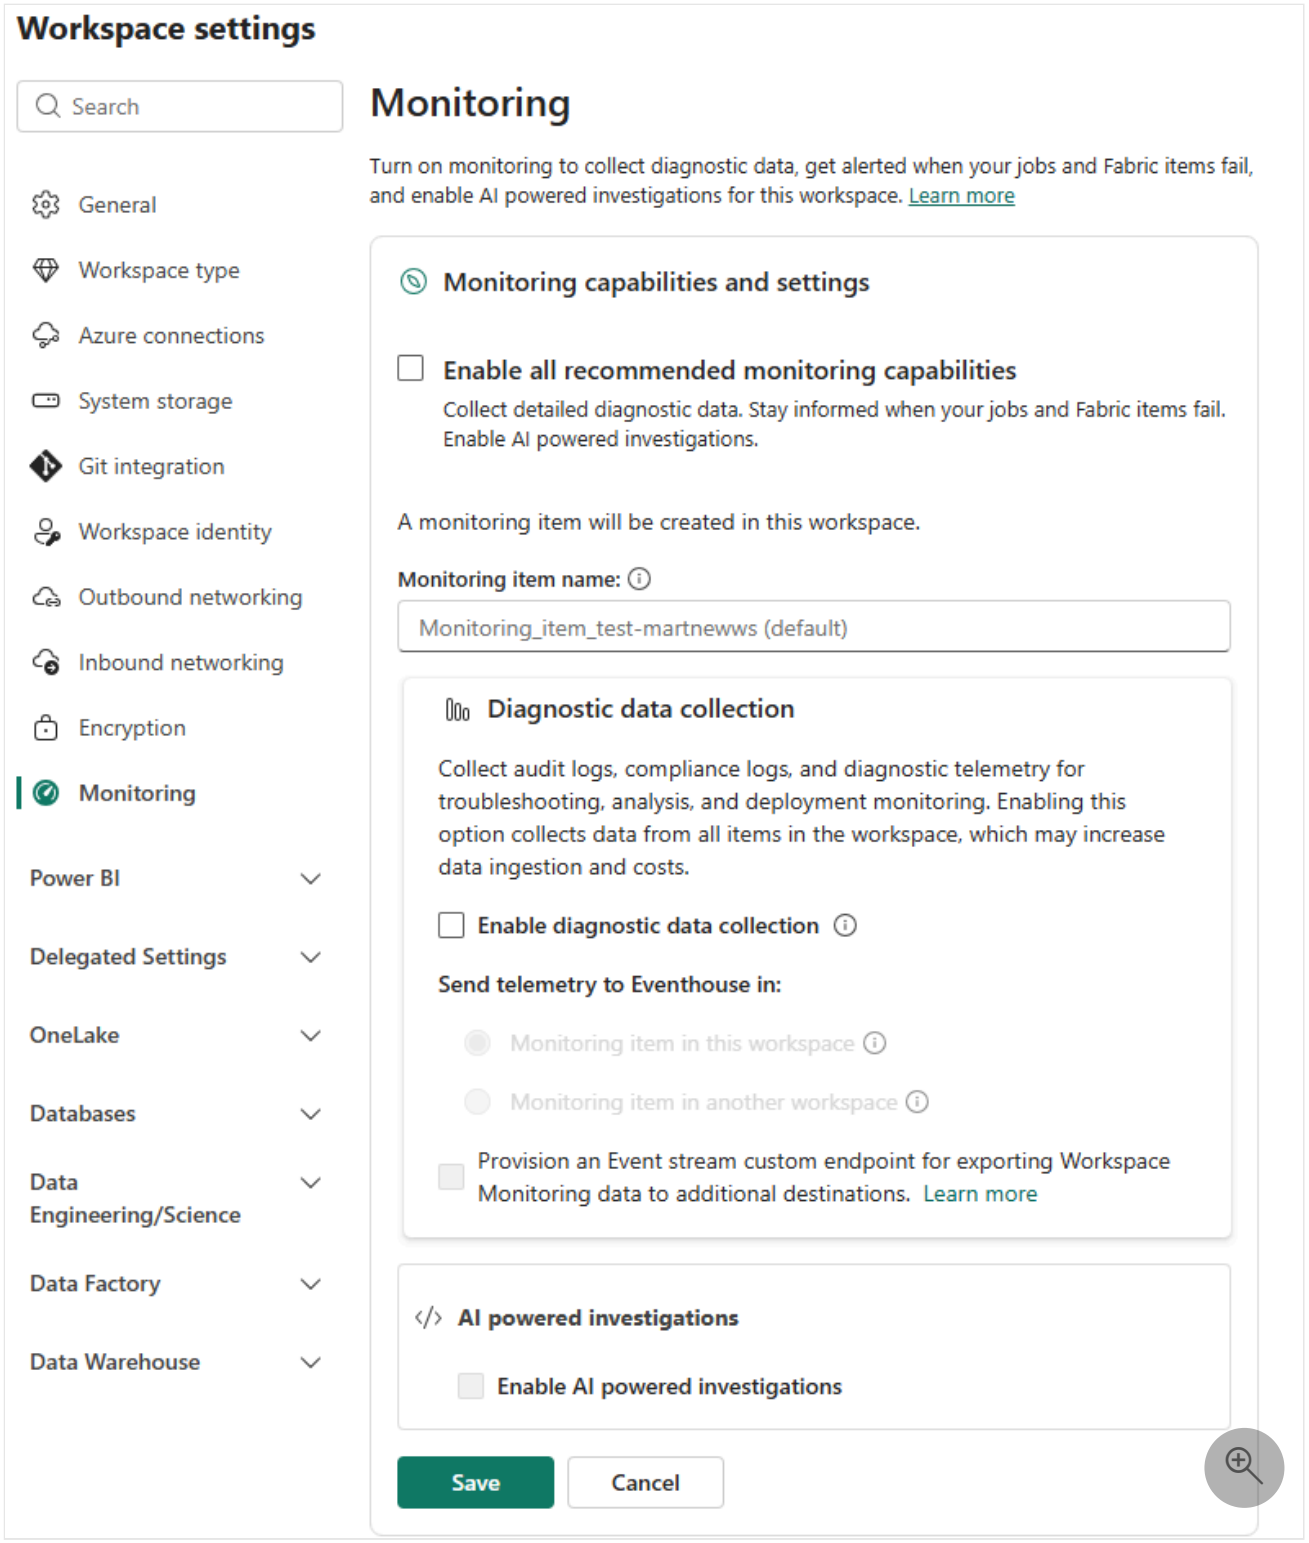

6. Select **Save**.

Create an alert rule

To set up failure notifications for an item, on the **Job runs** page, select the ellipses (...) next to the item, and then select **Create and manage alerts** > **Alerts (paid)**. The rule specifies the job events that trigger the alert, the recipients of the alert, and any automated actions to take when the alert is triggered.

Manage job alerts

You can manage both notifications for failed scheduled jobs and Activator-based alerts from the at-scale alert management experience in the Monitor hub. This experience gives you a single place to review and adjust alerts across your jobs. You can also manage Activator-based alerts from Real-Time hub and Activator.

Related content

- What is the Monitor hub?
- Monitor jobs in the Monitor hub
- Job Scheduler in Microsoft Fabric

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Monitor applications in the Monitor hub

The Monitor hub provides a guided way to monitor the health and usage of your applications in Microsoft Fabric. Use it to catch service issues early, track usage trends, and spot regressions after a release—without setting up monitoring for each application yourself.

This article shows you how to review baseline health and usage signals for an application, how the available metrics map to common monitoring scenarios, and how to reach richer diagnostics by onboarding to full observability.

> **Important**
>
> Application monitoring in the Monitor hub is currently in preview. See the Supplemental Terms of Use for Microsoft Azure Previews for legal terms that apply to Azure features that are in beta, preview, or otherwise not yet released into general availability.

Monitoring tiers

Monitoring comes in two tiers:

- **Always-on platform observability (default)**: Available out of the box, with no extra onboarding. This tier provides platform metrics for quick health checks, trend monitoring, and basic issue detection.
- **Full observability (onboarding-based)**: Available when you enable Workspace monitoring and advanced telemetry onboarding. This tier adds rich logs and traces, deeper dimensions and diagnostics, broader troubleshooting workflows, and more advanced alerting and root-cause analysis.

Prerequisites

Before you begin, make sure that you have the following:

- Access to Fabric and permission to open the Monitor hub.
- An application or workload that you can view in the Monitor hub.
- For full observability, Workspace monitoring and advanced telemetry onboarding enabled for your workspace.

Get started with monitoring

The default, always-on experience is low-friction and immediately useful. To start monitoring an application:

1. Sign in to Microsoft Fabric.
2. Open **Monitor** from the navigation pane. Switch to the new the Monitor hub experience if you're still in the classic view.
3. Select **Applications**.
4. Use the filters and search box to find your application.
5. Select the application to open the application experience view.
6. On the **Dashboard**, review the health signal metrics:
    - Sign-ins count
    - Query count
    - Query error count
    - Query error rate
    - Average duration
    - App load
7. Use the charts to identify regressions, then drill into deeper diagnostics if your workspace has full observability onboarding.

Map metrics to monitoring scenarios

The always-on metrics help you address common monitoring scenarios. Use the following table to identify the relevant metrics for each scenario.

Expand table

Scenario	Metrics to use
Is healthy right now?	Query error count, query error rate, and average duration trends.
Is usage stable and growing as expected?	Sign-ins count, app load, and query count.
Are regressions happening after release windows?	Correlate changes in query errors, query error rate, and average duration with deployment timing.

Considerations and limitations

The following limitations apply to the preview:

- Scope is intentionally backend-first. The preview defers some frontend and function-level scenarios.
- The always-on tier doesn't fully expose advanced trace-level workflows. Enable full observability through Workspace monitoring and advanced telemetry onboarding to access rich logs and traces.

Related content

- What is the Monitor hub?
- Monitor jobs in the Monitor hub

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Manage capacities in the Monitor hub (preview)

The Monitor hub gives Fabric administrators one place to monitor the health, capacity unit (CU) consumption, and throttling of every capacity in the tenant. From the **Manage capacities** page, you can spot capacities that are at risk, drill into a capacity's utilization and activity metrics, take action such as pausing or resizing a capacity, and review active issues.

> **Important**
>
> Monitor capacity in the Monitor hub is currently in preview. See the Supplemental Terms of Use for Microsoft Azure Previews for legal terms that apply to Azure features that are in beta, preview, or otherwise not yet released into general availability.

Prerequisites

- A Microsoft Fabric capacity in your tenant.
- Fabric administrator or capacity administrator permissions to view capacity details and take management actions.

Open capacity monitoring

Start from the Fabric portal and open the capacity monitoring view in the Monitor hub.

1. Sign in to the Fabric portal.
2. In the left navigation pane, select **Monitor**.
3. Select **Capacities**.

The **Manage capacities** page opens and shows all capacities for your tenant.

View capacity health across your tenant

The **Manage capacities** page summarizes the health of every capacity. Health tiles at the top show the total number of capacities and how many are **Healthy**, **At risk**, **Degraded**, or **Paused**.

To find and organize capacities, use the following controls:

- Use **Filter by keyword** to filter the list by the text you type.
- Select the **Filter** menu to narrow the list by capacity properties.
- Select **Sort by health state** to order capacities by health.
- Select **Column options** to choose which columns appear in table view.

Switch between two layouts with the view toggle:

- **Grid view** shows each capacity as a tile with its health status, **Current CU %**, and **Throttling**.

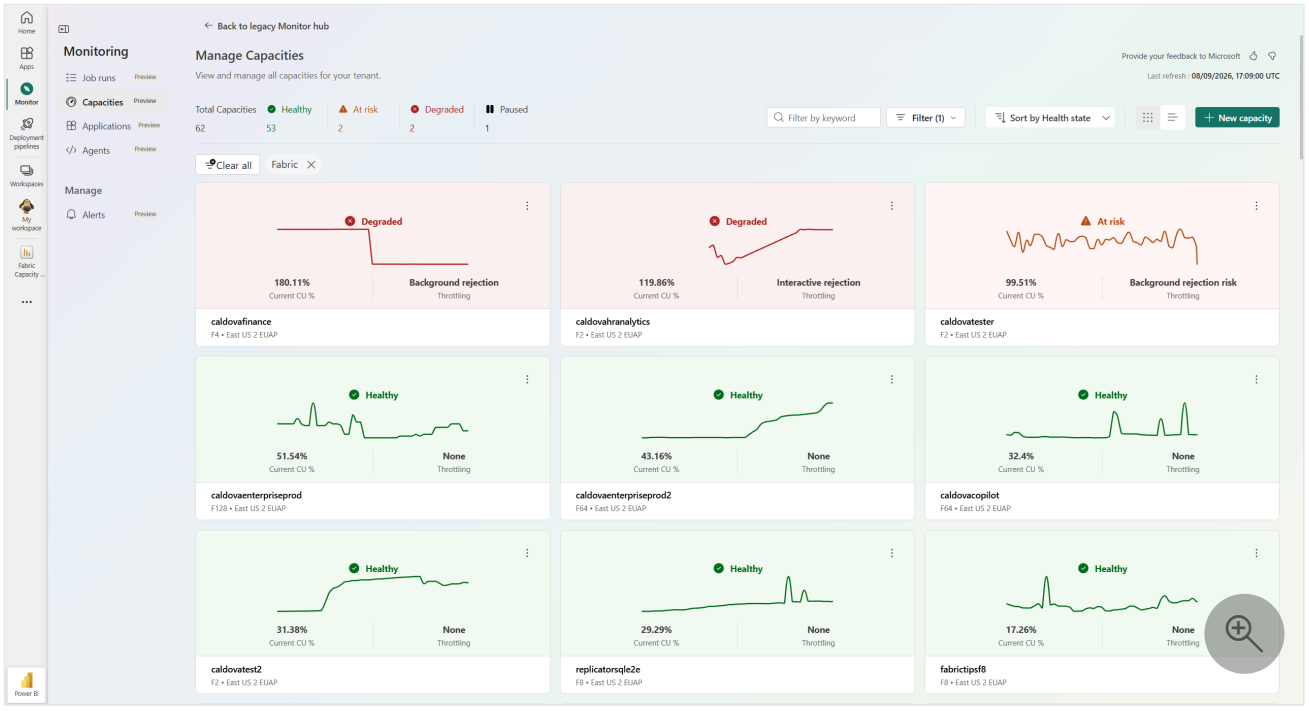

- **Table view** shows a list with the columns **Capacity Name**, **Current CU %**, **Throttling**, **Status**, **Size**, and **Region**.

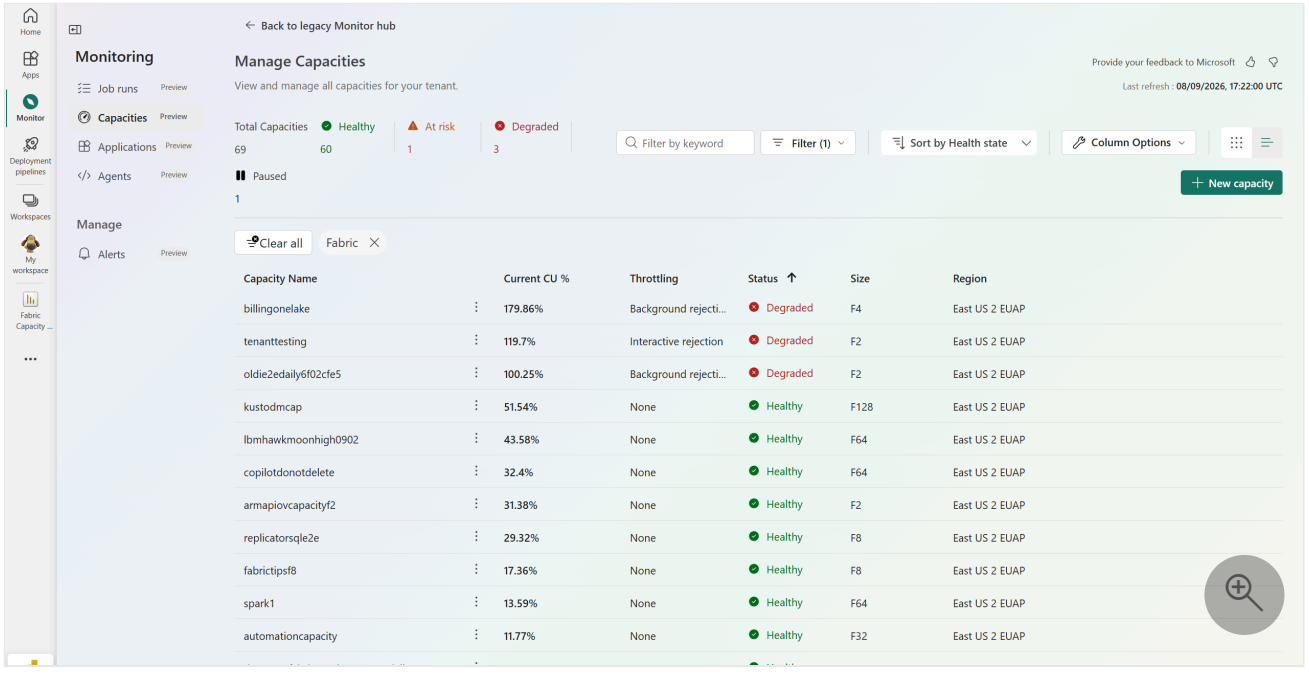

Create a new capacity

1. Sign in to the Fabric portal.
2. In the left navigation pane, select **Monitor**.
3. Select **Capacities**.
4. Select **New capacity**.
5. On the **New capacity** page, select a capacity type, and then enter the following information:

    - **Capacity type** - Select the type of capacity you want to create.
    - **Capacity name** - Enter a name for your capacity.
    - **Subscription** - Select the Azure subscription you want to use for the capacity.
    - **Resource group** - Select the Azure resource group you want to use for the capacity.
    - **Region** - Select the region where you want to create the capacity.
    - **Capacity size** - Select the number of capacity units to assign to the capacity.
    - **Capacity admins** - Select the admins for the new capacity.
    - Expand **Additional settings** to configure other settings if available, such as tags.

6. Select **Apply**.

View details for a capacity

On the **Manage capacities** page, select a capacity name to open its details page. The details page shows identifying information for the capacity: **SKU**, **Region**, **ID**, **Service admins**, and the number of **Assigned workspaces**.

Metric cards summarize the capacity's current state:

- **Utilization** shows total compute usage, split into interactive and background usage.
- **Current throttling** shows whether the capacity is throttled.
- **Current activity** shows interactive delay, interactive rejection, and background rejection rates.
- **Carry forward CUs** shows the capacity units carried forward.
- **Surge protection** shows the background rejection and recovery thresholds. Select **Configure** to adjust them.

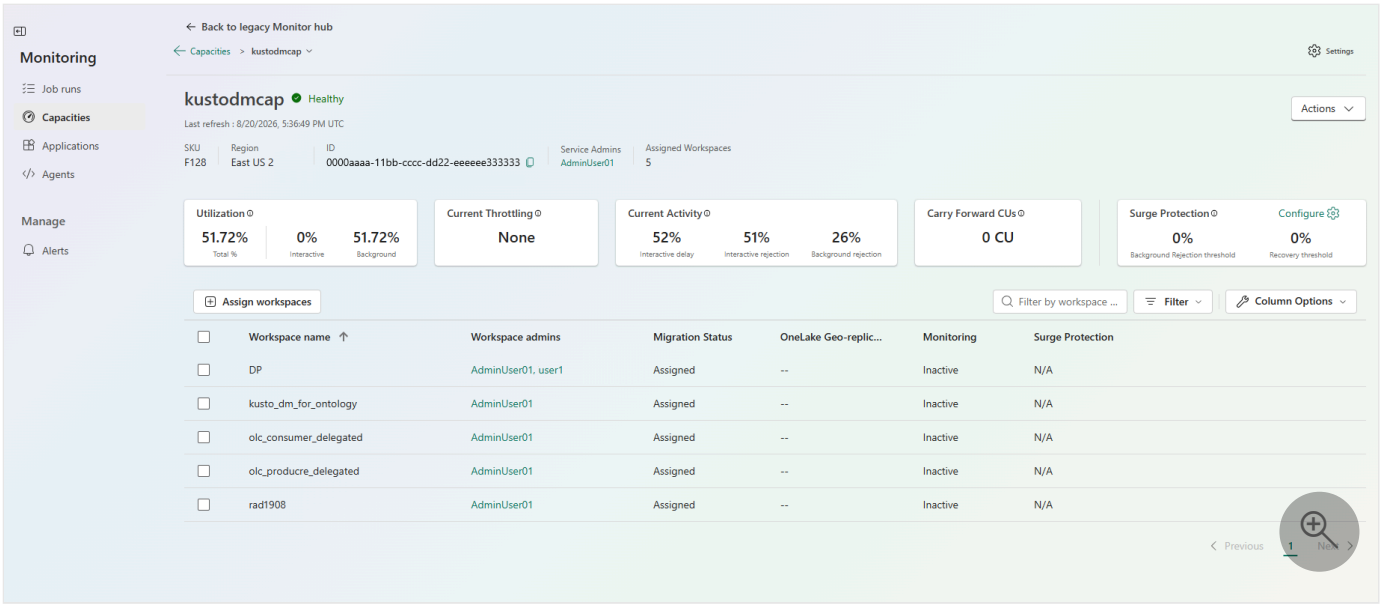

The list of workspaces that are assigned to the capacity shows details about each workspace, including the **Workspace admins**, **Migration status**, **OneLake geo-replication**, **Monitoring**, and **Surge protection** state. Select **Assign workspaces** to add workspaces to the capacity.

Take action on a capacity

On a capacity's details page, select **Actions** to manage the capacity. The following actions are available:

- **Pause** pauses the capacity.
- **Resize** changes the capacity SKU. To resize a capacity, you need Azure subscription owner permissions.
- **Reassign workspaces** moves workspaces to another capacity.
- **Configure surge protection** sets the background rejection and recovery thresholds.
- **Set capacity overage limit** limits how much overage the capacity can accrue.
- **Delete** deletes the capacity.

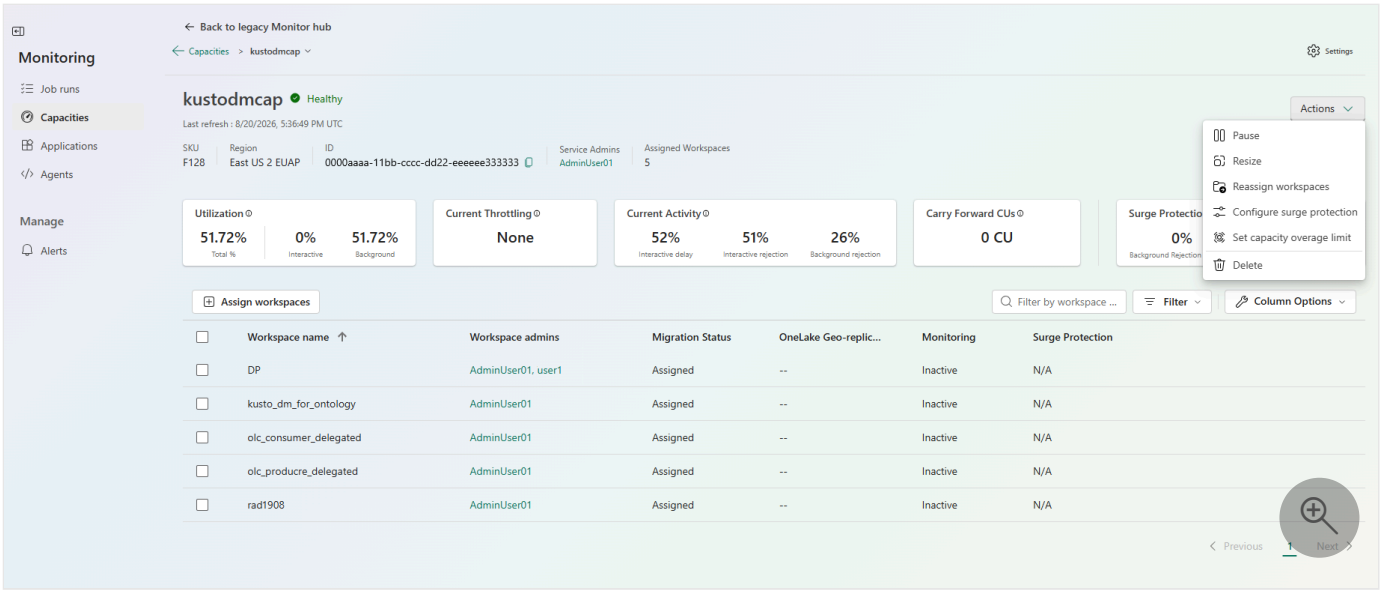

Related content

- Use the monitoring hub to track Fabric activity
- Monitor jobs in the Monitor hub

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Monitor agents in the Monitor hub (preview)

The **Agents** page in the Monitor hub gives you a centralized view of agent activity and health across your Microsoft Fabric items. Use it to quickly identify, filter, and resolve agent issues.

Any Fabric user can open the Monitor hub, but you only see activities for Fabric items you have permission to view.

> **Important**
>
> Monitor agents in the Monitor hub is currently in preview.
>
> See the Supplemental Terms of Use for Microsoft Azure Previews for legal terms that apply to Azure features that are in beta, preview, or otherwise not yet released into general availability.

Agent monitoring doesn't require any opt-in. All columns on the **Agents** page are available automatically, based on metrics collected when you create an agent.

Prerequisites

Before you begin, ensure that you have the following prerequisites:

- Access to Fabric and permission to open the Monitor hub.
- At least one Data Agent or Operations agent that you have permission to view.

Open agent monitoring

To view agent activity in your workspace:

1. Sign in to Microsoft Fabric.
2. Open **Monitor** from the navigation pane. Switch to the new Monitor hub experience if you're still in the classic view.
3. Select **Agents**. The page shows a list of agents, along with activity and health metrics for each one. Hover over a column header to see a tooltip that explains what the metric means.

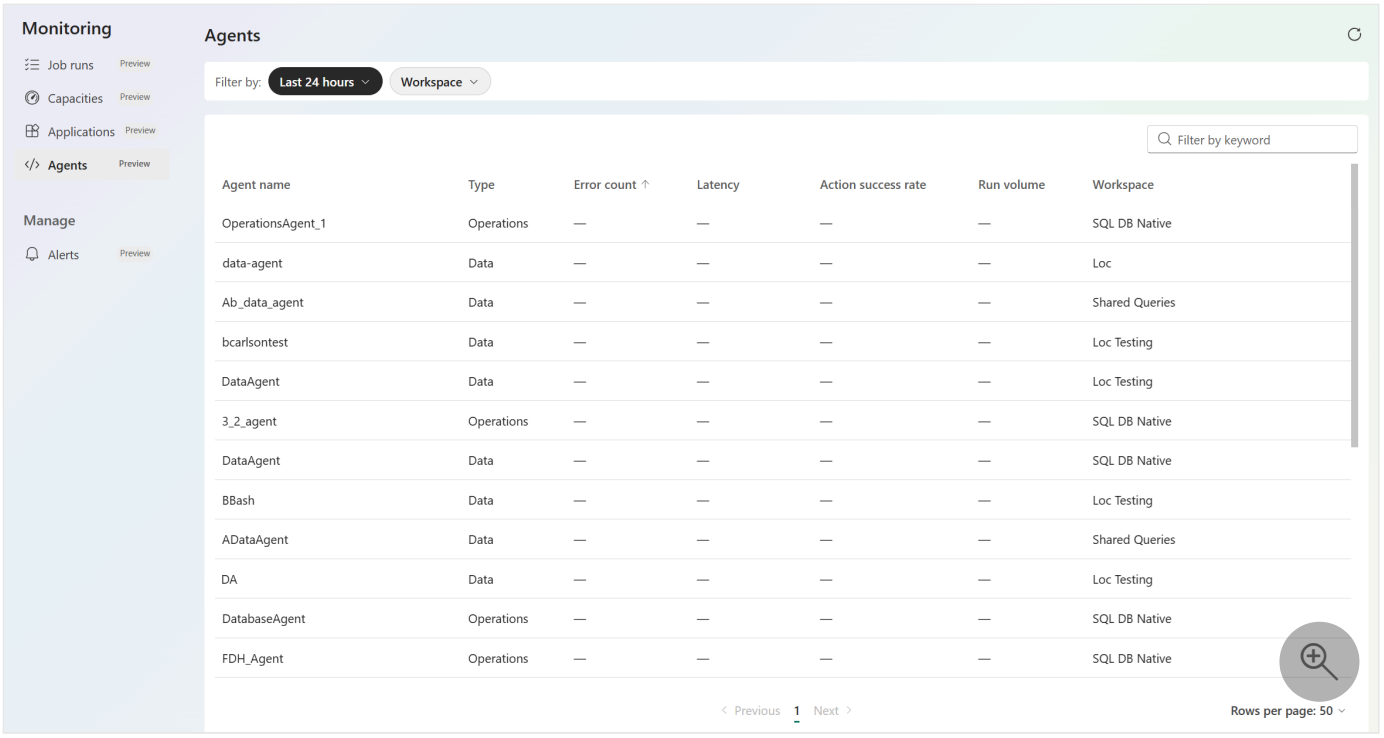

Sort and filter the Agents view

Use the following options to find the agents you're interested in.

- **Sort and customize columns:** Select a column header to sort the list. Use **Manage columns** to add, remove, or reorder columns. Hover over a column header to see a tooltip describing what that metric measures.
- **Use the filter drop-down menus:** Filter by **Item type** or **Workspace** to find a group of agents.
- **Filter by keyword:** Type text in the search box to list only those agents that contain the text.

  > **Tip**
  >
  > Keyword search checks only items loaded on the current page. The filter drop-down menus don't have this limitation. If your items span multiple pages, use the drop-down menus to narrow the list first.

- **Refresh the list:** The page retrieves agent data when it loads and doesn't update continuously. Select the **Reload** icon in the upper right to see the latest activity and health.

Get agent details

Selecting an agent in the list takes you to a different destination depending on the agent type:

- **Operations agent:** Select the agent to open the item-specific **Performance** view, where you can review detailed activity and health information for that agent.
- **Data Agent:** Select the agent to open the Data Agent item itself. There's currently no dedicated monitoring view for Data Agent, so you're taken back to the item instead.

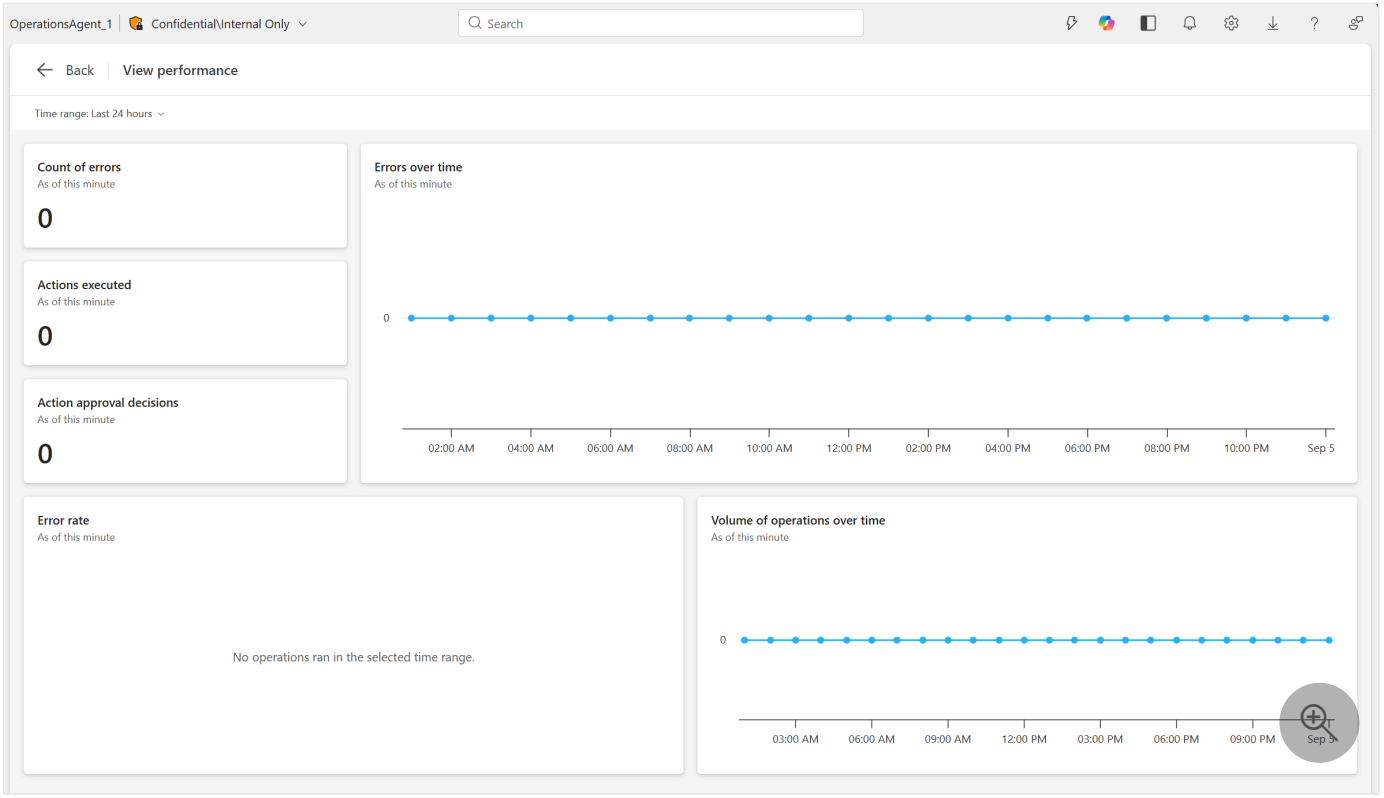

Supported item types in the Agents page

The **Agents** page currently supports these Fabric items:

- Data agent
- Operations agent

Support for additional agent types will expand over time to cover any agent you build on Fabric.

Related content

- Use the monitoring hub to track Fabric activity
- Monitor jobs in the Monitor hub

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

View audit logs

Article • 02/28/2024

Go to the Fabric admin portal to access this feature. For information about how to get to and use the admin portal, see What is the admin portal?

You manage Power BI audit logs in the Microsoft Purview compliance portal. The Audit logs tab provides a link to the Microsoft Purview compliance portal. To learn more, see Track user activities in Power BI.

Related content

- What is the admin portal?

Feedback

Was this page helpful? Yes No

Provide product feedback | Ask the community

Track user activities in Microsoft Fabric

Summarize this article for me

Knowing who is taking what action on which item in Microsoft Fabric, can be critical in helping your organization fulfill requirements such as meeting regulatory compliance and records management. This article discusses tracking user activities using the audit log.

Prerequisites

You must be assigned the Audit Logs role in Exchange Online to access the audit log. By default, the Compliance Management and Organization Management role groups have roles assigned on the **Admin roles** page in the Exchange admin center. For more information about the roles that can view audit logs, see Requirements to search the audit log.

Access

To access the audit logs, go to the Microsoft Purview portal.

Search the audit logs

You can search the audit logs using the Microsoft Purview filters. When you combine filters, the search results show only items that match all of the filter criteria. This section lists some of the available filters. For more information, review the Microsoft Purview documentation.

- **Date and time range** - Search the logs by date range using the *Start date* and *End date* fields. The default selection is the past seven days. The display presents the date and time in UTC format. For the maximum date range supported, please visit Microsoft Purview documentation.
- **Activities** - Your search returns the selected activities.
- **Users** - Search for activities performed by specific users. Enter one or more user names in the *Users* field. User names appear in email address format. Leave blank to return entries for all users (and service accounts) in your organization.
- **File, folder, or site** - Search by file name, folder name, or URL.

You can also use PowerShell to view audit logs. TO use PowerShell, Connect to Exchange Online PowerShell. You can also use the blog post Using Power BI Audit Log and PowerShell to assign Power BI Pro workspace types as a reference.

Audit log list

Operation list includes a list of all the audit log entries.

Considerations and limitations

When capacity ID and capacity name aren't available in the audit logs, you can view them in the Microsoft Fabric Capacity Metrics app.

Related content

- Operation list

Last updated on 02/13/2026

Operation list

The following operations are available in the audit logs.

To audit OneLake data access and storage operations, use OneLake diagnostics. The Fabric audit log isn't a complete source for OneLake data-plane activity.

Expand table

Friendly name	Operation name	Notes
Accept an invitation to external data share	AcceptExternalDataShare	Accept an invitation to external data
Activated policy set	ActivatedPolicySet	
Activated Spark Custom Live Pool	ActivatedSparkCustomLivePool	Whether a custom live pool activati
Add Admin Personal Workspace Access	AddAdminPersonalWorkspaceAccess	
Add Artifact To Pipeline	AddArtifactToPipeline	
Add Experiment Run	AddExperimentRun	
Add Model Version	AddModelVersion	
Add Tenant Key	AddTenantKey	
Add Tile	AddTile	Add Tile is dashboard activity, which as tiles to a Power BI dashboard
Add workload in a tenant or capacity	ExtensibilityActivationDynamic	Add a workload in a tenant or capac
Add workload version in a tenant	ExtensibilityActivationStatic	Add a specific version of a workload
Add workspace role	AddWorkspaceRoleViaAdminApi	Add workspace role
Added Power BI folder access	AddFolderAccess	Not currently used
Added Power BI group members	AddGroupMembers	
Added data source to Power BI gateway	AddDatasourceToGateway	
Added external resource	AddExternalResource	Registers or connects an external re
Added link to external resource	AddLinkToExternalResource	Links a Fabric item to an external re integration.
Added runtime lineage event subscription	AddedRuntimeLineageEventSubscription	
Added user to Power BI gateway cluster	AddUsersToGatewayCluster	Not currently used
Added user to Power BI gateway cluster datasource	AddUsersToGatewayClusterDatasource	
Admin Get Artifacts Byuser Id	AdminGetArtifactsByuserId	
Admin attached dataflow storage account to tenant	AdminAttachedDataflowStorageAccountToTenant	Not currently used
Analyze In Excel Dataset	AnalyzeInExcelDataset	
Analyze In Excel Report	AnalyzeInExcelReport	

Friendly name	Operation name	Notes
Analyzed Power BI report	AnalyzeInExcel	Generated when a user selects Analy semantic model in the service and s workbook
Analyzed Power BI semantic model	AnalyzedByExternalApplication	Generated when users interact with
Applied a PostgreSQL database schema	ApplyPostgreSQLDatabaseSchema	Generated when a user applies (plar database schema through the custo (pgschema-based ALM flow). The au operation result, and the affected Po
Applied a change to model in Power BI	ApplyChangeToPowerBIModel	A user makes a change to an existin any edit is made to the model (ex: w relationships, etc.)
Applied sensitivity label to Power BI item	SensitivityLabelApplied	
Apply CMK encryption to your workspace	ApplyWorkspaceEncryption	Triggered when CMK encryption is a
Assign Workspace To Alm Pipeline	AssignWorkspaceToAlmPipeline	
Assigned a workspace to a deployment pipeline	AssignWorkspaceToPipeline	Not currently used
Attach Notebook Environment	AttachNotebookEnvironment	Attach the environment for Noteboo
Attach source in GraphQL artifact	AttachSourceGraphQL	Attach source in graphql artifact
Attached dataflow storage account	AttachedDataflowStorageAccount	
Auto bind user credentials to git	AutoBindGitCredentials	Auto bind existing user credentials f connection
Binded Power BI semantic model to gateway	BindToGateway	
Binded monikers to Power BI datasources	BindMonikersToDatasources	
Branch Out in Git	BranchOutInGit	Branch out in Git is a workspace act new workspace is created, which is c branch based on the git branch that workspace.
Browsed PostgreSQL database objects	PgSQLDbObjectExplorer	Generated when a user browses Pos objects through Object Explorer. Th identity, operation result, and the af artifact.
Cancel Datamart Batch	CancelDatamartBatch	
Cancel Dataset Refresh	CancelDatasetRefresh	
Cancel Publish Environment	CancelPublishEnvironment	Cancel publish environment
Cancel Running Artifact	CancelRunningArtifact	
Cancel Spark Application	CancelSparkApplication	
Cancel Sql Analytics Endpoint Lakehouse Batch	CancelSqlAnalyticsEndpointLakehouseBatch	Canceled a lakehouse SQL analytics named *Canceled a default warehous* name: *CancelDefaultWarehouseBat
Cancel Warehouse Batch	CancelWarehouseBatch	
Cancel Workspace Upgrades As Admin	CancelWorkspaceUpgradesAsAdmin	
Cancel mounted warehouse	CancelMountedWarehouseBatch	Generated when a mounted wareho

Friendly name	Operation name	Notes
batch		
Canceled a Digital Operations Ontology Agent conversation	DigitalOperationsOntologyAgentConversationCancelled	A user canceled a Digital Operations
Canceled Power BI dataflow refresh	CancelDataflowRefresh	
Changed Power BI gateway admins	ChangeGatewayAdministrators	
Changed Power BI gateway data source users	ChangeGatewayDatasourceUsers	
Changed Power BI semantic model connections	SetAllConnections	
Changed capacity state	ChangeCapacityState	
Changed capacity user assignment	UpdateCapacityUsersAssignment	
Changed sensitivity label for Power BI item	SensitivityLabelChanged	
Check Scorecard Access	CheckScorecardAccess	Verifies the access permissions for a
Checkout Branch In Git	CheckoutBranchInGit	
Cleanup unused or corrupted files from dataflows' refreshes	CleanupDataflow	Cleanup unused or corrupted files fr
Clone Tile	CloneTile	
Co Author Notebook	CoAuthorNotebook	
Commit Notebook	CommitNotebook	
Commit To Git	CommitToGit	Commit To Git is an artifact activity, commit artifact to Git.
Completed an artifact access request action in Power BI	ArtifactAccessRequest	Tracks the notification sent to appro or denied.
Compute item sizes for workspace	ComputeItemsSize	Calculates the amount of data store to display in the item size report
Connect To Git	ConnectToGit	Connect To Git is a workspace activi users connect the workspace to Git.
Connected to Power BI semantic model from external app	ConnectFromExternalApplication	
Connected to a Warehouse or SQL analytics endpoint from an external app	ConnectWarehouseAndSqlAnalyticsEndpointLakehouseFromExternalApp	Previously named *Connected to a wa from an external app* (Operation name: *ConnectWarehouseAndDefa
Convert Personal Workspace To Workspace	ConvertPersonalWorkspaceToWorkspace	
Copied Power BI dashboard	CopyDashboard	
Copied Power BI report	CopyReport	
Copilot Interaction (Request Copilot features in Fabric)	CopilotInteraction	Request Copilot features in Fabric
Copy Scorecard	CopyScorecard	
Create an SQL query from a Warehouse	CreateSqlQueryFromWarehouse	
Create Artifact	CreateArtifact	

Friendly name	Operation name	Notes
Create Branch In Git	CreateBranchInGit	Create Branch In Git is a git provide when users create branch in Git.
Create Cloud Datasource From Kind Path	CreateCloudDatasourceFromKindPath	Creates a cloud-based data source u path.
Create Cloud Datasource	CreateCloudDatasource	Establishes a new cloud data connec services.
Create Data Sharing	CreateDataSharing	Create a new external data share
Create Datamart	CreateDatamart	Triggered when a new Datamart is c
Create Dataset By Quick Share	CreateDatasetByQuickShare	
Create Directory In Git	CreateDirectoryInGit	Create Directory In Git is a git provic when users create directory in Git.
Create Environment Resource	CreateEnvironmentResource	Create resources in environment
Create Fabric Identity for Workspace	CreateWorkspaceIdentityViaApi	Create a Fabric identity for a worksp
Create Folder	CreateFolder	Indicates a creation of a new worksp
Create Gateway Cluster User	CreateGatewayClusterUser	
Create Gateway Tenant Key	CreateGatewayTenantKey	
Create Goal Value Categories	CreateGoalValueCategories	
Create Hierarchy Goal Value	CreateHierarchyGoalValue	
Create Hierarchy Note	CreateHierarchyNote	
Create Lakehouse File	CreateLakehouseFile	
Create Lakehouse Folder	CreateLakehouseFolder	
Create Lakehouse Shortcut	CreateLakehouseShortcutLink	
Create Lakehouse Table	CreateLakehouseTable	
Create Link Goals	CreateLinkGoals	
Create Managed Private Endpoint	CreateManagedPrivateEndpoint	Create a private endpoint within a m workspace
Create Managed VNet	CreateManagedVNet	Create managed virtual network for
Create MetricSet Metric Endpoint	CreateMetricSetMetric	Create metric in a metricset endpoir
Create Notebook Resource	CreateNotebookResource	Create resources in notebook.
Create Org App	CreateOrgApp	
Create Personal Datasource	CreatePersonalDatasource	
Create Report From Lakehouse	CreateReportFromLakehouse	
Create Scorecard Hierarchy	CreateScorecardHierarchy	
Create Scorecard View	CreateScorecardView	
Create Semantic Metric Endpoint	CreateSemanticMetric	Create standalone semantic metric e
Create Service Principal Profile	CreateServicePrincipalProfile	
Create Template App Package	CreateTemplateAppPackage	Create Template App Package is an on creation of a Template App Packa

Friendly name	Operation name	Notes
Create Temporary Upload Location	CreateTemporaryUploadLocation	Generated when a temporary uploa
Created visual query from a Warehouse	CreateVisualQueryFromWarehouse	
Create Warehouse	CreateWarehouse	
Create Workspace	CreateWorkspace	
Create a shortcut	CreateShortcut	OneLake operation. Part of Shortcut
Create an external data share	CreateExternalDataShare	Create an external data share
Create file	CreateFile	OneLake operation. Creates a new f
Create gateway cluster data source from kind path JSON	CreateGatewayClusterDatasourceFromKindPath	Create gateway cluster data source
Create or update data access roles	CreateOrUpdateDataAccessRoles	Creates or updates data access role
Create subfolder	CreateSubfolder	
Create task flow	CreateTaskFlow	
Create virtual network data gateway proxy	CreateVirtualNetworkDataGatewayProxy	Create HTTP proxy for virtual netwo
Created Artifact Definition	CreatedArtifactDefinition	OneLake operation. Generated whe artifact definition is created (shared MirroredStorage and AWS Databrick
Created AWS Databricks Catalog	CreatedAWSDatabricksCatalog	OneLake operation. Generated whe Catalog item is created, capturing th workspace ConnectionId and the ini (catalogs/schemas/tables selected f
Created MirroredStorage	CreatedMirroredStorage	OneLake operation. Generated whe creates a Fabric MirroredStorage ite source to a workspace as OneLake s
Created Power BI app	CreateApp	
Created Power BI dashboard	CreateDashboard	
Created Power BI dataflow	CreateDataflow	
Created Power BI email subscription	CreateEmailSubscription	
Created Power BI gateway cluster datasource	CreateGatewayClusterDatasource	
Created Power BI gateway	CreateGateway	
Created Power BI group	CreateGroup	
Created Power BI report	CreateReport	
Created Power BI semantic model from external app	CreateDatasetFromExternalApplication	
Created Power BI semantic model	CreateDataset	
Created a Power BI metric value	CreateGoalValue	
Created a Power BI scorecard metric	CreateGoal	
Created a Power BI scorecard	CreateScorecard	
Created a visual query from	CreateVisualQueryFromSqlAnalyticsEndpointLakehouse	

Friendly name	Operation name	Notes
an SQL analytics endpoint		
Created a Power BI template or a workspace for a template app	CreateTemplateApp	
Created an install ticket for installing Power BI template app	CreateTemplateAppInstallTicket	
Created an organizational custom visual	InsertOrganizationalGalleryItem	
Created an SQL query from a SQL analytics endpoint	CreateSqlQueryFromSqlAnalyticsEndpointLakehouse	
Created deployment pipeline	CreateAlmPipeline	
Created Fabric policy rule	CreatedFabricPolicyRule	
Created a Digital Operations Ontology Agent conversation	DigitalOperationsOntologyAgentConversationCreated	A user created a Digital Operations
Created workload assignment using admin api	ExtensibilityCreatedWorkloadAssignmentAdminApi	
Custom visual requested Azure AD access token	GenerateCustomVisualAADAccessToken	
Custom visual requested Office Web Apps access token	CustomVisualWACAccessToken	Not currently used
D L P Info	DLPInfo	
D L P Rule Match	DLPRuleMatch	
D L P Rule Undo	DLPRuleUndo	
Dataflow migrated to external storage account	DataflowMigratedToExternalStorageAccount	Not currently used
Dataflow permissions added	DataflowPermissionsAdded	Not currently used
Dataflow permissions removed	DataflowPermissionsRemoved	Not currently used
Deactivated policy set	DeactivatedPolicySet	
Delete Alm Pipeline Access As Admin	DeleteAlmPipelineAccessAsAdmin	
Delete Alm Pipeline Access	DeleteAlmPipelineAccess	
Delete Artifact	DeleteArtifact	Generated when an item is deleted. in tenant admin settings, this operat user deletes the item. When item re is logged when the item is permane delete retention period expires or w from the recycle bin); the initial user `ArtifactSoftDeleted`.
Delete Capacity Delegation settings	DeleteCapacityTenantSettingDelegation	Delete Capacity delegation settings.
Delete Datamart	DeleteDatamart	
Delete Domain	DeleteDataDomainAsAdmin	Delete Domain
Delete Domain's Folder Relation As Folder Owner	DeleteDataDomainFolderRelationsAsFolderOwner	Delete Domain's Folder Relation As
Delete Environment Resource	DeleteEnvironmentResource	Delete resources in environment

Friendly name	Operation name	Notes
Delete Experiment Run	DeleteExperimentRun	
Delete Fabric Identity for Workspace	DeleteWorkspaceIdentityViaApi	Delete a Fabric identity for a worksp
Delete Folder	DeleteFolder	Indicated a deletion of a workspace
Delete Goal Current Value Rollup	DeleteGoalCurrentValueRollup	
Delete Goal Status Rules	DeleteGoalStatusRules	
Delete Goal Target Value Connection	DeleteGoalTargetValueConnection	
Delete Goal Target Value Rollup	DeleteGoalTargetValueRollup	
Delete Goal Value Categories	DeleteGoalValueCategories	
Delete Goal Value	DeleteGoalValue	
Delete Goals	DeleteGoals	
Delete Group Workspace	DeleteGroupWorkspace	
Delete Hierarchy Goal Value	DeleteHierarchyGoalValue	
Delete Hierarchy Note	DeleteHierarchyNote	
Delete Lakehouse File	DeleteLakehouseFile	
Delete Lakehouse Folder	DeleteLakehouseFolder	
Delete Lakehouse Table	DeleteLakehouseTable	
Delete Link Goals	DeleteLinkGoals	
Delete Managed Private Endpoint	DeleteManagedPrivateEndpoint	Delete a private endpoint within a n workspace
Delete Managed VNet	DeleteManagedVNet	Delete managedvirtual network for
Delete MetricSet Metric Endpoint	DeleteMetricSetMetric	Delete metric in a metricset endpoi
Delete Model Version	DeleteModelVersion	Triggered when a Machine Learning
Delete Notebook Resource	DeleteNotebookResource	Update resources in notebook.
Delete Scorecard Hierarchy	DeleteScorecardHierarchy	Triggered when a scorecard hierarch hierarchy's metadata is deleted, not
Delete Scorecard View	DeleteScorecardView	Triggered when a scorecard view (lis deleted.
Delete Semantic Metric Endpoint	DeleteSemanticMetric	Delete standalone semantic metric e
Delete Service Principal Profile As Admin	DeleteServicePrincipalProfileAsAdmin	
Delete Service Principal Profile	DeleteServicePrincipalProfile	
Delete Template App Package	DeleteTemplateAppPackage	Delete Template App Package is an on deletion of a Template App pack
Delete Tile	DeleteTile	Delete Tile is a dashboard activity, w tiles from a Power BI dashboard
Delete Warehouse	DeleteWarehouse	
Delete Workspace Via Admin Api	DeleteWorkspaceViaAdminApi	
Delete a shortcut	DeleteShortcut	OneLake operation. Part of Shortcut

Friendly name	Operation name	Notes
Delete admin monitoring folder via lockbox	DeleteAdminMonitoringFolderViaLockbox	
Delete admin usage dashboards via lockbox	DeleteAdminUsageDashboardsViaLockbox	
Delete all Domain's Folders Relations	DeleteAllDataDomainFoldersRelationsAsAdmin	Delete all Domain's Folders Relation
Delete blob	DeleteBlob	OneLake operation. Deletes a folder
Delete configured credentials	DeleteGitProviderCredentials	Delete configured git provider crede
Delete file	DeleteFileOrBlob	OneLake operation. Deletes a folder
Delete source in GraphQL artifact	DeleteSourceGraphQL	Delete source in graphql item
Delete subfolder	DeleteSubfolder	
Delete usage metrics v2 package via lockbox	DeleteUsageMetricsv2PackageViaLockbox	
Delete virtual network data gateway proxy	DeleteVirtualNetworkDataGatewayProxy	Delete HTTP proxy for virtual netwo
Delete workload package	ExtensibilityDeletePackage	Delete a workload package
Delete workspace role	DeleteWorkspaceRoleViaAdminApi	Delete workspace role
Delete Workspace Delegation settings	DeleteWorkspaceTenantSettingDelegation	Delete workspace delegation setting
Deleted Power BI comment	DeleteComment	
Deleted Power BI dashboard	DeleteDashboard	
Deleted Power BI dataflow	DeleteDataflow	
Deleted Power BI email subscription	DeleteEmailSubscription	
Deleted Power BI folder access	DeleteFolderAccess	Not currently used
Deleted Power BI gateway cluster datasource	DeleteGatewayClusterDatasource	
Deleted Power BI gateway cluster	DeleteGatewayCluster	
Deleted Power BI gateway	DeleteGateway	
Deleted Power BI group	DeleteGroup	
Deleted Power BI metric	DeleteGoal	
Deleted Power BI note	DeleteNote	
Deleted Power BI report	DeleteReport	
Deleted Power BI scorecard	DeleteScorecard	
Deleted Power BI semantic model from external app	DeleteDatasetFromExternalApplication	
Deleted Power BI semantic model rows	DeleteDatasetRows	Indicates that the Push Datasets - D called
Deleted Power BI semantic model	DeleteDataset	
Deleted Power BI template app or a workspace for a template app	DeleteTemplateApp	

Friendly name	Operation name	Notes
Deleted Publish to Web Embed Code	DeleteEmbedCode	
Deleted a SQL query from a SQL analytics endpoint	DeleteSqlQueryFromSqlAnalyticsEndpointLakehouse	This audit event covers both deletin the SQL analytics endpoint of the la
Deleted SQL query from a Warehouse	DeleteSqlQueryFromWarehouse	Covers both deleting SQL and visua
Deleted an organizational custom visual	DeleteOrganizationalGalleryItem	
Deleted current value connection of Power BI metric	DeleteGoalCurrentValueConnection	
Deleted deployment pipeline	DeleteAlmPipeline	
Deleted Fabric policy rule	DeletedFabricPolicyRule	
Deleted link to external resource	DeleteLinkToExternalResource	
Deleted member of Power BI gateway cluster	DeleteGatewayClusterMember	
Deleted organizational Power BI content pack	DeleteOrgApp	
Deleted runtime lineage event subscription	DeletedRuntimeLineageEventSubscription	
Deleted sensitivity label from Power BI item	SensitivityLabelRemoved	
Deleted snapshot for user in Power BI tenant	DeleteSnapshot	Generated when a user deletes a sn model
Deleted workload assignment using admin api	ExtensibilityDeletedWorkloadAssignmentAdminApi	
Deploy Model Version	DeployModelVersion	
Deploy user application in FunctionSet	DeployUserAppFunctionSet	Deploy user application through Fu
Deployed to a pipeline stage	DeployAlmPipeline	
Detect Customizations For Solution	DetectCustomizationsForSolution	
Determine if the user can share a datasource	DeterminePrincipalCanShareDatasource	Get the policy decision for the user
Disconnect From Git	DisconnectFromGit	Disconnect From Git is a workspace when users disconnect the workspa
Discovered Power BI semantic model data sources	GetDatasources	
Downgrade Workspace	DowngradeWorkspace	
Download Environment Resource	DownloadEnvironmentResource	Download resources in environment
Download Notebook Resource	DownloadNotebookResource	Delete resources in notebook.
Download Spark App Log	DownloadSparkAppLog	
Downloaded Power BI report	DownloadReport	

Friendly name	Operation name	Notes
Downloaded Exported Power BI artifact file	ExportArtifactDownload	An export of the .pptx or .pdf file is
Drop Lakehouse File	DropLakehouseFile	
Drop Lakehouse Folder	DropLakehouseFolder	
Drop Lakehouse Table	DropLakehouseTable	
Edit Artifact Endorsement	EditArtifactEndorsement	
Edit Report Description	EditReportDescription	
Edit Sql Analytics Endpoint Lakehouse Endorsement	EditSqlAnalyticsEndpointLakehouseEndorsement	Edited a lakehouse SQL analytics en
Edit Tile	EditTile	Edit Tile is a dashboard activity, whi edits to settings for tiles in a Power
Edit Warehouse Endorsement	EditWarehouseEndorsement	
Edit Widget Tile	EditWidgetTile	
Edit mounted warehouse endorsements	EditMountedWarehouseEndorsement	Generated when mounted warehou
Edited Power BI app endorsement	EditContentProviderProperties	
Edited Power BI certification permission	EditCertificationPermission	Not currently used
Edited Power BI dashboard	EditDashboard	Not currently used
Edited Power BI dataflow endorsement	EditDataflowProperties	
Edited Power BI report endorsement	EditReportProperties	
Edited Power BI report	EditReport	
Edited Power BI semantic model endorsement	EditDatamartEndorsement	
Edited Power BI semantic model from external app	EditDatasetFromExternalApplication	
Edited Power BI semantic model properties	EditDatasetProperties	
Edited Power BI semantic model	EditDataset	
Enable/Disable AllowedInUntrustedContexts settings	AllowedInUntrustedContextsViaApi	Enable or disable AllowedInUntruste
Encrypted credentials for Power BI gateway datasource	EncryptCredentials	
Encrypted credentials using Power BI gateway cluster	EncryptClusterCredentials	
Evaluate Diagnostics Download	EvaluateDiagnosticsDownload	
Evaluate Diagnostics Query	EvaluateDiagnosticsQuery	
Evaluate chat response based on the data gateway diagnostics data	EvaluateDiagnosticsChat	Evaluate chat response based on th
Evaluate data sources against DMTS data policies	EvaluateDataSourcesAgainstTenantDlpPolicies	Evaluate data sources against DMTS

Friendly name	Operation name	Notes
Executed a tenant relocation	TenantRelocationExecuted	Executed tenant relocation.
Explore Dataset	ExploreDataset	
Export Package For Solution	ExportPackageForSolution	
Export Power BI activity events	ExportActivityEvents	
Export Scorecard	ExportScorecard	Exports a scorecard
Exported Power BI dataflow	ExportDataflow	
Exported Power BI item to another file format	ExportArtifact	
Exported Power BI report to another file format or exported report visual data	ExportReport	
Exported Power BI tile data	ExportTile	
Exported a PostgreSQL database schema	ExportPostgreSQLDatabaseSchema	Generated when a user dumps a Po through the custom PG schema serv flow). The audit log records caller id affected PostgreSQL database artifa
Extract Template App Package	ExtractTemplateAppPackage	Extract Template App Package is an when users extract an existing temp template app workspace
Fetch a shortcut's metadata	GetShortcut	OneLake operation. Part of Shortcut
Finish Publish Environment	FinishPublishEnvironment	Finish publish environment
Follow Goal	FollowGoal	
Gateway Cluster S S O Test Connection	GatewayClusterSSOTestConnection	
Generate Custom Visual W A C Access Token	GenerateCustomVisualWACAccessToken	
Generate Multi Resource Embed Token	GenerateMultiResourceEmbedToken	
Generate screenshot	GenerateScreenshot	
Generated Power BI Embed Token	GenerateEmbedToken	
Generated Power BI dataflow SAS token	GenerateDataflowSasToken	
Get All Private Link Services For Tenant	GetAllPrivateLinkServicesForTenant	
Get All Scorecards	GetAllScorecards	
Get Alm Pipeline Users As Admin	GetAlmPipelineUsersAsAdmin	
Get Alm Pipelines As Admin	GetAlmPipelinesAsAdmin	
Get Artifact Access By User Id Via Admin Api	GetArtifactAccessByUserIdViaAdminApi	
Get Artifact Link Shared To Whole Org As Admin	GetArtifactLinkSharedToWholeOrgAsAdmin	
Get Artifact Users Via Admin Api	GetArtifactUsersViaAdminApi	
Get Artifacts By Id Via Admin Api	GetArtifactsByIdViaAdminApi	
Get Artifacts Via Admin Api	GetArtifactsViaAdminApi	

Friendly name	Operation name	Notes
Get Batch Lakehouse Table Details	GetBatchLakehouseTableDetails	
Get Cloud Supported Datasources	GetCloudSupportedDatasources	
Get Dashboards In Group As Admin	GetDashboardsInGroupAsAdmin	
Get Dataflow Users As Admin	GetDataflowUsersAsAdmin	
Get Dataflows In Group As Admin	GetDataflowsInGroupAsAdmin	
Get Dataset Info	GetDatasetInfo	Get the info of the semantic model
Get Dataset Query Scale-Out Sync Status	GetDatasetQueryScaleOutSyncStatus	Get Dataset Query Scale-Out Sync S activity, which is generated when us scale out-enabled Power BI semanti
Get Dataset Users As Admin	GetDatasetUsersAsAdmin	
Get Datasets In Group As Admin	GetDatasetsInGroupAsAdmin	
Get Datasource Details With Credentials Async	GetDatasourceDetailsWithCredentialsAsync	
Get Dax Capabilities	GetDaxCapabilities	
Get Default Scorecard View	GetDefaultScorecardView	
Get Domain Delegation settings	DeleteDomainTenantSettingDelegation	Delete Domain Delegation settings.
Get Fabric Identity Token for workspace	GetworkspaceIdentityTokenViaApi	Get a Fabric identity token for a wor
Get Fabric Identity for Workspace	GetWorkspaceIdentityViaApi	Get a Fabric identity for a workspace
Get Followed Goals	GetFollowedGoals	
Get Gateway Cluster	GetGatewayCluster	
Get Gateway Clusters With Role Options	GetGatewayClustersWithRoleOptions	
Get Gateway Container I P Configuration	GetGatewayContainerIPConfiguration	
Get Gateway Container N S Lookup Details	GetGatewayContainerNSLookupDetails	
Get Gateway Container Test Net Connection Details	GetGatewayContainerTestNetConnectionDetails	
Get Gateway datasource limit status	GetGatewayDatasourceLimitStatus	Gets the datasource limit status of t gateway
Get Goal By Hierarchy Item Ids	GetGoalByHierarchyItemIds	
Get Goal Status Rules	GetGoalStatusRules	
Get Goal Value Categories	GetGoalValueCategories	
Get Groups As Admin	GetGroupsAsAdmin	Get Groups as Admin is a workspace retrieving list of Power BI workspace
Get Hierarchy Goal Values	GetHierarchyGoalValues	
Get item sizes for workspace	GetItemsSize	Retrieves the computed sizes of all i refresh the data
Get Lakehouse Table Details	GetLakehouseTableDetails	

Friendly name	Operation name	Notes
Get MetricSet Metric Endpoint	GetMetricSetMetric	Read metric in a metricset endpoint
Get Model Diagram Layouts	GetPowerBIDataModelDiagramLayouts	Get diagram layouts when open dat
Get Model SAS Token via Lockbox	GetModelSASTokenViaLockbox	Gets the SAS Token for a given mod
Get My Goals	GetMyGoals	
Get Pending Change Status	GetPendingChangeStatus	
Get Power BI group users	GetGroupUsers	
Get Publish To Web Artifacts As Admin	GetPublishToWebArtifactsAsAdmin	
Get Relevant Measures	GetRelevantMeasures	
Get Reports As Admin	GetReportsAsAdmin	
Get Reports In Group As Admin	GetReportsInGroupAsAdmin	Get Reports in Group as Admin is a generated on retrieving the reports workspace using an API call.
Get Roles For Scorecard	GetRolesForScorecard	Gets the roles for a scorecard
Get Scorecard Additional Properties	GetScorecardAdditionalProperties	Gets the additional properties on a
Get Scorecard By Hierarchy Item Ids	GetScorecardByHierarchyItemIds	
Get Scorecard Hierarchies	GetScorecardHierarchies	
Get Scorecard Hierarchy Items	GetScorecardHierarchyItems	
Get Scorecard Hierarchy	GetScorecardHierarchy	
Get Scorecard View	GetScorecardView	
Get Scorecard Views	GetScorecardViews	
Get Scorecards As Admin	GetScorecardsAsAdmin	Get all the scorecards as an admin
Get Semantic Metric Endpoint	GetSemanticMetric	Read standalone semantic metric er
Get Service Principal Profile	GetServicePrincipalProfile	
Get Service Principal Profiles As Admin	GetServicePrincipalProfilesAsAdmin	
Get Service Principal Profiles	GetServicePrincipalProfiles	
Get Subscription Users By Dashboard Id As Admin	GetSubscriptionUsersByDashboardIdAsAdmin	
Get Subscription Users By Report Id As Admin	GetSubscriptionUsersByReportIdAsAdmin	
Get Subscriptions For User As Admin	GetSubscriptionsForUserAsAdmin	
Get Tenant Settings Via Admin Api	GetTenantSettingsViaAdminApi	
Get USEC roles for an artifact	GetArtifactRoles	OneLake operation.
Get Unused Artifacts	GetUnusedArtifacts	
Get User Datasource Limit Status	GetUserDatasourceLimitStatus	
Get User Datasources By Data Source Reference	GetUserDatasourcesByDataSourceReference	

Friendly name	Operation name	Notes
Get Virtual Network	GetVirtualNetwork	
Get Workspace Users Via Admin Api	GetWorkspaceUsersViaAdminApi	
Get Workspaces By Id Via Admin Api	GetWorkspacesByIdViaAdminApi	
Get Workspaces Via Admin Api	GetWorkspacesViaAdminApi	
Get a task flow	GetTaskFlow	
Get all connections	GetAllConnections	Get all connections
Get all metric sets Endpoint	GetAllMetricSets	Get all metric sets endpoint
Get data artifact table details	GetDataArtifactTableDetails	Get the details (for example, schema
Get datasource share policy	GetDatasourceShareTenantPolicy	Retrieve the datasource share policy
Get delegated capacity tenant setting overrides	GetCapacityDelegatedTenantSettingOverridesViaAdminApi	Get capacity delegated tenant settir
Get delegated domain tenant setting overrides	GetDomainDelegatedTenantSettingOverridesViaAdminApi	Get domain delegated tenant settin
Get delegated workspace tenant setting overrides	GetWorkspaceDelegatedTenantSettingOverridesViaAdminApi	Get workspace delegated tenant set
Get list of users part of the datasource share policy	GetDatasourceSharePrincipalsPolicy	Retrieve the datasource share princi the tenant
Get on-demand billing limits	GetOnDemandBillingLimits	Generated when a capacity admin o page and the current billing categor limits are returned for a Fabric capa
Get on-demand billing quota	GetOnDemandBillingQuota	Generated when the On-demand bi subscription-level on-demand billin retrieved.
Get query text from secured telemetry store via Lockbox	GetQueryTextTelemetryViaLockbox	Retrieved query text from secured te
Get refresh history via lockbox	GetRefreshHistoryViaLockbox	
Get single connection by ID	GetConnection	Get single connection by ID
Get the Details of a Moniker	DiscoverSystemDetailsOfMoniker	Get Moniker and related Data Sourc
Get the current set of DLP policies applied on the Tenant	GetTenantDlpPolicies	Get the current set of DLP policies a
Get virtual network data gateway proxy	GetVirtualNetworkDataGatewayProxy	Get HTTP proxy info for virtual netw
Goals Create Role	GoalsCreateRole	
Goals Delete Role	GoalsDeleteRole	
Goals Get Role	GoalsGetRole	
Goals Update Role	GoalsUpdateRole	
Granted consent to tenant relocation	TenantRelocationConsentGranted	Tenant relocation consent granted.
Hard-deleted a folder by cleanup scheduler (Preview)	HardDeletedSubfolderByCleanupScheduler	Generated when the cleanup schedu folder after the retention period exp recovery is enabled in tenant admin
Import Package For Solution	ImportPackageForSolution	

Friendly name	Operation name	Notes
Import file to Power BI ended	ImportArtifactEnd	Generated when importing Power B `ImportSource` indicates `PowerBI` or tells you if the file is new (`Publish`)
Import file to Power BI started	ImportArtifactStart	Generated when importing Power B `ImportSource` is `PowerBI`, the file im client or API. When `ImportSource` is import originated from OneDrive or
Imported file to Power BI	Import	
Imported OneLake lifecycle policy	ImportedLifecyclePolicy	
Imported PostgreSQL sample data	ImportSampleDataToPostgreSQLDatabase	Generated when a user imports sup Native PostgreSQL database. The au operation result, and the affected Po
Initiate Cloud O Auth Login	InitiateCloudOAuthLogin	
Initiated Power BI gateway cluster authentication process	InitiateGatewayClusterOAuthLogin	
Insert Domain	InsertDataDomainAsAdmin	Insert Domain
Inserted Power BI note	InsertNote	
Inserted or updated current value connection of Power BI metric	UpsertGoalCurrentValueConnection	
Inserted or updated target value connection of Power BI metric	UpsertGoalTargetValueConnection	
Inserted snapshot for user in Power BI tenant	InsertSnapshot	Generated when user uploads a sna semantic model
Install Teams Analytics Report	InstallTeamsAnalyticsReport	
Installed Power BI app	InstallApp	
Installed Power BI template app	InstallTemplateApp	
Instantiate App	InstantiateApp	
List Lakehouse Tables	ListLakehouseTables	
List data access roles	ListDataAccessRoles	Returns a list of data access roles
Load Lakehouse Table	LoadLakehouseTable	
Load Spark App Log	LoadSparkAppLog	
Manage Relationships	ManageRelationships	
Map Upn	MapUpn	
Migrated dataflow storage location	MigratedDataflowStorageLocation	Not currently used
Migrated workspace to a capacity	MigrateWorkspaceIntoCapacity	
Modified OneLake default tier	ModifiedDefaultTier	
Modified OneLake file soft-delete settings	ModifiedOneLakeFileSoftDeleteSettings	Generated when a workspace admi soft-delete retention settings (enabl
Modify Workspace Capacity	ModifyWorkspaceCapacity	Modify Workspace Capacity is a cap on assigning a Power BI workspace the UI.

Friendly name	Operation name	Notes
Move Goals	MoveGoals	Moves goals within a scorecard
Move Scorecard	MoveScorecard	
Move items into subfolder	MoveItemsIntoSubfolder	
No Activity	NoActivity	
Opt In For P P U Trial	OptInForPPUTrial	
Override Sjd Spark Settings	OverrideSjdSparkSettings	
Patch Gateway Cluster	PatchGatewayCluster	
Patch Goal Value Categories	PatchGoalValueCategories	
Patched Power BI metric value	PatchGoalValue	
Patched Power BI metric	PatchGoal	
Patched Power BI note	PatchNote	
Patched Power BI scorecard	PatchScorecard	
Permanently delete Workspaces	DeleteWorkspacesPermanentlyAsAdmin	Generated when workspaces are pe
Pin Report Get Channels In Team	PinReportGetChannelsInTeam	
Pin Report Get User Joined Teams	PinReportGetUserJoinedTeams	
Pin Report To Teams Channel	PinReportToTeamsChannel	
Pin Tile	PinTile	
Pin Widget Tile	PinWidgetTile	
Planning Session Upgraded	PlanningSessionUpgraded	Session type is upgraded in plan ite
Post Dataset Rows	PostDatasetRows	
Post Notebook Comment	PostNotebookComment	
Post configure credentials	PostGitProviderCredentials	Configure git provider credentials fc
Posted Power BI comment	PostComment	
Preview Lakehouse Table	PreviewLakehouseTable	
Printed Power BI Dashboard	PrintDashboard	
Printed Power BI report page	PrintReport	
Promoted Power BI template app	PromoteTemplateAppPackage	
Provision Scorecard	ProvisionScorecard	
Publish/Unpublish a workload	ExtensibilityUpdatePublishingState	Publish or unpublish a workload, ma in Workloads Hub
PublishDataflow	PublishDataflow	Publish Dataflow
Published Power BI report to web	PublishToWebReport	
Put Table	PutTable	
Ran Power BI email subscription	RunEmailSubscription	
Re-encrypted credentials using Power gateway cluster	ReencryptCredentials	

Friendly name	Operation name	Notes
Read active Fabric policy	ReadActiveFabricPolicy	
Read Artifact	ReadArtifact	
Read Environment Resource	ReadEnvironmentResource	Read resources in environment
Read Experiment Run	ReadExperimentRun	
Read Fabric policy rules	ReadFabricPolicyRules	
ReadDataflow	ReadDataflow	Read Dataflow
Rebind Report	RebindReport	
Received Power BI dataflow secret from Key Vault	ReceiveDataflowSecretFromKeyVault	
Recovered an item (Preview)	ArtifactRecovered	Generated when a soft-deleted item item recovery is enabled in tenant a
Recovered folder (Preview)	RecoveredSubfolder	Generated when one or more soft-d part of artifact recovery. Available w tenant admin settings.
Refresh Datamart	RefreshDatamart	
Refresh Goal Current Value Rollup	RefreshGoalCurrentValueRollup	
Refresh Goal Target Value Rollup	RefreshGoalTargetValueRollup	
Refresh Lakehouse Data	RefreshLakehouseData	
Refresh Sql Analytics Endpoint Lakehouse Metadata	RefreshSqlAnalyticsEndpointLakehouseMetadata	Refreshed metadata for a lakehouse Previously named *Refreshed metada* (Operation name: *RefreshDefaultW
Refresh mounted warehouse metadata	RefreshMountedWarehouseMetadata	Generated when mounted warehou
Refreshed AWS Databricks Catalog shortcuts	RefreshedAWSDatabricksCatalogShortcuts	OneLake operation. Generated whe operation re-syncs the already-mirr underlying AWS Databricks Unity Ca
Refreshed current value of Power BI metric	RefreshGoalCurrentValue	
Refreshed target value of Power BI metric	RefreshGoalTargetValue	
Register workload development instance	ExtensibilityRegisterDevInstance	Register a development instance of
Remove Admin Personal Workspace Access	RemoveAdminPersonalWorkspaceAccess	
Remove workload	ExtensibilityDeactivation	Remove a workload from a tenant o
Removed Power BI group members	DeleteGroupMembers	
Removed a workspace from a deployment pipeline	UnassignWorkspaceFromPipeline	Not currently used
Removed data source from Power BI gateway	RemoveDatasourceFromGateway	
Removed user from Power BI gateway cluster datasource	RemoveGatewayClusterDatasourceUser	
Removed user from Power BI gateway cluster	RemoveGatewayClusterUser	
Removed workspace from a capacity	RemoveWorkspacesFromCapacity	

Friendly name	Operation name	Notes
Rename Datamart	RenameDatamart	
Rename Lakehouse File	RenameLakehouseFile	
Rename Lakehouse Folder	RenameLakehouseFolder	
Rename Lakehouse Table	RenameLakehouseTable	
Rename Report	RenameReport	Rename Report is a report activity, v the name of a Power BI Report throu
Rename Warehouse	RenameWarehouse	
Renamed Power BI dashboard	RenameDashboard	
Request Cognitive Service	RequestCognitiveService	Request Cognitive Service in ML wo
Request Copilot features in Fabric	RequestCopilot	Request Copilot features in Fabric
Request OpenAI by Spark code in Fabric	RequestSparkCodeFirst	Request OpenAI by Spark code in Fa
Request OpenAI model	RequestOpenAI	Request OpenAI models in ML work
Requested Power BI dataflow refresh	RequestDataflowRefresh	A dataflow refresh starts. The refresl manually from the portal or an API.
Requested Power BI semantic model refresh from external app	RefreshDatasetFromExternalApplication	
Requested Power BI semantic model refresh	RefreshDataset	
Requested SAS token for Power BI storage	AcquireStorageSASFromExternalApplication	
Requested account key for Power BI storage	AcquireStorageAccountKey	
Reserve workload name	ExtensibilityNameReserved	Reserves a workload name
Restore deleted workspace	RestoreWorkspaceViaAdminApi	Restores the deleted workspace
Restored Power BI workspace	RestoreWorkspace	
Resume Suspended Datamart	ResumeSuspendedDatamart	
Resume Suspended Sql Analytics Endpoint Lakehouse	ResumeSuspendedSqlAnalyticsEndpointLakehouse	Resumed a suspended lakehouse SC named *Resumed a suspended defaul* name: *ResumeSuspendedDefaultW
Resume Suspended Warehouse	ResumeSuspendedWarehouse	
Resume suspended mounted warehouse	ResumeSuspendedMountedWarehouse	Generated when a suspended mour
Retrieved PostgreSQL SQL audit policy	GetSqlAuditPolicyOnDatabase	Generated when a user retrieves the PostgreSQL database artifact. The au operation result, and the affected Pc
Retrieved Power BI app users	GetAppUsersAsAdmin	
Retrieved Power BI apps for user	GetUserAppsAsAdmin	Not currently used
Retrieved Power BI apps	GetAppsAsAdmin	
Retrieved Power BI capacities for user	GetUserCapacitiesAsAdmin	Not currently used

Friendly name	Operation name	Notes
Retrieved Power BI capacity users	GetCapacityUsersAsAdmin	
Retrieved Power BI dashboard tiles	GetDashboardTilesAsAdmin	
Retrieved Power BI dashboard users	GetDashboardUsersAsAdmin	
Retrieved Power BI dashboards for user	GetUserDashboardsAsAdmin	Not currently used
Retrieved Power BI dashboards	GetDashboardsAsAdmin	
Retrieved Power BI data sources for user	GetUserDatasourcesAsAdmin	Not currently used
Retrieved Power BI dataflows for user	GetUserDataflowsAsAdmin	Not currently used
Retrieved Power BI dataflows	GetDataflowsAsAdmin	
Retrieved Power BI gateway cluster datasource	GetGatewayClusterDatasource	
Retrieved Power BI gateway cluster datasources	GetGatewayClusterDatasources	
Retrieved Power BI gateway datasource users	GetDatasourceUsersAsAdmin	Not currently used
Retrieved Power BI gateway tenant key	GetGatewayTenantKeys	
Retrieved Power BI gateway tenant policy	GetGatewayTenantPolicy	
Retrieved Power BI gateway users	GetGatewayUsersAsAdmin	Not currently used
Retrieved Power BI gateways for user	GetUserGatewaysAsAdmin	Not currently used
Retrieved Power BI group users	GetGroupUsersAsAdmin	
Retrieved Power BI groups for user	GetUserGroupsAsAdmin	Not currently used
Retrieved Power BI imports	GetImportsAsAdmin	
Retrieved Power BI metric value	GetGoalValue	
Retrieved Power BI metric	GetGoal	
Retrieved Power BI refresh history	GetRefreshHistory	
Retrieved Power BI refreshable by ID	GetRefreshablesForRefreshIdAsAdmin	
Retrieved Power BI refreshables for capacity	GetRefreshablesForCapacityAsAdmin	
Retrieved Power BI refreshables	GetRefreshablesAsAdmin	
Retrieved Power BI report users	GetReportUsersAsAdmin	
Retrieved Power BI reports for user	GetUserReportsAsAdmin	Not currently used

Friendly name	Operation name	Notes
Retrieved Power BI scorecard by using report ID	GetScorecardByReportId	
Retrieved Power BI scorecard	GetScorecard	
Retrieved Power BI semantic models for user	GetUserDatasetsAsAdmin	Not currently used
Retrieved Power BI semantic models	GetDatasetsAsAdmin	
Retrieved Power BI tenant keys	GetTenantKeysAsAdmin	
Retrieved Power BI workspaces	GetWorkspaces	
Retrieved a model from Power BI	GetPowerBIDataModel	A user opens the Open data model model.
Retrieved all Power BI gateway cluster datasources	GetAllGatewayClusterDatasources	
Retrieved all supported datasources for Power BI gateway cluster	GetGatewayClusterSupportedDatasources	
Retrieved allowed Power BI gateway regions	GetGatewayRegions	
Retrieved authentication details for Power BI gateway cluster datasource	GetGatewayClusterDatasourceOAuthDetails	
Retrieved data sources from Power BI dataflow	GetDataflowDatasourcesAsAdmin	
Retrieved data sources from Power BI semantic model	GetDatasetDatasourcesAsAdmin	
Retrieved links between semantic models and dataflows	GetDatasetToDataflowsLinksAsAdmin	
Retrieved list of Power BI gateway installer principals	GetGatewayInstallerPrincipals	
Retrieved list of datasource users for Power BI gateway cluster	GetGatewayClusterDatasourceUsers	
Retrieved list of modified workspaces in Power BI tenant	GetModifiedWorkspacesAPI	
Retrieved member status of Power BI gateway cluster	GetGatewayClusterMemberStatus	
Retrieved metrics of Power BI scorecard	GetGoalsForScorecard	
Retrieved multiple Power BI gateway clusters	GetGatewayClusters	
Retrieved multiple Power BI metric values	GetGoalValues	
Retrieved multiple Power BI scorecards	GetScorecards	
Retrieved scan result in Power BI tenant	GetWorkspacesInfoResult	
Retrieved snapshots for user in Power BI tenant	GetSnapshots	Generated when user retrieves snap model such as when a user visits the

Friendly name	Operation name	Notes
Retrieved status of Power BI gateway cluster datasource	GetGatewayClusterDatasourceStatus	
Retrieved status of Power BI gateway cluster	GetGatewayClusterStatus	
Retrieved upstream dataflows from Power BI dataflow	GetDataflowUpstreamDataflowsAsAdmin	
Retry lakehouse SQL analytics endpoint creation for a Lakehouse	RetryLakehouseSqlEndpointCreation	Retry SQL analytics endpoint creatic
Revoke an external data share	RevokeExternalDataShare	Revoke an external data share
Revoked consent to tenant relocation	TenantRelocationConsentRevoked	Tenant relocation consent revoked.
Rotate Tenant Key	RotateTenantKey	
Rotated Power BI gateway tenant key	RotateTenantKeyEncryptionKey	
Run Artifact	RunArtifact	
Save Model Diagram Layouts	SavePowerBIDataModelDiagramLayouts	Save diagram layouts after update c
Saved an autogenerated report to Power BI	SaveAutogeneratedReport	After exploring an autogenerated Pc application, a user saved it to the Pc
Saved an autogenerated semantic model to Power BI	SaveAutogeneratedDataset	After exploring an autogenerated Pc external application, a user saved it
Schedule Artifact	ScheduleArtifact	
Sent a scan request in Power BI tenant	GetWorkspacesInfoAPI	
Sent a message in a Digital Operations Ontology Agent conversation	DigitalOperationsOntologyAgentMessageSent	A user sent a message in a Digital O conversation
Set Capacity Tenant Key	SetCapacityTenantKey	
Set D Q Refresh Schedule Of Dataset	SetDQRefreshScheduleOfDataset	
Set DirectQuery Refresh Schedule Of Dateset	SetDQRefreshScheduleOfDateset	
Set Lakehouse Endorsement	SetLakehouseEndorsement	
Set Lakehouse Sensitivity Label	SetLakehouseSensitivityLabel	
Set Model Refresh Schedule Of Dateset	SetModelRefreshScheduleOfDateset	
Set Model Refresh Schedule Of Dataset	SetModelRefreshScheduleOfDataset	
Set Notebook Default Lakehouse	SetNotebookDefaultLakehouse	Set default lakehouse for notebook.
Set PostgreSQL SQL audit policy	SetSqlAuditPolicyOnDatabase	Generated when a user updates SQI PostgreSQL database artifact. The au operation result, and the affected Pc
Set Sjd Retry Policy	SetSjdRetryPolicy	
Set dataflow storage location for a workspace	SetDataflowStorageLocationForWorkspace	

Friendly name	Operation name	Notes
Set scheduled refresh on Power BI dataflow	SetScheduledRefreshOnDataflow	
Set scheduled refresh on Power BI semantic model	SetScheduledRefresh	
Set warehouse hardware acceleration	SetWarehouseHardwarePlatform	Changed the current hardware acce in a workspace.
Share Artifact	ShareArtifact	
Share Datamart	ShareDatamart	
Share Lakehouse Table	ShareLakehouseTable	
Share Warehouse	ShareWarehouse	
Shared Power BI dashboard	ShareDashboard	
Shared Power BI report	ShareReport	
Shared Power BI semantic model	ShareDataset	
Soft-deleted an item (Preview)	ArtifactSoftDeleted	Generated when an item is soft-dele bin. Available when item recovery is settings.
Start Notebook Session	StartNotebookSession	
Start Publish Environment	StartPublishEnvironment	Start publish environment
Started Power BI extended trial	OptInForExtendedProTrial	Not currently used
Started Power BI trial	OptInForProTrial	
Stop Notebook Session	StopNotebookSession	
Switch Branch Git	SwitchBranchInGit	Switch Branch Git is a workspace act the user changes what git branch is
Sync Dataset Query Scale-Out Replicas	SyncDatasetQueryScaleOutReplicas	Sync Dataset Query Scale-Out Repli which is generated when users requ replicas of a scale out-enabled Powe read/write replica.
Take Over Email Subscription	TakeOverEmailSubscription	
Tested Power BI gateway datasource connection with single sign-on	GatewayClusterDatasourceSSOTestConnection	
Took over Power BI semantic model	TakeOverDataset	
Took over a Power BI datasource	TakeOverDatasource	
Took ownership of Power BI dataflow	TookOverDataflow	
Take over a Fabric item	TakeOverArtifact	
Trial License Extension	TrialLicenseExtension	Extend user trials by user list or tena
Unassign Workspace From Alm Pipeline	UnassignWorkspaceFromAlmPipeline	
Undelete blob	UndeleteBlob	OneLake operation. Restore a soft d
Undo Git	UndoGit	Undo Git is an item activity, which is changes done to item.
Unfollow Goal	UnfollowGoal	

Friendly name	Operation name	Notes
Unpublished Power BI app	UnpublishApp	
Update Alm Pipeline Access As Admin	UpdateAlmPipelineAccessAsAdmin	
Update Alm Pipeline	UpdateAlmPipeline	
Update Artifact	UpdateArtifact	
Update Byo Resource Access	UpdateByoResourceAccess	
Update Capacity Delegation settings	UpdateCapacityTenantSettingDelegation	Update Capacity delegation settings
Update Data Sharing	UpdateDataSharing	Update an existing external data sha
Update Datamart Metadata	UpdateDatamartMetadata	
Update Datamart Settings	UpdateDatamartSettings	
Update Datamart	UpdateDatamart	
Update Dataset Parameters	UpdateDatasetParametersForSolution	Update Dataset Parameters is a sem generated when updates are made parameters
Update Dataset	UpdateDataset	Update Dataset is a semantic mode when users updated the properties
Update Default Domain	UpdateDefaultDataDomainAsAdmin	Update Default Domain
Update Default Personal Workspace Capacity	UpdateDefaultPersonalWorkspaceCapacity	
Update Domain Access Permissions	UpdateDataDomainAccessAsAdmin	Update Data Domain Access Permis
Update Domain Branding	UpdateDataDomainBrandingAsAdmin	Update Domain Branding
Update Domain Contributors Scope	UpdateDataDomainContributorsScopeAsAdmin	Update Domain Contributors Scope
Update Domain Delegation settings	UpdateDomainTenantSettingDelegation	Update Domain Delegation settings
Update Domain	UpdateDataDomainAsAdmin	Update Domain
Update Domain's Folders Relations As Contributor	UpdateDataDomainFoldersRelationsAsContributor	Update Domain's Folders Relations
Update Domain's Folders Relations	UpdateDataDomainFoldersRelationsAsAdmin	Update Domain's Folders Relations
Update Environment Resource	UpdateEnvironmentResource	Update resources in environment
Update Environment Spark Settings	UpdateEnvironmentSparkSettings	Update environment spark settings
Update Experiment Run	UpdateExperimentRun	
Update Folder Access	UpdateFolderAccess	Indicates an update to workspace a
Update Folder	UpdateFolder	Indicates a workspace update
Update From Git	UpdateFromGit	Update From Git is an item activity, update item from Git.
Update Gateway Cluster Member	UpdateGatewayClusterMember	
Update Gateway Installer Principals	UpdateGatewayInstallerPrincipals	
Update Gateway Tenant Policy	UpdateGatewayTenantPolicy	

Friendly name	Operation name	Notes
Update Goal Connection Settings	UpdateGoalConnectionSettings	
Update Goal Current Value Connection Owner	UpdateGoalCurrentValueConnectionOwner	
Update Goal Target Value Connection Owner	UpdateGoalTargetValueConnectionOwner	
Update Goals	UpdateGoals	
Update Group	UpdateGroup	
Update Hierarchy Goal Value	UpdateHierarchyGoalValue	
Update Hierarchy Note	UpdateHierarchyNote	
Update In Place Sharing Settings	UpdateInPlaceSharingSettings	
Update MetricSet Metric Endpoint	UpdateMetricSetMetric	Update metric in a metricset endpoi
Update Notebook Library	UpdateNotebookLibrary	
Update Notebook Resource	UpdateNotebookResource	Update resources in notebook.
Update Notebook Spark Property	UpdateNotebookSparkProperty	
Update Notification Settings	UpdateNotificationSettings	Update user notification settings for
Update Publish to Web Embed Code Owner	UpdateEmbedCodeOwner	
Update Report Content	UpdateReportContent	
Update Scorecard Dataset	UpdateScorecardDataset	
Update Scorecard Hierarchy	UpdateScorecardHierarchy	
Update Scorecard View	UpdateScorecardView	
Update Semantic Metric Endpoint	UpdateSemanticMetric	Update standalone semantic metric
Update Service Principal Profile	UpdateServicePrincipalProfile	
Update Sql Analytics Endpoint Lakehouse Settings	UpdateSqlAnalyticsEndpointLakehouseSettings	Updated settings for a lakehouse SC named *Updated a default warehouse* name: *UpdateDefaultWarehouse*)
Update Sql Analytics Endpoint Lakehouse	UpdateSqlAnalyticsEndpointLakehouse	Updated a lakehouse SQL analytics
Update USEC roles for an artifact	UpdateArtifactRoles	OneLake operation.
Update Virtual Network	UpdateVirtualNetwork	
Update Warehouse Metadata	UpdateWarehouseMetadata	
Update Warehouse Settings	UpdateWarehouseSettings	
Update Warehouse	UpdateWarehouse	
Update Workspace Delegation settings	UpdateWorkspaceTenantSettingDelegation	Update Workspace Delegation setti
Update capacity resource governance settings	UpdateCapacityResourceGovernanceSettings	Not currently in Microsoft 365 admi
Update datasource share policy	UpdateDatasourceShareTenantPolicy	Set the datasource share policy set

Friendly name	Operation name	Notes
Update list of users part of the datasource share policy	UpdateDatasourceSharePrincipalsPolicy	Set the datasource share principals t tenant
Update mounted warehouse settings	UpdateMountedWarehouseSettings	Generated when mounted warehou
Update mounted warehouse	UpdateMountedWarehouse	Generated when mounted warehou
Update on-demand billing limits	UpdateOnDemandBillingLimits	Generated when a capacity admin e category, or changes its 24-hour co One audit event is emitted per billin
Update source in GraphQL artifact	UpdateSourceGraphQL	Update source in GraphQL item
Update subfolder	UpdateSubfolder	
Update task flow	UpdateTaskFlow	
Update the current set of DLP policies applied on the Tenant	UpdateTenantDlpPolicies	Update the current set of DLP polici
Update virtual network data gateway proxy	UpdateVirtualNetworkDataGatewayProxy	Update HTTP proxy for virtual netwo
Update workspace role	UpdateWorkspaceRoleViaAdminApi	Update workspace role
Updated Artifact Definition	UpdatedArtifactDefinition	OneLake operation. Generated whe artifact definition is updated via the (shared across item types such as M Databricks Catalog).
Updated AWS Databricks Catalog Definition	UpdatedAWSDatabricksCatalogDefinition	OneLake operation. Generated whe workspace connection or mirrored s is updated via the update-item-defi
Updated MirroredStorage definition	UpdatedMirroredStorageDefinition	OneLake operation. Generated whe modifies the mirroring scope of an e changing which external paths are e
Updated Power BI access request settings	UpdateAccessRequestSettings	
Updated Power BI app	UpdateApp	
Updated Power BI dataflow	UpdateDataflow	
Updated Power BI discoverable model settings	UpdateDiscoverableModelSettings	Generated when a report is set to fe
Updated Power BI email subscription	UpdateEmailSubscription	
Updated Power BI gateway cluster datasource	UpdateGatewayClusterDatasource	
Updated Power BI gateway data source credentials	UpdateDatasourceCredentials	
Updated Power BI semantic model data sources	UpdateDatasources	
Updated Power BI semantic model parameters	UpdateDatasetParameters	
Updated Power BI workspace access	UpdateWorkspaceAccess	
Updated Power BI workspace	UpdateWorkspace	
Updated an organizational custom visual	UpdateOrganizationalGalleryItem	
Updated capacity admin	UpdateCapacityAdmins	

Friendly name	Operation name	Notes
Updated capacity custom settings	UpdateCapacityCustomSettings	
Updated capacity display name	UpdateCapacityDisplayName	
Updated credentials for Power BI gateway cluster	UpdateGatewayClusterDatasourceCredentials	
Updated dataflow storage assignment permissions	UpdatedDataflowStorageAssignmentPermissions	
Updated deployment pipeline access	UpdateAlmPipelineAccess	
Updated deployment pipeline configuration	SetConfigurationAlmPipeline	
Updated Fabric policy rules	UpdatedFabricPolicyRules	
Updated featured tables	UpdateFeaturedTables	
Updated organization's Power BI settings	UpdatedAdminFeatureSwitch	
Updated parameters for installed Power BI template app	UpdateInstalledTemplateAppParameters	
Updated settings for Power BI template app	UpdateTemplateAppSettings	
Updated snapshots for user in Power BI tenant	UpdateSnapshot	Generated when user updates snaps semantic models
Updated testing permissions for Power BI template app	UpdateTemplateAppTestPackagePermissions	
Updated the Power BI datasource	UpdateDatasource	
Updated the Power BI gateway	UpdateGateway	
Updated workspace Analysis Services settings	SetASSeverPropertyOnWorkspaceFromExternalApplicationDetailedInfo	Not currently used
Upgrade Workspace	UpgradeWorkspace	
Upgrade Workspaces As Admin	UpgradeWorkspacesAsAdmin	
Upload workload package	ExtensibilityUploadPackage	Upload a workload package in a ten
Upsert Datamart Parameters	UpsertDatamartParameters	
Upsert Goal Current Value Rollup	UpsertGoalCurrentValueRollup	
Upsert Goal Status Rules	UpsertGoalStatusRules	
Upsert Goal Target Value Rollup	UpsertGoalTargetValueRollup	
Upsert Goal Values	UpsertGoalValues	
Upsert Sql Analytics Endpoint Lakehouse Parameters	UpsertSqlAnalyticsEndpointLakehouseParameters	Upserted parameters for a lakehous Previously named *Updated paramet* (Operation name: *UpsertDefaultWa
Upsert Warehouse Parameters	UpsertWarehouseParameters	
Upsert mounted warehouse parameters	UpsertMountedWarehouseParameters	Generated when mounted warehou updated

Friendly name	Operation name	Notes
Used Power BI to explore data in an external application	ExploreDataExternally	Someone used Power BI to explore application.
View Datamart	ViewDatamart	
View Spark App Input Output	ViewSparkAppInputOutput	
View Spark App Log	ViewSparkAppLog	
View Spark Application	ViewSparkApplication	
View SQL Analytics Endpoint Lakehouse	ViewSqlAnalyticsEndpointLakehouse	Viewed a lakehouse SQL analytics e *Viewed a default warehouse* (Operat *ViewDefaultWarehouse*)
View Warehouse	ViewWarehouse	
View mounted warehouse	ViewMountedWarehouse	Generated when mounted warehou
Viewed Power BI dashboard	ViewDashboard	Some fields such as CapacityID and the report or dashboard is viewed fr a Power BI workspace
Viewed Power BI dataflow	ViewDataflow	
Viewed Power BI metadata	ViewMetadata	
Viewed Power BI report	ViewReport	A report is also generated per page fields such as CapacityID and Capac report or dashboard is viewed from Power BI workspace.
Viewed Power BI tile	ViewTile	
Viewed Power BI usage metrics	ViewUsageMetrics	
View or modify CMK encryption settings	GetWorkspaceEncryption	Logged when the encryption setting modified, or applied to items.
disable workspace	DisableWorkspaceViaAdminApi	Disables the workspace
Disable CMK encryption for your workspace	DisableWorkspaceEncryption	Triggered when CMK encryption is d
Disable Workspace Outbound Access Protection	DisableWorkspaceOutboundAccessProtection	Outbound Access Protection Disabl
Enable Workspace Outbound Access Protection	EnableWorkspaceOutboundAccessProtection	Outbound Access Protection Enable
Get Workspace Resource Instance Rules	GetWorkspaceResourceInstanceRules	
Set Workspace Resource Instance Rules	SetWorkspaceResourceInstanceRules	
Enable workspace inbound access protection	EnableWorkspaceInboundAccessProtection	Enable workspace inbound access p
Disable workspace inbound access protection	DisableWorkspaceInboundAccessProtection	Disable workspace inbound access p
Update workspace outbound access protection connection rules	UpdateWorkspaceOutboundConnections	Update workspace outbound access
Update workspace outbound access protection gateway rules	UpdateWorkspaceOutboundGateways	Update workspace outbound access
Enable Git for workspace with restricted outbound access	EnableGitForWorkspaceWithRestrictedOutboundAccess	Enable Git for workspace with restri

Friendly name	Operation name	Notes
Disable Git for workspace with restricted outbound access	DisableGitForWorkspaceWithRestrictedOutboundAccess	Disable Git for workspace with restr
Get workspace IP firewall rules	GetFolderFirewallRules	Get workspace IP firewall rules
Update workspace IP firewall rules	SetFolderFirewallRules	Update workspace IP firewall rules
Enabled external data share access for inbound access restricted workspaces	ExternalDataSharesBypassForWorkspaceEnabled	
Disabled external data share access for inbound access restricted workspaces	ExternalDataSharesBypassForWorkspaceDisabled	
Updated Git connection settings	GitConnectionSettingsUpdated	Updated Git connection settings
Branch workspace configuration updated	GitBranchWorkspaceConfigured	Branch workspace configuration up

Considerations and limitations

When capacity ID and capacity name aren't available in the audit logs, you can view them in the Microsoft Fabric Capacity Metrics app.

> **Important**
>
> **Starting July 7, 2025**, As part of our ongoing efforts to simplify audit logging and ensure consistency across Fabric, we're **standardizing redundant item operations** into a single set of unified operations. This migration aligns warehouse, warehouse, and SQL Analytics Endpoint lakehouse actions under common names used across the Fabric platform.

What's Changing?

Expand table

Friendly Name	Old Operation Name	New Operation Name
Create Datamart	CreateDatamart	CreateArtifact
Create Warehouse	CreateWarehouse	CreateArtifact
Delete Datamart	DeleteDatamart	DeleteArtifact
Delete Warehouse	DeleteWarehouse	DeleteArtifact
Rename Datamart	RenameDatamart	UpdateArtifact
Rename Warehouse	RenameWarehouse	UpdateArtifact
Share Datamart	ShareDatamart	ShareArtifact
Share Warehouse	ShareWarehouse	ShareArtifact
Update Datamart	UpdateDatamart	UpdateArtifact
Update Datamart Metadata	UpdateDatamartMetadata	UpdateArtifact
Update Datamart Settings	UpdateDatamartSettings	UpdateArtifact
Update Warehouse	UpdateWarehouse	UpdateArtifact
Update Warehouse Metadata	UpdateWarehouseMetadata	UpdateArtifact
Update Warehouse Settings	UpdateWarehouseSettings	UpdateArtifact

Friendly Name	Old Operation Name	New Operation Name
Upsert Datamart Parameters	UpsertDatamartParameters	UpdateArtifact
Upsert SQL Analytics Endpoint Lakehouse Parameters	UpsertSqlAnalyticsEndpointLakehouseParameters	UpdateArtifact
Upsert Warehouse Parameters	UpsertWarehouseParameters	UpdateArtifact
View Datamart	ViewDatamart	ReadArtifact
View SQL Analytics Endpoint Lakehouse	ViewSqlAnalyticsEndpointLakehouse	ReadArtifact
View Warehouse	ViewWarehouse	ReadArtifact

Related content

Track user activities in Fabric

Last updated on 09/08/2026

Find users who signed in

10/09/2025

If you're an admin for your organization, you can use the Microsoft Entra admin center to view sign-in logs.

This guide shows you how to find sign-in logs specifically for Power BI. For more information about sign-in logs for other users and applications, see the Microsoft Entra sign-in logs documentation.

> **Note**
>
> The *Sign-in logs* report provides useful information, but it doesn't identify the type of license for each user. Use the Microsoft 365 admin center to view licenses.

Requirements

Any user can view a report of their own sign-ins. To see a report for all users, you must have a Fabric administrator role.

Use the Microsoft Entra admin center to view sign-ins

To view sign-in activity, follow these steps:

1. Sign in to the Microsoft Entra admin center, and then expand **Entra ID** from the left navigation pane.

2. From the left navigation pane under **Entra ID**, select **Monitoring & health** > **Sign-in logs**.

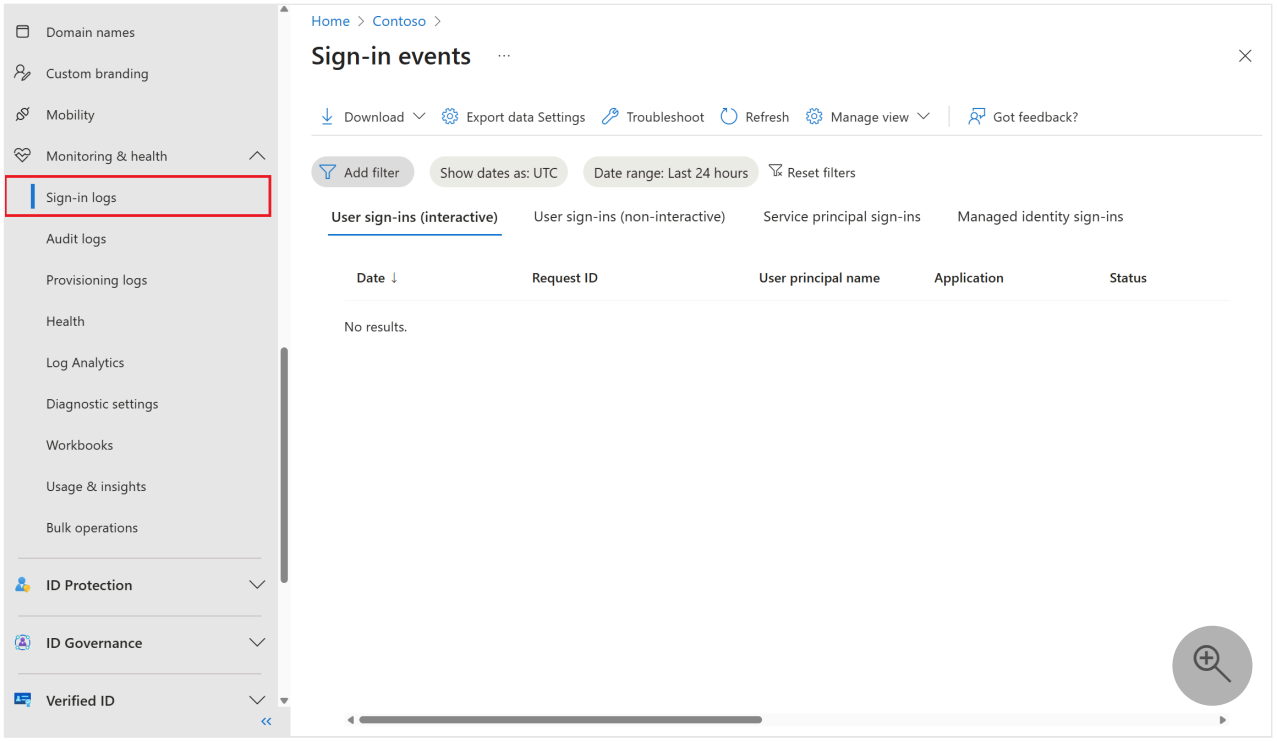

By default, all sign-ins from the last 24 hours for all users and all applications are shown.

3. To select a different time period, select **Date** in the working pane and choose from the available time intervals. For more information about available time intervals, see data retention.

4. To see only sign-ins to specific applications, add filters.

5. Select **Add filter** and then select **Application** as the field to filter by.

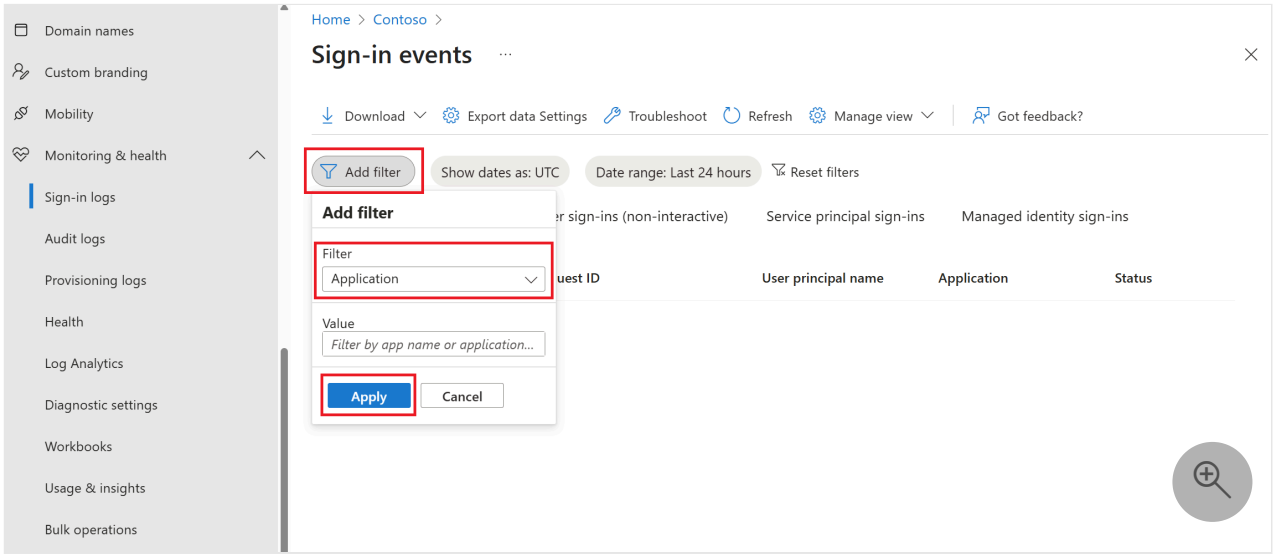

- To see only sign-in activity that's specific to Power BI service, enter **Microsoft Power BI**.

- To see only sign-in activity that's specific to the on-premises data gateway, enter **Power BI Gateway**.

6. Select **Apply**.

Export the data

You can download a sign-in report in either of two formats: a CSV file, or a JSON file. Use the following steps to download your report:

1. From the command bar for the **Sign-in events** page, select **Download** and then select one of the following options:

   - **Download JSON** to download a JSON file for the currently filtered data.
   - **Download CSV** to download a CSV file for the currently filtered data.

2. Decide what type of sign-ins you want to export, and then select **Download**.

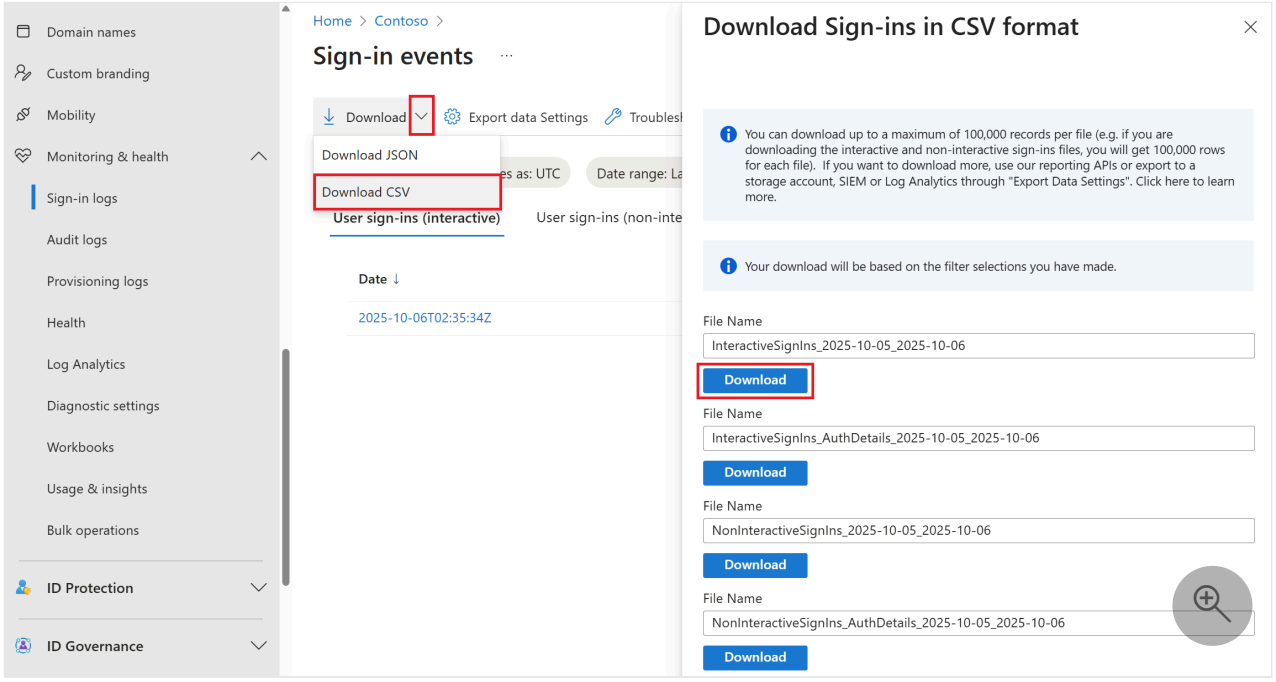

> **Note**
>
> You can download up to a maximum of 100,000 records per file. For example, if you're downloading the interactive and non-interactive sign-ins files, you get 100,000 rows for each file. If you want to download more, use our reporting APIs or export to a storage account, SIEM, or Log Analytics through **Export Data Settings**.

Data retention

Sign-in-related data is available for up to seven days, unless your organization has a Microsoft Entra ID P1 or P2 license. If you use Microsoft Entra ID P1 or Microsoft Entra ID P2, you can see

data for the past 30 days. For more information, see How long does Microsoft Entra ID store reporting data?.

Related content

- Use the Monitoring hub
- Microsoft Entra sign-in logs

View refresh summary

Article • 11/21/2023

The refresh summary admin settings page lets you view your capacity's refresh history. You can also export the refresh history, and view details related to a specific refresh. The information in this page can help you investigate refresh errors, and establish a refresh schedule for the Power BI items that reside on your capacities.

To learn how to access the Fabric admin portal settings, see What is the admin portal?

Schedule

The schedule tab lists all the refreshes that took place in a specific capacity. Select the capacity you want to review from the *choose a capacity* dropdown menu. Use the *refresh* button to refresh the table's results, and the *export* button to export a .csv file.

To view details for a specific refresh instance, select the instance and then select **Details**.

History

The history tab lists all the refreshes that took place in all the capacities you're an admin of. The table headers allow you to sort the information and apply filters. Use the *refresh* button to refresh the table's results, and the *export* button to export a .csv file.

Related content

- What is the admin portal?

Feedback

Was this page helpful? Yes No

Provide product feedback | Ask the community

Manage workspaces

As a Fabric administrator, you can view and manage the workspaces in your organization from **Workspaces** in the Govern section of the OneLake catalog.

To open **Workspaces**:

1. Sign in to Fabric by using your admin account credentials.
2. Select **OneLake catalog**.
3. Select **Govern**, and then select **Workspaces** under **Manage**.

The **Workspaces** page lists the workspaces in your tenant. Select a workspace to display the commands available for its type and state.

The following table describes the columns of the list of workspaces.

Expand table

Column	Description
Name	The name given to the workspace.
Description	The information that is given in the description field of the workspace settings.
Type	The workspace type: **Workspace**, **Personal Group** (My workspace), or **Admin Workspace**.
State	The state lets you know if the workspace is available for use. There are five states, **Active**, **Orphaned**, **Deleted**, **Removing**, and **Not found**. For more information, see Workspace states.
Capacity name	The name of the capacity assigned to the workspace.
Capacity SKU Tier	The SKU tier of the workspace's assigned capacity. For more information, see Configure and manage capacities in Power BI Premium.

For more information about workspaces, see Workspaces. To retrieve workspace metadata programmatically, use the admin REST API.

Find and filter workspaces

Use **Filter by keyword** to search the workspace list. Select **Filter** to refine the list by:

- **Type**: Admin Workspace, Personal Group, or Workspace.

- **Capacity SKU Tier**.
- **Capacity name**.
- **State**.

Select **Clear all** to remove the applied filters.

Workspace states

The following table describes the possible workspace states.

Expand table

State	Description
Active	A normal workspace. It doesn't indicate anything about usage or what's inside, only that the workspace itself is "normal."
Orphaned	A workspace with no admin user. You need to assign an admin.
Deleted	A deleted workspace. When a workspace is deleted, it enters a retention period. During the retention period, a Fabric administrator can restore the workspace. See Retention and recovery for detail. When the retention period ends, the workspace enters the *Removing* state.
Removing	At the end of a deleted workspace's retention period, it moves into the *Removing* state. During this state, the workspace is permanently removed. Permanently removing a workspace takes a short while, and depends on the service and folder content.
Not found	If the customer's API request includes a workspace ID for a workspace that doesn't belong to the customer's tenant, "Not found" is returned as the status for that ID.

Retention and recovery

Fabric provides retention and recovery capabilities at both the workspace and item levels. When workspaces or items are deleted, they enter a retention period during which administrators can restore them.

Workspace retention:

- Personal workspaces (*My workspaces*) have a fixed 30-day retention period
- Collaborative workspaces have a configurable retention period (default 7 days, adjustable from 7 to 90 days)
- Administrators can restore deleted workspaces or permanently delete them before the retention period expires

Item retention:

- Item recovery is turned on by default, and individually deleted items have a configurable retention period (default 3 days, adjustable from 3 to 90 days)
- Administrators and workspace admins can restore deleted items using REST APIs
- Administrators can permanently delete items before the retention period expires

For detailed instructions on setting up retention periods, restoring workspaces and items, and permanently deleting resources, see Retention and recovery in Fabric.

Workspace options

When you select a workspace, a command bar appears above the list. You can also select **More options (...)** next to a workspace to open the same actions in a menu. The available commands depend on the workspace type and state. For example, an active workspace can provide edit, access, export, and reassignment commands, while a deleted workspace can provide edit, export, restore, and permanent-delete commands.

Expand table

Option	Description
Edit	Opens **Workspace settings** for the selected workspace.
Access	Opens the **Manage access** pane for the selected workspace.
Export	Exports the workspace list as a *.csv* file.
Reassign workspace	Opens a pane where you can change the workspace type or assigned capacity.
Restore	Restores a deleted My workspace or collaborative workspace. For details, see Restore a deleted My workspace as an app workspace or Restore a deleted collaborative workspace.
Permanently delete	Permanently deletes a deleted collaborative workspace before its retention period ends. For details, see Permanently delete a deleted collaborative workspace during the retention period.

> **Note**
>
> Admins can also manage and recover workspaces using PowerShell cmdlets.

Admins can also control users' ability to create new workspace experience workspaces and classic workspaces. See Workspace settings in this article for details.

Manage workspace access

To manage who can access a workspace:

1. Select the workspace in the list.
2. Select **Access** on the command bar, or select **More options (...)** and then select **Access**.
3. In the **Manage access** pane, you can:

    - Select **Add people or groups** to grant workspace access.
    - Search for a person or group that already has access.
    - Review or change the role assigned to a person or group.

Edit workspace settings

To open the settings for a workspace:

1. Select the workspace in the list.
2. Select **Edit** on the command bar.
3. Update the available settings in the **Workspace settings** pane.

The settings available in the pane depend on the workspace and your permissions.

Workspace item limits

Workspaces can contain up to 1,000 Fabric and Power BI items, including both parent and child items.

Users who try to create new items after reaching this limit receive an error during the item creation process. To develop a plan for managing item counts in workspaces, Fabric admins can review workspace inventory in the Govern report in the OneLake catalog.

> **Note**
>
> If specific item types have lower limits, those limits still apply. For item-specific limits, review the documentation for that item type.

Reassign a workspace to a different capacity

Capacities host workspaces and the data they contain. Reassign a workspace to change its workspace type or assigned capacity.

1. Select the workspace in the list.
2. Select **Reassign workspace** on the command bar.
3. In the **Reassign workspace** pane, select an available workspace type.
4. If required, select the capacity under **Details**.
5. Select **Apply**.

> **Note**
>
> - The types of items in the workspace can affect your ability to change workspace types or move the workspace to a capacity in a different region.
> - Moving a workspace to a different capacity might start successfully but finish with errors, which could affect some or all items in the workspace. For details, see Capacity reassignment restrictions and common issues.

Govern My workspaces

Every Fabric user has a personal workspace called My workspace where they can work with their own content. While only My workspace owners have access to their My workspaces, Fabric admins can use a set of features to help them govern these workspaces. With these features, Fabric admins can:

- Gain access to the contents of any user's My workspace
- Designate a default capacity for all existing and new My workspaces
- Prevent users from moving My workspaces to a different capacity that might reside in noncompliant regions
- Restore deleted My workspaces as app workspaces

These features are described in the following sections.

Gain access to any user's My workspace

Fabric admins can temporarily access any user's My workspace to view and manage its contents. This access automatically revokes after 24 hours. The My workspace owner's access remains intact.

When you gain temporary access to a My workspace, you can:

- See the My workspace in the list of workspaces accessible from the navigation pane. The icon indicates that it's a My workspace.
- Perform any actions in the My workspace as if it's your own My workspace. You can view and make any changes to the contents, including sharing or unsharing. But you can't grant anyone else access to the My workspace.

You can manage access to My workspaces by using the Govern section of the OneLake catalog or the Fabric Admin APIs.

To manage access by using the OneLake catalog Govern section:

1. Sign in to Fabric using your admin account credentials.
2. Open the **OneLake catalog**, select the **Govern** section, and then select **Workspaces**.
3. Select the personal workspace you want to access, and then select **Get Access** on the command bar.
4. To remove access, select the workspace, and then select **Remove Access** on the command bar.

   > **Note**
   >
   > If you don't remove access, it automatically revokes after 24 hours.

To manage access by using the Fabric Admin APIs:

- Use Workspaces - Grant Admin Temporary Access to gain access to the My workspace.
- Use Workspaces - Remove Admin Temporary Access to remove access to the My workspace.

Designate a default capacity for My workspaces

A Fabric admin or capacity admin can designate a capacity as the default capacity for My workspaces. To configure a default capacity for My workspaces, go to the details section in your capacity settings.

For details, see Designate a default capacity for My workspaces

Prevent My workspace owners from reassigning their My workspaces to a different capacity

Fabric admins can designate a default capacity for My workspaces. However, even if a My workspace is assigned to Power BI Premium capacity, the owner of the workspace can still move it back to Power BI Pro workspace type. Moving a workspace from Power BI Premium workspace type to Power BI Pro workspace type might cause the content contained in the workspace to be become noncompliant with respect to data-residency requirements, since it might move to a different region. To prevent this situation, the Fabric admin can block My workspace owners from moving their My workspace to a different workspace type by turning on the **Block users from reassigning personal workspaces (My Workspace)** tenant setting. See Workspace settings for detail.

Moving data around

Workspaces and the data they contain reside on capacities. Workspace admins can move the data contained in a workspace by reassigning the workspace to a different capacity. The capacity can be in the same region or a different region.

For details, see Capacity reassignment restrictions and common issues.

Related content

- Administration overview

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Workspace tenant settings

As a Fabric administrator, you can configure workspace settings in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

Create workspaces

Workspaces are places where users collaborate on dashboards, reports, and other content. Microsoft Fabric admins can use the **Create workspaces** setting to designate which users in the organization can create workspaces. Admins can let everybody or nobody in an organization create workspaces. Workspace creation can also be limited to members of specific security groups. Learn more about workspaces.

List of workspaces

The **Workspaces** section in **OneLake catalog** > **Govern** has settings for the workspaces in your tenant. In that section, you can sort and filter the list of workspaces and display the details for each workspace. See Manage workspaces for details.

Publish apps

In the tenant settings, you also control which users have permissions to distribute apps to the organization. See Publish apps to the entire organization for details.

Use semantic models across workspaces

Admins can control which users in the organization can use semantic models across workspaces. When this setting is enabled, users still need the required Build permission for a specific semantic model.

For more information, see Intro to semantic models across workspaces.

Identify your workspace ID

The easiest way to find your workspace ID is in the URL of the Fabric site for an item in a workspace. As in Power BI, the Fabric URL contains the workspace ID, which is the unique identifier after `/groups/` in the URL, for example: `https://powerbi.com/groups/11aa111-a11a-1111-`

`1abc-aa1111aaaa/...`. Alternatively, you can find the workspace ID in **Admin portal** > **Workspaces** by selecting the **Actions** menu > **Details** next to the workspace name.

Block users from reassigning personal workspaces (My Workspace)

Personal workspaces are the My workspaces that every user has for their personal content. Microsoft Fabric and capacity admins can designate a preferred capacity for My workspaces. By default, however, My workspace owners can still change the capacity assignment of their workspace. If a Microsoft Fabric or capacity admin designates a Premium capacity as the default capacity for My workspaces, but a My workspace owner then changes that capacity assignment back to shared capacity, this could result in non-compliance with data residency requirements.

To prevent such a scenario, the Microsoft Fabric admin can turn on the **Block users from reassigning personal workspaces (My Workspace)** tenant setting. When this setting is on, My workspace owners can't change the capacity assignment of their My workspace.

To turn on the setting, sign in to Fabric using your admin account credentials. Open the **OneLake catalog** and go to the **Govern** tab > **Configurations** > **Tenant settings**. Then scroll to the **Workspace settings** section, and find **Block users from reassigning personal workspaces (My Workspace)**.

For more information, see Prevent My workspace owners from reassigning their My workspaces to a different capacity.

Automatically convert and store reports using Power BI enhanced metadata format (PBIR) (preview)

Turn on this setting to automatically convert reports to Power BI enhanced metadata format (PBIR) after they are edited and save them in the workspace.

When this setting is enabled, new reports are also created in PBIR format. PBIR provides a source-control-friendly file structure that can improve co-development workflows and report lifecycle management.

This setting is on by default, but it doesn't take effect before early 2026. If you opt out before then, use of PBIR in the service remains off until general availability.

For more information, see Power BI enhanced report format (PBIR).

Related content

- About tenant settings
- CI/CD workflow options in Fabric
- Set up and manage workspace retention

Last updated on 09/29/2026

Capacity reassignment restrictions and common issues

Workspaces and the data they contain reside on capacities. Workspace admins can move the data contained in a workspace by reassigning the workspace to a different capacity. The capacity can be in the same region or a different region.

Reassigning a workspace to another capacity

In the Fabric UI, workspaces can be moved to other capacities in the following ways:

- Fabric admins can reassign workspaces to a different capacity individually via the Workspaces page in the Fabric Admin portal.
- Fabric admins and capacity admins can reassign workspaces to a capacity in bulk via the **Workspaces assigned to this capacity** option in the capacity's settings.
- Workspace admins can reassign their workspace to a different capacity via the Workspace type option of the workspace settings.

Permission Requirements by capacity type to reassign workspaces

- **Power BI Premium (P) capacity** admin can effectively "claim" any workspace in the organization and pull it into their capacity by searching for the workspace name or assigning by user/group.
- **Fabric (F) Capacity** admin **must** also be a **Workspace Admin** of that specific workspace.
- **A capacity** admin **must** also be a **Workspace Admin** of that specific workspace.
- **Embedded (EM) capacity** admin, similar to P capacity, can "claim" any workspace.
- **Global Admin or Fabric Admin** can move any workspace to any capacity type.

Capacity Administrator can perform following operations by Capacity type:

Expand table

Capacity Type	Can Move Any Workspace?	Workspace Admin Role Required?	Management Tool
P (Premium)	**Yes**	No	Fabric Admin Portal

Capacity Type	Can Move Any Workspace?	Workspace Admin Role Required?	Management Tool
F (Fabric)	**No**	**Yes**	Fabric Portal, User API, or Admin API
A (Azure)	**No**	**Yes**	Fabric Portal, User API, or Admin API
EM (Office)	**Yes**	**No**	Fabric Admin Portal

Restrictions on moving workspaces around

Moving workspaces from one capacity to another has the following restrictions:

- When you move a workspace, all jobs related to items in the workspace get canceled.
- Only movable item types can move between regions. **If you're reassigning a workspace to a capacity located in a different region, you must remove all non-movable items first, otherwise reassignment will fail**.

  The following items types are movable:
  - Report
  - Semantic model (small storage format)
  - Dashboard
  - Dataflow Gen1
  - Paginated Report
  - Datamart
  - Scorecard

  All other item types can't be moved between regions and must be removed from the workspace before you can migrate the workspace to a capacity in another region.

  After you remove the non-movable items and the workspace is migrated to a different region, you can create new items of the non-movable type. It can take up to an hour after the migration before can do so.
- Only Power BI items can move from Premium capacity or Fabric to Power BI Pro or Power BI Premium Per-User (PPU) workspace type (with exceptions as noted here). If you're changing a workspace from Premium capacity or Fabric to Power BI Pro or Power BI Premium Per-User (PPU) workspace type, you must remove all non-Power BI items and any Power BI items that can't be moved first. Otherwise the workspace type change fails.

The following item types are considered Power BI items from the perspective of the workspace type.

- Report
- Semantic model (small storage format and large storage format)
- Dashboard
- Org app**
- Dataflow Gen1
- Paginated Report
- Metric set*
- Exploration**
- Datamart*
- Scorecard

*Can't move to Power BI Pro
**Can't move to Power BI Pro or Power BI Premium Per-User (PPU)

All other item types must be removed from the workspace before you can change its workspace type from Power BI Premium capacity or Fabric to Power BI Pro or Power BI Premium Per-User (PPU).

> **Note**
>
> If you have Dataflow Gen2 items in your workspace, their underlying staging lakehouse and staging warehouse items only become visible in the workspace UI after **all** Dataflow Gen2 items in the workspace are deleted. These staging items are Fabric items as well, and as such their existence can prevent the workspace from being successfully migrated from one region to another. To ensure that your workspace can be successfully migrated across regions, first delete all Dataflow Gen2 items in the workspace. Then delete all the staging lakehouses and warehouses in the workspace that become visible.

Discovering capacity reassignment issues

In some cases, a workspace might appear to be successfully reassigned to a new capacity, but some or all of its items remain attached to the original capacity. In this case, a warning banner appears when you open an affected item.

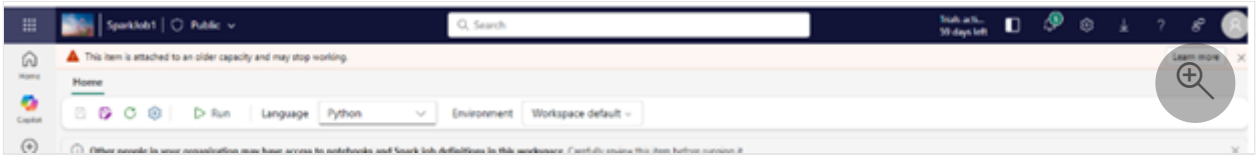

If you have access to the workspace settings, a similar warning banner appears when you view the **Workspace type** settings.

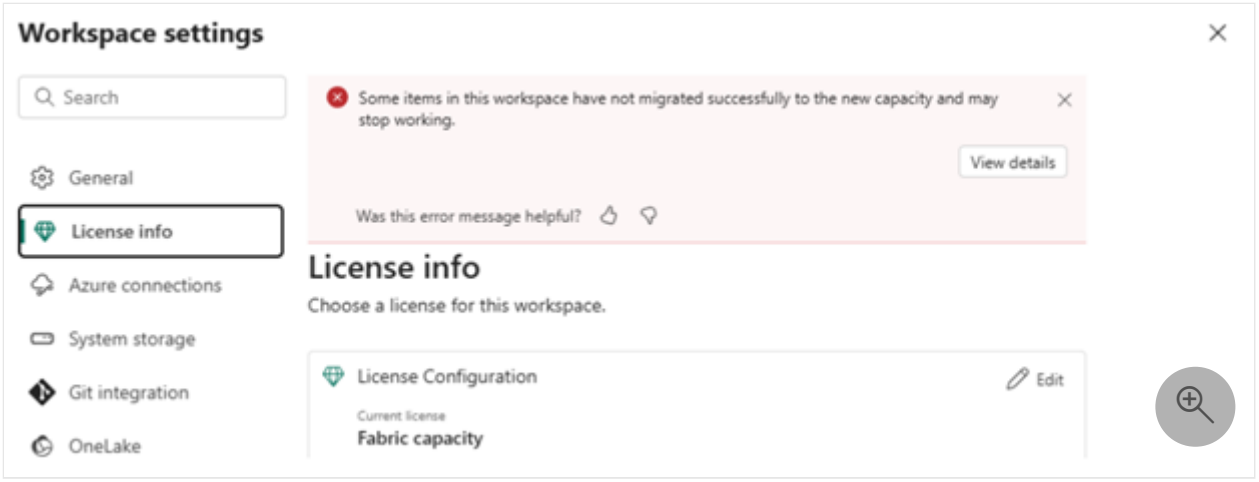

To see which items are affected, select **View details** in the warning banner. A detailed view opens showing all items that remain attached to a different capacity than the workspace. For each item, the capacity name, its region, and the reason the migration failed appears.

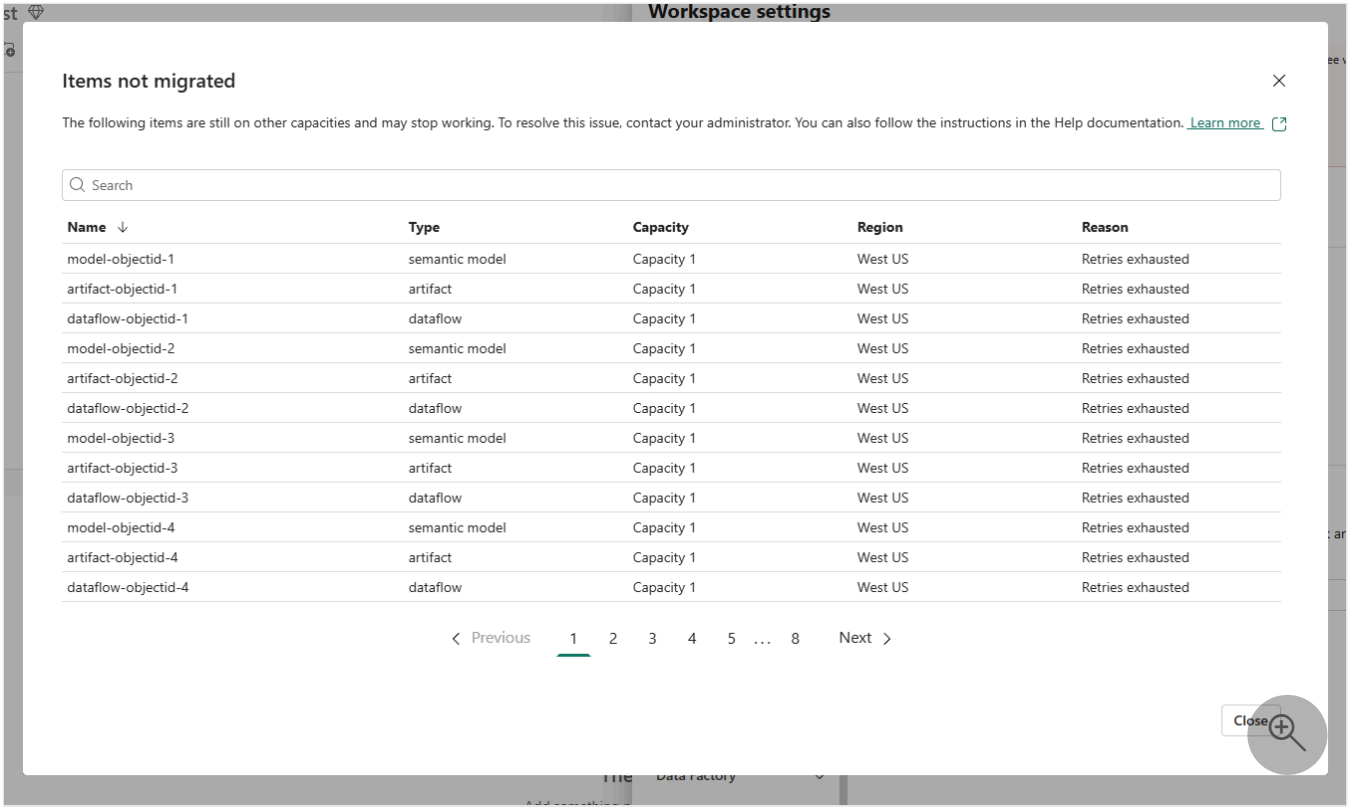

These items are at risk of becoming unavailable if their attached capacity is deleted or paused. To prevent disruption and resolve these issues, a workspace administrator should take the steps described in the next section.

Resolving common capacity reassignment issues

When a workspace reassignment fails for some or all items, the detailed view shows a reason for each failed migration. This section helps you understand and resolve these issues.

General troubleshooting steps

Before addressing specific errors, follow these general steps:

1. **Don't delete or pause the source capacity**. Items that failed to migrate remain attached to their original capacity (shown in the "Capacity" column). Deleting or pausing this capacity makes those items unavailable.
2. **Retry the reassignment**. Many failures result from transient errors. Try reassigning the workspace to the destination capacity again.
3. **Consider same-region migration**. If you're moving between regions, try moving to a different capacity in the same region first to rule out cross-region migration limitations.

Error types

The following sections describe common error types you might encounter during capacity reassignment, along with their causes and resolutions.

Unsupported migration

Why this happens: The workspace contains items or configurations that can't be migrated automatically. Common causes include:

- Private Link networking is enabled for the workspace.
- Customer-managed key (CMK) configuration issues.

How to resolve:

1. **Check Private Link settings:**
    - Go to **Workspace settings** > **Inbound networking**.
    - If Private Link is enabled, temporarily disable it.
    - Reassign the workspace.
    - Re-enable Private Link after successful migration.
2. **Review customer-managed key (CMK) configuration:**
    - Verify that the destination capacity supports the workspace's CMK settings.
    - Ensure your CMK permissions are correctly configured for the target region.
    - For more information, see Customer-managed keys.

Retries exhausted

Why this happens: The system attempted to migrate the item multiple times but encountered persistent failures, such as network timeouts, resource constraints, or capacity state changes during migration.

How to resolve:

1. **Verify capacity stability:**
    - Ensure both source and destination capacities exist, are active, and not paused.
    - Keep capacities in a stable state throughout the entire migration process.
    - Avoid scaling or pausing operations during reassignment.
2. **Retry the reassignment**. After confirming capacity stability, attempt the workspace reassignment again.
3. **Try a different capacity**. To prevent cross-region migration limitations from causing migration failures, try moving the workspace to a different capacity in the same region.

DB Migration error

Why this happens: This error is specific to semantic models (datasets) and indicates that the database migration component failed to transfer the model data. This error can occur due to network timeouts, resource constraints, or capacity state changes during migration.

How to resolve:

1. **Verify capacity stability:**
    - Ensure both source and destination capacities exist, are active, and not paused.
    - Keep capacities in a stable state throughout the entire migration process.
    - Avoid scaling or pausing operations during reassignment.
2. **Retry the reassignment**. After confirming capacity stability, attempt the workspace reassignment again.
3. **Try a different capacity**. Try moving the workspace to a different capacity in the same region to prevent cross-region migration limitations from causing migration failures.

Related articles

- Manage workspaces

ⓘ **Note:** The author created this article with assistance from AI. Learn more

Last updated on 02/12/2026

Retention and recovery in Fabric

Fabric has retention and recovery capabilities at both the workspace and item levels. When you delete workspaces or items, they enter a retention period during which you can restore them. These capabilities help protect against accidental deletions by giving you an opportunity to recover deleted content before it's permanently removed.

This article explains how retention and recovery work at both the workspace and item levels, so you can choose the right protection strategy for your Fabric content.

Workspace retention vs. item retention

The following table compares workspace retention and item-level soft-delete and recovery:

Expand table

Feature	Collaborative workspace retention	My workspace retention	Item retention
Applies when	User deletes a collaborative workspace	*My workspace* is deleted	User deletes a supported item
Default retention state	Enabled	Enabled	Enabled
Can an admin disable retention?	No	No	Yes
Minimum retention period	7 days	30 days (fixed)	3 days (when enabled)
Maximum retention period	90 days	30 days (fixed)	90 days

> **Note**
>
> Personal workspaces (*My workspaces*) have a fixed 30-day retention period. You can't disable or change this retention period.

Workspace retention

When you delete a workspace, Fabric doesn't permanently remove it immediately. Instead, it enters a configurable retention period during which you can restore it. At the end of the retention period, Fabric permanently removes the workspace and irreversibly deletes its contents.

The retention period for personal workspaces (*My workspaces*) is fixed at 30 days.

The default retention period for collaborative workspaces is seven days. You can change the retention period (from 7 to 90 days) by using the **Define workspace retention period** setting in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**.

Changing this setting also applies to collaborative workspaces that are already deleted. Fabric calculates the remaining retention from the current setting value rather than from the value that was in effect when the workspace was deleted. Decreasing the retention period can make an already deleted workspace eligible for immediate permanent deletion. For more information, see How changing the retention period affects workspaces that are already deleted.

During the retention period, you can restore a deleted workspace or permanently delete it before the retention period expires. For step-by-step instructions, see Set up and manage workspace retention.

Item soft-delete and recovery

> **Note**
>
> Item recovery is now enabled by default, with a retention period of three days. This default applies to tenants that never explicitly configured the **Fabric Item Recovery** setting and to newly created tenants. Fabric preserves settings that an admin explicitly configured, whether item recovery was explicitly turned on or off.

Fabric supports soft-delete and recovery for individual items within workspaces. When you delete a supported item, it enters a retention period during which workspace contributors, members, and admins can recover it. This feature provides an additional layer of data protection at the item level.

You can change the retention period (from 3 to 90 days) or turn item recovery off by using the **Fabric Item Recovery** setting in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**.

For step-by-step instructions on setting up item recovery and restoring or permanently deleting items, see Recover or permanently delete items.

Supported item types

The following item types support soft-delete and recovery:

- API for GraphQL
- Copy job
- Cosmos database
- Data agent *(preview)*
- Data pipeline
- Environment
- Graph index *(preview)*
- Graph queryset *(preview)*
- Healthcare data solutions
- KQL queryset
- Lakehouse
- Map *(preview)*
- Mirrored Azure Databricks catalog
- Mirrored database
- Mirrored Snowflake
- ML experiment
- ML model
- Notebook
- Ontology *(preview)*
- Operational agent *(preview)*
- Real-Time Dashboard
- Spark job definition
- SQL database *(preview)*
- User data function
- Variable library
- Warehouse
- Warehouse snapshot

For supported item types only, when the Item Recovery setting is on and someone deletes an item in a workspace through the Fabric UI or API, Fabric moves the item to a soft-deleted state. You can't recover unsupported item types because Fabric irreversibly deletes them.

Fabric plans to add support for more item types.

Permissions for item recovery and deletion

The following table shows the permissions required to recover or permanently delete items:

Expand table

Action	Workspace Viewer	Workspace Contributor	Workspace Member	Workspace Admin
Recover a deleted item	No	Yes	Yes	Yes
Permanently delete a recoverable item	No	No	No	Yes

> **Note**
>
> Different item types have different permissions for the delete operation. To learn more about the permissions for each item, see Item permissions.

What is restored when an item is recovered

When you recover an item, Fabric restores it to its original state before deletion. The recovery process preserves all properties, configurations, and relationships except shared item permissions.

OneLake catalog

Soft-deleted items don't appear in the catalog. After recovery, items reappear in the catalog with full functionality. The actions menu, edit capabilities, and hover cards work normally. Fabric restores Data Loss Prevention icons and policy tips. Favorite status and parent-child relationships are preserved. For more information, see OneLake catalog overview.

Lineage and relations

Soft-deleted items don't appear in lineage views. After recovery, Fabric restores full lineage relations. For more information, see Lineage.

Governance and administration

When you soft-delete or recover items, Fabric also disables or restores governance and administration features:

- **Endorsement and tags** - Fabric restores endorsement status after recovery. Tags are fully restored in all views, including the workspace list and workspace lineage. Fabric also restores tag-based filtering and all tag associations. For more information, see Endorsement overview.

- **Information protection** - Fabric reapplies Data Loss Prevention (DLP) policies after recovery and restores role-based access controls. Protection policies remain intact, sensitivity labels are preserved on recovered items, and last modified user information (UPN) is maintained. For more information, see Data loss prevention for Fabric and Sensitivity labels in Fabric.

- **Metadata scanning** - Metadata scanning doesn't return soft-deleted items. After recovery, scan results include the recovered items with all properties. For more information, see Metadata scanning.

Item share permissions

Fabric restores shared data after recovery, but it doesn't restore share permissions. After recovery, you need to re-share the item to make the shared data accessible again. For more information about sharing data warehouses, see Share your data and manage permissions.

Warehouse-specific recovery

When you recover a warehouse, Fabric restores all the metadata and data, but Fabric doesn't support snapshot restore. If you delete a warehouse, Fabric also deletes all the snapshots, and you can't recover them as part of the warehouse recovery process.

Billing during the retention period

Item recovery doesn't introduce a separate billing meter. During the retention period, soft-deleted items continue to incur standard OneLake storage costs, which are billed at the same rate as active data. In addition, soft-deleted items might consume small amounts of Capacity Units (CUs) from workload background maintenance activity. For more information, see OneLake compute and storage consumption.

To stop incurring costs for a soft-deleted item, permanently delete it before the retention period expires. For more information, see Recover or permanently delete items.

Limitations of soft-delete and recovery

- Only specific item types support soft-delete and recovery. See Supported item types for details.
- Shared item permissions aren't restored after soft-delete and recovery. You need to re-share the item to restore access for other users.
- You can't delete workspace folders that have soft-deleted items until you permanently delete all items in the folder.
- Only tenant admins can set up item retention settings. Workspace admins can't change these settings at the workspace level.
- When item recovery is enabled, admin insights on the Govern tab in the OneLake catalog reflect deleted items only after they're permanently deleted. Soft-deleted items that are still in the retention period don't appear as deleted in the insights.
- When item recovery is enabled, Fabric workspace item events are generated throughout the item lifecycle. Soft-delete operations generate `Microsoft.Fabric.ItemSoftDeleteSucceeded` or `Microsoft.Fabric.ItemSoftDeleteFailed` events, recovery operations generate `Microsoft.Fabric.ItemRecoverSucceeded` or `Microsoft.Fabric.ItemRecoverFailed` events, and permanent delete operations continue to generate `Microsoft.Fabric.ItemDeleteSucceeded` or `Microsoft.Fabric.ItemDeleteFailed` events.

Known issues

- If you use Git integration or deployment pipelines, a Git sync or pipeline deployment might re-create an item that was previously soft-deleted. The version created by Git contains only the item definition and doesn't include data. If you need to restore the item's data, recover the soft-deleted item from the recycle bin, delete the copy created by Git, and then re-sync with Git to reconcile the workspace state.

Related content

- Set up and manage workspace retention
- Recover or permanently delete items
- Manage workspaces
- Admin portal overview

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Set up and manage workspace retention

When you delete a workspace, it enters a retention period during which you can restore it. This retention capability helps protect against accidental deletions so that workspace content isn't permanently lost before you have a chance to recover it.

This article describes how to set up the retention period, restore deleted workspaces, and permanently delete workspaces before the retention period expires.

For an overview of how retention works for workspaces and items, see Retention and recovery in Fabric.

Prerequisites

- You need a tenant admin role to set up retention settings and restore or permanently delete workspaces.

Set up the retention period for deleted collaborative workspaces

By default, Fabric retains deleted collaborative workspaces for seven days. You can change the length of the retention period (from 7 to 90 days) by using the **Define workspace retention period** tenant setting.

1. Go to **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings** > **Workspace settings** > **Define workspace retention period**.
2. Turn on the setting and enter the number of days for the retention period. You can choose anywhere from 7 to 90 days.
3. Select **Apply**.

> **Note**
>
> When the **Define workspace retention period** setting is off, deleted collaborative workspaces still have a retention period of seven days.
>
> This setting doesn't affect the retention period of *My workspaces*. *My workspaces* always have a 30-day retention period.

How changing the retention period affects already deleted workspaces

The **Define workspace retention period** setting applies to every deleted collaborative workspace in the tenant, including workspaces that you already deleted before you changed the setting. Fabric doesn't store a copy of the retention period at the time of deletion. Instead, Fabric calculates the remaining retention for each deleted workspace by using the current value of the setting:

```
Remaining retention = current retention period - days since the workspace was deleted
```

Because Fabric recalculates the remaining retention from the current setting value, changing the setting takes effect immediately for workspaces that are already in a retention period:

- **When you increase the retention period**, deleted workspaces stay recoverable for longer. For example, if the setting is 30 days and you deleted a workspace 10 days ago, that workspace has 20 days of retention left. If you then change the setting to 60 days, the same workspace has 50 days of retention left.
- **When you decrease the retention period**, deleted workspaces become recoverable for a shorter time. Using the same example, if you change the setting from 30 days to 15 days, the workspace that you deleted 10 days ago has 5 days of retention left instead of 20.

> ⚠ **Warning**
>
> If you decrease the retention period so that a deleted workspace is already in its retention period for longer than the new value, that workspace becomes eligible for permanent deletion. Fabric permanently deletes it the next time it checks for expired workspaces, and there's no extra grace period. For example, if you deleted a workspace 30 days ago and you change the setting to seven days, Fabric permanently deletes that workspace and irreversibly deletes its contents.
>
> Before you decrease the retention period, go to **OneLake catalog** > **Govern** > **Workspaces** and restore any deleted workspace that you still need.

Fabric evaluates deleted workspaces for permanent deletion periodically rather than scheduling the permanent deletion in advance. Increasing the retention period extends a deleted workspace as long as the workspace is still in its retention period and shows the status **Deleted**. After Fabric begins permanently deleting a workspace, increasing the setting doesn't recover it.

Restore a deleted collaborative workspace

While a deleted collaborative workspace is in a retention period, you can restore it and its contents.

1. In **OneLake catalog** > **Govern** > **Workspaces**, find the deleted collaborative workspace you want to restore. Collaborative workspaces are of type **Workspace**. A workspace that is in a retention period has the status **Deleted**.
2. Select the workspace and then choose **Restore** from the ribbon, or select **More options (...)** and choose **Restore**.
3. In the **Restore workspaces** panel that appears, enter a new name for the workspace and assign at least one user the Admin role in the workspace.
4. Select **Restore**.

Permanently delete a deleted collaborative workspace during the retention period

While a deleted collaborative workspace is in a retention period, you can permanently delete it before the end of its retention period.

1. In **OneLake catalog** > **Govern** > **Workspaces**, find the deleted collaborative workspace you want to permanently delete. Collaborative workspaces are of type **Workspace**. A workspace that is in a retention period has the status **Deleted**.
2. Select the workspace and then choose **Permanently delete** from the ribbon, or select **More options (...)** and choose **Permanently delete**.

Fabric asks you to confirm the permanent deletion. After you confirm, the workspace and its contents are no longer recoverable.

> ⓘ **Note**
>
> If the workspace has a workspace identity, Fabric doesn't delete the identity while the workspace is in its retention period. Its service principal and app registration stay visible in the Azure portal until the workspace is permanently deleted. For more information, see Deleting the identity.

Restore a deleted My workspace as an app workspace

When your organization removes users from Microsoft Entra ID, their My workspaces show up as **Deleted** in the **State** column in **OneLake catalog** > **Govern** > **Workspaces**. You can restore deleted My workspaces as app workspaces that other users can collaborate in.

During this restoration process, you need to assign at least one workspace admin and give the new workspace a name. After you restore the workspace, it shows up as **Workspace** in the **Type** column in **OneLake catalog** > **Govern** > **Workspaces**.

1. In **OneLake catalog** > **Govern** > **Workspaces**, find the deleted personal workspace you want to restore.
2. Select the workspace and then choose **Restore** from the ribbon, or select **More options (...)** and choose **Restore**.
3. In the **Restore workspaces** panel that appears, enter a new name for the workspace and assign at least one user the Admin role in the workspace.
4. Select **Restore**.

After you restore the deleted workspace as an app workspace, it's just like any other app workspace.

Related content

- Retention and recovery in Fabric
- Recover or permanently delete items
- Manage workspaces
- Admin portal overview

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Recover or permanently delete items

Fabric supports soft delete and recovery for individual items within workspaces. When you delete a supported item, it enters a retention period during which workspace contributors, members, and admins can recover it. This protection helps prevent accidental data loss by giving you an opportunity to restore items before they're permanently removed.

In this article, you learn how to set up item retention, recover deleted items, and permanently delete items before the retention period expires.

For an overview of how retention works for workspaces and items, see Retention and recovery in Fabric.

Prerequisites

- To set up item retention settings, you need a tenant admin role.
- To recover deleted items, you need at least **Workspace Contributor** permissions. See Permissions for the full permissions matrix.
- To permanently delete recoverable items, you need **Workspace Admin** permissions.

Set up the retention period for deleted items

> **Note**
>
> Item recovery is now enabled by default, with a retention period of three days. The default applies to tenants that never explicitly configured the **Item Recovery** setting and to newly created tenants. Fabric preserves settings that an admin explicitly configured: if you already turned on item recovery, your current retention value stays the same, and if you explicitly turned it off, it stays off.

You can change the retention period (from 3 to 90 days) or turn item recovery off by using the **Item Recovery** tenant setting.

1. In the admin portal, go to **Tenant settings** > **Item Recovery**.
2. Make sure the setting is turned on, and then enter the number of days for the retention period (3 to 90 days).
3. Select **Apply**.

> **Note**
>
> When the **Item Recovery** setting is off, Fabric doesn't retain deleted items when you delete an individual item.
>
> The Item Recovery setting doesn't affect the retention period of workspaces. See Retention and recovery in Fabric for details about workspace retention.

Restore a deleted item

While a deleted item is in a retention period, workspace contributors, members, and admins can restore it and its properties.

Use the workspace recycle bin to restore an item

The workspace recycle bin is a dedicated view within each workspace where you can browse, restore, or permanently delete soft-deleted items.

1. Open the workspace that contained the deleted item.
2. Select **Recycle bin** from the workspace navigation.
3. Find the item you want to recover. You can use the filter and search options to locate specific items.
4. Select the item, then choose **Restore** from the toolbar.

The restored item reappears in its original workspace with its properties, configurations, and relationships intact.

Use the REST API to restore an item

Use the REST API to programmatically recover deleted items.

Request:

HTTP

```
POST
https://api.fabric.microsoft.com/v1/workspaces/{workspaceId}/recoverableItems/{itemId}
/recover
```

Parameters:

- `workspaceId` - The unique identifier of the workspace containing the deleted item.
- `itemId` - The unique identifier of the item to recover.

Example:

Bash

```
curl --request POST \
  --url https://api.fabric.microsoft.com/v1/workspaces/aaaabbbb-0000-cccc-1111-dddd2222eeee/recoverableItems/b1b1b1b1-cccc-dddd-eeee-f2f2f2f2f2f2/recover \
  --header 'Authorization: Bearer {token}'
```

Response:

- **200 OK** - Item successfully recovered.
- **404 Not Found** - Item not found or no longer in recovery period.
- **403 Forbidden** - User doesn't have permission to recover the item.

> ⓘ **Note**
>
> Recovered items retain their original properties, settings, and permissions. If the item depends on other deleted items, you might need to restore those dependencies first.

> ⓘ **Important**
>
> If a new item with the same name as the soft-deleted item exists in the same workspace, the recovery fails until you rename the existing item. To recover the deleted item, rename the new item, then retry the recovery operation.

Permanently delete a recoverable item

While a deleted item is in its retention period, you can permanently delete it if you have workspace admin permissions.

Use the workspace recycle bin to permanently delete an item

1. Open the workspace that contained the deleted item.
2. Select **Recycle bin** from the workspace navigation.
3. Find the item you want to permanently delete.
4. Select the item, then choose **Delete permanently** from the toolbar.

You're asked to confirm the permanent deletion.

> ⊗ **Caution**
>
> After you confirm permanent deletion, the item and its contents are no longer recoverable.
>
> After you permanently delete an item, OneLake retains the data for an additional seven days. You can't restore data during this period. The deletion is final and can't be reversed.

Use the REST API to permanently delete an item

Use the REST API to programmatically delete items during the retention period.

Request:

HTTP

```
DELETE
https://api.fabric.microsoft.com/v1/workspaces/{workspaceId}/recoverableItems/{itemId}
```

Parameters:

- `workspaceId` - The unique identifier of the workspace containing the deleted item.
- `itemId` - The unique identifier of the item to permanently delete.

Example:

Bash

```
curl --request DELETE \
  --url https://api.fabric.microsoft.com/v1/workspaces/aaaabbbb-0000-cccc-1111-
dddd2222eeee/recoverableItems/b1b1b1b1-cccc-dddd-eeee-f2f2f2f2f2f2 \
  --header 'Authorization: Bearer {token}'
```

> ⊗ **Caution**
>
> After you confirm permanent deletion, the item and its contents are no longer recoverable.
>
> After you permanently delete an item, OneLake retains the data for an additional seven days. You can't restore data during this period. The deletion is final and can't be reversed.

Related content

- Retention and recovery in Fabric
- Set up and manage workspace retention
- Manage workspaces
- Admin portal overview

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 08/31/2026

Microsoft Fabric tenant settings

Fabric tenant settings are configured in the tenant settings section of the Admin portal. For information about how to get to and use tenant settings, see About tenant settings.

Users can create Fabric items

Users can use production-ready features to create Fabric items. Turning off this setting doesn't affect users' ability to create Power BI items. This setting can be managed at both the tenant and the capacity levels.

For more information, see Enable Fabric for your organization.

Users can create Ontology (preview) items

Users can create ontologies to unify enterprise semantics across data, models, and logic to operationalize decision intelligence with context-aware AI agents.

For more information, see Ontology item overview.

Users can create digital twin builder (preview) items

Users can create digital twin builder items to build comprehensive digital twins of real-world environments and processes and support broad operational analysis.

For more information, see What is digital twin builder (preview)?.

Users can discover and create org apps (preview)

Turn on this setting to let users create org apps as items. Users with access can view them. If this setting is turned off, org app items are hidden until the setting is turned on again. The prior version of workspace apps remains available.

For more information, see Get started with org apps (preview).

Product Feedback

This setting allows Microsoft to prompt users for feedback through in-product surveys within Fabric and Power BI. User participation is voluntary.

For more information, see Learn about Fabric feedback.

Users can be informed of upcoming conferences featuring Microsoft Fabric when they are logged in to Fabric

When this setting is enabled, signed-in users in your organization can receive in-product notifications about upcoming conferences that feature Fabric.

ML models can serve real-time predictions from API endpoints (preview)

With this setting on, users can create real-time predictions from model and version endpoints. Even with real-time endpoints turned off, batch predictions can still be generated.

For more information, see Serve real-time predictions with ML model endpoints.

Detect anomalies in Real-Time Intelligence (Preview)

This setting allows users to use statistical detection algorithms to detect anomalies in real-time data.

For more information, see Detect anomalies in Real-Time Intelligence.

Users can create dbt job items (preview)

Users can import, author, and execute dbt projects directly within Fabric, enabling SQL-based transformation workflows without setting up a separate CLI environment.

For more information, see dbt job items in Fabric.

Enable operations agents

Users can create operations agents that use Azure OpenAI to create operations plans and recommend actions in response to real-time data. By turning on this setting, you agree to the preview terms for this feature. This setting can be managed at both the tenant and the capacity levels.

For more information, see Operations Agents.

All Power BI users can see "Set alert" button to create Fabric Activator alerts

When this setting is enabled, all Power BI users see the **Set alert** button in reports. Only users with permission to create Fabric items can actually create Fabric Activator alerts.

For more information, see Create Activator alerts from Power BI.

Related content

- About tenant settings
- Enable Fabric for your organization

Last updated on 08/31/2026

Product feedback

Product feedback allows users to give Microsoft feedback about Fabric. Product feedback is enabled by default.

Disable product feedback

Configure product feedback in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

To disable product feedback, follow the following steps:

1. From the tenant settings, expand **Product feedback**.
2. Change the *Enabled* switch to **Disabled**.
3. Select, **Apply**.

Related content

- Admin overview
- Learn about Fabric feedback

Last updated on 09/29/2026

Help and support tenant settings

These settings are configured in the tenant settings section of the Admin portal. For information about how to get to and use tenant settings, see About tenant settings.

Publish "Get Help" information

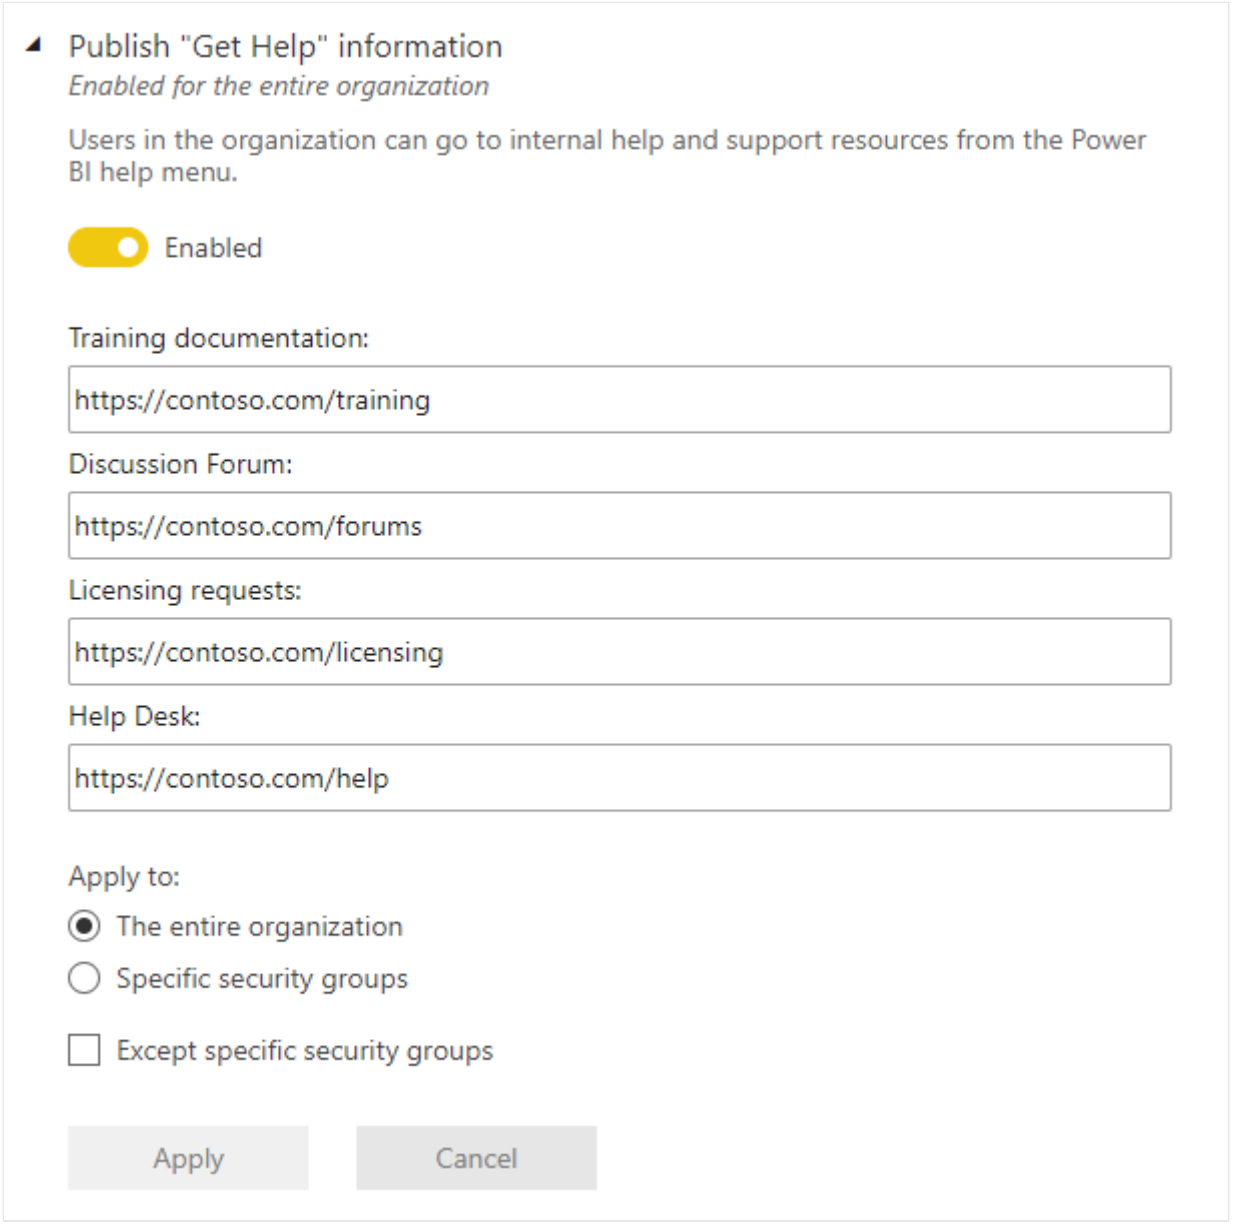

Admins can specify internal URLs to override the destination of links on the Power BI help menu and for workspace type upgrades. If custom URLs are set, users in the organization go to internal help and support resources instead of the default destinations. The following resource destinations can be customized:

- **Learn**. By default, this help menu link targets a list of all our Power BI learning paths and modules. To direct this link to internal training resources instead, set a custom URL for

Training documentation.

- **Community**. To take users to an internal forum from the help menu, instead of to the Fabric and Power BI Community, set a custom URL for **Discussion forum**.
- **Workspace type upgrades**. Users with a Fabric Trial workspace type can be presented with the opportunity to upgrade to Power BI Pro (Pro) or Power BI Premium Per-User (PPU). Users with a Fabric Trial or Power BI Pro workspace type can be presented with the opportunity to upgrade their account to a Power BI Premium Per-User (PPU) workspace type. If you specify an internal URL for **Licensing requests**, you redirect users to an internal request and purchase flow and prevent self-service purchase. You might want to prevent users from buying licenses, but are okay with letting users start a Power BI individual trial or a trial of a Fabric capacity. For this scenario, see Users can try Fabric paid features to separate the buy and try experiences.
- **Get help**. To take users to an internal help desk from the help menu, instead of to Fabric In-product Support, set a custom URL for **Help Desk**.

> **Note**
>
> The Fabric In-product Support center and the option to open support cases to Microsoft (Get Microsoft Help) will always be available for Admins.

Receive email and Teams notifications for service outages or incidents

If this tenant is impacted by a service outage or incident, mail-enabled security groups receive email and Teams notifications.

Learn more about Service interruption notifications.

Users can try Microsoft Fabric paid features

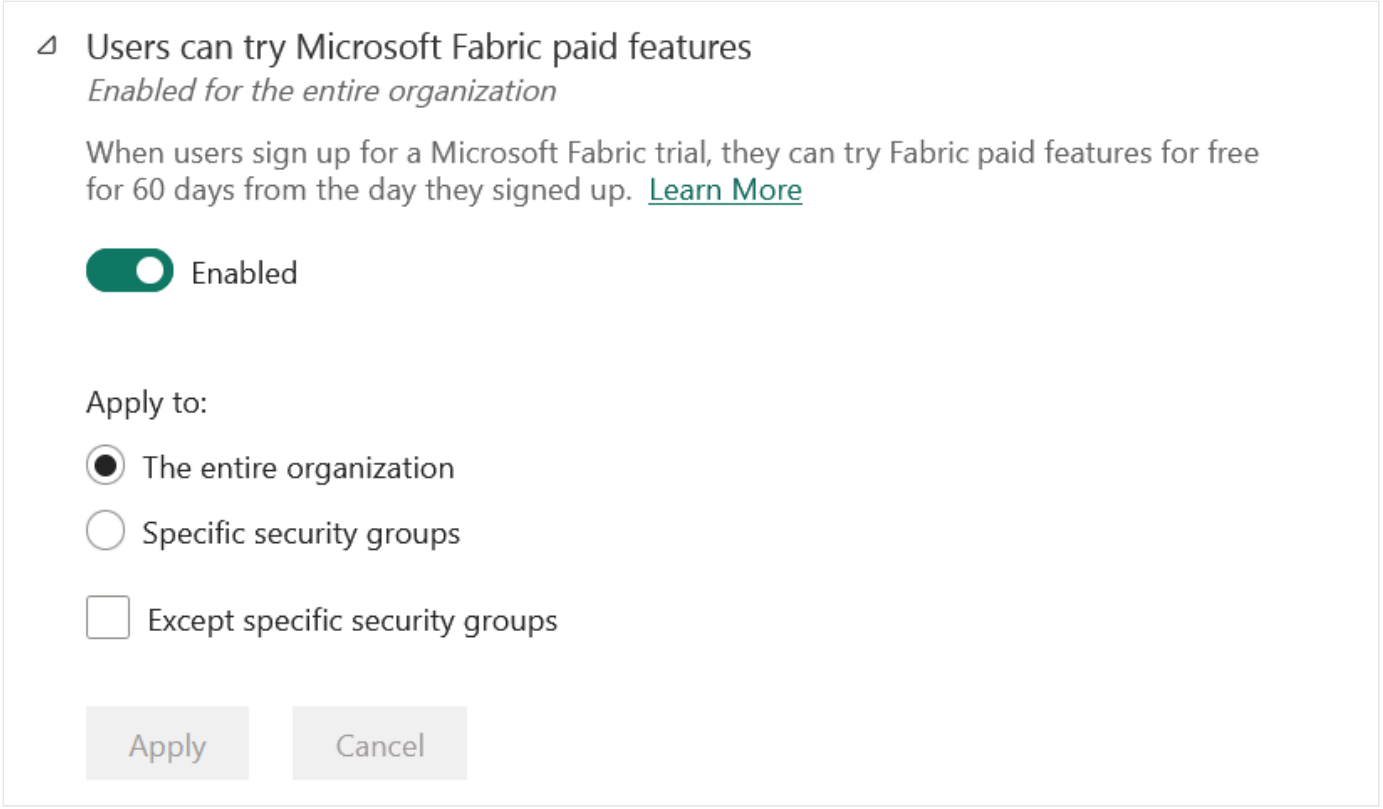

The setting to **Users can try Fabric paid features** is enabled by default. This setting increases your control over how users get license upgrades. In scenarios where you block self-service purchase, this setting lets users use more features free for 60 days. Users can start a Power BI individual trial or a trial of a Fabric capacity. Changing **Users can try Fabric paid features** from **enabled** to **disabled** blocks self-service trials of new licenses and of the Fabric capacity trial. It doesn't impact purchases that were already made.

The user's license upgrade and trial experience depends on how you combine license settings. The following table shows how the upgrade experience is affected by different setting combinations:

Expand table

Self-service purchase setting	Users can try Fabric paid features	End-user experience
Enabled	Disabled	User can buy an upgraded license, but can't start a trial
Enabled	Enabled	User can start a free trial and can upgrade to a paid license
Disabled	Disabled	User sees a message to contact the IT admin to request a license
Disabled	Enabled	User can start a trial, but must contact the IT admin

Self-service purchase setting	Users can try Fabric paid features	End-user experience
		to get a paid license

> **Note**
>
> You can add an internal URL for licensing requests in **Help and support settings**. If you set the URL, it overrides the default self-service purchase experience. It doesn't redirect sign-up for a trial. Users who can buy a license in the scenarios described in the table are redirected to your internal URL.

To learn more, see Enable or disable self-service sign-up and purchasing.

Show a custom message before publishing reports

Admins can provide a custom message that appears before a user publishes a report from Power BI Desktop. After you enable the setting, you need to provide a **custom message**. The custom message can be plain text or follow markdown syntax, as in the following example:

markdown

```
### Important Disclaimer

Before publishing the report to a workspace, be sure to validate that the appropriate users or groups have access to the destination workspace. If some users or groups should *not* have access to the content and underlying artifacts, remove or modify their access to the workspace, or publish the report to a different workspace. Learn about [giving access to workspaces](/power-bi/collaborate-share/service-give-access-new-workspaces).
```

The **custom message** text area supports scrolling, so you can provide a message up to 5,000 characters.

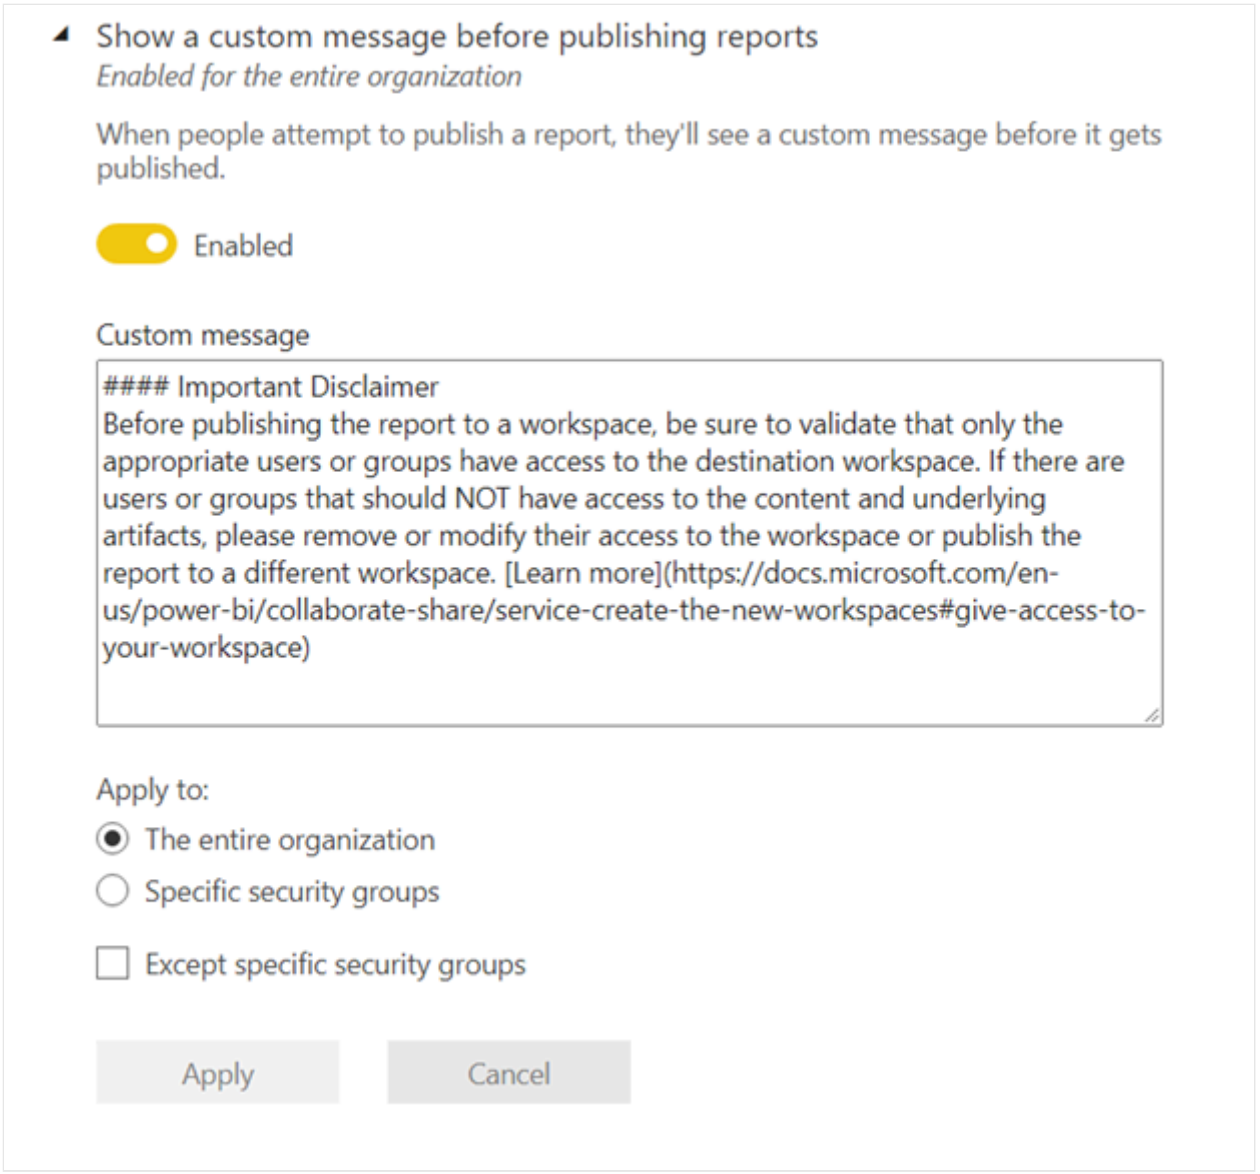

When your users publish reports to workspaces in Power BI, they see the message you wrote.

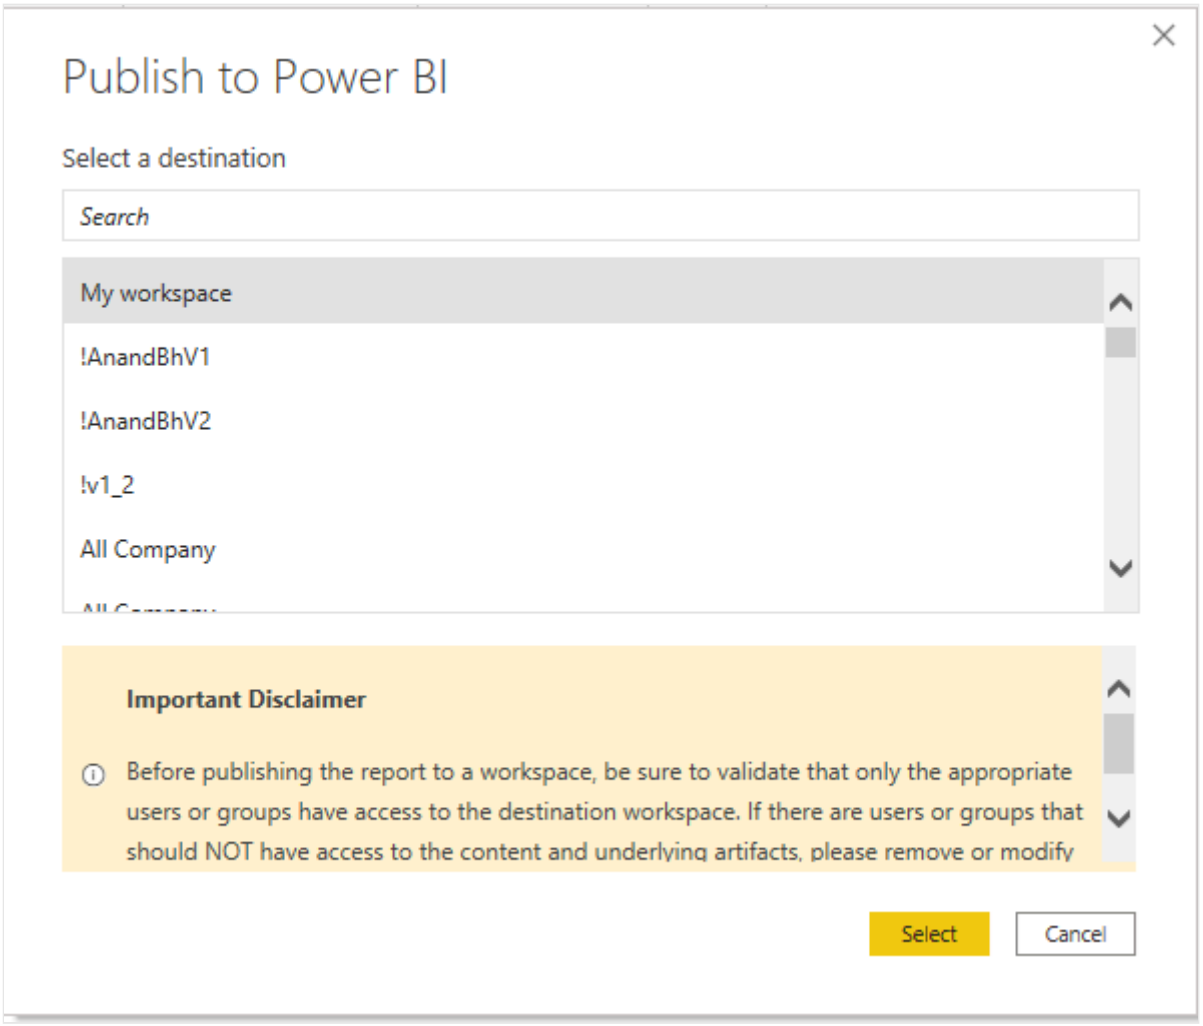

As with other tenant settings, you can choose who the **custom message** applies to:

- **The entire organization**.
- **Specific security groups**.
- Or **Except specific security groups**.

Related content

- About tenant settings
- Service interruption notifications

Last updated on 07/01/2026

Domain management tenant settings

Domain management tenant settings are configured in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

Allow tenant and domain admins to override workspace assignments (preview)

This setting controls whether tenant and domain admins can override existing workspace domain assignments. When disabled, tenant and domain admins cannot reassign a workspace that is already assigned to a domain to another domain. When enabled, they can override such assignments. The setting is enabled by default. The domain REST APIs respect this setting.

To enable or disable the setting, go to **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings** > **Domain management settings**, expand **Allow tenant and domain admins to override workspace assignments (preview)**, and set the toggle as desired.

Related content

- Fabric domains
- About tenant settings

Last updated on 09/29/2026

Workspace tenant settings

As a Fabric administrator, you can configure workspace settings in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

Create workspaces

Workspaces are places where users collaborate on dashboards, reports, and other content. Microsoft Fabric admins can use the **Create workspaces** setting to designate which users in the organization can create workspaces. Admins can let everybody or nobody in an organization create workspaces. Workspace creation can also be limited to members of specific security groups. Learn more about workspaces.

List of workspaces

The **Workspaces** section in **OneLake catalog** > **Govern** has settings for the workspaces in your tenant. In that section, you can sort and filter the list of workspaces and display the details for each workspace. See Manage workspaces for details.

Publish apps

In the tenant settings, you also control which users have permissions to distribute apps to the organization. See Publish apps to the entire organization for details.

Use semantic models across workspaces

Admins can control which users in the organization can use semantic models across workspaces. When this setting is enabled, users still need the required Build permission for a specific semantic model.

For more information, see Intro to semantic models across workspaces.

Identify your workspace ID

The easiest way to find your workspace ID is in the URL of the Fabric site for an item in a workspace. As in Power BI, the Fabric URL contains the workspace ID, which is the unique identifier after `/groups/` in the URL, for example: `https://powerbi.com/groups/11aa111-a11a-1111-`

`1abc-aa1111aaaa/...`. Alternatively, you can find the workspace ID in **Admin portal** > **Workspaces** by selecting the **Actions** menu > **Details** next to the workspace name.

Block users from reassigning personal workspaces (My Workspace)

Personal workspaces are the My workspaces that every user has for their personal content. Microsoft Fabric and capacity admins can designate a preferred capacity for My workspaces. By default, however, My workspace owners can still change the capacity assignment of their workspace. If a Microsoft Fabric or capacity admin designates a Premium capacity as the default capacity for My workspaces, but a My workspace owner then changes that capacity assignment back to shared capacity, this could result in non-compliance with data residency requirements.

To prevent such a scenario, the Microsoft Fabric admin can turn on the **Block users from reassigning personal workspaces (My Workspace)** tenant setting. When this setting is on, My workspace owners can't change the capacity assignment of their My workspace.

To turn on the setting, sign in to Fabric using your admin account credentials. Open the **OneLake catalog** and go to the **Govern** tab > **Configurations** > **Tenant settings**. Then scroll to the **Workspace settings** section, and find **Block users from reassigning personal workspaces (My Workspace)**.

For more information, see Prevent My workspace owners from reassigning their My workspaces to a different capacity.

Automatically convert and store reports using Power BI enhanced metadata format (PBIR) (preview)

Turn on this setting to automatically convert reports to Power BI enhanced metadata format (PBIR) after they are edited and save them in the workspace.

When this setting is enabled, new reports are also created in PBIR format. PBIR provides a source-control-friendly file structure that can improve co-development workflows and report lifecycle management.

This setting is on by default, but it doesn't take effect before early 2026. If you opt out before then, use of PBIR in the service remains off until general availability.

For more information, see Power BI enhanced report format (PBIR).

Related content

- About tenant settings
- CI/CD workflow options in Fabric
- Set up and manage workspace retention

Last updated on 09/29/2026

Information protection tenant settings

Information protection tenant settings help you to protect sensitive information in your Power BI tenant. Allowing and applying sensitivity labels to content ensures that information is only seen and accessed by the appropriate users. These settings are configured in the tenant settings section of the Admin portal. For information about how to get to and use tenant settings, see About tenant settings.

Allow users to apply sensitivity labels for content

With this setting enabled, specified users can apply sensitivity labels from Microsoft Purview Information Protection.

All prerequisite steps must be completed before enabling this setting.

Sensitivity label settings, such as encryption and content marking for files and emails, aren't applied to content. Sensitivity labels and protection are only applied to files exported to Excel, PowerPoint, or PDF files that are controlled by **Export to Excel** and **Export reports as PowerPoint presentation or PDF documents** settings. All other export and sharing options don't support the application of sensitivity labels and protection.

To learn more, see Sensitivity labels in Power BI.

To view sensitivity label settings for your organization, visit the Microsoft Purview portal.

Apply sensitivity labels from data sources to their data in Power BI

When this setting is enabled, Power BI semantic models that connect to sensitivity-labeled data in supported data sources can inherit those labels, so that the data remains classified and secure when brought into Power BI.

To learn more about sensitivity label inheritance from data sources, see Sensitivity label inheritance from data sources.

Automatically apply sensitivity labels to downstream content

When a sensitivity label is applied to a semantic model or report in the Power BI service, it's possible to have the label trickle down and be applied to content that's built from that semantic model or report.

To learn more, see Sensitivity label downstream inheritance.

Allow workspace admins to override automatically applied sensitivity labels

Fabric admins can enable the **Allow workspace admins to override automatically applied sensitivity labels** tenant setting. This makes it possible for workspace admins to override automatically applied sensitivity labels without regard to label change enforcement rules.

To learn more, see Relaxations to accommodate automatic labeling scenarios.

Restrict content with protected labels from being shared via link with everyone in your organization

When this setting is enabled, users can't generate a sharing link for **People in your organization** for content with protection settings in the sensitivity label.

> **Note**
>
> This setting is disabled if you haven't enabled both the **Allow users to apply sensitivity labels for Power BI content** setting and the **Allow shareable links to grant access to everyone in your organization** setting. Additionally, this setting only applies to sharing via **People in your organization** links; it does not apply to sharing of Apps to the entire organization.

Sensitivity labels with protection settings include encryption or content markings. For example, your organization might have a *Highly Confidential* label that includes encryption and applies a *Highly Confidential* watermark to content with this label. Therefore, when this tenant setting is enabled and a report has a sensitivity label with protection settings, then users can't create sharing links for **People in your organization**:

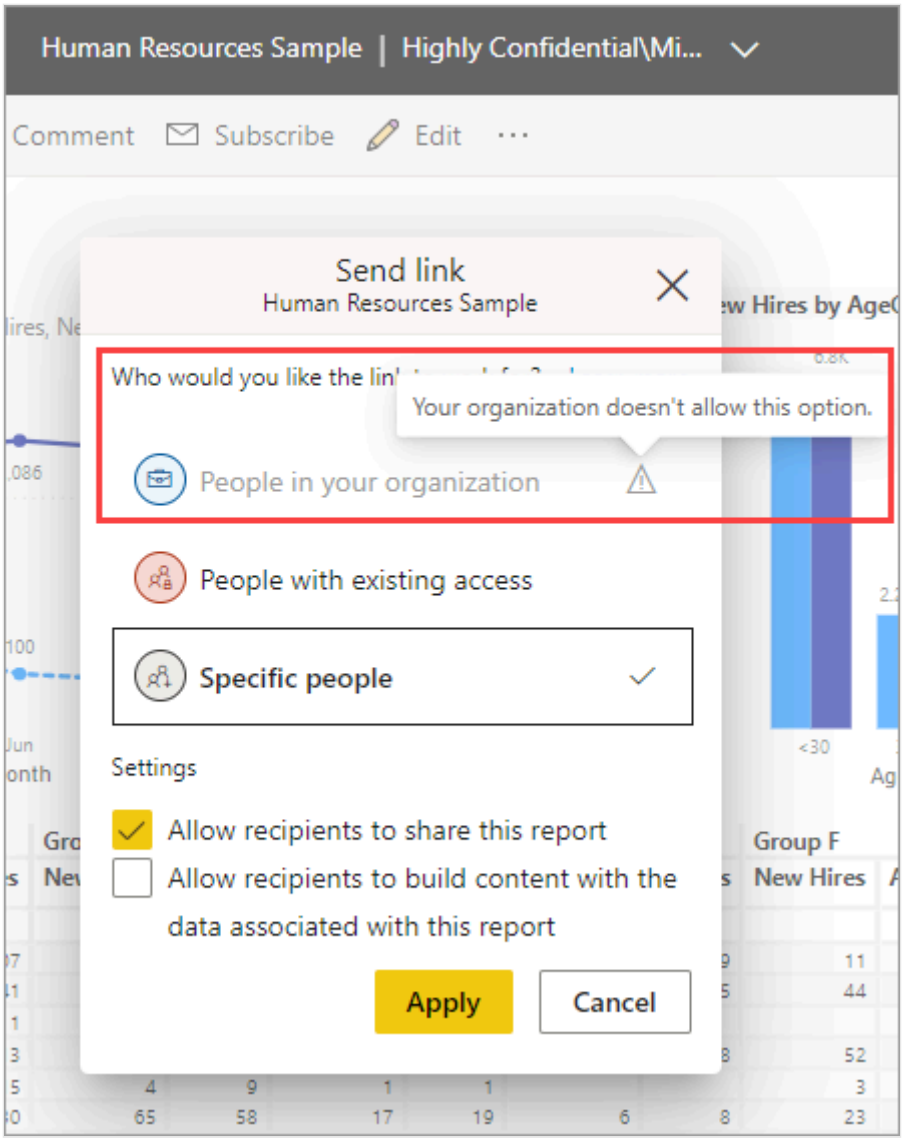

To learn more about protection settings for sensitivity labels, see Restrict access to content by using sensitivity labels to apply encryption.

Domain admins can set default sensitivity labels for their domains (preview)

Domain admins can set a default sensitivity label for their domains. The label they set overrides your organization's default labels in Microsoft Purview as long as it has a higher priority than the existing default labels set for your tenant.

A domain's default label automatically applies to new Fabric items created within the domain. Reports, semantic models, dataflows, dashboards, scorecards, and some additional item types aren't currently supported.

For more information, see Domain default sensitivity labels.

Allow Microsoft Purview to secure AI interactions

Turn on this setting to allow Microsoft Purview to access, process, and store prompts and responses, including metadata, for data security and compliance scenarios such as sensitive info type classification, reporting in Microsoft Purview Data Security Posture Management for AI, Audit, Insider Risk Management, Communication Compliance, and eDiscovery.

This capability is part of Microsoft Purview and isn't included in Copilot in Fabric pricing.

For more information, see Microsoft Purview for AI interactions.

Related content

- About tenant settings
- Protected sensitivity labels in Fabric and Power BI

Last updated on 04/09/2026

Azure Maps Services settings

Tenant settings control whether members of your organization can use Map items. Administrators can enable or disable this feature to allow or restrict access. If enabled, users can create, view, and interact with Map items in their Fabric environment. If disabled, Map item functionality is unavailable to all users in the tenant.

Fabric includes three tenant settings that determine whether Azure Maps services are available in Map items and Map visuals. These settings affect user access to map-powered features, including Map items and Map visuals within notebooks.

Access to Azure Maps Services

When the **Users can use Azure Maps services** setting is enabled, users can access both the Map item and the Map visual in notebooks. These features, powered by Azure Maps, allow users to create location-aware reports, monitor real-time telemetry, and explore spatial data directly within Fabric.

This setting determines whether members of your organization, or specific security groups, can access Azure Maps-powered experiences within Fabric. When enabled, users can utilize features such as Interactive Maps and geospatial analytics across Fabric components like Map items, Map visuals, and notebooks. Disabling this setting restricts access to these location-aware capabilities, helping administrators manage data usage and compliance based on organizational needs.

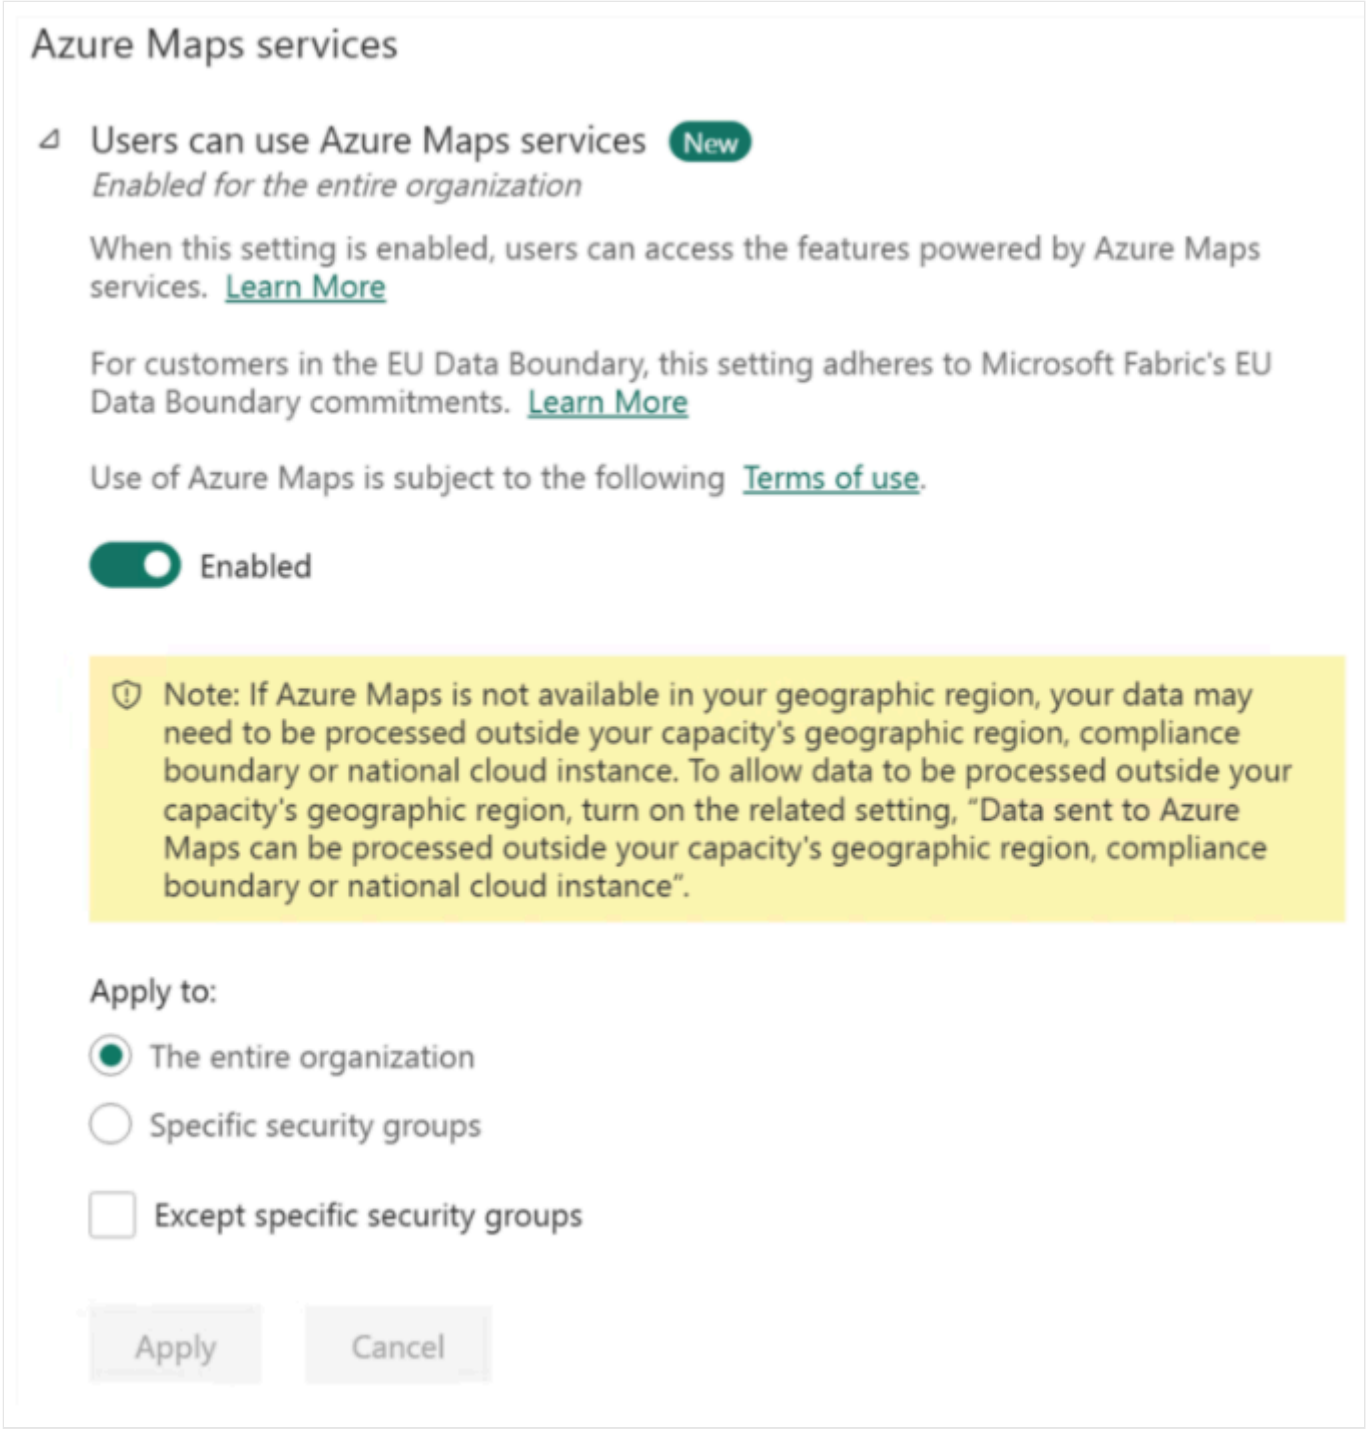

Default: Enabled

> **Note**
>
> Even when this setting is turned off, users can still open Map items. In this limited experience, the map background is blank because Azure Maps basemaps aren't available, but added data layers continue to display. Turn the setting back on to restore the full Azure Maps experience.

Configure global data processing

The **Data sent to Azure Maps can be processed outside your capacity's geographic region, compliance boundary or national cloud instance** setting is relevant only for customers who plan to use Map items and Map visuals in notebooks powered by Azure Maps, and whose capacity's geographic region is outside the supported Azure Maps service regions.

When enabled, this setting allows data sent to Azure Maps to be routed to the nearest available region, which may reside outside your organization's compliance boundary or national cloud instance. This ensures uninterrupted access to Azure Maps capabilities in unsupported regions.

> ⓘ **Note**
>
> This setting alone does not enable Azure Maps. You must also enable the setting **Users can use Azure Maps services**.

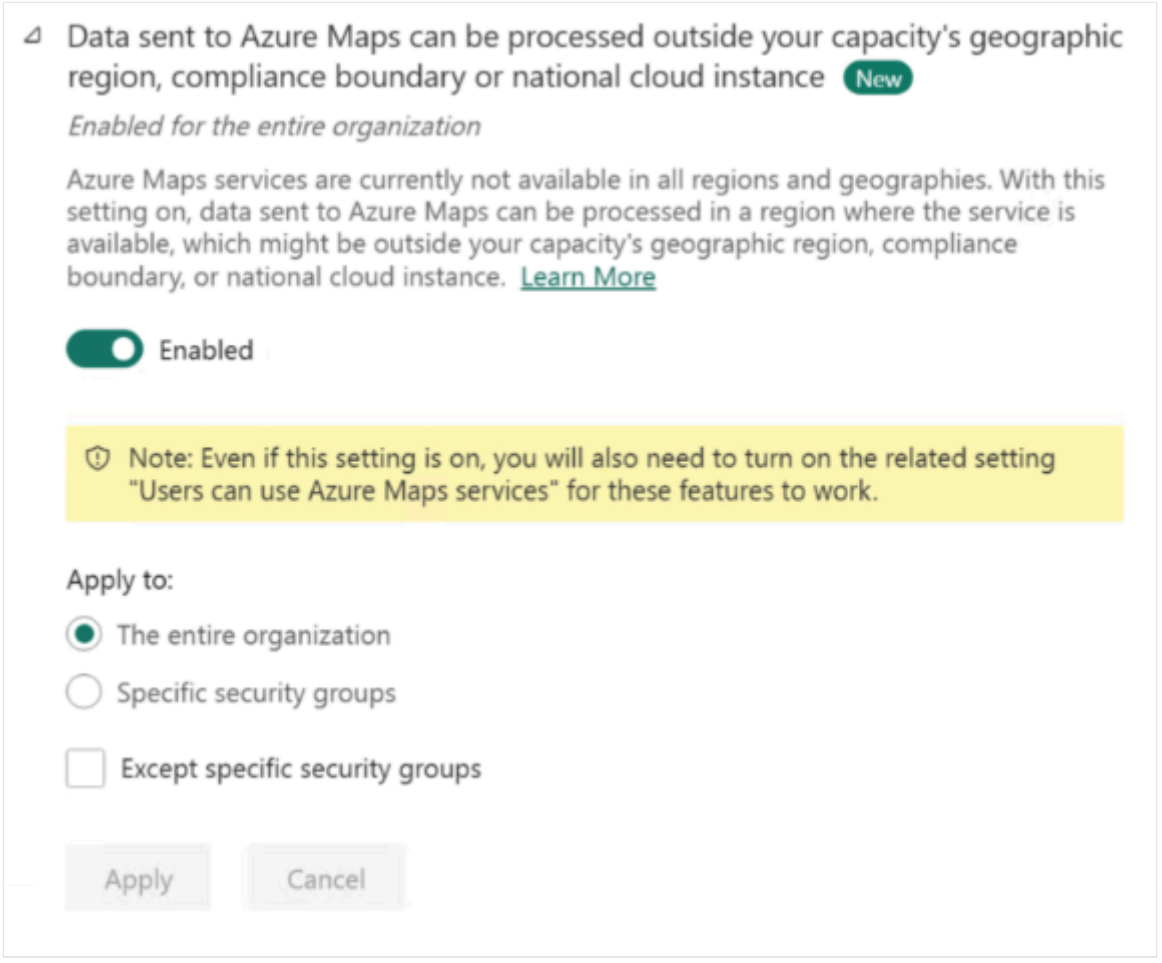

Default: Disabled

> ⓘ **Note**
>
> If your capacity's geographic region isn't a supported Azure Maps region and this setting is disabled, users can still open Map items in a limited experience. The map displays a blank

background instead of an Azure Maps basemap, but added data layers remain visible. Re-enable this setting to restore the full Azure Maps experience.

> **Note**
>
> Azure Maps does not process or transmit any customer names or personally identifiable information (PII).

Access to Azure Maps Weather Services (Preview)

The **Users can use Azure Maps Weather Services (Preview)** setting controls whether users can access weather data in Map experiences that use Azure Maps Weather, sourced from AccuWeather.

This setting is relevant only for organizations that already use Azure Maps-powered features in Fabric. It doesn't turn on Azure Maps by itself. To use Azure Maps Weather Services, you must also enable **Users can use Azure Maps services**.

For more information, see Azure Maps Weather services.

Note: The author created this article with assistance from AI. Learn more

---

Last updated on 08/20/2026

Export and sharing settings

The export and sharing settings allow the Fabric administrator the flexibility to determine and allow Power BI content to export to formats within their organization's security and compliance guidelines. These settings also allow you to keep unauthorized export formats from being exported by users.

Sharing settings are also managed through these settings. You can determine how and who can share Power BI content in your organization, as well as determine settings for sharing content with users outside your organization. These settings are configured in the tenant settings section of the Admin portal. For information about how to get to and use tenant settings, see About tenant settings.

External data sharing

When you turn on this setting, the users you specify will be able to share read-only links to data stored in your organization's OneLake storage with collaborators both outside and inside your organization. Collaborators who receive the link will be able to view, build on, and share the data both within and beyond their own Fabric tenants, using their organization's licenses and capacities.

External data sharing has important security considerations. For more information about external data sharing, see External data sharing.

To turn on external data sharing, open the **Admin portal**, go to **Tenant settings**, expand **External data sharing** under **Export and sharing settings**, set the toggle to **Enabled**, and then specify which users can create external data shares.

Users can accept external data shares

When you turn on this setting, the users you specify will be able to accept read-only links to data from another organization’s Fabric tenant. Users who accept an external share link can view, build on, and share this data, both inside and outside of your organization. For more information about external data sharing and its security considerations, see External data sharing.

To allow users to accept external data shares, open the **Admin portal**, go to **Tenant settings**, expand **Users can accept external data shares** under **Export and sharing settings**, set the

toggle to **Enabled**, and then specify which users can accept external data shares.

This setting is unrelated to **Allow specific users to turn on external data sharing**, which refers to sharing Power BI semantic models via Entra B2B.

Guest users can access Microsoft Fabric

When you turn on this setting, Microsoft Entra Business-to-Business (Microsoft Entra B2B) guest users can access Fabric. If you turn off this setting, B2B guest users receive an error when trying to access Fabric and any Fabric items they have permissions to. Disabling this setting for the entire organization also prevents users from inviting guests via item sharing and permissions. Use the specific security groups option to control which B2B guest users can access Fabric.

To learn more, see Distribute Power BI content to external guest users with Microsoft Entra B2B.

Users can invite guest users to collaborate through item sharing and permissions

This setting helps organizations choose whether new guest users can be invited to the organization through Fabric sharing, permissions, and subscription experiences.

To invite external users to your organization, the user must also have the Microsoft Entra Guest Inviter role. Once invited, external users become Microsoft Entra B2B guest users. This setting only controls the ability to invite through Fabric.

To learn more, see Invite guest users.

> ⓘ **Important**
>
> This setting was previously called **Share content with external users**.

Guest users can browse and access Fabric content

This setting allows Microsoft Entra B2B guest users to have full access to the browsing experience using the left-hand navigation pane in the organization. Guest users who have been assigned workspace roles or specific item permissions continue to have those roles and/or permissions, even if this setting is disabled.

To learn more about sending Fabric content to Microsoft Entra B2B guest users, read Distribute Power BI content to external guest users with Microsoft Entra B2B.

Users can see guest users in lists of suggested people

This setting helps organizations limit visibility of external users in sharing experiences. When disabled, Microsoft Entra guest users aren't shown in people picker suggested users lists. This helps prevent accidental sharing to external users and seeing which external users have been added to your organization through Power BI sharing UIs.

> **Important**
>
> When the setting is disabled, you can still give permission to a guest user by providing their full email address in people pickers.

Publish to web

People in your organization can publish public reports on the web. Publicly published reports don't require authentication to view them.

Only admins can allow the creation of new publish-to-web embed codes. Go to Embed codes in the admin portal to review and manage public embed codes. If any of the codes contain private or confidential content remove them. Review embed codes regularly to make sure no confidential information is live on the web.

The **Publish to web** setting in the admin portal gives options for which users can create embed codes. Admins can set **Publish to web** to **Enabled** and **Choose how embed codes work** to **Allow only existing embed codes**. In that case, users can create embed codes, but they have to contact the admin to allow them to do so.

Users see different options in the UI based on the **Publish to web** setting.

Expand table

Feature	Enabled for entire organization	Disabled for entire organization	Specific security groups
Publish to web under report **More options**	Enabled for all	Not visible for all	Only visible for authorized users or groups.

Feature	Enabled for entire organization	Disabled for entire organization	Specific security groups
(...) menu			
Manage embed codes under **Settings**	Enabled for all	Enabled for all	Enabled for all - **Delete** option only for authorized users or groups. - **Get codes** enabled for all.
Embed codes within admin portal	Status has one of the following values: - Active - Not supported - Blocked	Status displays **Disabled**	Status has one of the following values: - Active - Not supported - Blocked If a user isn't authorized based on the tenant setting, status displays **infringed**.
Existing published reports	All enabled	All disabled	Reports continue to render for all.

Learn more about publishing to the web.

Copy and paste visuals

Turn on this setting to allow users in the organization to copy visuals from a tile or report visual and paste them as static images into external applications.

Export to Excel

Users in the organization can export the data from a visualization to an Excel file.

To learn more, see Export the data that was used to create a visualization.

> **Note**
>
> Fabric automatically applies a sensitivity label on the exported file and protects it according to the label's file encryption settings.

Export to .csv

Users in the organization can export data from a tile, visualization, or paginated report to a *.csv* file.

To change this setting, go to **Tenant settings**, find **Export to .csv** under **Export and sharing settings**, choose whether the setting is on or off, select the user scope under **Apply to**, and then select **Apply**.

Download reports

Users in the organization can download .pbix files and paginated reports.

To learn more, see Download a report from the Power BI service to Power BI Desktop.

Users can work with semantic models in Excel using a live connection

Turn this setting on to allow users to export data to Microsoft Excel from a Power BI visual or semantic model, or export a semantic model to an Excel workbook with Analyze in Excel, both options with a live connection to the XMLA endpoint.

To learn more, see Create Excel workbooks with refreshable Power BI data.

Export reports as PowerPoint presentations or PDF documents

This setting lets users export reports as PowerPoint presentations or PDF documents.

To learn more, see export PowerPoint presentations and export PDF documents.

Export reports as MHTML documents

Users in the organization can export paginated reports as MHTML documents when this setting is turned on.

Export reports as Word documents

This setting lets users in the organization export paginated reports as Microsoft Word documents.

To learn more, see Export Power BI paginated report to Microsoft Word.

Export reports as XML documents

This setting lets users in the organization export paginated reports as XML documents.

To learn more, see Export Power BI paginated report to XML.

Export reports as image files

Users in the organization can use the *export report to file* API to export reports as image files.

To learn more, see Export Power BI paginated report to an Image File.

Print dashboards and reports

This setting lets users in the organization print dashboards and reports.

To learn more, see Print from the Power BI service.

Certification

Choose whether people in your organization or specific security groups can certify items like apps, reports, or datamarts as trusted sources for the wider organization.

> **Important**
>
> When a user certifies an item, their contact details are visible along with the certification badge.

Read Enable content certification for more details.

Endorse master data

Choose whether people in your organization or specific security groups can endorse items such as lakehouses, warehouses, or datamarts as one of the core sources for your organization's data records.

When someone endorses an item as master data, their name and email are shown with the endorsement badge.

For more information, see Enable master data endorsement.

Users can set up email subscriptions

This setting lets users create email subscriptions to reports and dashboards. Read Email subscriptions for reports and dashboards in the Power BI service to learn more.

B2B guest users can set up and be subscribed to email subscriptions

There may be instances that admin may want B2B guest users to receive email subscriptions but not other external users. Use this setting to allow B2B guest users to set up and subscribe themselves to email subscriptions.

If this setting is off, only users in your organization can create and receive email subscriptions.

The **Users can send email subscriptions to external users** switch is automatically turned off if the **B2B guest users can set up and be subscribed to email subscriptions** switch is turned off. This is because B2B users are external users that have been granted elevated permissions to get content. Since B2B guest users have higher permissions than other external users, if they can't get the email subscription neither can the other external users.

Users can send email subscriptions to external users

External users are users who haven't been added to your Microsoft Entra ID. Turn this setting off to prevent users from subscribing external users to email subscriptions.

Featured content

This setting lets you enable or disable the ability of users in your organization to promote their published content to the **Featured** section of the Power BI Home page. By default, anyone with the Admin, Member, or Contributor role in a workspace in your organization can feature content on Power BI Home.

To learn more, see Feature content on colleagues' Power BI Home page.

You can also manage featured content on the **Featured content** page in the Admin portal. Go to Manage featured content for more details.

Allow connections to featured tables

This setting lets Fabric admins control who in the organization can use featured tables in the Excel Data Types Gallery. Read more about Power BI featured tables in Excel.

Connections to featured tables are also disabled if the **Allow live connections** setting is set to Disabled.

Allow shareable links to grant access to everyone in your organization

This tenant setting is available for admins looking to disable creating shareable links to **People in your organization**.

If this setting is turned off for a user with permissions to share a report, that user can only share the report via link to **Specific people** or **People with existing access**. The following image shows what that user sees if they attempt to share the report via link:

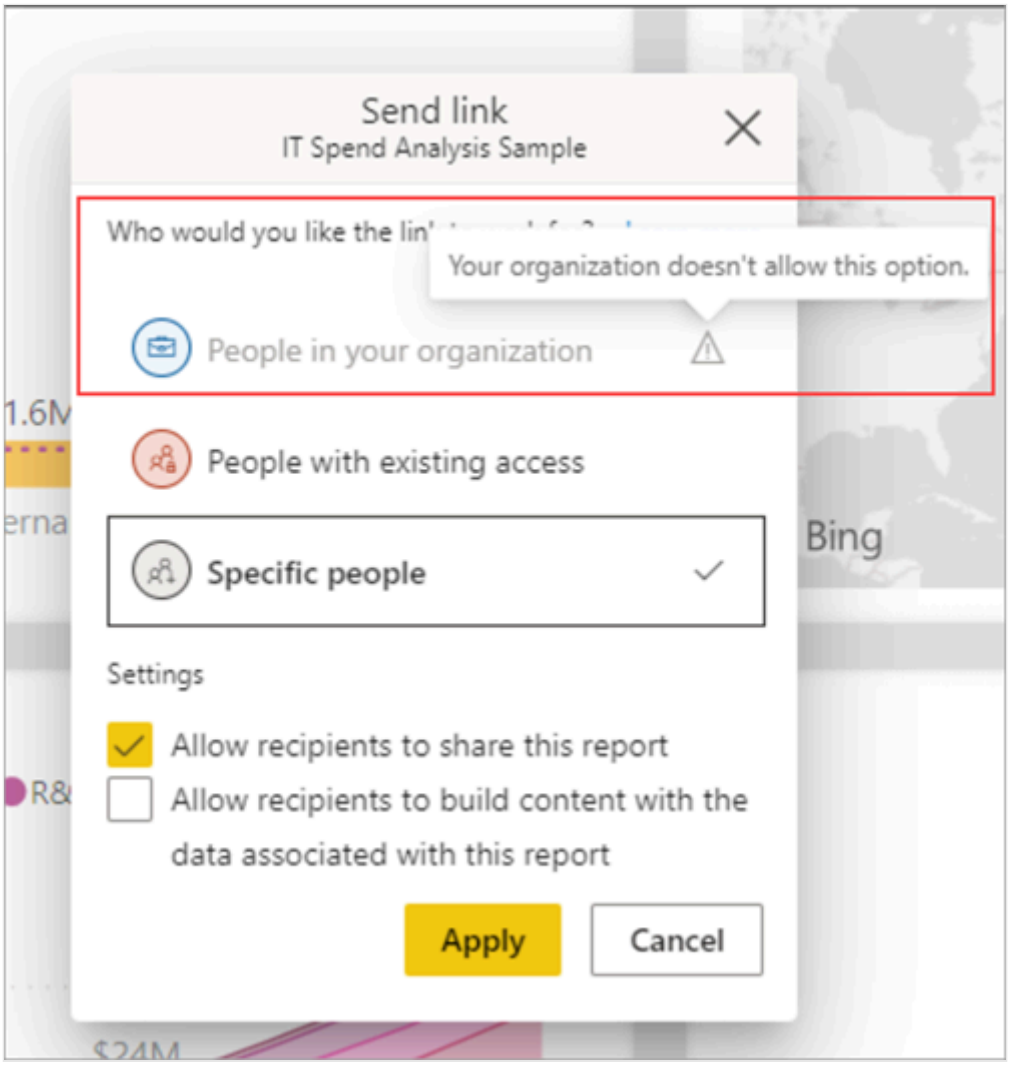

To learn more, see Link settings.

Enable Microsoft Teams integration

This setting allows organizations to access features that work with Microsoft Teams and the Power BI service. These features include launching Teams experiences from Power BI like chats, the Power BI app for Teams, and getting Power BI notifications from Teams. To completely enable or disable Teams integration, work with your Teams admin.

Read more about collaborating in Microsoft Teams with Power BI.

Install Power BI app for Microsoft Teams automatically

Automatic installation makes it easier to install the Power BI app for Microsoft Teams, without needing to change Microsoft Teams app setup policies. This change speeds up the installation and removes admin hassles of configuring and maintaining infrastructure needed by an app setup policy.

When the app is installed, users receive notifications in Teams and can more easily discover and collaborate with colleagues. The Power BI app for Teams provides users with the ability to open all Fabric content.

Automatic installation happens for a user under the following conditions:

- The Power BI app for Microsoft Teams is set to **Allowed** in the Microsoft Teams admin portal.
- The Power BI tenant setting **Install Power BI app for Microsoft Teams automatically** is **Enabled**.
- The user has a Microsoft Teams license.
- The user opens the Power BI service in a web browser.

When the app is installed, users receive notifications in Teams and can more easily discover and collaborate with colleagues. The Power BI app for Teams provides users with the ability to open all Fabric content.

To learn more, see Add the Power BI app to Microsoft Teams.

Enable Power BI add-in for PowerPoint

The Power BI add-in for PowerPoint makes it possible for users to add live, interactive data from Power BI to a PowerPoint presentation. See About the Power BI add-in for PowerPoint for

more detail.

When this setting is on (default), entry points for opening a new PowerPoint presentation with the add-in already loaded are available in Power BI. When this setting is off, the entry points in Power BI are unavailable.

This integration requires that your organization's Microsoft Office admin has enabled support for add-ins.

If you turn this setting off, that doesn't prevent people from using the add-in starting from PowerPoint. To completely block adding live Power BI report pages to PowerPoint slides using the add-in, the add-in must be turned off in both Power BI and PowerPoint.

Allow DirectQuery connections to Power BI semantic models

When this setting is turned on (default), users can use DirectQuery to connect to Azure Analysis Services or Power BI datasets.

To learn more about DirectQuery, see Use DirectQuery in Power BI Desktop.

If you turn this switch off, it effectively stops users from publishing new composite models on Power BI semantic models to the service. Existing reports that leverage a composite model on a Power BI semantic model continue to work, and users are still able to create composite models using Desktop, but they can't publish to the service.

To learn more about composite models, see Use composite models in Power BI Desktop.

> **Note**
>
> Live connections to Power BI semantic models aren't affected by this switch, nor are live or DirectQuery connections to Azure Analysis Services. These continue to work regardless of whether the setting is on or off. In addition, any published reports that leverage a composite model on a Power BI semantic model continue to work even if the setting has been turned off after they were published.

Guest users can work with shared semantic models in their own tenants

When this setting is turned on, Microsoft Entra B2B guest users of semantic models shared with them by users in your organization can access and build on those semantic models in their own tenant.

This setting is off by default for customers. If this setting is disabled, a guest user can still access the semantic model in the provider tenant but not in their own tenant.

Allow specific users to turn on external data sharing

As a Fabric admin, you can specify which users or user groups in your organization can share semantic models externally with guests from a different tenant through the in-place mechanism. Authorized guest users can then discover, connect to, and work with these shared semantic models in their own tenants.

Disabling this setting prevents any user from sharing semantic models externally by blocking the ability of users to turn on external sharing for semantic models they own or manage.

> **Note**
>
> This setting relates to sharing Power BI semantic models via Entra B2B capabilities. It is unrelated to the **External data sharing (preview)** and **Users can accept external data shares (preview)** tenant settings, which control the external data sharing feature. The external data sharing feature enables sharing data from an organization's OneLake storage locations to external Fabric tenants, and uses secure Fabric-to-Fabric communication channels rather than Entra B2B.

Users with read or write permission can download data from notebooks

With this setting on, users with read or write permission can download or export data from tables in notebook outputs.

For more information, see Download data from notebook outputs.

Related content

- About tenant settings

- Tenant settings index

Last updated on 04/09/2026

Discovery tenant settings

These settings are configured in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

Discoverability is a feature that semantic model owners can use to make their endorsed content discoverable by users who don't yet have access to it.

Make promoted content discoverable

Allow users in this organization who can promote content to make content they promote discoverable by users who don't have access to it. You can also specify users and/or groups to exclude from the permitted groups.

To learn more, see Semantic model discoverability.

Make certified content discoverable

Allow users in the organization who can certify content to make content they certify discoverable by users who don't have access to it.

Discover content

Allow users to find and request access to content they don't have access to if it was made discoverable by its owners.

To learn more, see Find recommended items.

Related content

- About tenant settings

Last updated on 09/29/2026

App tenant settings

These settings are configured in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

Create template organizational apps

Users in the organization can create template apps that use semantic models built on one data source in Power BI Desktop.

To learn more, see Create a template app in Power BI.

Push apps to end users

Admins can allow report creators to share apps directly with end users, without requiring installation from AppSource. In **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**, the setting is **Push apps to end users**.

To learn more, see Automatically install apps for end users.

Publish apps to the entire organization

Admins use this setting to decide which users can publish apps to the entire organization, rather than specific groups. The following image shows the **Entire organization** option when creating an app.

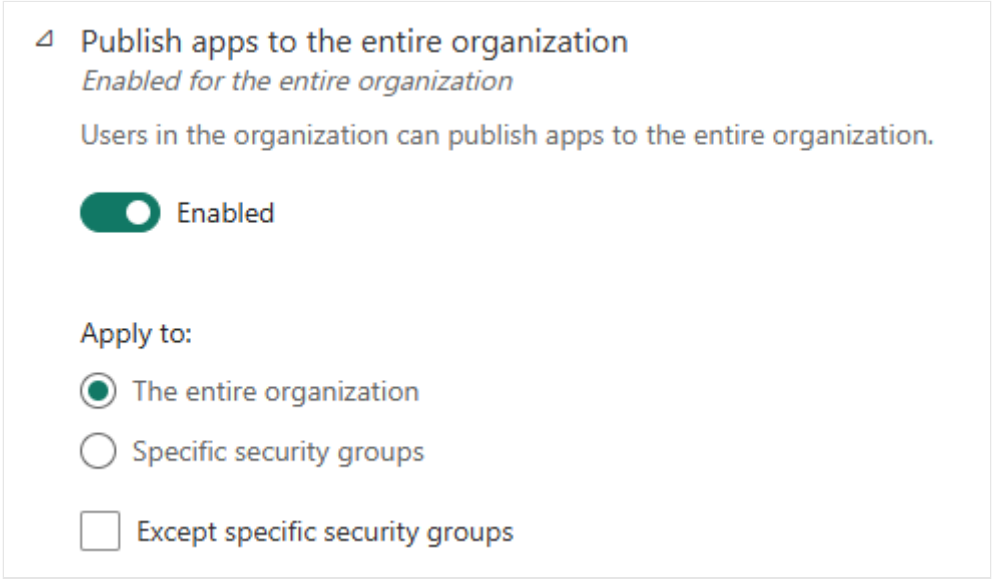

To learn more, see Publish an app in Power BI.

Related content

- About tenant settings

Last updated on 09/29/2026

Configure Fabric data agent tenant settings

To use a data agent in Microsoft Fabric, configure the required tenant settings. This guide walks you through the necessary configurations for a seamless setup.

> **Important**
>
> Users may configure Fabric data agents to be consumed from other services such as Microsoft Foundry, Microsoft Copilot Studio, M365 Copilot or as an MCP server ("non-Fabric services"). When users connect to these non-Fabric services, responses returned by Fabric data agents may be sent outside of Fabric's compliance boundary or geographic region, and processed and/or stored according to the non-Fabric service(s) applicable terms and data handling policies.

Access tenant settings

To configure the required settings, you need administrative privileges to access tenant settings in Microsoft Fabric.

1. **Sign in to Microsoft Fabric** with an admin account.
2. Go to **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**.

When you're in **Tenant Settings**, enable the necessary configurations.

> **Note**
>
> The tenant settings might take up to one hour to take effect after you enable them.

Enable Copilot and Azure OpenAI tenant switch

For a Fabric data agent to function properly, enable the **Copilot and Azure OpenAI Service** tenant settings. These settings control user access and data processing policies.

Required settings

- **Users can use Copilot and other features powered by Azure OpenAI**:

- Enable this setting to allow users to access Copilot-powered features, including Fabric data agent. You can manage this setting at both the tenant and the capacity levels. For more information, see Overview of Copilot in Fabric.
- To enable this setting, check the option in **Tenant Settings** as shown in the following screenshot:

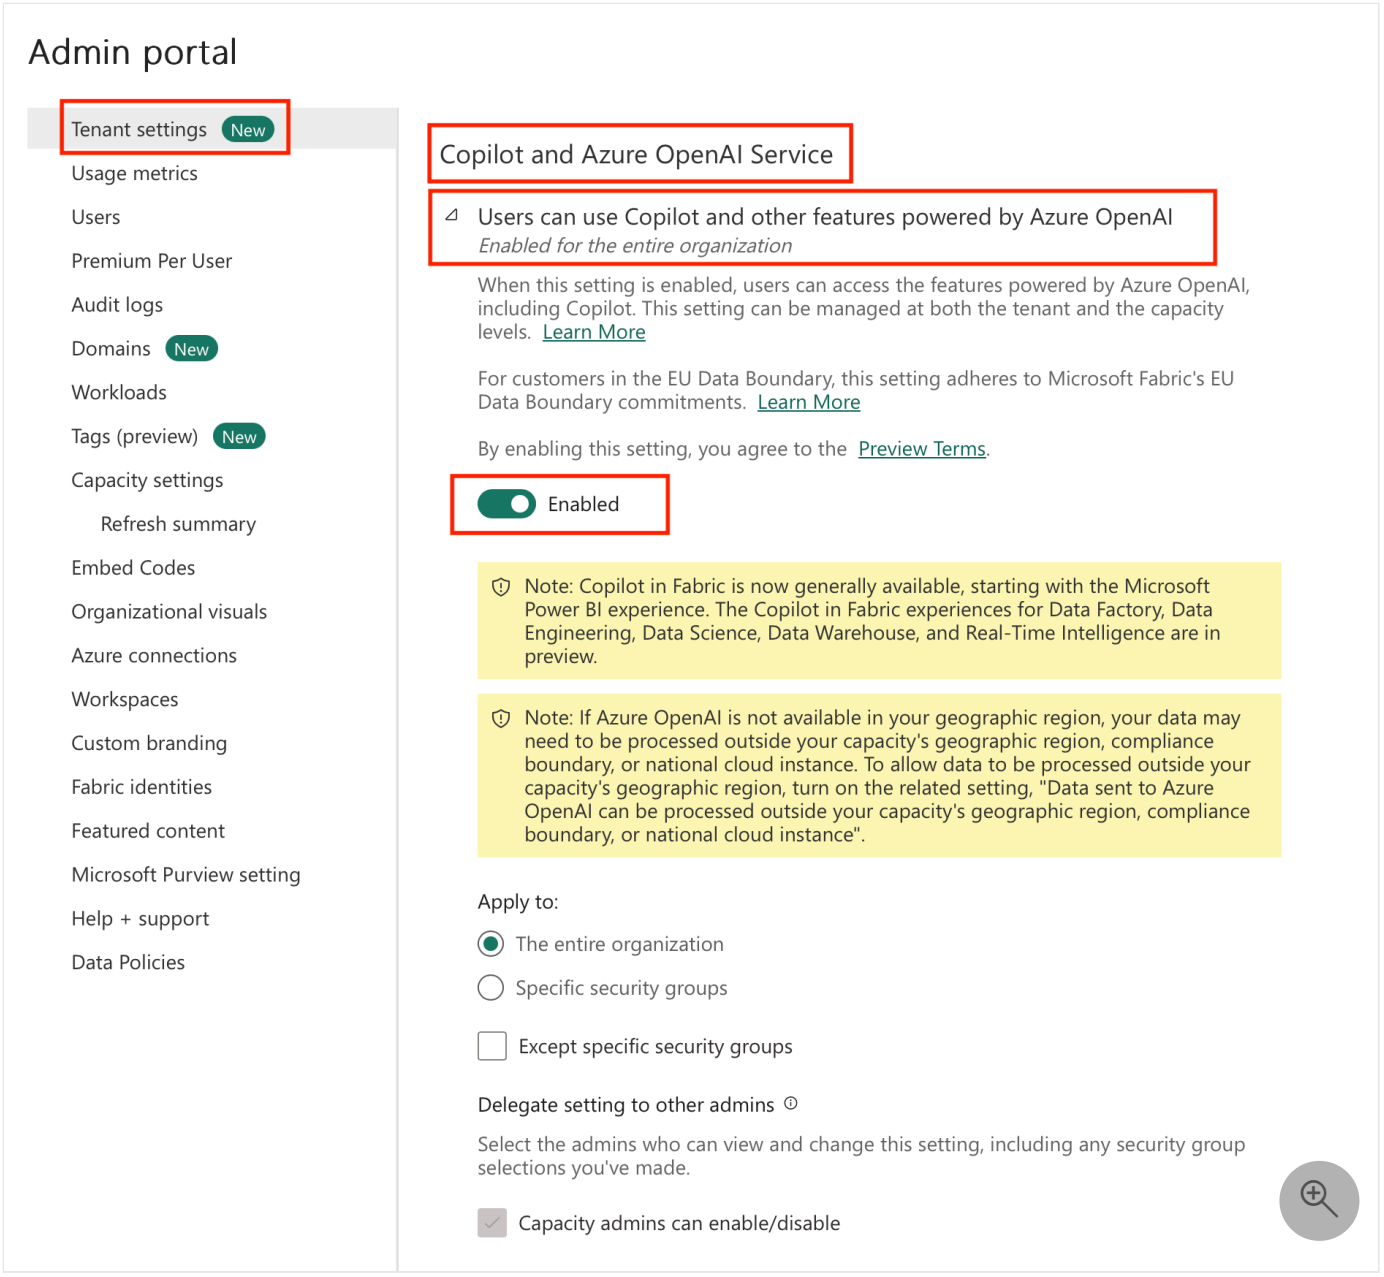

- **Capacities can be designated as Fabric Copilot capacities**:
  - Enable this setting to allow capacity administrators to designate capacities as Fabric Copilot capacities for Copilot usage, including Fabric data agent.
  - For more information, see Capacities can be designated as Fabric Copilot capacities.
- **Data sent to Azure OpenAI can be processed outside your capacity's geographic region, compliance boundary, or national cloud instance**

- Required for customers using Fabric data agent whose capacity's geographic region is outside of the EU data boundary and the US.
- To enable this setting, check the option in **Tenant Settings** as shown in the following screenshot:

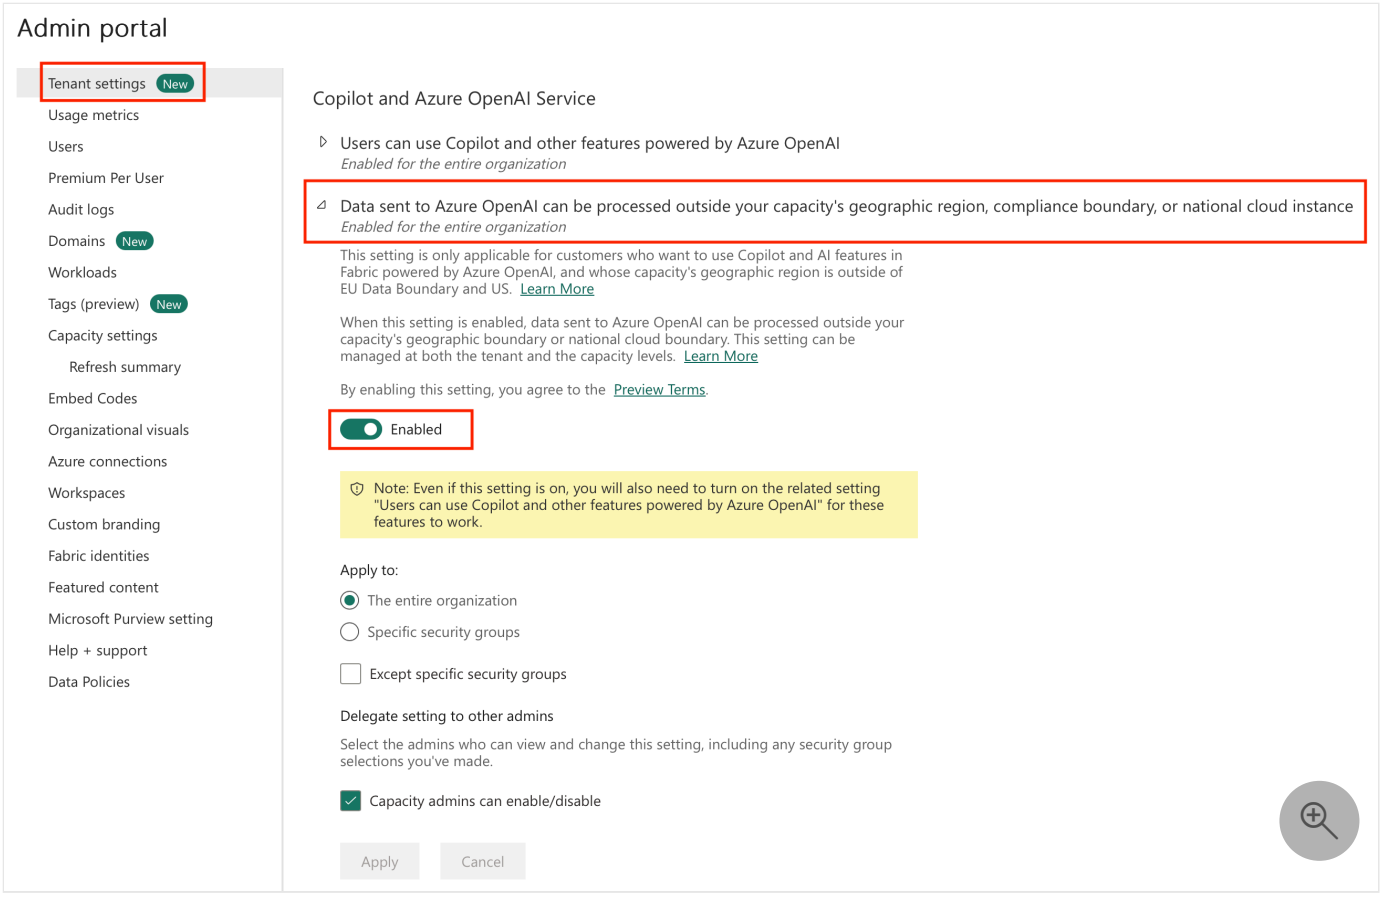

- **Data sent to Azure OpenAI can be stored outside your capacity's geographic region, compliance boundary, or national cloud instance**
  - Required for customers using Copilot in Notebooks and the Data agent Feature in Fabric powered by Azure OpenAI whose capacity's geographic region is outside of the EU data boundary and the US.
  - To enable this setting, check the option in **Tenant Settings** as shown in the following screenshot:

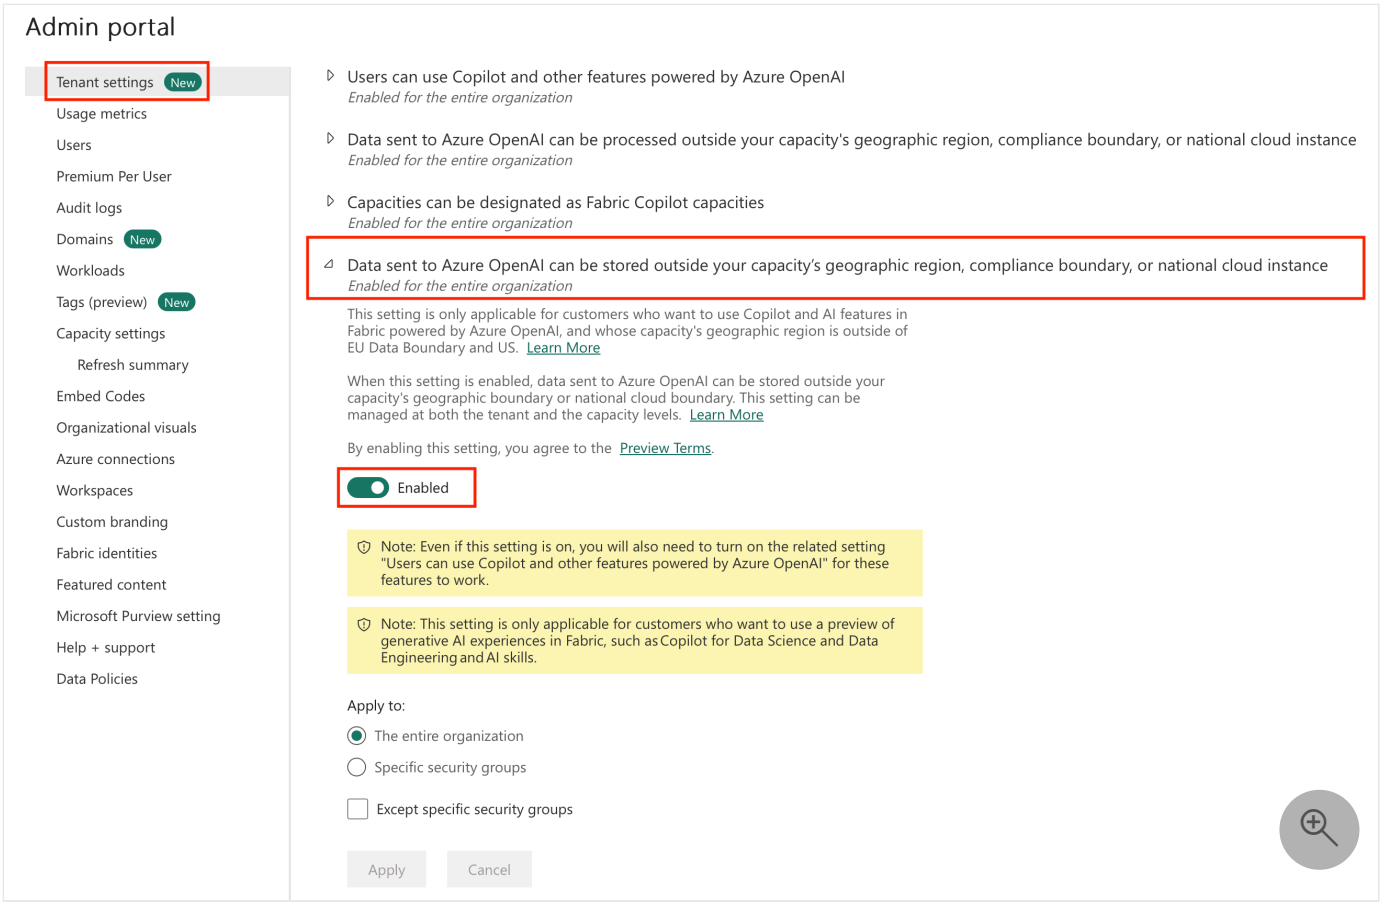

- This setting is only applicable for customers who want to use Copilot in Notebooks and Fabric data agents powered by Azure OpenAI whose capacity's geographic region is outside of the EU data boundary and the US.
- In order to use fully conversational agentic AI experiences, the agent needs to store conversation history across user sessions. This ensures that the AI agent keeps context about what a user asked in previous sessions. Conversation history is stored for as long as the user allows, up to 28 days if not manually removed.
- Users can delete their conversation history at any time by clearing the chat.

Related content

- Data agent concept
- About tenant settings

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 09/29/2026

Copilot and Agent tenant settings

Copilot in Fabric and Agent settings are controlled by the **Copilot and AI** tenant settings group. There are multiple settings governing user access and data processing policies, and some of them are enabled by default whereas others require the Fabric administrator to enable them.

For information about how to get to the Fabric tenant settings, see About tenant settings - How to get to the tenant settings.

Settings enabled by default

- Users can use Copilot and other features powered by Azure OpenAI
- Users can use Copilot, AI Agents and other AI experiences powered by OpenAI as a Microsoft Subprocessor
- Capacities can be designated as Copilot in Fabric capacities

Settings disabled by default

- Data sent to Azure OpenAI can be processed outside your capacity's geographic region, compliance boundary, or national cloud instance
- Data sent to OpenAI as a Microsoft Subprocessor can be processed outside your capacity's geographic region, compliance boundary, or national cloud instance
- Data sent to Azure OpenAI can be stored outside your capacity's geographic region, compliance boundary, or national cloud instance
- Conversation history stored outside your capacity's geographic region, compliance boundary, or national cloud instance
- Users can access a standalone, cross-item Copilot in Power BI experience (preview)
- Only show approved items in the standalone Copilot in Power BI experience (preview)

> **Important**
>
> Users may configure Fabric data agents to be consumed from other services such as Microsoft Foundry, Microsoft Copilot Studio, Microsoft 365 Copilot or as an MCP server ("non-Fabric services"). When users connect to these non-Fabric services, responses returned by Fabric data agents may be sent outside of Fabric's compliance boundary or geographic region, and processed and/or stored according to the non-Fabric service(s) applicable terms and data handling policies.

Users can use Copilot and other features powered by Azure OpenAI

When this setting is enabled, users can access the features powered by Azure OpenAI, including Copilot and Agents, as shown in the following screenshot:

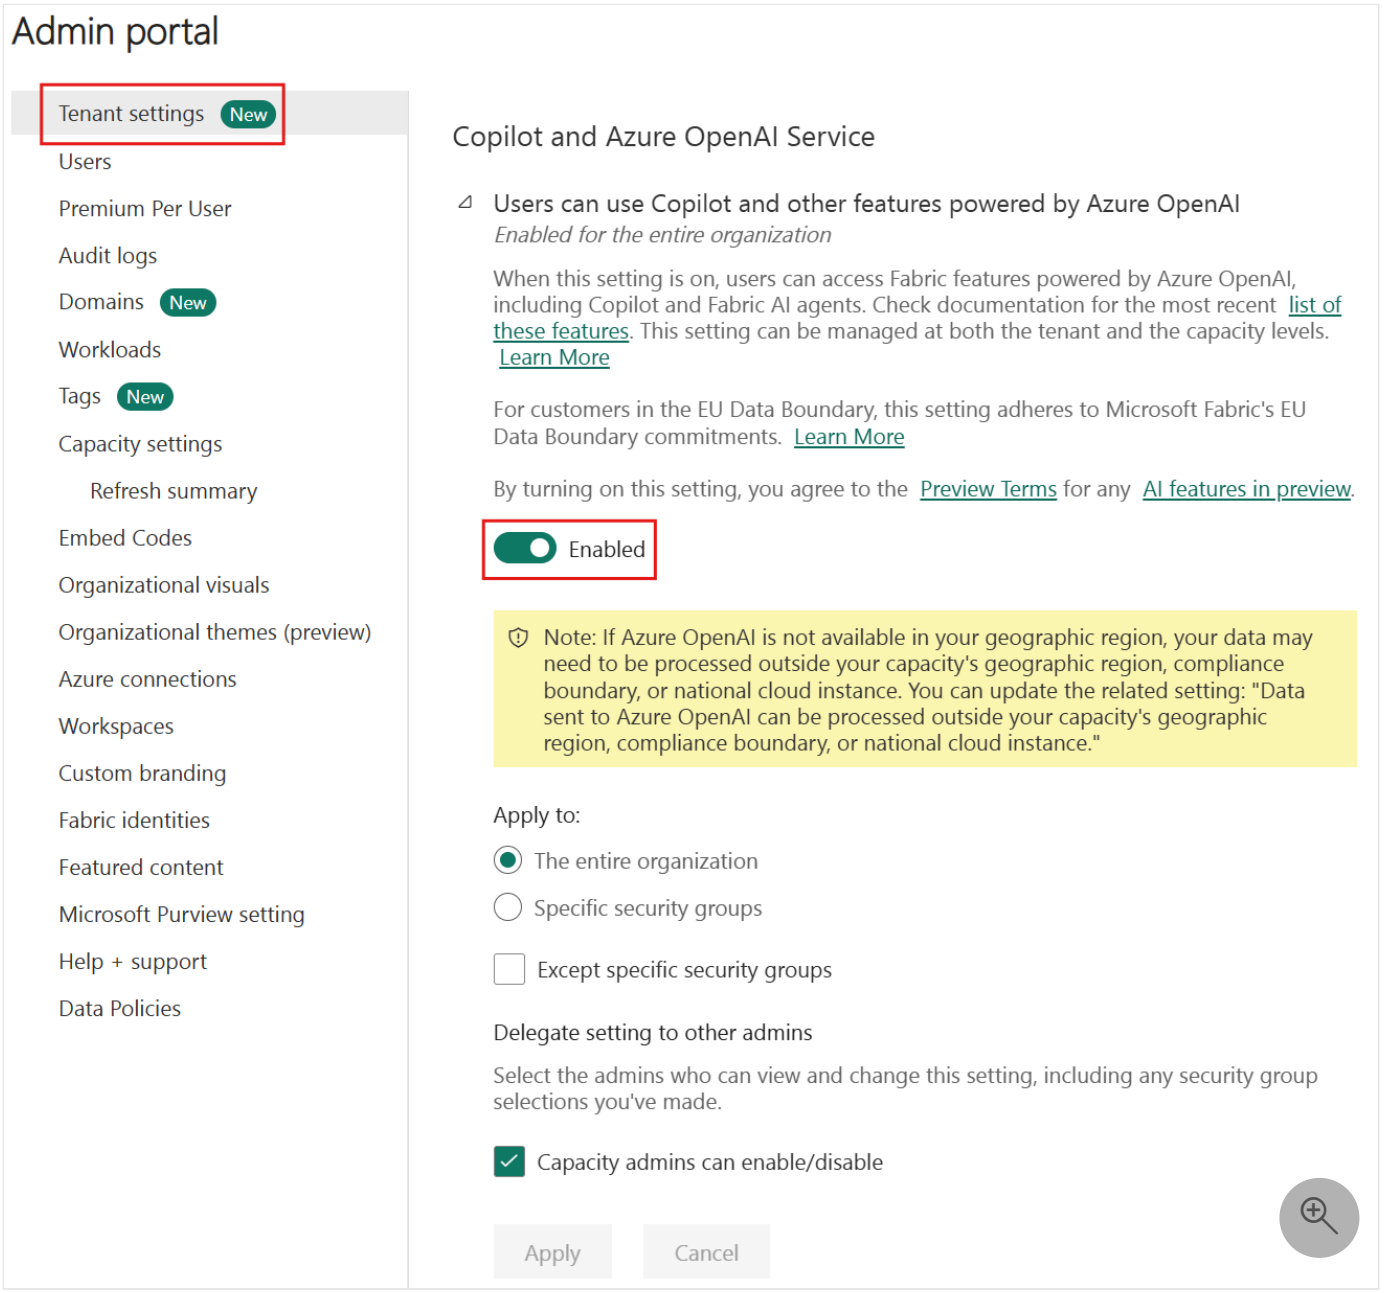

This setting can be managed at both the tenant and the capacity levels. When this setting is enabled, the service may execute background jobs at no charge to the tenant capacity to support end user experiences. For more information, see Overview of Copilot in Fabric.

Default: Enabled

Users can use Copilot, AI Agents and other AI experiences powered by OpenAI as a

Microsoft Subprocessor

When this setting is enabled, users can access Copilot, AI Agents, and other AI experiences powered by OpenAI-operated models.

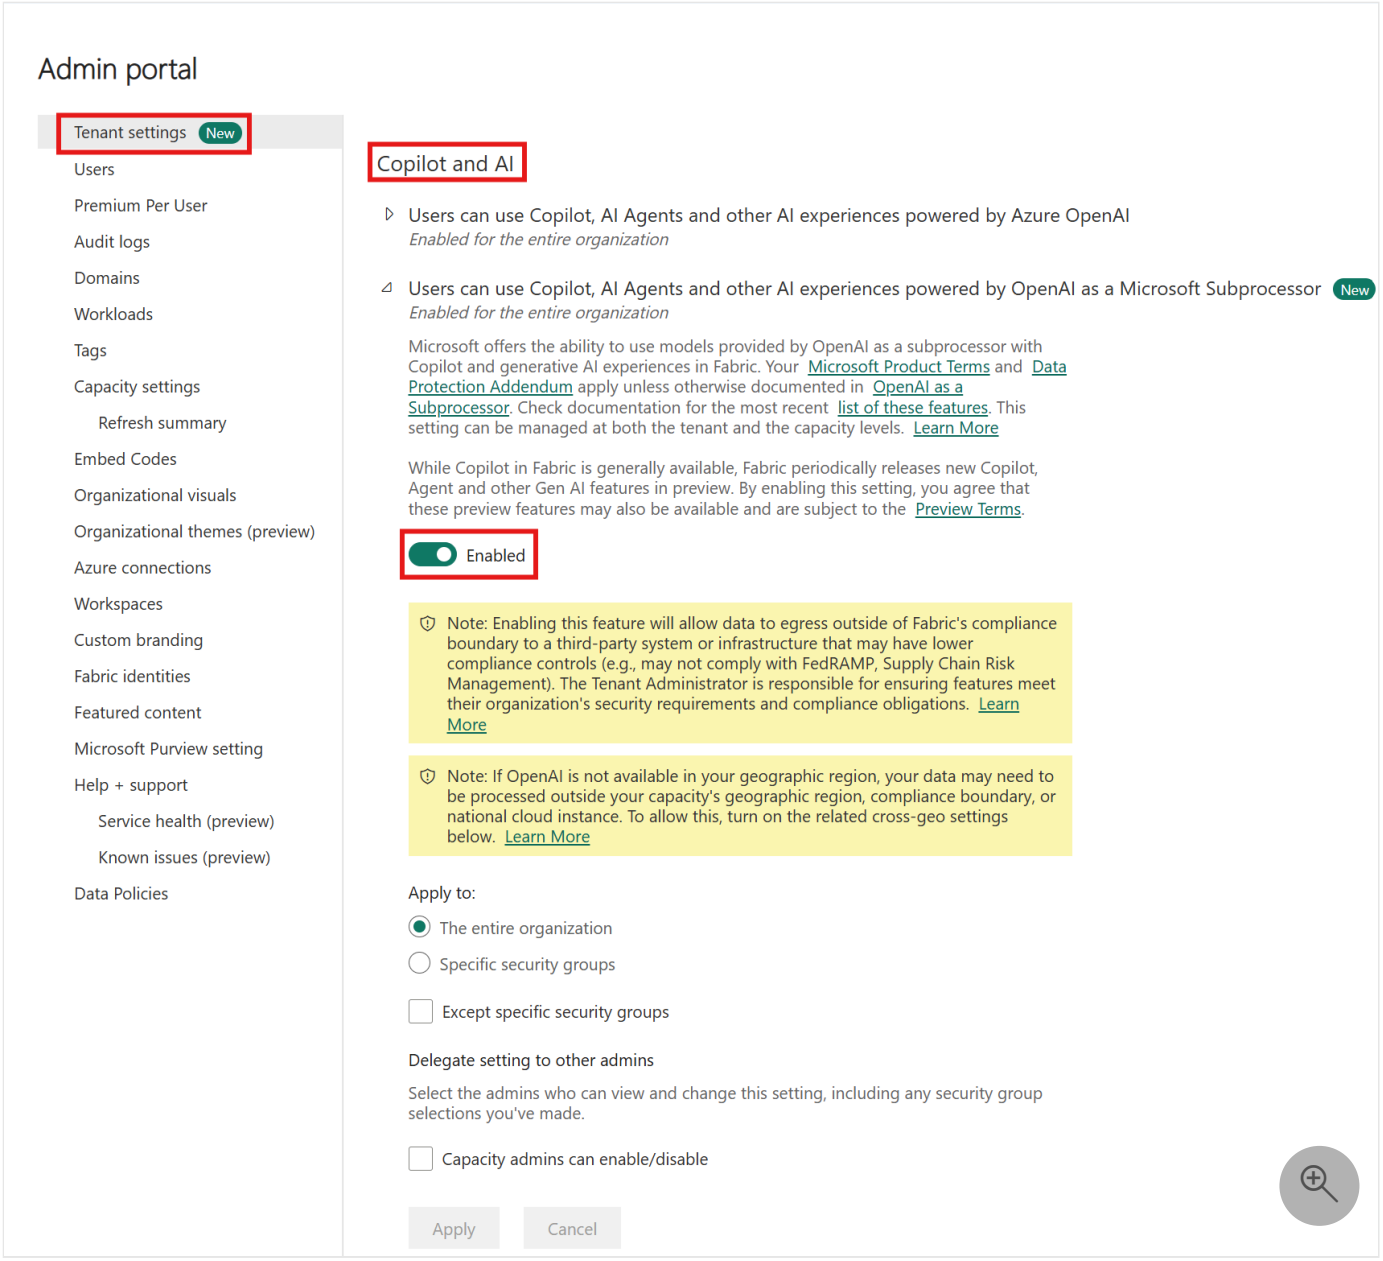

This setting can be managed at both the tenant and the capacity levels. For more information about data processing commitments, availability limitations, and compliance exclusions, see OpenAI as a subprocessor in Microsoft Fabric.

Default: Disabled

Data sent to Azure OpenAI can be processed outside your capacity's geographic region, compliance boundary, or national cloud instance

This setting is only applicable for customers who want to use Copilot and AI features in Fabric powered by Azure OpenAI, and whose capacity's geographic region is outside of the EU data boundary and the US. When this setting is enabled, service background jobs may execute across geographic boundaries at no charge to the tenant capacity to support end user experiences.

The following screenshot shows how to configure this setting:

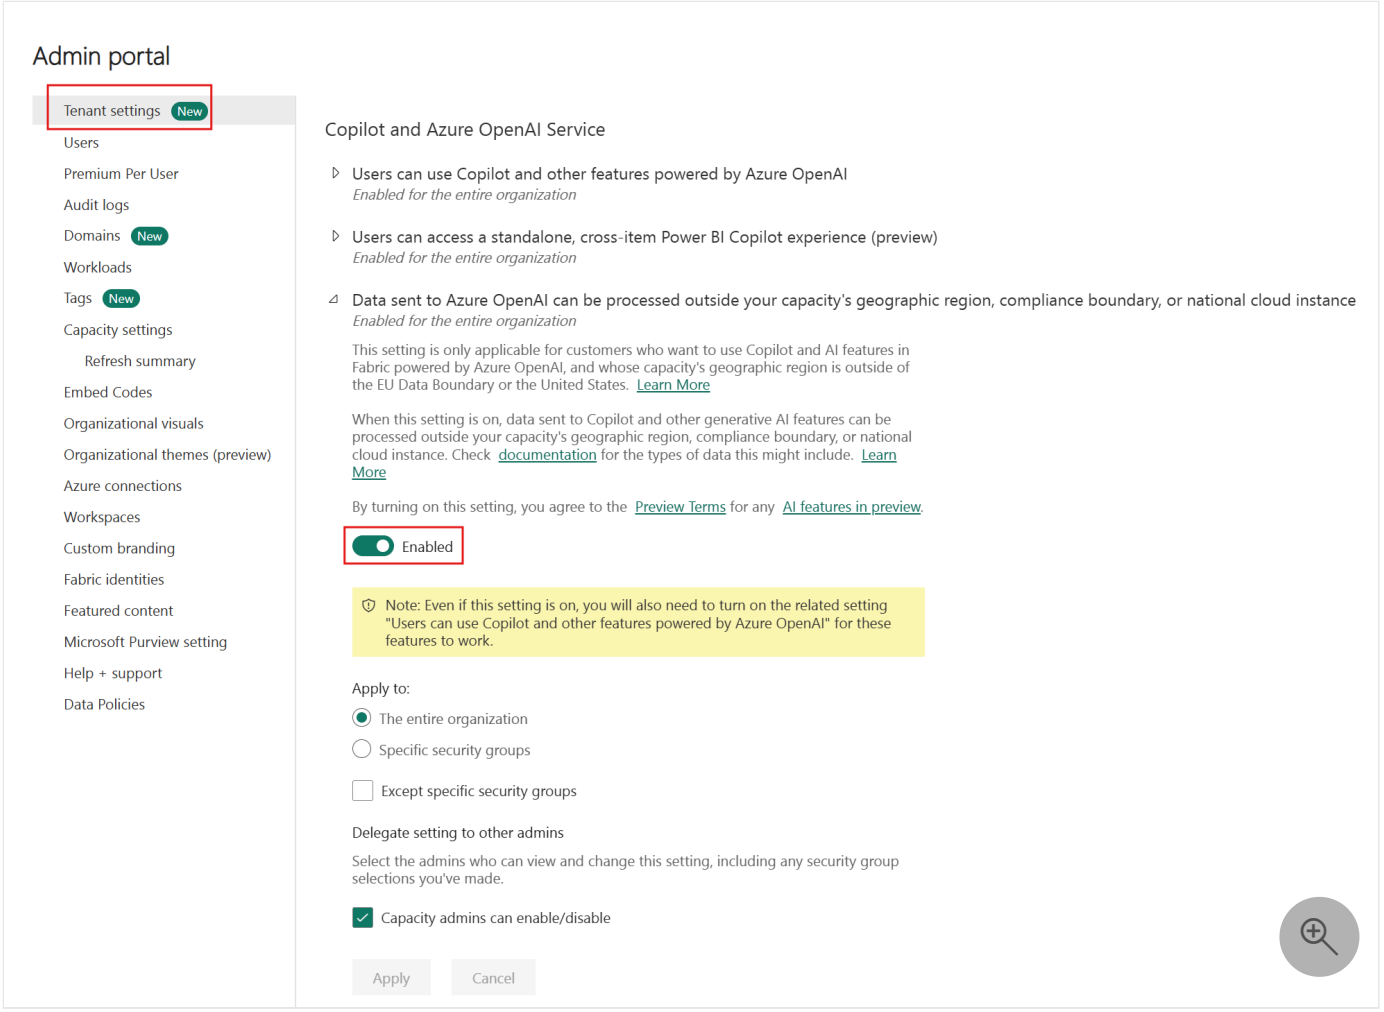

For more information, see Available regions.

Default: Disabled

Data sent to OpenAI as a Microsoft Subprocessor can be processed outside your capacity's geographic region, compliance boundary, or national cloud instance

This setting is only applicable for customers who want to use Copilot and AI features in Fabric powered by OpenAI-operated models, and whose capacity's geographic region is outside of the EU data boundary and the US. When this setting is enabled, service background jobs may

execute across geographic boundaries at no charge to the tenant capacity to support end user experiences.

The following screenshot shows how to configure this setting:

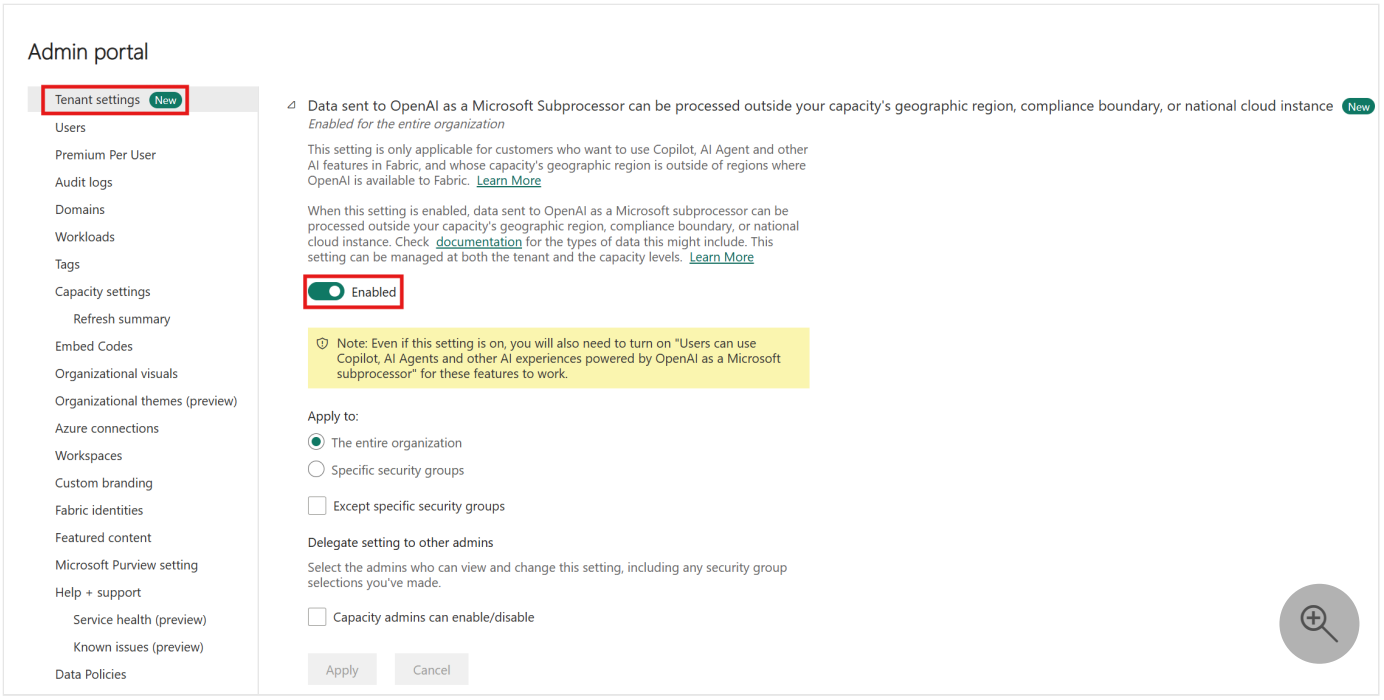

For more information, see OpenAI as a subprocessor in Microsoft Fabric.

Default: Disabled

Data sent to Azure OpenAI can be stored outside your capacity's geographic region, compliance boundary, or national cloud instance

This setting is only applicable for customers who want to use Copilot in notebooks and data agents in Fabric powered by Azure OpenAI, and whose capacity's geographic region is outside of the EU data boundary and the US. The following screenshot shows how to configure this setting:

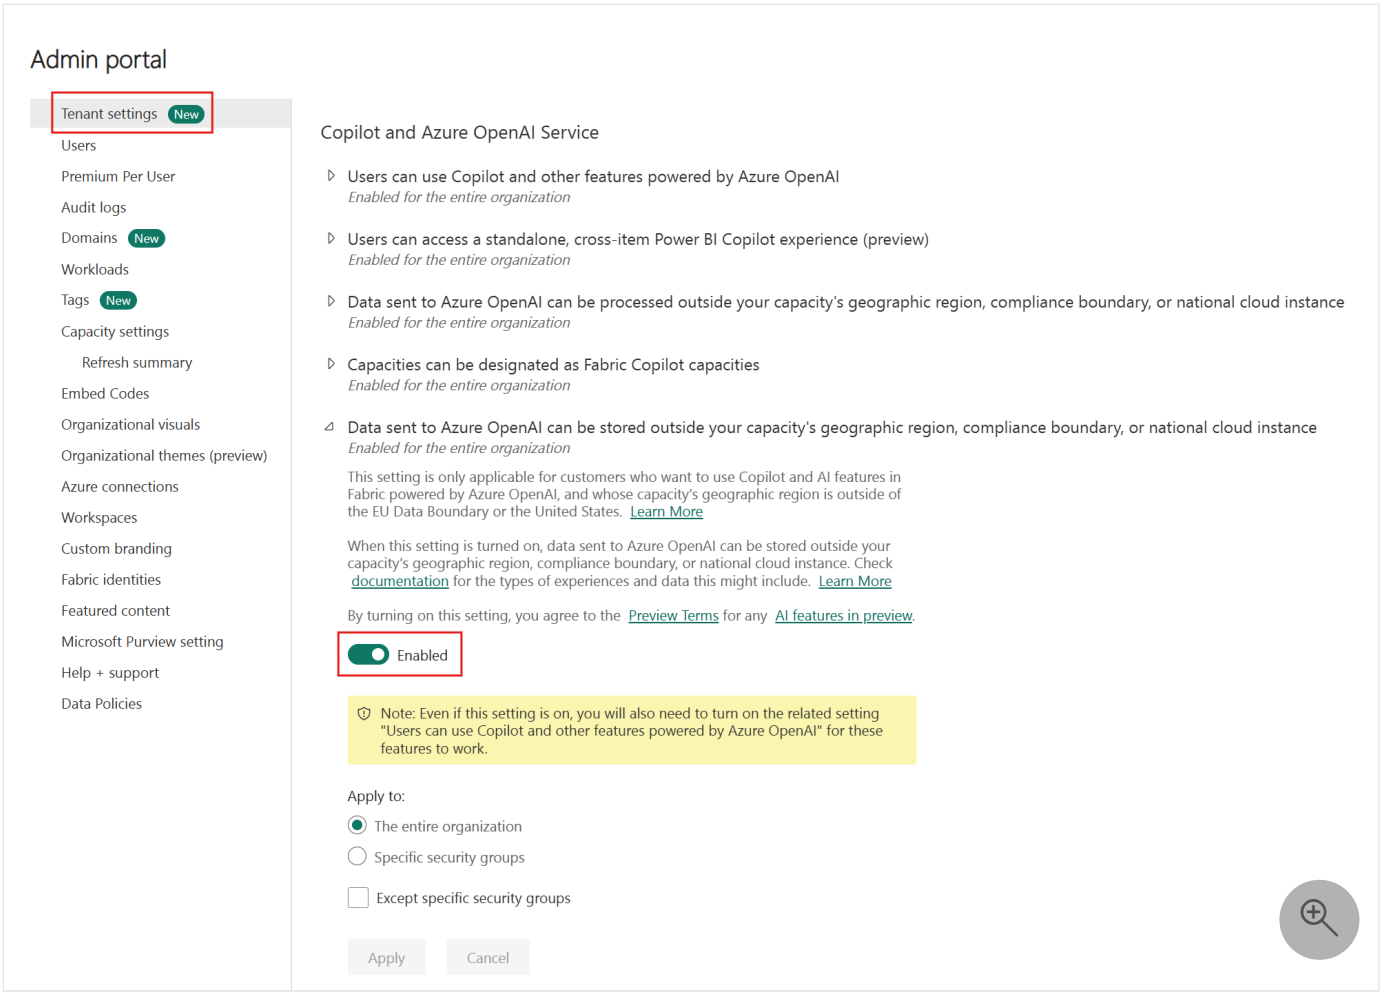

For more information, see Available regions.

Default: Disabled

Conversation history stored outside your capacity's geographic region, compliance boundary, or national cloud instance

This setting is only applicable for customers who want to use Copilot in notebooks and Fabric data agents powered by Azure OpenAI, and whose capacity's geographic region is outside of the EU data boundary and the US.

To use fully conversational agentic AI experiences, the agent needs to store conversation history across user sessions. This ensures that the AI agent keeps context about what a user asked in previous sessions. Experiences such as Copilot in notebooks and Fabric data agents store conversation history across the user's sessions. This history is stored inside the Azure security boundary, in the same region and in the same Azure OpenAI resources that process all your Fabric AI requests. The difference in this case is that the conversation history is stored for as long

as the user allows. For experiences that don't store conversation history across sessions, no data is stored. Prompts are only processed by Azure OpenAI resources that Fabric uses.

Users can delete their conversation history at any time by clearing the chat. This option exists both for Copilot in notebooks and data agents. If the conversation history isn't manually removed, it's stored for 28 days.

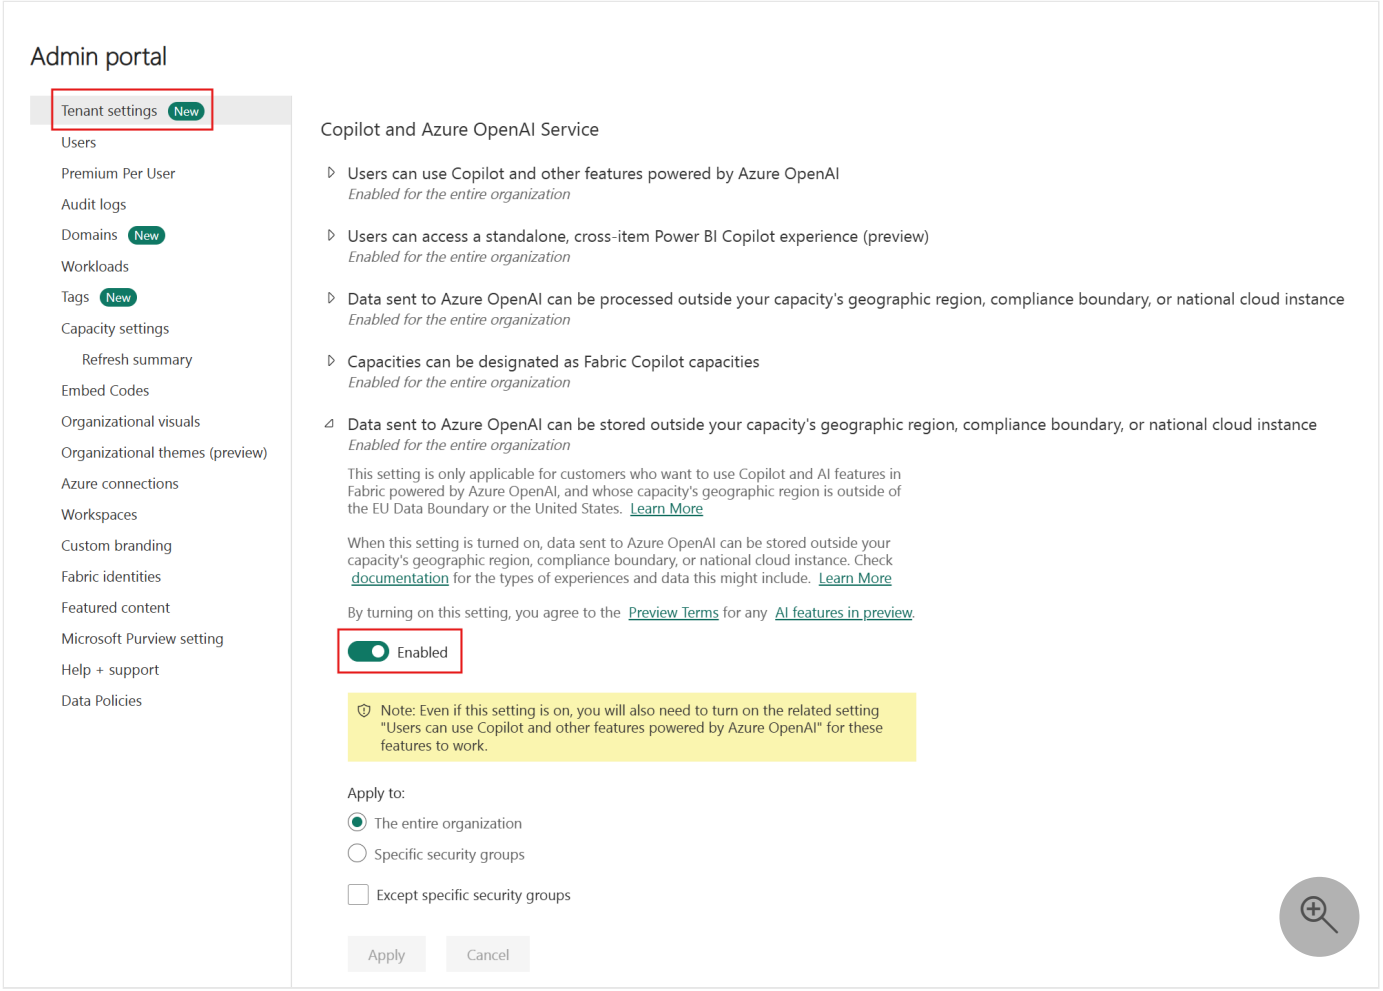

For more information, see Available regions.

Default: Disabled

Capacities can be designated as Copilot in Fabric capacities

Copilot capacities enable users' usage and billing to be consolidated under a single capacity. Fabric administrators can assign specific groups or the entire organization to manage capacities as Copilot in Fabric capacities. Capacity administrators must designate user access to each Copilot capacity and can view item names linked to users' Copilot activity in the Microsoft Fabric Capacity Metrics app.

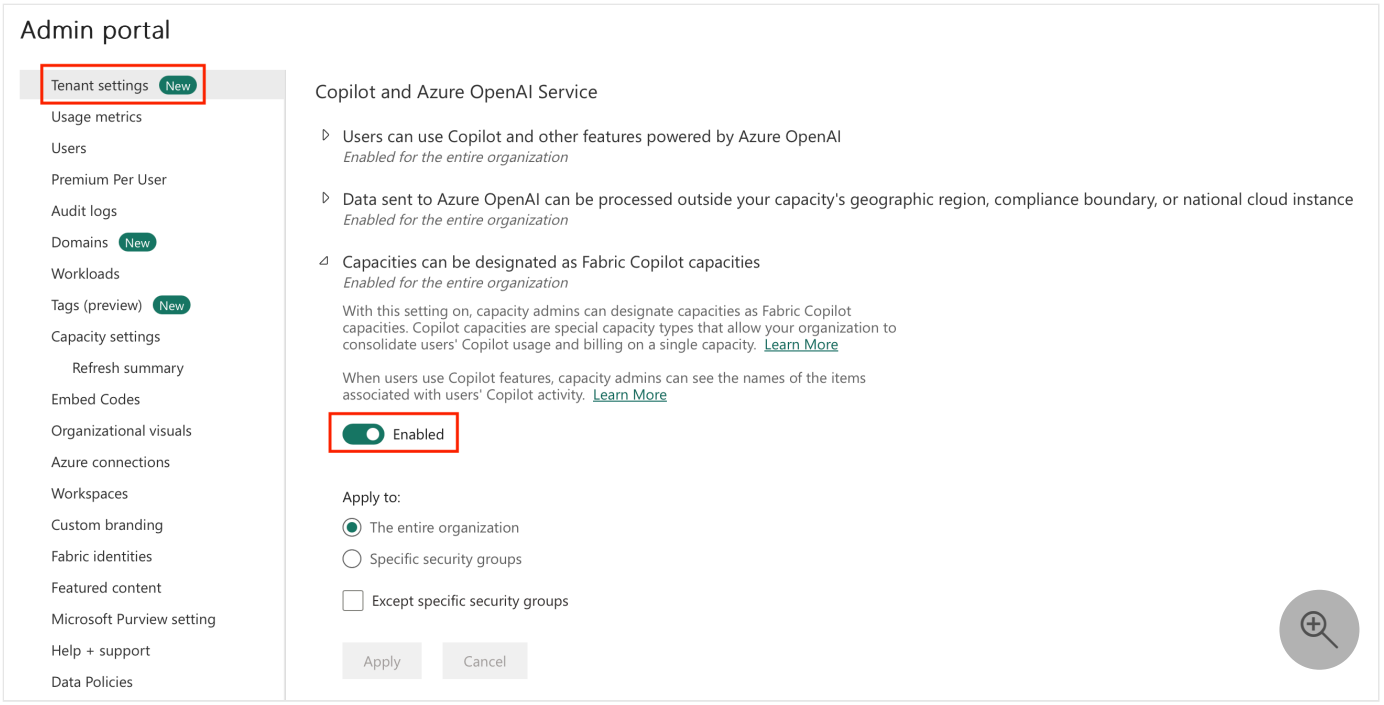

Default: Enabled

Users can access a standalone, cross-item Copilot in Power BI experience (preview)

You can enable Copilot as a standalone experience for Fabric. Enabling this setting allows users to access the standalone Copilot experience from the left navigation. The Azure OpenAI setting must be enabled at the tenant level to use the standalone experience. This setting also affects the Power BI agent. Turn on this setting to enable the Power BI agent in Microsoft 365. To learn more, see standalone Copilot experience.

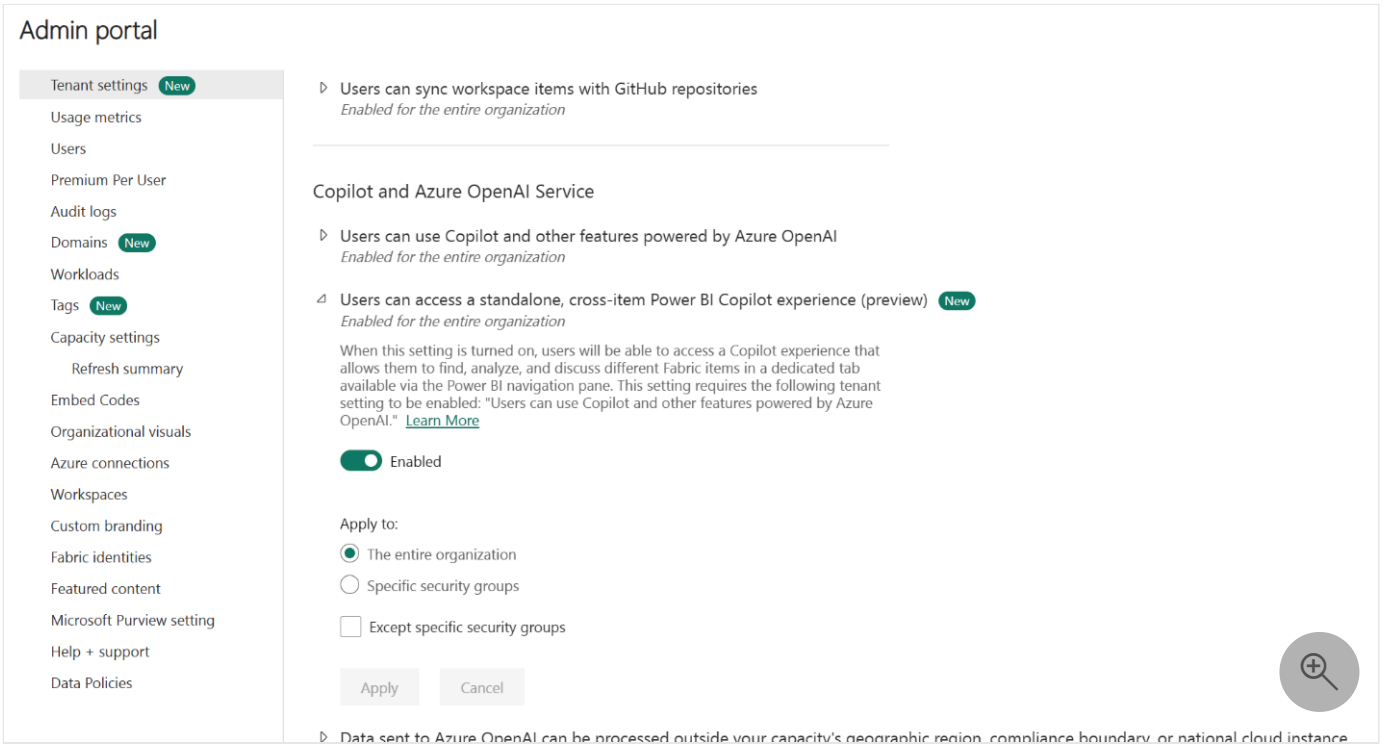

Default: Disabled

Only show approved items in the standalone Copilot in Power BI experience (preview)

Tenant admins can default Copilot search to be limited to items that are marked as approved for Copilot. This setting is delegated to workspace admins by default, allowing workspace admins to make broader content findable by Copilot search when appropriate.

> **Note**
>
> This setting was previously called *Only show AI-prepped items in the standalone Copilot in Power BI experience (preview)*, but has been updated to reflect the new name of the **approved for Copilot setting**.

This setting is applicable in the standalone Copilot in Power BI as well as the Power BI agent. If this setting is turned on for standalone Copilot, it's also mandatory for the Power BI agent. To learn more, see standalone Copilot experience.

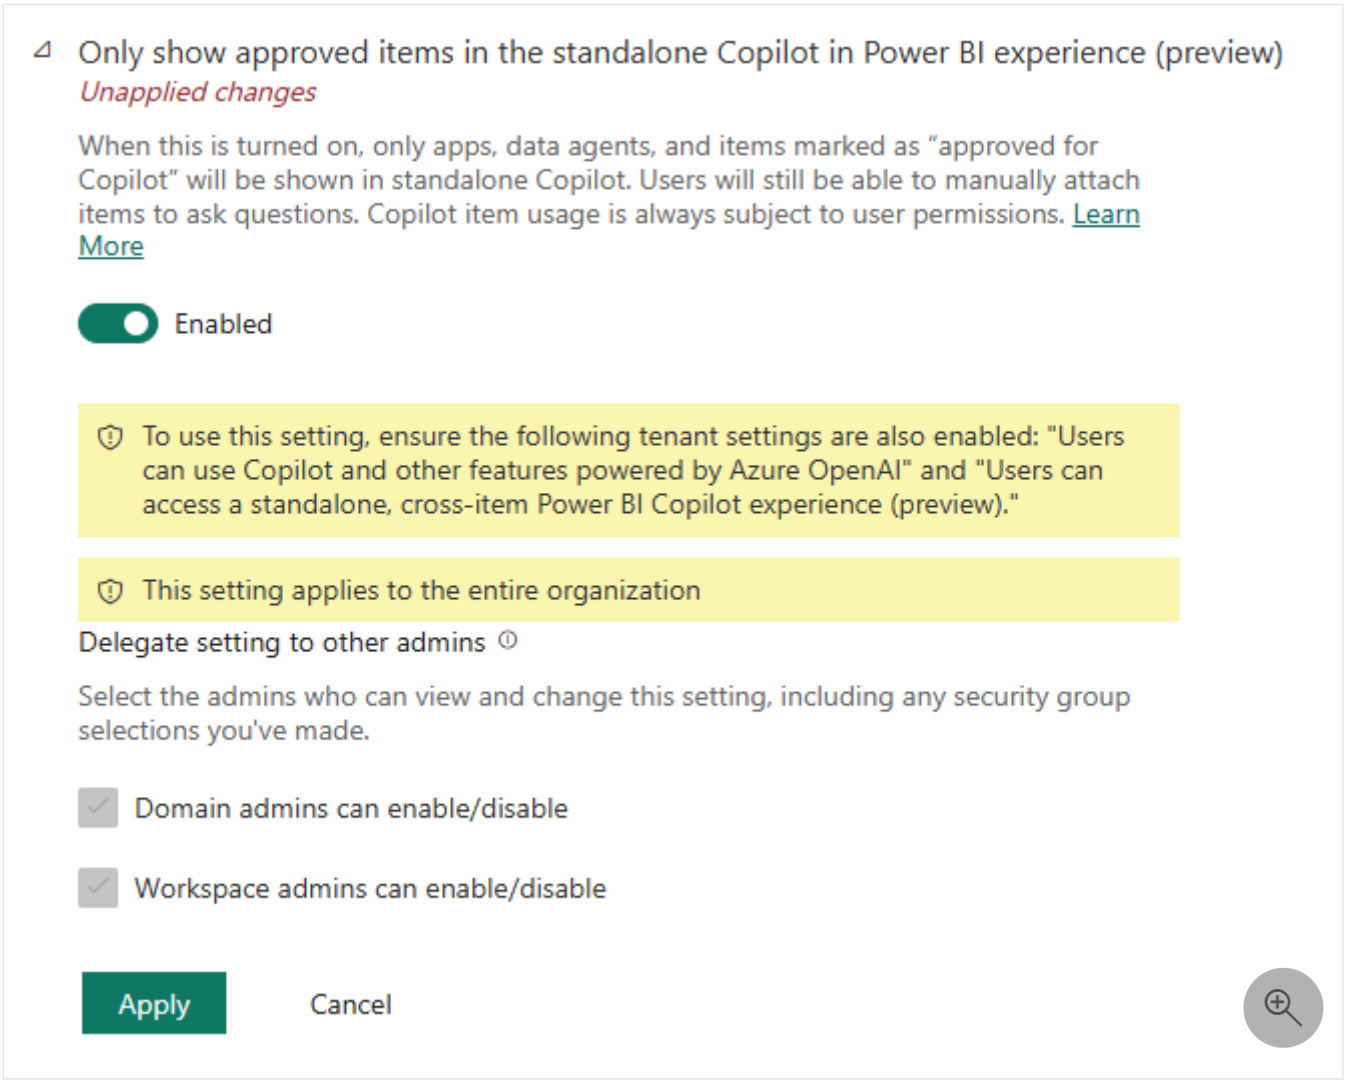

Default: Disabled

Related content

- Copilot in Fabric and Power BI overview
- OpenAI as a subprocessor in Microsoft Fabric
- About tenant settings

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 08/31/2026

Tenant settings for additional workloads

These settings are configured in the tenant settings section of the Admin portal. For information about how to get to and use tenant settings, see About tenant settings.

Additional workloads let organizations extend Fabric with partner and custom workloads. Tenant admins can control who can add, develop, and use those workloads.

Workspace admins can add and remove additional workloads (preview)

Turn on this setting to allow workspace admins to add and remove workloads in their workspaces.

If you turn this setting off, existing workloads remain in place and items created with those workloads continue to work normally.

When users interact with a workload, their data and access tokens, including name and email, are sent to the publisher. Sensitivity labels and protection settings, including encryption, aren't applied to items created with workloads.

For more information, see Workload Development Kit.

Capacity admins and contributors can add and remove additional workloads

Turn on this setting to allow capacity admins and users with Contributor permission in capacity settings to add and remove additional workloads in capacities.

If you turn this setting off, existing workloads remain in place and items created with those workloads continue to work normally.

When users interact with a workload, their data and access tokens, including name and email, are sent to the publisher. Sensitivity labels and protection settings, including encryption, aren't applied to items created with workloads.

For more information, see Enable the development tenant setting.

Workspace admins can develop partner workloads

Turn on this setting to allow workspace admins to develop partner workloads by using a local machine development environment.

If you turn this setting off, developers can no longer upload partner workloads to that workspace.

For more information, see Enable the development tenant setting.

Users can see and work with additional workloads not validated by Microsoft

Turn on this setting to allow users to see and work with additional workloads that Microsoft hasn't validated.

Before you enable this setting, make sure that you trust the publishers of any workloads you add and that those workloads meet your organization's policies.

For more information, see Partner workload publishing requirements.

Related content

- About tenant settings
- Workload Development Kit overview

Last updated on 04/09/2026

Integration tenant settings

These settings are configured in the tenant settings section of the Admin portal. For information about how to get to and use tenant settings, see About tenant settings.

Allow XMLA endpoints and Analyze in Excel with on-premises semantic models

When enabled, users in the organization can use Excel to view and interact with on-premises Power BI semantic models. This also allows connections to XMLA endpoints.

To learn more, see Create Excel workbooks with refreshable Power BI data.

Semantic Model Execute Queries REST API

When enabled, users in the organization can query semantic models by using Data Analysis Expressions (DAX) through Power BI REST APIs.

To learn more, see Datasets - Execute Queries.

Users can use the Power BI Model Context Protocol server endpoint (preview)

When enabled, the Power BI Model Context Protocol (MCP) server endpoint allows users to connect MCP clients to Power BI and use MCP tools to interact with Power BI artifacts they have permission to access. GitHub Copilot in Visual Studio operates by default. Other MCP clients should be configured with service principals.

To allow service principals to connect to the MCP server, also turn on **Service principals can call Fabric public APIs**.

For more information, see Power BI Model Context Protocol server endpoint.

Use ArcGIS Maps for Power BI

When enabled, users in the organization can use the ArcGIS Maps for Power BI visualization provided by Esri.

To learn more, see Create ArcGIS maps in Power BI.

Use global search for Power BI

When enabled, users in the organization can use external search features that rely on Azure Search.

To learn more, see Navigation for Power BI business users: global search.

Users can use the Azure Maps visual

When enabled, users in the organization can create and view reports that use the Azure Maps visual.

- Microsoft temporarily stores and processes your data for essential services like translating locations into latitudes and longitudes.
- If Azure Maps services are unavailable in your region, your data can be processed outside your tenant's geographic region, compliance boundary, or national cloud instance.

To learn more, see Get started with Azure Maps Power BI visual.

Data sent to Azure Maps can be processed outside your tenant's geographic region, compliance boundary, or national cloud instance

Azure Maps services aren't available in all regions and geographies. When this setting is on, data sent to Azure Maps can be processed in a region where the service is available, which can be outside your tenant's geographic region, compliance boundary, or national cloud instance.

For more information, see Azure Maps geographic scope.

Data sent to Azure Maps can be processed by Microsoft Online Services Subprocessors

Some Azure Maps visual services, including the selection tool and processing location names in some regions, require mapping capabilities provided in part by Microsoft Online Services subprocessors. Microsoft shares only the data needed to provide those services.

For more information, see Azure Maps subprocessors.

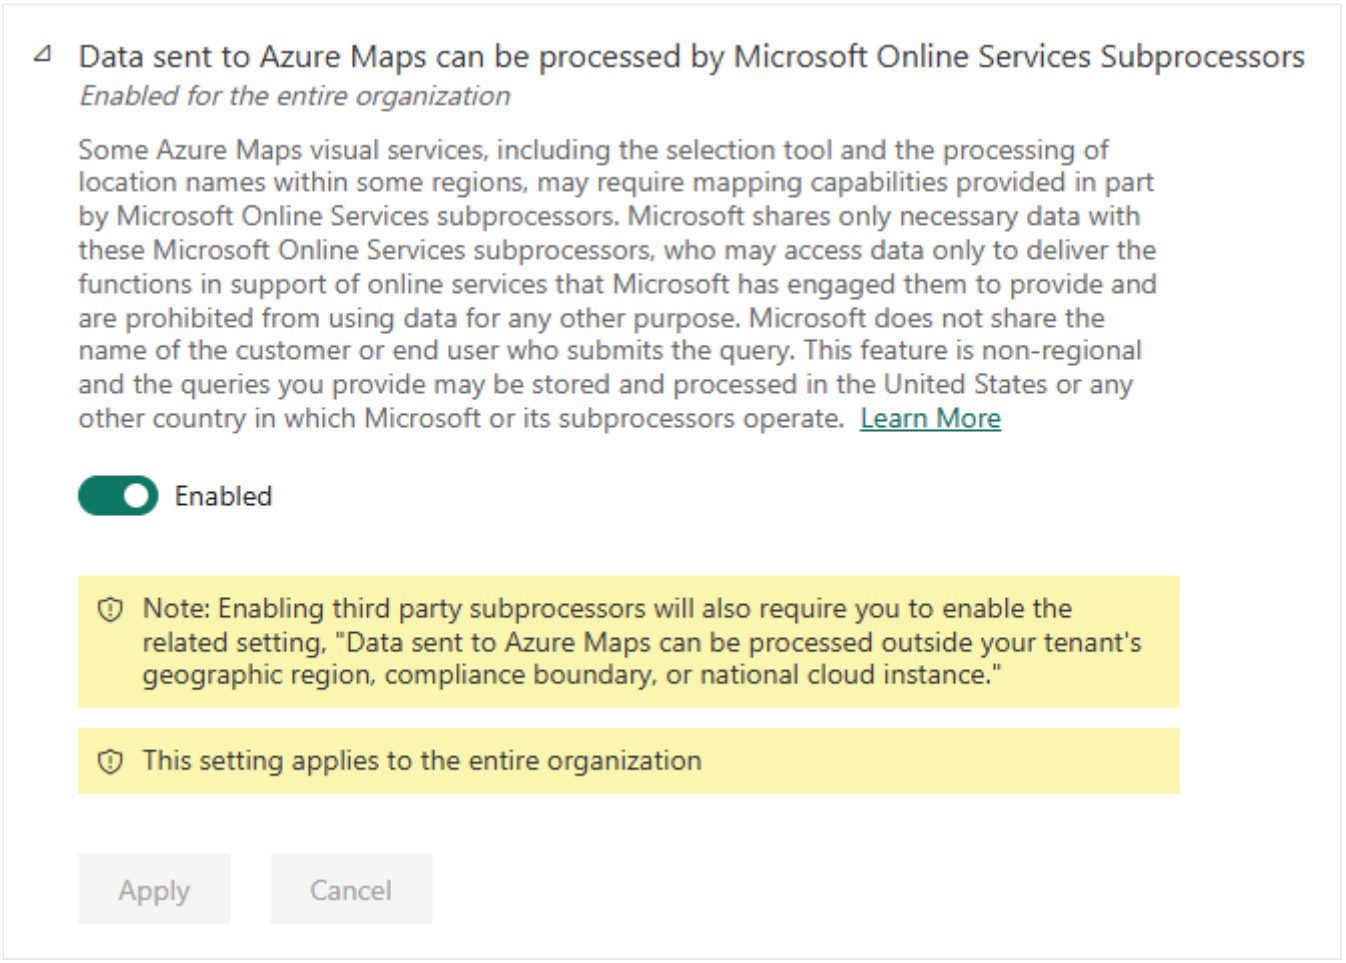

Map and filled map visuals

When enabled, users in the organization can use map and filled map visualizations in their reports.

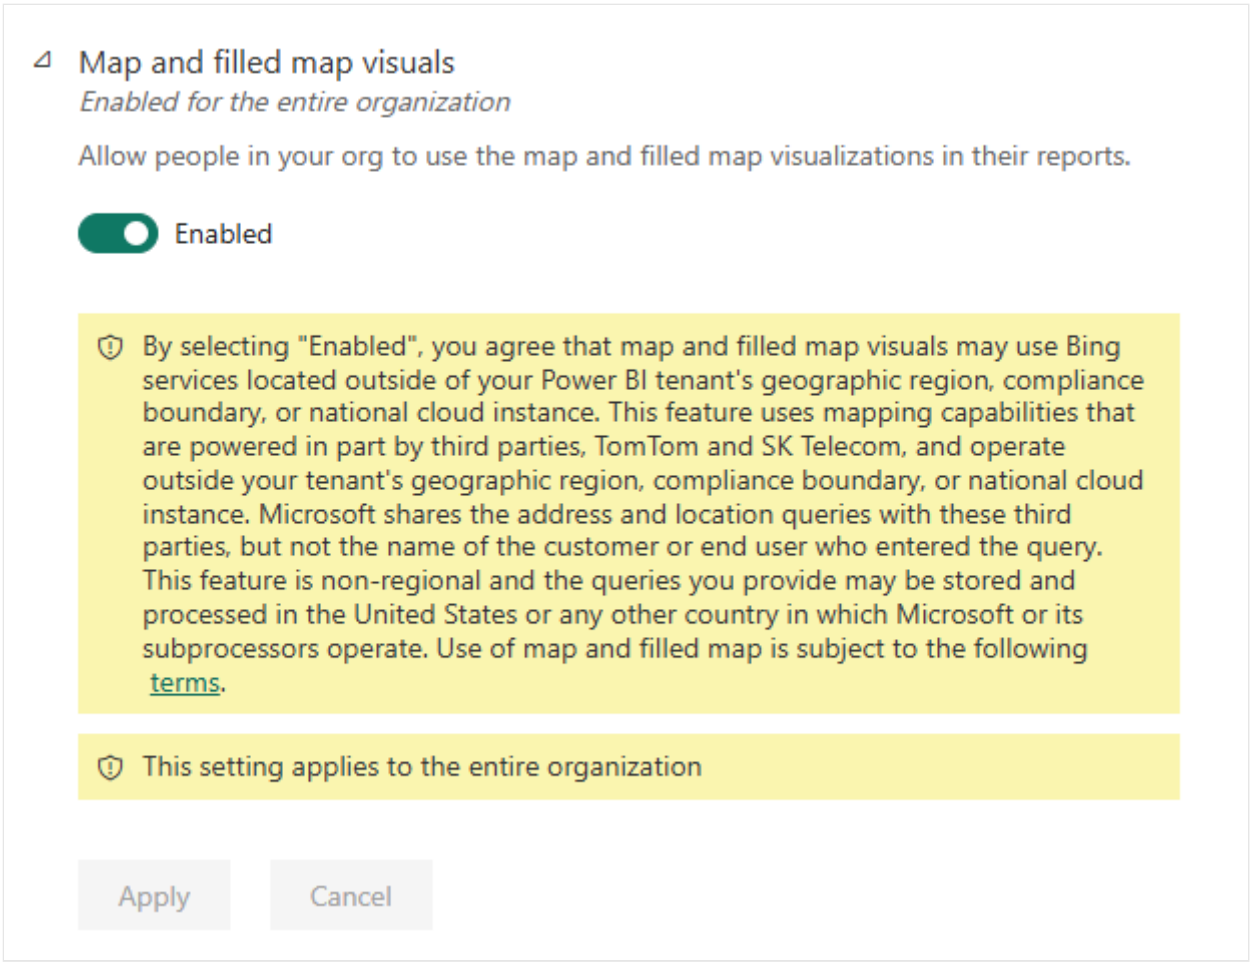

> **Note**
>
> In a future release, Power BI plans to deprecate older map visuals and migrate existing reports to Azure Maps. Learn about converting to Azure Maps.

Integration with SharePoint and Microsoft Lists

Users in the organization can create Fabric reports directly from SharePoint and Microsoft Lists. Then they can build Fabric reports on the data in those lists and publish them back to the lists, to be visible to others who can access the list.

This setting is enabled by default. Even if the feature is disabled, in SharePoint and Microsoft Lists users can still see **Power BI** > **Visualize the list**, and any existing reports, on the **Integrate** menu. If they select **Visualize the list**, they go to an error page explaining that their admin disabled the feature.

Learn more about creating reports from SharePoint and Microsoft Lists.

Dremio SSO

Enable SSO capability for Dremio. By enabling, user access token information, including name and email, is sent to Dremio for authentication.

To learn more, see Microsoft Entra ID-based single sign-on for Dremio Cloud and Power BI.

Snowflake SSO

For semantic model owners to be able to enable single sign-on for DirectQuery connections to Snowflake in semantic model settings, a Fabric admin must enable the **Snowflake SSO** setting. This setting approves sending Microsoft Entra credentials to Snowflake for authentication for the entire organization.

To learn more, see Connect to Snowflake in the Power BI Service.

Redshift SSO

Enable SSO capability for Redshift. By enabling, user access token information, including name and email, is sent to Redshift for authentication.

To learn more, see Overview of single sign-on for on-premises data gateways in Power BI.

Google BigQuery SSO

Enable SSO capability for Google BigQuery. By enabling, user access token information, including name and email, is sent to Google BigQuery for authentication.

To learn more, see Google BigQuery (Azure AD).

Microsoft Entra single sign-on (SSO) for Gateway

This setting enables Microsoft Entra SSO through gateways to cloud data sources that rely on Microsoft Entra ID-based authentication. It gives seamless Microsoft Entra SSO connectivity to Azure-based data sources, such as Azure Synapse Analytics (SQL DW), Azure Data Explorer, Snowflake on Azure, and Azure Databricks through an on-premises data gateway.

This feature is important for users who work with reports that require SSO connectivity in DirectQuery mode to data sources deployed in an Azure virtual network (Azure VNet). When you configure SSO for an applicable data source, queries execute under the Microsoft Entra identity of the user that interacts with the Power BI report.

An important security-related consideration is that gateway owners have full control over their on-premises data gateways. This means that it's theoretically possible for a malicious gateway owner to intercept Microsoft Entra SSO tokens as they flow through an on-premises data gateway (this isn't a concern for VNet data gateways because they're maintained by Microsoft).

A related consideration applies to VNet data gateways, although through a different vector. While VNet data gateways themselves are maintained by Microsoft (so the OPDG interception vector doesn't apply), a VNet gateway operates within the network perimeter of a VNet owned by the customer. As a result, a malicious VNet owner could theoretically intercept SSO tokens at the network layer — traffic flowing in and out of the gateway is subject to the network security posture of the hosting VNet. Customers should ensure their VNet network controls (NSGs, route tables, peering, and monitoring) follow least-privilege and zero-trust principles to mitigate this risk.

Because of this possible threat, the Microsoft Entra SSO feature is disabled by default for on-premises data gateways. As a Fabric admin, you must enable the **Microsoft Entra single sign-on (SSO) for Gateway** tenant setting in the Fabric admin portal before data sources can be enabled for Microsoft Entra SSO on an on-premises data gateway or VNet gateway. Before enabling the feature, make sure to restrict the ability to deploy on-premises data gateways in your organization to appropriate administrators, and apply equivalent governance to who can deploy VNet data gateways and manage the underlying VNets.

To learn more, see Microsoft Entra SSO.

Users can view Power BI files saved in OneDrive and SharePoint (Preview)

This setting allows users to view Power BI files saved in OneDrive and SharePoint Online document libraries in their browser without needing to download the file and open in Power BI Desktop on their local machine. When enabled, the setting applies to all users in your organization. This setting is on by default.

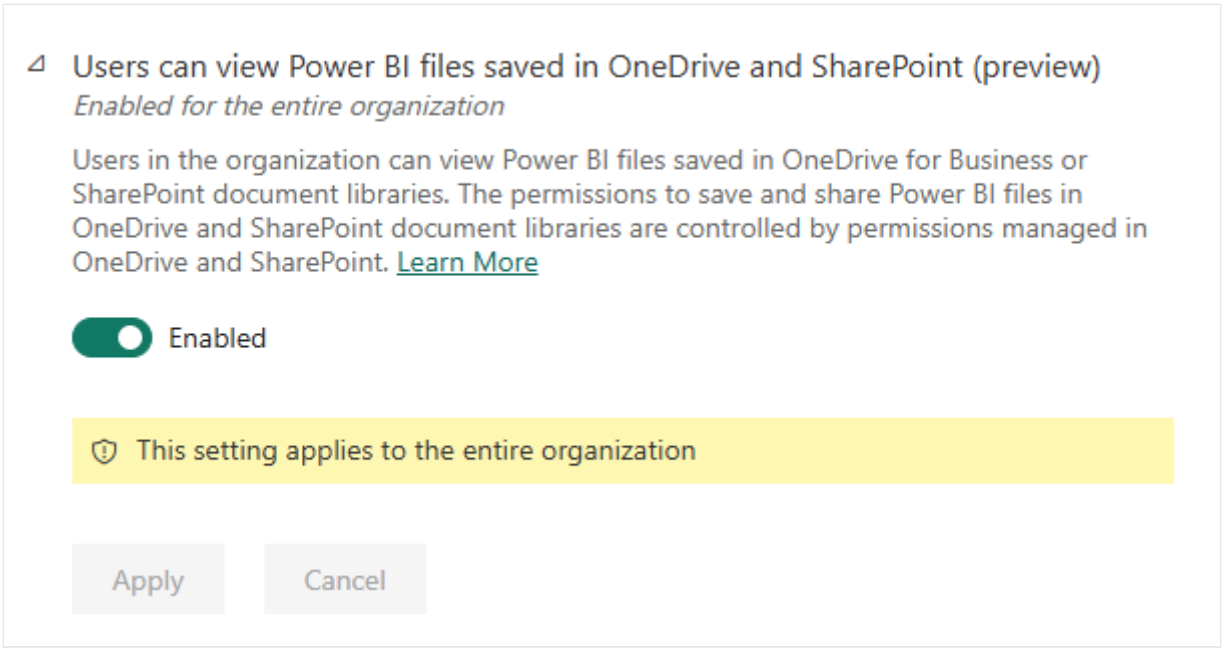

Learn more about viewing Power BI files saved in OneDrive and SharePoint.

Users can share links to Power BI files stored in OneDrive and SharePoint through Power BI Desktop (preview)

Users can share links to Power BI Desktop files (.pbix) saved to OneDrive and SharePoint through Power BI Desktop. Sharing uses standard OneDrive and SharePoint sharing functionality. When enabled, this setting applies to all users in your organization.

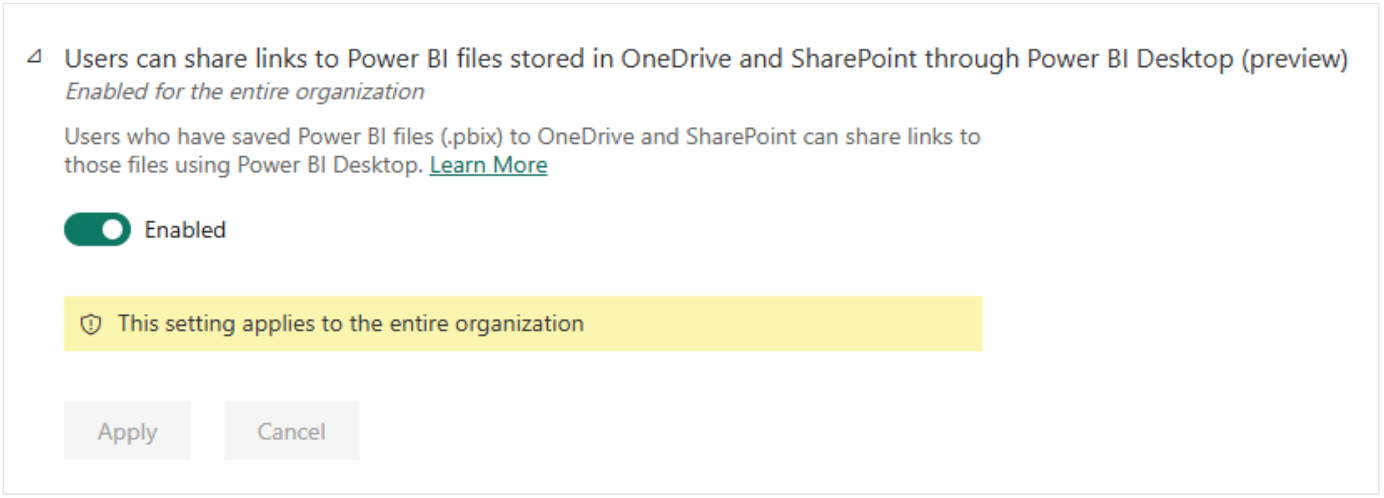

During public preview, if a user enables share through the Power BI Desktop menu, but the admin setting is disabled for the tenant, a **Share** button still appears in Power BI Desktop, but the user is notified that the capability is disabled when they attempt to share.

Learn more about sharing links through Power BI Desktop.

Enable granular access control for all data connections

Turn on this setting to enforce strict access control for all data connection types. When this setting is on, shared items are disconnected from data sources if they're edited by users who don't have permission to use those data connections.

For more information, see Manage semantic model credentials and connections.

Semantic models can export data to OneLake

Semantic models configured for OneLake integration can export tables in import mode to OneLake. Once exported, users can use the data in other Fabric items, including lakehouses and warehouses.

For more information, see OneLake integration overview.

Semantic model owners can choose to automatically update semantic models from files imported from OneDrive or SharePoint

Turn on this setting to let semantic model owners choose to automatically update semantic models with changes made to corresponding Power BI files stored in OneDrive or SharePoint.

For more information, see Refresh from OneDrive or SharePoint.

ArcGIS GeoAnalytics for Fabric Runtime

Users in your organization can use Esri's ArcGIS GeoAnalytics for Fabric Runtime in Microsoft Fabric Spark Runtime.

For more information, see ArcGIS GeoAnalytics for Fabric Runtime.

Allow non-Entra ID auth in Eventstream

This setting lets users keep key-based authentication enabled in Eventstream custom endpoints. Turning it off restricts Eventstream custom endpoint authentication to Microsoft Entra ID only.

Related content

- About tenant settings

- Tenant settings index

Last updated on 05/17/2026

Power BI visuals tenant settings

These settings are configured in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

All the Power BI visuals admin settings, including Power BI visuals tenant settings, are described in Manage Power BI visuals admin settings.

Allow visuals created using the Power BI SDK

Users in the organization can add, view, share, and interact with visuals imported from AppSource or from a file. Visuals allowed in the *Organizational visuals* page aren't affected by this setting.

To learn more, see Visuals from AppSource or a file.

Add and use certified visuals only (block uncertified)

Users in the organization with permissions to add and use visuals can add and use certified visuals only. Visuals allowed in the *Organizational visuals* page aren't affected by this setting, regardless of certification.

To learn more, see Certified Power BI visuals.

Allow downloads from custom visuals

Enabling this setting lets custom visuals download any information available to the visual (such as summarized data and visual configuration) upon user consent. It's not affected by download restrictions applied in your organization's Export and sharing settings.

To learn more, see Export data to file.

AppSource Custom Visuals SSO

Enable this setting to allow AppSource custom visuals to use single sign-on. This feature allows custom visuals from AppSource to get Microsoft Entra ID access tokens for signed-in users through the Authentication API.

Microsoft Entra ID access tokens include personal information, including users' names and email addresses, and might be sent across regions and compliance boundaries.

To learn more, see AppSource custom visuals SSO.

Allow access to the browser's local storage

When this setting is on, custom visuals can store information in the user's browser local storage.

To learn more, see AppSource custom visuals SSO.

Related content

- About tenant settings
- Manage Power BI visuals admin settings

Last updated on 09/29/2026

R and Python visuals tenant settings

These settings are configured in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

Interact with and share R and Python visuals

Users in the organization can interact with and share visuals created with R or Python scripts.

- Create and use R visuals in Power BI.
- Create Power BI visuals with Python.

> **Note**
>
> This setting applies to the entire organization and can't be limited to specific groups.

Related content

- About tenant settings

Last updated on 09/29/2026

Audit and usage tenant settings

These settings are configured in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

Usage metrics for content creators

When this setting is enabled, users in the organization can see usage metrics for dashboards, reports and semantic models that they have appropriate permissions to.

To learn more, see Monitor usage metrics in the workspaces.

Per-user data in usage metrics for content creators

When this setting is enabled, content creator account information (such as user name and email address) will be exposed in the metrics report. If you don't wish to gather this information for all users, you can disable the feature for specified security groups or for an entire organization. Account information for the excluded users then shows in the report as *Unnamed*.

Per-user data is enabled for usage metrics by default.

Show user data in the Microsoft Fabric Capacity Metrics app and reports

When this setting is enabled, active user data, including names and email addresses, is displayed in the Microsoft Fabric Capacity Metrics app and reports. This setting is enabled by default.

Azure Log Analytics connections for workspace administrators

Power BI integration with Azure Log Analytics enables Fabric administrators and Premium workspace owners to connect their Premium workspaces to Azure Log Analytics to monitor the connected workspaces.

When the setting is enabled, administrators and Premium workspace owners can configure **Azure Log Analytics for Power BI**.

Workspace admins can turn on monitoring for their workspaces (preview)

When this setting is enabled, workspace admins can turn on monitoring for their workspaces. When a workspace admin turns on monitoring, a read-only Eventhouse that includes a KQL database is created. After the Eventhouse and KQL database are added to the workspace, logging is turned on and data is sent to the database. Enable workspace monitoring, a feature that allows workspace admins to monitor their workspace.

Microsoft can store query text to aid in support investigations

When this setting is enabled, Microsoft can store the query text generated when users use Fabric items such as reports and dashboards. This data is sometimes necessary for debugging and resolving complex issues related to the performance and functionality of Fabric items such as semantic models. The setting is enabled by default.

Storing and retaining query text data can have implications for data security and privacy. While it is recommended to leave the setting on to facilitate support, if there are organizational requirements that don't permit storing query text, or if you wish to opt out of this feature for any other reason, you can turn off the feature as follows:

Go to **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**, find **Microsoft can store query text to aid in support investigations** in the **Audit and usage** section, and set the toggle to **Disabled**.

For more information about the diagnostic query text storage feature, see Diagnostic query text storage.

Related content

- About tenant settings
- Diagnostic query text storage

Last updated on 09/29/2026

Dashboard tenant settings

These settings are configured in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

Web content on dashboard tiles

If this setting is enabled, users in the organization can add and view web content tiles on Power BI dashboards.

To learn more, see Add images, videos, and more to your dashboard.

> **Note**
>
> Enabling this feature could expose your organization to security risks via malicious web content. Due to these security risks, the setting is disabled by default.

Related content

- About tenant settings

Last updated on 09/29/2026

Git integration tenant settings

Configure the Git integration tenant admin settings in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. The tenant admin can choose to delegate control of these switches to the workspace admin or capacity admin. If the tenant admin enables delegation, the capacity admin can override the tenant admin's decision to enable or disable the switch. The workspace admin can override the tenant and the capacity settings.

For information about how to get to and use tenant settings, see About tenant settings.

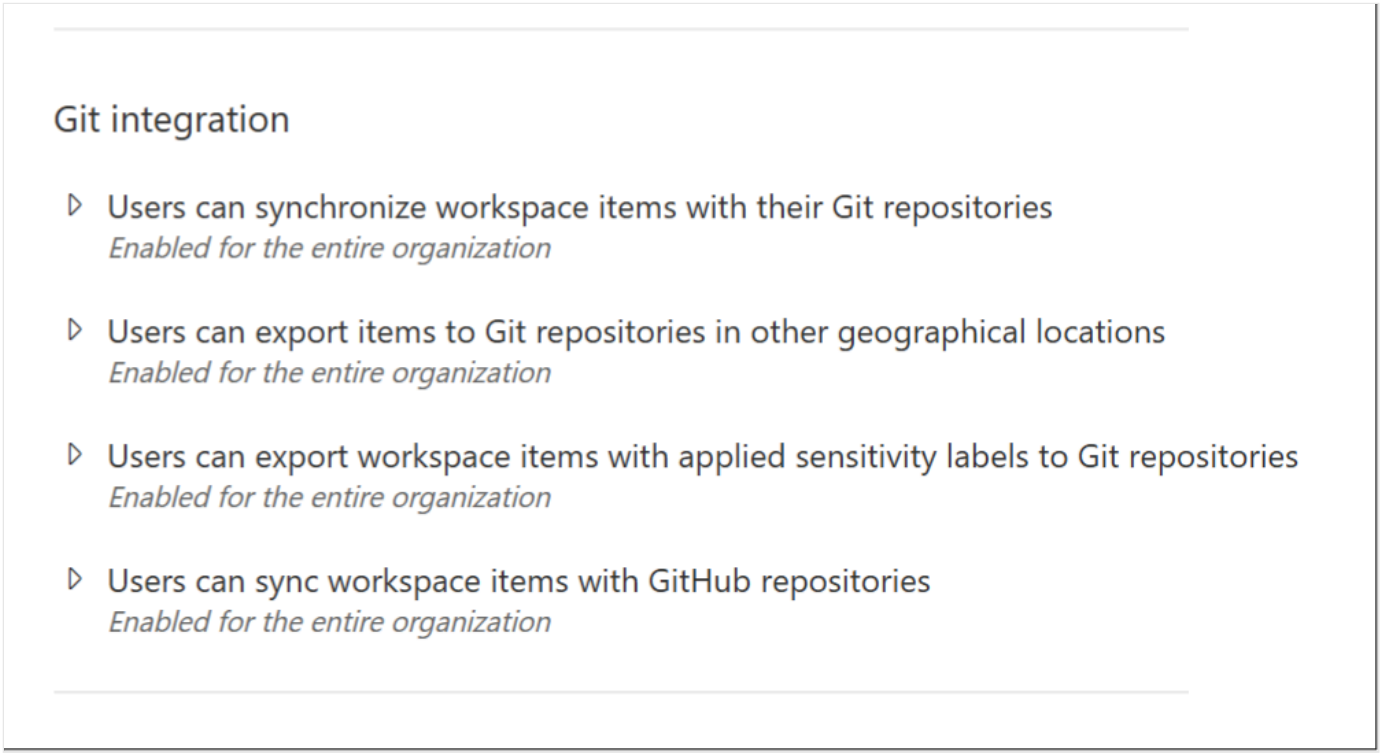

> **Important**
>
> The switches that control Git integration are part of Fabric and only work if the Fabric admin switch is turned on. If Fabric is disabled, Git integration will work for workspaces that contain only Power BI items.

Shared Git integration tenant admin settings

The following section provides a brief overview of the general settings found in each of the Git integration tenant admin settings.

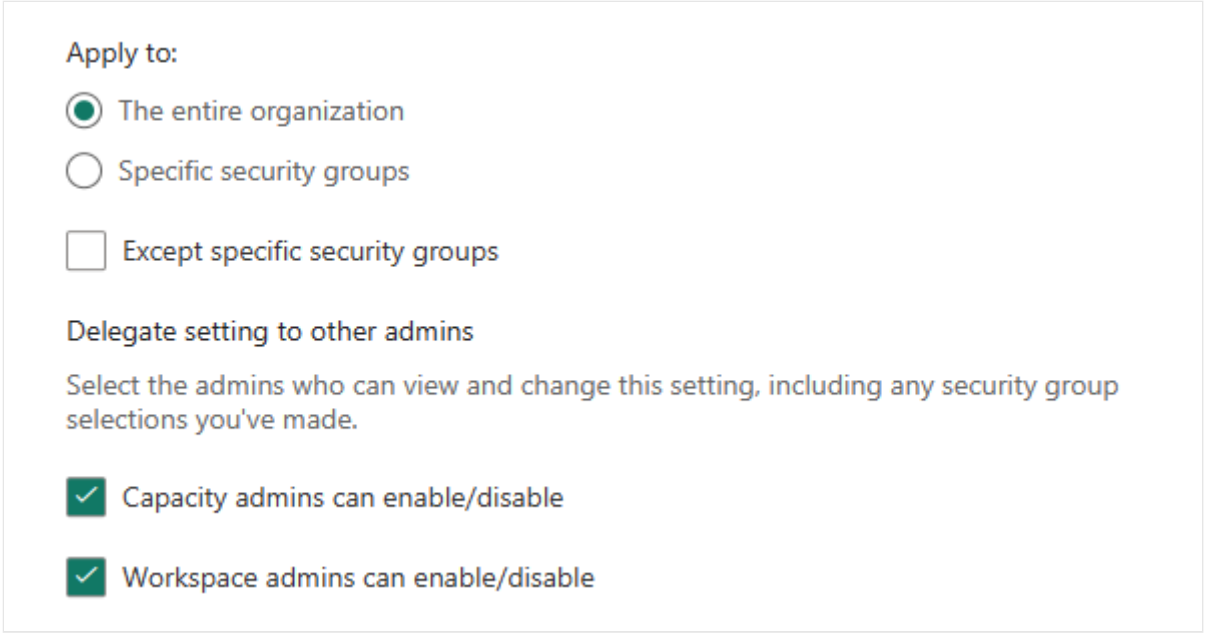

Specify which users a setting applies to

When you enable a Git integration tenant setting, use the **Apply to** options to control which users the setting affects:

- **The entire organization**: The setting applies to every user in the tenant.
- **Specific security groups**: The setting applies only to members of the security groups you enter. Users who aren't members of those groups aren't affected. The security groups you specify are included by default.
- **Except specific security groups**: Use this option to exclude one or more security groups. It's optional and can be combined with the options in the previous two bullet points, so you can enable a setting broadly and still carve out groups to exclude. Enter the security groups you want to exclude in the box that appears.

For example, you can enable a switch for the entire organization but exclude a security group whose members shouldn't have the capability. Or you can enable it for a specific security group while excluding a subgroup within it.

Delegate settings to other admins

By default, only the Fabric (tenant) admin controls a Git integration tenant setting. To let other admins manage the setting for their own scope, expand **Delegate settings to other admins** and select who can override it:

- **Capacity admins can enable/disable**: A capacity admin can override the tenant admin's choice for their capacity. The override applies only to the workspaces assigned to that

capacity. Capacity admins make the change from the capacity's **Delegated tenant settings** page by selecting **Override tenant admin selection**.

- **Workspace admins can enable/disable**: A workspace admin can override both the tenant and capacity settings for their own workspace.

If you don't select either option, the setting stays locked at the value set at the tenant level. When you delegate, the result is the layered precedence described earlier: tenant, then capacity, then workspace, with each lower level able to override only when delegation is turned on.

For more information on delegation, see Delegation and how it integrates with Git.

Users can synchronize workspace items with their Git repositories

Users can synchronize a workspace with an Azure Git repository, edit their workspace, and update their Git repos using the Git integration tool. You can enable Git integration for the entire organization, or for a specific group.
This switch is **enabled** by default. Disable it to prevent users from syncing workspace items with their Git repositories.

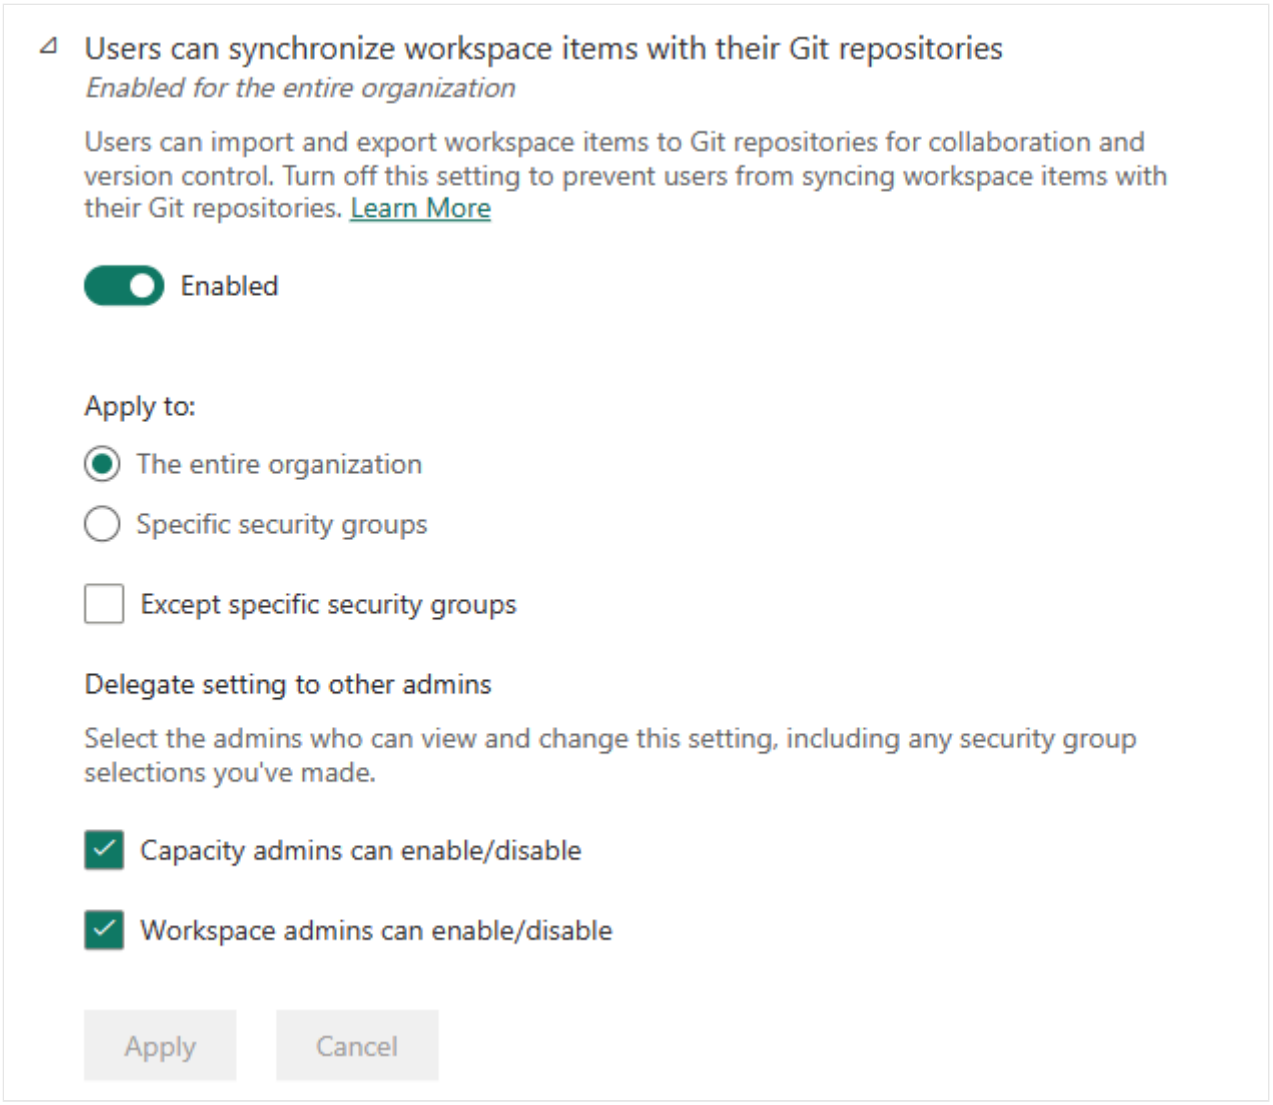

To learn more, see Introduction to Git integration.

To get started with Git integration, see Manage a workspace with Git.

Users can export items to Git repositories in other geographical locations

If a workspace capacity is in one geographic location (for example, Central US) while the *Azure DevOps* repo is in another location (for example, West Europe), the Fabric admin can decide whether to allow users to commit metadata (or perform other Git actions) to another geographical location. Only the metadata of the item is exported. Item data and user related information are not exported.

Enable this setting to allow all users, or a specific group of users, to export metadata to other geographical locations.

> **Important**

Tenant settings apply to the Fabric user who performs the Git operation, not the account used to connect to the repository. If you enable this setting for specific security groups, include every Fabric user who needs this capability in an allowed group, even when a different account is used to connect the Git repository.

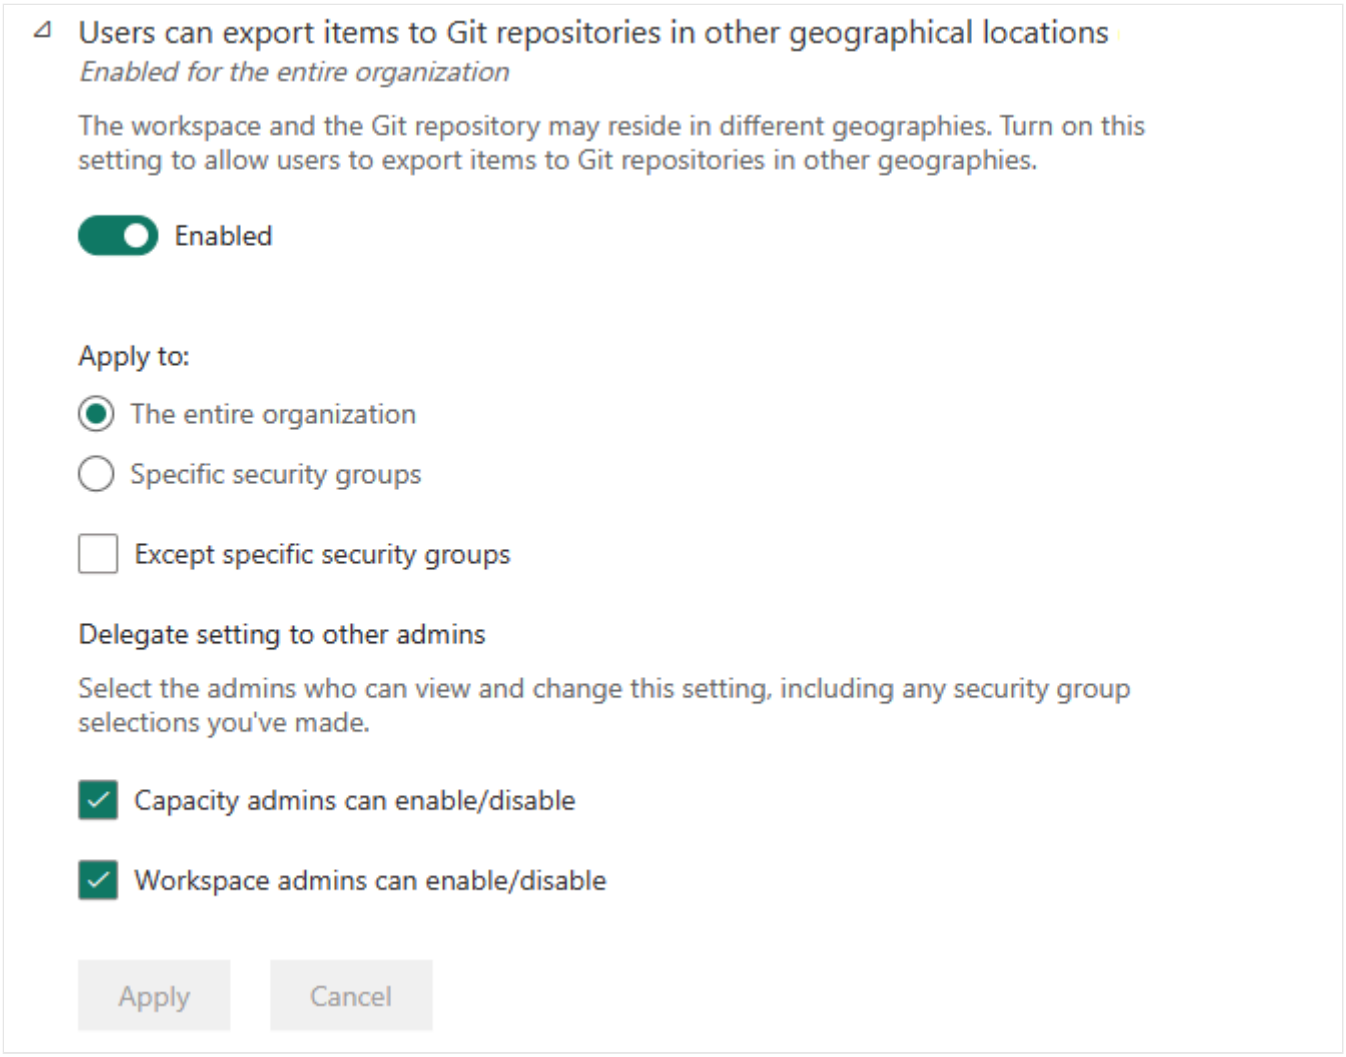

Note

GitHub doesn't support enforcement of this switch.

Users can export workspace items with applied sensitivity labels to Git repositories

Sensitivity labels aren't included when exporting an item. Therefore, the Fabric admin can choose whether to block the export of items that have sensitivity labels, or to allow it even though the sensitivity label won't be included.

Enable this setting to allow all users, or a specific group of users, to export items without their sensitivity labels.

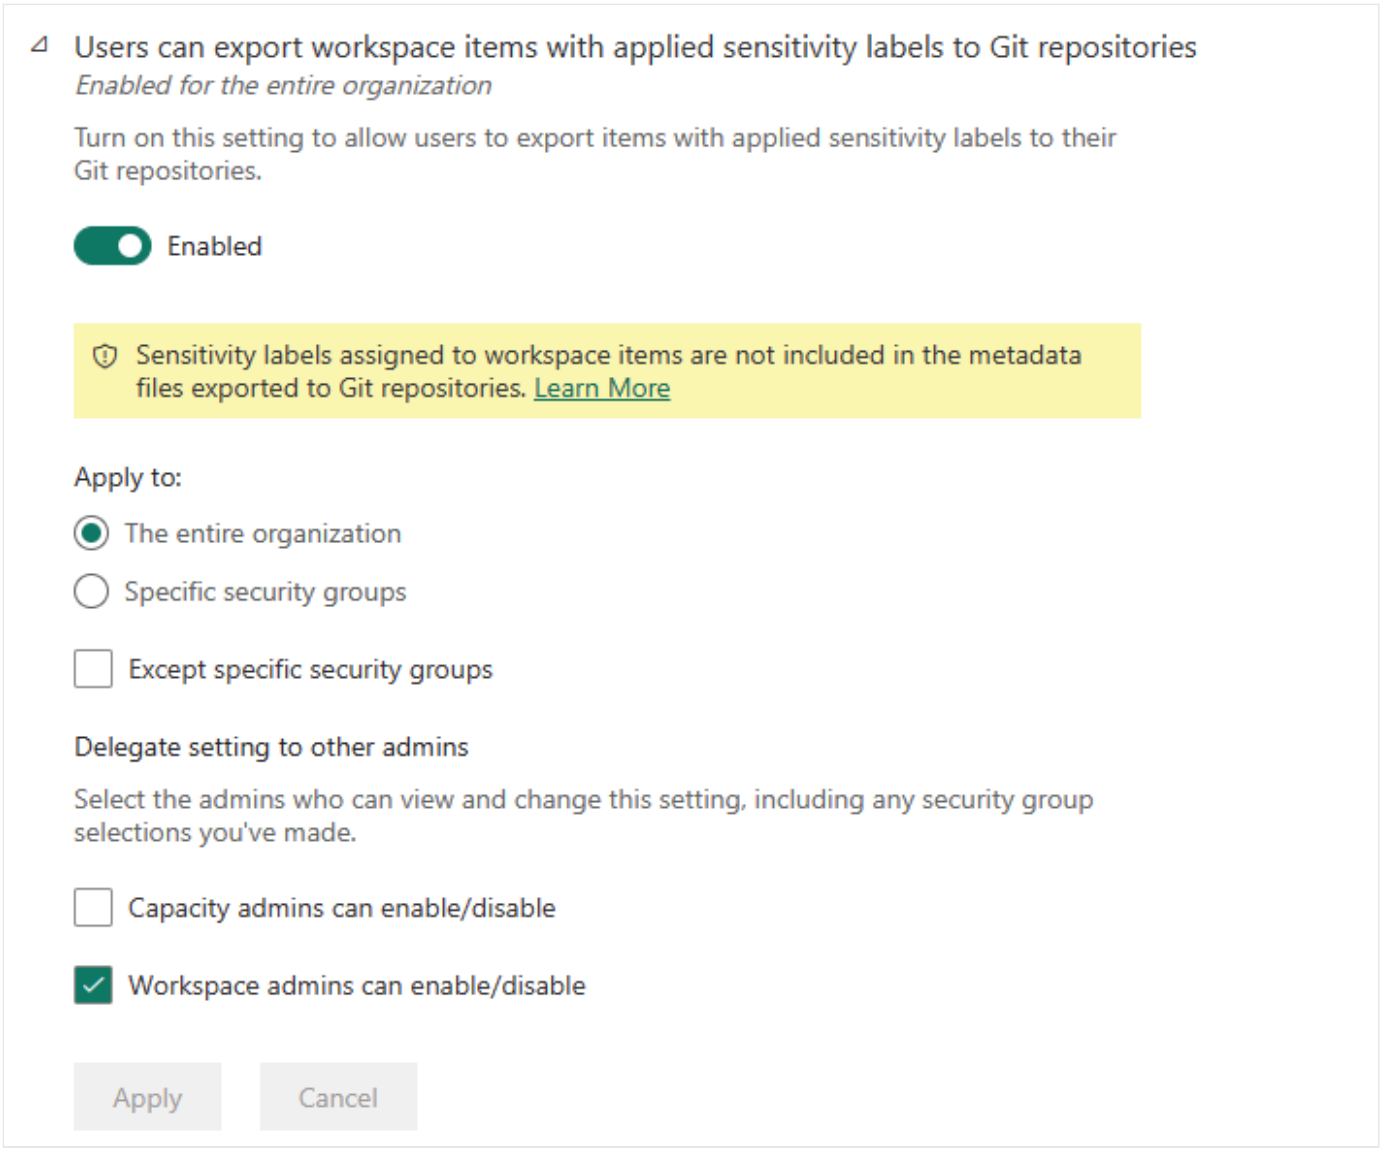

Learn more about sensitivity labels.

Users can sync workspace items with GitHub repositories

Users can synchronize a workspace with their GitHub repository, edit their workspace, and update their GitHub repos using the Git integration tool. You can enable Git integration for the entire organization, or for a specific group.

This switch is **disabled** by default. Enable it to allow users to sync workspace items with their Git repositories.

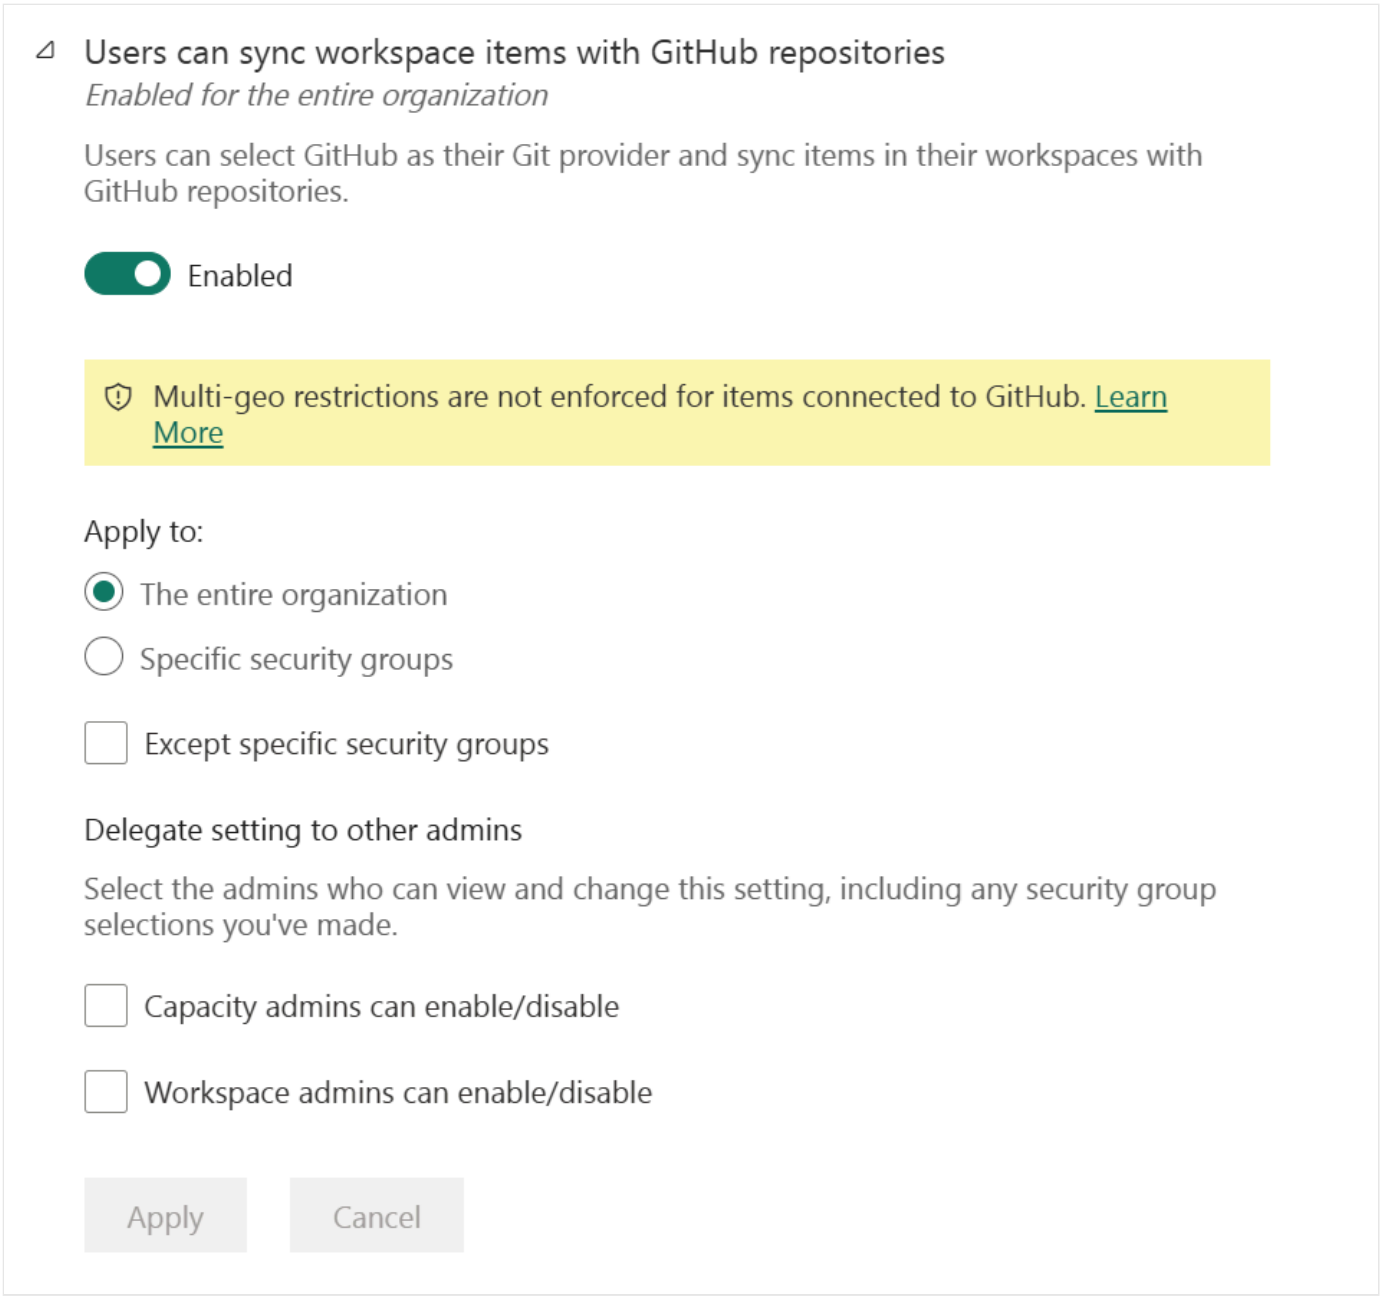

To learn more, see Introduction to Git integration.

To get started with Git integration, see Manage a workspace with Git.

Delegation and how it works with Git integration

Delegation is the admin model that lets the tenant admin push control of a tenant setting down to lower administrative scopes instead of locking every decision at the org level.

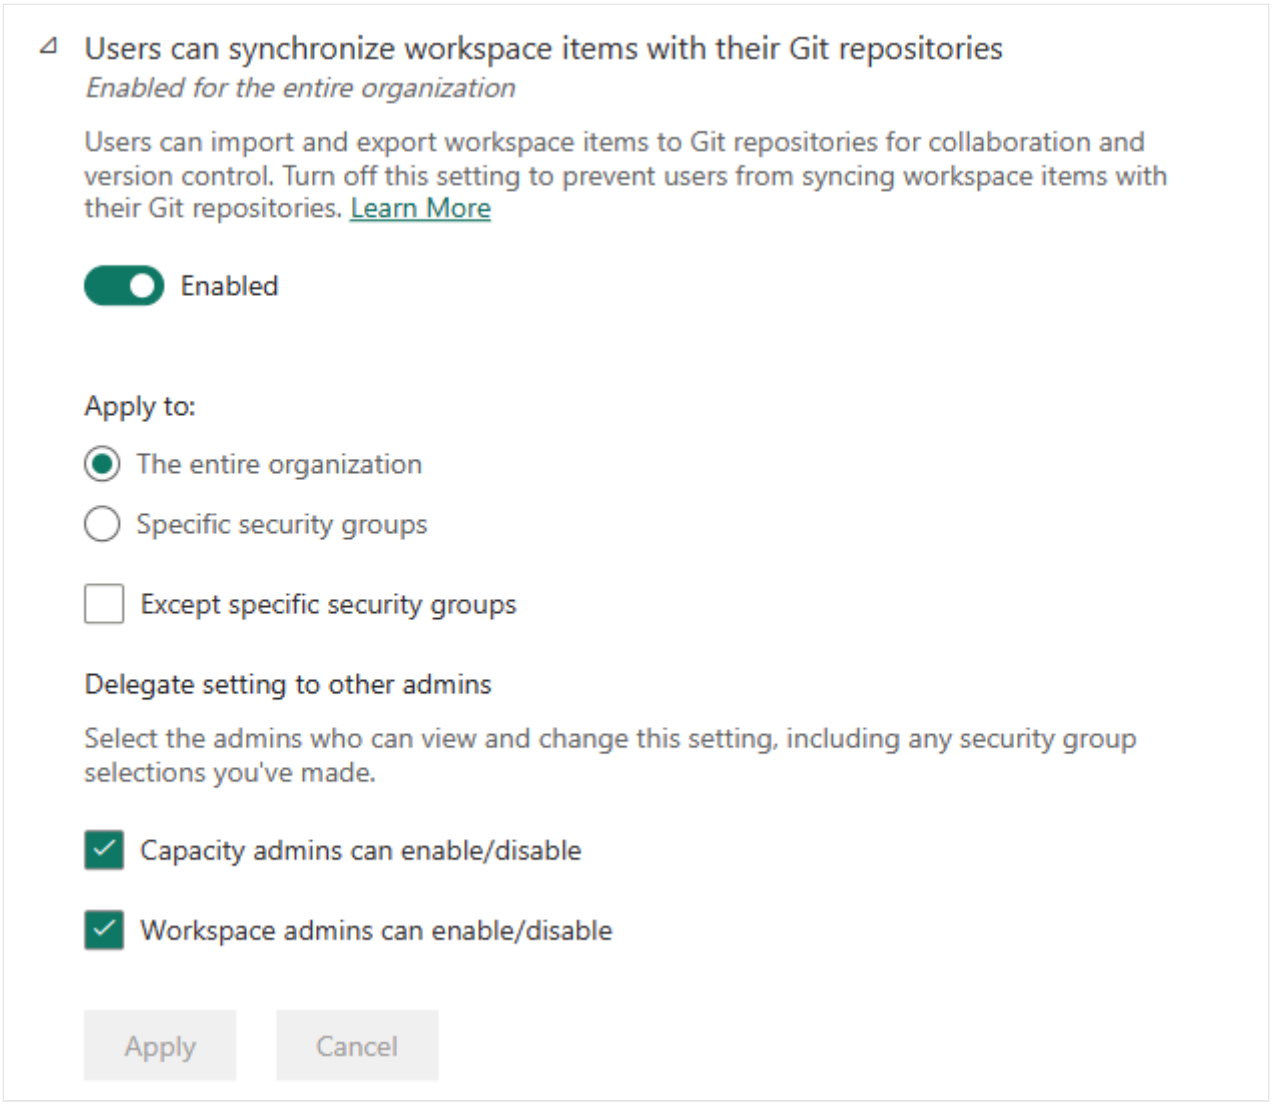

For the Git integration switches, the layered override model is:

- Tenant admin sets the switch in the tenant settings and chooses whether to delegate it.
- If delegation is enabled, the capacity admin can override the tenant admin's enable or disable decision for their capacity.
- The workspace admin can override both the tenant and capacity settings for their workspace.

So the precedence is tenant → capacity → workspace, with each lower level able to override when delegation is turned on. For more information, see Power BI implementation planning: Workspaces at the workspace level.

Related content

- About tenant settings

ⓘ**Note:** The author created this article with assistance from AI. Learn more

Last updated on 10/07/2026

Developer tenant settings

These settings are configured in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

To manage Power BI developer settings, you must be a Fabric administrator. For more information about the Fabric administrator role, see Understand Fabric admin roles.

> **Note**
>
> The developer settings in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings** are different from and not related to the developer mode setting for debugging visuals.

Embed content in apps

Users in the organization can embed Power BI dashboards and reports in software as a service (SaaS) applications. Disabling this setting prevents users from being able to use the REST APIs to embed Power BI content within their application.

To learn more, see What is Power BI embedded analytics?.

Learn about the Embed for your customers method to build an app that uses non-interactive authentication against Power BI.

Service principals can create workspaces, connections, and deployment pipelines

Use a service principal to access these Fabric APIs that aren't protected by a Fabric permission model.

- Create Workspace
- Create Connection
- Create Deployment Pipeline

To allow an app to use service principal authentication, its service principal must be included in an allowed security group. You can control who can access service principals by creating

dedicated security groups and using these groups in other tenant settings.

This setting is disabled by default for new customers.

Service principals can call Fabric public APIs

Use a service principal to access Fabric public APIs that include create, read, update, and delete (CRUD) operations, and are protected by a Fabric permission model.

To allow an app to use service principal authentication, its service principal must be included in an allowed security group. You can control who can access service principals by creating dedicated security groups and using these groups in other tenant settings.

This setting is enabled by default for new customers.

Allow service principals to create and use profiles

An app owner with many customers can use service principal profiles as part of a multitenancy solution to enable better customer data isolation and establish tighter security boundaries between customers.

To learn more, see Service principal profiles for multitenancy apps.

Block ResourceKey Authentication

For extra security, you can block the use of resource key-based authentication. The Block ResourceKey Authentication setting applies to streaming and PUSH datasets. If disabled, users aren't allowed to send data to streaming and PUSH datasets using the API with a resource key.

This setting applies to the entire organization. You can't apply it only to a select security group.

Define maximum number of Fabric identities in a tenant

This setting specifies the maximum number of Fabric identities that can be created in a Fabric tenant. If this setting is disabled, a default value of 10,000 Fabric identities can be created in a tenant. If you enable this setting, you can set a custom limit of how many Fabric identities can be created in your Fabric tenant. You can also manage this value programmatically by using the Update Tenant Setting REST API.

The maximum number of Fabric identities defined and created for a Fabric tenant can't exceed the maximum number of resources allowed in a Microsoft Entra tenant. Fabric doesn't validate that your custom limit falls within your Microsoft Entra ID resource quota. Before setting a custom limit, check your tenant's Microsoft Entra service limits. To increase your resource limit in your Microsoft Entra tenant, contact Microsoft support.

Related content

- About tenant settings
- Enable service principal authentication for developer settings

---

Last updated on 09/29/2026

Admin API tenant settings

These settings are configured in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

Service principals can access read-only admin APIs

Web apps registered in Microsoft Entra ID use an assigned service principal to access read-only admin APIs without a signed-in user. To allow an app to use service principal authentication, its service principal must be included in an allowed security group. By including the service principal in the allowed security group, you're giving the service principal read-only access to all the information available through admin APIs (current and future). For example, user names and emails, semantic model, and report detailed metadata.

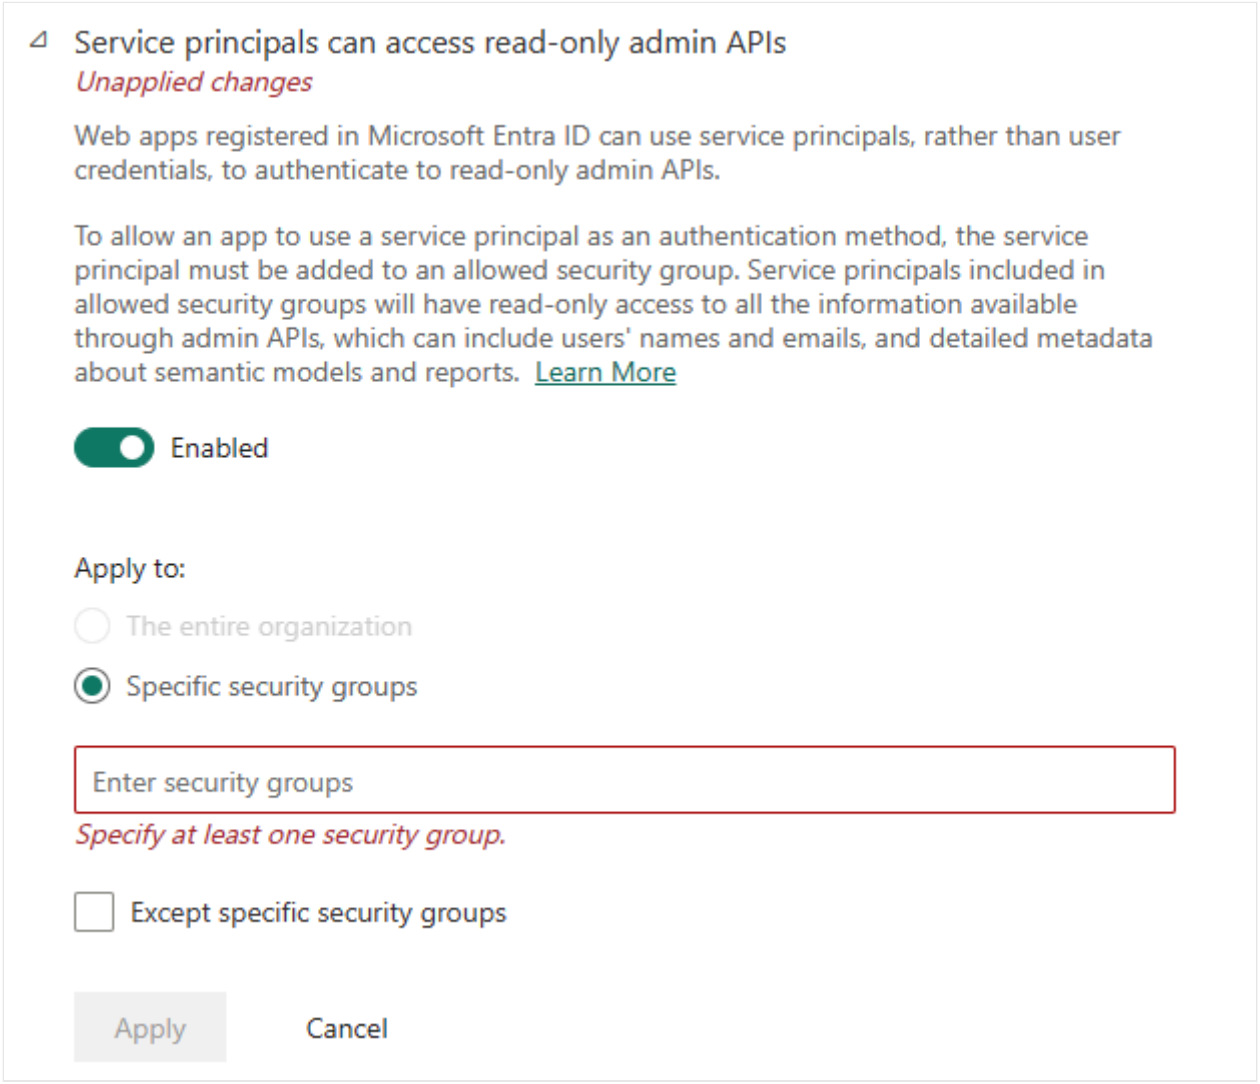

To learn more, see Allow service principals to use read-only admin APIs

> **Note**
>
> This setting is also required to run Fabric data risk assessments in Microsoft Purview Data Security Posture Management (DSPM) for AI.

Service principals can access admin APIs used for updates

Web apps registered in Microsoft Entra ID use an assigned service principal to access Fabric admin APIs used for updates without a signed-in user. To allow an app to use service principal authentication, include its service principal in an allowed security group.

> **Note**
>
> This setting is also required to run Fabric data risk assessments in Microsoft Purview DSPM for AI.

Enhance admin APIs responses with detailed metadata

Users and service principals allowed to call Power BI admin APIs might get detailed metadata about Power BI items. For example, responses from GetScanResult APIs contain the names of semantic model tables and columns.

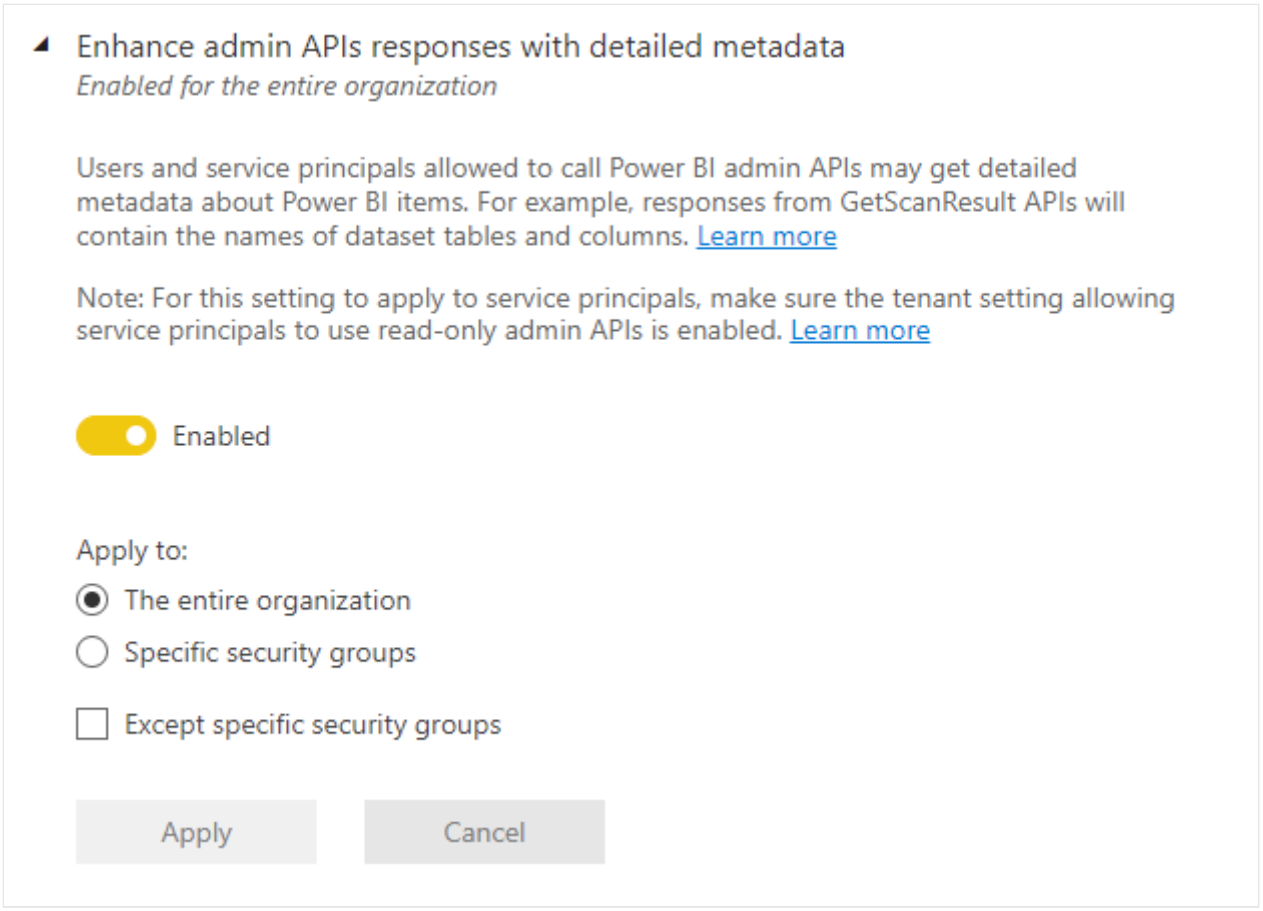

To learn more, see Metadata scanning.

> ⓘ **Note**
>
> For this setting to apply to service principals, make sure the tenant setting **Allow service principals to use read-only admin APIs** is enabled. To learn more, see Set up metadata scanning.

Enhance admin APIs responses with DAX and mashup expressions

Users and service principals eligible to call Power BI admin APIs get detailed metadata about queries and expressions comprising Power BI items. For example, responses from GetScanResult API contain DAX and mashup expressions.

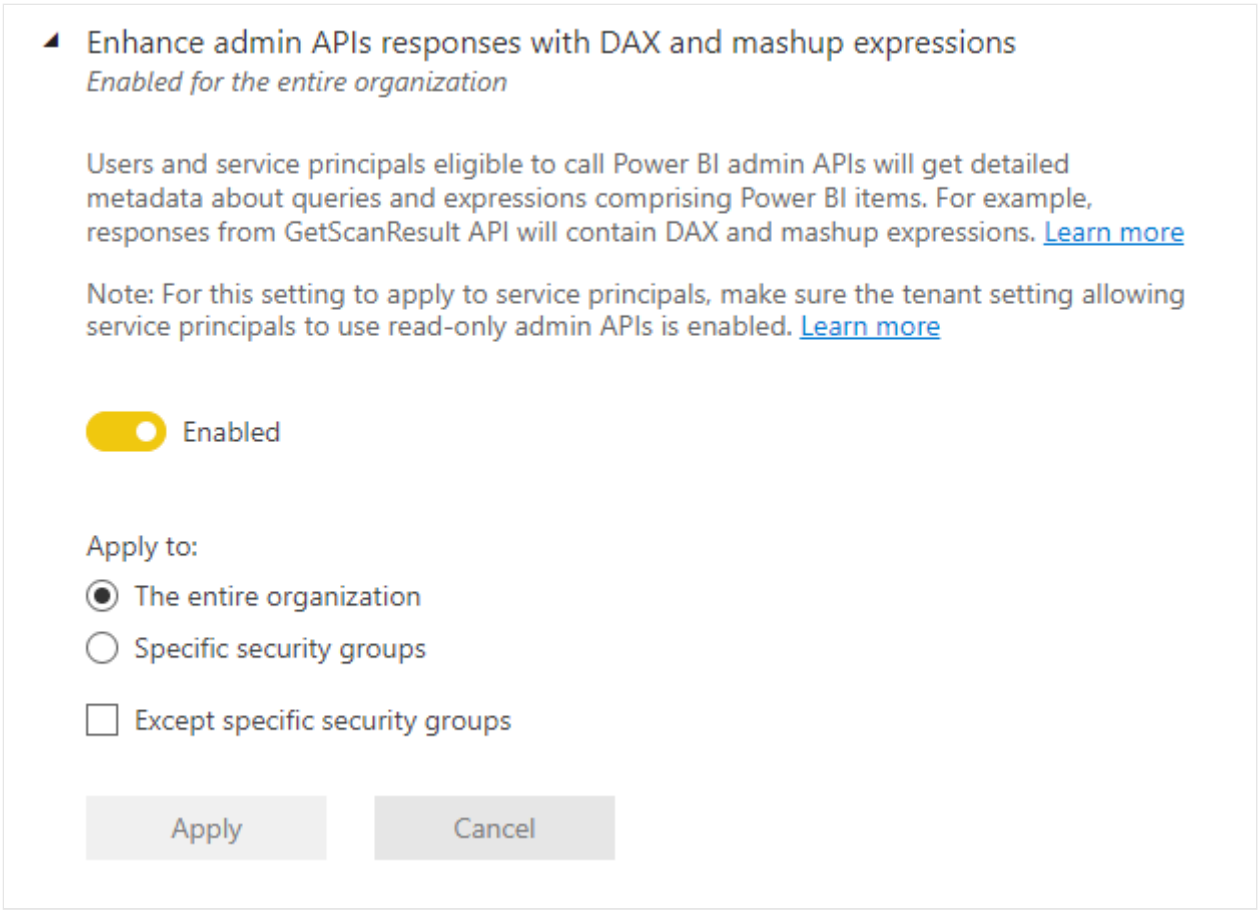

To learn more, see Metadata scanning.

> **Note**
>
> For this setting to apply to service principals, make sure the tenant setting **Allow service principals to use read-only admin APIs** is enabled. To learn more, see Set up metadata scanning.

Related content

- About tenant settings

Last updated on 09/29/2026

Gen 1 dataflow tenant settings

These settings are configured in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

Create and use Gen1 dataflows

Users in the organization can create and use dataflows. For an overview of dataflows, see Introduction to dataflows and self-service data prep. To enable dataflows in a Premium capacity, see Configure workloads.

> **Note**
>
> This setting applies to the entire organization and can't be limited to specific groups.

Related content

- About tenant settings

Last updated on 09/29/2026

Template app tenant settings

These settings are configured in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

Publish template apps

Users in the organization can create template apps workspaces. Control which users can publish template apps or distribute them to clients outside your organization by way of Microsoft AppSource or other distribution methods.

To learn more about template apps, see What are Power BI template apps?.

Install template apps

Users in the organization can download and install template apps **only** from Microsoft AppSource. Control which specific users or security groups can install template apps from AppSource.

To learn about AppSource, see What is Microsoft AppSource?.

Install template apps not listed in AppSource

Control which users in the organization can download and install template apps **not** listed on Microsoft AppSource.

Related content

- About tenant settings

Last updated on 09/29/2026

Q&A tenant settings

These settings are configured in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

Review questions

When this setting is enabled, semantic model owners can review questions end-users ask about their data.

To learn more, see Intro to Q&A tooling to train Power BI Q&A.

Synonym sharing

When this setting is enabled, users can share Q&A synonyms as suggested terms with everyone in your organization.

To learn about synonyms, see Field synonyms.

> **Note**
>
> If you disable this setting and apply the changes, and then later re-enable synonym sharing, it might take a few weeks to reshare all the synonyms within your organization.

Related content

- About tenant settings

Last updated on 09/29/2026

Tenant settings for Explore

These settings are configured in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

Users with view permission can launch Explore

Explore is a lightweight visual data exploration experience that lets users perform ad hoc analysis on semantic models and on items connected to those semantic models.

Turn on this setting to allow users with view permission on a semantic model to launch Explore from that model and from items connected to it.

For more information, see Explore data in the Power BI service.

Related content

- About tenant settings
- Explore data in the Power BI service

Last updated on 09/29/2026

Semantic model security tenant setting

These settings are configured in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

Block republish and disable package refresh

Disable package refresh, and only allow the semantic model owner to publish updates. When using the XMLA Endpoint, a user can discover only those semantic models for which they are the owner.

To learn more about semantic model security, see Semantic model permissions.

Related content

- About tenant settings

Last updated on 09/29/2026

Overview of advanced networking settings

These settings are configured in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

Tenant-level Private Link

Increase security by allowing people to use a Private Link to access your Fabric tenant. Someone must finish the setup process in Azure. If that's not you, grant permission to the right person or group by entering their email address.

Review the considerations and limitations before enabling private endpoints.

To learn more, see Private links for Fabric tenants.

Block Public Internet Access

For extra security, block access to your Fabric tenant via the public internet. This means people who don't have access to the Private Link won't be able to connect. Turning this on can take 10 to 20 minutes to take effect.

To learn more, see Private links for Fabric tenants.

Configure workspace-level inbound network rules

When this setting is on, workspace admins can configure inbound private link access protection in workspace settings. When a workspace is configured to restrict inbound network access, existing tenant-level private links can no longer connect to that workspace.

If you turn this setting off, all workspaces revert to their previous inbound network configuration.

For more information, see Workspace-level inbound network rules.

Configure workspace-level outbound network rules

When this setting is on, workspace admins can configure outbound access protection in workspace settings.

If you turn this setting off, outbound access protection is also turned off in all workspaces in the tenant.

For more information, see Private links for Fabric workspaces.

Configure workspace IP firewall rules (preview)

When this setting is enabled, workspace admins can configure IP firewall rules in workspace settings.

If tenant-level public access is blocked, turning on this setting still allows workspace admins to permit access from specific IP rules.

For more information, see Set up workspace IP firewall rules.

Related content

- About tenant settings
- Private links for Fabric tenants

Last updated on 09/29/2026

Encryption tenant setting

These settings are configured in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

Apply customer-managed keys

Customer-managed keys (CMK) are encryption keys that you create, own, and manage in your Azure Key Vault (AKV). By using a CMK, you can supplement default encryption with an extra encryption layer. You can use customer-managed keys for greater flexibility to manage access controls or to meet specific regulatory compliance.

Enable CMK for Fabric workspaces tenant setting

By default, the CMK feature is disabled at the tenant level. This means that workspace administrators cannot enable CMK until the tenant administrator enables the setting. CMK can be enabled and disabled for the workspace while the tenant setting is on. Once the tenant setting is turned off, you can no longer enable CMK for workspaces in that tenant or disable CMK for workspaces that already have CMK turned on in that tenant.

To learn how to set up CMK encryption, see customer-managed keys for Fabric workspaces.

Related content

- About tenant settings

Last updated on 09/29/2026

Scorecards tenant settings

These settings are configured in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

Create and use scorecards

Users in the organization can create and use scorecards in Power BI.

To learn more, see Get started with scorecards in Power BI.

Turning on this setting shows the Scorecards hub and allows users to create, browse, and consume scorecards and goals.

Use the Scorecards tenant setting in Fabric tenant settings to choose whether the feature is enabled and to define which users or groups can access it.

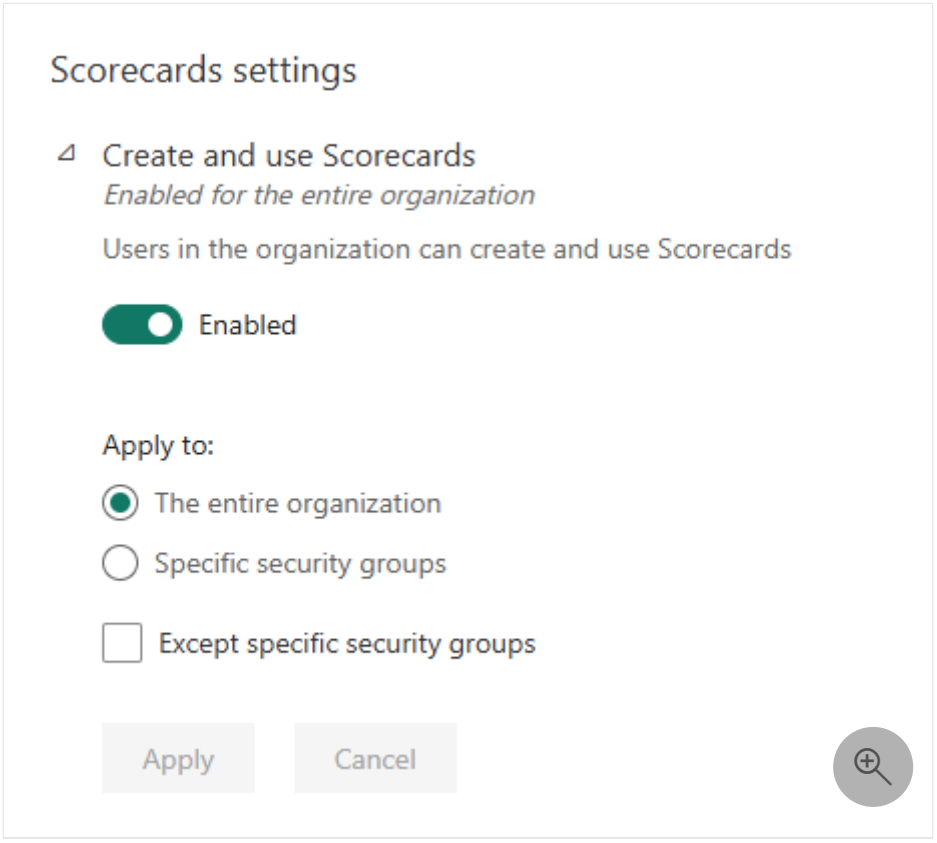

Related content

- About tenant settings
- Get started with scorecards in Power BI

Last updated on 09/29/2026

User experience experiments tenant settings

These settings are configured in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

Help Power BI optimize your experience

Enabling this feature allows the Power BI team to gather early feedback and to make data-driven decisions as to which in-product experience is received more positively by users.

When this feature is enabled, individual users in the same organization might get minor variations in the user experience, including content, layout, and design, before these variations go live for all users. This means that different users in the same tenant might have slightly different experiences.

Related content

- About tenant settings

Last updated on 09/29/2026

Share data with your Microsoft 365 services

This article helps Microsoft Fabric administrators and decision makers understand how and where Fabric metadata is used.

The **Share Fabric data with your Microsoft 365 services** tenant setting controls whether Power BI and Fabric automatically send information about your organization's Fabric content to Microsoft 365. When an admin turns on this setting, Fabric sends details about your reports, dashboards, and other content to Microsoft 365 in the background. No user action is required for this sharing to occur. Microsoft 365 uses this information to help users find and return to their Fabric content through search results, the Quick Access list on Office.com, and personalized recommendations.

This information flows through Microsoft Graph. The setting is turned on by default when your Fabric tenant and Microsoft 365 tenant are in the same geographic region.

> **Important**
>
> Users in Microsoft 365 experiences only see metadata for Power BI items that they have permission to view in Power BI. Power BI continues to enforce the user's permissions, including row-level security (RLS) and object-level security (OLS), when the user opens or interacts with an item.

What information is shared

When you turn on this setting, Fabric sends the following information to Microsoft 365:

Expand table

Content type	What's shared
Power BI reports	Detailed item context including the report name, description, web address, who has access, workspace, who created and last modified the report, dates of creation and last update, page names, chart titles, and column or measure names
Power BI workspace apps	Basic details only: name, web address, who has access, and workspace
Power BI dashboards	Basic details only: name, web address, who has access, and workspace

Content type	What's shared
Power BI paginated reports (RDL)	Basic details only: name, web address, who has access, and workspace

Power BI also shares information about how users interact with content when you turn on this setting. For example, Fabric reports which reports a user viewed. Microsoft 365 uses these signals to show recently accessed and recommended Fabric content on Office.com.

For the complete list of shared properties, see Data shared with Microsoft 365.

> **Note**
>
> Power BI reports shared using organization-wide ("org-wide") links aren't discoverable in Microsoft 365 experiences (such as search) for users who only have access through those links, even when you enable this setting. Users can still access the report directly through the link, but the report doesn't appear in discovery surfaces unless it's shared using another access method.

Supported Microsoft 365 experiences

When you turn on this setting, Microsoft 365 uses the shared Power BI metadata in the following experiences:

- **Semantic Index for Microsoft Copilot:** Shared metadata contributes to the Semantic Index for Microsoft Copilot. The index helps Copilot understand relationships between users and the Power BI items they can access so it can return relevant results.
- **Item attachment menus in Copilot Chat:** Users can search for and attach Power BI reports as context for a Copilot conversation. The attachment menu only shows reports that the signed-in user has permission to view in Power BI.
- **Copilot search:** Copilot uses the shared metadata to help users find relevant Power BI items that they have permission to view.

The following example shows Power BI reports in the **Add content** menu in Copilot Chat. A user can select **Power BI**, search for a report, and attach it as context for the conversation.

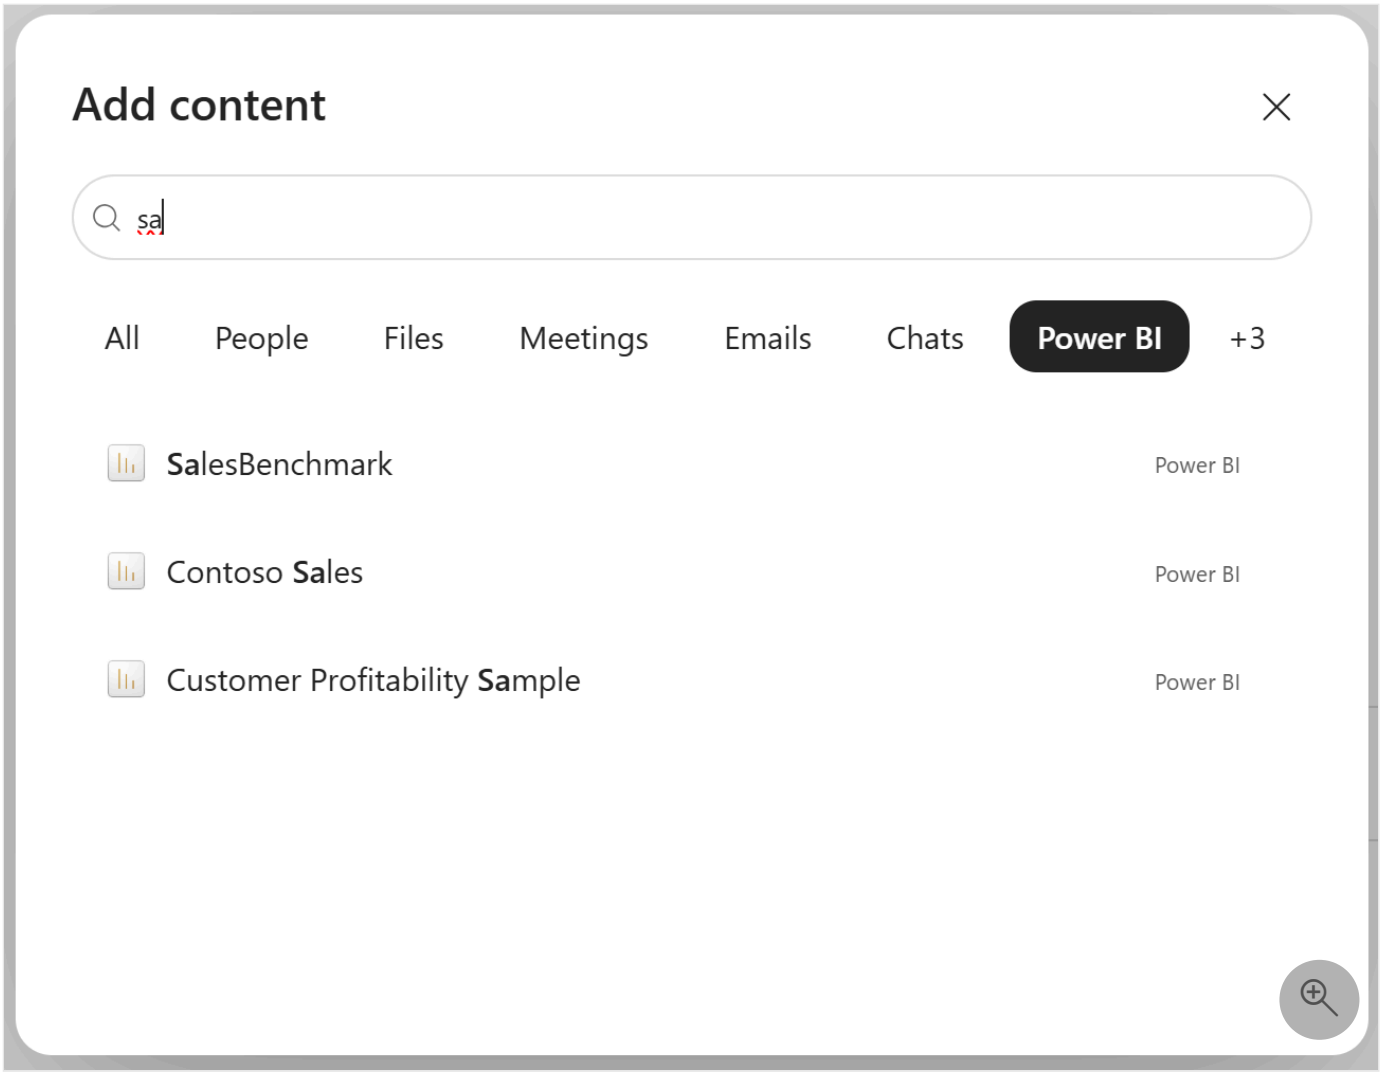

The following example shows Copilot search results filtered to Power BI content. Copilot uses the shared metadata, such as item names, types, descriptions, and modification dates, to help the user identify relevant Power BI items.

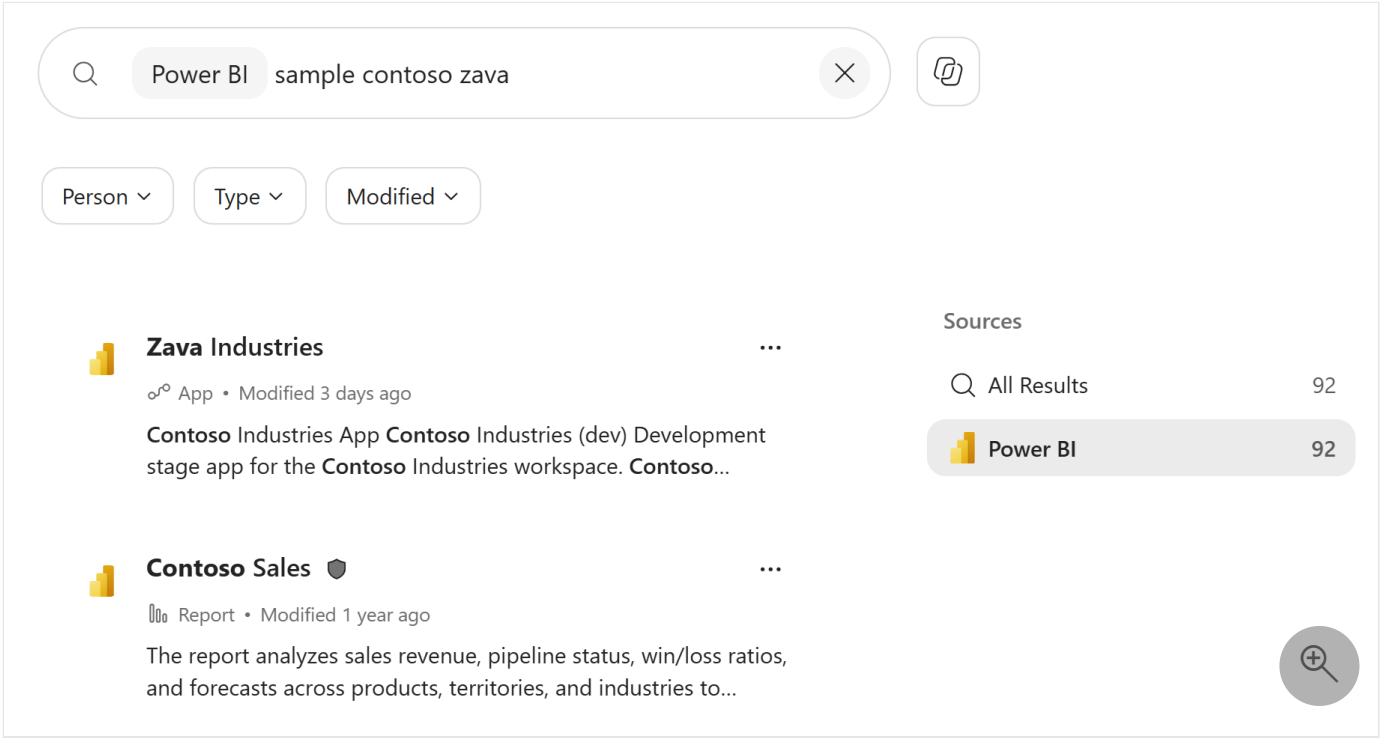

What this setting doesn't control

Turning off this setting doesn't block Fabric content from appearing in Microsoft 365. Users who have access to Fabric content can still take explicit action to use that content in Microsoft 365 products. The following experiences continue to work even when this setting is turned off, because they depend on a user actively doing something:

- **Excel pivot tables connected to Power BI semantic models:** In Excel, users can discover all the Power BI models they have access to and explore that data by using PivotTables and other Excel capabilities. The **Share Fabric data with your Microsoft 365** tenant setting doesn't affect this behavior. Connecting Excel to Power BI models requires Fabric admin approval through a separate tenant setting. For more information, see Power BI semantic model experience in Excel.

- **Link previews in Teams and Outlook:** When a user pastes a link to a Power BI report in a Teams chat or Outlook email, people who have access to that report always see a preview with the report name and other metadata. The **Share Fabric data with your Microsoft 365** tenant setting doesn't affect this behavior. For more information, see Power BI link previews in Teams.

- **Fabric data agents in the Microsoft 365 Agent Store:** If your organization publishes Fabric data agents to the Microsoft 365 Agent Store, users with access can chat with that agent in Microsoft 365. The **Share Fabric data with your Microsoft 365** tenant setting doesn't affect this behavior. For more information, see Use a Fabric data agent in Microsoft 365 Copilot.

- **[Frontier only] Asking Microsoft 365 Copilot about Power BI data:** Microsoft 365 Copilot can answer questions by using Power BI reports and semantic models that the user has permission to view, regardless of this setting. Access to this feature requires Microsoft 365 admin approval. For more information, see Fabric IQ in Microsoft 365 Copilot (Frontier).

- **[Frontier only] Asking Microsoft 365 Copilot Cowork about Power BI data:** Power BI reports and semantic models that the user has permission to view can feed into Copilot Cowork workflows. Users can request emails or schedule meetings triggered by the latest metrics. Users can also produce deep analysis that combines Work IQ context (such as emails, Teams chats, and documents) with insights from Power BI data. The **Share Fabric data with your Microsoft 365** tenant setting doesn't affect this behavior. Microsoft 365 admins can disable access to the Copilot Cowork agent for their organization, and access to the Fabric IQ plugin requires Microsoft 365 admin approval. For more information, see Fabric IQ in Cowork (Frontier).

All of these experiences rely on the user being signed in and having permissions for the content. None of them rely on Fabric sending information to Microsoft 365 in the background.

Data residency

Fabric and Microsoft 365 are separate cloud services that might be hosted in different geographic regions, even if you purchase them together.

By default, Fabric only shares information within the geography where your Fabric tenant is located. If your Microsoft 365 tenant is in a different geography, you need to allow cross-geography sharing for the integration to work. You can enable cross-geography sharing by using a second toggle in the tenant setting. For more information, see How to turn sharing with Microsoft 365 services on and off.

Where is Fabric data stored?

For more information about data storage locations, see Find your Fabric home region and Product Availability by Geography.

Where is Microsoft 365 data stored?

For more information about data storage for Microsoft 365, see Where your Microsoft 365 customer data is stored and Multi-Geo Capabilities in Microsoft 365.

How to turn sharing with Microsoft 365 services on and off

The **Share Fabric data with your Microsoft 365 services** tenant setting is on by default. To change this setting:

1. Go to **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings** > **Share Fabric data with your Microsoft 365 services**.
2. Set the toggle to **Enabled** or **Disabled** as needed.

Changes can take up to 24 hours to take effect.

To allow sharing across geographies when your Fabric and Microsoft 365 tenants are in different regions, set the second toggle to **Enabled**. Turning on this toggle means that you acknowledge that Fabric data might flow outside its current geographic region.

> **Note**
>
> You can see the second toggle only when the main sharing toggle is enabled.

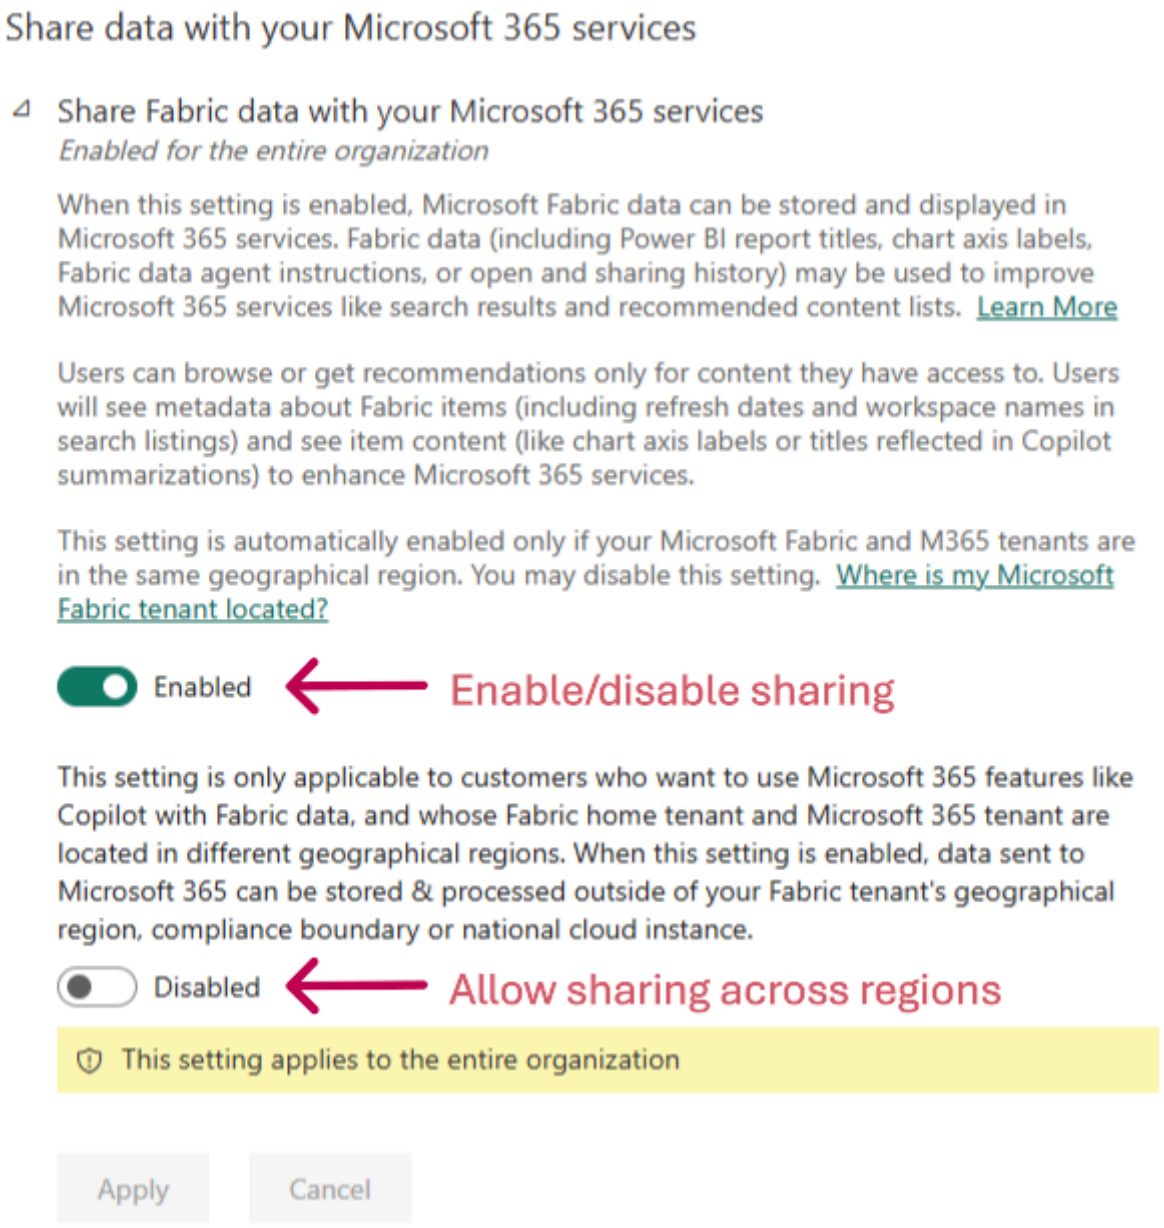

Data shared with Microsoft 365

The following tables show what specific information Fabric sends to Microsoft 365 when sharing is turned on.

Item details (used for search and discovery in Microsoft 365)

Expand table

Property	What is shared	Example
TenantID	Microsoft Entra tenant identifier	aaaabbbb-0000-cccc-1111-dddd2222eeee
ItemType	Fabric category for the item	Report
DisplayName	Display name for the item	Retail Analysis Sample
Description	Content description in the service. See Report settings for more detail.	Sample containing retail sales data
URL	Content item URL	`https://powerbi-df.analysis-df.windows.net/groups/8b5ac04e-89c1-4fc6-a364-e8411dfd8d17/reports/aaaabbbb-0000-cccc-1111-dddd2222eeee/ReportSection2`
ACL	Access Control List (ACL) with permissions and Microsoft Entra user, security group, and distribution list identifiers	`{"accessType": "grant", "id": "aaaaaaaa-bbbb-cccc-dddd-eeeeeeeeeeee", "type": "read"}`
WorkspaceName	Workspace name. See Create a workspace.	Retail workspace
WorkspaceURL	Link to the workspace in the service	`https://powerbi-df.analysis-df.windows.net/groups/8b5ac04e-89c1-4fc6-a364-e8411dfd8d17`
Creator	Microsoft Entra user principal name (UPN) of the person who created the content	user1@fourthcoffee.com
CreatedDate	Date the content was created	2011-06-30T23:32:46Z
LastModifiedUser	Microsoft Entra UPN of the last person who modified the content	user1@fourthcoffee.com
LastModifiedDate	Last modified date for the content	2011-06-30T23:32:46Z
PageNames	Display names for pages within the report	Sales Summary, Regional details, Returns
ChartTitles	Display names for visuals in the report layout	Regional sales over time
FieldNames	Names of columns and measures used in the report	revenue, date, product_category

User activity (used to show recently accessed and recommended content on Office.com)

Expand table

Property	What is shared	Example
LastRefreshDate	Last refresh date for the content	2011-06-30T23:32:46Z
UserID	Microsoft Entra UPN of the user who acted on the item	user1@fourthcoffee.com
SignalType	The type of action the user took on the content	Viewed
ActorID	Microsoft Entra ID of the user who acted on the item	aaaaaaaa-bbbb-cccc-dddd-eeeeeeeeeeee
StartTime/EndTime	Date and time the user performed the action	2011-06-30T23:32:46Z

> **Note**
>
> Power BI reports shared using organization-wide ("org-wide") links aren't discoverable in Microsoft 365 experiences (such as search) for users who only have access through those links, even when the **Share Fabric data with your Microsoft 365 services** setting is enabled. This behavior is expected: users can still access the report directly via the link, but it doesn't appear in discovery surfaces unless the report is shared via another access method.

Related content

- About tenant settings

Note: The author created this article with assistance from AI. Learn more

Last updated on 09/15/2026

Insights tenant settings

These settings are configured in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

Receive notifications for top insights (preview)

Users in the organization can enable notifications for top insights in report settings.

To learn more about insights, see Find Insights in your reports.

Show entry points for insights (preview)

Users in the organization can use entry points for requesting insights inside reports.

To learn more about insights, see Find Insights in your reports.

Related content

- About tenant settings

Last updated on 09/29/2026

Datamart tenant settings

Datamart tenant settings are configured in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

Create Datamarts (Preview)

When this setting is on, specified users in the organization can create datamarts.

For more information, see Administration of datamarts.

Related content

- About tenant settings
- Administration of datamarts

Last updated on 09/29/2026

Tenant settings for semantic models

Fabric administrators can enable or disable semantic model editing in the service for the entire organization or for specific security groups, using the setting described in this article. This setting is configured in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

Users can edit semantic models in the Power BI service

Turn on this setting to allow users to edit semantic models in the service. This setting doesn't apply to Direct Lake semantic models or editing a semantic model through an API or XMLA endpoint.

To learn more, see Enabling semantic model editing.

Related content

- About tenant settings
- Edit semantic models in the Power BI service

Last updated on 09/29/2026

Tenant settings for scale-out

Scale-out tenant settings are configured in **OneLake catalog** > **Govern** > **Configurations** > **Tenant settings**. For information about how to get to and use tenant settings, see About tenant settings.

Scale out queries for large semantic models

For semantic models that use the large semantic model storage format, Power BI Premium can automatically distribute queries across additional semantic model replicas when query volume is high. Scale-out is enabled on a tenant by default.

For more information, see Power BI semantic model scale-out.

Related content

- About tenant settings
- Power BI semantic model scale-out

Last updated on 09/29/2026

OneLake tenant settings

OneLake tenant settings are configured in **Configurations** > **Tenant settings** in the Govern section of the OneLake catalog. For information about how to get to and use tenant settings, see About tenant settings.

Users can access data stored in OneLake with apps external to Fabric

Users can access data stored in OneLake with apps external to the Fabric environment, such as custom applications created with Azure Data Lake Storage (ADLS) APIs, OneLake file explorer, and Databricks. Users can already access data stored in OneLake with apps internal to the Fabric environment, such as Spark, Data Engineering, and Data Warehouse.

To learn more, see Allow apps running outside of Fabric to access data via OneLake.

Users can sync data in OneLake with the OneLake file explorer app

Turn on this setting to allow users to use OneLake file explorer. This app will sync OneLake items to Windows File Explorer, similar to OneDrive.

To learn more, see OneLake file explorer.

Use short-lived user-delegated SAS tokens

This setting allows users to request OneLake user delegation keys for this tenant, which are used to sign OneLake SAS tokens. OneLake SAS tokens enable applications to access data in OneLake through short-lived SAS tokens, based on a Fabric user's Entra identity. These token's permissions can be further limited to provide least privileged access and cannot exceed a lifetime of one hour. This setting allows users to request OneLake user delegation keys for this tenant.

To learn more, see OneLake shared access signatures.

Authenticate with OneLake user-delegated SAS tokens

Allow applications to authenticate using a OneLake SAS token. Fabric users can create OneLake SAS by requesting a user delegation key. The tenant setting, Use short-lived user delegated SAS tokens, must be turned on to generate user delegation keys. The lifetimes of the user delegation keys and SAS tokens cannot exceed one hour. This setting is delegated to workspace admins by default, allowing them to override the default tenant setting.

To learn more, see OneLake shared access signatures.

Include end-user identifiers in OneLake diagnostic logs

Control whether OneLake diagnostic logs capture end user identifiable information (EUII), such as email addresses and IP addresses. When enabled, these fields are recorded to support diagnostics, investigations, and usage analysis across your tenant. When disabled, these fields are redacted from new events.

To learn more, see OneLake diagnostics.

Related content

- About tenant settings

Last updated on 09/29/2026